

CAS Cloud and Platform Manager

Microservices, Containers und Kubernetes

Niklaus Hirt

Consulting IT Specialist
Cloud and DevOps Architect

nikh@ch.ibm.com

Rotkreuz, 26.05.2023

Who am I?

Niklaus Hirt

Passionate about tech for 40 years

- High-school in Berne
- Degree in Computer Science at EPFL
- ELCA
- CAST
- IBM



✉ nikh@ch.ibm.com

🐦 @nhirt

Agenda – Microservices and Containers – A crash course

Module 1: Microservices

Module 2: Microservices - Let's get real

Module 3: Containers

Module 4: Kubernetes - Basics

Lunch

Agenda – Microservices and Containers – A crash course

Module 5: Kubernetes - Applied

Module 6: Kubernetes - Advanced Concepts

Module 7: Kubernetes - State

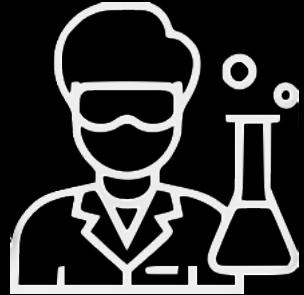
Module 8: Mesh Networking (optional)

Self Study - Optional

Lab 1: Containers Hands-On

Lab 2: Kubernetes Hands-On

Instructions can be found in the Annex



Sources and documentation are available here:

Course Documents

https://github.com/niklaushirt/k8s_training_public

Course Files

<https://github.com/niklaushirt/training>

Online Course

https://niklaushirt.github.io/k8s_training_web/

Kubernetes Workshop Series

Microservices

01



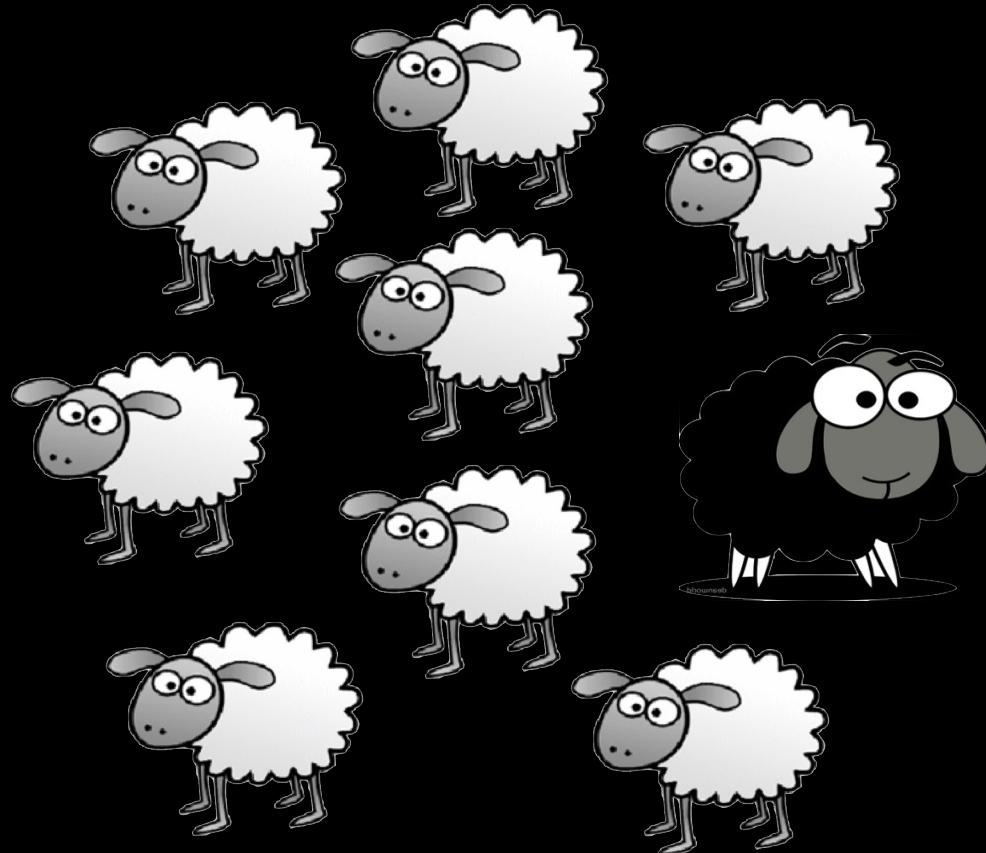


disruption

dɪs'rʌpʃn/
noun

Business. a radical change in an industry, business strategy, etc., especially involving the introduction of a new product or service that creates a new market

Disruption is new reality



disruption

dɪs'rʌpʃn/

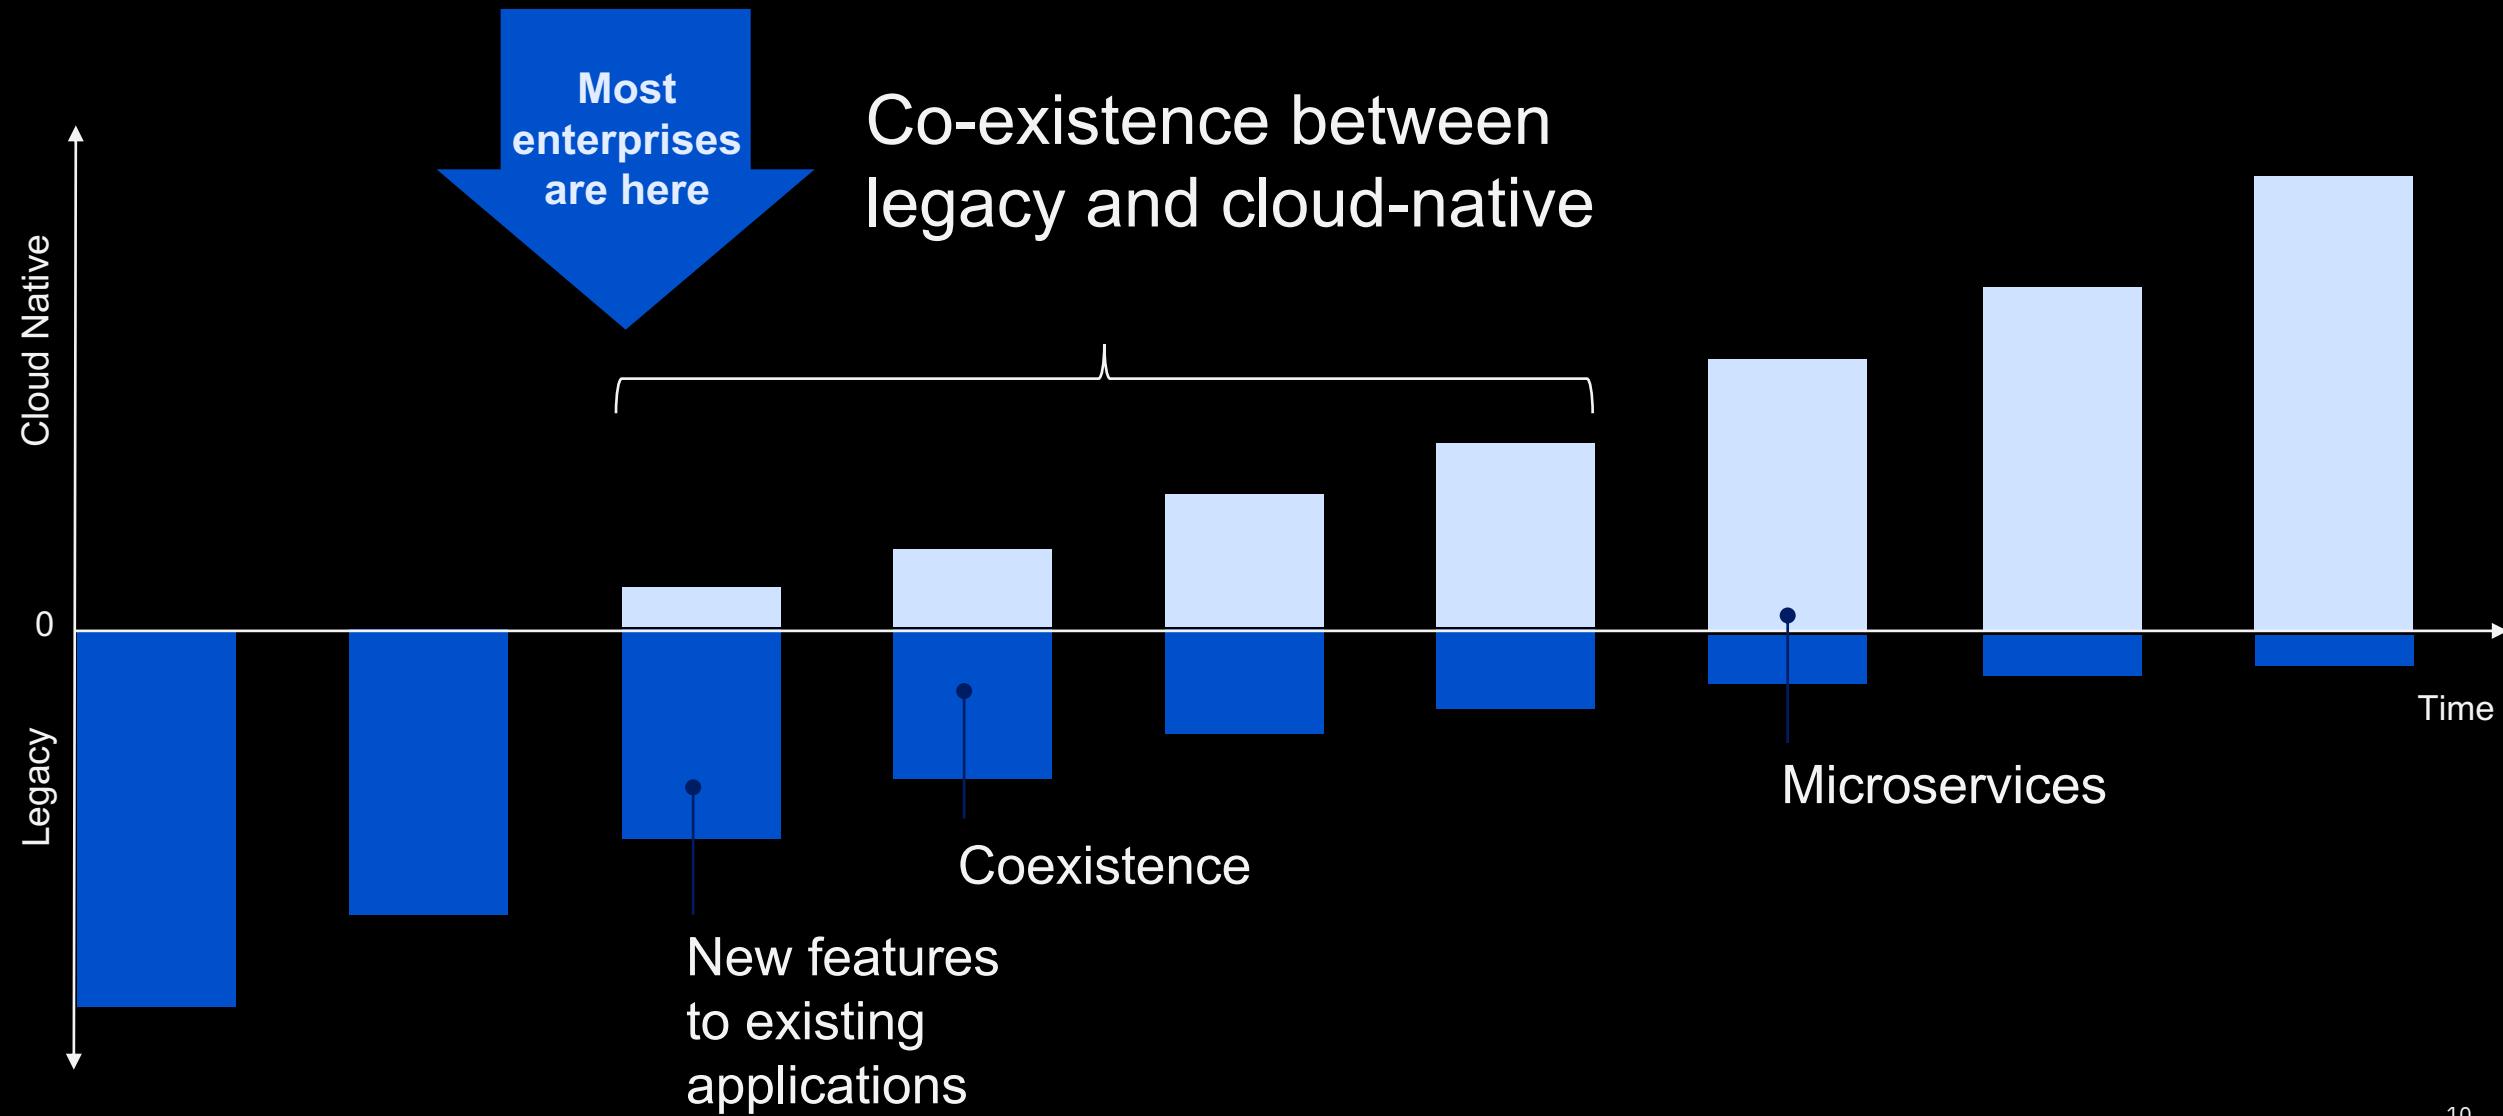
noun

Business. a radical change in an industry, business strategy, etc., especially involving the introduction of a new product or service that creates a new market

Digital Transformation is the new normal

Digital Transformation - Way to go

Cloud native and legacy apps will co-exist for the next 10+ years



How do you achieve digital transformation?

Adopt new Processes

- Continuous Integration
- Continuous Build
- Continuous Deploy
- Agile Dev Models
- DevOps

Adopt new Technology

- Architectural patterns (Microservices)
- Frameworks (Java MicroProfile, Spring ...)
- Node.js , Swift, GoLang, Python(!)

Adopt new Tools

- Git, Github, Gitlab
- CI/CD (ArgoCD, ...)
- AIOps

Adopt Cloud-Native

- Cloud
- Containers
- Kubernetes

How do you achieve digital transformation?

Adopt new Processes

- Continuous Integration
- Continuous Build
- Continuous Deploy
- Agile Dev Models
- DevOps

Adopt new Technology

- Architectural patterns (Microservices)**
- Frameworks
(Java MicroProfile, Spring ...)
- Node.js , Swift, GoLang, Python(!)

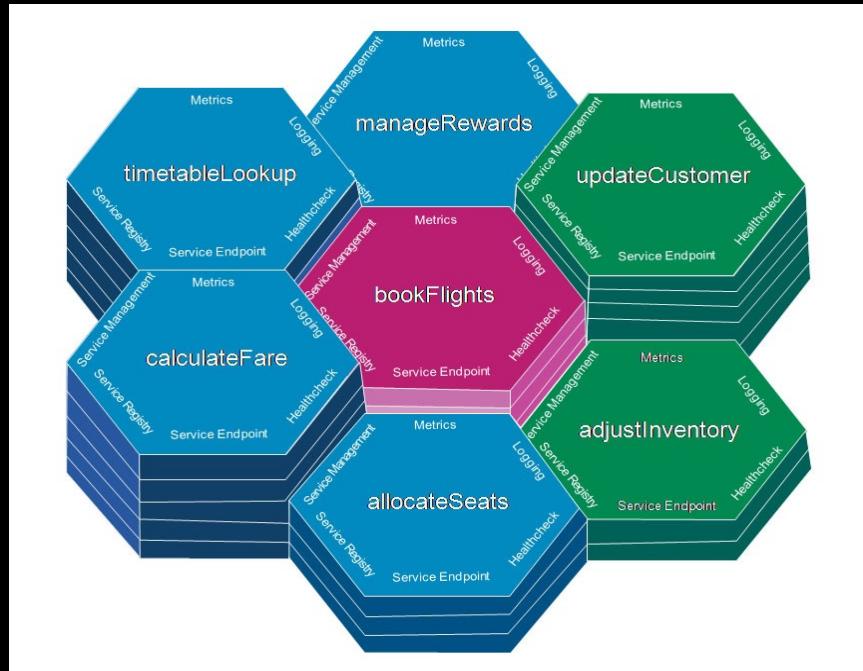
Adopt new Tools

- Git, Github, Gitlab
- CI/CD (ArgoCD, ...)
- AIOps

Adopt Cloud-Native

- Cloud
- Containers**
- Kubernetes**

Microservices

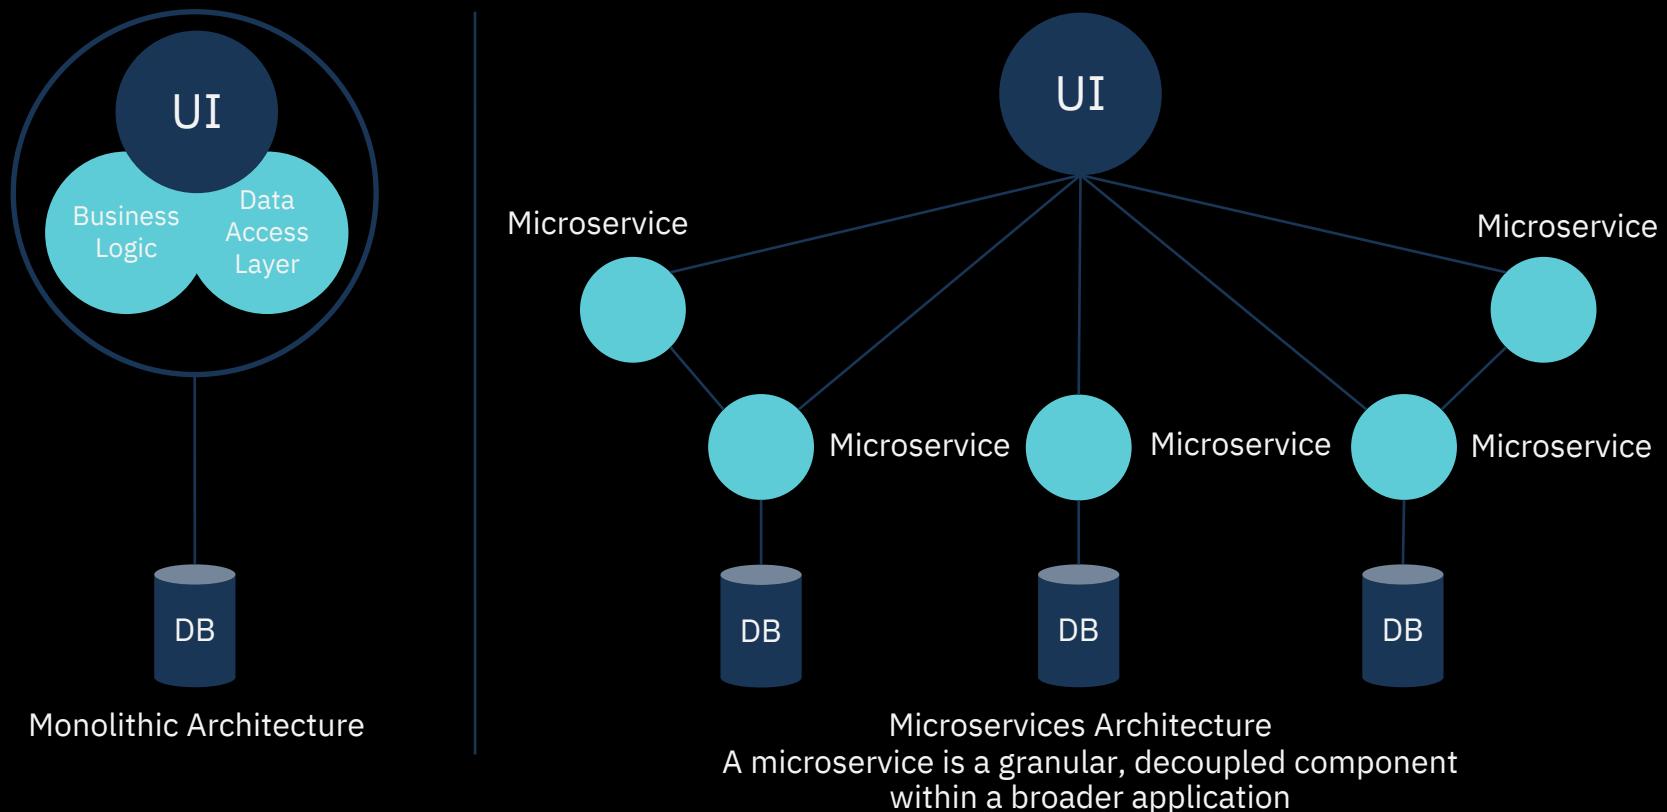


Decomposing an application into single function modules which are independently deployed and operated

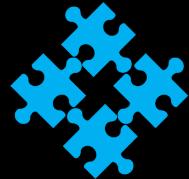
Accelerate delivery by minimizing communication and coordination between people

Microservices architecture

Simplistically, microservices architecture is about breaking down large silo applications into more manageable, fully decoupled pieces



Microservices – key tenets



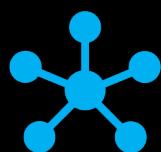
Large monoliths are **broken down** into many small services

- Each service runs its own process
- There is one service per container



Services are **optimized for a single function**

- There is only one business function per service
- The Single-responsibility Principle: A microservice should have one, and only one, reason to change



Communication via **REST API** and **message brokers**

- Avoid tight coupling introduced by communication through a database



Per-service continuous delivery (CI/CD)

- Services evolve at different rates
- Let the system evolve, but set architectural principles to guide that evolution



Per-service high availability and scaling decisions

- One size or scaling policy is not appropriate for all
- Not all services need to scale; others require autoscaling up to large numbers

Microservices – advantages

In a microservices architecture each component:

- Is developed independently and has limited, explicit dependencies on other services
- Is developed by a single, small team in which all team members can understand the entire code base
- Is developed on its own timetable so new versions are delivered independently of other services
- Scales and fails independently which better isolates problems
- Can be developed in a different language
- Manages its own data to select the best technology and schema

That said....

Microservices

are

hard

Microservices, when
implemented incorrectly,
can make poorly written
applications even more
dysfunctional

QUESTIONS?



Kubernetes Workshop Series

Let's get real

02



Microservices – The Pitfall

Containerizing your WAR file
doesn't mean you're doing
microservices.

- What is the domain? What is reality?
- Where are the transactional boundaries?
- How should microservices communicate across boundaries?
- What if we just turn the database inside out?

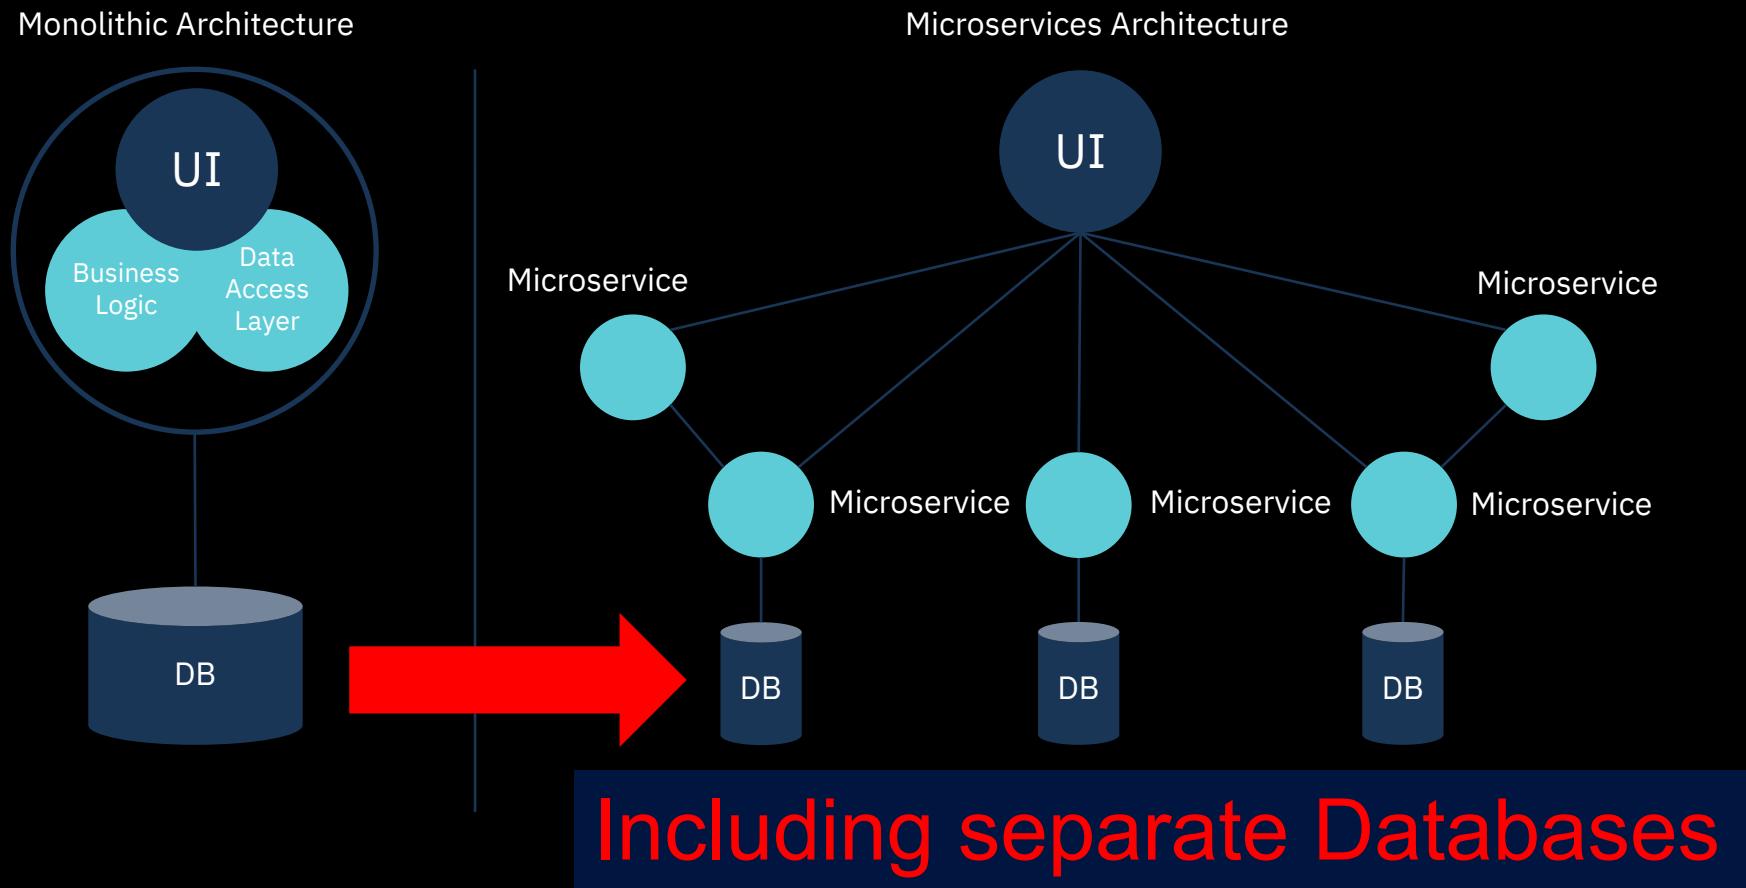
Microservices – breaking down the ~~monolith~~ beast

How to identify candidates

- Is there anything that is **scaling differently** than the rest of the system?
- Is there anything that feels “**tacked-on**”?
- Is there anything **changing much faster** than the rest of the system?
- Is there anything requiring more **frequent deployments** than the rest of the system?
- Is there a part of the system that a small team, **operates independently**?
- Is there a **subset of tables** in your datastore that isn’t connected to the rest of the system?

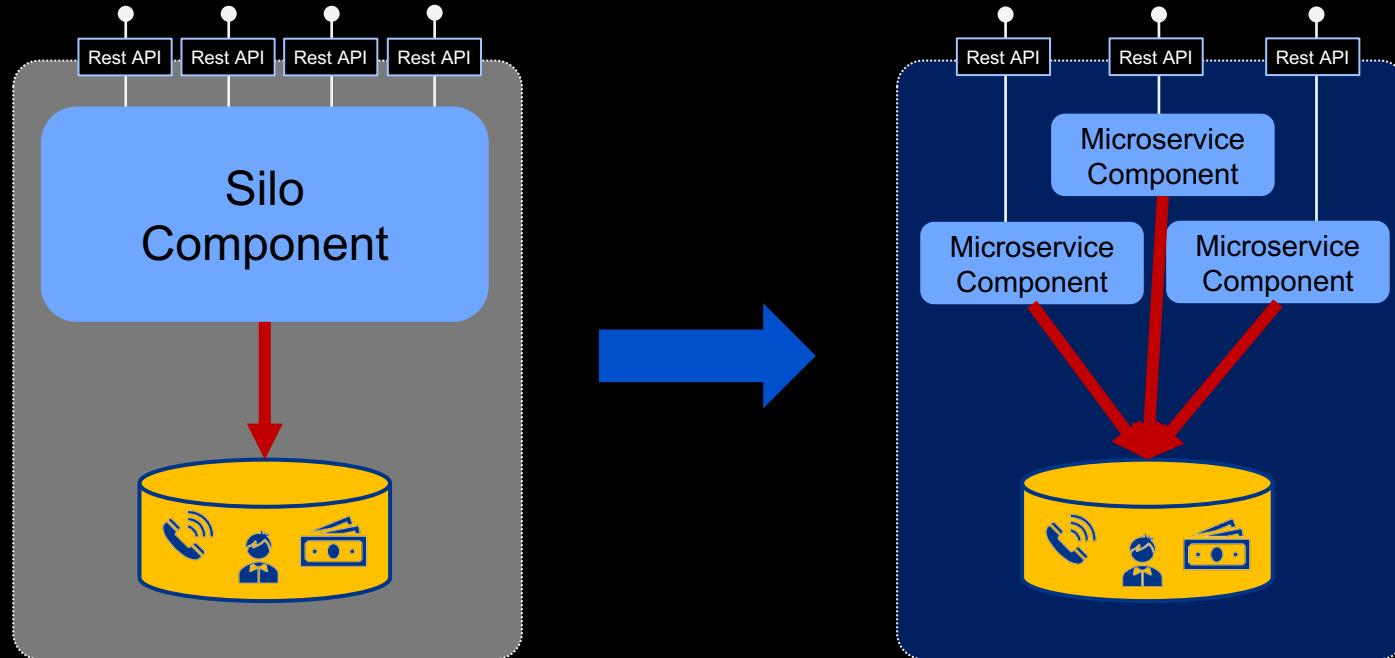
Microservices – Database refactoring

Simplistically, microservices architecture is about breaking down large silo applications into more manageable, fully decoupled pieces



Microservices – Database refactoring - Antipattern

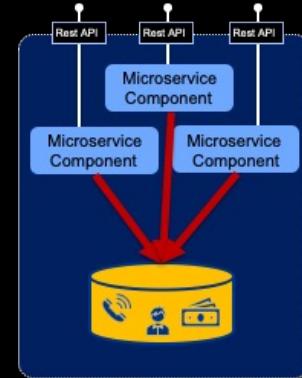
Common mistake



Microservices – monolithic database

Challenges

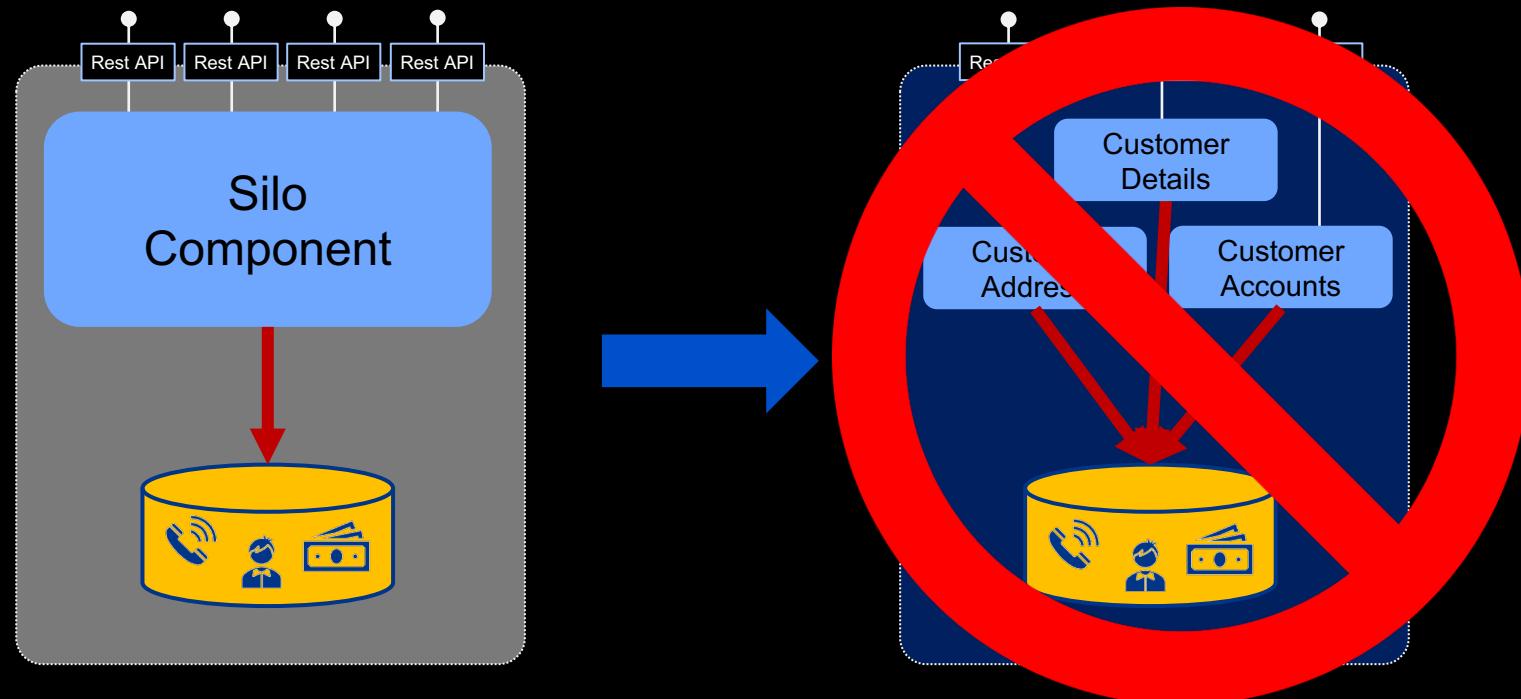
- Inability to deploy your service changes independently because of **tight coupling through the Database**.
- Schema changes need to be coordinated amongst all the services
- **Difficult to scale** individual services
- All your services have to use a **relational database** even when a no-SQL datastore would bring benefits
- Difficult to improve application performance (DB/Table size)



Microservices – Database refactoring

Data Isolation

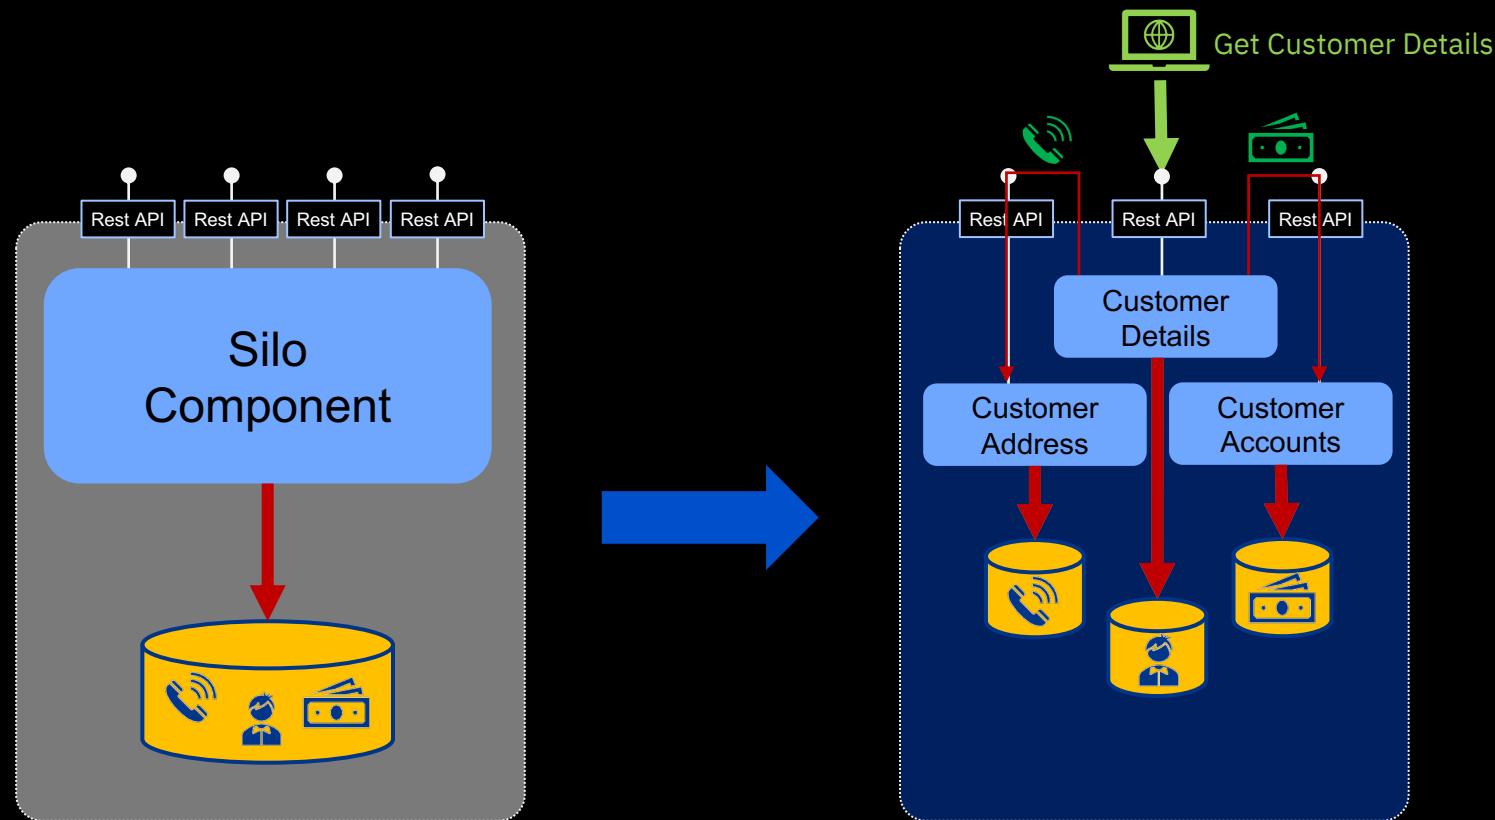
Each microservice must keep its own data. One typical error that happens in the beginning of a microservice transformation is the use of a centralized database, the same way it was used in the monolithic application. This creates a single point of failure and dependency between the microservices.



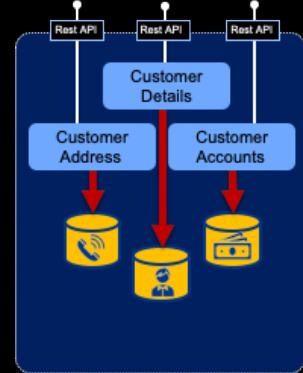
Microservices – Database refactoring

Data Isolation

Each microservice must keep its own data. One typical error that happens in the beginning of a microservice transformation is the use of a centralized database, the same way it was used in the monolithic application. This creates a single point of failure and dependency between the microservices.



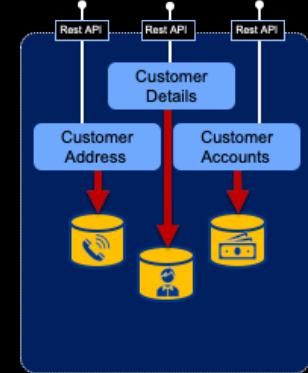
Microservices – per-service database



Benefits

- Avoid unexpected data modification – the API enforces consistency
- Make changes to our system without blocking progress of other parts of the system
- Make changes to the schema/databases at our leisure
- Store the data in a project-owned databases, using the appropriate technology
- Become much more scalable, fault tolerant, and flexible

Microservices – per-service database



Drawbacks

- Costly to refactor
- Introduces latency
- Needs robust transaction handling by the service (cost!)
- More difficult to debug
- More difficult to operationalize
- More difficult to test (needs well designed API tests)

Microservices

Good or bad idea?

Microservices – reasons not to use them

Some thoughts

- If speed is your goal, microservices might not be the solution (real-time processing, ...)
- Applications that require tight integration between individual components and services
- At what point is it in its lifecycle? Mission-critical?
- Not all applications are large enough to break down into microservices

Ask yourself: What **Business Problem** does this solve?

Microservices – reasons not to use them

What works for **large and complex** applications
does not always work at a **smaller scale**

What makes sense for a **new application**
does not always make sense when maintaining or
updating **existing applications**

Microservices – the Monolith is NOT dead

Why Microservices Fail? ★

File Edit View Insert Format Data Tools Add-ons Help

View only

https://blog.runscope.com/posts/5-reasons-not-to-use-microservices

A	B	C	D
5 Monolith First	Short positioning	https://twitter.com/martinfowler/	https://martinfowler.com/bliki/MonolithFirst.html
6 Microservices For Greenfield?	General analysis	https://twitter.com/samnewman	https://samnewman.io/blog/2015/04/07/microservices-for-greenfield/
7 In Defence of the Monolith part 1	Deep technical analysis	https://twitter.com/dkhaywood	https://www.infoq.com/articles/monolith-defense-part-1
8 In Defence of the Monolith part 2	Deep technical analysis	https://twitter.com/dkhaywood	https://www.infoq.com/articles/monolith-defense-part-2
9 The fast-moving monolith	Use case	Raffaele Spazzoli	https://developers.redhat.com/blog/2016/10/27/the-fast-moving-monolith-how-we-sped-up-delivery-from-every-three-months-to-every-week/
10 Majestic Modular Monoliths - video	Technical analysis	https://twitter.com/axelfontaine	https://vimeo.com/233980163
11 Majestic Modular Monoliths - slides	Technical analysis	https://twitter.com/axelfontaine	https://speakerdeck.com/axelfontaine/majestic-modular-monoliths
12 When not to use microservices	General analysis	https://twitter.com/cfe84	https://www.feval.fr/posts/microservices/
13 Presentation: Complex Event Flows in Distributed Systems	Technical use case analysis	https://twitter.com/berndruecker	https://www.infoq.com/presentations/event-flow-systems
14 Monitoring and Managing Workflows Across Collaborating Microservices	Technical use case analysis	https://twitter.com/berndruecker	https://www.infoq.com/articles/monitor-workflow-collaborating-microservices
15 The rise of non-microservices architectures	Reflection	https://twitter.com/bibryam	https://developers.redhat.com/blog/2018/09/10/the-rise-of-non-microservices-architectures/
16 The Majestic Monolith	Old but gold	https://twitter.com/dhh	https://m.signalnoise.com/the-majestic-monolith-29166d022228
17 Goodbye Microservices	Use case	Alexandra Noonan	https://segment.com/blog/goodbye-microservices/
18 Long Live and Prosper To Monolith	Motivational	https://twitter.com/alexstob	https://www.slideshare.net/alexstob/long-live-and-prosper-to-monolith
19 Github: Moduliths	Mono framework	https://twitter.com/olivergierke/	https://github.com/odrotbohm/moduliths
20 Why You Shouldn't Use Microservices	Personal rant	https://hackernoon.com/@jensbole	https://hackernoon.com/dont-use-microservices-c3b5484b329a
21 Self-contained System	Alternative architecture	https://twitter.com/lnnog	https://scs-architecture.org/
22 About When Not to Do Microservices	Short positioning	https://twitter.com/christianposta	http://blog.christianposta.com/microservices/when-not-to-do-microservices/
23 5 Reasons Not to Use Microservices	Short analysis	Michael Churchman	https://blog.runscope.com/posts/5-reasons-not-to-use-microservices
24 Pattern: Monolithic Architecture	Pattern	https://twitter.com/crichardson	https://microservices.io/patterns/monolithic.html
25 It's time to stop making "Microservices" the goal of modernization.	Short	https://twitter.com/rbarcia	https://medium.com/@rbarcia/its-time-to-stop-making-microservices-the-goal-of-modernization-71758b400287
26 The Death of Microservice Madness in 2018	Long analysis	https://twitter.com/dwmkerr	https://dwmkerr.com/the-death-of-microservice-madness-in-2018/

https://docs.google.com/spreadsheets/d/1vjnjAII_8TZBv2XhFHra7kEQzQpOHSZpFIWDjynYYf0

QUESTIONS?



Kubernetes Workshop Series

~~Docker~~ Containers

03



Everybody Loves Containers



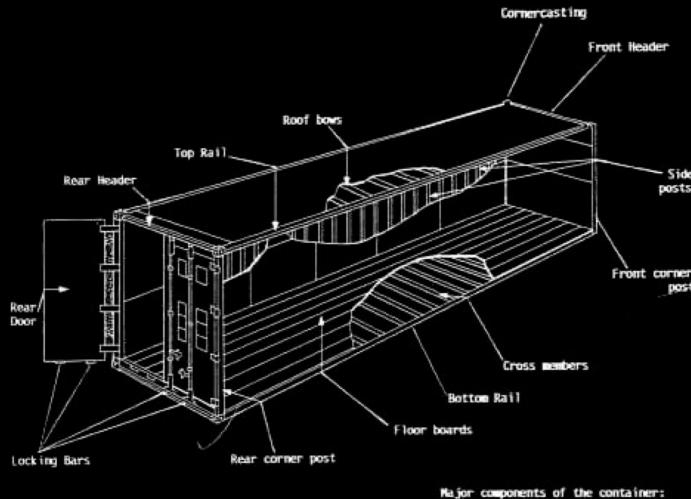
A standard way to package an application and all its dependencies so that it can be moved between environments and run without changes.

Microservices implementation with Containers

Why it works – separation of concerns

Development

- Worries about what's “**inside**” the container
 - Code
 - Libraries
 - Package Manager
 - Apps
 - Data
- All Linux servers look the same

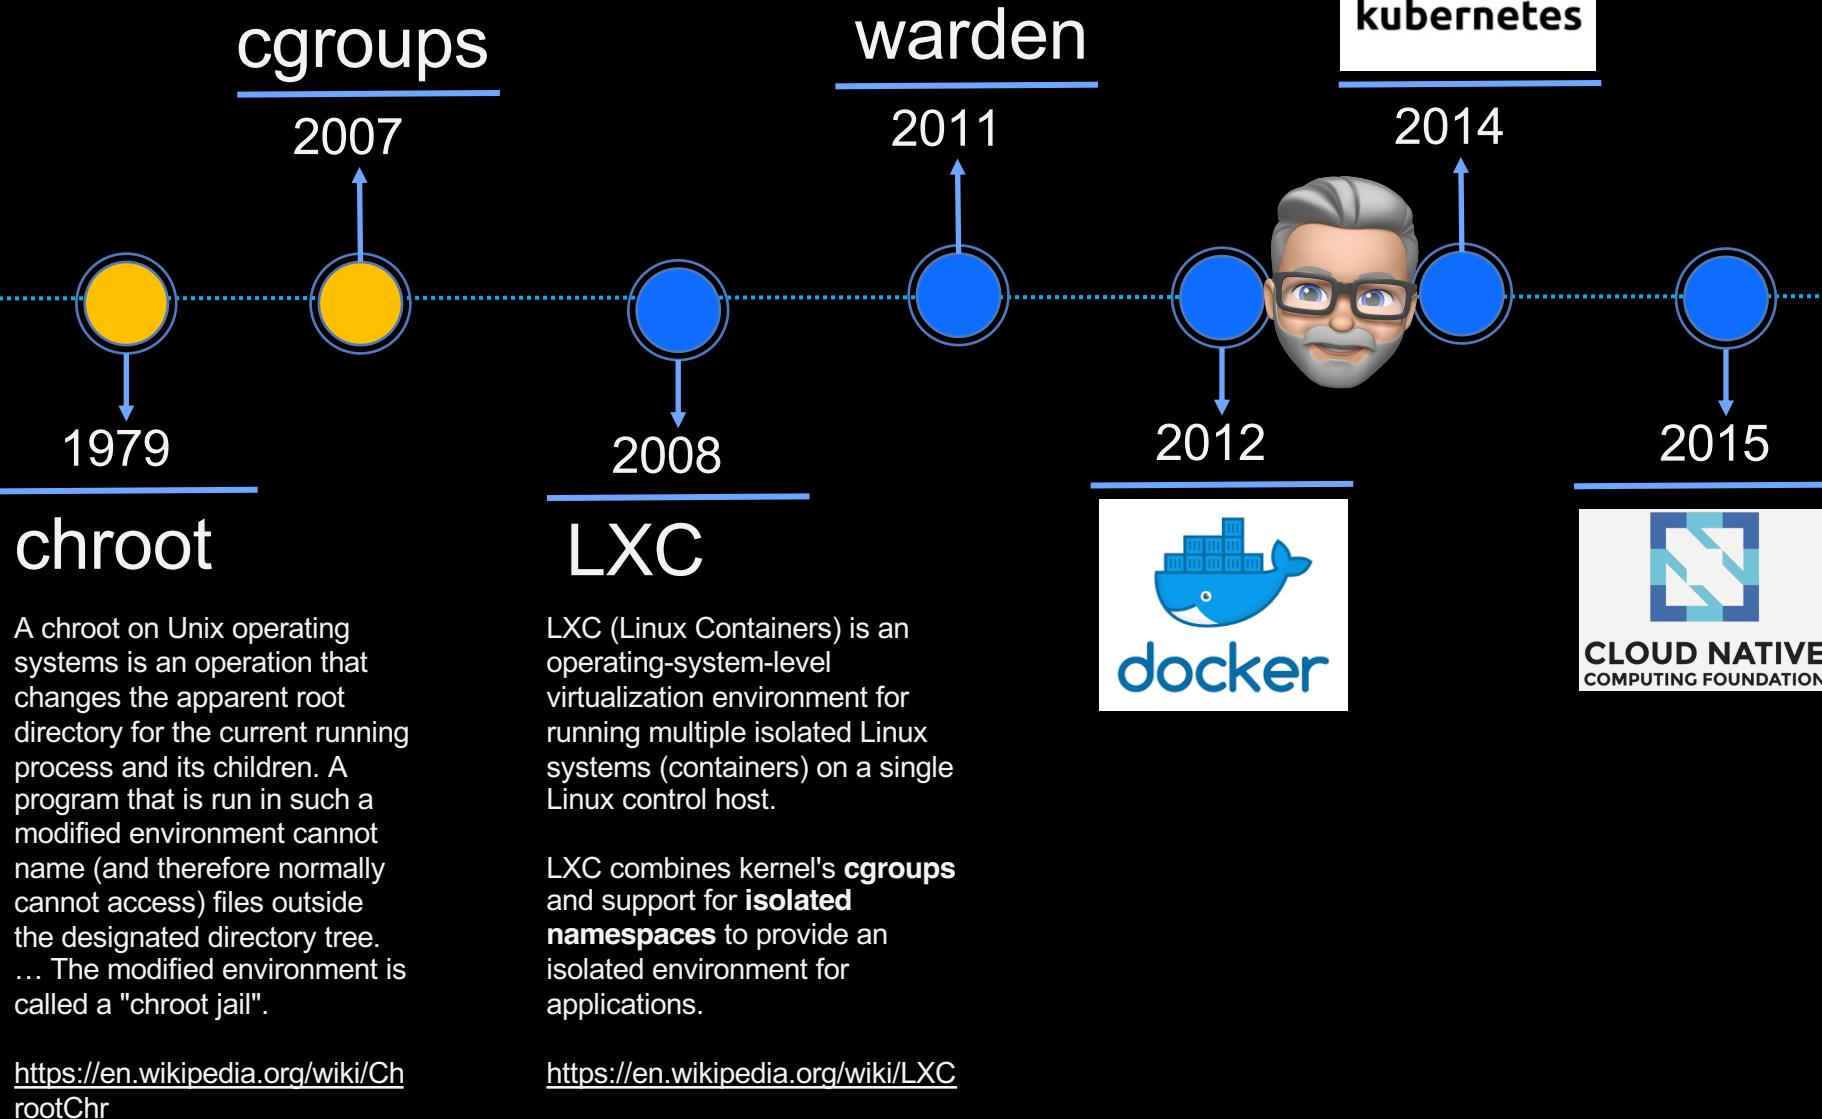


Operations

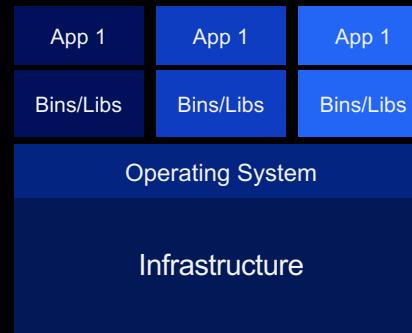
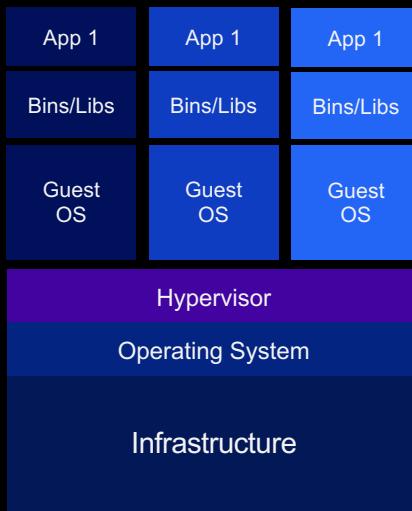
- Worries about what's “**outside**” the container
 - Logging
 - Remote Access
 - Monitoring
 - Network Config
- All containers start, stop, copy, attach, migrate, etc... the same way

Clear ownership boundary between
Dev and IT Ops drives DevOps adoption and fosters agility

Container History



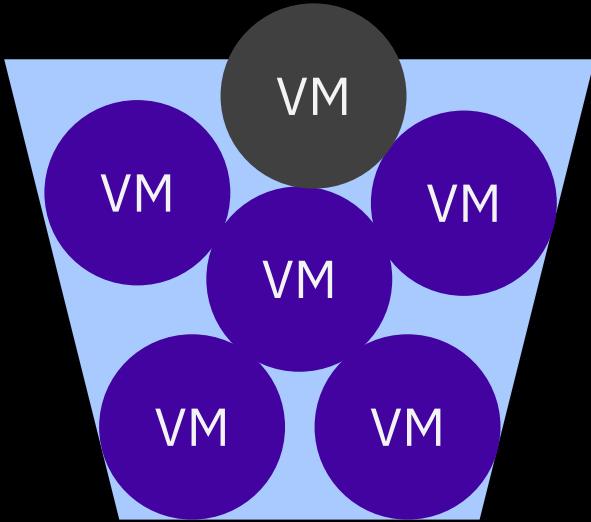
VMs vs. Containers



- + VM Isolation (OS)
- Complete OS
- Static Compute
- Static Memory

- + Container Isolation (Kernel)
- + Shared Kernel
- + Burstable Compute
- + Burstable Memory

VMs vs. Containers



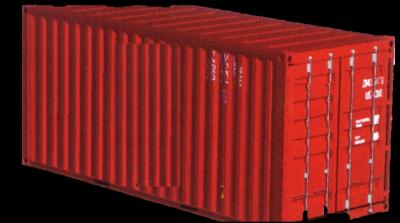
- ++ VM Isolation (OS)**
- Complete OS
- Static Compute
- Static Memory
- Low Resource Utilization

- + Container Isolation (Kernel)
- + Shared Kernel
- + Burstable Compute
- + Burstable Memory
- + High Resource Utilization

Advantages of Containers



Containers provide “just enough” isolation



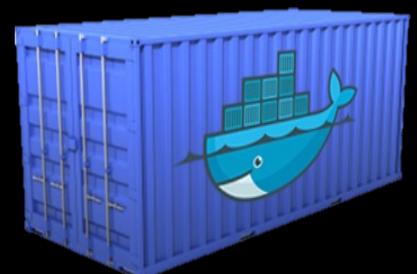
Containers use hardware more efficiently



Containers are immutable



Containers are portable



Container Universe

Components

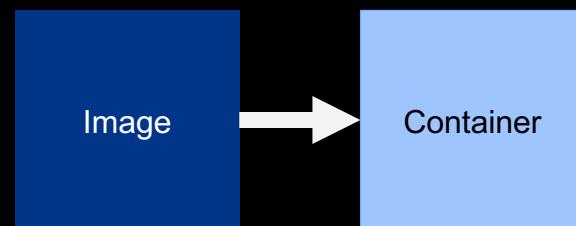
Container

Smallest compute unit



Components

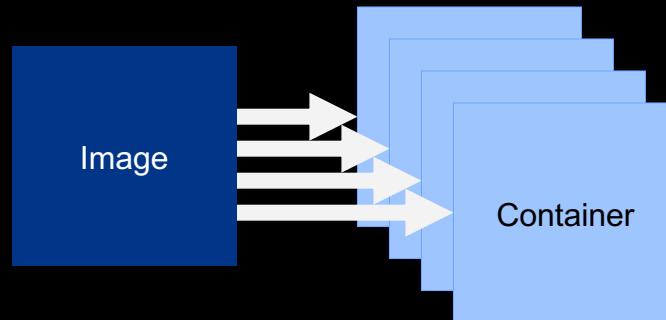
Containers
are created from
Images



one process per container

Components

As **many Containers**
as needed can be created from
Images



Components

Images are versioned

`webapp:v2.0`

Image Name Image Tag

latest is a special tag and designates the latest built image. Should be used with caution to ensure immutability.

Components

Images are Open Container Initiative - OCI Compliant

The screenshot shows the official website of the Open Container Initiative. At the top, there's a navigation bar with links for About, Community, Join, Blog, News, FAQ, Release notices, and Contact us, along with social media icons for GitHub and Twitter.

The main content area features a large dark blue box containing the following text:

Open Container Initiative

The **Open Container Initiative** is an open governance structure for the express purpose of creating open industry standards around container formats and runtimes.

Established in June 2015 by Docker and other leaders in the container industry, the OCI currently contains three specifications: the Runtime Specification (runtime-spec), the Image Specification (image-spec) and the Distribution Specification (distribution-spec). The Runtime Specification outlines how to run a "filesystem bundle" that is unpacked on disk. At a high-level an OCI implementation would download an OCI Image then unpack that image into an OCI Runtime filesystem bundle. At this point the OCI Runtime Bundle would be run by an OCI Runtime.

[Learn more ↗](#)

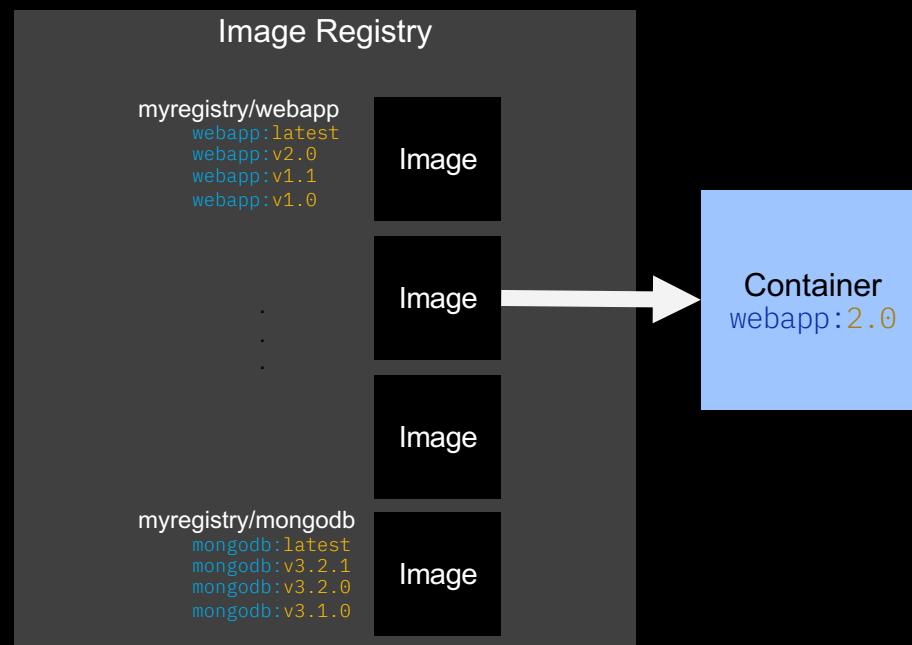
Below this text is a large 3D cube icon, colored blue and purple, with teal arrows pointing to its height, width, and depth.

At the bottom of the page are three white callout boxes with blue icons:

- Participate** in the technical community
- Become a **Member Organization** and support the Open Container Initiative
- Use the tooling and apply to be **OCI Certified**

Components

The **Image Registry**
stores the **versioned**
Images
to create
Containers



Components



Image

A read-only snapshot of a container stored in a registry to be used as a template for building containers



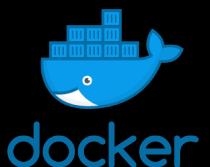
Container

The running, standard unit in which the application service resides or transported



Registry

Available in SaaS or Enterprise to deploy anywhere you choose
Stores, distributes, and shares container images

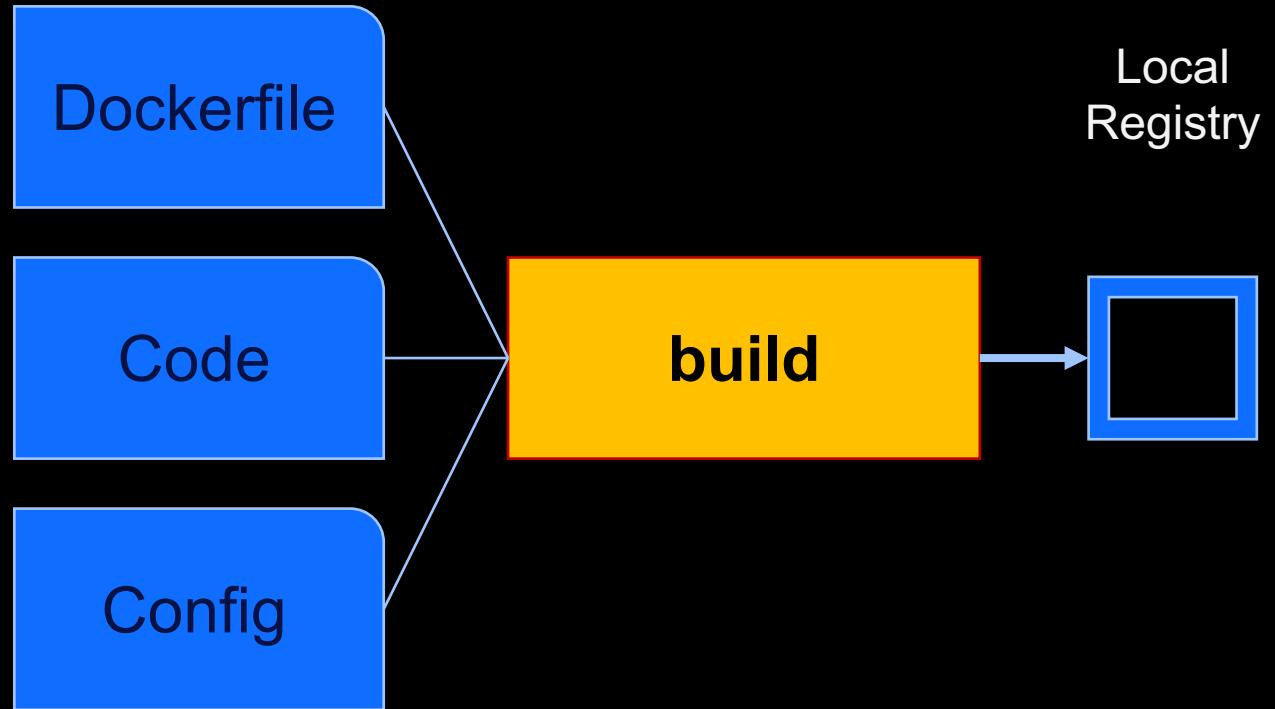


Docker Container Engine

A program that creates, ships, and runs application containers
Runs on any physical and virtual machine or server locally, in private or public cloud. Client communicates with Engine to execute commands

Create Images

Basics – Build



Basics – Build - Dockerfile

- Recipe - A text file that builds an image using directives

- FROM
- RUN
- COPY
- ENTRYPOINT
- LABEL
- ENV
- ARG
- USER
- EXPOSE

```
# Pull base image
FROM tomcat:8-jre8

# Maintainer
MAINTAINER "youremailaddress"

# Run command
RUN apt-get update && apt-get -y upgrade

# Set variables
ENV myName John Doe

# Copy to images tomcat path
ADD web.xml /usr/local/tomcat/conf/
ADD yourwarfile.war /usr/local/tomcat/webapps/

# Expose the port of the web application
EXPOSE 80

# Run server
CMD ["catalina.sh", "run"]
```

Base Image from:

- hub.docker.com
- quay.io
- your own

A base image is the image that is used to create all of your container images.

Start with an empty image:

- **FROM** scratch

Or use an official image, such as:

- **Ubuntu**: This is one of the most downloaded Docker images in the industry.
- **UBI**: Red Hat Universal Base images allow commercial and open source developers to build containers based on RHEL
- **Alpine**: Extremely small, for systems with constrained resources
- **Nginx**: Tiny web and proxy server
- ...

Basics – Build - Dockerfile

- Recipe - A text file that builds an image using directives

- FROM
- RUN
- COPY
- ENTRYPOINT
- LABEL
- ENV
- ARG
- USER
- EXPOSE

```
# Pull base image
FROM tomcat:8-jre8

# Maintainer
MAINTAINER "youremailaddress"

# Run command
RUN apt-get update && apt-get -y upgrade

# Set variables
ENV myName John Doe

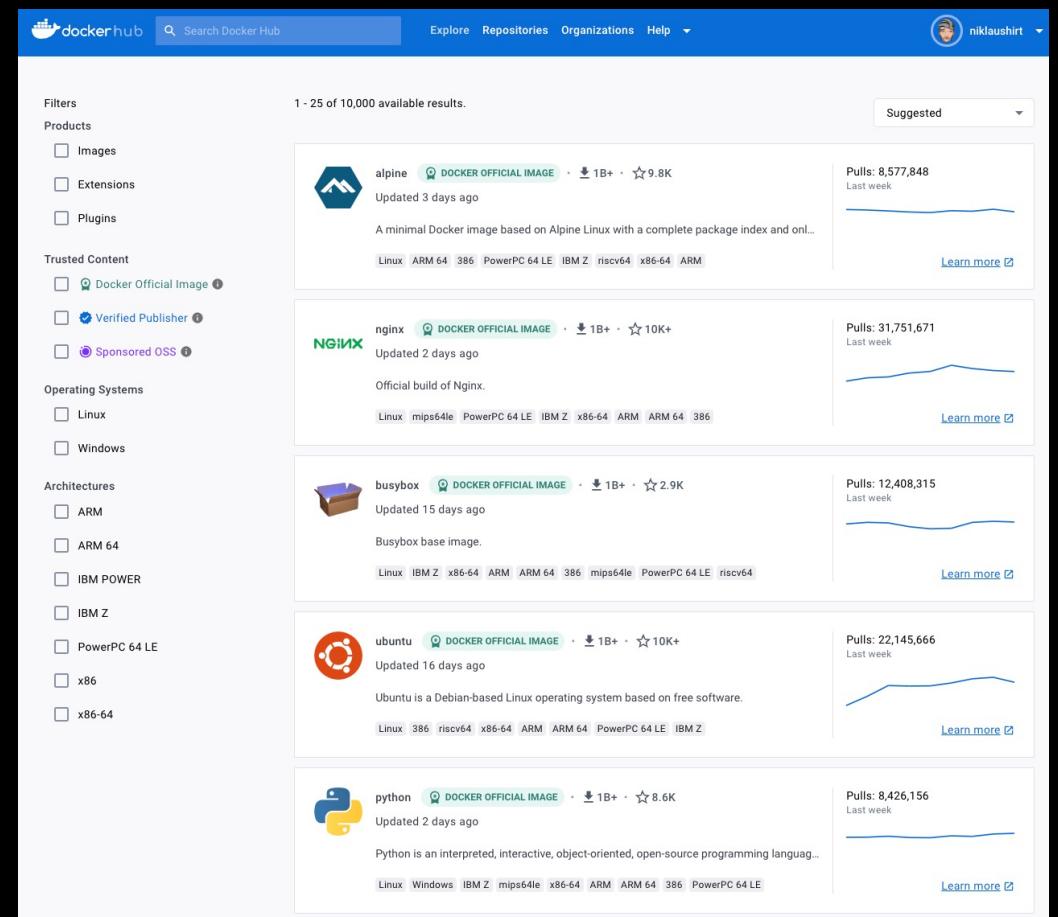
# Copy to images tomcat path
ADD web.xml /usr/local/tomcat/conf/
ADD yourwarfile.war /usr/local/tomcat/webapps/

# Expose the port of the web application
EXPOSE 80

# Run server
CMD ["catalina.sh", "run"]
```

Base Image from:

- hub.docker.com
- quay.io
- your own



Basics – Build - Dockerfile

- **Recipe** - A text file that builds an image using directives

- FROM
- RUN
- COPY
- ENTRYPOINT
- LABEL
- ENV
- ARG
- USER
- EXPOSE

```
# Pull base image
FROM tomcat:8-jre8

# Maintainer
MAINTAINER "youremailaddress"

# Run command
RUN apt-get update && apt-get -y upgrade

# Set variables
ENV myName John Doe

# Copy to images tomcat path
ADD web.xml /usr/local/tomcat/conf/
ADD yourwarfile.war /usr/local/tomcat/webapps/

# Expose the port of the web application
EXPOSE 80

# Run server
CMD ["catalina.sh", "run"]
```

Base Image from:

- hub.docker.com
- quay.io
- your own

Run commands

apt-get, yum, chown, mv, cp, ...

Set environment variables

Can be overridden when running

Add assets to the image

Jar, sources, golang app, ...

Expose a port of the process
running in the container

Startup command when image is
being run

Basics – Build - Dockerfile – USER

```
# Pull base image
FROM tomcat:8-jre8

# -----
# Switch to root user (not really needed, you are root by default)
USER root

# Do stuff that requires root rights
# Add an unprivileged user
RUN groupadd -r mygroup -g 433 && \
      useradd -u 431 -r -g mygroup -s /sbin/nologin myuser

# Do system stuff
RUN apt-get update && apt-get -y upgrade

# -----
# Switch to non-root user
USER myuser

# Do stuff that doesn't require root rights
RUN apt-get update && apt-get -y upgrade
...

# Run server
CMD ["catalina.sh", "run"]
```

By default, containers run as root

A container running as root has **full control of the host system**.

Requiring root can be dangerous for others and **may not be available in all environments**.

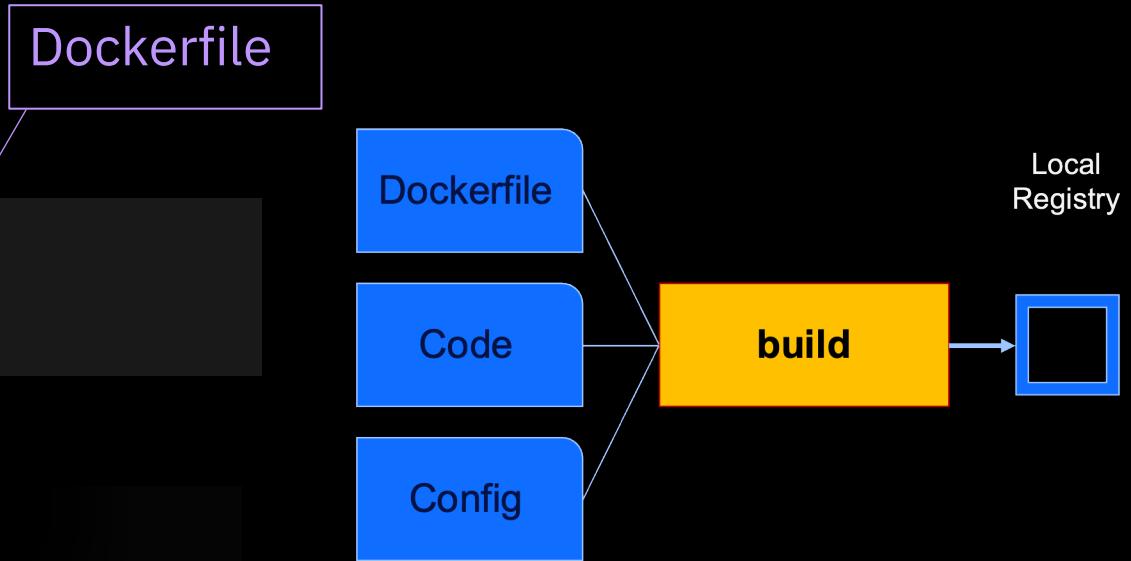
Use the **USER** instruction to specify a non-root user for containers to run as.

If your software does not create its own user, you can create a user and group in the Dockerfile.

Basics – Build - Dockerfile

```
> docker/podman build -t myimage:1.0.0 :
```

```
> podman build -t myimage:1.0.0 .
STEP 1/6: FROM tomcat:8-jre8
STEP 2/6: MAINTAINER "youremailaddress"
--> Using cache 360f554b3a76949389ffd63ecfde40bff00fb1705b387406a3fab5294438de47
--> Pushing cache []:f5fdb8f90292956a6fcc07763609d24be78c762ec7eecf448fc094a93322701a
--> 360f554b3a7
STEP 3/6: RUN apt-get update && apt-get -y upgrade
--> Using cache f050872e53e599219ff94c6db56237e6a73369675c7765bc4c953ee5cdd481cf
--> Pushing cache []:96bc0a0287094c44b2251ea2b967e1bcdcae8d499ede59a1db773a0bb8902ccb
--> f050872e53e
STEP 4/6: ENV myName John Doe
--> Using cache 809a87c3671b47d172ebbb5c2d0ca9944006d95a867d31b12edba0fad4e9880
--> Pushing cache []:005b263c3639ad8a692775141337cff489579082287a1c644679eea02bd7110f
--> 809a87c3671
STEP 5/6: EXPOSE 80
--> Pushing cache []:4fe61da890ff36b64a2cbe9f97929b80d715d938e1a65f04653c02276c1f1011
--> 52a988679c8
STEP 6/6: CMD ["catalina.sh", "run"]
COMMIT myimage:1.0.0
--> Pushing cache []:b2047af1e4593c24be782895a5878f76f1154e83f3be054cdb1a13f13f50419a
--> 9dcc2afafdf2
Successfully tagged localhost/myimage:1.0.0
9dcc2afafdf2a9fc97ed2daff4df34322f82ac2aabfeed0ce31d7c0c90bf3b2c
```



Images Filesystem

Basics – Layered File System

Layers Inheritance

- Build on the work of those who came before you

```
# Pull base image
FROM tomcat:8-jre8

# Maintainer
MAINTAINER "youremailaddress"

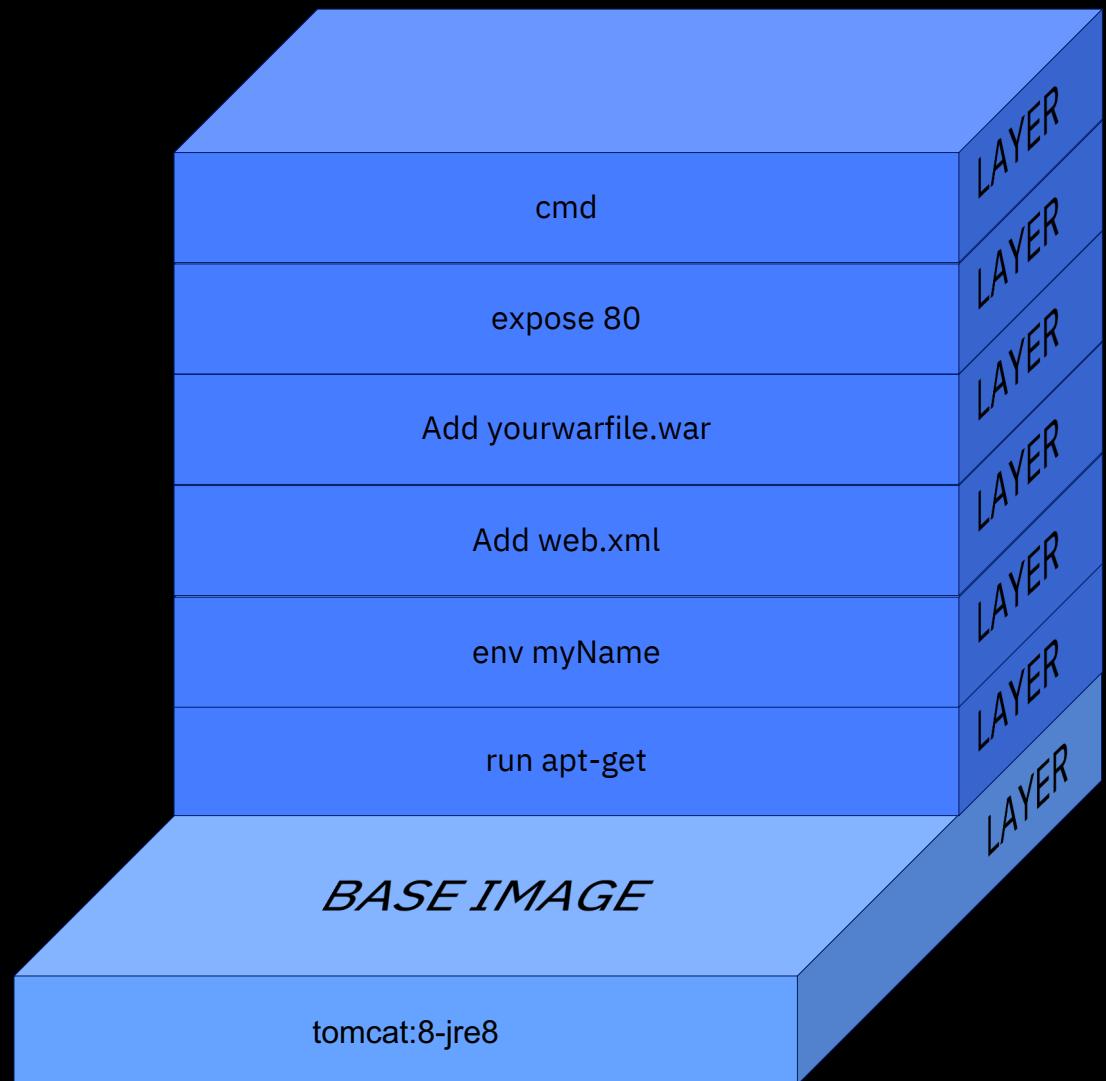
# Run command
RUN apt-get update && apt-get -y upgrade

# Set variables
ENV myName John Doe

# Copy to images tomcat path
ADD web.xml /usr/local/tomcat/conf/
ADD yourwarfile.war /usr/local/tomcat/webapps/

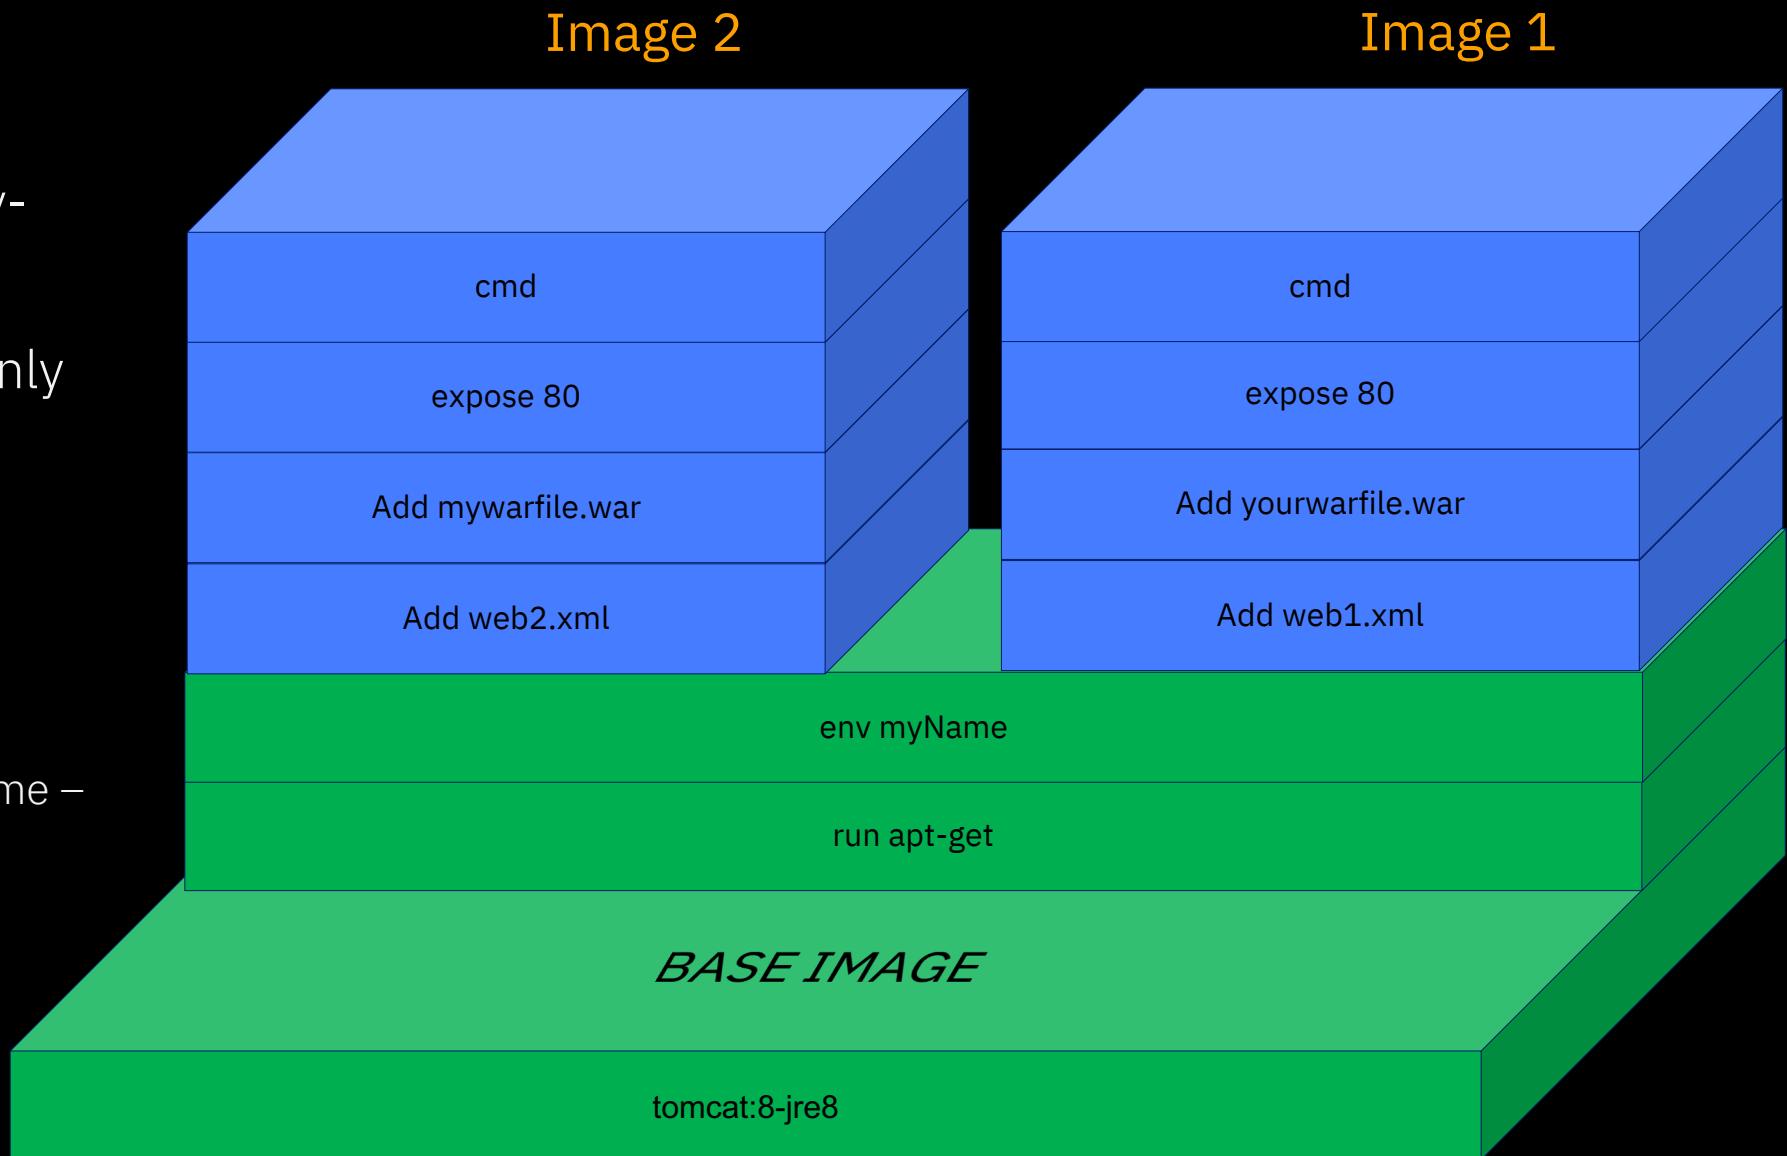
# Expose the port of the web application
EXPOSE 80

# Run server
CMD ["catalina.sh", "run"]
```



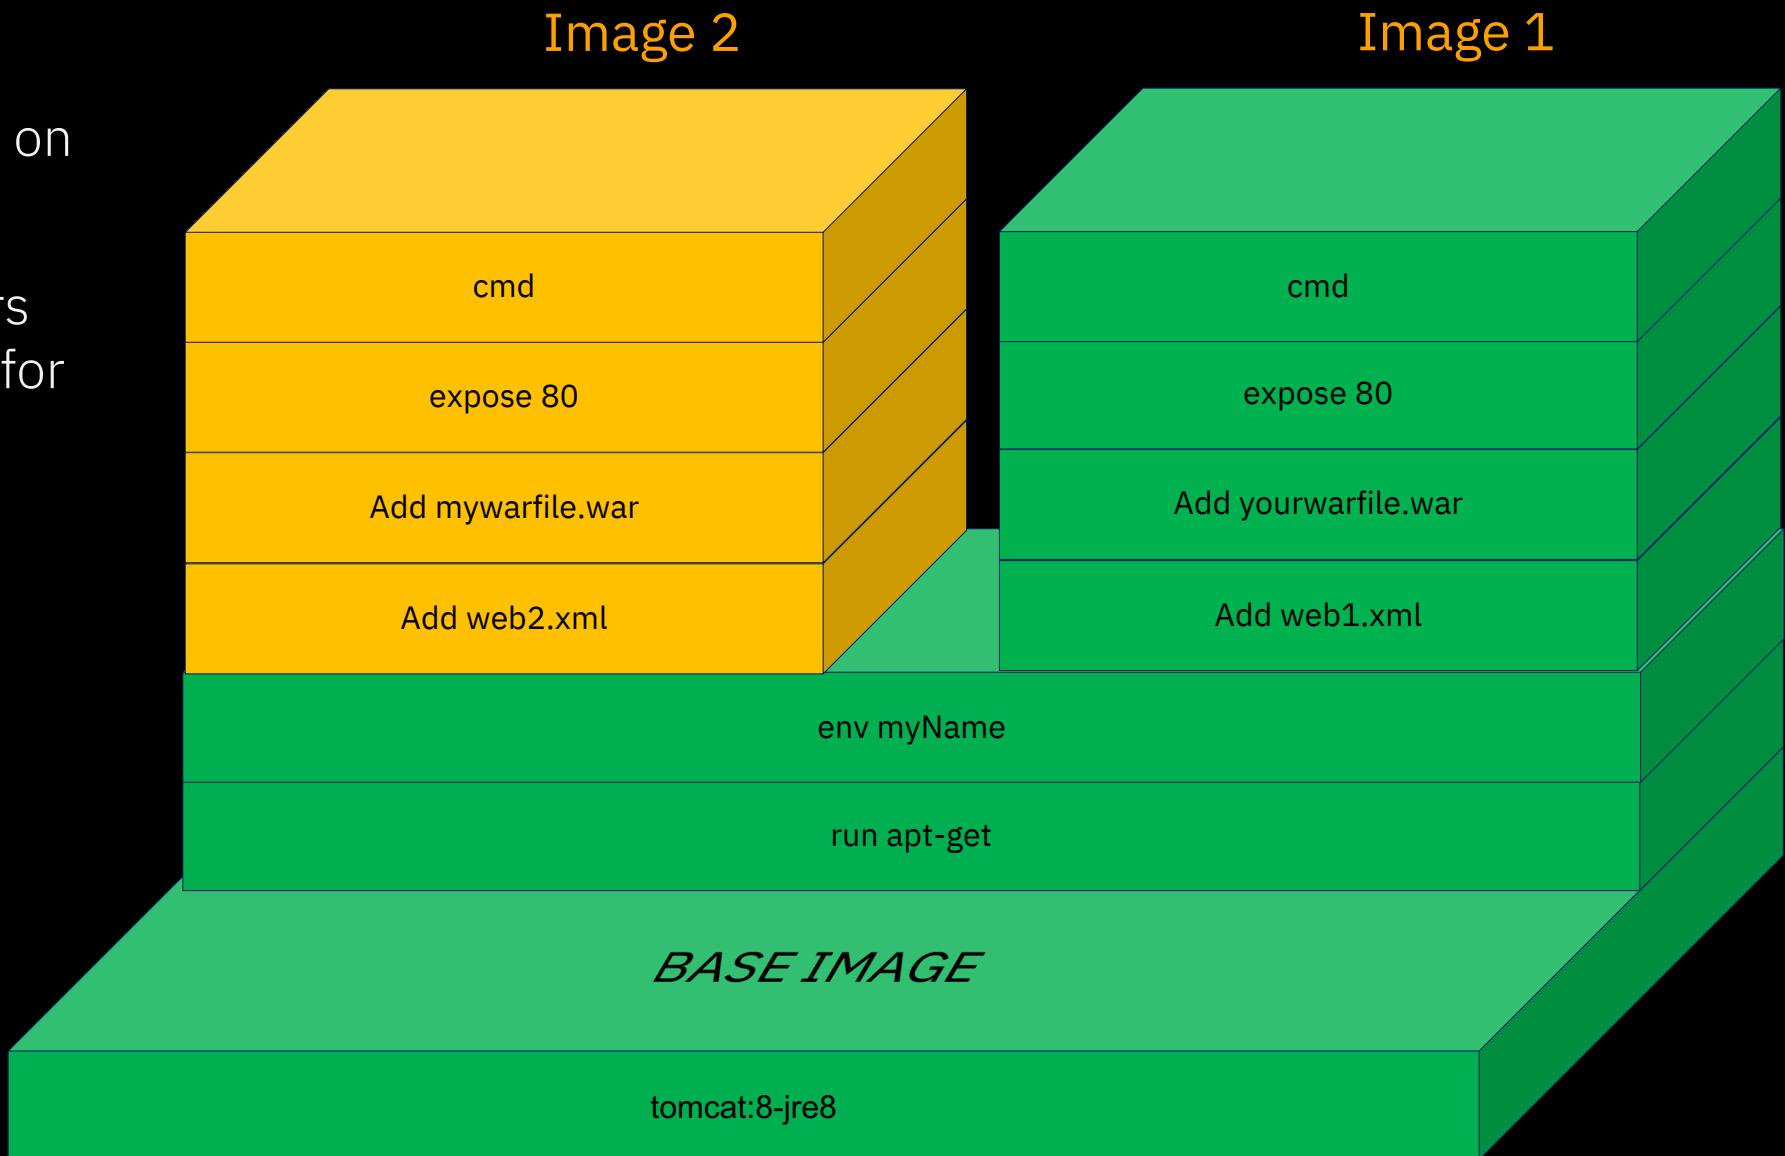
Basics – Layered File System

- Containers (OCI) use a [layered filesystem](#) (copy-on-write)
- New files (& edits) are only visible to current/above layers
- Layers allow for [reuse](#)
 - More containers per host
 - Faster start-up/download time – base layers are "cached"



Basics – Layered File System

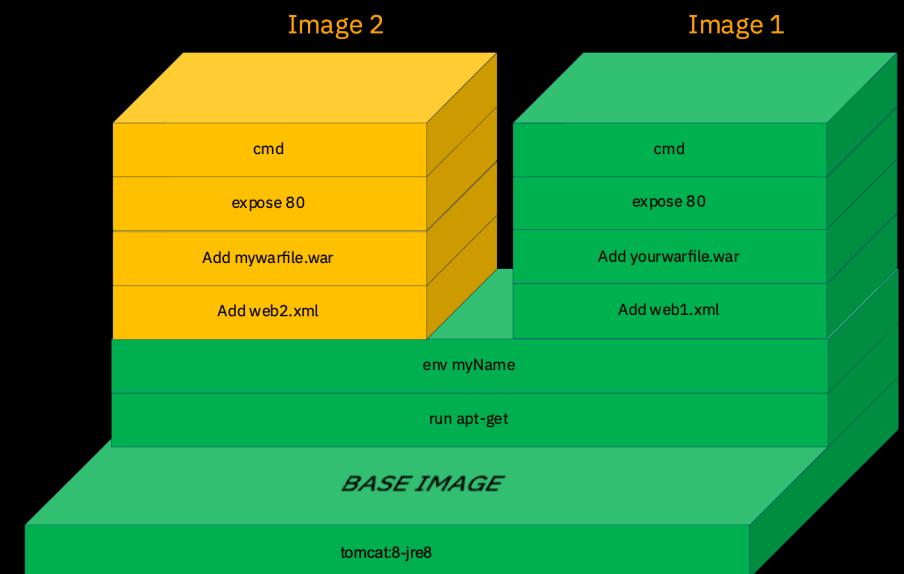
- If I already had **Image 1** on my local system
- Only the additional layers need to be downloaded for **Image 2**



Basics – Build - Dockerfile

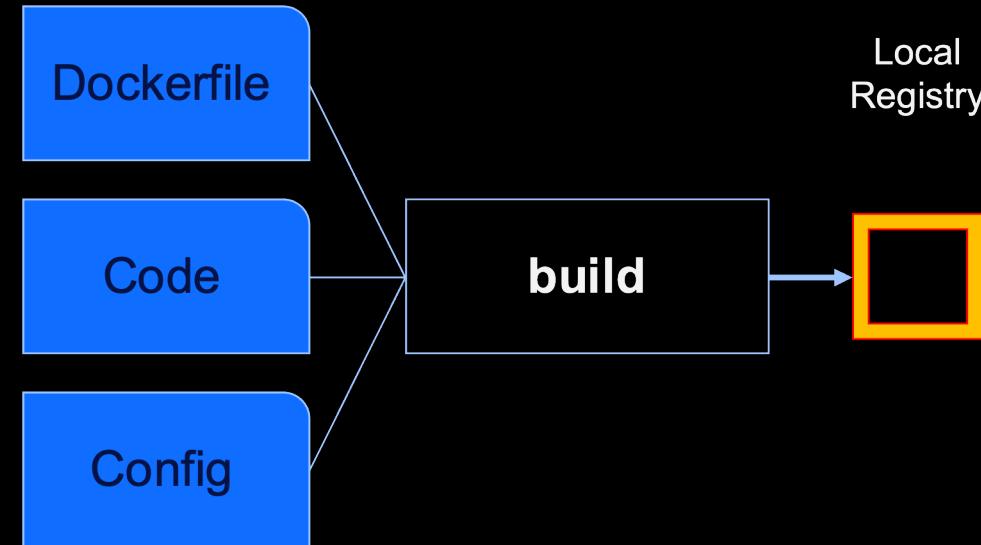
```
> podman build -t myimage:1.0.0 .
```

```
> podman build -t myimage:1.0.0 .
STEP 1/6: FROM tomcat:8-jre8
STEP 2/6: MAINTAINER "youremailaddress"
--> Using cache 360f554b3a76949389ffd63ecfde40bff00fb1705b387406a3fab5294438de47
--> Pushing cache []:f5fdb8f90292956a6fcc07763609d24be78c762ec7eecf448fc094a93322701a
--> 360f554b3a7
STEP 3/6: RUN apt-get update && apt-get -y upgrade
--> Using cache f050872e53e599219ff94c6db56237e6a73369675c7765bc4c953ee5cdd481cf
--> Pushing cache []:96bc0a0287094c44b2251ea2b967e1bcd8ae8d499ede59a1db773a0bb8902ccb
--> f050872e53e
STEP 4/6: ENV myName John Doe
--> Using cache 809a87c3671b47d172ebbb5c2d0ca9944006d95a867d31b12edba0fad84e9880
--> Pushing cache []:005b263c3639ad8a692775141337cff489579082287a1c644679eea02bd7110f
--> 809a87c3671
STEP 5/6: EXPOSE 80
--> Pushing cache []:4fe61da890ff36b64a2cbe9f97929b80d715d938e1a65f04653c02276c1f1011
--> 52a988679c8
STEP 6/6: CMD ["catalina.sh", "run"]
COMMIT myimage:1.0.0
--> Pushing cache []:b2047af1e4593c24be782895a5878f76f1154e83f3be054cdb1a13f13f50419a
--> 9dcc2afafdf2
Successfully tagged localhost/myimage:1.0.0
9dcc2afafdf2a9fc97ed2daff4df34322f82ac2aabfeed0ce31d7c0c90bf3b2c
```



Basics – Build – Local Images

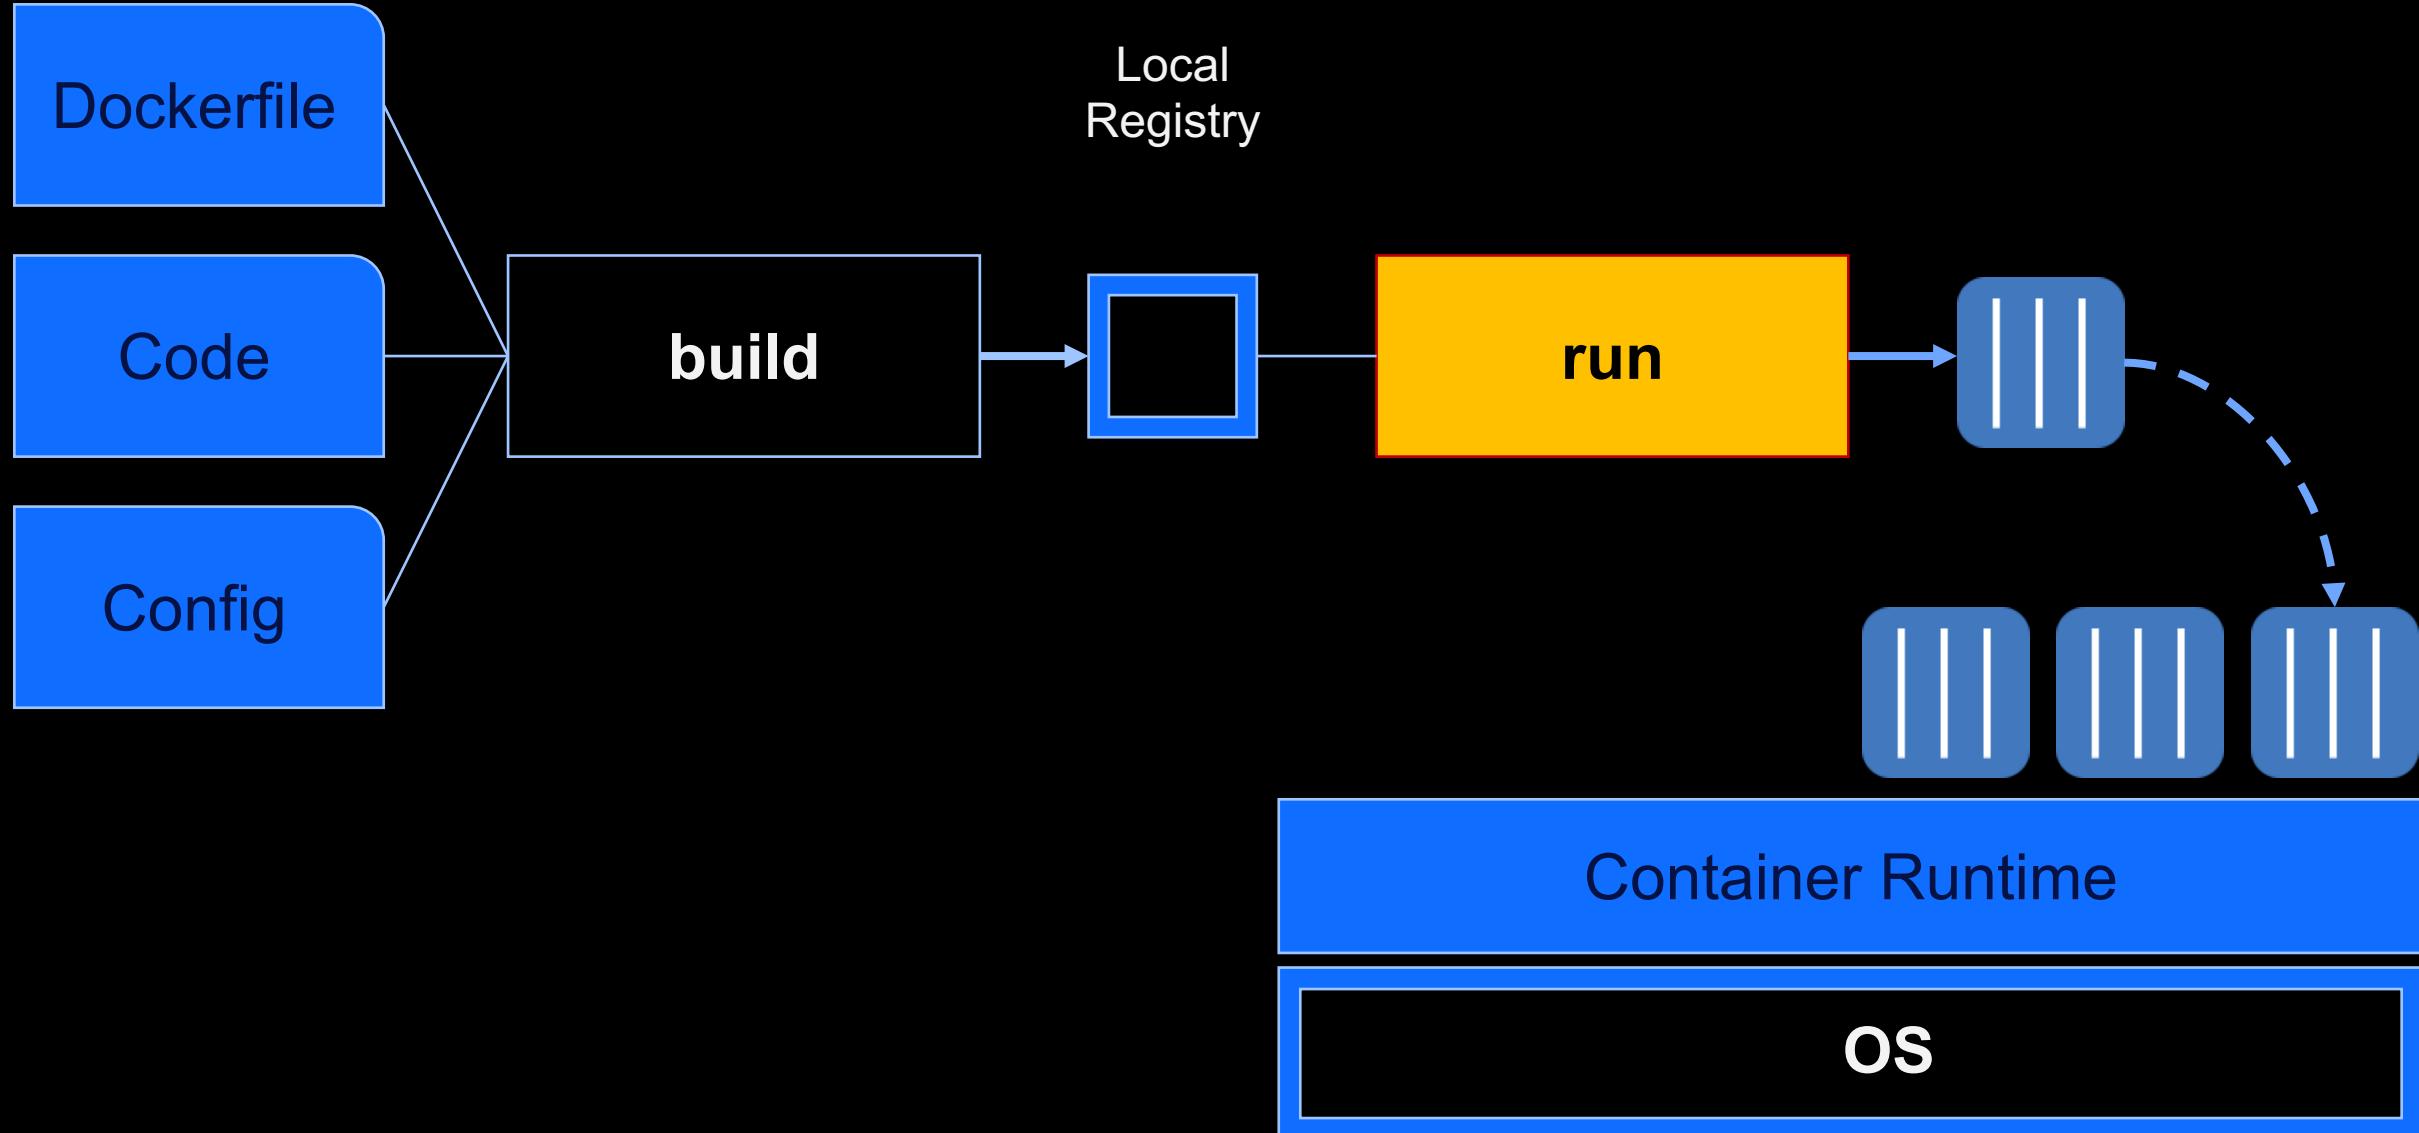
```
> podman images  
> podman rmi 569118d95a3e
```



REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
localhost/myimage	1.0.0	569118d95a3e	16 minutes ago	276 MB
localhost/niklaushirt/myimage	1.0.0	569118d95a3e	16 minutes ago	276 MB
localhost/k8sdemo-backend	lab	a1e30bbb2f45	41 minutes ago	1.36 GB
docker.io/library/tomcat	8-jre8	6802686d7abd	33 hours ago	235 MB
docker.io/library/node	8-stretch	50c57a9369c7	3 years ago	879 MB

Run Containers

Basics – Run

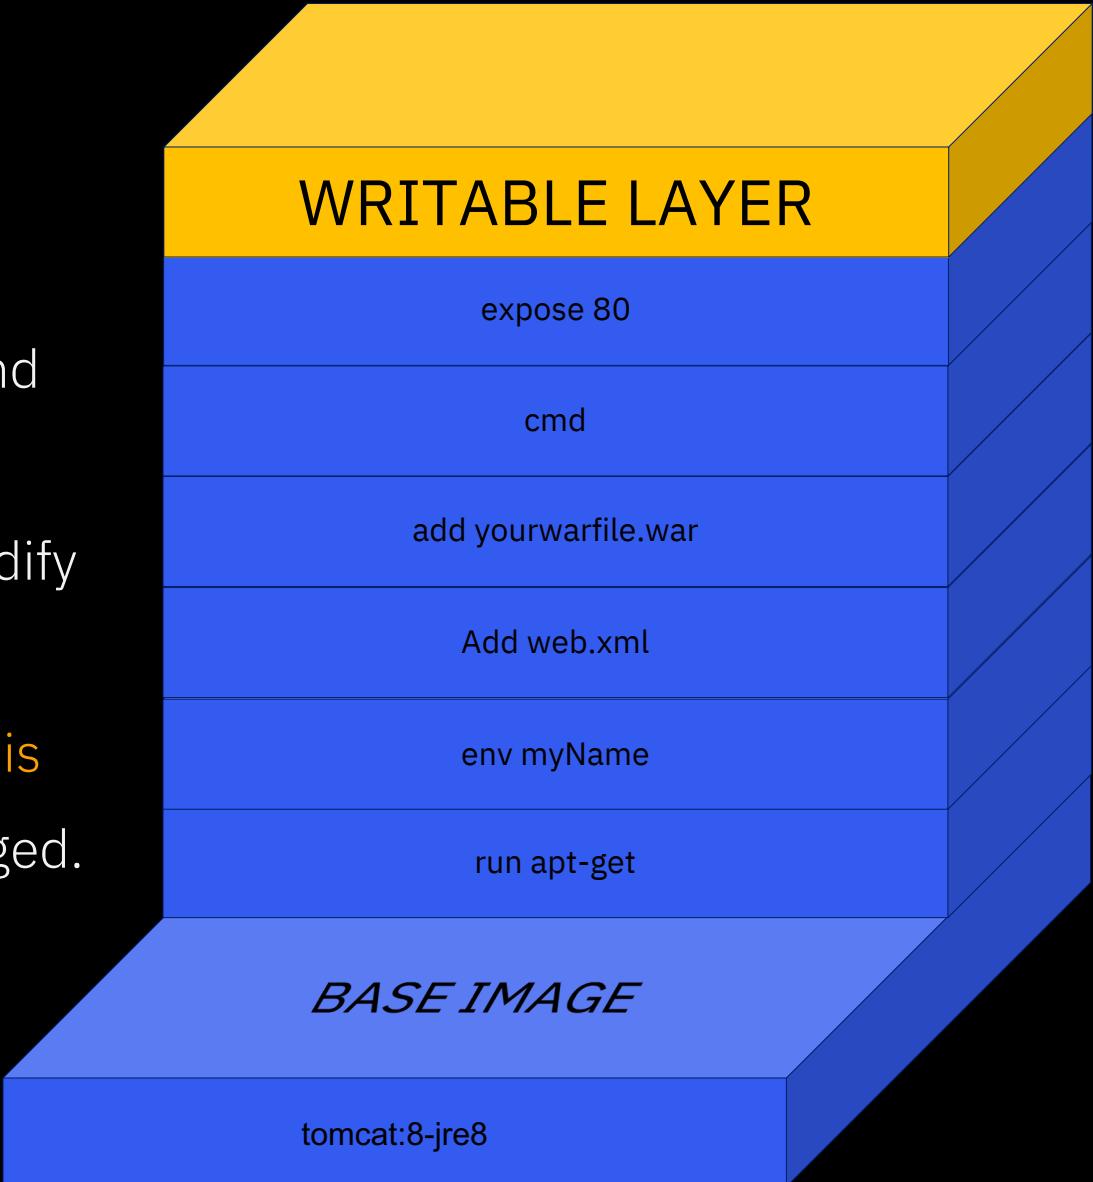


Basics – Run

Layers

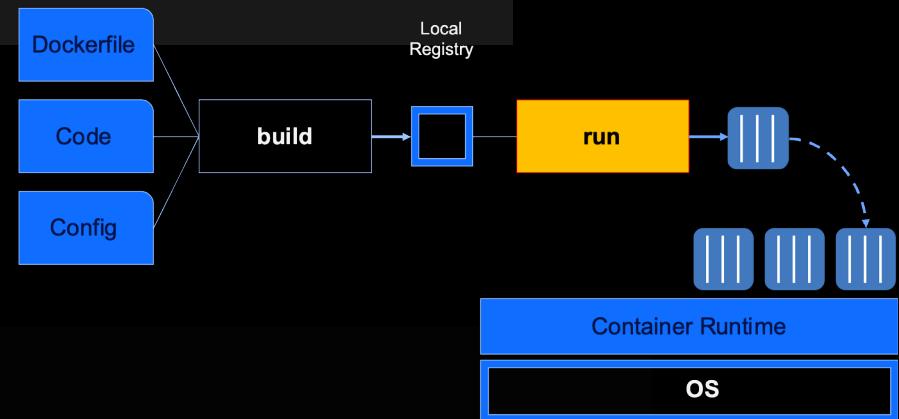
Inheritance

- The biggest difference between a [container](#) and an [image](#) is the [top writable layer](#).
- All writes to the container that add new or modify existing data are stored in this writable layer.
- When a container is deleted, its writable layer is also deleted. The underlying image is unchanged.



Basics – Run

```
> podman run --rm -ti myimage:1.0.0 /bin/bash  
  
> podman run --rm -d -e MYVAR=foo -p 8080:80 myimage:1.0.0  
  
> podman ps [-a]  
  
> podman stop/kill <CONTAINER_ID>
```



```
> podman run --rm -d --name myimage -p 8080:8080 myimage:1.0.0
```

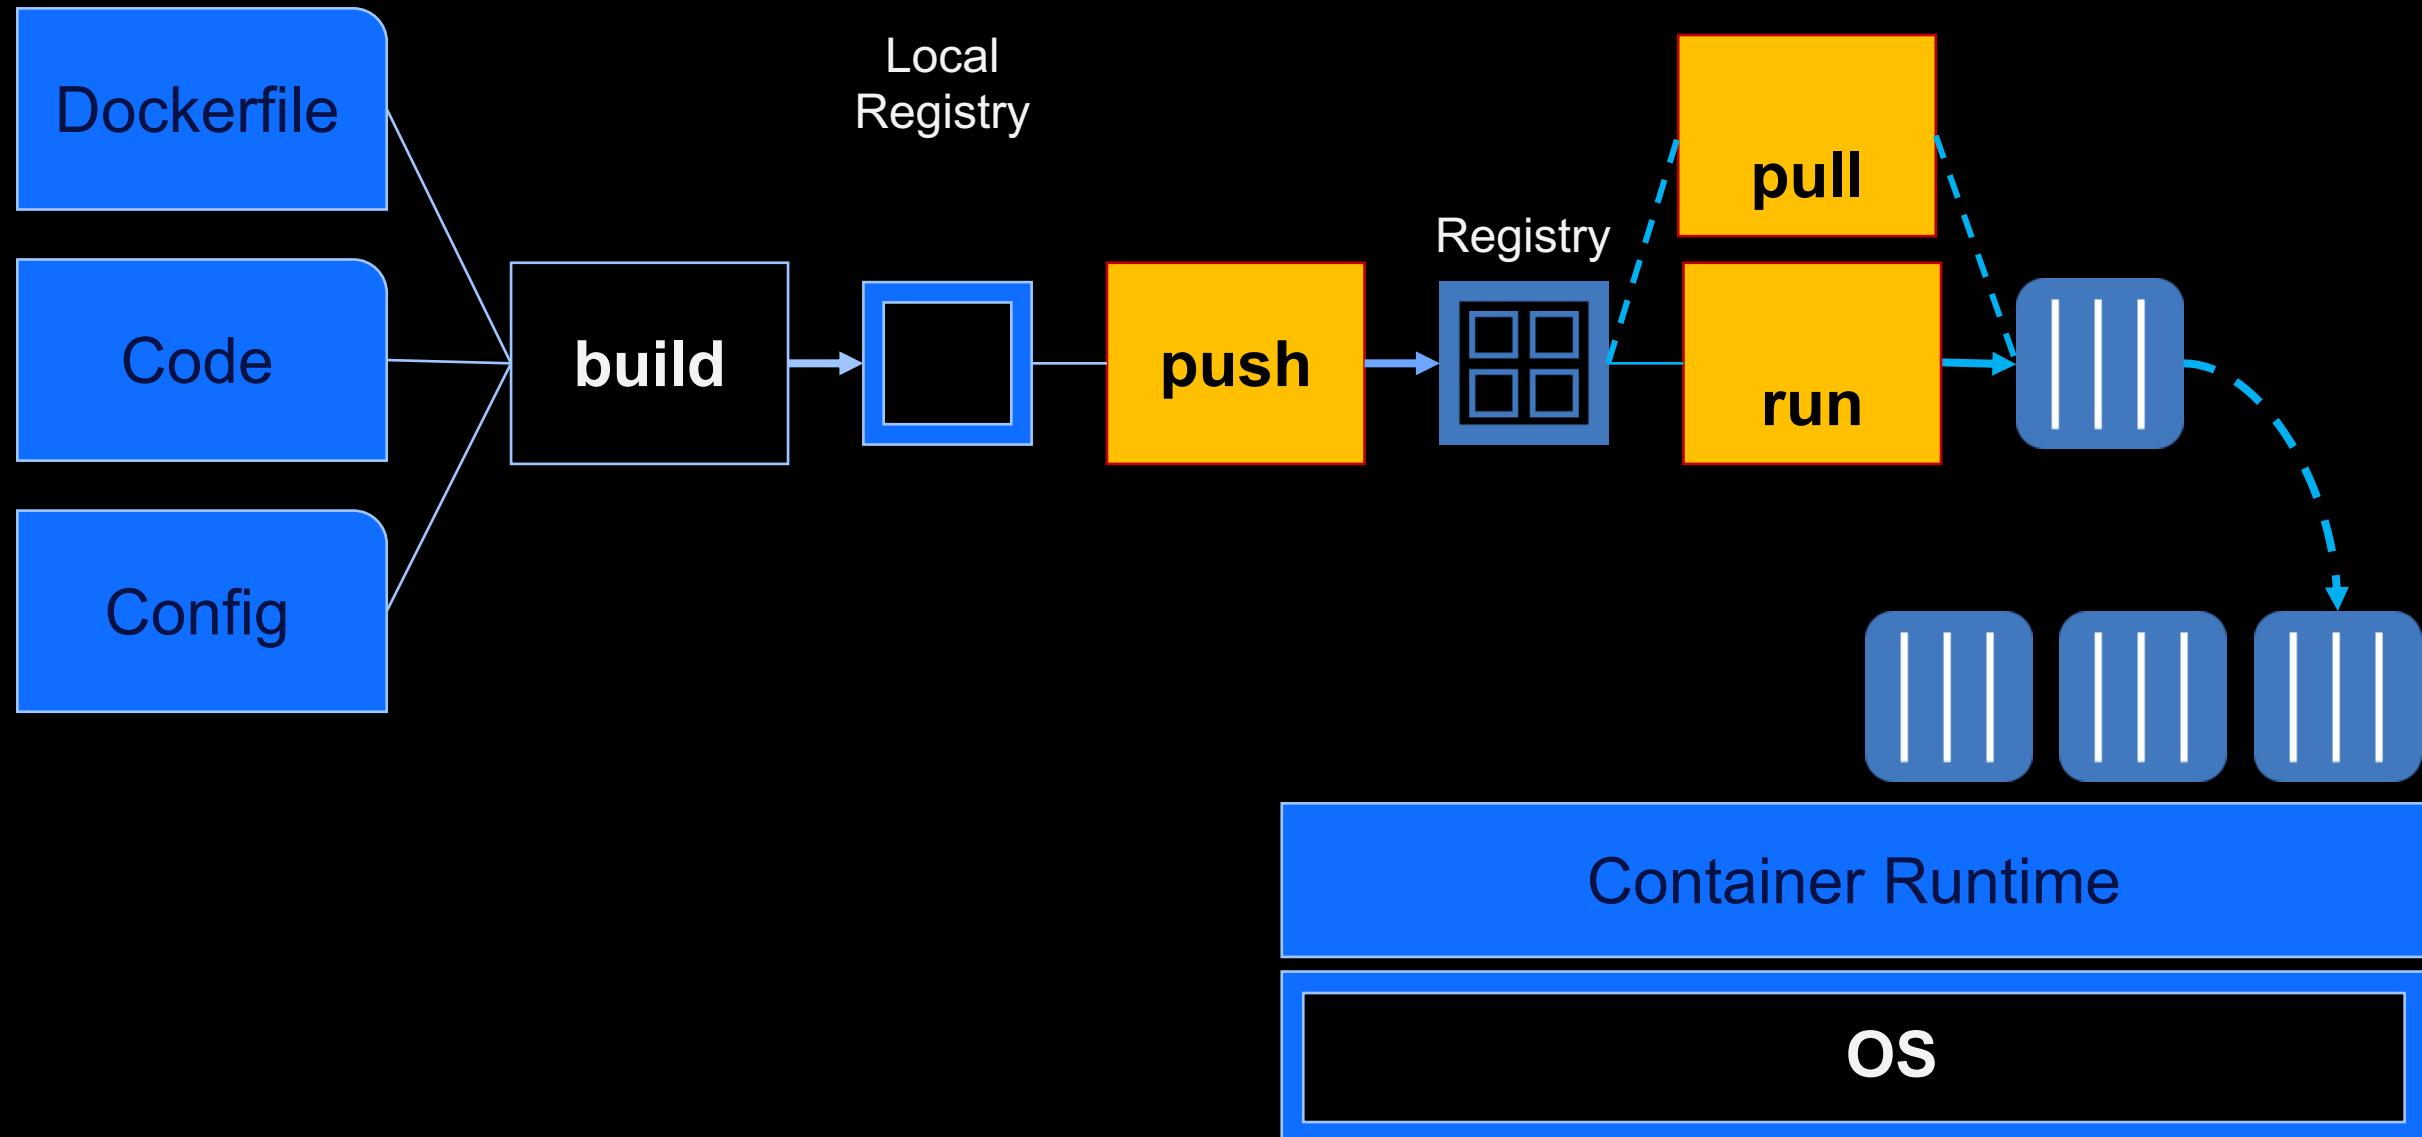
```
b3812ee218eb64bde9c8cd06a4a8cf89b0c499af1a62004ccb109f0d831934a8
```

```
> podman ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
b3812ee218eb	localhost/myimage:1.0.0	catalina.sh run	4 seconds ago	Up 4 seconds	0.0.0.0:8080->8080/tcp	myimage
> podman kill b3812ee218eb	b3812ee218eb					

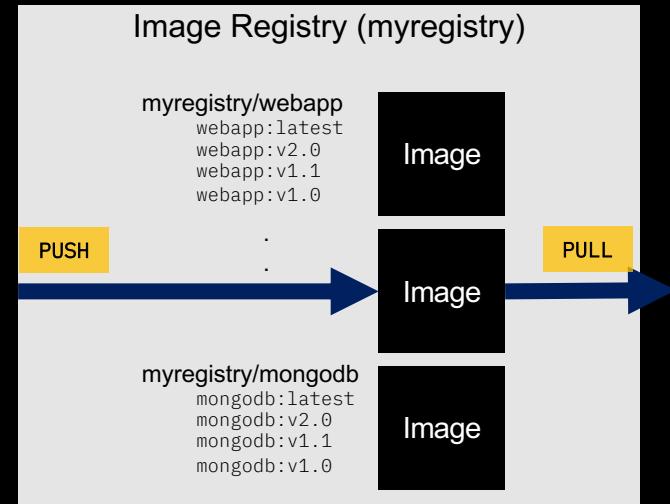
Store Images

Basics – Store, Retrieve & Run with registry



Basics – Store & Retrieve with registry

```
> podman login docker.io/quay.io  
> podman tag myimage:1.0.0 myregistry/myimage:1.0.0  
> podman push myregistry/myimage:1.0.0  
  
> podman pull myregistry/myimage:latest
```



```
> podman login docker.io  
Authenticating with existing credentials for docker.io  
Existing credentials are valid. Already logged in to docker.io  
> podman tag myimage:1.0.0 niklaushirt/myimage:1.0.0  
> podman images  


| REPOSITORY                    | TAG       | IMAGE ID     | CREATED        | SIZE    |
|-------------------------------|-----------|--------------|----------------|---------|
| localhost/myimage             | 1.0.0     | 569118d95a3e | 20 minutes ago | 276 MB  |
| localhost/niklaushirt/myimage | 1.0.0     | 569118d95a3e | 20 minutes ago | 276 MB  |
| localhost/k8sdemo-backend     | lab       | a1e30bbb2f45 | 46 minutes ago | 1.36 GB |
| docker.io/library/tomcat      | 8-jre8    | 6802686d7abd | 33 hours ago   | 235 MB  |
| docker.io/library/node        | 8-stretch | 50c57a9369c7 | 3 years ago    | 879 MB  |

  
> podman push niklaushirt/myimage:1.0.0  
Getting image source signatures  
Copying blob sha256:27951c26c895e6b2d302273f91369f3e7bd21430d5d6678ba7c322bc08fdfb60  
Copying blob sha256:874b048c963ab55b06939c39d59303fb975d323822a4ea48a02ac8dc635ea371
```

Basics – Registries

Hosting image repositories

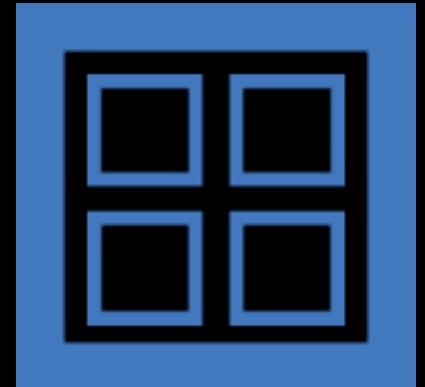
- You can create your own registry
- A registry is managed by a registry container

Public and Private registries

- Public Registry like Quay or Docker Hub
- <https://quay.io>
- <https://hub.docker.com>

Login into the registry

- docker login domain:port
- podman login domain:port



Recap

- Containers are not VMs
- Containers provide many benefits:
 - Efficiency
 - Portability
 - Consistency
- New challenges with containers:
 - Production apps dependent on open-source projects
 - Existing tools may not be sufficient for containers
 - Manage containers when you get to have more than a handful

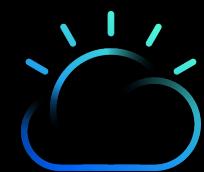
Control Groups

For limiting resources (cgroups)

+ Namespaces

For filesystem isolation (chroot)

= Containers



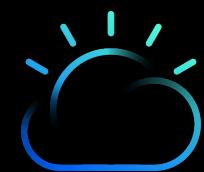
Notes from last course

Build images for other platforms (on M1 for example)

```
> podman buildx build --platform linux/amd64 -t myimage:x86 --load .
```

Build multi architecture images

```
> podman manifest rm myimage-mf  
  
> podman build --platform linux/arm64/v8 --platform linux/amd64  
      -t myimage:1.0.0 --manifest myimage-mf .  
  
> podman manifest inspect myimage-mf  
> podman manifest push myimage-mf myimage:1.0.0
```



Notes from last course

Run Portainer

```
> podman run -d --rm  
  -p 8000:8000 -p 9001:9000 --name=portainer  
  -v portainer_data:/data  
  -v /run/podman/podman.sock:/var/run/docker.sock:Z  
  --privileged  
  portainer/portainer-ce:2.17.0-alpine
```

QUESTIONS?



Kubernetes Workshop Series

Kubernetes - Basics

04



What happened so far... Everybody Loves Containers



A standard way to package an application and all its dependencies so that it can be moved between environments and run without changes.

Everybody Loves Containers

But when you go



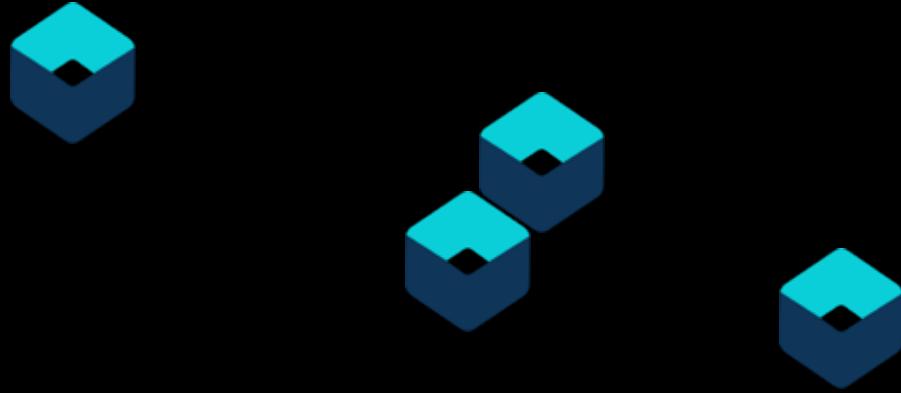
From this....



To this....



Everyone's container journey starts with one container....



At first the growth is easy to handle....



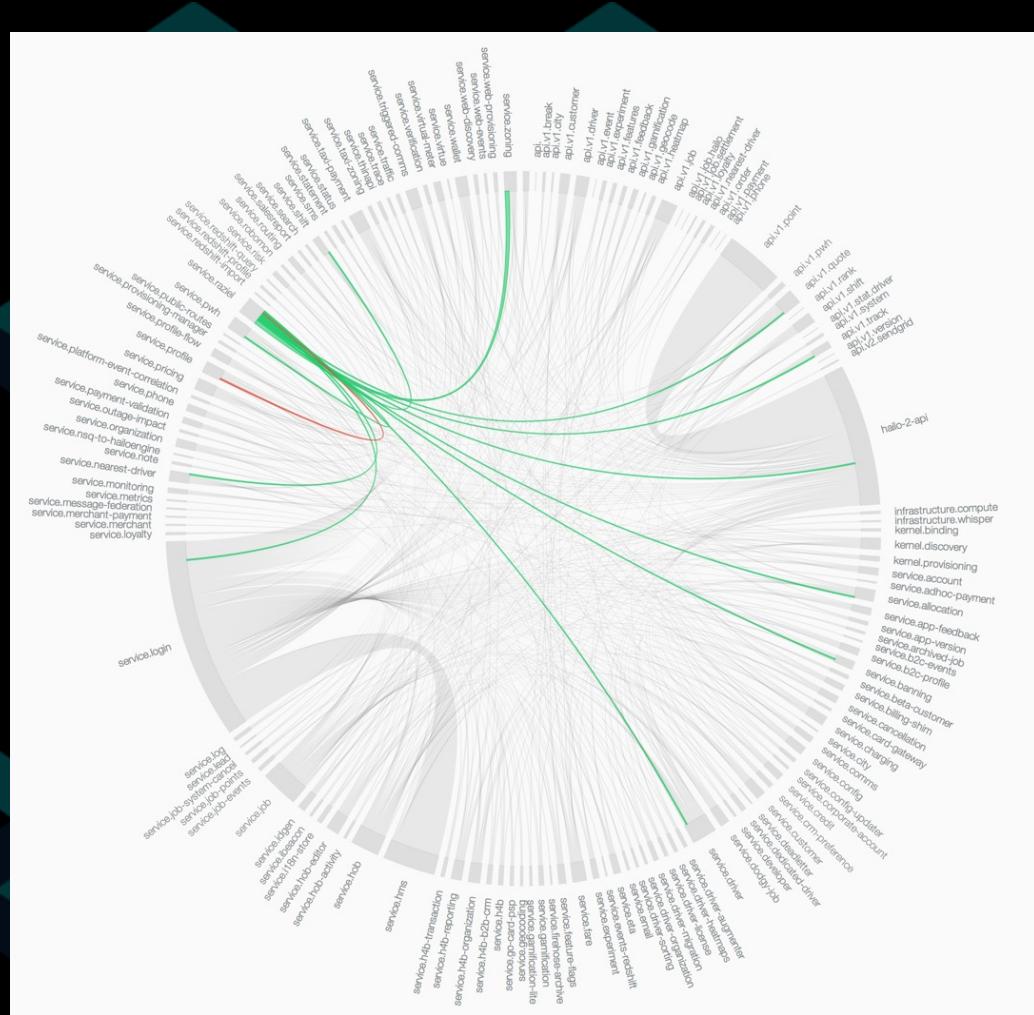
Pets vs Cattle

That's how it always starts.

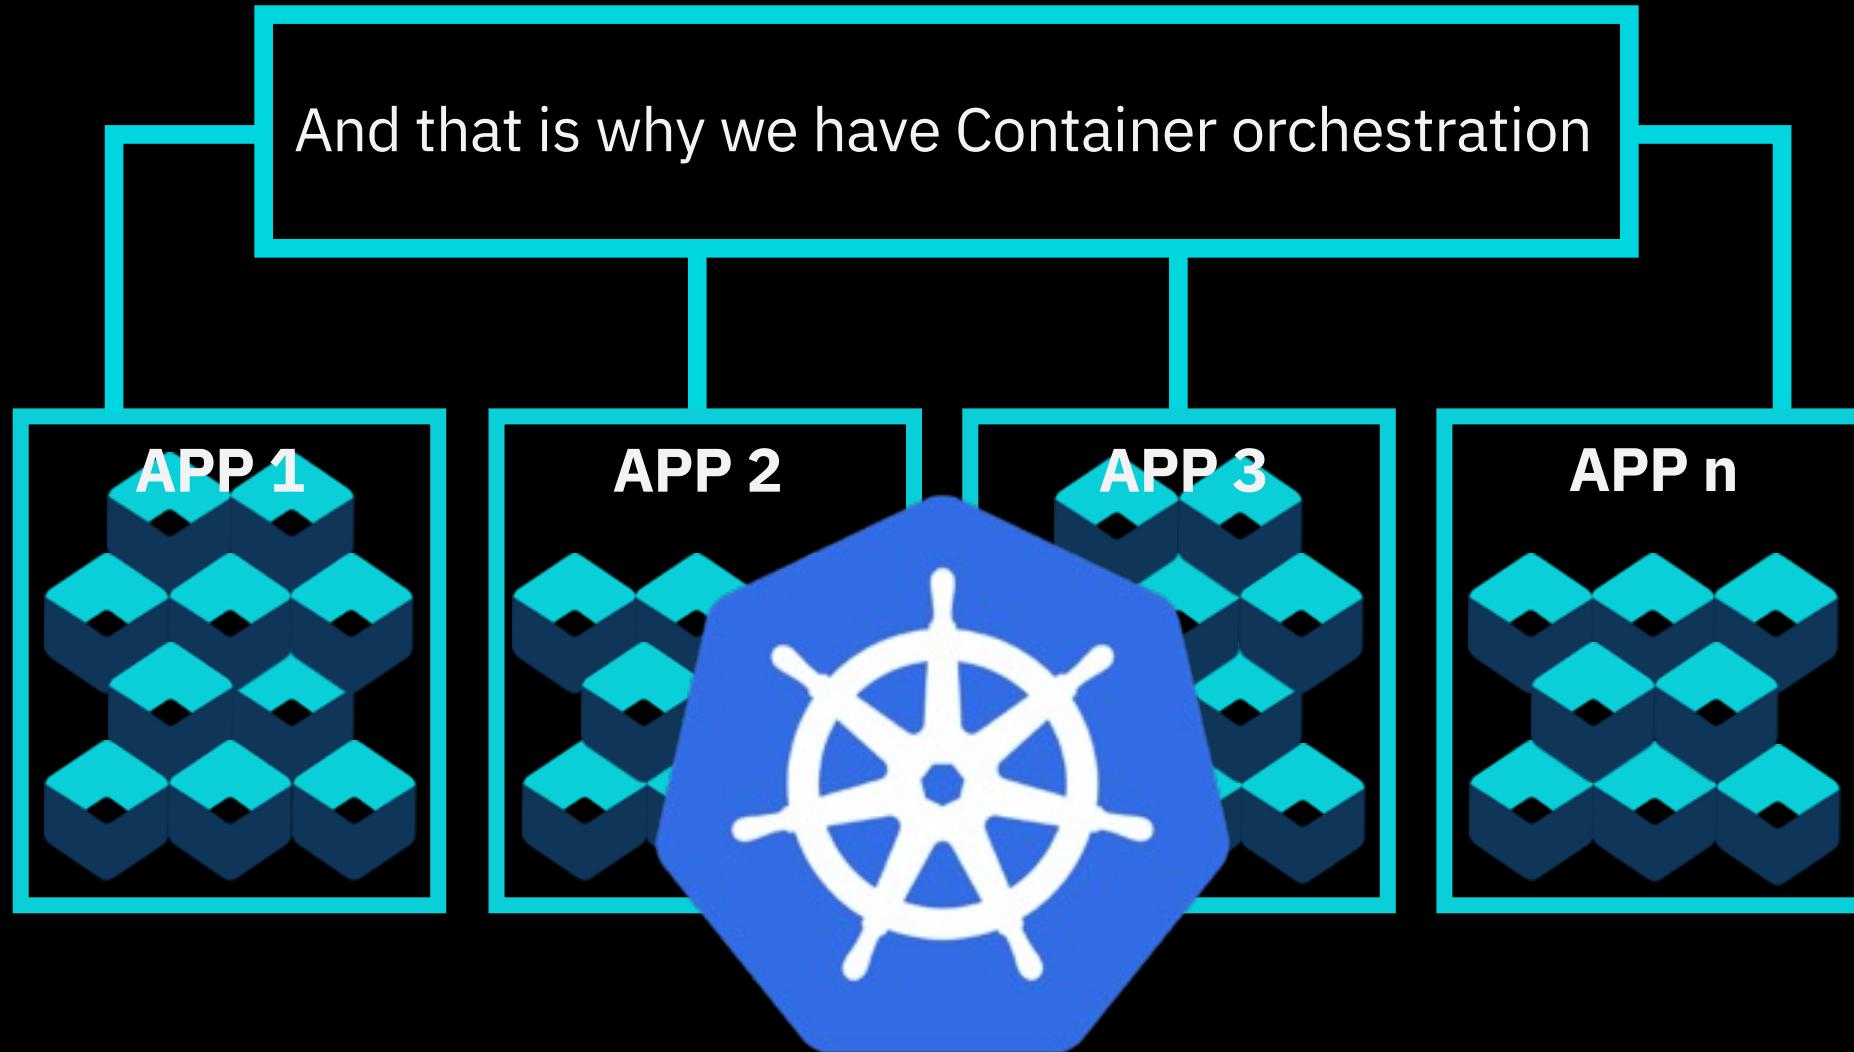
Then later there's running and screaming...

The trade off

Improved delivery velocity
in exchange for
increased operational
complexity



Enter Kubernetes (K8s)



Kubernetes – Declarative System



Imperative Systems

In an imperative system, the user knows the desired state and the user determines the sequence of commands to transition the system to the desired state.

Declarative Systems

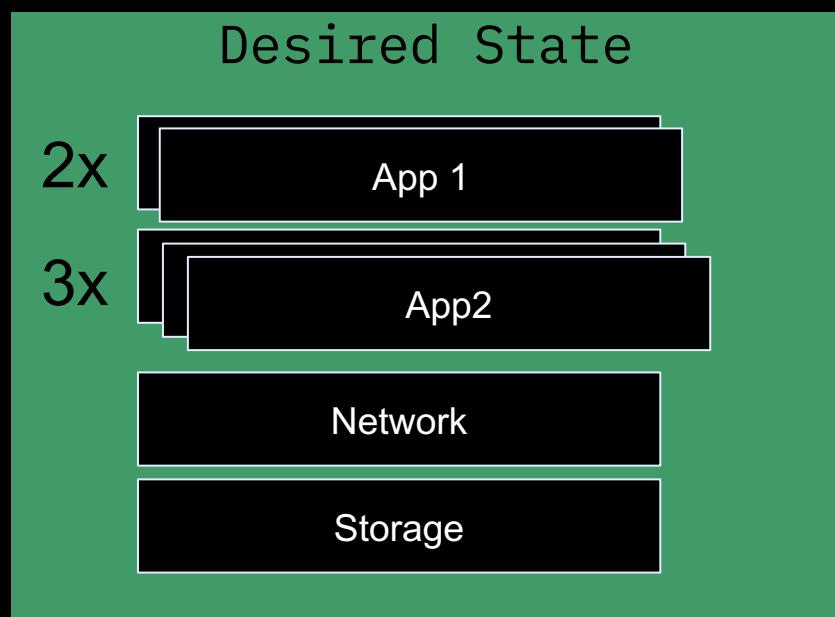
By contrast, in a declarative system, the user knows the desired state, supplies a representation of the desired state to the system, then the system reads the current state and determines the sequence of commands to transition the system to the desired state.

Kubernetes – Declarative System

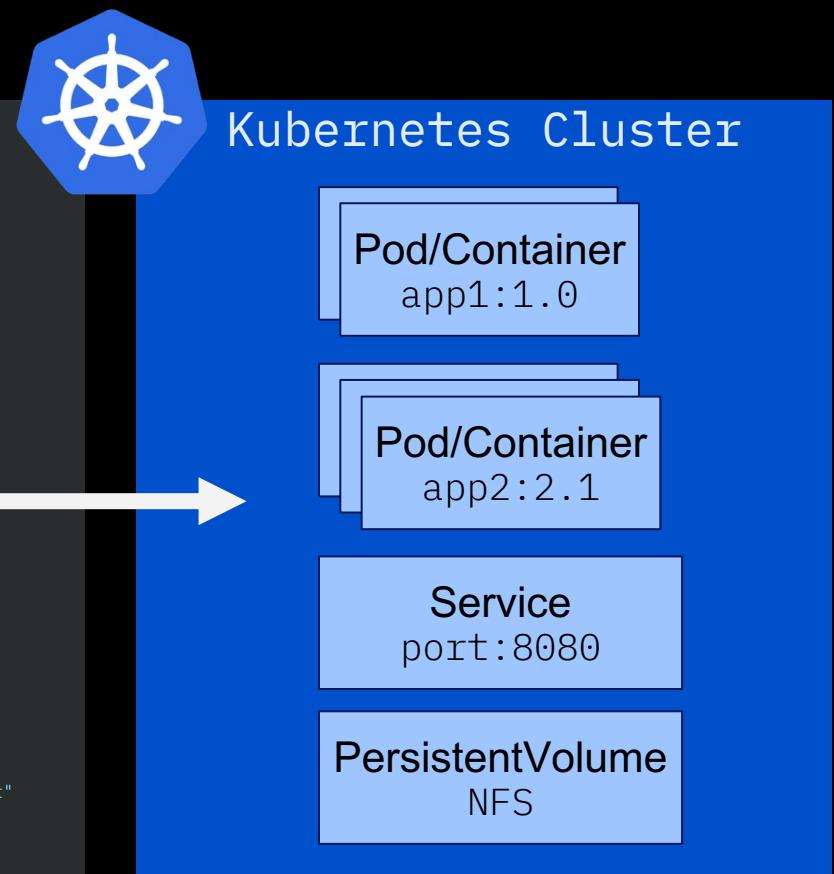


The Desired State

Kubernetes ensures that all the containers running across the cluster are in the desired state at any moment.



```
kind: Deployment
apiVersion: extensions/v1beta1
metadata:
  name: mcmk-ibm-mcmk-prod-klusterlet
  labels:
    app: ibm-mcmk-prod
    chart: ibm-mcmk-prod-3.1.2
    component: "klusterlet"
    release: mcmk
    heritage: Tiller
spec:
  replicas: 1
  selector:
    matchLabels:
      app: ibm-mcmk-prod
      component: "klusterlet"
      release: mcmk
  template:
    metadata:
      labels:
        app: ibm-mcmk-prod
        component: "klusterlet"
        release: mcmk
        chart: ibm-mcmk-prod-3.1.2
      annotations:
        productName: "IBM Multi-cloud Manager - Klusterlet"
        productID: "354b8990aab44c9988a0edfda101b128"
        productVersion: "3.1.2"
    spec:
```



Kubernetes Strengths



- Kubernetes has a **clear governance model** managed by the Linux Foundation
- A growing and **vibrant** Kubernetes **ecosystem**
- Kubernetes **avoids** dependencies
- Kubernetes supports a **wide range** of deployment options

K8s Architecture

Kubernetes Management Architecture



Nodes – hosts (Baremetal/VMs) that run Kubernetes applications

Master nodes

- Controls and manages the cluster
- Scheduling and replication logic

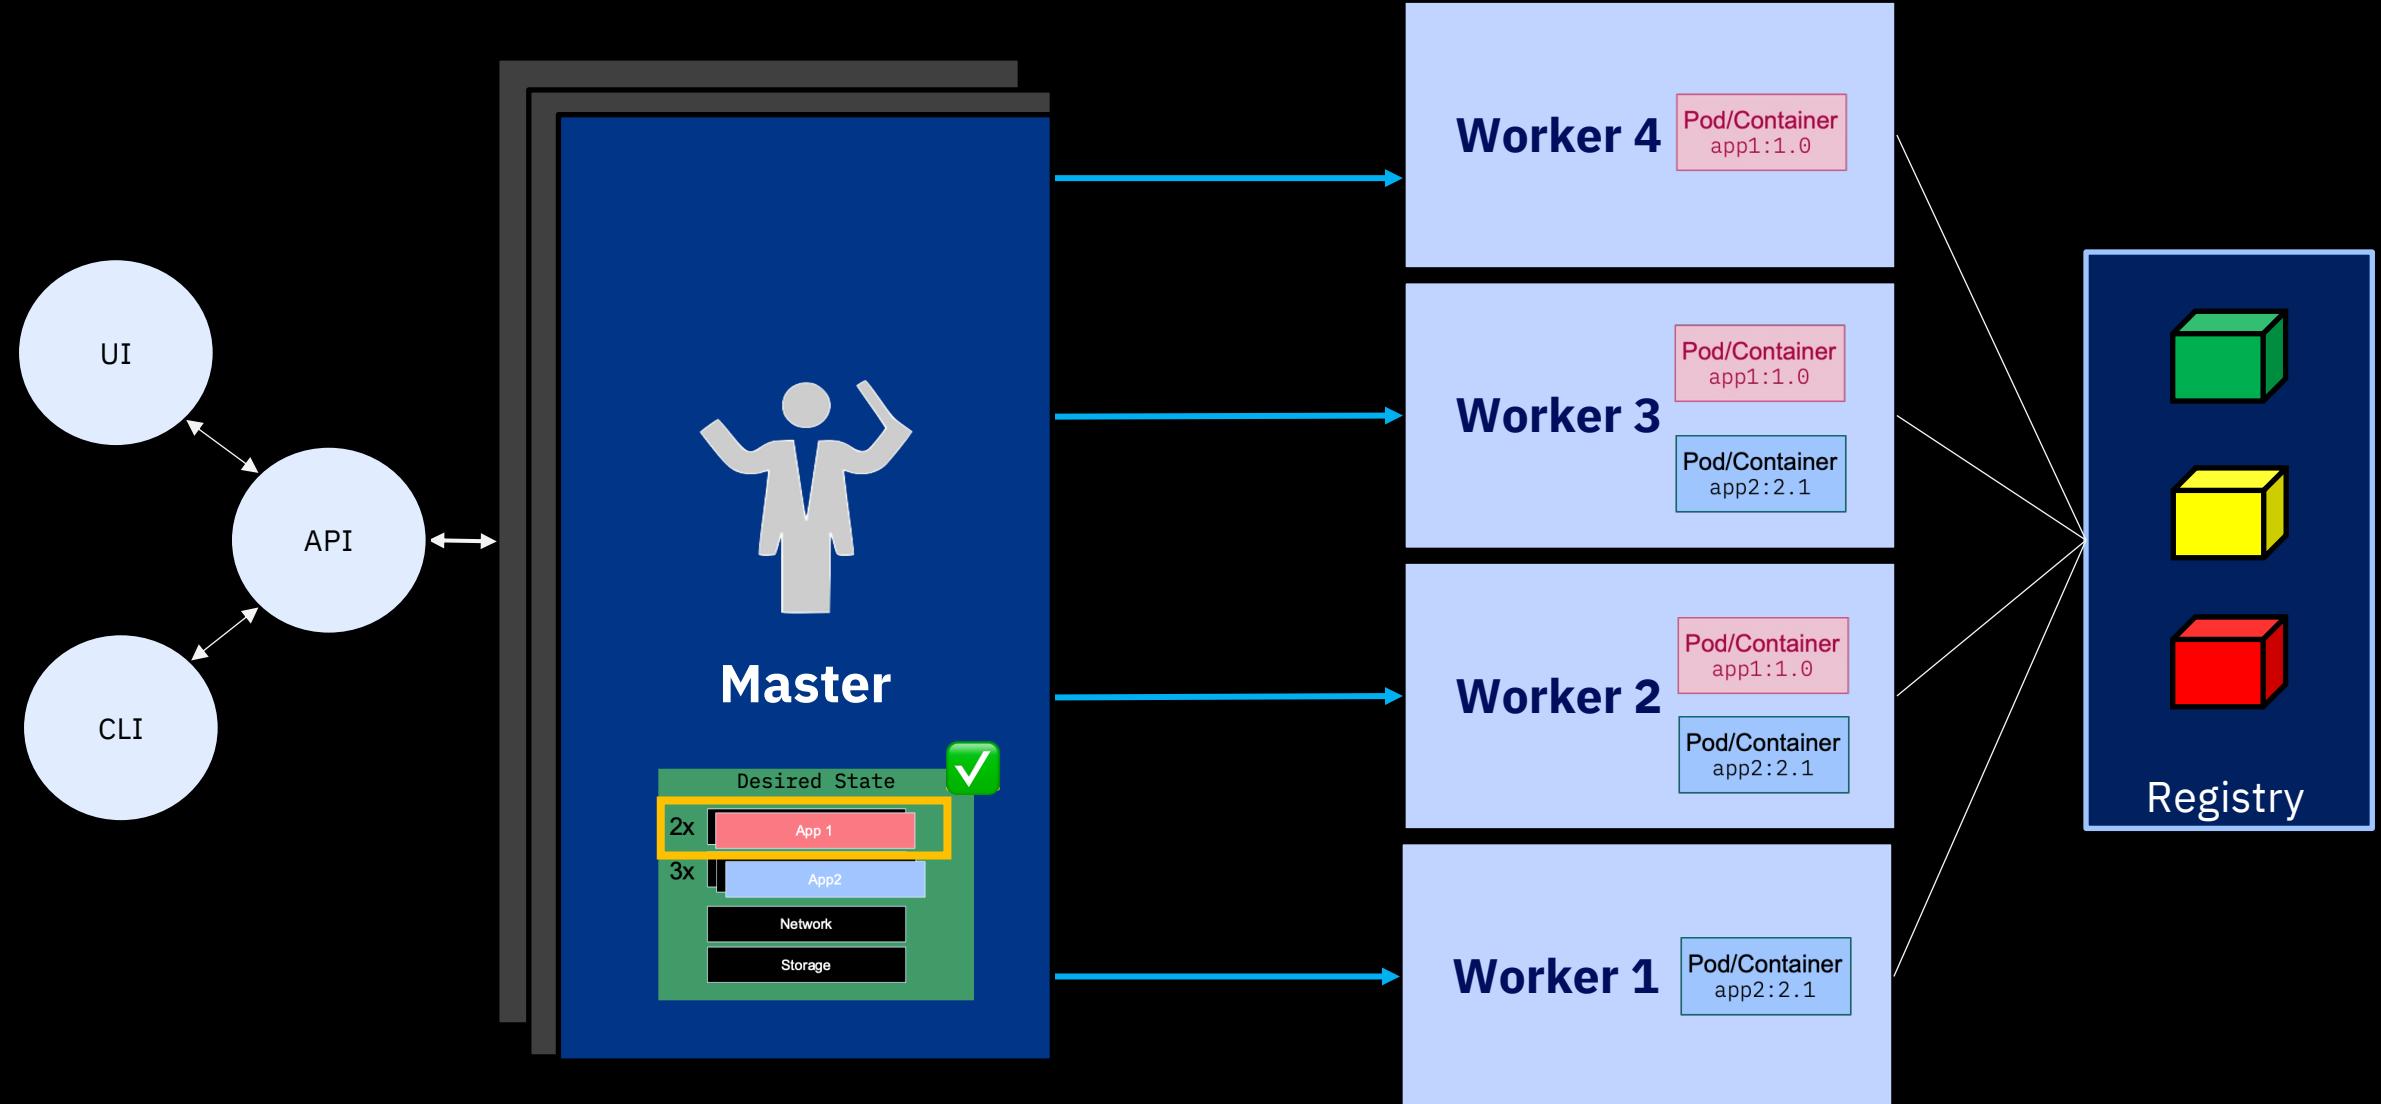
Worker nodes

- Hosts the containers
- Container Runtime

Worker

Pod/Container
app2:2.1

Kubernetes Management Architecture - Example





Deployment

- A set of pods to be deployed together
- Declarative - creates a **ReplicaSet** describing the desired state
- Scale: A Deployment can be scaled

Pods

- Smallest deployment unit in K8s
- Collection of **Containers** that run on a worker node

ReplicaSet

- Ensures availability and scalability
- Maintains the number of **Pods** as requested by user

Deployment

ReplicaSet

Pod

Container A

Pod

Container A

Kubernetes Cluster Architecture

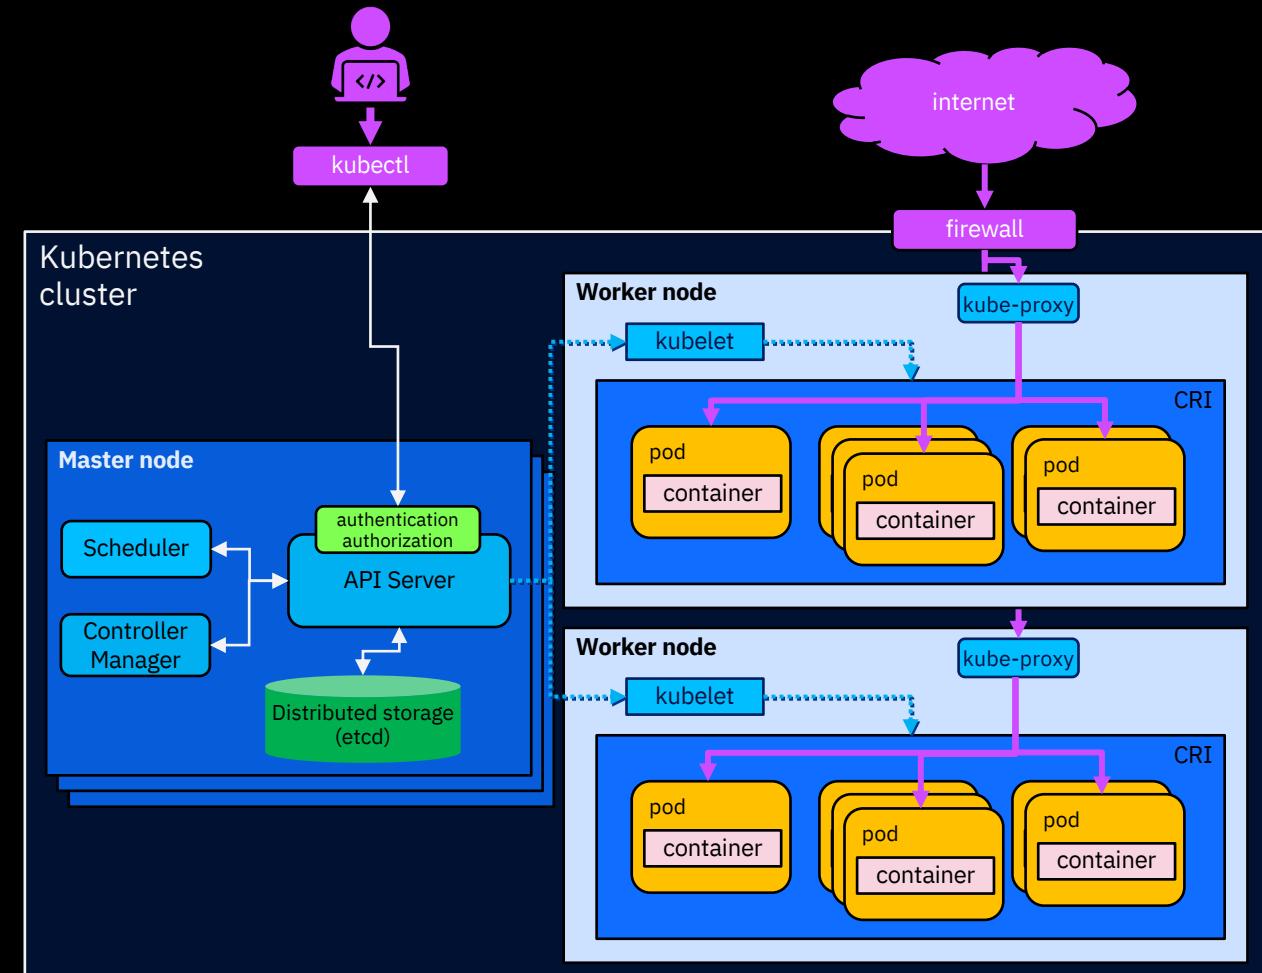


Master node

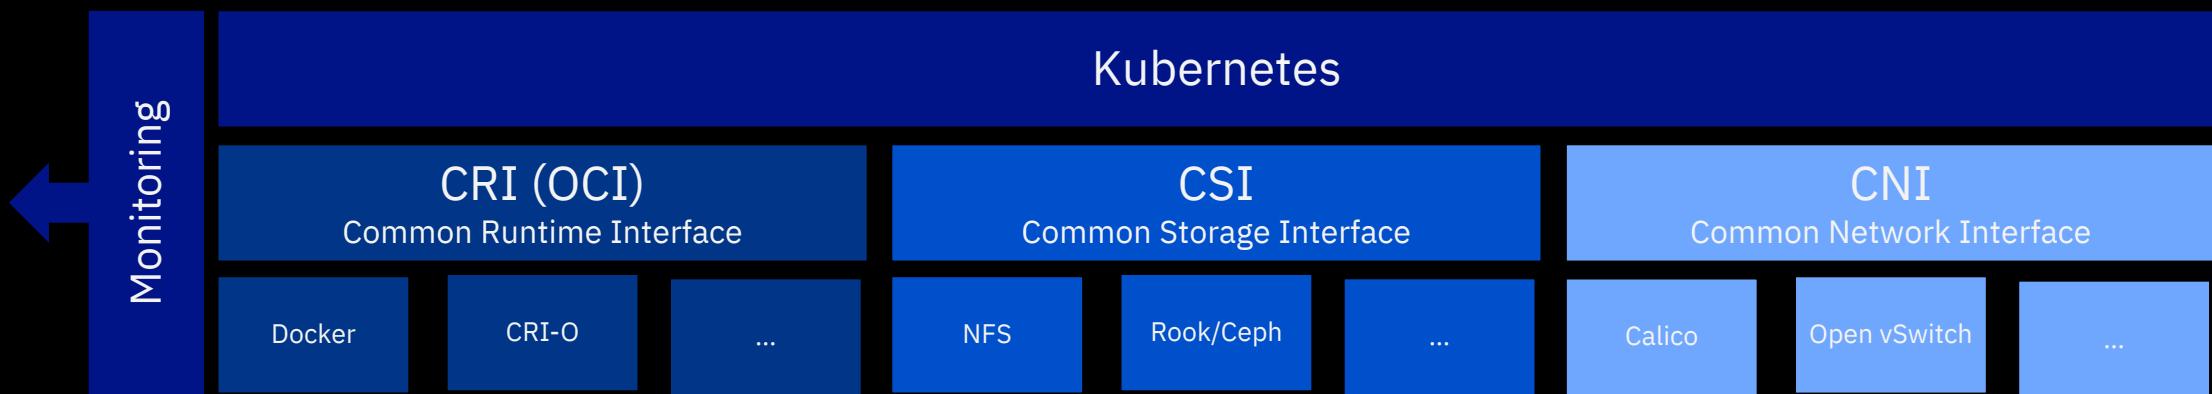
- Node that manages the cluster
- Scheduling, replication & control
- Multiple nodes for HA

Worker nodes

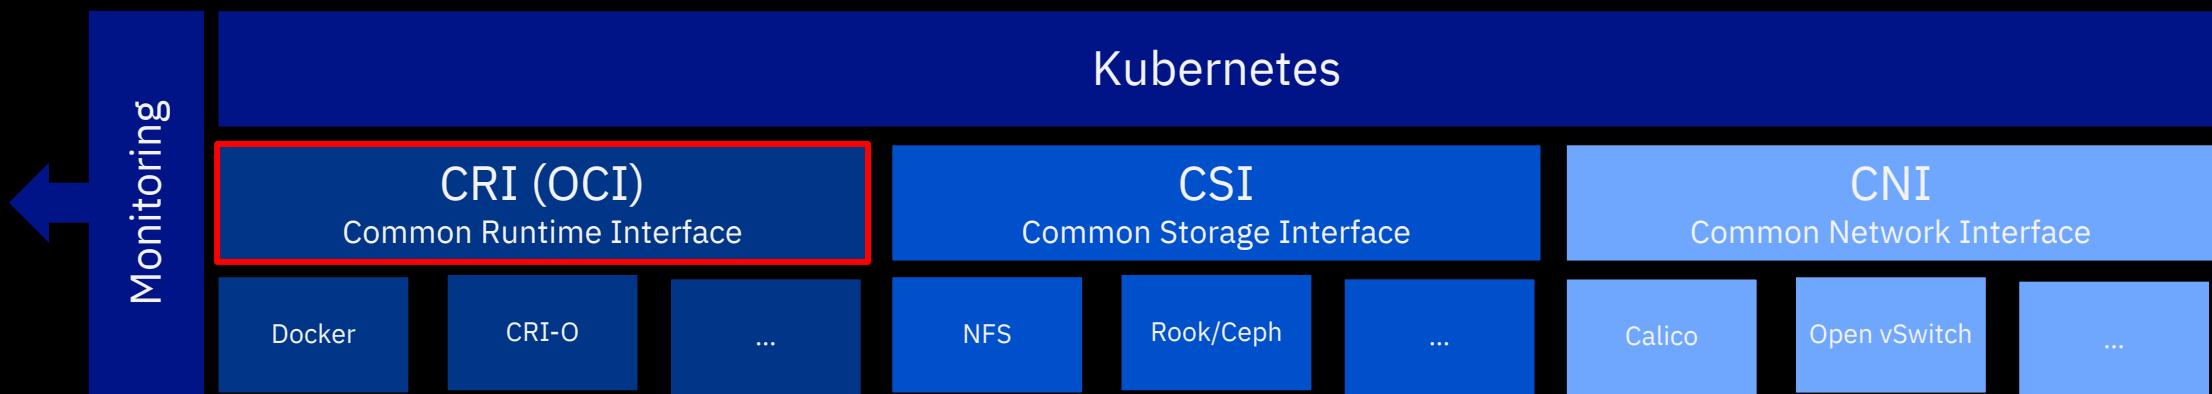
- Node where pods are run
- Docker engine
- kubelet agent accepts & executes commands from the master to manage pods
- kube-proxy – routes inbound or ingress traffic



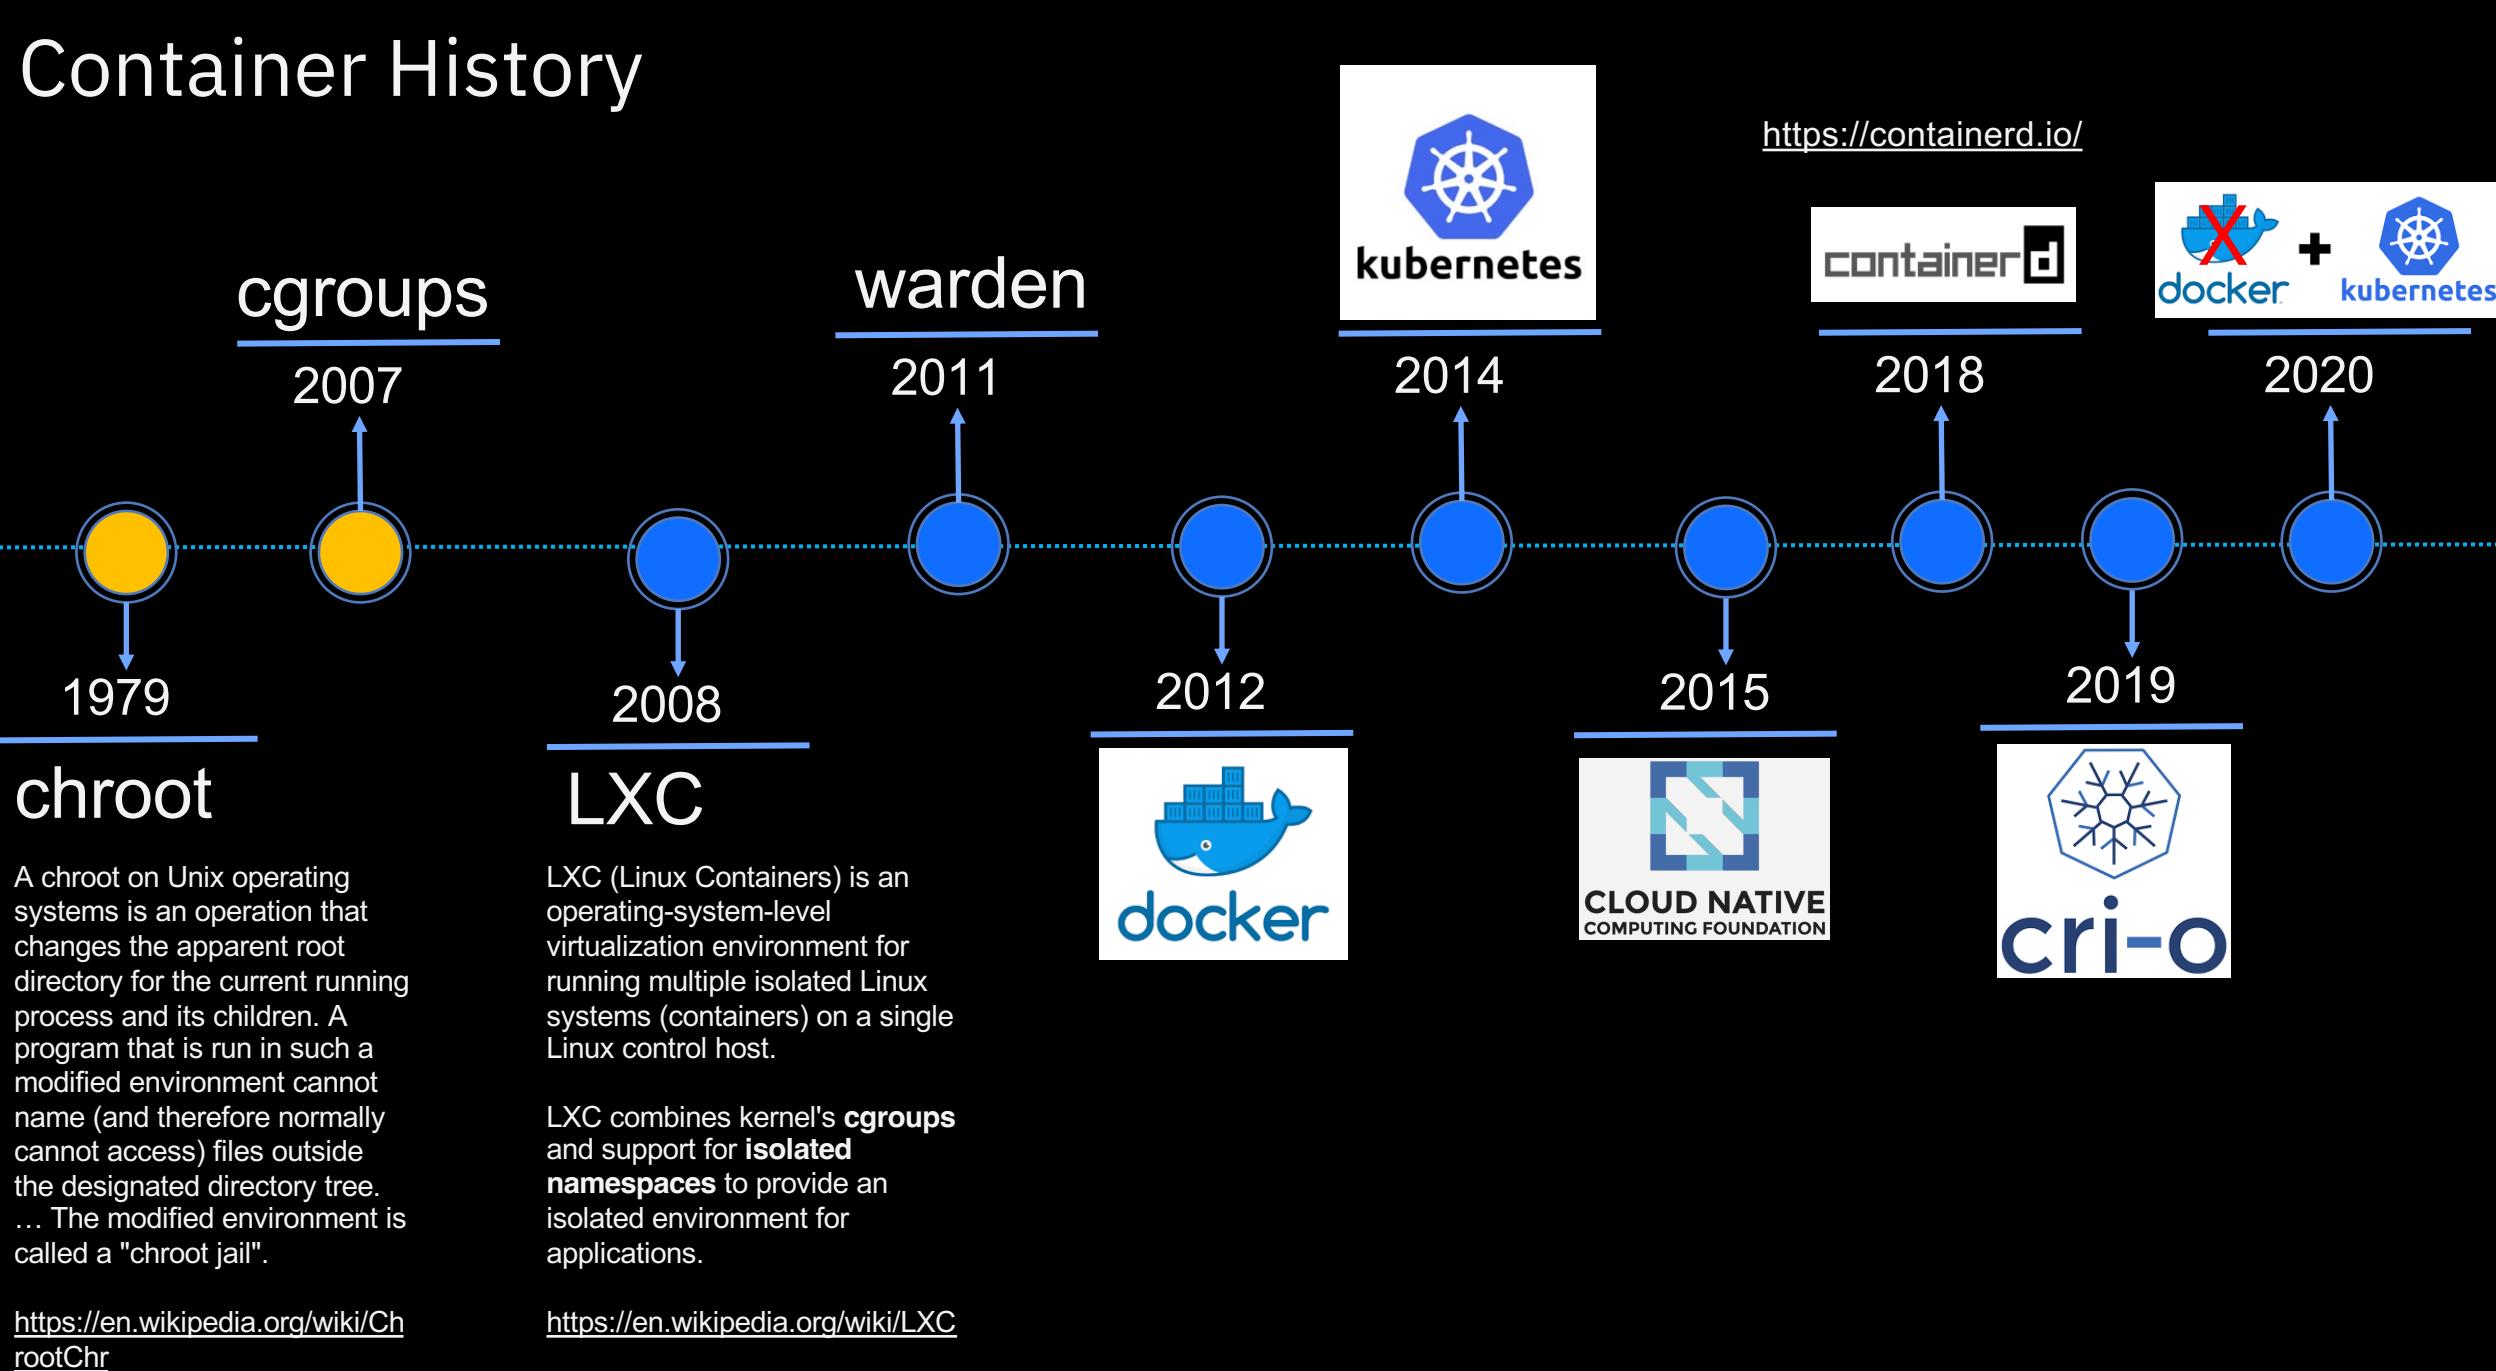
Kubernetes – Common Interfaces



Kubernetes – Common Interfaces



Container History



A chroot on Unix operating systems is an operation that changes the apparent root directory for the current running process and its children. A program that is run in such a modified environment cannot name (and therefore normally cannot access) files outside the designated directory tree. ... The modified environment is called a "chroot jail".

<https://en.wikipedia.org/wiki/ChrootChr>

LXC combines kernel's **cgroups** and support for **isolated namespaces** to provide an isolated environment for applications.

<https://en.wikipedia.org/wiki/LXC>

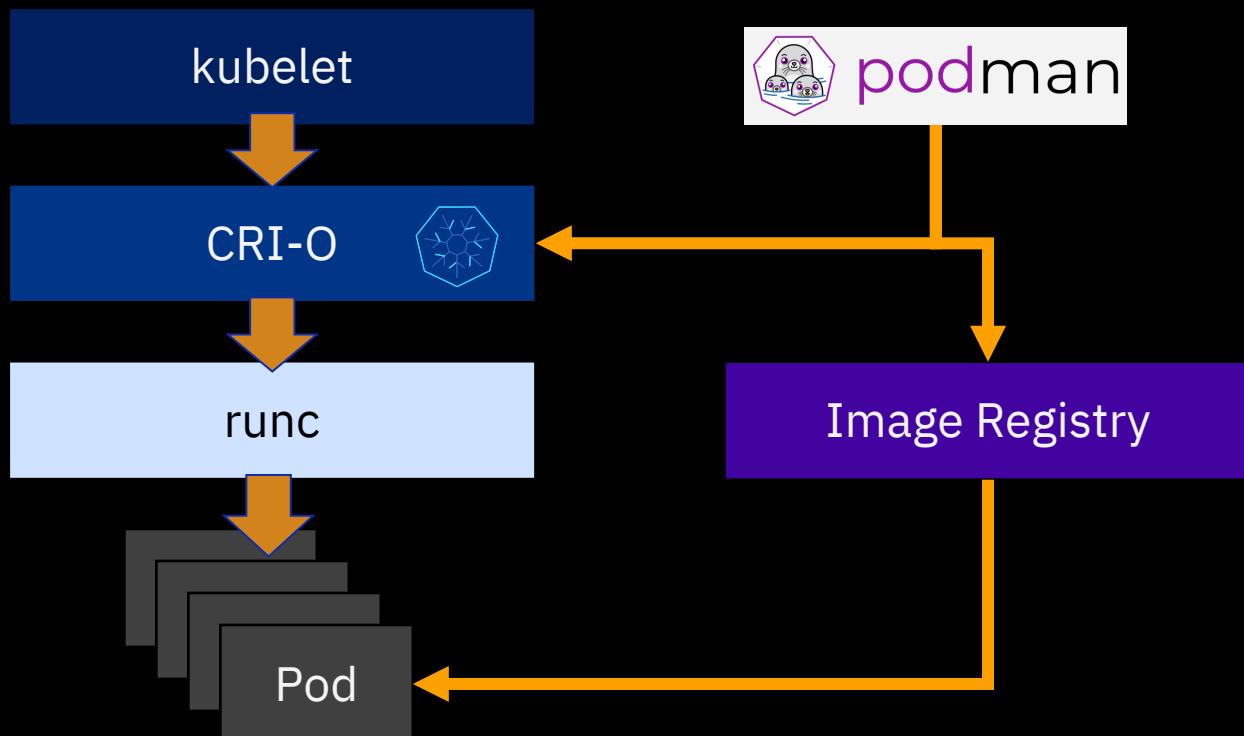


Kubernetes – Why CRI-O?



- **A Truly Open Project:** Operated as part of the **broader Kubernetes community**. Contributors from companies including Red Hat, SUSE, Intel, Google, Alibaba, IBM and more.
- **Lightweight:** CRI-O is made of lots of small components. In comparison, the **Docker Engine** is a **heavyweight daemon** which is communicated to using the docker CLI tool in a client/server fashion.
- **More Securable:** As CRI-O containers are children of the process that spawned it (not the daemon) they're fully compatible with tooling like cgroups & security constraints to **provide an extra layer of protection** to your containers. When using the Docker, **every container** is a 'child' of that large **Docker Daemon**. This complicates or outright prevents the use of this tooling.
- **Aligned with Kubernetes:** As an official Kubernetes project, CRI-O releases **in lock step** with **Kubernetes**, with similar version numbers. ie. CRI-O 1.11.x works with Kubernetes 1.11.x.

Kubernetes – CRI-O



A screenshot of a Twitter thread from Alan Moran (@alanmoran). The first tweet is from May 29, 2018, at 11:49 PM:

I completely forgot that ~2 months ago I set up "alias docker='podman'" and it has been a dream. **#nobigfatdaemons** @projectatomic

This tweet has 7 Retweets and 15 Likes. The second tweet is from May 30, 2018, at 11:49 AM:

Only downside is no Mac OS support (Main dev machine)

The third tweet is a reply from Joe Thompson (@caffeinepresent) on May 30, 2018:

Replies to @alanmoran @projectatomic So, what reminded you?

The fourth tweet is another reply from Alan Moran (@alanmoran) on May 30, 2018:

docker help 😊

Accessing the Cluster

kubectl – talking to the Cluster

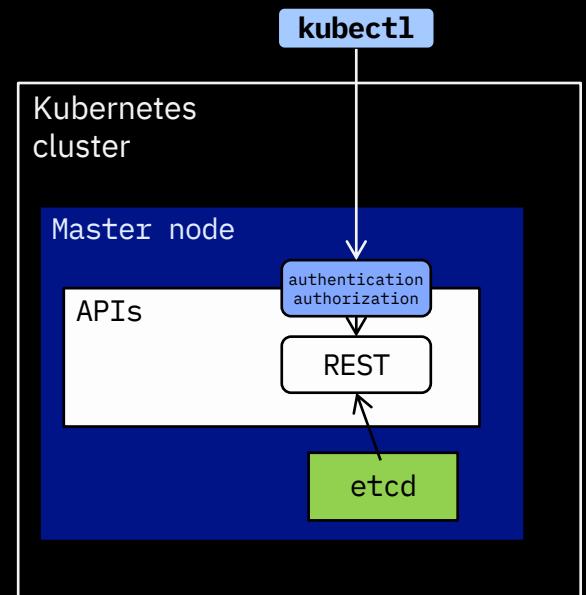


`kubectl` is a command line interface for running commands against Kubernetes clusters (read state, create objects, ...).

`kubectl` looks for a file named `config` in the `$HOME/.kube` directory.

It contains all the information needed to communicate with your cluster.

```
training@ubuntu:~/training/demo-app/k8sdemo$ kubectl config view
apiVersion: v1
clusters:
- cluster:
  certificate-authority: /home/training/.minikube/ca.crt
  server: https://172.17.0.4:8443
  name: minikube
contexts:
- context:
  cluster: minikube
  user: minikube
  name: minikube
current-context: minikube
kind: Config
preferences: {}
users:
- name: minikube
  user:
    client-certificate: /home/training/.minikube/profiles/minikube/client.crt
    client-key: /home/training/.minikube/profiles/minikube/client.key
To start, open Terminal and run ./welcome.sh
```



kubectl – talking to the Cluster

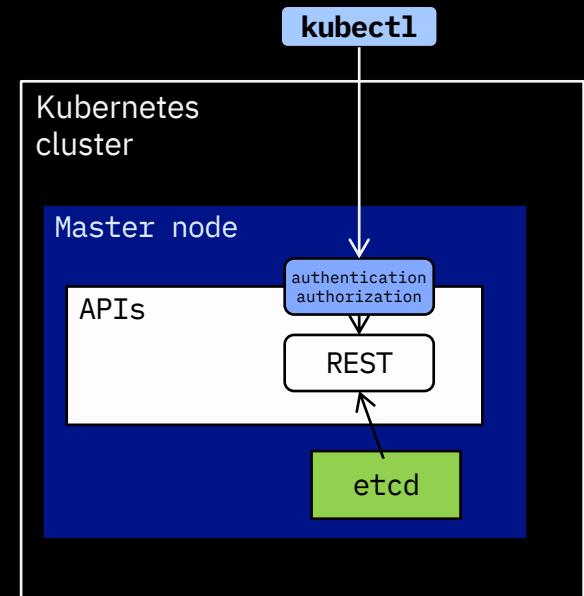


`kubectl` is a command line interface for running commands against Kubernetes clusters.

```
> kubectl [command] [TYPE] [NAME]  
  
> kubectl create -f example.yaml  
> kubectl apply -f example.yaml  
> kubectl delete -f example.yaml  
  
> kubectl get pods  
> kubectl describe nodes <node-name>  
> kubectl logs <pod-name>
```

Create objects in yaml file
Modify objects in yaml file
Delete objects in yaml file

List Pods in Namespace
Details about K8s object
Get the logs for a Pod

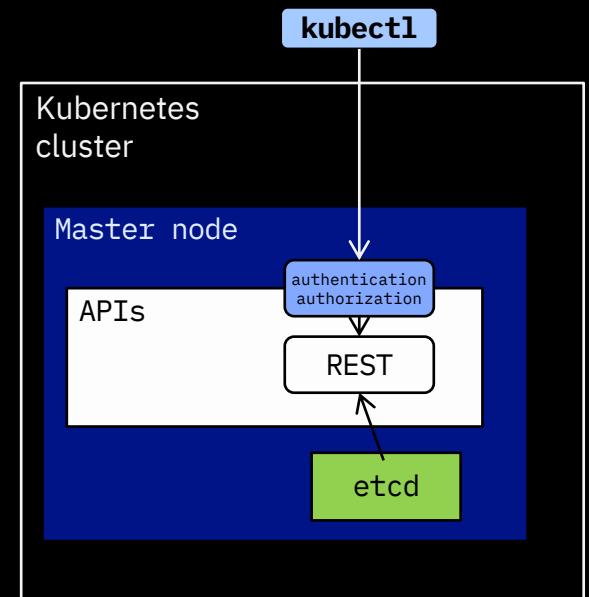


kubectl – talking to the Cluster



`kubectl` is a command line interface for running commands against Kubernetes clusters (read state, create objects, ...).

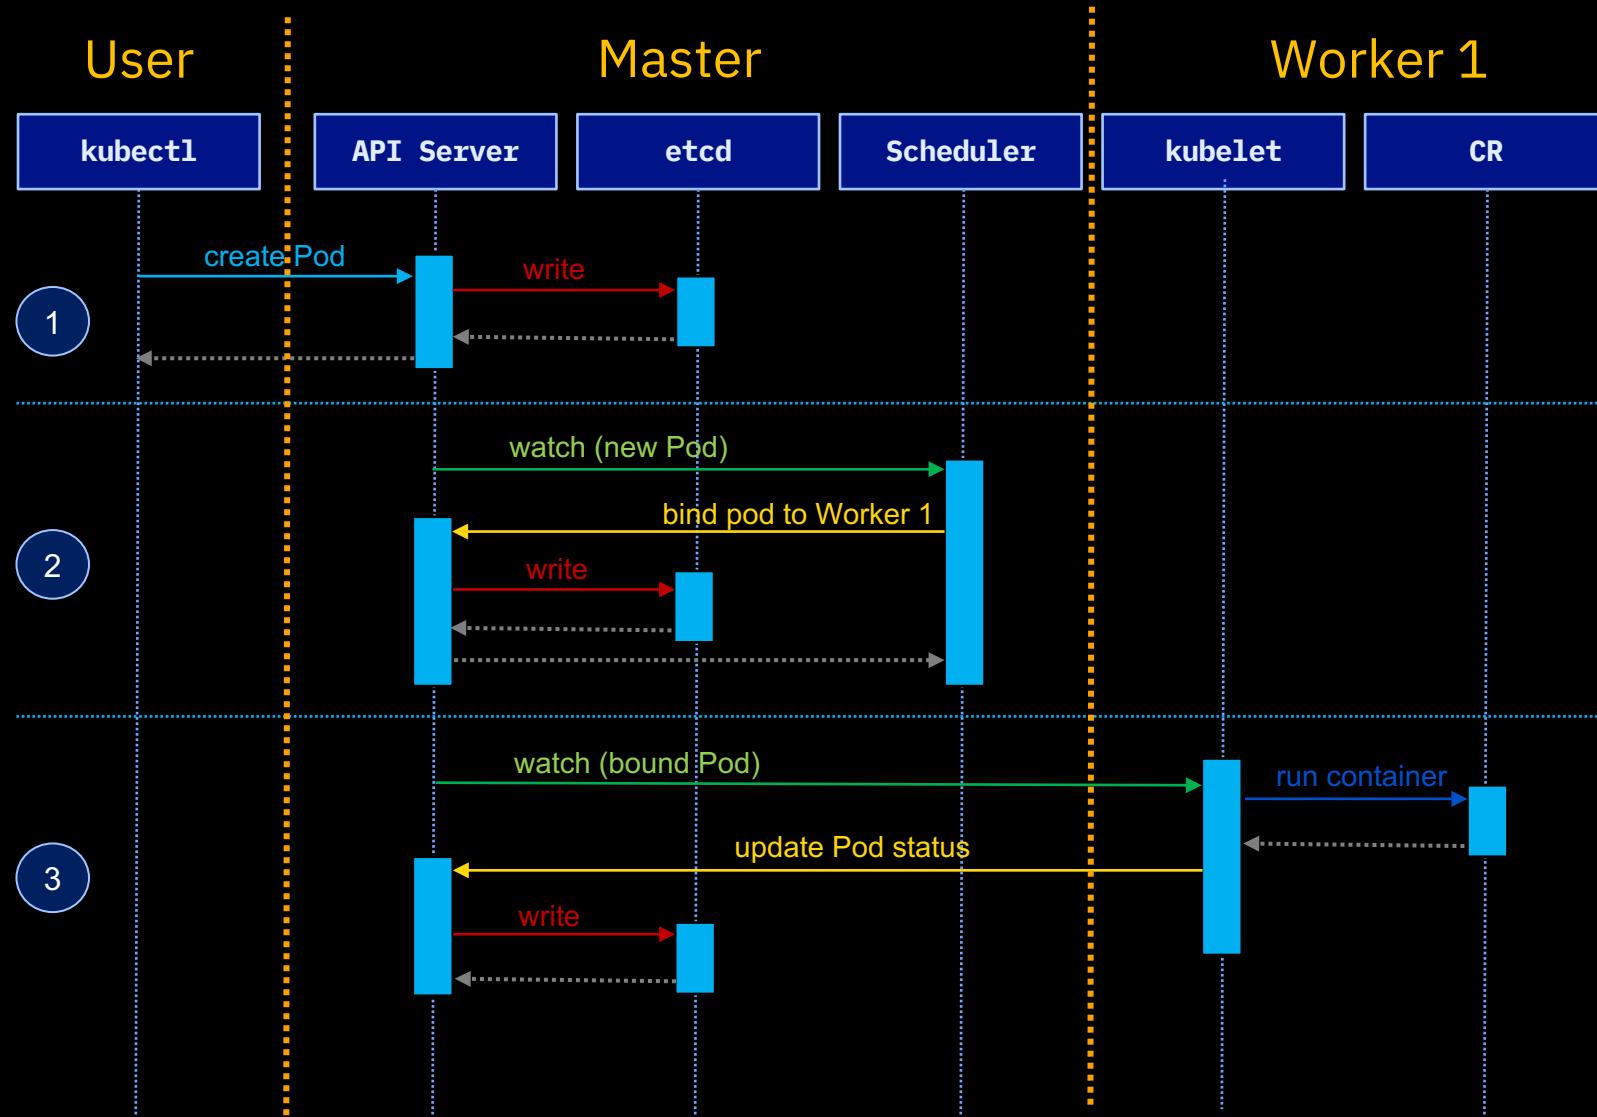
- Kubernetes uses `etcd` as a key-value database store.
- It stores the `configuration` of the Kubernetes cluster in etcd.
- It also stores the **actual state** of the system and the **desired state** of the system in etcd.
- Anything you might read from a `kubectl get xyz` command is coming from etcd.
- Any change you make via `kubectl create` will cause an entry in etcd to be updated.



kubectl – talking to the Cluster



What does this actually do: `kubectl run nginx --image=nginx:1.7.9`



K8s Objects

Kubernetes Objects with YAML Manifests



<pre>apiVersion: v1 kind: Pod</pre>	What kind of object you want to create
<pre>metadata: name: nginx namespace: default labels: app: k8s-demo</pre>	Data that helps uniquely identify the object (Name, Namespace, Labels, ...)
<pre>spec: containers: - name: nginx image: nginx:1.7.9 ports: - containerPort: 80 env: - name: APPLICATION_NAME value: k8sdemo</pre>	Description of the characteristics you want the resource to have: its DESIRED STATE
<pre>status: phase: Running hostIP: 192.168.252.150 podIP: 10.131.1.51</pre>	Describes the CURRENT STATE of the object. This is generated once the object has been deployed.



Pod

- A group of **one or more containers** is called a Pod.
- Containers in a pod are **deployed together**, and are started, stopped, and replicated as a group.
- Containers in pod share the **same network interface**.
- Applications in the same pod
 - Share IP Address and port space
 - Share the same hostname
 - Can communicate using native IPC
 - Can share mounted storage
 - IP Address can change when restarting the Pod





Naming and Structuring

Name

- Each resource object by type has a **unique name**

Namespace

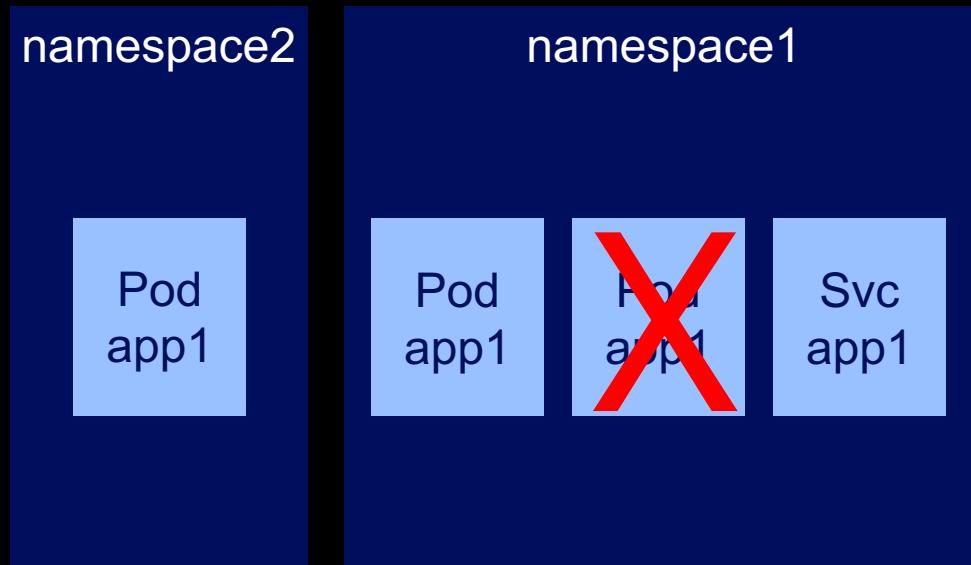
- **Logical grouping:** group objects by application, tier, environment, ...

Example:

- app1-frontend-prod
- app2-datalayer-dev

- **Resource isolation:** Each namespace is a virtual cluster within the physical cluster

- Resource objects are scoped within namespaces
- Names of resources need to be **unique within a namespace**, but not **across namespaces**
- Low-level resources are not in namespaces: nodes, persistent volumes, and namespaces themselves
- **Resource quotas:** Namespaces can divide cluster resources
- **Network isolation:** Namespaces can isolate cluster resources



- Initial namespaces
 - **default** – The default namespace for objects with no other namespace
 - **kube-system** – The namespace for objects created by the Kubernetes system
 - **openshift-xxx** – The namespaces for objects created by OpenShift

Creating a Pod



Imperative

```
> kubectl run nginx --image=nginx:1.7.9
```

Or

```
> kubectl create -f example.yaml
```

Declarative

example.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: default
spec:
  containers:
  - name: nginx
    image: nginx:1.7.9
  ports:
  - containerPort: 80
```

Creating a Pod



```
> kubectl create -f example.yaml
```

```
> kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-5bd87f76c-vxc79	1/1	Running	0	29s

```
> kubectl get pods nginx-5bd87f76c-vxc79 -o yaml
```

```
apiVersion: v1
kind: Pod
metadata:
...
status:
...
```

example.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: default
spec:
  containers:
  - name: nginx
    image: nginx:1.7.9
    ports:
    - containerPort: 80
```



Deleting a Pod



```
> kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-5bd87f76c-vxc79	1/1	Running	0	29s

```
> kubectl delete pods nginx-5bd87f76c-vxc79
```

```
> kubectl delete -f example.yaml
```

example.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: default
spec:
  containers:
  - name: nginx
    image: nginx:1.7.9
    ports:
    - containerPort: 80
```



10.0.0.1

Configuring Resources and Containers



Label

- **Metadata** assigned to Kubernetes resources (pods, services, etc.)
- Key-value pairs for identification
- Critical to Kubernetes as it relies on querying the cluster for resources that have certain labels

Selector

- An expression that **matches labels** to identify related resources

```
> kubectl get pods -l release=april2023
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-5bd87f76c-vxc79	1/1	Running	0	29s

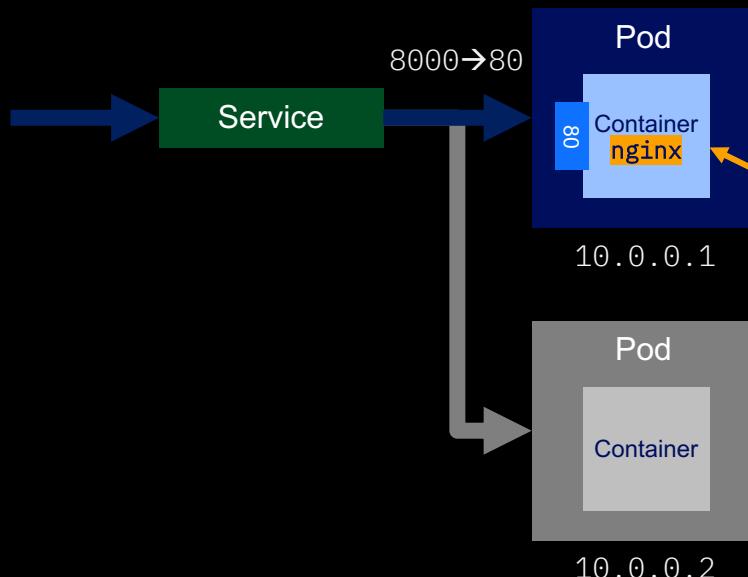
```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: default
  labels:
    app: my-nginx
    release: april2023
spec:
  containers:
  - name: nginx
    image: nginx:1.7.9
    ports:
    - containerPort: 80
```

Basic Networking

Service



- A service provides a way to refer to a set of Pods (selected by labels) with a single static IP address.
- Creates an entry in the Kubernetes DNS



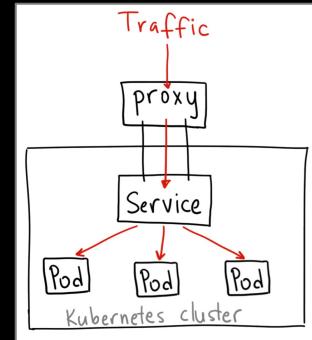
```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    app: my-nginx
    release: april2023
spec:
  containers:
  - name: nginx
    image: nginx:1.7.9
    ports:
    - containerPort: 80
```

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  ports:
  - port: 8000
    targetPort: 80
    protocol: TCP
  selector:
    app: my-nginx
```

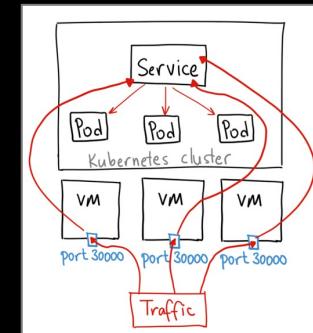


Service

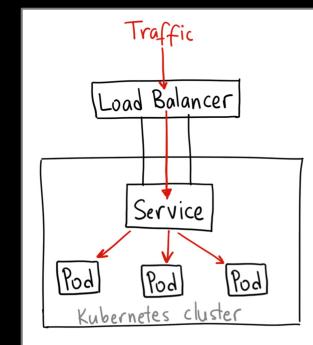
ClusterIP (default): This type exposes the service on the cluster internal IP. This means that the service is only reachable from within the cluster.



NodePort: This type exposes the service on each Node's static IP address.



LoadBalancer: This service type exposes a service using a provided external load balancer.



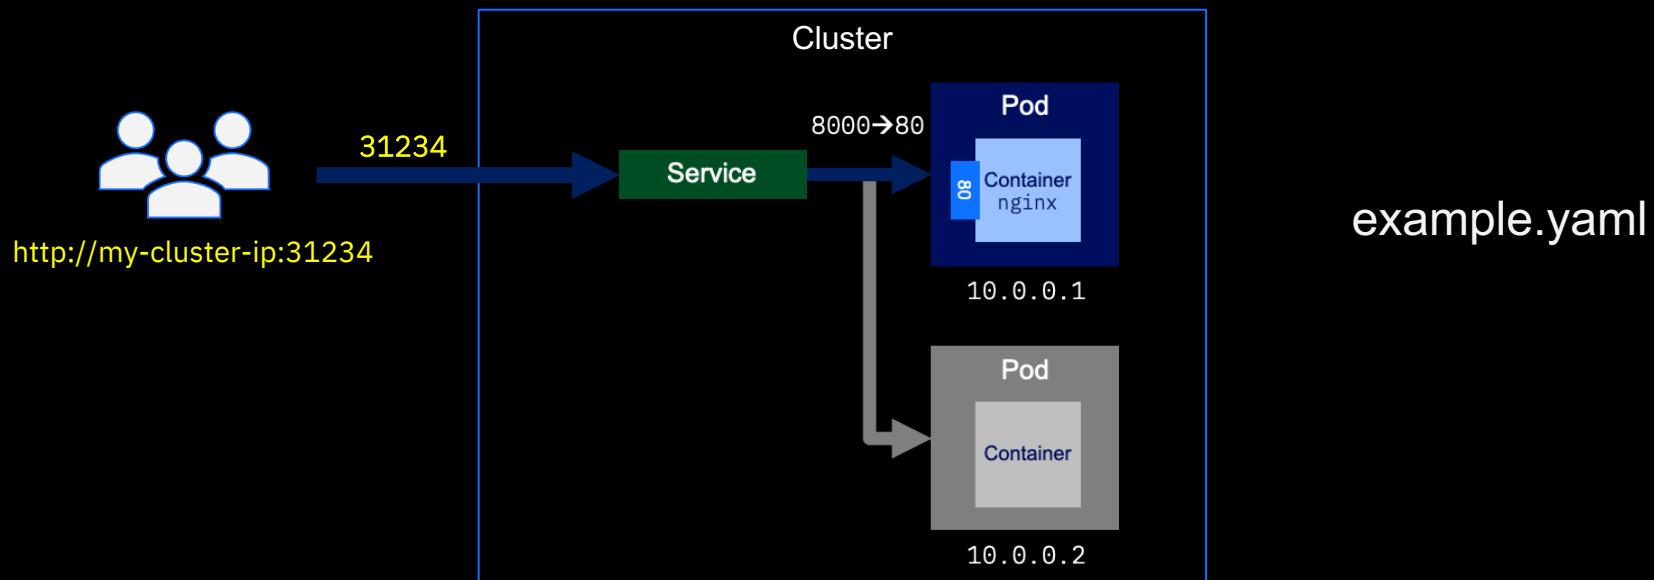


Creating a Service

```
> kubectl expose deployment nginx --type="NodePort" --port=8000 --target-port=80  
or  
> kubectl create -f example.yaml
```

```
> kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
nginx	NodePort	10.106.115.135	<none>	8000:32499/TCP	6s



```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  type: NodePort
  ports:
  - port: 8000
    targetPort: 80
    protocol: TCP
  selector:
    app: nginx
```

Service Access



So, a Kubernetes service with the name **k8s-demo-service** in namespace **default** could be addressed:

When referred from the same namespace

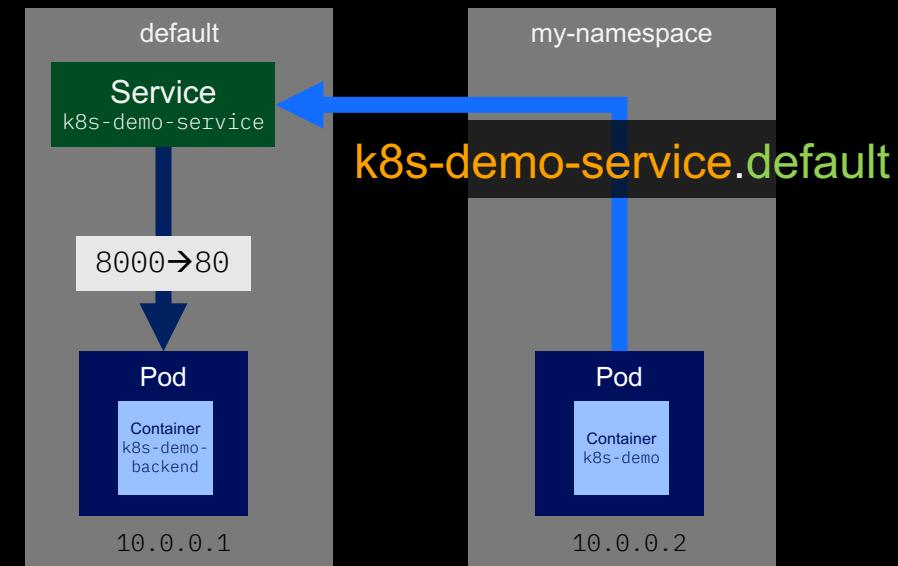
k8s-demo-service

When referred from another namespace

k8s-demo-service.default

Fully qualified name

k8s-demo-service.default.svc.cluster.local

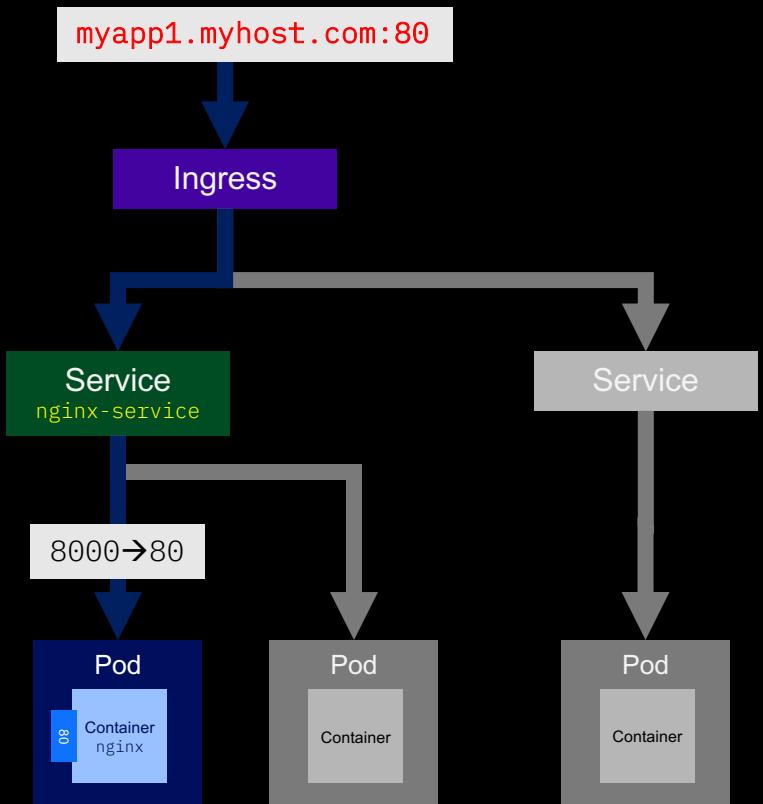


Ingress



An ingress can be configured to give services externally-reachable URLs, load balance traffic, terminate SSL, and offer name based virtual hosting. An ingress controller is responsible for fulfilling the ingress, usually with an [additional Loadbalancer](#).

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: nginx-ingress
  annotations:
    ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - host: myapp1.myhost.com
    http:
      paths:
      - backend:
          serviceName: nginx-service
          servicePort: 80
```

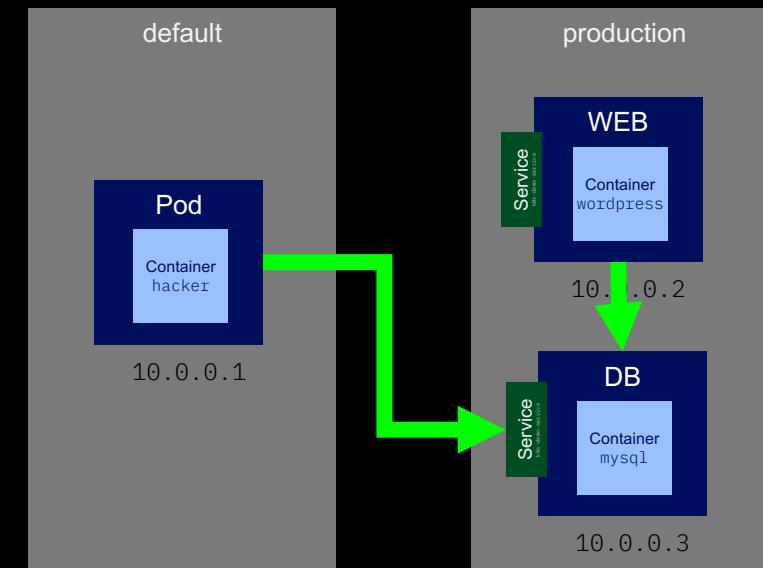


Securing Service Access



NetworkPolicies allow you to specify how a pod is allowed to communicate with "endpoints" and "services" over the network.

⚠ By default, Pods can access **ALL** other Pods.



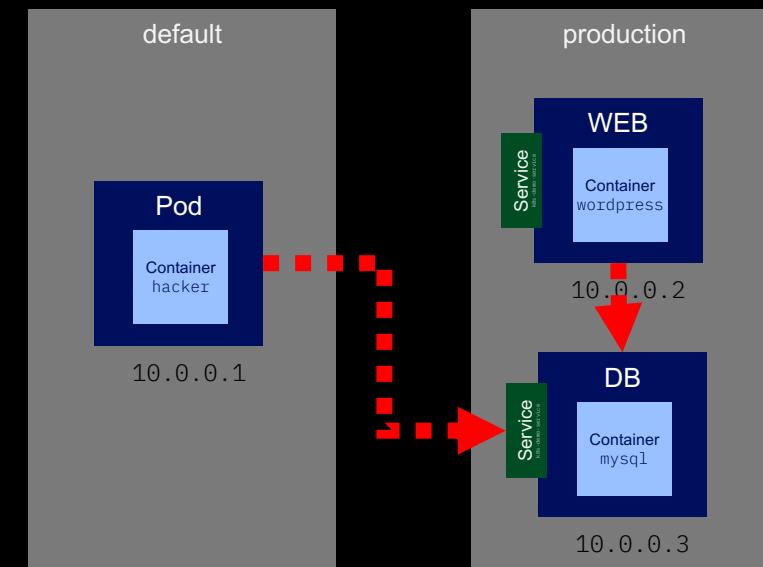


Securing Service Access

NetworkPolicies allow you to specify how a pod is allowed to communicate with "endpoints" and "services" over the network.

⚠ By default, Pods can access **ALL** other Pods.

```
kind: NetworkPolicy
metadata:
  name: namespace-deny
  namespace: production
spec:
  podSelector:
    matchLabels: {}
```



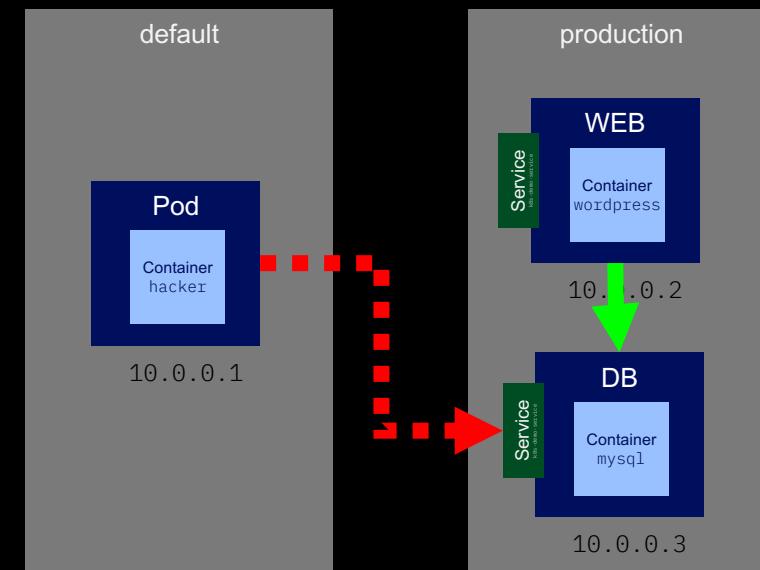
Securing Service Access



NetworkPolicies allow you to specify how a pod is allowed to communicate with "endpoints" and "services" over the network.

⚠ NetworkPolicies are cumulative!!!

```
kind: NetworkPolicy
metadata:
  name: namespace-allow-db
  namespace: production
spec:
  podSelector:
    matchLabels:
      run: DB
  ingress:
    - from:
        - podSelector:
            matchLabels:
              run: WEB
```



Deployments

Deployment



- A Deployment object defines a Pod **creation template** and **desired replica count**.
- Create or delete Pods as needed to meet the replica count.
- Manage safely **rolling out changes** to your running Pods.

```
› kubectl create -f example.yaml
```

example.yaml

```
apiVersion: app/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 2
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.7.9
      ports:
      - containerPort: 80
```



Deployment

Deployment

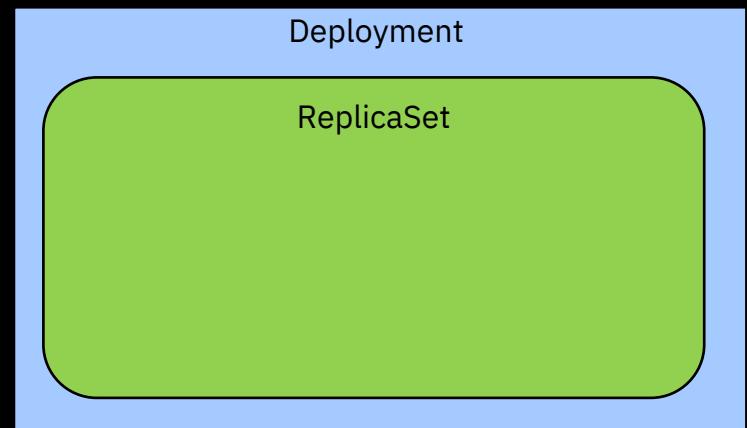
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: k8s-demo-deployment
  labels:
    app: k8s-demo
spec:
...
  replicas: 2
  selector:
    matchLabels:
      app: k8s-demo
  template:
    metadata:
      labels:
        app: k8s-demo
    spec:
      containers:
        - name: k8s-demo
          imagePullPolicy: Always
          image: niklaushirt/k8s-demo:1.0
          ports:
            - containerPort: 80
```

Deployment

Deployment



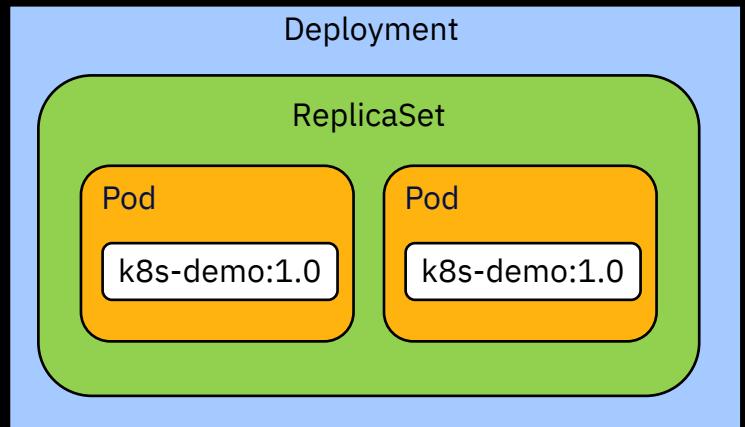
```
Deployment
ReplicaSet
apiVersion: apps/v1
kind: Deployment
metadata:
  name: k8s-demo-deployment
  labels:
    app: k8s-demo
spec:
...
  replicas: 2
  selector:
    matchLabels:
      app: k8s-demo
  template:
    metadata:
      labels:
        app: k8s-demo
    spec:
      containers:
        - name: k8s-demo
          imagePullPolicy: Always
          image: niklaushirt/k8s-demo:1.0
          ports:
            - containerPort: 80
```



Deployment



```
Deployment
  ReplicaSet
    Pod
      Deployment
        apiVersion: apps/v1
        kind: Deployment
        metadata:
          name: k8s-demo-deployment
          labels:
            app: k8s-demo
        spec:
          ...
          replicas: 2
          selector:
            matchLabels:
              app: k8s-demo
          template:
            metadata:
              labels:
                app: k8s-demo
            spec:
              containers:
                - name: k8s-demo
                  imagePullPolicy: Always
                  image: niklaushirt/k8s-demo:1.0
                  ports:
                    - containerPort: 80
```





StatefulSet, DaemonSet, Job...

StatefulSet

- Intended to be used with stateful applications and distributed systems.
 - Pods are created sequentially
 - Ordinal index and stable network identity
 - Can create storage for each replica

DaemonSet

- Ensures that all (or some) nodes run a copy of a pod
 - running a cluster storage daemon
 - running a logs collection daemon
 - running a node monitoring daemon

Job

- Creates one or more pods and ensures that a specified number of them successfully terminate

Cron Job

- Manage time-based jobs, once at a specified in time, or repeatedly at a specified time point

Resource Quota



- Limits resource consumption **per namespace**
- Limit can be number of resource objects by type (pods, services, etc.)
- Limit can be total amount of compute resources (CPU, memory, etc.)
- Overcommit is allowed; contention is handled on a first-come, first-served basis

Configuration

Configuring Resources and Containers

Implicit Environment Variables



allowed: true
enemies: aliens
lives: 3

```
kind: Deployment
spec:
  containers:
    - name: test-container
      image: xxx
      ...
  env:
    - name: allowed
      value: true
    - name: enemies
      value: aliens
    - name: lives
      value: 3
```

Can be used as normal environment variable (here **enemies=aliens**)

Configuring Resources and Containers

Environment variables from external sources



ConfigMap

- Configuration values to be used by containers in a pod
- Stores configuration outside of the container image, making containers more reusable

Secret (*not so secret*)

- Sensitive info that containers need to read or consume
- Not encrypted! Only base64 encoded!!!
- Must be protected from read access by unauthorized users

Configuring Resources and Containers

Environment variables from ConfigMaps or Secrets



```
apiVersion: v1
kind: ConfigMap
metadata:
  name: example-config
data:
  allowed: true
  enemies: aliens
  lives: 3
immutable: true
```

```
apiVersion: v1
kind: Secret
metadata:
  name: my-tls-secret
data:
  tls.crt: >-
    LS0tLS1CRUdJTiBDSUVSakND...
  tls.key: >-
    LS0tLS1CRUdJTi0tLQpNSUlF...
type: Opaque
immutable: true
```

Configuring Resources and Containers

Environment variables from ConfigMaps or Secrets



```
apiVersion: v1
kind: ConfigMap
metadata:
  name: example-config
data:
  allowed: true
  enemies: aliens
  lives: 3
```

```
kind: Deployment
spec:
  containers:
    - name: test-container
      image: xxx
      ...
      ...
```

Can be used as normal environment variable (here **SPECIAL_LEVEL_KEY=aliens**)

Configuring Resources and Containers

Mapping files from ConfigMaps or Secrets



```
apiVersion: v1
kind: Secret
metadata:
  name: my-tls-secret
data:
  tls.crt: >-
    LS0tLS1CR0dJTiBDSUVSakND...
  tls.key: >-
    LS0tLS1CR0dJTi0tLQpNSU1F...
type: Opaque
```

```
kind: Deployment
...
spec:
  containers:
    - name: test-container
      image: xxx
      ...
      volumeMounts:
        - name: my-tls-secret-volume
          mountPath: /workdir/home/certs/public.crt
          subPath: tls.crt
      ...
      volumes:
        - name: my-tls-secret-volume
          secret:
            secretName: my-tls-secret
            items:
              - key: tls.crt
                path: tls.crt
```

The Certificate will be mapped to the `/workdir/home/certs/public.crt` file

Configuring Resources and Containers

Mapping files from ConfigMaps or Secrets



```
apiVersion: v1
kind: ConfigMap
metadata:
  name: example-config
data:
  ...
  
```

```
apiVersion: v1
kind: Secret
metadata:
  name: my-tls-secret
data:
  ...
  
```

```
kind: Deployment
...
spec:
  containers:
    - name: test-container
      image: xxx
      ...
  envFrom:
    - configMapRef:
        name: example-config
        optional: true
    - secretRef:
        name: my-tls-secret
        optional: true
  
```



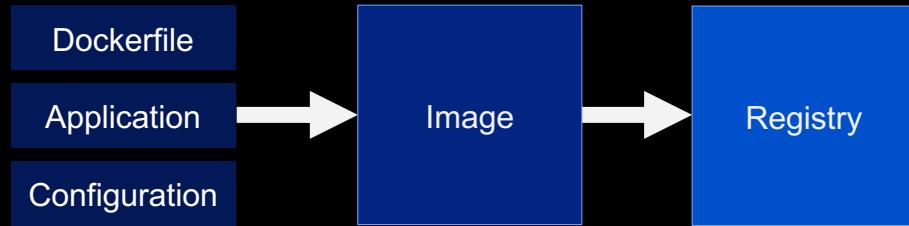
◦◦ Kubernetes Workshop Series
Kubernetes - Applied

05



..but how do I
use this for
real?

Minimal Application – Container Image



```
FROM node:8-stretch

# Change working directory
WORKDIR "/app"

# Update packages and install dependency packages for services
RUN apt-get update \
&& apt-get dist-upgrade -y \
&& apt-get clean \
&& echo 'Finished installing dependencies'

# Install npm production packages
COPY package.json /app/
RUN cd /app; npm install --production

COPY ./myapp /app

ENV NODE_ENV production
ENV BACKEND_URL https://api.nasa.gov/planetary/apod?api_key=_KEY
ENV PORT 3000

EXPOSE 3000

CMD ["npm", "start"]
```

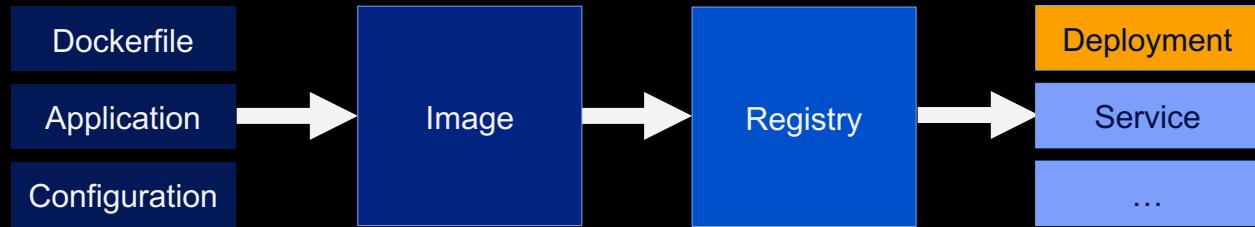
Minimum components you need:

- Built Application or source
(NodeJS, Java, Golang, ...)
- Dockerfile
 - Use the right base image for your app

```
> podman build -t niklaushirt/k8s-demo:1.0 .
> podman push niklaushirt/k8s-demo:1.0
```



Minimal Application – Kubernetes Manifests



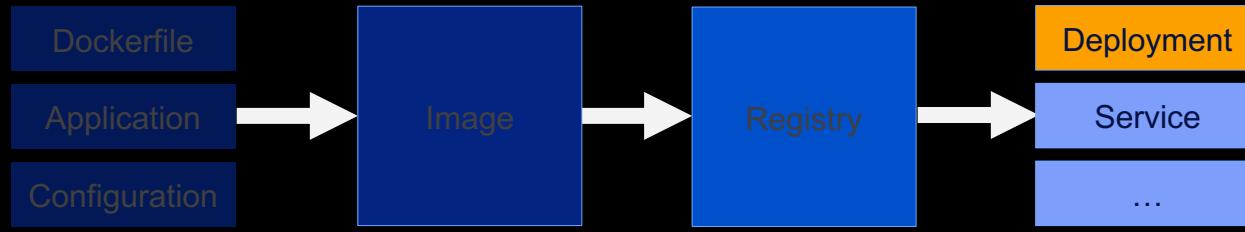
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: k8s-demo-deployment
  labels:
    app: k8s-demo
spec:
  replicas: 2
  selector:
    matchLabels:
      app: k8s-demo
  template:
    metadata:
      labels:
        app: k8s-demo
    spec:
      containers:
        - name: k8s-demo
          imagePullPolicy: Always
          image: niklaushirt/k8s-demo:1.0
          ports:
            - containerPort: 80
```

Minimum components you need:

- Deployment
- Service



Minimal Application – Kubernetes Manifests



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: k8s-demo-deployment
  labels:
    app: k8s-demo
spec:
  replicas: 2
  selector:
    matchLabels:
      app: k8s-demo
  template:
    metadata:
      labels:
        app: k8s-demo
    spec:
      containers:
        - name: k8s-demo
          imagePullPolicy: Always
          image: niklaushirt/k8s-demo:1.0
          ports:
            - containerPort: 80
```

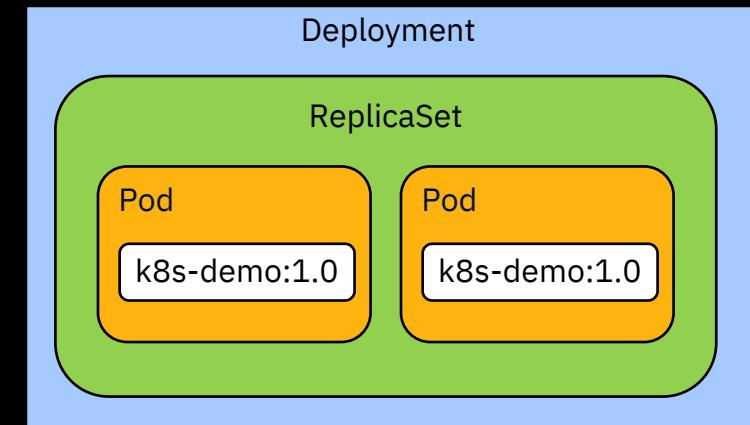
Deployment

ReplicaSet

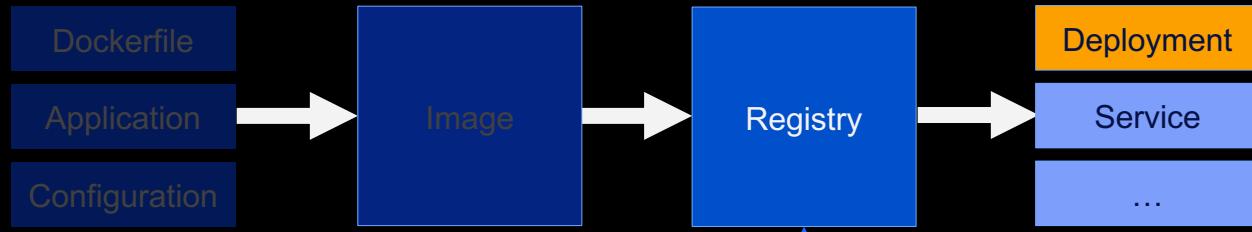
Pod

Minimum components you need:

- Deployment
- Service



Minimal Application – Kubernetes Manifests



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: k8s-demo-deployment
  labels:
    app: k8s-demo
spec:
  replicas: 2
  selector:
    matchLabels:
      app: k8s-demo
  template:
    metadata:
      labels:
        app: k8s-demo
    spec:
      containers:
        - name: k8s-demo
          imagePullPolicy: Always
          image: niklaushirt/k8s-demo:1.0
        ports:
          - containerPort: 80
```

Minimum components you need:

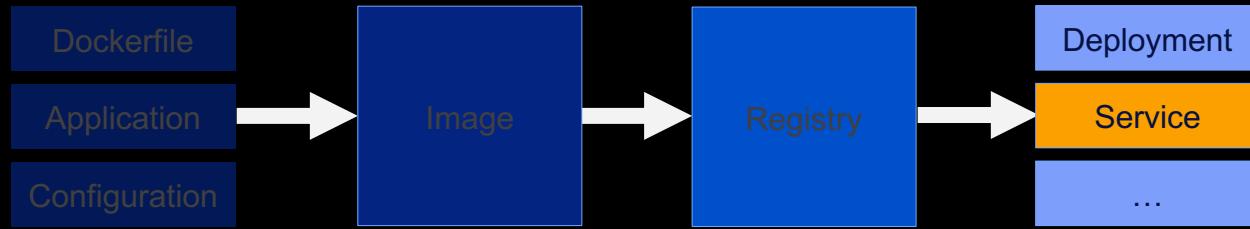
- Deployment
- Service

imagePullPolicy:

IfNotPresent: only pulled if not already present locally
Always: every time the kubelet launches a container



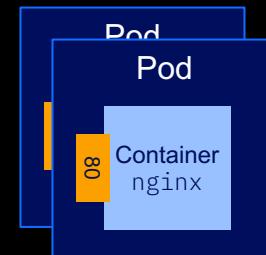
Minimal Application – Kubernetes Manifests



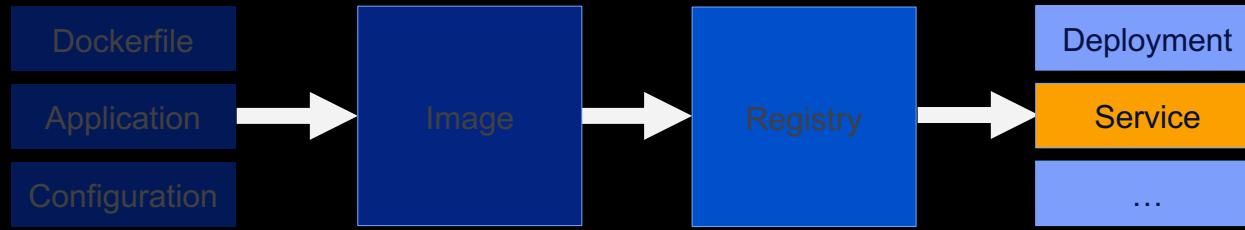
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: k8s-demo-deployment
  labels:
    app: k8s-demo
spec:
  replicas: 2
  selector:
    matchLabels:
      app: k8s-demo
  template:
    metadata:
      labels:
        app: k8s-demo
    spec:
      containers:
        - name: k8s-demo
          imagePullPolicy: Always
          image: niklaushirt/k8s-demo:1.0
          ports:
            - containerPort: 80
```

Minimum components you need:

- Deployment
- Service



Minimal Application – Kubernetes Manifests

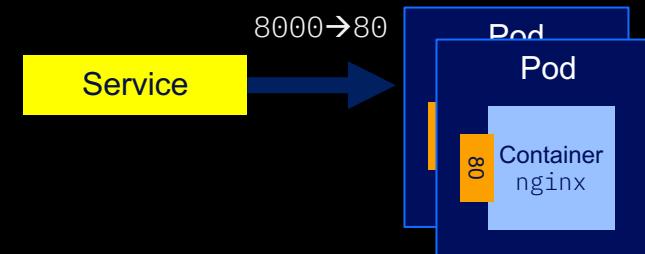


```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: k8s-demo-deployment
  labels:
    app: k8s-demo
spec:
  replicas: 2
  selector:
    matchLabels:
      app: k8s-demo
  template:
    metadata:
      labels:
        app: k8s-demo
    spec:
      containers:
        - name: k8s-demo
          imagePullPolicy: Always
          image: niklaushirt/k8s-demo:1.0
          ports:
            - containerPort: 80
```

```
apiVersion: v1
kind: Service
metadata:
  name: k8s-demo-service
spec:
  ports:
    - port: 8000
      targetPort: 80
      protocol: TCP
    selector:
      app: k8s-demo
```

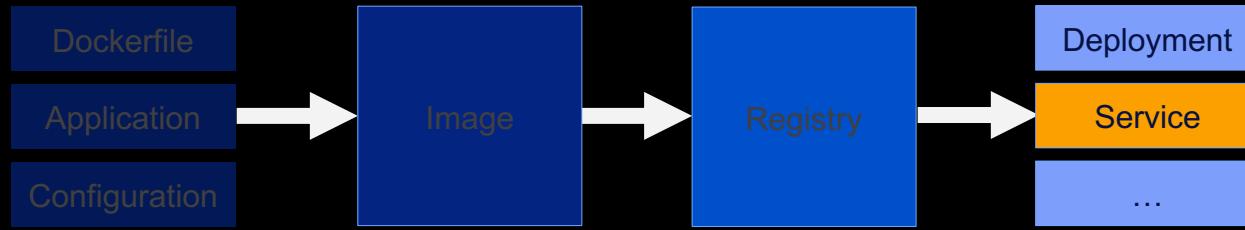
Minimum components you need:

- Deployment
- Service





Minimal Application – Kubernetes Manifests

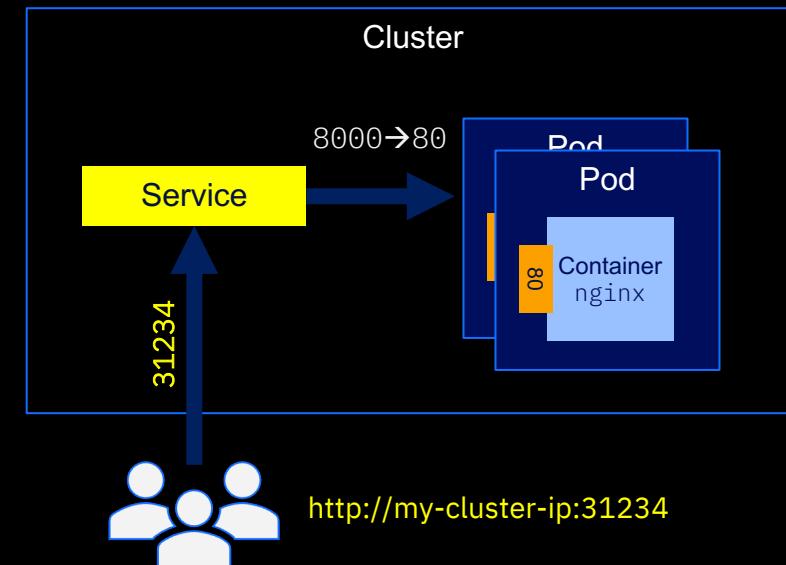


```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: k8s-demo-deployment
  labels:
    app: k8s-demo
spec:
  replicas: 2
  selector:
    matchLabels:
      app: k8s-demo
  template:
    metadata:
      labels:
        app: k8s-demo
    spec:
      containers:
        - name: k8s-demo
          imagePullPolicy: Always
          image: niklaushirt/k8s-demo:1.0
          ports:
            - containerPort: 80
```

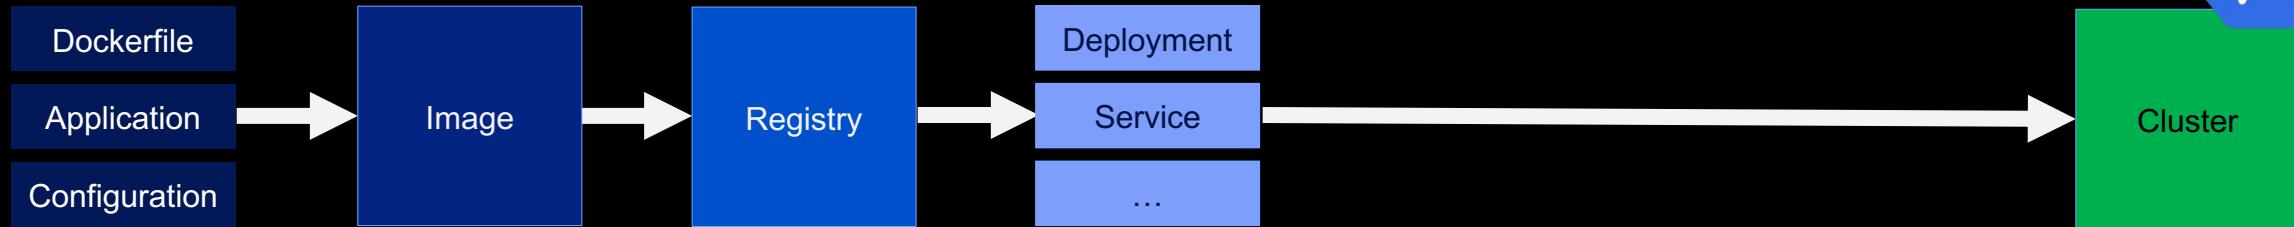
```
apiVersion: v1
kind: Service
metadata:
  name: k8s-demo-service
spec:
  type: NodePort
  ports:
    - port: 8000
      targetPort: 80
      nodePort: 31234
      protocol: TCP
    selector:
      app: k8s-demo
```

Minimum components you need:

- Deployment
- Service



Minimal Application – Deploying



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: k8s-demo-deployment
  labels:
    app: k8s-demo
spec:
  replicas: 3
  selector:
    matchLabels:
      app: k8s-demo
  template:
    metadata:
      labels:
        app: k8s-demo
    spec:
      containers:
        - name: k8s-demo
          image: niklaushirt/k8s-demo:1.0
          ports:
            - containerPort: 80
```

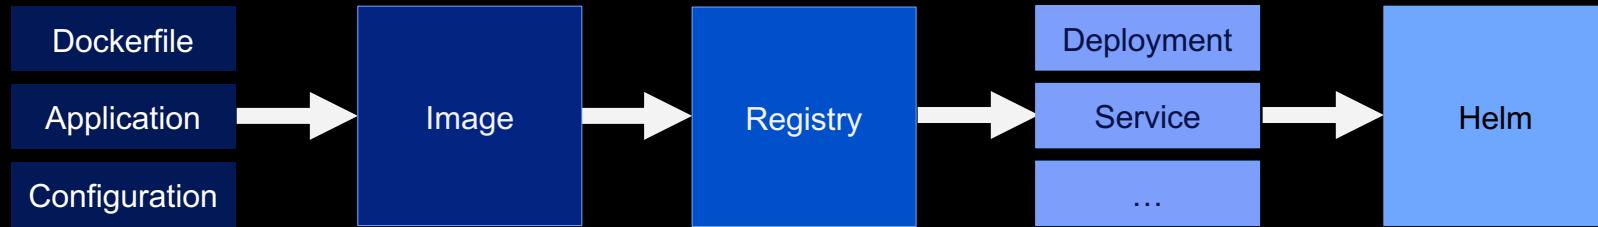
```
apiVersion: v1
kind: Service
metadata:
  name: k8s-demo-service
spec:
  ports:
    - port: 8000
      targetPort: 80
      protocol: TCP
  selector:
    app: k8s-demo
```

Minimum components you need:

- Deployment
- Service

```
> kubectl apply -f deployment.yaml
> kubectl apply -f service.yaml
```

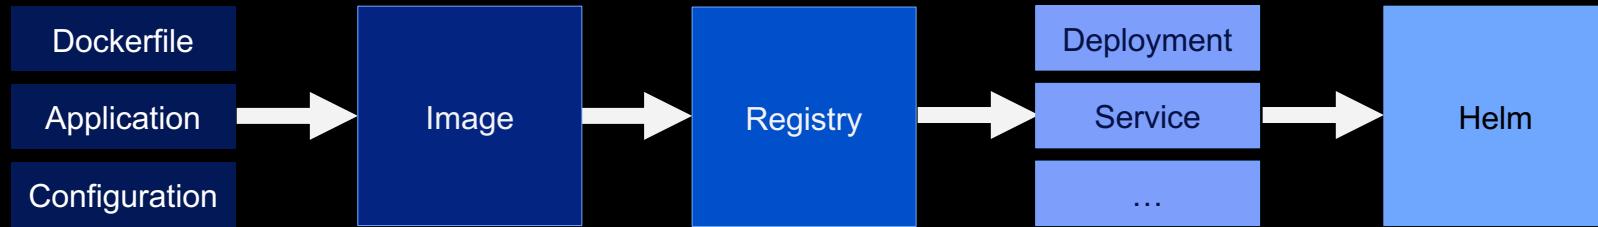
Minimal Application – Helm



The package manager for Kubernetes

Helm helps you manage Kubernetes applications — Helm Charts help you **define, install, and upgrade** even the most complex Kubernetes application.

Minimal Application – Helm



```
apiVersion: v1
name: k8s-demo
description: Demo App for the training.
keywords:
- Training
maintainers:
- name: Niklaus Hirt

version: 1.10.0
appVersion: 19.0.0.6

home: https://github.com/niklaushirt
icon: https://github.com/niklaushirt/logo-01.png
sources:
- https://github.com/niklaushirt/charts/tree/master/stable
```

Minimum components we need:

Chart.yaml

Basic information about the chart

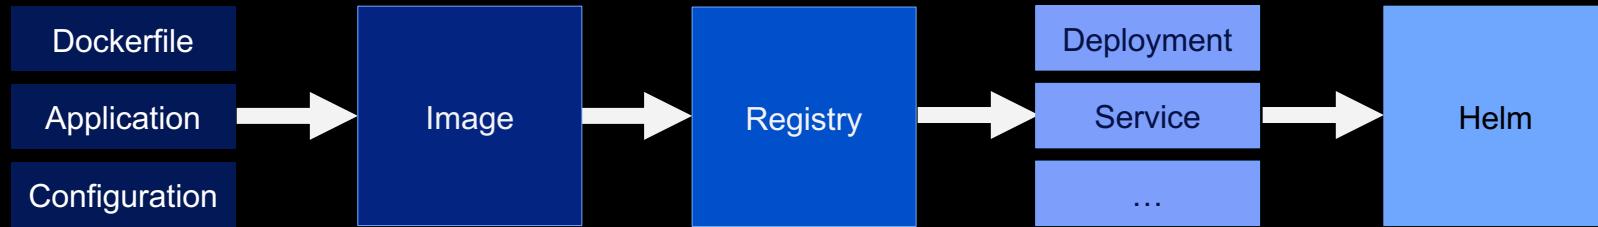
Templates (yaml)

Kubernetes deployment manifests (yaml) with placeholders for external parameters

values.yaml

A set of variables that will be passed on at deploy time to the deployment manifest templates

Minimal Application – Helm



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: k8s-demo-deployment
  labels:
    app: k8s-demo
spec:
  replicas: {{ .Values.replicas }}
  selector:
    matchLabels:
      app: k8s-demo
  template:
    metadata:
      labels:
        app: k8s-demo
    spec:
      containers:
        - name: k8s-demo
          image: {{ .Values.image.name }}:{{ .Values.image.tag }}
          ports:
            - containerPort: 80
```

Minimum components we need:

Chart.yaml

Basic information about the chart

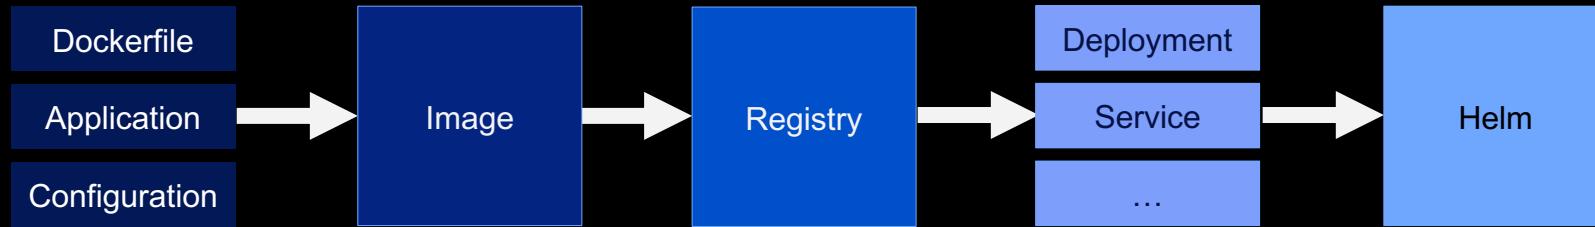
Templates (yaml)

Kubernetes deployment manifests (yaml) with placeholders for external parameters

values.yaml

A set of variables that will be passed on at deploy time to the deployment manifest templates

Minimal Application – Helm



```
apiVersion: v1
image:
  name:
    - niklaushirt/k8s-demo
  tag: 1.0
  pullPolicy: IfNotPresent
replicas: 1
```

Minimum components we need:

Chart.yaml

Basic information about the chart

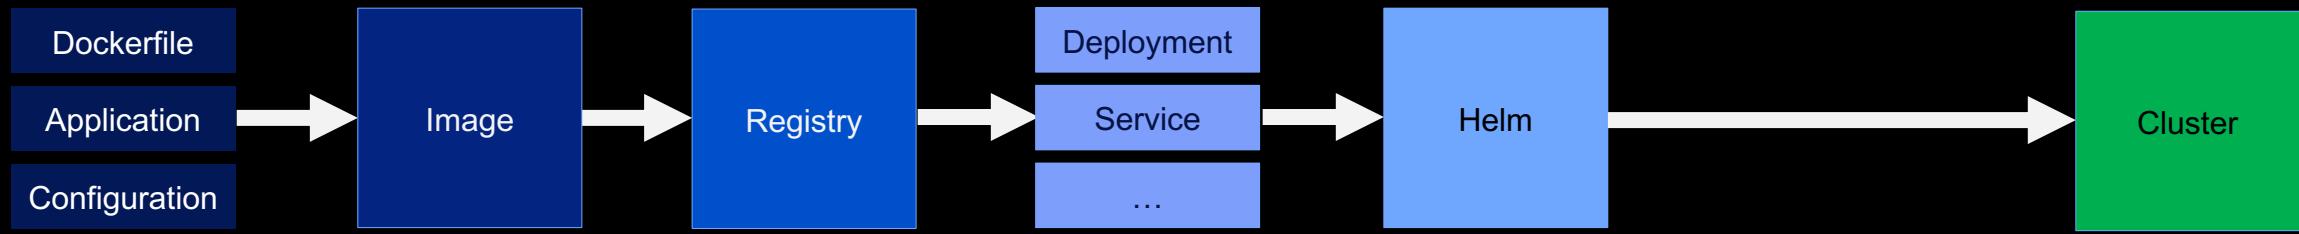
Templates (yaml)

Kubernetes deployment manifests (yaml) with placeholders for external parameters

values.yaml

A set of variables that will be passed on at deploy time to the deployment manifest templates

Minimal Application – Helm



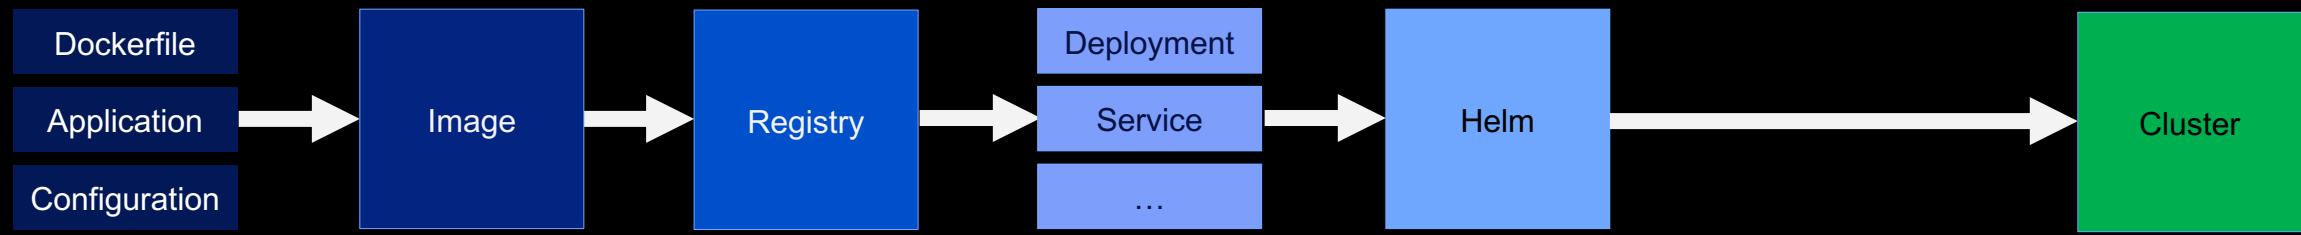
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: k8s-demo-deployment
  labels:
    app: k8s-demo
spec:
  replicas: {{ .Values.replicas }}
  selector:
    matchLabels:
      app: k8s-demo
  template:
    metadata:
      labels:
        app: k8s-demo
    spec:
      containers:
        - name: k8s-demo
          image: {{ .Values.image.name }}:{{ .Values.image.tag }}
```

helm install k8s-demo

```
apiVersion: v1
image:
  name:
    - niklaushirt/k8s-demo
  tag: 1.0
  pullPolicy: IfNotPresent
replicas: 1
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: k8s-demo-deployment
  labels:
    app: k8s-demo
spec:
  replicas: 1
  selector:
    matchLabels:
      app: k8s-demo
  template:
    metadata:
      labels:
        app: k8s-demo
    spec:
      containers:
        - name: k8s-demo
          image: niklaushirt/k8s-demo:1.0
          ports:
            - containerPort: 80
```

Minimal Application – Helm



```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: k8s-demo-deployment
  labels:
    app: k8s-demo
spec:
  replicas: {{ .Values.replicas }}
  selector:
    matchLabels:
      app: k8s-demo
  template:
    metadata:
      labels:
        app: k8s-demo
    spec:
      containers:
        - name: k8s-demo
          image: {{ .Values.image.name }}:{{ .Values.image.tag }}
          ports:
            - containerPort: 80
  
```

```

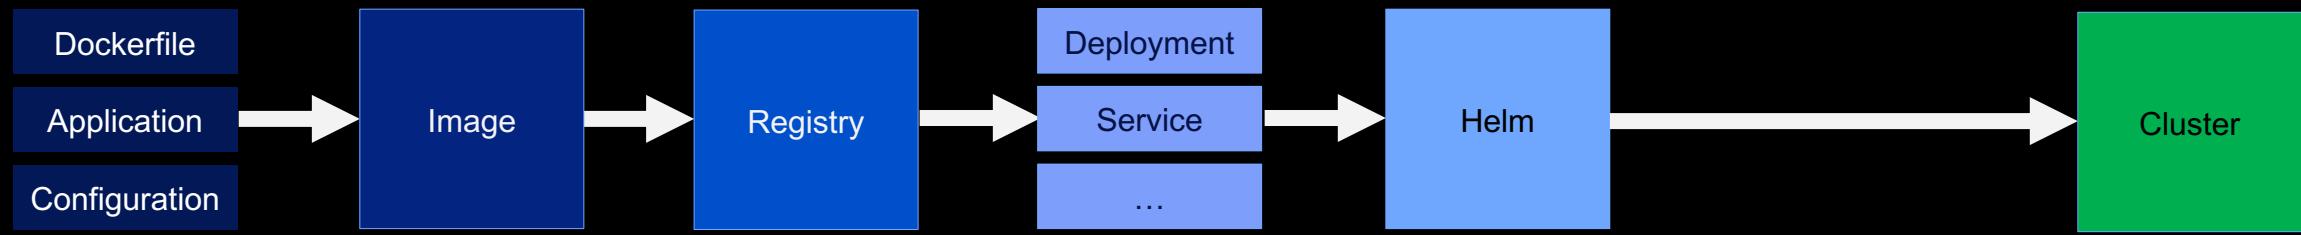
apiVersion: v1
image:
  name:
    - niklaushirt/k8s-demo
  tag: 1.1
  pullPolicy: IfNotPresent
replicas: 3
  
```

`helm install k8s-demo`

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: k8s-demo-deployment
  labels:
    app: k8s-demo
spec:
  replicas: 3
  selector:
    matchLabels:
      app: k8s-demo
  template:
    metadata:
      labels:
        app: k8s-demo
    spec:
      containers:
        - name: k8s-demo
          image: niklaushirt/k8s-demo:1.1
          ports:
            - containerPort: 80
  
```

Minimal Application – Helm



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: k8s-demo-deployment
  labels:
    app: k8s-demo
spec:
  replicas: {{ .Values.replicas }}
  selector:
    matchLabels:
      app: k8s-demo
  template:
    metadata:
      labels:
        app: k8s-demo
    spec:
      containers:
        - name: k8s-demo
          image: {{ .Values.image.name }}:{{ .Values.image.tag }}
          ports:
            - containerPort: 80
```

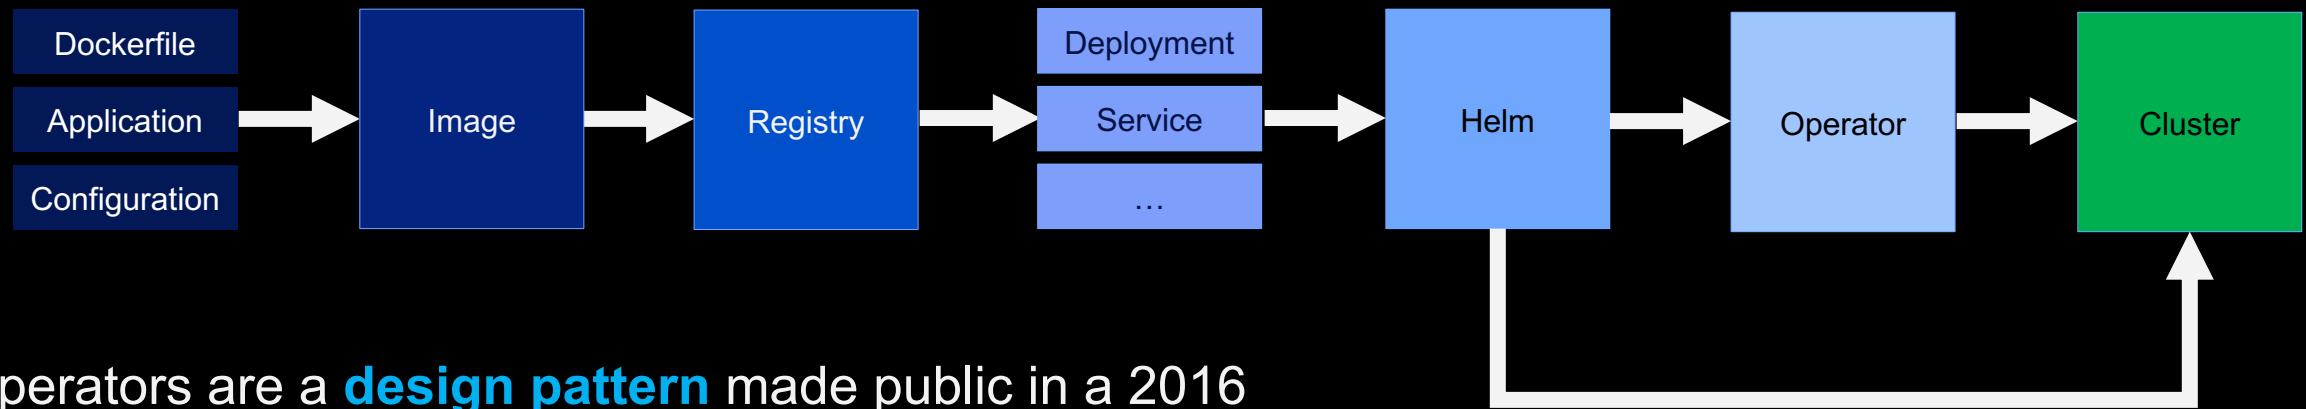
```
apiVersion: v1
image:
  name:
    - niklaushirt/k8s-demo
  tag: 1.0
  pullPolicy: IfNotPresent
replicas: 1
```

helm install --set **replicas=3** k8s-demo

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: k8s-demo-deployment
  labels:
    app: k8s-demo
spec:
  replicas: 3
  selector:
    matchLabels:
      app: k8s-demo
  template:
    metadata:
      labels:
        app: k8s-demo
    spec:
      containers:
        - name: k8s-demo
          image: niklaushirt/k8s-demo:1.0
          ports:
            - containerPort: 80
```



Minimal Application – Operators



Operators are a **design pattern** made public in a 2016 CoreOS [blog](#) post.

The goal of an Operator is to put **operational knowledge into software**.

Operators implement and automate common **Day-1** (installation, configuration, etc) and **Day-2** (re-configuration, update, backup, failover, restore, etc.) **activities** in a piece of software running inside your Kubernetes cluster, by integrating natively with Kubernetes concepts and APIs.

K9s – the Norton Commander of Kubernetes



```
Context: services/c107-e-us-south-containers-cloud-ibm-com:3_ <0> all      <ctrl-d> Delete      <y> YAML
Cluster: c107-e-us-south-containers-cloud-ibm-com:30958   <1> services    <d> Describe
User: IAMnikh@ch.ibm.com/c107-e-us-south-containers-clou... <2> default    <e> Edit
K9s Rev: 0.14.0                                         <e> Edit
K8s Rev: v1.16.2                                         <k> Kill
CPU: 24%                                                 <l> Logs
MEM: 55%                                                 <shift-l> Logs Previous
                                                        <s> Shell

Pod(default)[8]
NAME          READY   STATUS    RS   CPU   MEM   %CPU/R   %MEM/R   %CPU/L   %MEM/L IP           NODE   QOS   AGE
dashboard-78dcdbfd49-wp2r2   1/1   Terminated   0   0   21   0   21   0   0   172.30.153.228   10.94.80.44   BU   6m58s
dashboard-78dcdbfd49-wrwzz  1/1   Running    0   n/a   n/a   n/a   n/a   n/a   n/a   172.30.211.203   10.94.80.49   BU   7s
kubetoy-deployment-866fc8c687-gv9vl  1/1   Running   23   0   16   0   25   0   0   172.30.212.45   10.94.80.49   BU   35d
nginx-policy-deployment-9dd74c6bf-ggnfm  1/1   Running   0   0   1   0   0   0   0   172.30.212.60   10.94.80.49   BE   5d14h
nginx-policy-deployment-9dd74c6bf-t6nrs  1/1   Running   0   0   1   0   0   0   0   172.30.31.248   10.94.80.27   BE   5d14h
openldap-96d6c5f88-fw7lm   1/1   Running   0   0   19   0   0   0   0   172.30.153.207  10.94.80.44   BE   43d
openldap-admin-78cfcbb5d5-bf68p  1/1   Running   0   1   83   0   0   0   0   172.30.153.205  10.94.80.44   BE   43d
student-ui-68c59b4cd4-vcsd5  1/1   Running   0   0   21   0   21   0   0   172.30.211.199   10.94.80.49   BU   27d

<pod>
```

<https://github.com/derailed/k9s>

Preinstalled in the training VM

For dumb people like me
(who just don't get vi)

`KUBE_EDITOR="nano" k9s`

I did create an alias in .bashrc/.zshrc
`alias k9s='KUBE_EDITOR="nano" k9s'`

QUESTIONS?





◦◦ Kubernetes Workshop Series
Kubernetes – Advanced concepts

06

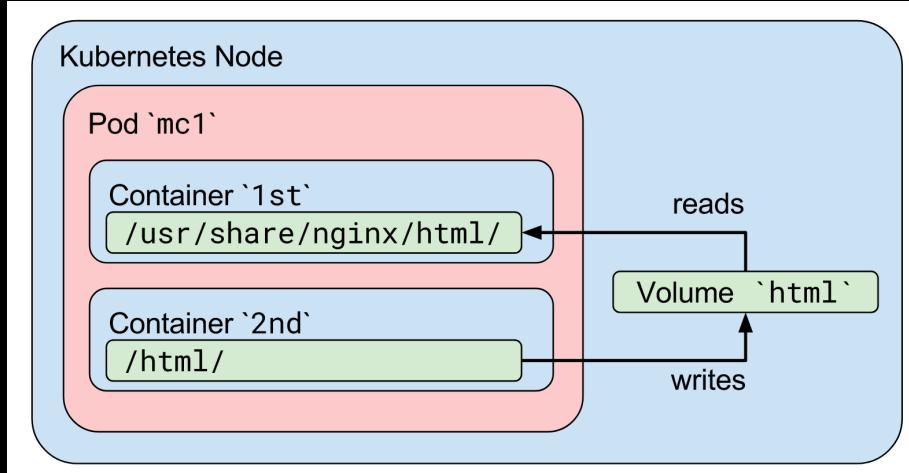


One container per Pod

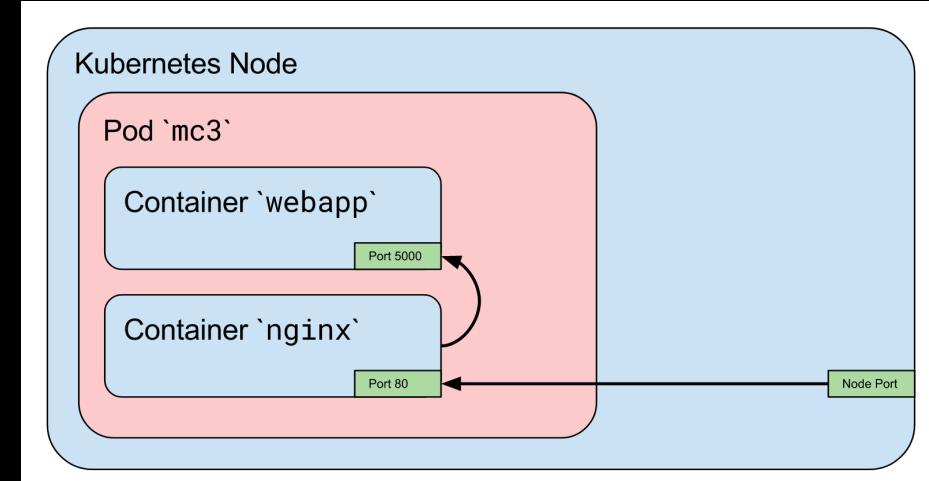
Multi-Container Pod Design Patterns



Shared Volumes

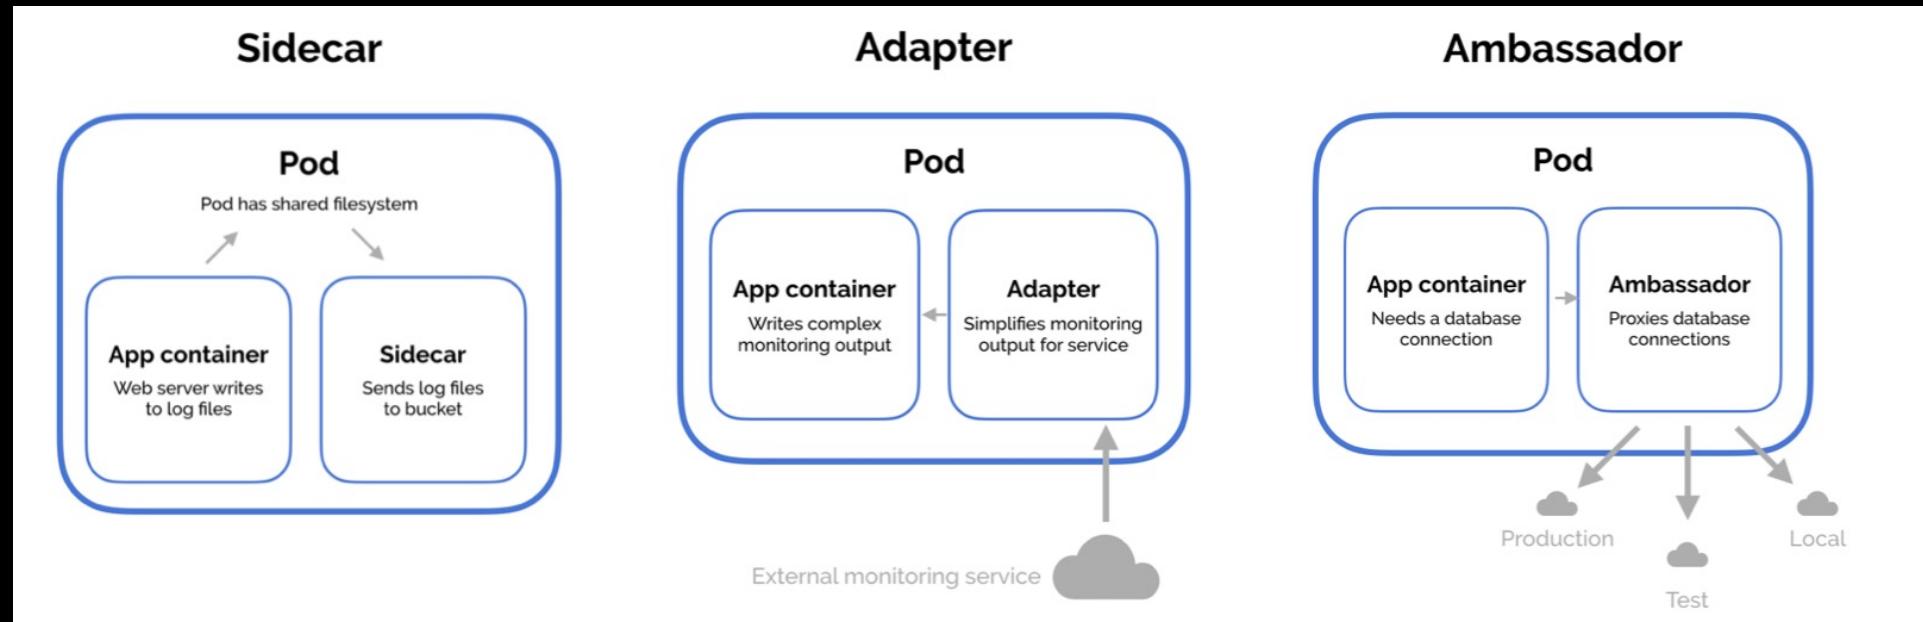


No Network isolation



**In what order containers are being started in a Pod?
The containers scale together...**

Multi-Container Pod Design Patterns



Sidecar pattern

The sidecar pattern consists of a main application—i.e. your web application—plus a helper container with a **responsibility that is essential to your application**, but is not necessarily part of the application itself.

Adapter pattern

The adapter pattern is used to standardize and normalize application output or **monitoring data** for aggregation.

Ambassador pattern

The ambassador pattern is a useful way to connect containers with the outside world (**proxy**).

Requests and Limits

Resources: Requests and Limits



```
apiVersion: apps/v1
kind: Deployment
metadata:
...
spec:
...
template:
...
resources:
  requests:
    memory: 512Mi
    cpu: 100m
  limits:
    memory: 1900Mi
    cpu: 300m
```

When working with containers in Kubernetes, it's important to know what are the resources involved and how they are needed.

Some processes will require more CPU or memory than others. Some are critical and should never be starved.

Knowing that, we should configure our containers and Pods properly in order to get the best of both.



Resources: Requests and Limits

```
apiVersion: apps/v1
kind: Deployment
metadata:
...
spec:
...
template:
...
resources:
  requests:
    memory: 512Mi
    cpu: 100m
  limits:
    memory: 1900Mi
    cpu: 300m
```

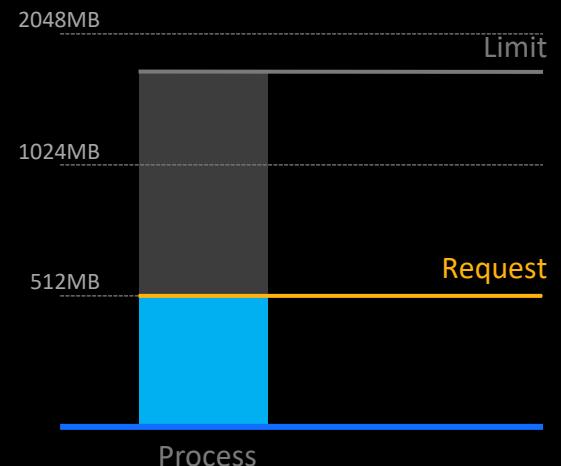
Requests affect how the pods are scheduled in Kubernetes.

When a Pod is created, the scheduler finds the nodes which can accommodate the Pod according to the resource request.

Requests define the **minimum amount of resources that containers need**.

If you think that your app requires at least 256MB of memory to operate, this is the request value.

The application can use more than 256MB, but Kubernetes guarantees a minimum of 256MB to the container.



Error event: **FailedScheduling**



Resources: Requests and Limits

```
apiVersion: apps/v1
kind: Deployment
metadata:
...
spec:
...
template:
...
resources:
  requests:
    memory: 512Mi
    cpu: 100m
limits:
  memory: 1900Mi
  cpu: 300m
```

Limits define the max amount of resources that the container can consume.

They protect the Cluster from runaway Pods that eat up all CPU and/or Memory.

The **CPU limit** defines how much CPU time the container can use.

The Linux kernel (cgroup) checks to see if this limit is exceeded; if so, the kernel waits before allowing that cgroup to resume execution.

Error event:

CPU: **NONE → throttling!**

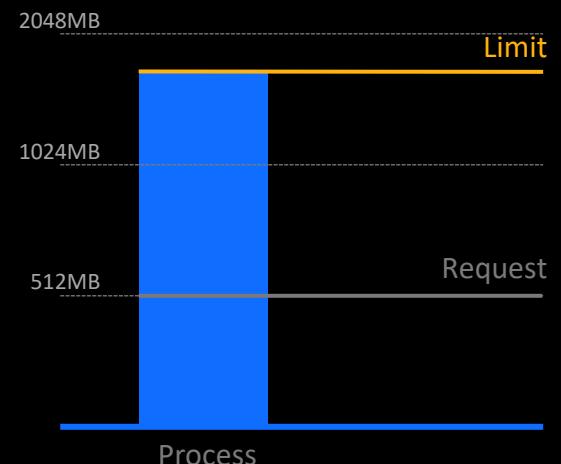
The **memory limit** defines a memory limit for that cgroup.

If the container tries to allocate more memory than this limit, the Linux kernel stops the processes in the container that tried to allocate memory.

If that process is the container's PID 1, and the container is marked as restartable, Kubernetes restarts the container.

Error events:

Memory: **OOM Killed (out of memory)** – only for PID1, otherwise invisible



Probes

Readiness and Liveness Probes

Running

The Pod has been bound to a node, and all of the containers have been created.

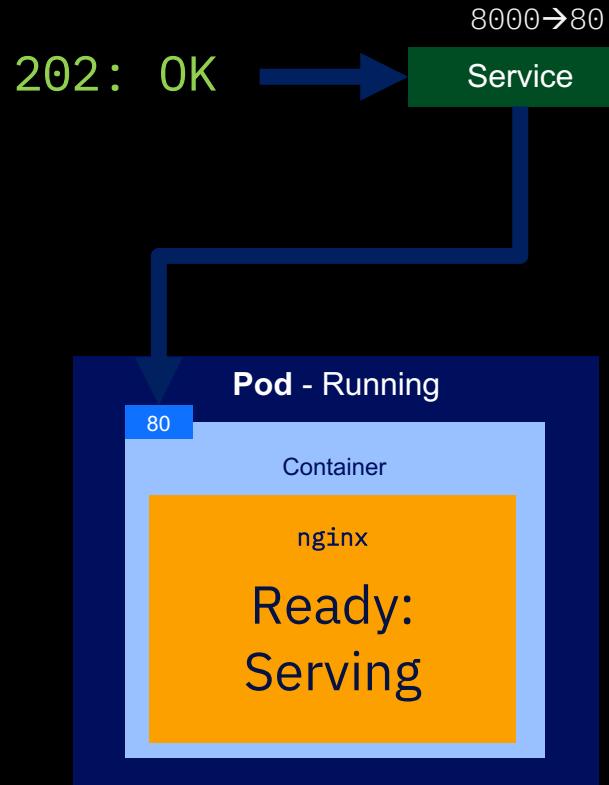
Does not guarantee that the process in the container is fully initialized and Ready to receive traffic.

As soon as a Pod is Running it is also considered Ready and traffic is sent to it.

HTTP Error

Readiness and Liveness Probes

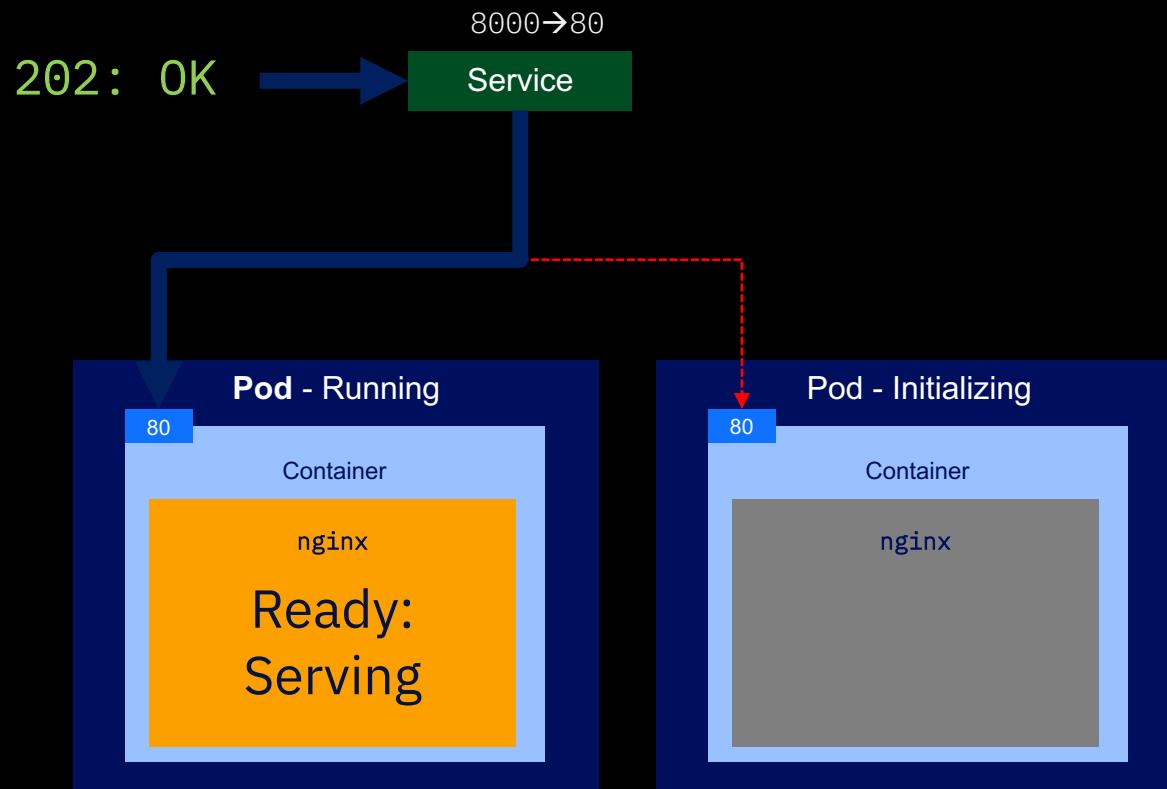
- NGINX Application running



```
apiVersion: v1
kind: Deployment
metadata:
  name: nginx
spec:
  replicas: 1
  ...
```

Readiness and Liveness Probes

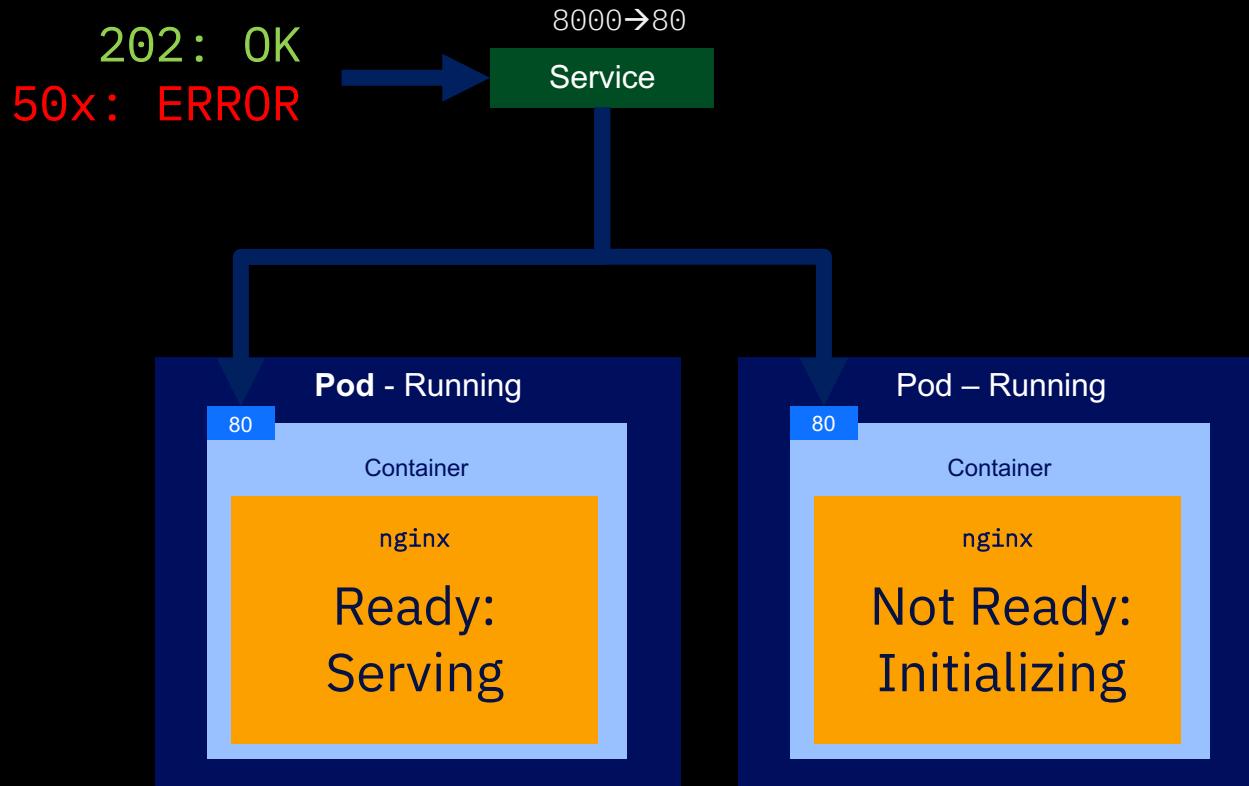
- NGINX Application running
- Scale-up



```
apiVersion: v1
kind: Deployment
metadata:
  name: nginx
spec:
  replicas: 2
  ...
```

Readiness and Liveness Probes

- NGINX Application running
- Scale-up
- Traffic routed to the new Pod as soon as it's Running, even if the process is not ready



```
apiVersion: v1
kind: Deployment
metadata:
  name: nginx
spec:
  replicas: 2
  ...
  readinessProbe:
    httpGet:
      path: /healthz
      port: 80
    initialDelaySeconds: 30
    periodSeconds: 10
    successThreshold: 1
    failureThreshold: 5
  livenessProbe:
    httpGet:
      path: /liveness
      port: 80
    initialDelaySeconds: 30
    periodSeconds: 10
    successThreshold: 1
    failureThreshold: 5
```

Readiness and Liveness Probes

RUNNING ≠ READY

The image displays two screenshots illustrating the difference between a pod's running state and its readiness.

Kubernetes Pods Table: On the left, a screenshot of a Kubernetes dashboard shows a table of pods. The columns are Name, Status, Ready, and Restart Count. A specific pod named "cp4waiops-demo-ui-7fbbf8cc5d-7tsmn" is highlighted. The "Status" column shows "Running" and the "Ready" column shows "0/1". Both the "Status" and "Ready" columns are enclosed in red boxes.

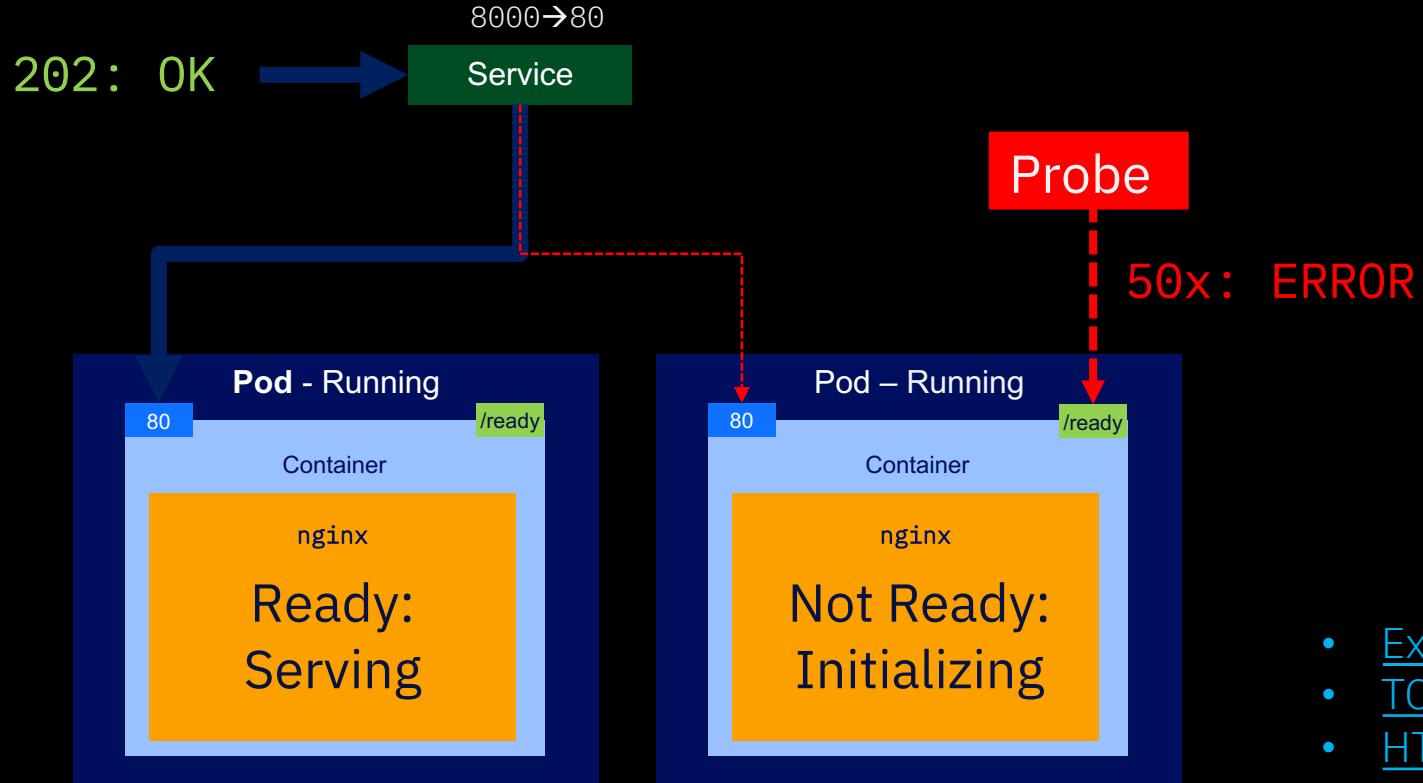
Name	Status	Ready	Restart Count
cp4waiops-demo-ui-7fbbf8cc5d-7tsmn	Running	0/1	0

Service Status Page: On the right, a screenshot of a service status page titled "Départ Abfahrt Partenza" (Departure Departure Departure) shows an error message: "Application is not available. The application is currently not serving requests at this endpoint. It may not have been started or is still starting." Below this, a list of possible reasons is provided:

- Possible reasons you are seeing this page:
 - The host doesn't exist. Make sure the hostname was typed correctly and that a route matching this hostname exists.
 - The host exists, but doesn't have a matching path. Check if the URL path was typed correctly and that the route was created using the desired path.
 - Route and path matches, but all pods are down. Make sure that the resources exposed by this route (pods, services, deployment configs, etc) have at least one pod running.

Readiness Probes

- NGINX Application running
- Readiness Probe checks /ready Endpoint
- If no response, keep status Initializing

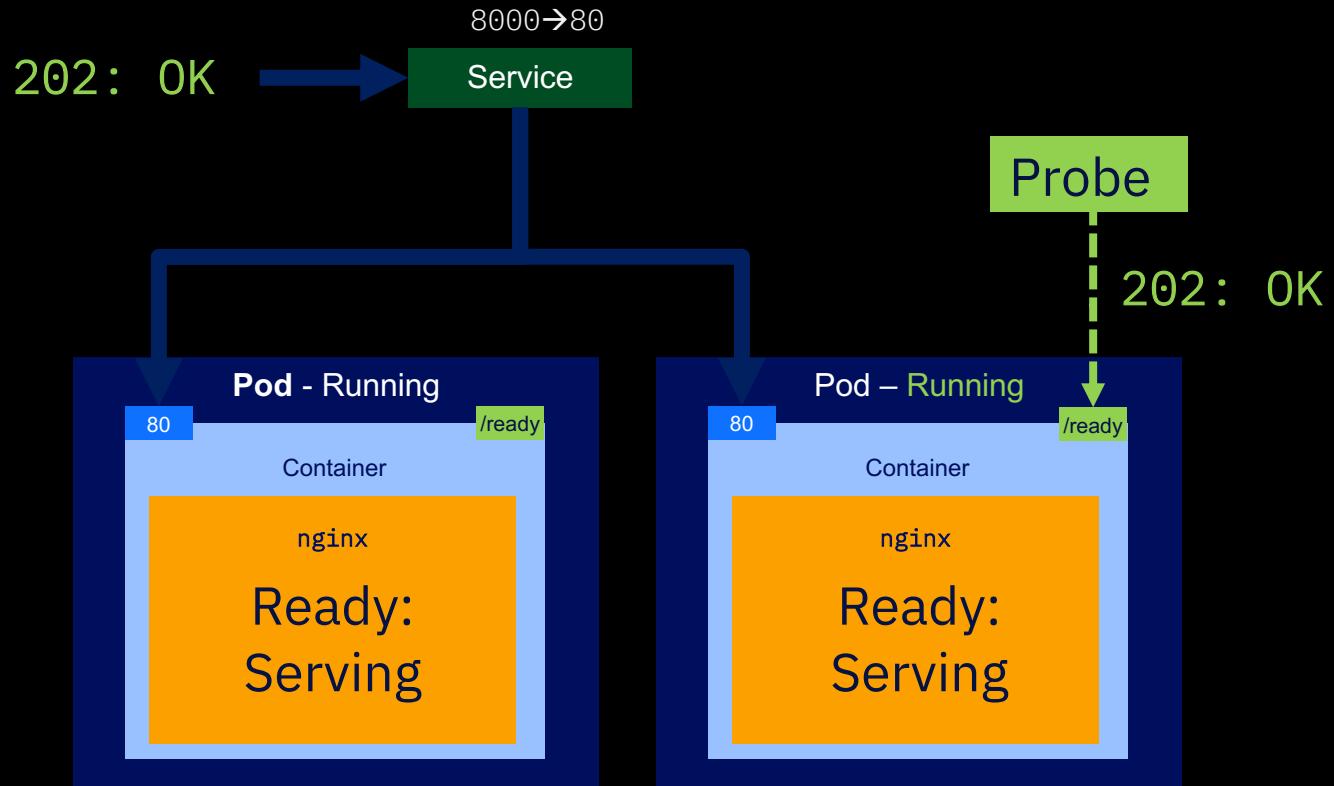


```
apiVersion: v1
kind: Deployment
metadata:
  name: nginx
spec:
  replicas: 2
...
readinessProbe:
  # an http probe
  httpGet:
    path: /ready
    port: 8080
  initialDelaySeconds: 20
  periodSeconds: 5
```

- ExecAction: Executes a specified command
- TCPSocketAction: Performs a TCP check
- HTTPGetAction: Performs an HTTP GET request

Readiness Probes

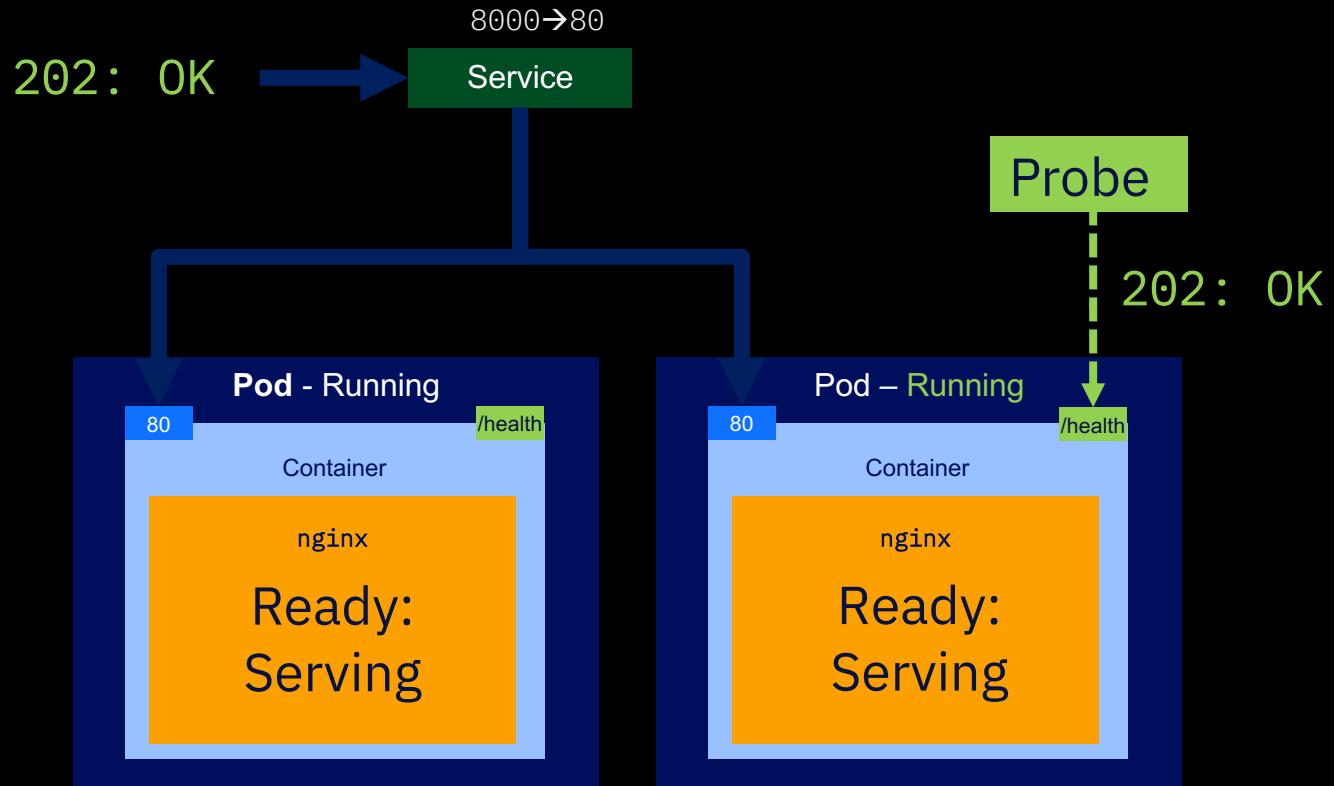
- NGINX Application running
- Readiness Probe checks /ready Endpoint
- If response, set status Ready



```
apiVersion: v1
kind: Deployment
metadata:
  name: nginx
spec:
  replicas: 2
...
readinessProbe:
  # an http probe
  httpGet:
    path: /ready
    port: 8080
  initialDelaySeconds: 20
  periodSeconds: 5
```

Liveness Probes

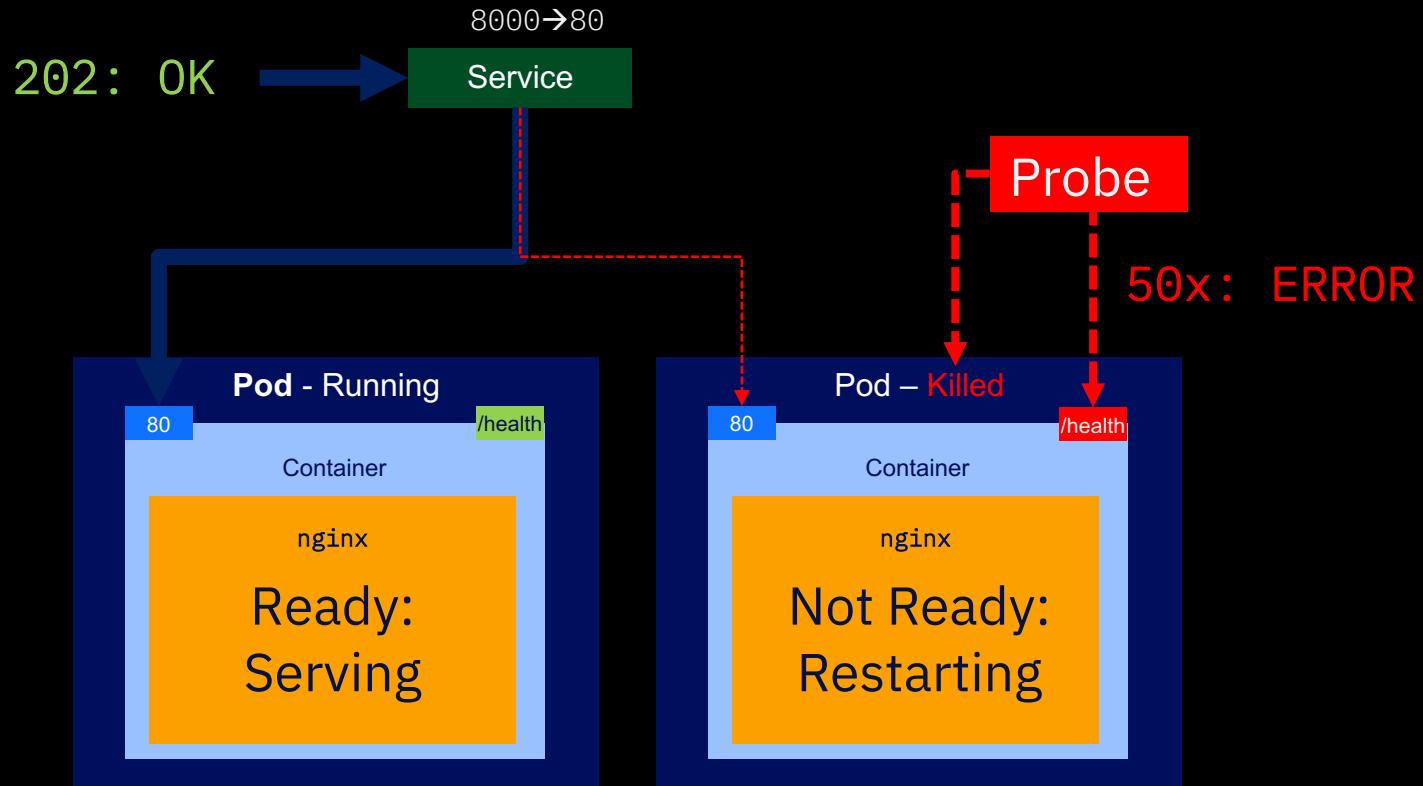
- NGINX Application running
- Liveness Probe checks /health Endpoint
- If no response, restarts Pod



```
apiVersion: v1
kind: Deployment
metadata:
  name: nginx
spec:
  replicas: 2
...
livenessProbe:
  # an http probe
  httpGet:
    path: /health
    port: 8080
  initialDelaySeconds: 20
  periodSeconds: 5
```

Liveness Probes

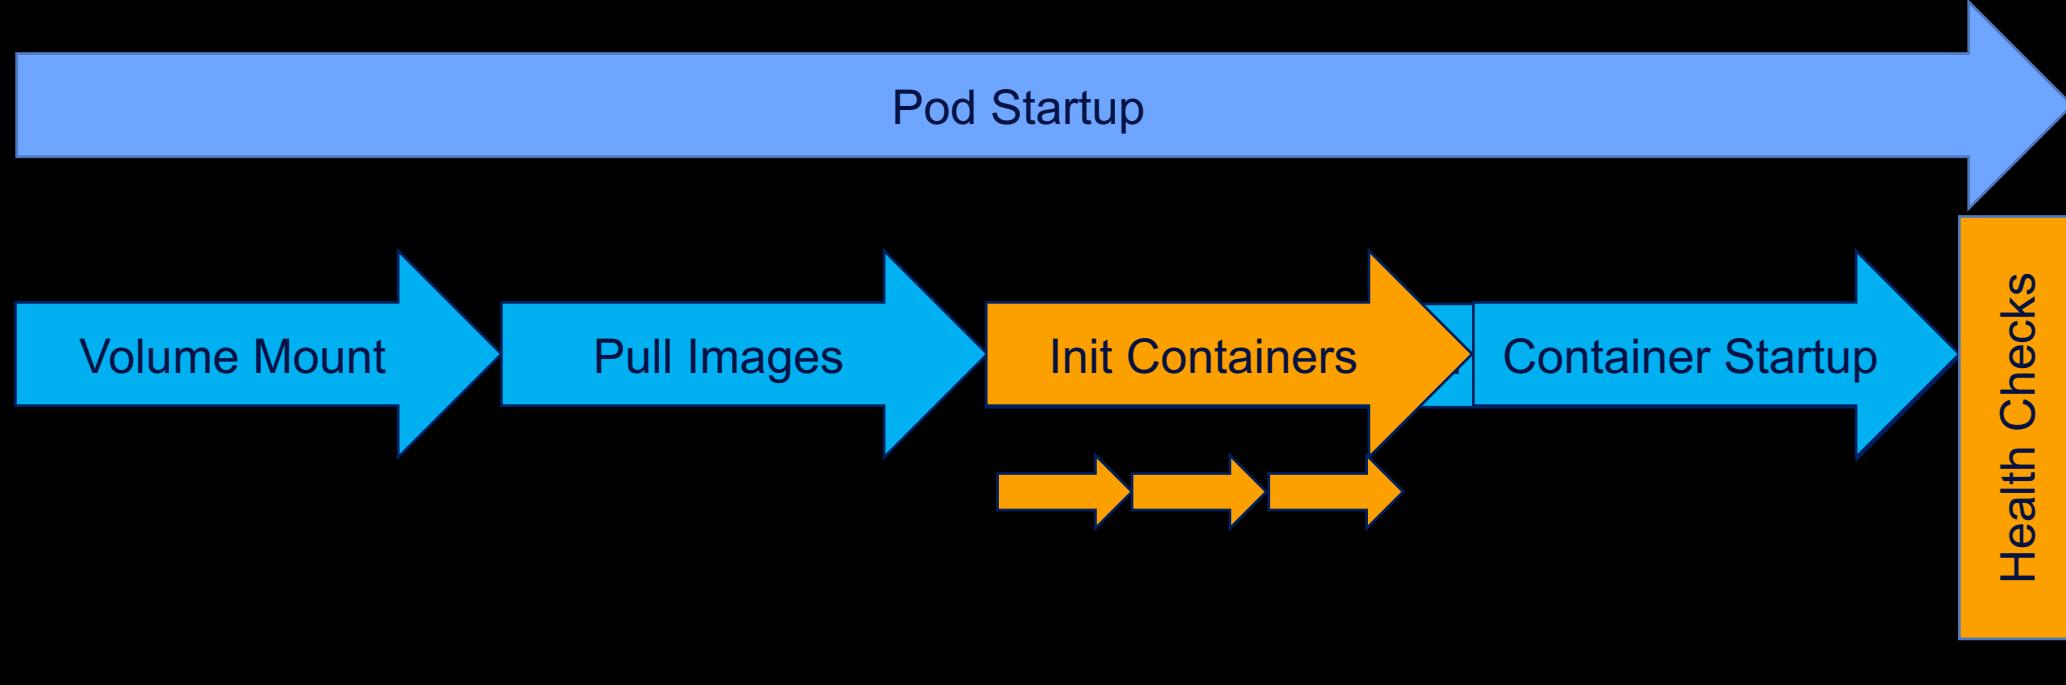
- NGINX Application running
- Liveness Probe checks /health Endpoint
- If no response, restarts Pod



```
apiVersion: v1
kind: Deployment
metadata:
  name: nginx
spec:
  replicas: 2
...
livenessProbe:
  # an http probe
  httpGet:
    path: /health
    port: 8080
  initialDelaySeconds: 20
  periodSeconds: 5
```

Init Containers

Init Containers - Pod Design Pattern



- Init Containers can contain **utilities** or custom code for setup that are not present in an app image.
- Init Containers can run with a **different view of the filesystem** than app containers in the same Pod.
- Init Containers offer a mechanism to **block or delay app container startup** until a set of preconditions are met.
- Init Containers can **securely run utilities or custom code** that would otherwise make an app container image less secure.

Init Containers - Pod Design Pattern



Examples:

- Wait for a Service to be created, using a shell one-line command
- Register the Pod with a remote server from the downward API with a command
- Wait for some time before starting the app container
- Clone a Git repository into a Volume
- Place values into a configuration file and run a template tool to dynamically generate a configuration file for the main app container.

Init Containers

```
apiVersion: app/v1
kind: Deployment
metadata:
  name: nginx-deployment
...
```

Init Containers

Example

- Sleep before starting NGINX

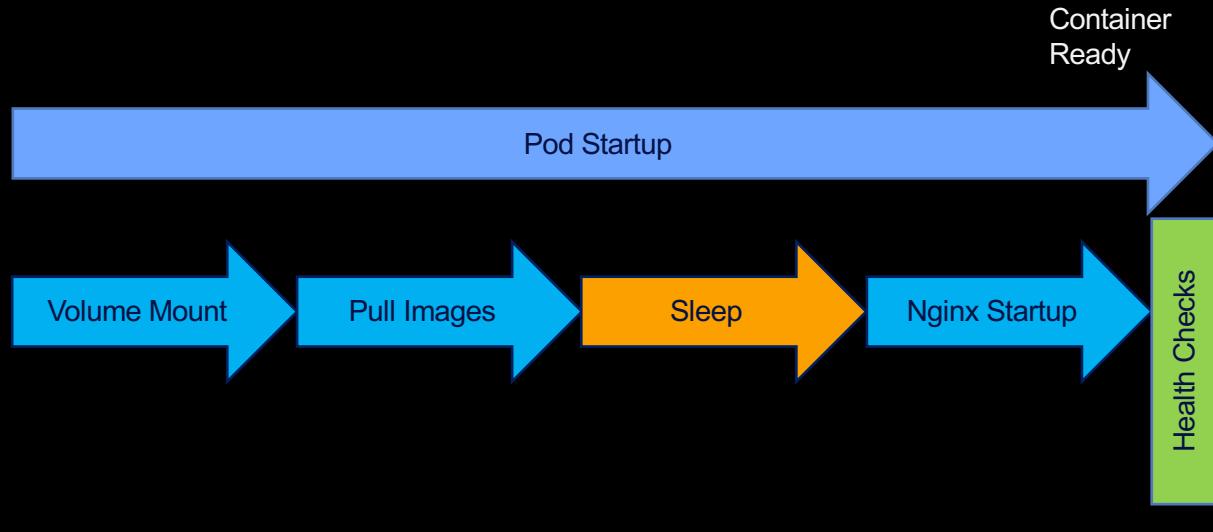
example.yaml

```
apiVersion: app/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 2
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.9
          ports:
            - containerPort: 80
```

Init Containers

Example

- Sleep before starting NGINX



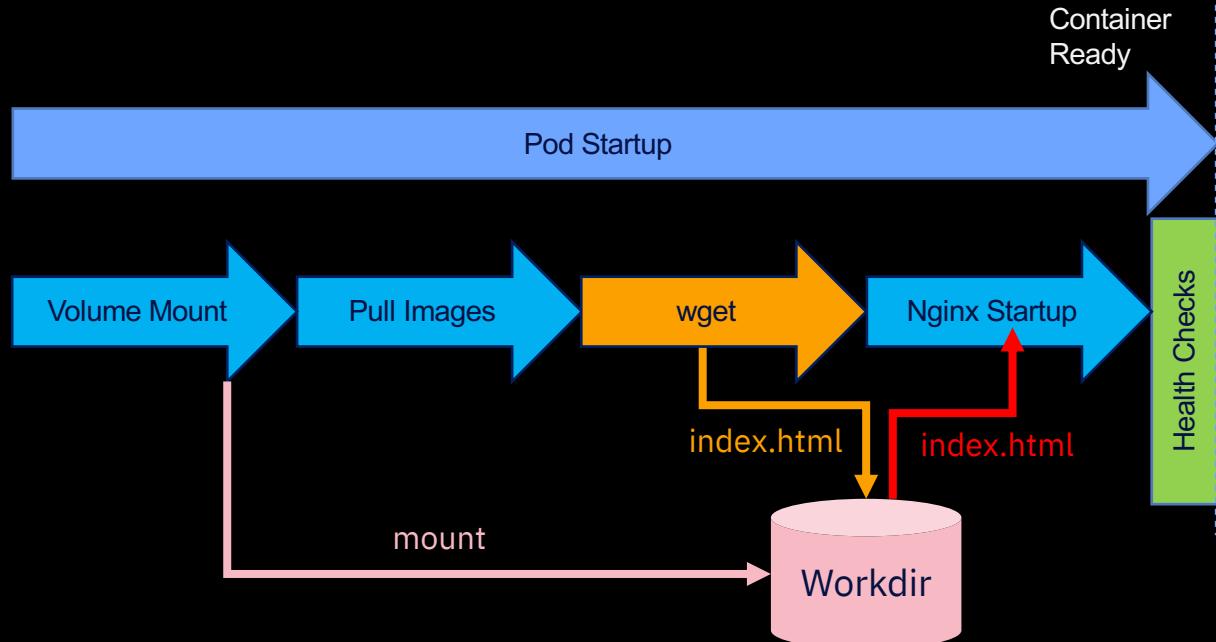
example.yaml

```
apiVersion: app/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 2
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.7.9
        ports:
        - containerPort: 80
#This container is run during pod initialization
initContainers:
#Sleep for 2 Seconds
- name: sleep-1
  image: busybox
  imagePullPolicy: IfNotPresent
  command: ['sh', '-c', 'sleep 2;echo sleep complete;']
```

Init Containers

Example

- NGINX Data initialization

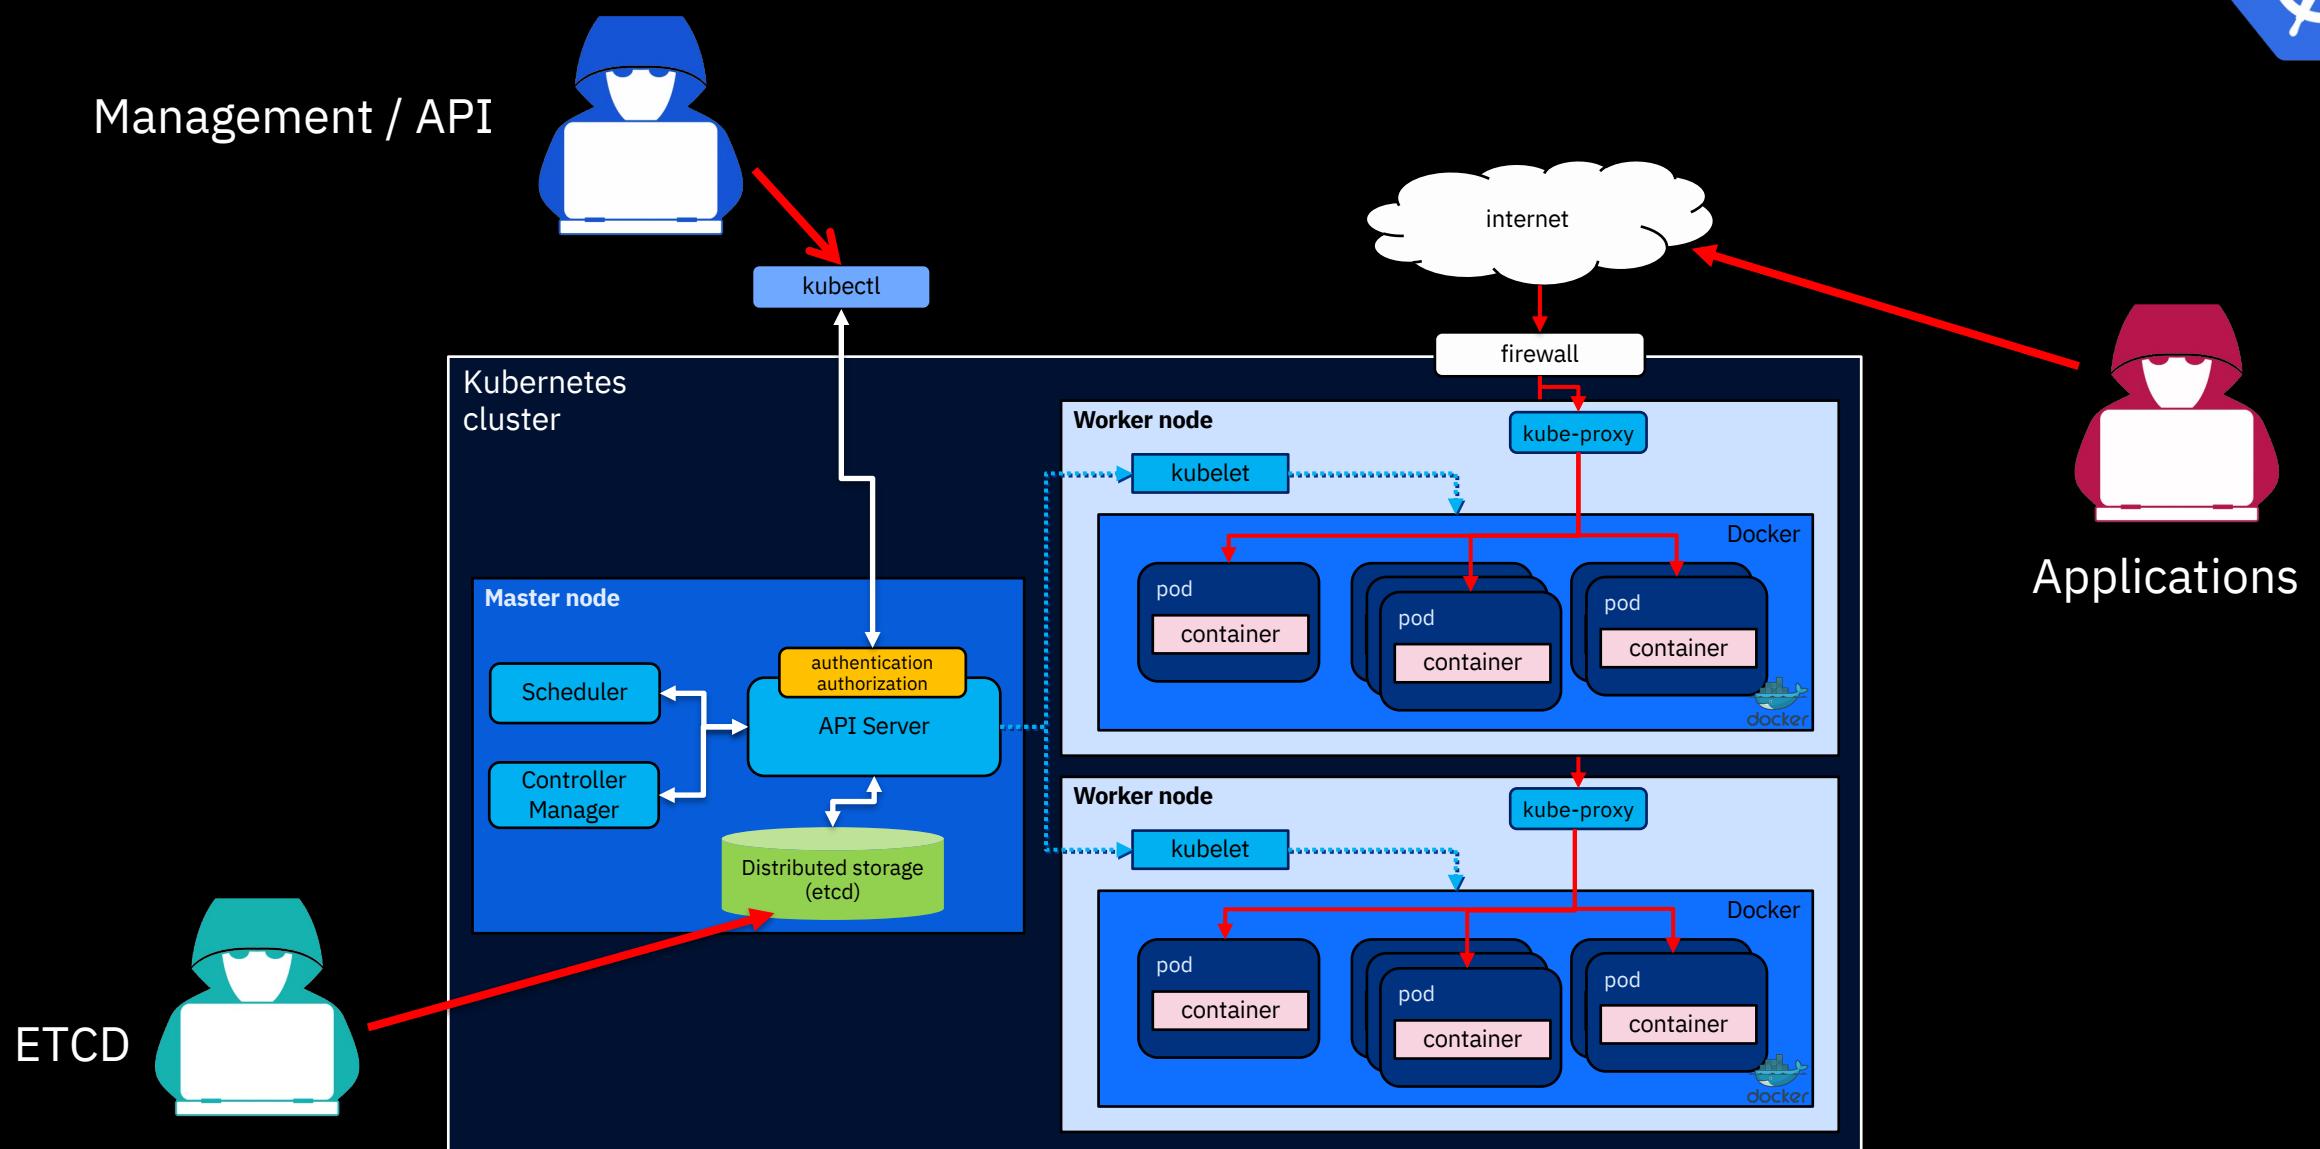


example.yaml

```
apiVersion: app/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 2
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.9
          ports:
            - containerPort: 80
      #Volume mapped into the nginx pod
      volumeMounts:
        - name: workdir
          mountPath: /usr/share/nginx/html
      #This container is run during pod initialization
      initContainers:
        - name: install
          image: busybox
          command:
            - wget
            - "-O"
            - "/work-dir/index.html"
            - "http://kubernetes.io"
          volumeMounts:
            - name: workdir
              mountPath: "/work-dir"
      volumes:
        - name: workdir
          emptyDir: {}
```

Basic Security

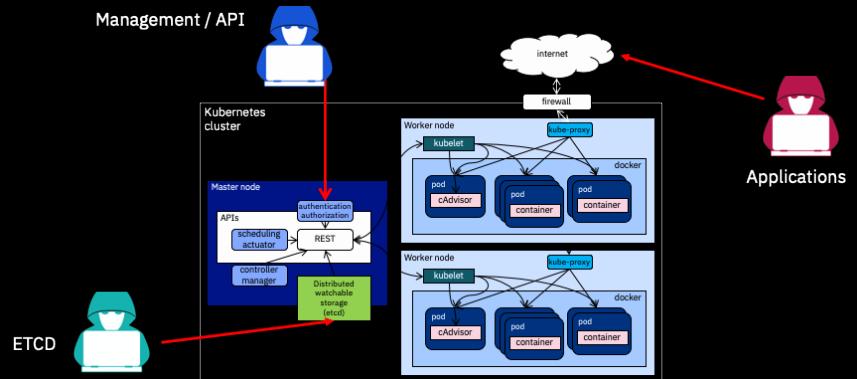
Kubernetes – Security – Attack surface



Kubernetes – Security Basic Topics

Reduce Kubernetes Attack Surfaces

- Secure access to etcd
- Controlling access to the Kubernetes API
- Controlling access to the Kubelet
- Enforce resource usage limits for workloads
- Rotate infrastructure credentials frequently
- Enable audit logging
- Use Linux security features
- Controlling what privileges containers run with
- Enforcing Network Policies
- Image Scanning of your containers
- Use Kubernetes secrets



Source: <https://kubernetes.io/docs/tasks/administer-cluster/securing-a-cluster>
<https://kubernetes.io/blog/2018/07/18/11-ways-not-to-get-hacked/>

RBAC Basic Elements

Rules

Operations (verbs) that can be carried out on a group of resources which belong to different API Groups.

Roles and ClusterRoles

Both consist of rules.

- Role: applicable to a single namespace
- ClusterRole: is cluster-wide

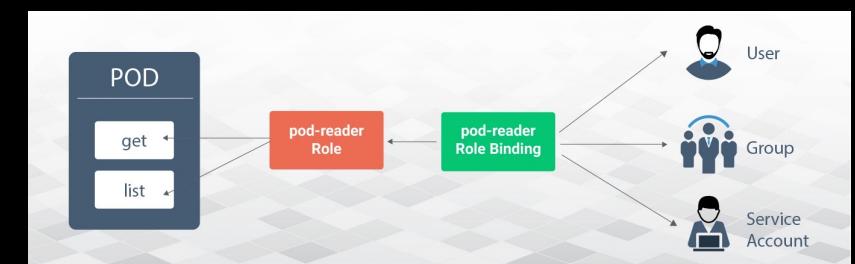
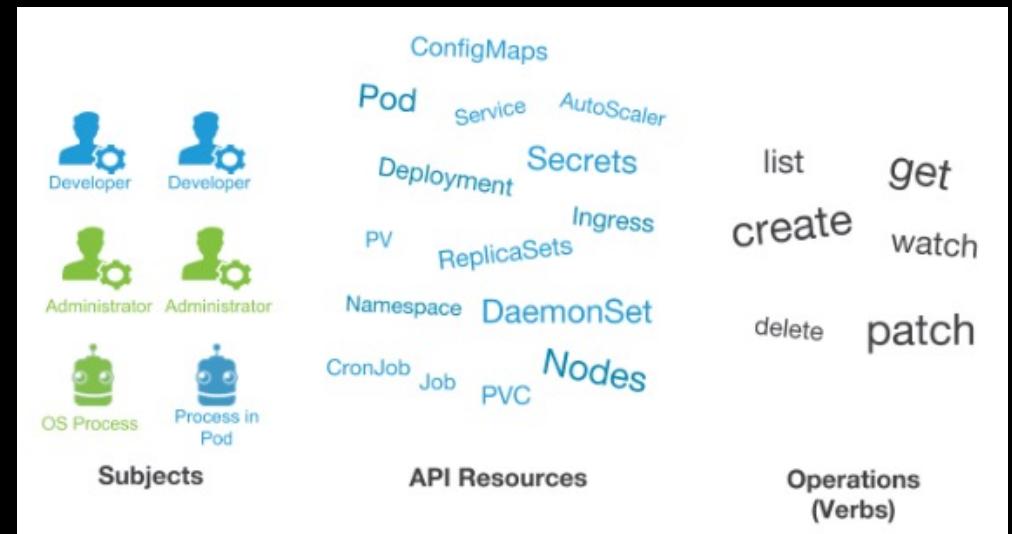
Subjects

Entity that attempts an operation in the cluster

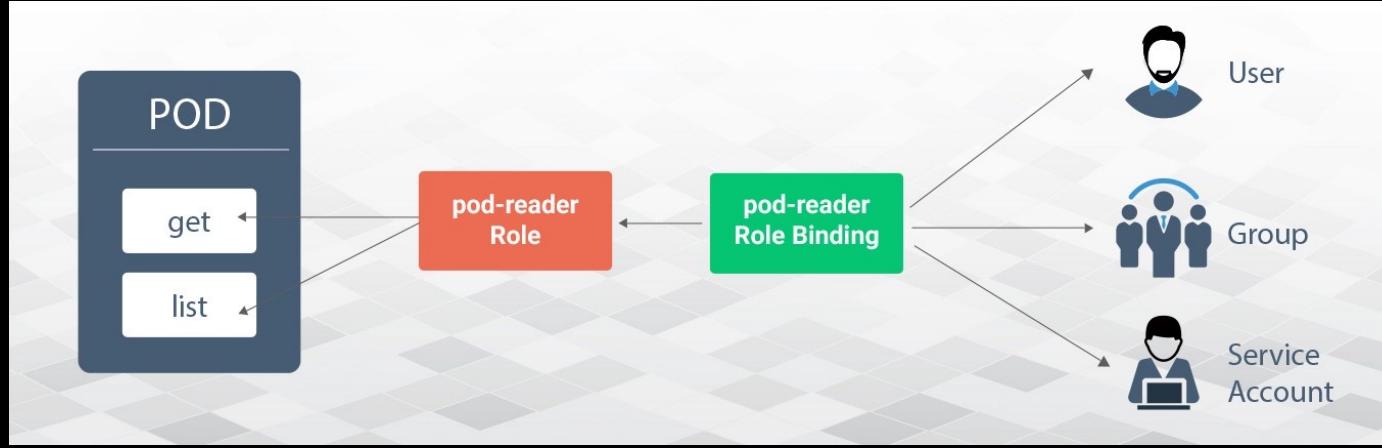
- User Accounts: Humans or processes living outside the cluster.
- Service Accounts: Namespaced account.
- Groups: Reference multiple accounts. (default groups like cluster-admin)

RoleBindings and ClusterRoleBindings

Bind subjects to roles



RBAC Example



```
kind: Role
metadata:
  name: pod-reader
rules:
- apiGroups: []
  resources: ["pods"]
  verbs: ["get","list"]
```

```
kind: RoleBinding
metadata:
  name: pod-reader
subjects:
- kind: Serviceaccount
  name: my-account
roleRef:
  kind: Role
  name: pod-reader
```

Troubleshooting Debugging

Most common errors

```
> kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-pull-error	0/1	ImagePullBackOff	0	31s
nginx-crashloop	0/1	CrashLoopBackOff	2	27s

Most common errors

```
> kubectl get pods
NAME          READY   STATUS        RESTARTS   AGE
nginx-pull-error  0/1    ImagePullBackOff  0          31s
nginx-crashloop  0/1    CrashLoopBackOff  2          27s
```

ImagePullBackOff

- 1.The specified image is not present in the registry.
- 2.A typo in the image name or tag.
- 3.Image pull access denied from the given registry due to credential issue.

```
apiVersion: v1
kind: Deployment
metadata:
  name: nginx-xyz
spec:
...
  containers:
    - name: nginx-xyz
      image: 'nginx-fail:1.7.9'
  serviceAccountName: default
  imagePullSecrets:
    - name: default-dockercfg-8q42l
...
```

Most common errors

```
> kubectl get pods
NAME          READY   STATUS        RESTARTS   AGE
nginx-pull-error  0/1    ImagePullBackOff  0          31s
nginx-crashloop  0/1    CrashLoopBackOff  2          27s
```

CrashLoopBackOff

- 1.Check Logs
- 2.Check Probes
- 3.kubectl debug
- 4....

exec /usr/sbin/nginx: exec format error

```
java.lang.NullPointerException: null
at
org.springframework.boot.actuate.metrics.web.servlet.LongTaskTimingHandlerInterceptor.stop
TaskTimers(LongTaskTimingHandlerInterceptor.java:123) ~[spring-boot-actuator-
2.3.3.RELEASE.jar!/:2.3.3.RELEASE]
at
org.springframework.boot.actuate.metrics.web.servlet.LongTaskTimingHandlerInterceptor.after
completion(LongTaskTimingHandlerInterceptor.java:79) ~[spring-boot-actuator-
2.3.3.RELEASE.jar!/:2.3.3.RELEASE]
...
at org.apache.tomcat.util.net.SocketProcessorBase.run(SocketProcessorBase.java:49) [tomcat-
embed-core-9.0.37.jar!/:9.0.37]
at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1149)
[na:1.8.0_302]
at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:624)
[na:1.8.0_302]
at org.apache.tomcat.util.threads.TaskThread$WrappingRunnable.run(TaskThread.java:61) [tomcat-
embed-core-9.0.37.jar!/:9.0.37]
at java.lang.Thread.run(Thread.java:748) [na:1.8.0_302]
2023-04-03 14:20:09.012 INFO 1 --- [extShutdownHook] j.LocalContainerEntityManagerFactoryB
: Closing JPA EntityManagerFactory for persistence unit 'default'
```

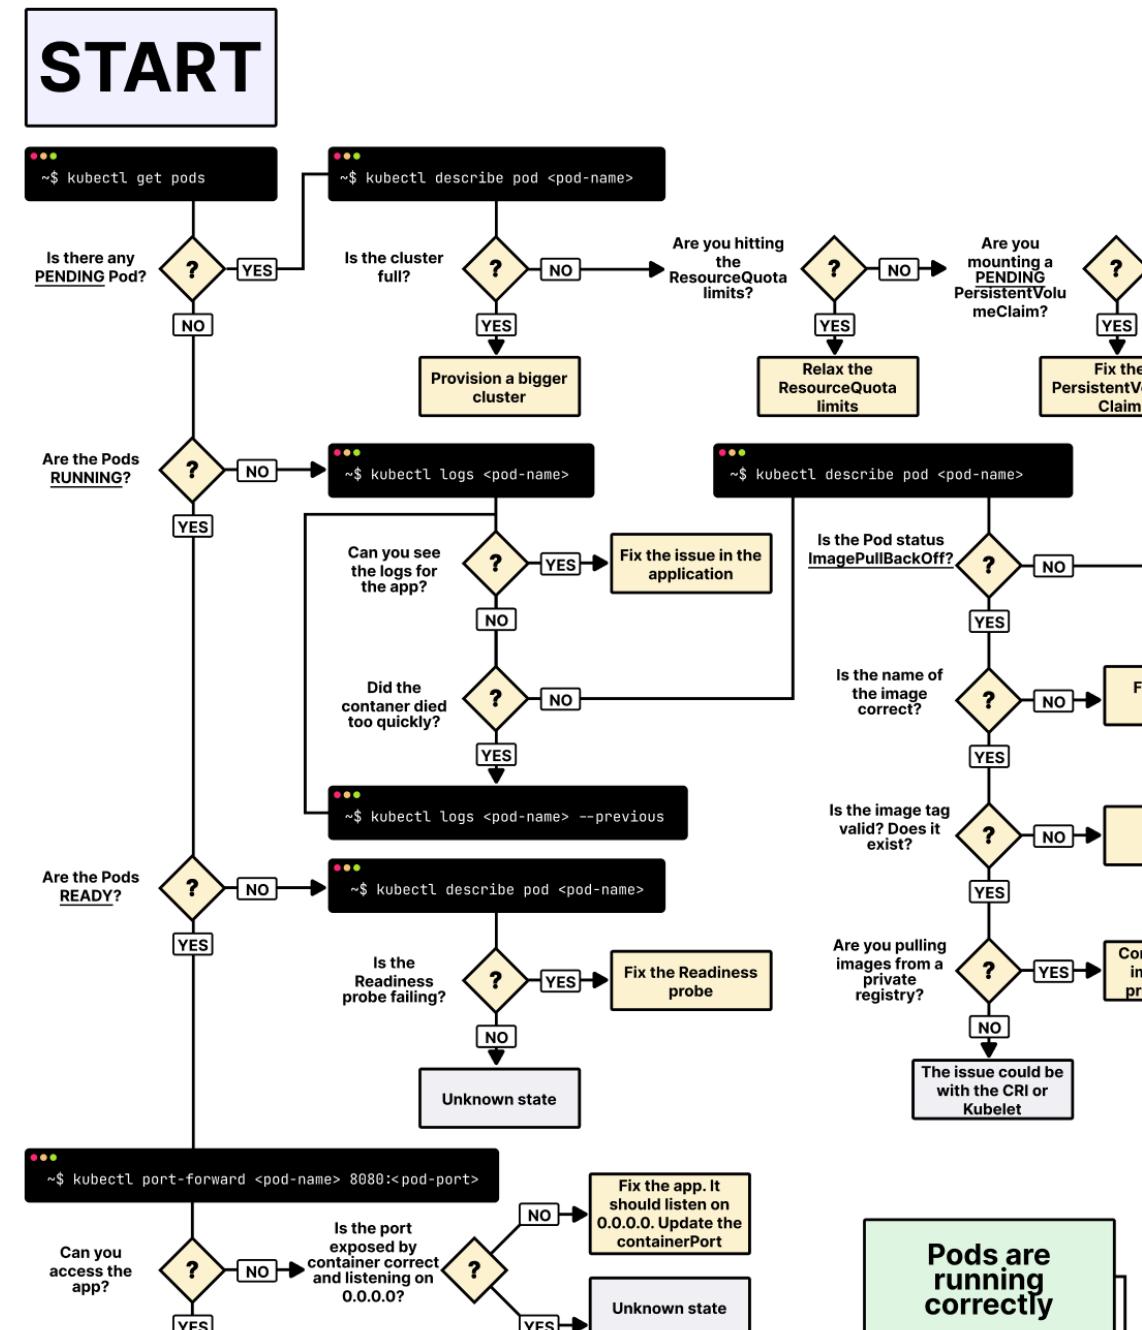
Troubleshooting

A visual guide on troubleshooting Kubernetes deployments

<https://learnk8s.io/troubleshooting-deployments>

How To Troubleshoot Kubernetes Pods: Beginners

<https://devopscube.com/troubleshoot-kubernetes-pods/>



Dev/GitOps

DevOps – **Git**Ops and Dev**Sec**Ops

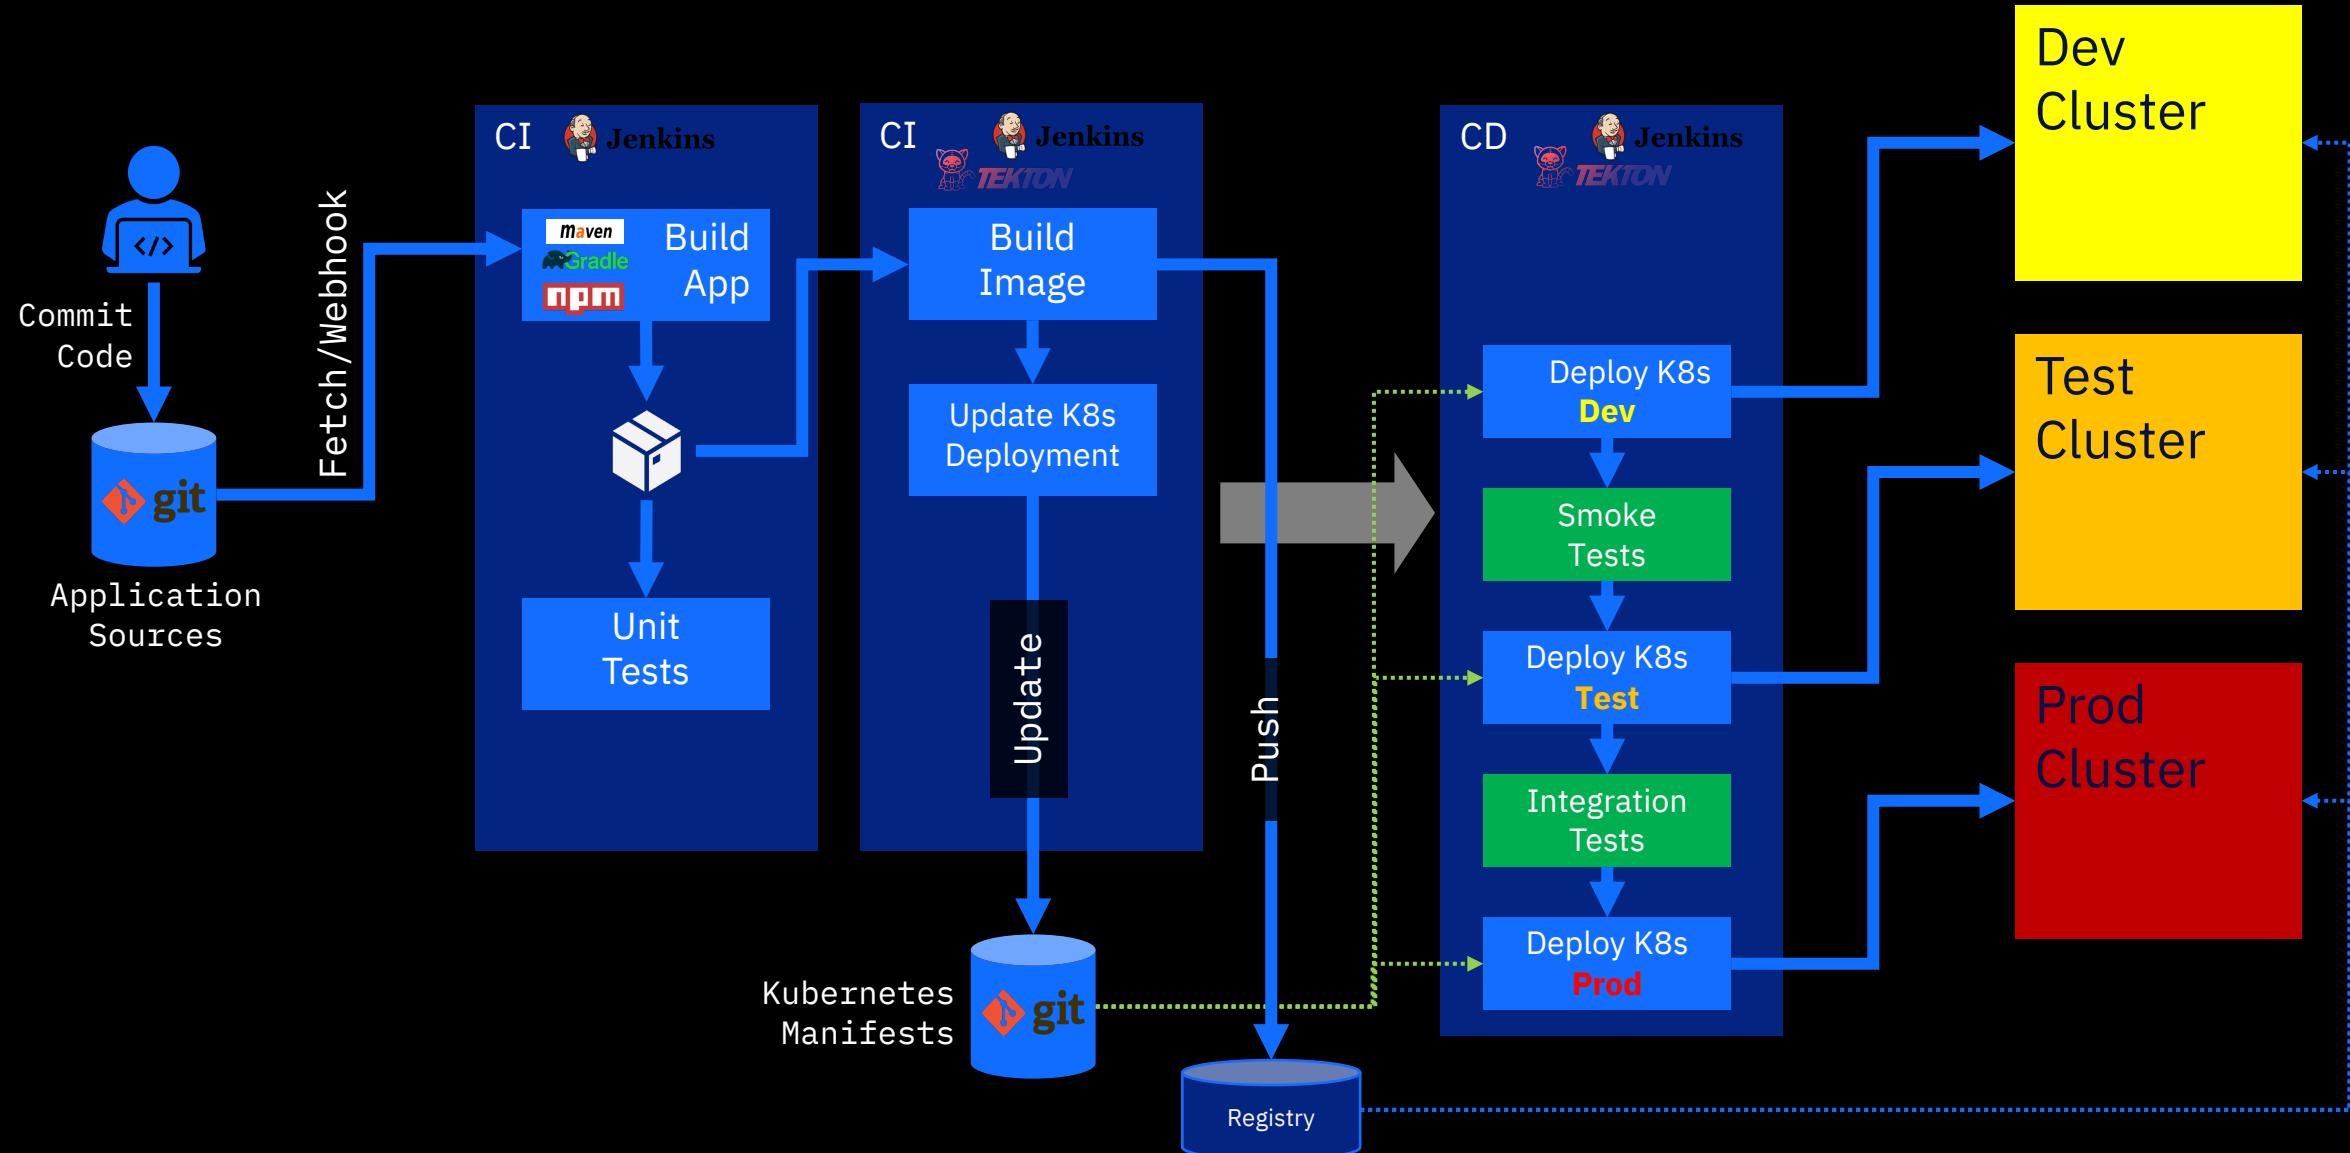
GitOps builds on DevOps with **Git** as a single source of truth for declarative infrastructure and applications - the whole system.

1. Having a completely automated delivery pipeline that can roll out changes to your infrastructure when changes are made to Git.
2. Everything has to be version controlled and stored in a single source of truth from which you can recover.

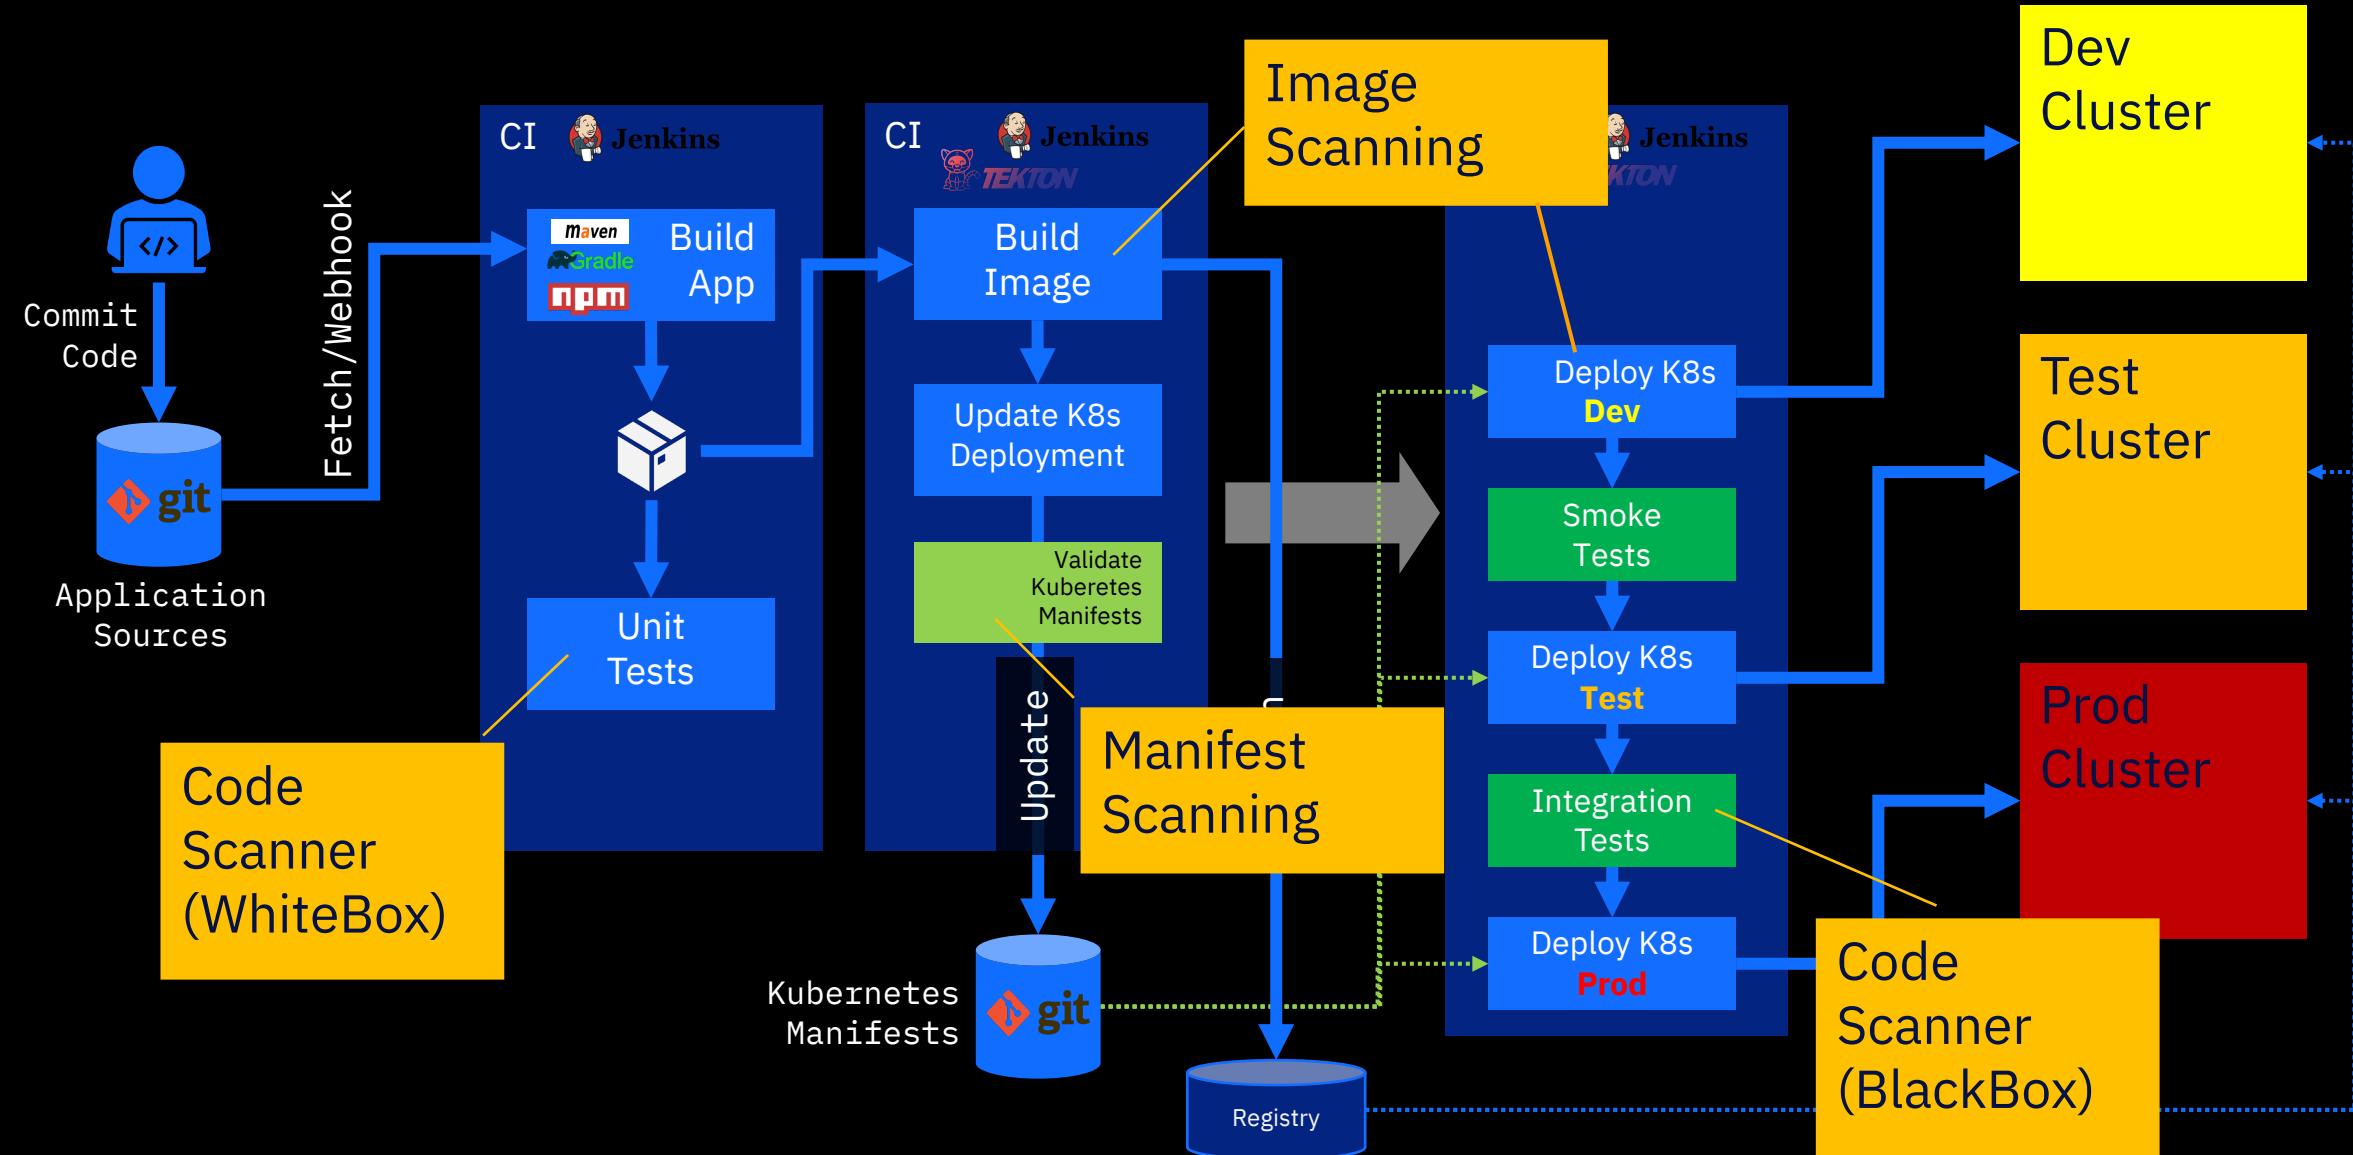
Dev**Sec**Ops means thinking about application and infrastructure security from the start.

1. The idea is to get security back in to the lifecycle → shifting security left
2. The goal is to make security as silent and seamless as possible
3. «Everyone is responsible for security» mindset
 - Developers write secure code
 - SREs secure the environments

Kubernetes – CI/CD Pipeline - GitOps



Kubernetes – CI/CD Pipeline - GitOps



QUESTIONS?



Kubernetes Workshop Series

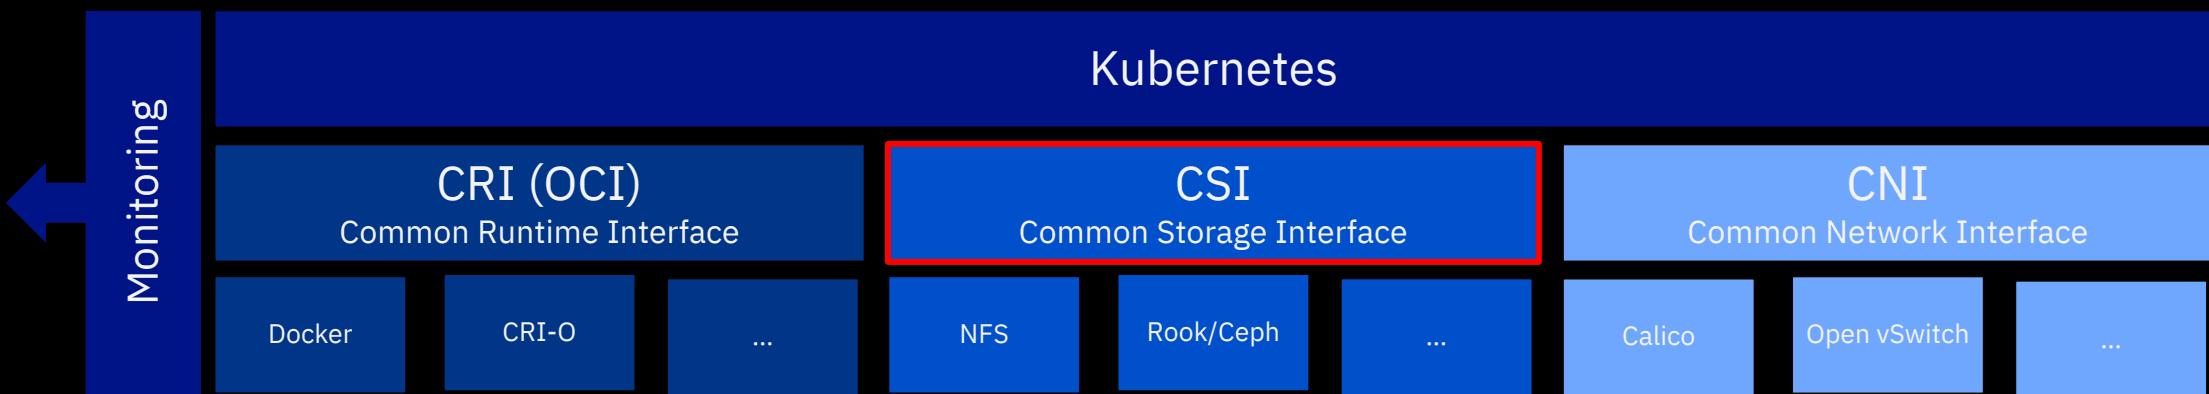
Container State

07



Kubernetes – Storage

The Common Storage Interface (CSI) provides an API for users and administrators that **abstracts details of how storage is provided from how it is consumed**.

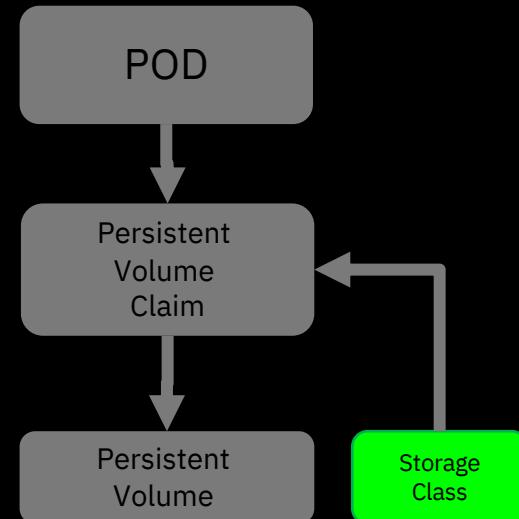


Kubernetes – PersistentVolume subsystem

Storage Class

Provides a way for administrators to describe the “classes” of storage they offer.

Different classes might have different quality-of-service levels, backup policies, ...



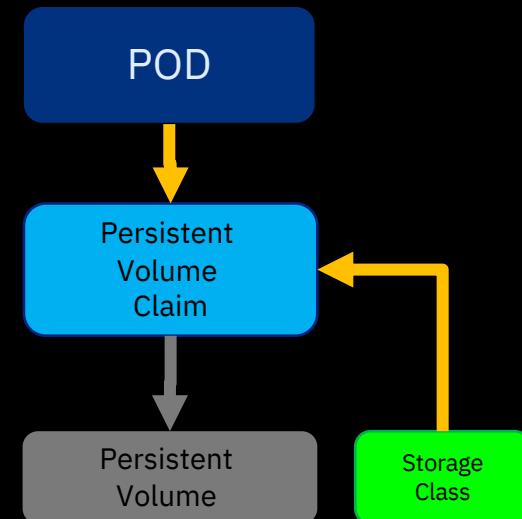
```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: my-storage-class
  annotations:
    storageclass.beta.kubernetes.io/is-default-class: "true"
  labels:
    addonmanager.kubernetes.io/mode: Reconcile
provisioner: k8s.io/minikube-hostpath
```

Kubernetes – PersistentVolume subsystem

PersistentVolumeClaim (PVC)

Is a **request for storage** by a user.

- Pods attach to PVCs to **persist** data.
- PVCs can request specific **size and access modes** (e.g., they can be mounted once read/write or many times read-only).



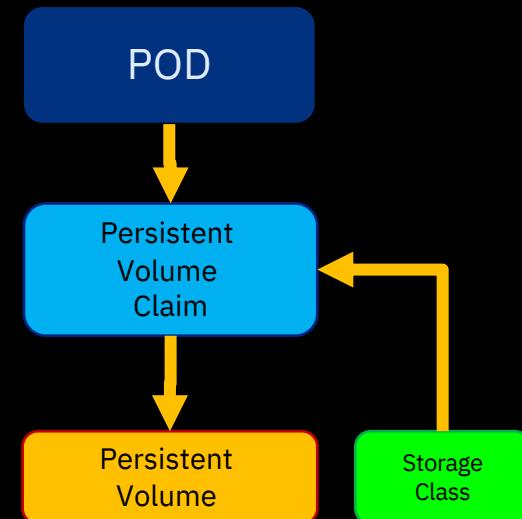
```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-pv-claim
spec:
  storageClassName: my-storage-class
  accessModes:
  - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
```

Kubernetes – PersistentVolume subsystem

PersistentVolume (PV)

Is a **piece of storage** in the cluster that has been **provisioned** by an administrator or dynamically provisioned using Storage Classes

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-pv-volume
spec:
  capacity:
    storage: 1Gi
  accessModes:
  - ReadWriteOnce
  hostPath:
    path: "/mnt/data"
```



Usually created dynamically except for NFS, hostPath and iSCSI.

There is a dynamic provisioner for NFS:

<https://github.com/kubernetes-sigs/nfs-subdir-external-provisioner>

Kubernetes – PersistentVolume subsystem

Provisioning Persistent Volumes

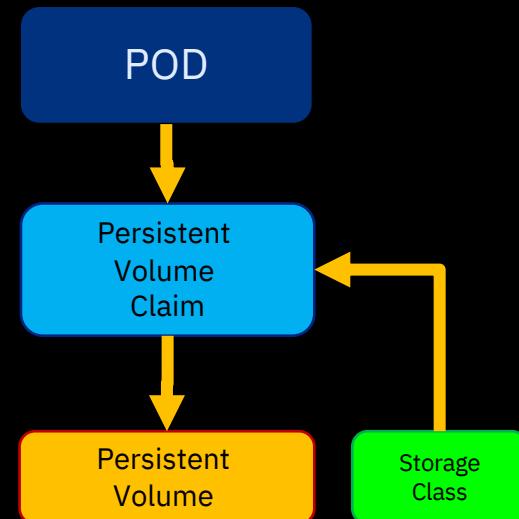
There are two ways PVs can be provisioned:

- **Static**

A cluster administrator manually creates a number of PVs. They carry the details of the real storage, which is available for use by cluster users.

- **Dynamic**

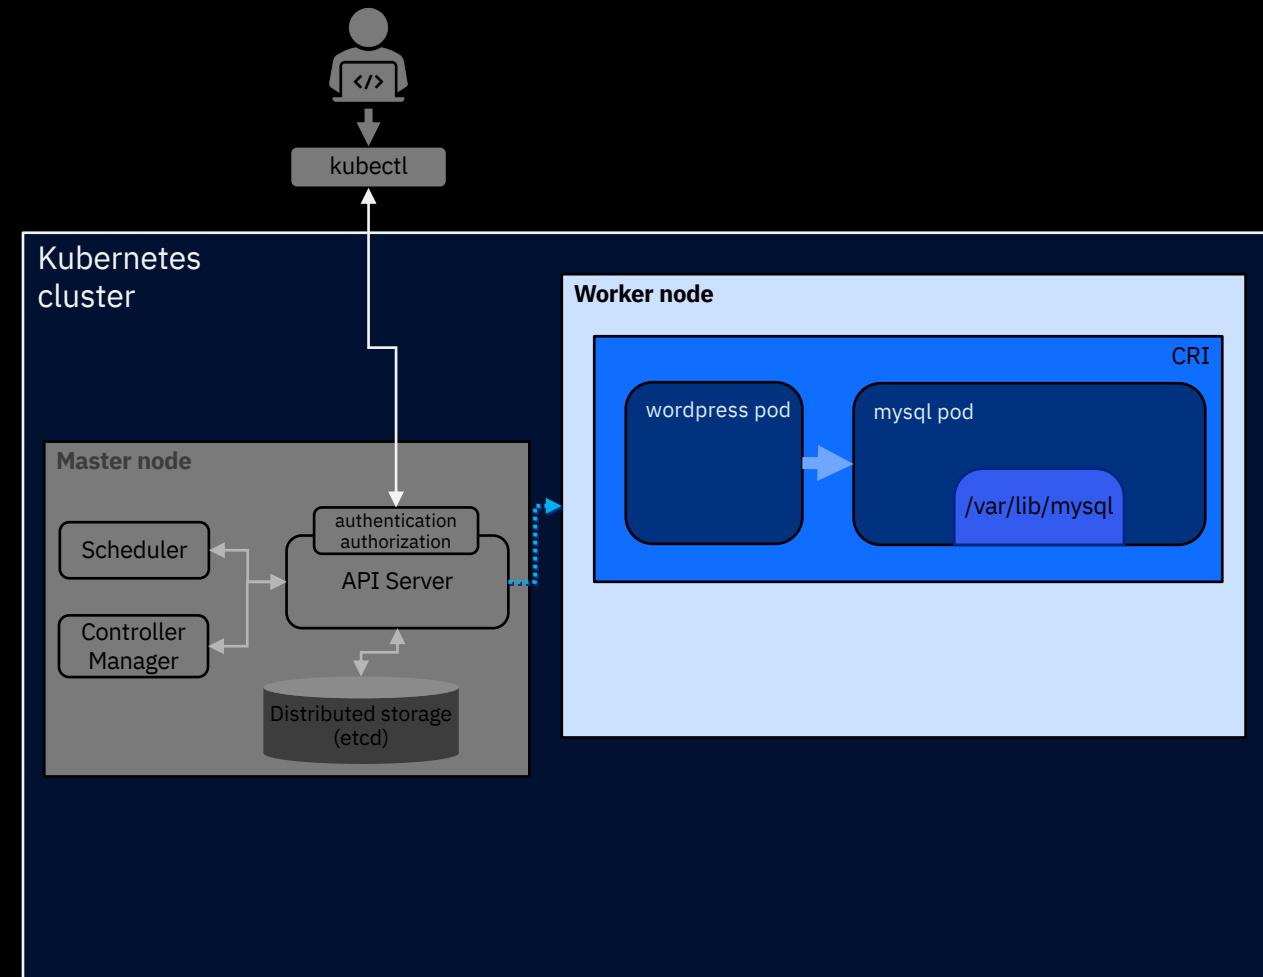
When none of the static PVs the administrator created match a user's PersistentVolumeClaim, the cluster may try to dynamically provision a volume specially for the PVC if there is a storage solution that supports it.



Persistence - Example

I have a [Wordpress](#) server with a [mysql](#) backend.

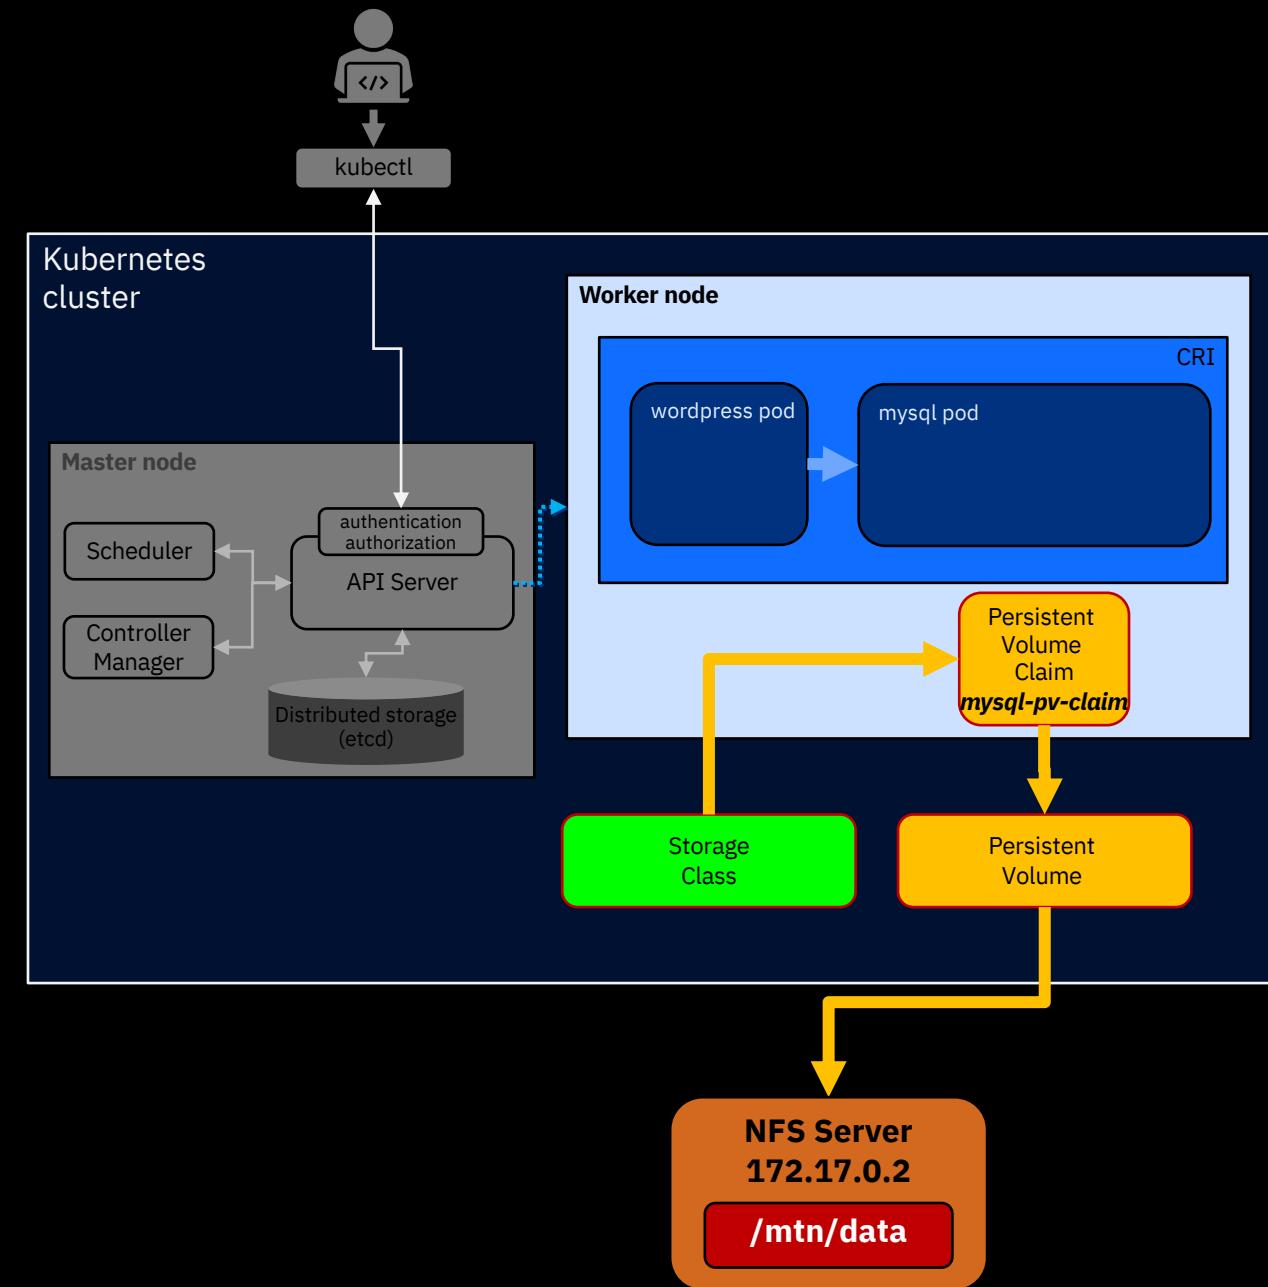
Each time the mysql pod restarts I lose all my data.



Persistence - Example

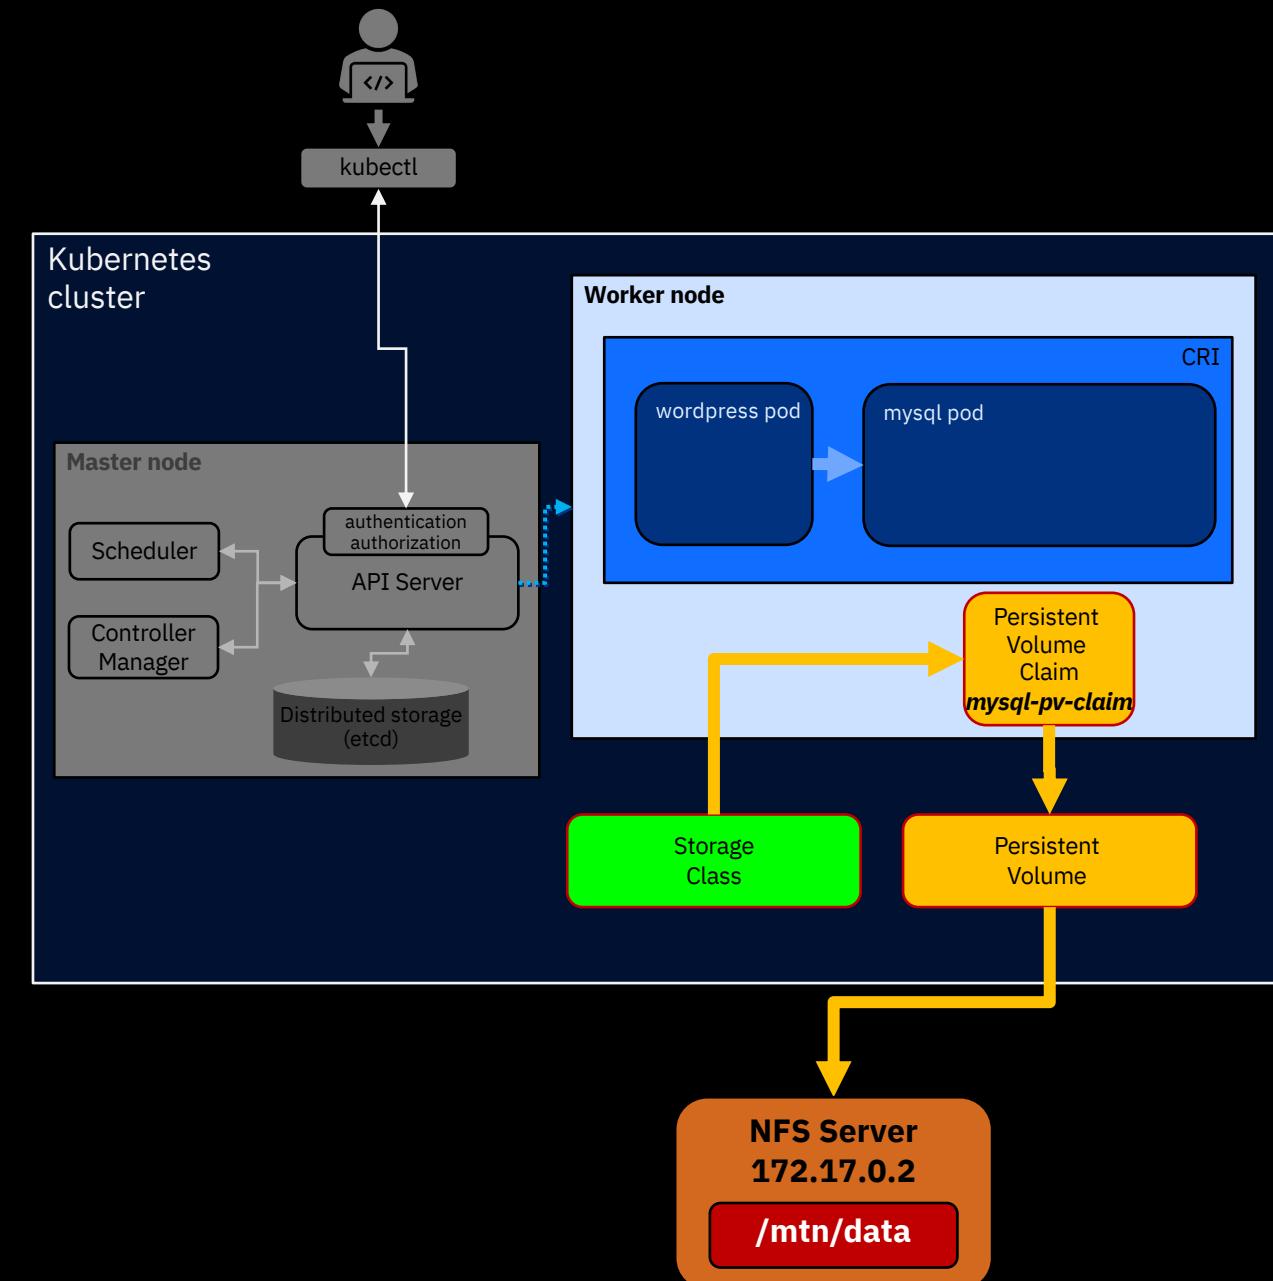
```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-pv-volume
spec:
  capacity:
    storage: 20Gi
  accessModes:
  - ReadWriteOnce
  nfs:
    path: "/mnt/data"
    server: 172.17.0.2
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mysql-pv-claim
  labels:
    app: wordpress
spec:
  storageClassName: nfs-client
  accessModes:
  - ReadWriteOnce
  resources:
    requests:
      storage: 20Gi
```



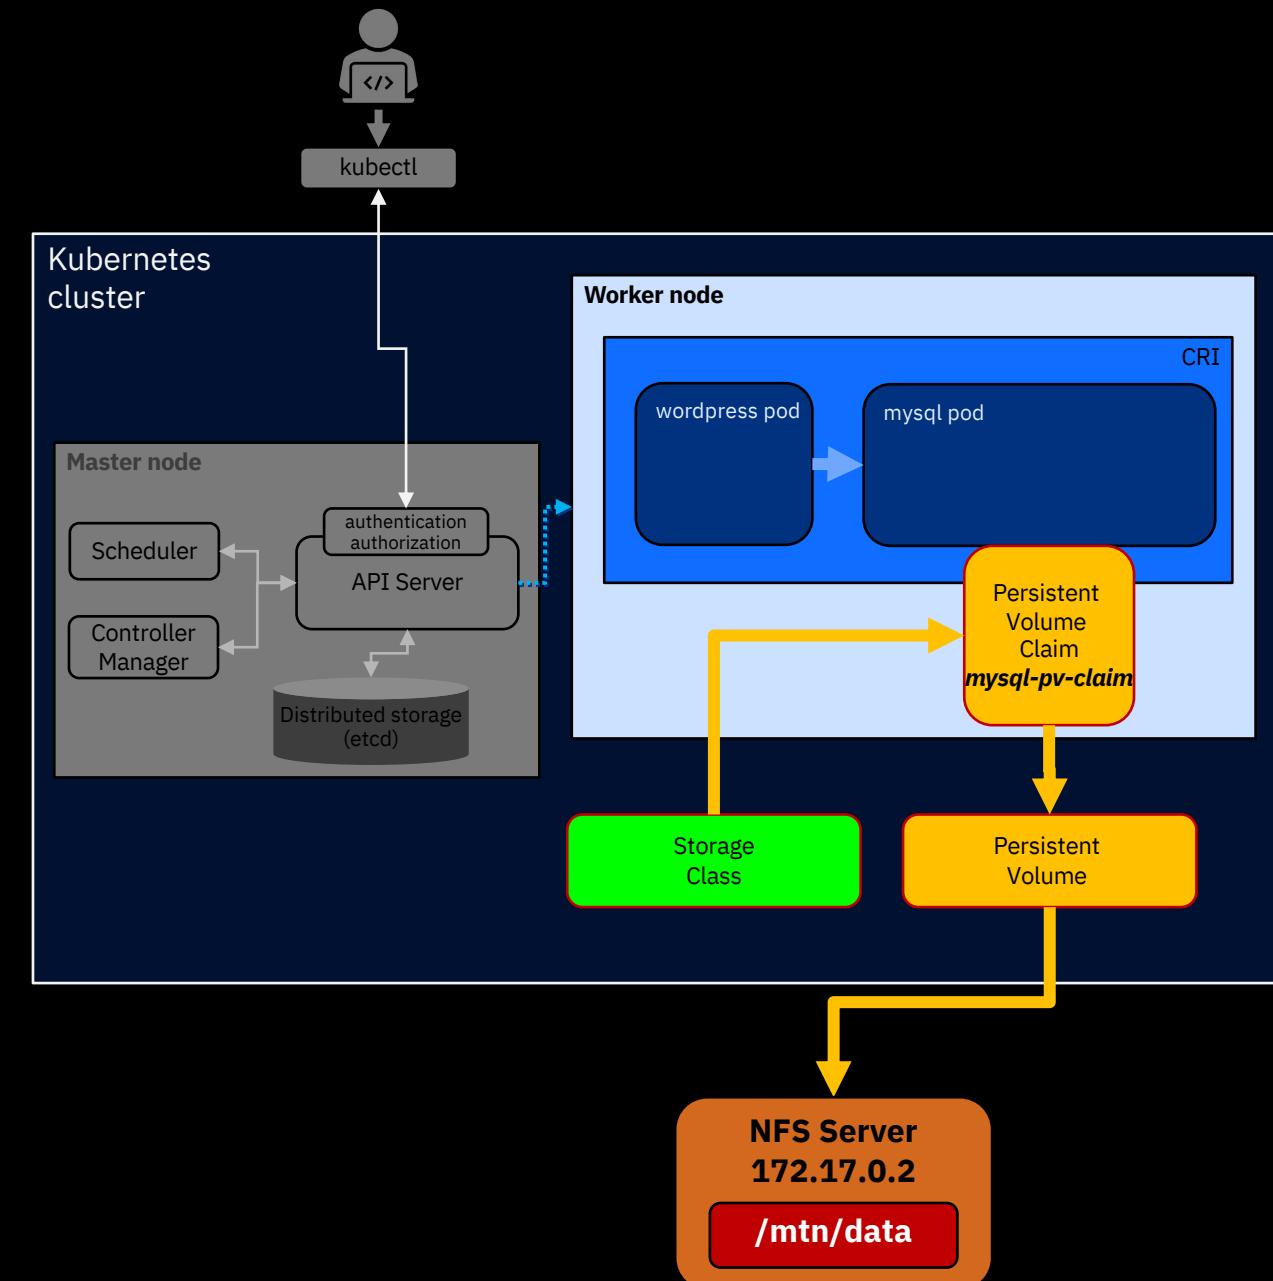
Persistence - Example

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: wordpress-mysql
spec:
  template:
    spec:
      containers:
        - image: mysql:5.6
          name: mysql
      volumeMounts:
        - name: mysql-persistent-storage
          mountPath: /var/lib/mysql
      volumes:
        - name: mysql-persistent-storage
          persistentVolumeClaim:
            claimName: mysql-pv-claim
```



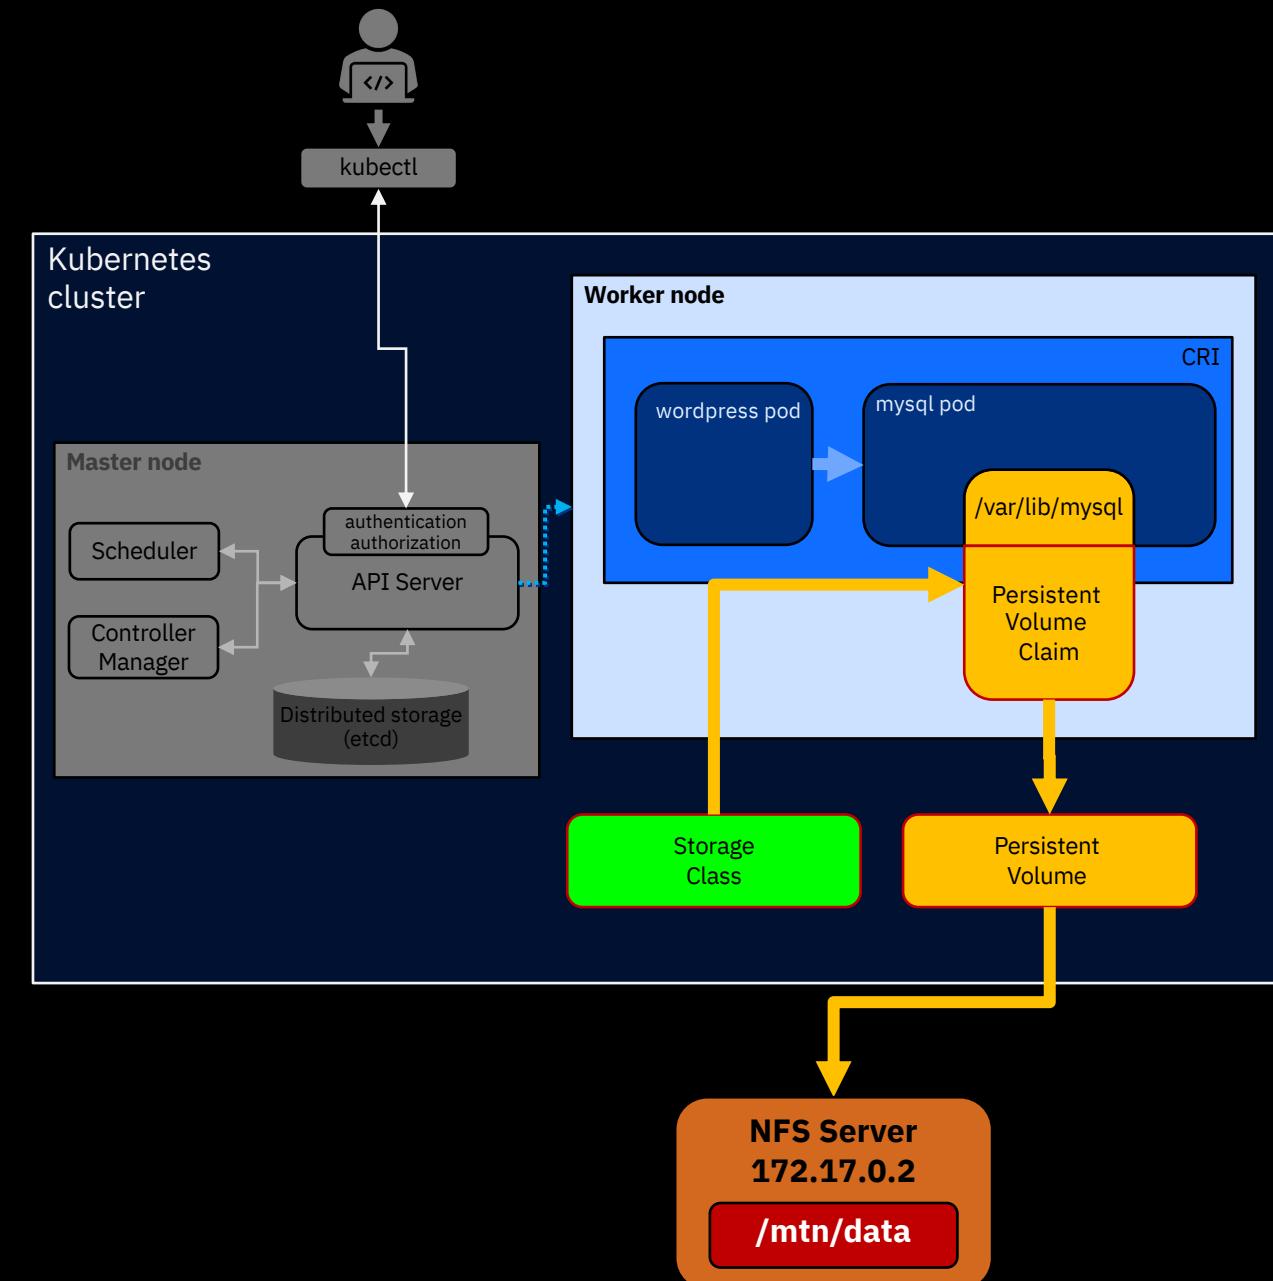
Persistence - Example

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: wordpress-mysql
spec:
  template:
    spec:
      containers:
        - image: mysql:5.6
          name: mysql
      volumeMounts:
        - name: mysql-persistent-storage
          mountPath: /var/lib/mysql
      volumes:
        - name: mysql-persistent-storage
          persistentVolumeClaim:
            claimName: mysql-pv-claim
```



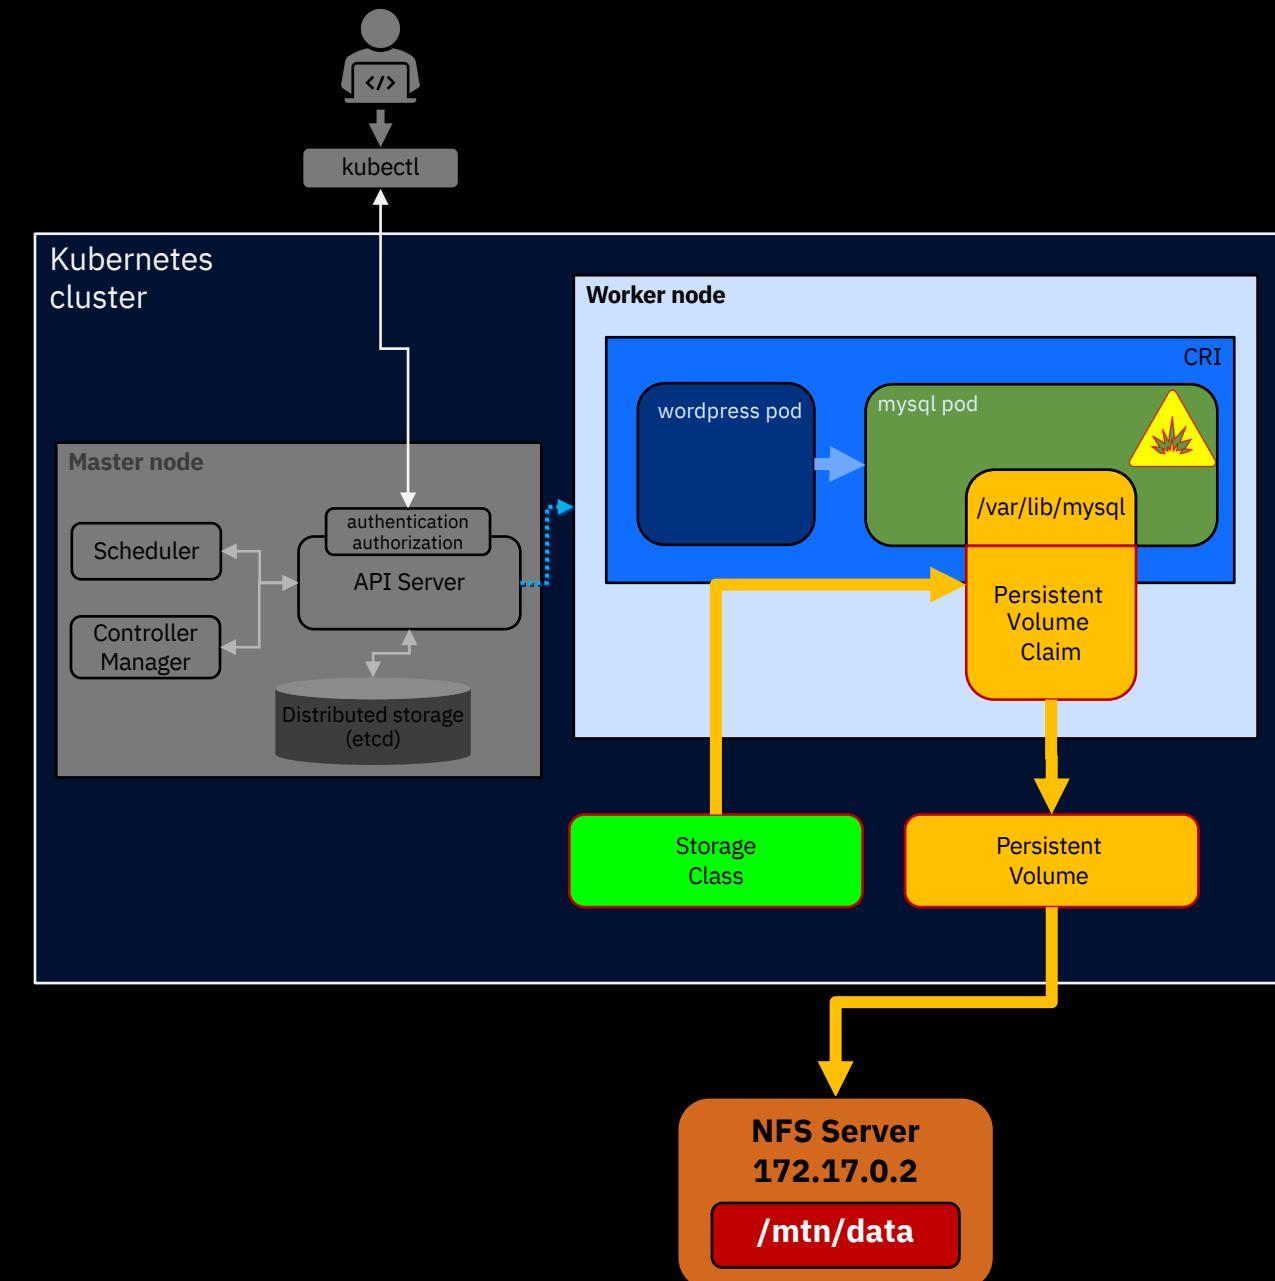
Persistence - Example

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: wordpress-mysql
spec:
  template:
    spec:
      containers:
        - image: mysql:5.6
          name: mysql
      volumeMounts:
        - name: mysql-persistent-storage
          mountPath: /var/lib/mysql
      volumes:
        - name: mysql-persistent-storage
          persistentVolumeClaim:
            claimName: mysql-pv-claim
```



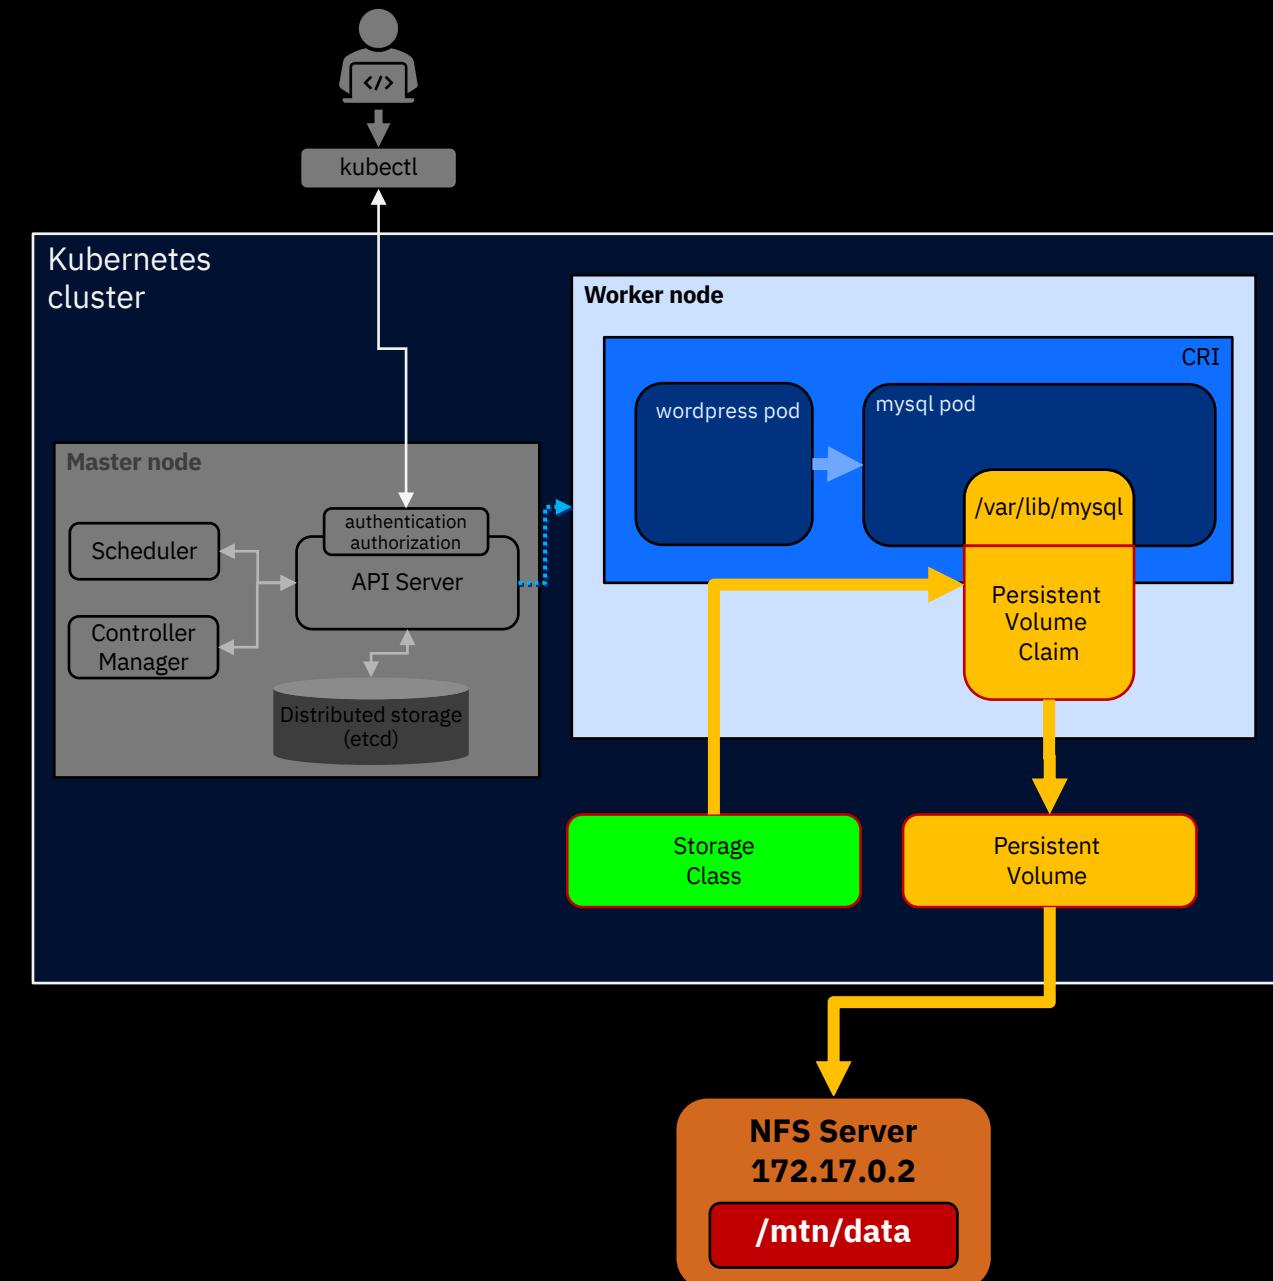
Persistence - Example

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: wordpress-mysql
spec:
  template:
    spec:
      containers:
        - image: mysql:5.6
          name: mysql
      volumeMounts:
        - name: mysql-persistent-storage
          mountPath: /var/lib/mysql
      volumes:
        - name: mysql-persistent-storage
          persistentVolumeClaim:
            claimName: mysql-pv-claim
```



Persistence - Example

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: wordpress-mysql
spec:
  template:
    spec:
      containers:
        - image: mysql:5.6
          name: mysql
      volumeMounts:
        - name: mysql-persistent-storage
          mountPath: /var/lib/mysql
      volumes:
        - name: mysql-persistent-storage
          persistentVolumeClaim:
            claimName: mysql-pv-claim
```



Kubernetes – PersistentVolume subsystem

Access Modes

ReadWriteOnce

The volume can be mounted as read-write by a single node

ReadOnlyMany

The volume can be mounted read-only by many nodes

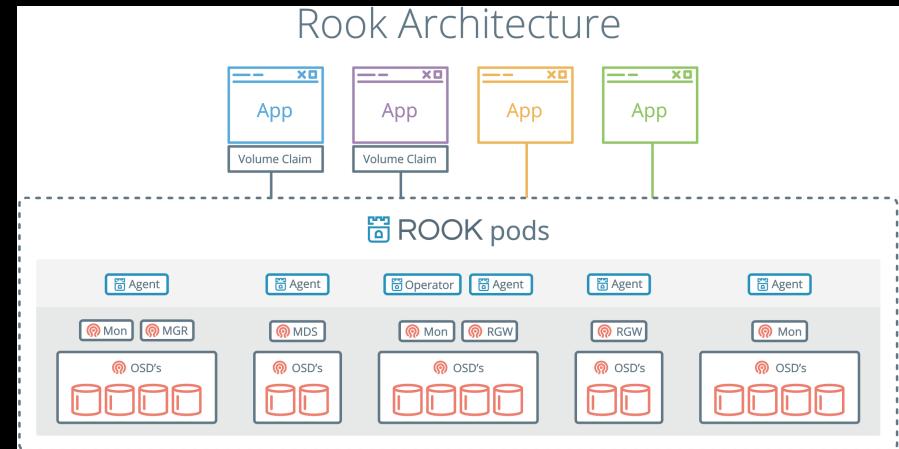
ReadWriteMany

The volume can be mounted as read-write by many nodes

Rook/Ceph / ODF

Rook turns distributed storage systems into self-managing, self-scaling, self-healing storage services.

Ceph is a highly scalable distributed storage solution for **block storage**, **object storage**, and **shared file systems** with years of production deployments.

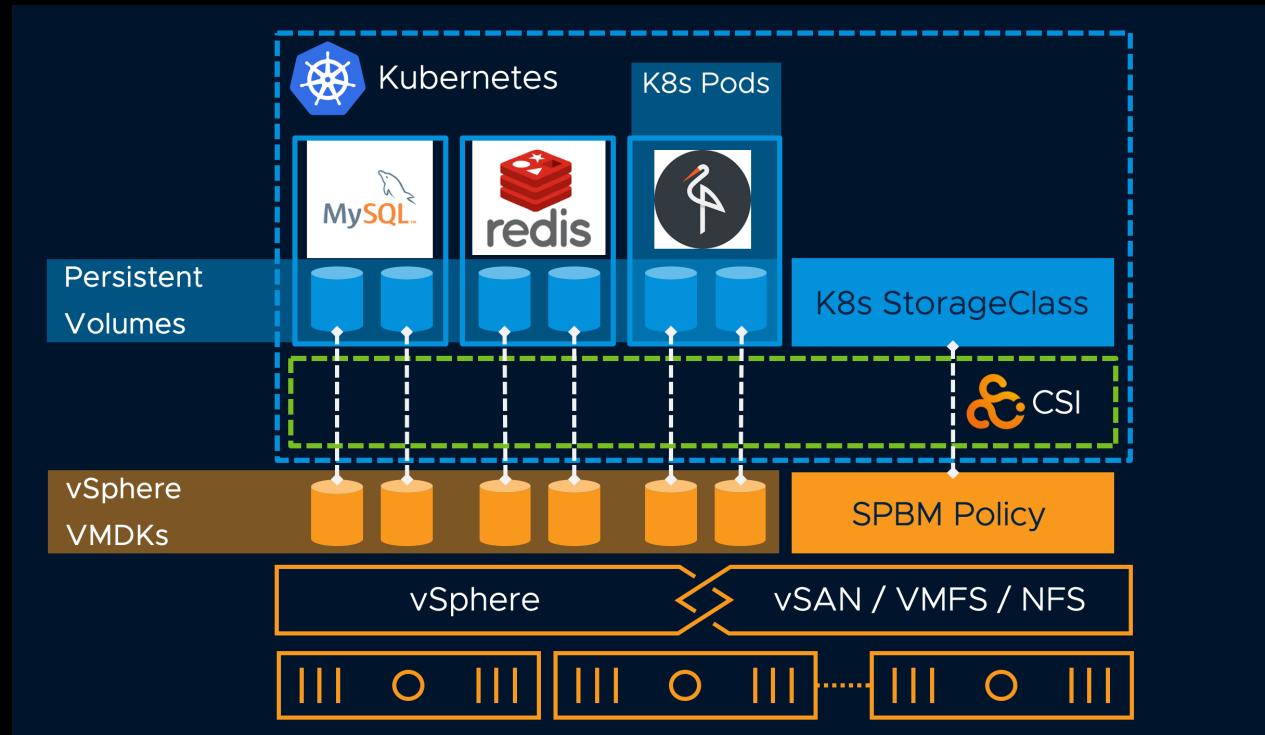


cassandra	crds: Set annotations pods, deployments and so on
ceph	We were defaulting to `ext4` at first and then moved to
cockroachdb	crds: Set annotations pods, deployments and so on
edgefs	edgefs image version update in examples and docs
minio	minio: add necessary update verb in minio RBAC
nfs	NFS: Fix operator.yaml line endings
noobaa	Rook-NooBaa Design Doc
yugabytedb	yugabytedb: Documentation, user guides & examples

vSphere Container Storage



Cloud Native Storage (CNS) is a vSphere and Kubernetes (K8s) feature that makes K8s aware of how to provision storage on vSphere on-demand

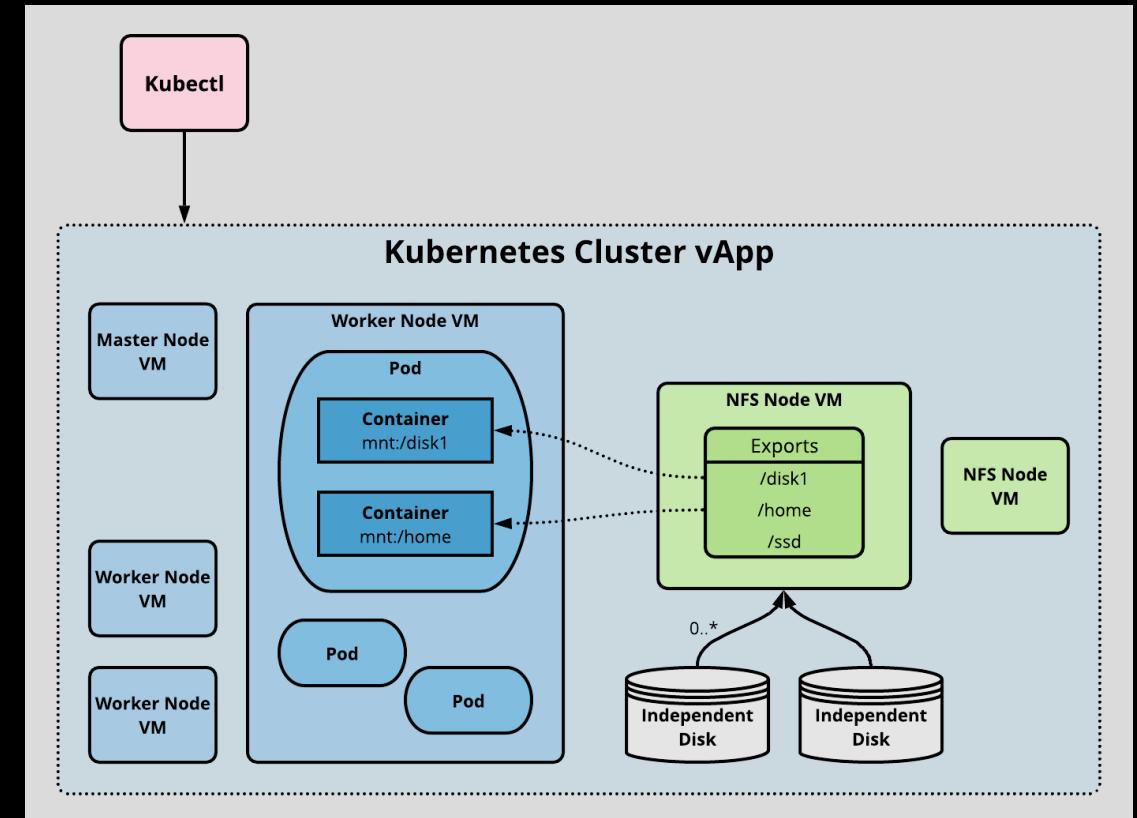


<https://blogs.vmware.com/virtualblocks/2019/08/14/introducing-cloud-native-storage-for-vsphere/>

NFS

NFS stands for Network File System – it's a shared filesystem that can be accessed over the network.

- The NFS server must already exist – Kubernetes doesn't run the NFS, pods in just access it.
- NFS provisioning is not dynamic by default: You have to create the PV manually and then the PVC.



QUESTIONS?



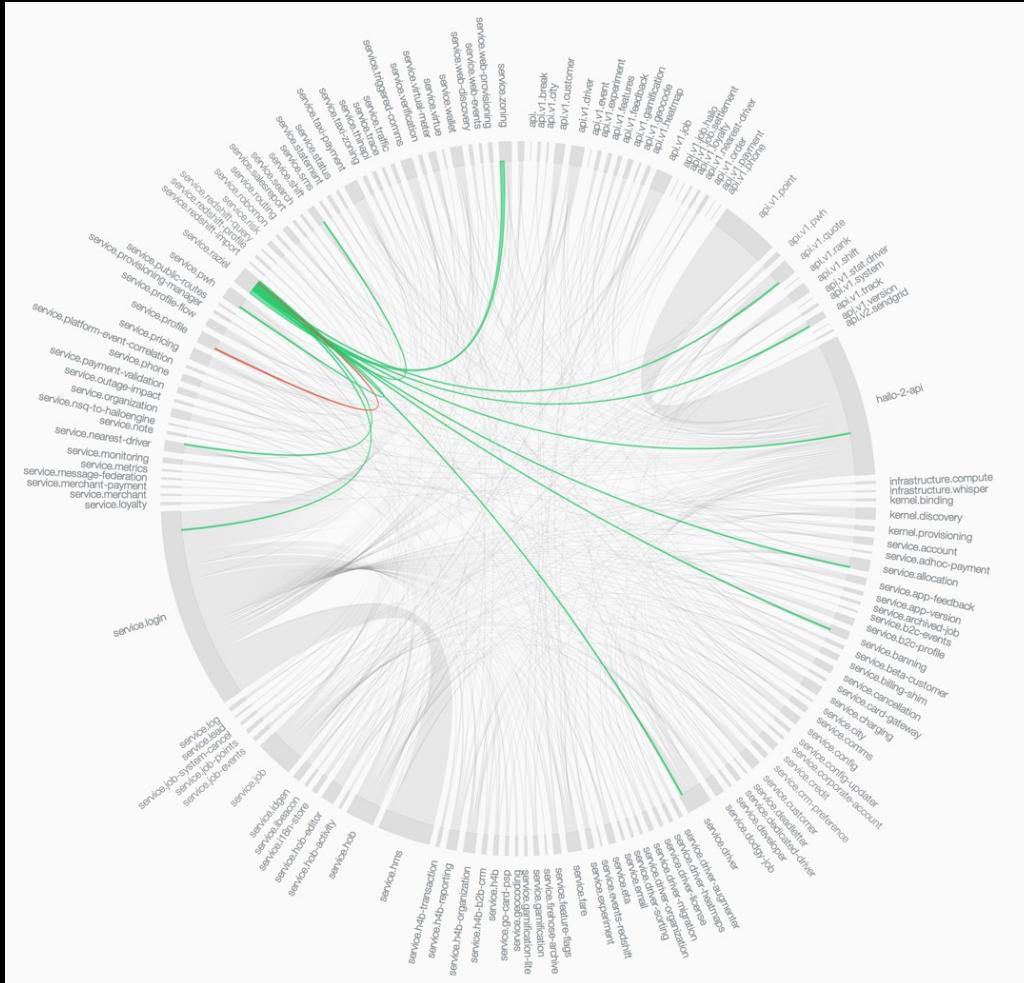
Kubernetes Workshop Series

Mesh Networking



The trade off

Improved delivery velocity
in exchange for
increased operational complexity



Addressing DevOps Challenges



#	CHALLENGE
CHALLENGE 1	ROLL OUT NEW VERSION WITHOUT DOWNTIME OR CHANGING CODE
CHALLENGE 2	HOW TO DO CANARY TESTING
CHALLENGE 3	HOW TO DO A/B TESTING
CHALLENGE 4	THINGS DON'T ALWAYS GO CORRECTLY IN PRODUCTION...
CHALLENGE 5	HOW CAN I LIMIT RATE FOR SOME OF MY SERVICES?
CHALLENGE 6	I NEED TO VIEW AND MONITOR WHAT IS GOING ON WHEN CRISIS ARISES
CHALLENGE 7	HOW CAN I SECURE MY SERVICES?

Service Mesh

**Dedicated infrastructure layer
to make
service-to-service communication
fast, safe and reliable**

Istio



A service mesh designed to connect, manage and secure micro services

= Forbes

3,859 views | May 25, 2017, 01:39am

Google, IBM And Lyft Want To Simplify Microservices Management With Istio

ZDNet

EDITION: EU

SCANDINAVIA AFRICA UK ITALY SPAIN MORE NEWSLETTERS ALL WRITERS

MUST READ: Meet the new Microsoft Phone, powered by Android (No Windows required)

Google, IBM, and Lyft launch open source project Istio

Istio gives developers a vendor-neutral way to connect, secure, manage, and monitor networks of different microservices on cloud platforms.

Google Cloud

Blog Latest Stories Product News Topics

GOOGLE CLOUD PLATFORM

Istio: a modern approach to developing and managing microservices

Varun Talwar
Product Manager, Cloud Service Platform

May 24, 2017

f t in e

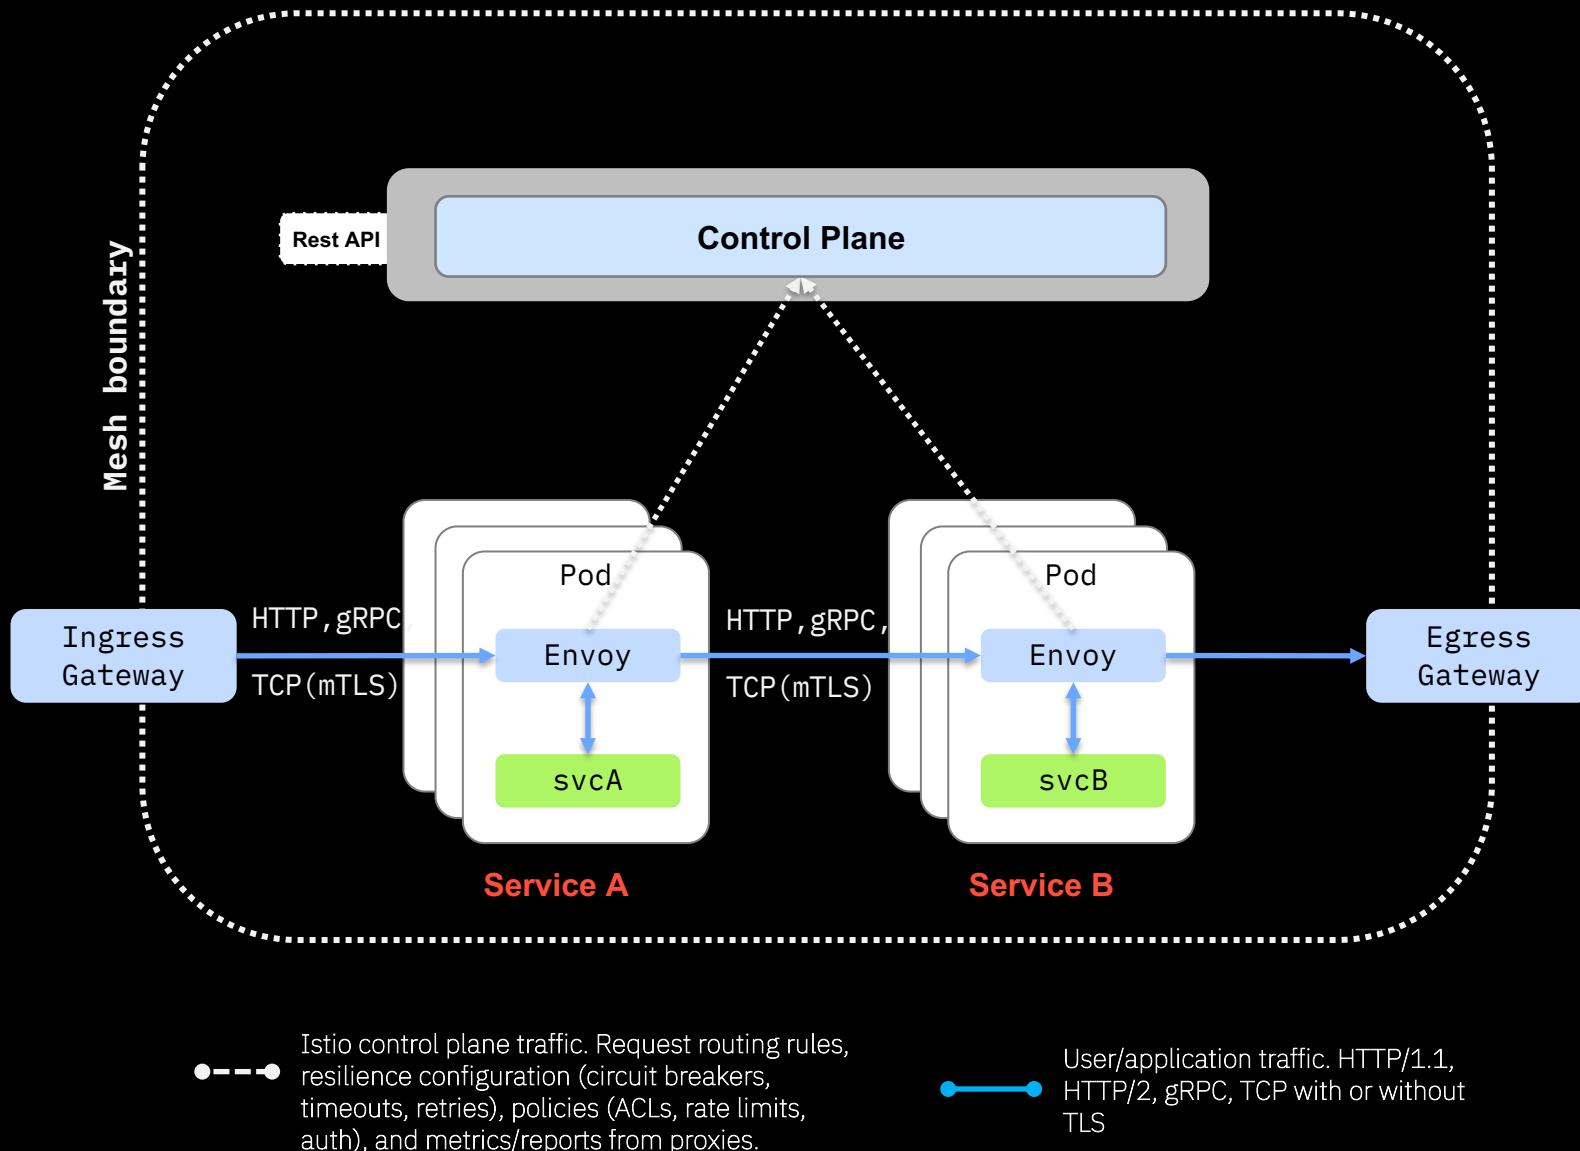
Today Google, IBM and Lyft announced the alpha release of Istio: a new open-source project that provides a uniform way to help connect, secure, manage and monitor microservices.

Launched in May 2017 by Google, Lyft and IBM

Open Source

Zero Code Changes

ISTIO - Architecture



ISTIO Gateway



Load balancer **operating at the edge of the mesh** receiving incoming HTTP/TCP connections

- Configures ports to expose externally
- Maps each exposed port to a request destination
- Each gateway can have one or more Virtual Services that defines these request destinations

Gateway is **attached to Istio Ingress Controller** (which can be the Kubernetes Ingress Controller)

Request destinations

- Ports for the gateway to expose and hosts for the corresponding services
- Attributes: **servers**

```
kind: Gateway
metadata:
  name: helloworld-gateway
spec:
  selector:
    istio: ingressgateway
  servers:
    - hosts:
        - myapp.demo.com
      port:
        name: http
        number: 80
        protocol: HTTP
```

ISTIO

Custom resource definitions

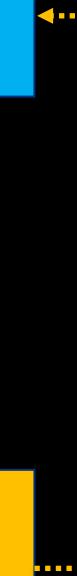
```
kind: Gateway
metadata:
  name: helloworld-gateway
spec:
  selector:
    istio: ingressgateway
  servers:
  - hosts:
    - myapp.demo.com
  port:
    name: http
    number: 80
    protocol: HTTP
```



<http://myapp.demo.com>

POD
helloworld
version = v1

SERVICE
helloworld
selector
app: helloworld



ISTIO

Virtual Service

Request sources

- Hosts that sources can invoke
- Attributes: **hosts** and **gateways**

Route destinations

- Subset of the destination
- Attributes: **route** and **destination**

Protocol selection

- How to connect to the destination subset
- Attributes: **http**, **tcp**, **tls**

Routing rules

- Additional routing attributes, applied for the route destinations
- Attributes: **weight** and **match**

HTTP traffic policy

- Protocol-specific connection quality of service
- Attributes: **timeout**, **retries**, **fault**, **rewrite**, and **redirect**

```
kind: VirtualService
metadata:
  name: helloworld
spec:
  hosts:
    - myapp.demo.com
  gateways:
    - helloworld-gateway
  http:
    - match:
        - uri:
            exact: /demo
      route:
        - destination:
            host: helloworld
```

ISTIO

Custom resource definitions

Ingress Configuration

```
kind: Gateway
metadata:
  name: helloworld-gateway
spec:
  selector:
    istio: ingressgateway
  servers:
    - hosts:
      - myapp.demo.com
    port:
      name: http
      number: 80
      protocol: HTTP
```

URL Routing

```
kind: VirtualService
metadata:
  name: helloworld
spec:
  hosts:
    - myapp.demo.com
  gateways:
    - helloworld-gateway
  http:
    - match:
      - uri:
          exact: /demo
    route:
      - destination:
          host: helloworld
```



<http://myapp.demo.com/demo>

POD
helloworld
version = v1

SERVICE
helloworld
selector
app: helloworld

helloworld.namespace.svc.cluster.local

ISTIO

Destination Rule



Destination host

- Host that route destinations can select
- Attribute: **host**

Host subset

- Identify a subset of service endpoints
- Attribute: **labels**

Traffic policy

- Influence expected quality of service for destinations
- Attributes: **trafficPolicy**
loadBalancer, connectionPool, outlierDetection, and tls

Traffic policy can be applied to a port

- Attribute: **portLevelSettings**

```
kind: DestinationRule
metadata:
  name: helloworld-destination
spec:
  host: helloworld
  subsets:
    - name: v1
      labels:
        version: v1
    - name: v2
      labels:
        version: v2
  trafficPolicy:
    tls:
      mode: SIMPLE
      connectionPool:
        tcp:
          maxConnections: 100
```

ISTIO

Custom resource definitions

Ingress Configuration

```
kind: Gateway
metadata:
  name: helloworld-gateway
spec:
  selector:
    istio: ingressgateway
  servers:
    - hosts:
        - myapp.demo.com
    port:
      name: http
      number: 80
      protocol: HTTP
```



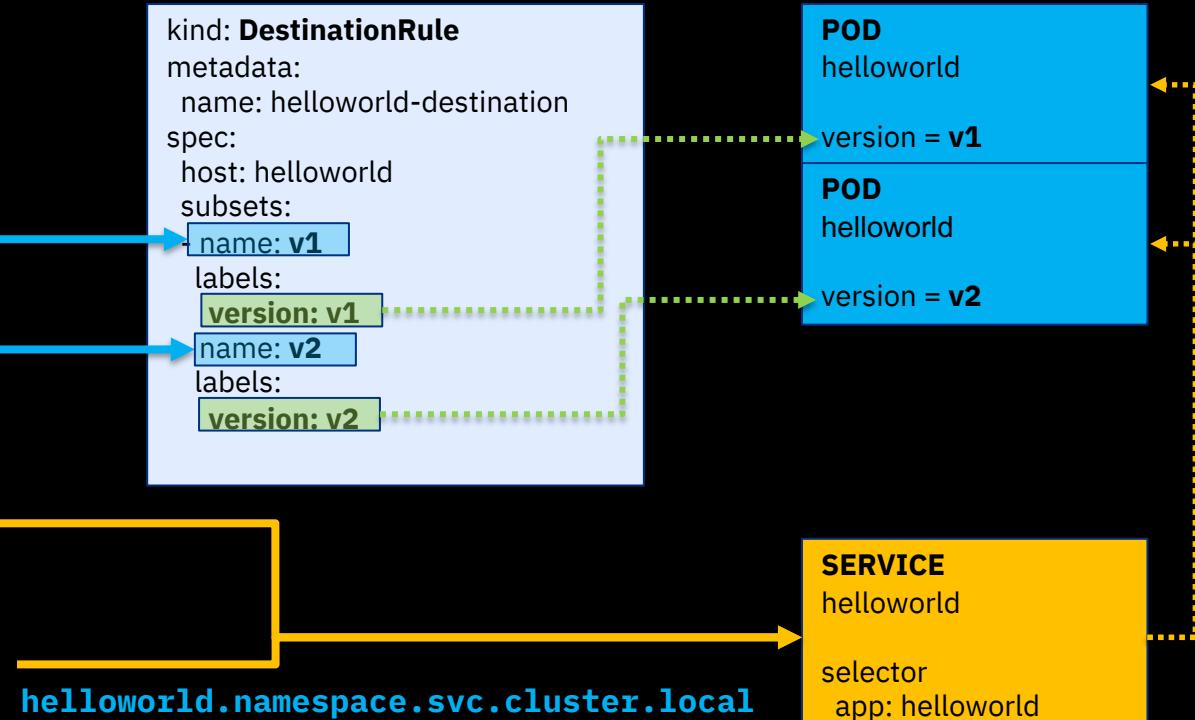
<http://myapp.demo.com/demo>

URL Routing

```
kind: VirtualService
metadata:
  name: helloworld
spec:
  hosts:
    - myapp.demo.com
  gateways:
    - helloworld-gateway
  http:
    - match:
        - uri:
            exact: /demo
      route:
        - destination:
            host: helloworld
```

Traffic Splitting

```
kind: DestinationRule
metadata:
  name: helloworld-destination
spec:
  host: helloworld
  subsets:
    - name: v1
      labels:
        version: v1
    - name: v2
      labels:
        version: v2
```



ISTIO

Virtual Service

Request sources

- Hosts that sources can invoke
- Attributes: **hosts** and **gateways**

Route destinations

- Subset of the destination
- Attributes: **route** and **destination**

Protocol selection

- How to connect to the destination subset
- Attributes: **http**, **tcp**, **tls**

Routing rules

- Additional routing attributes, applied for the route destinations
- Attributes: **weight** and **match**

HTTP traffic policy

- Protocol-specific connection quality of service
- Attributes: **timeout**, **retries**, **fault**, **rewrite**, and **redirect**

```

kind: VirtualService
metadata:
  name: helloworld
spec:
  hosts:
    - myapp.demo.com
  gateways:
    - helloworld-gateway
  http:
    - match:
        - uri:
            exact: /demo
      route:
        - destination:
            host: helloworld
            subset: v1
            weight: 90
        - destination:
            host: helloworld
            subset: v2
            weight: 10
  
```

ISTIO

Custom resource definitions

Ingress Configuration

```
kind: Gateway
metadata:
  name: helloworld-gateway
spec:
  selector:
    istio: ingressgateway
  servers:
    - hosts:
        - myapp.demo.com
      port:
        name: http
        number: 80
        protocol: HTTP
```



<http://myapp.demo.com/demo>

URL Routing

```
kind: VirtualService
metadata:
  name: helloworld
spec:
  hosts:
    - myapp.demo.com
  gateways:
    - helloworld-gateway
  http:
    - match:
        - uri:
            exact: /demo
      route:
        - destination:
            host: helloworld
            subset: v1
            weight: 90
        - destination:
            host: helloworld
            subset: v2
            weight: 10
```

Traffic Splitting

```
kind: DestinationRule
metadata:
  name: helloworld-destination
spec:
  host: helloworld
  subsets:
    - name: v1
      labels:
        version: v1
    - name: v2
      labels:
        version: v2
```

helloworld.namespace.svc.cluster.local



ISTIO

Custom resource definitions

Ingress Configuration

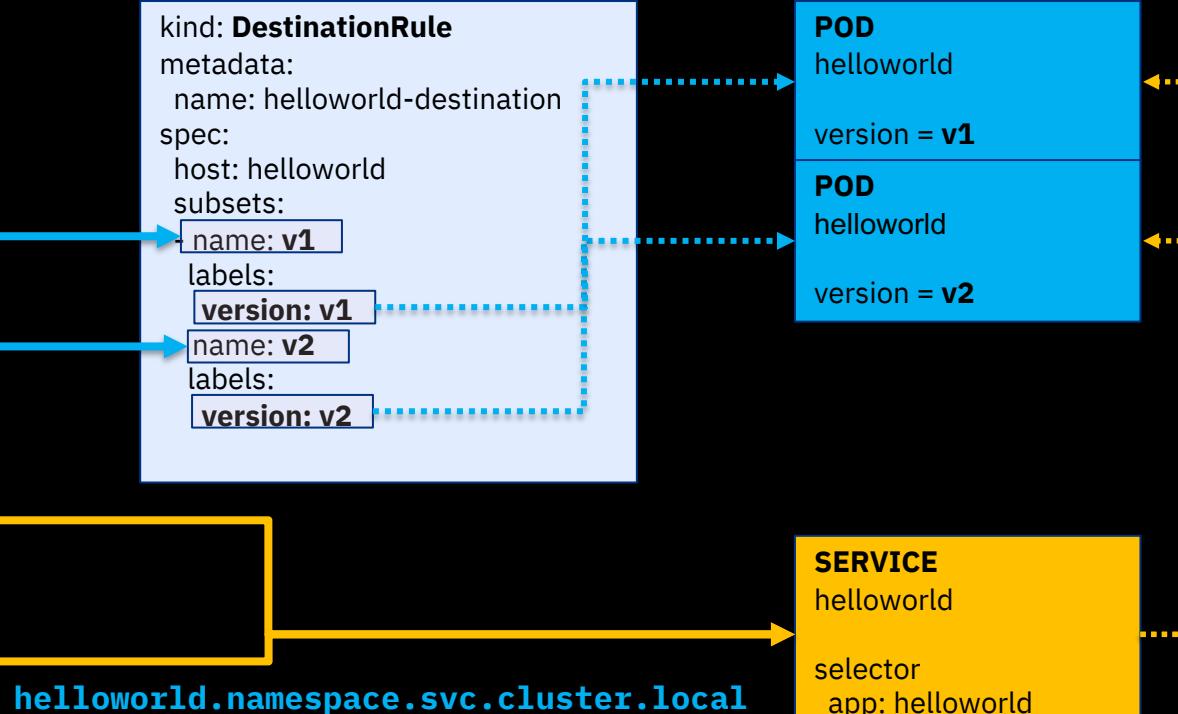
```
kind: Gateway
metadata:
  name: helloworld-gateway
spec:
  selector:
    istio: ingressgateway
  servers:
    - hosts:
        - myapp.demo.com
      port:
        name: http
        number: 80
        protocol: HTTP
```

URL Routing

```
kind: VirtualService
metadata:
  name: helloworld
spec:
  hosts:
    - myapp.demo.com
  gateways:
    - helloworld-gateway
  http:
    - match:
        - uri:
            exact: /
      route:
        - destination:
            host: reviews
            subset: v1
            weight: 100
        - destination:
            host: reviews
            subset: v2
            weight: 100
    - match:
        - headers:
            end-user:
              exact: jason
      route:
        - destination:
            host: reviews
            subset: v1
        - destination:
            host: reviews
            subset: v2
```

Traffic Splitting

```
kind: DestinationRule
metadata:
  name: helloworld-destination
spec:
  host: helloworld
  subsets:
    - name: v1
      labels:
        version: v1
    - name: v2
      labels:
        version: v2
```

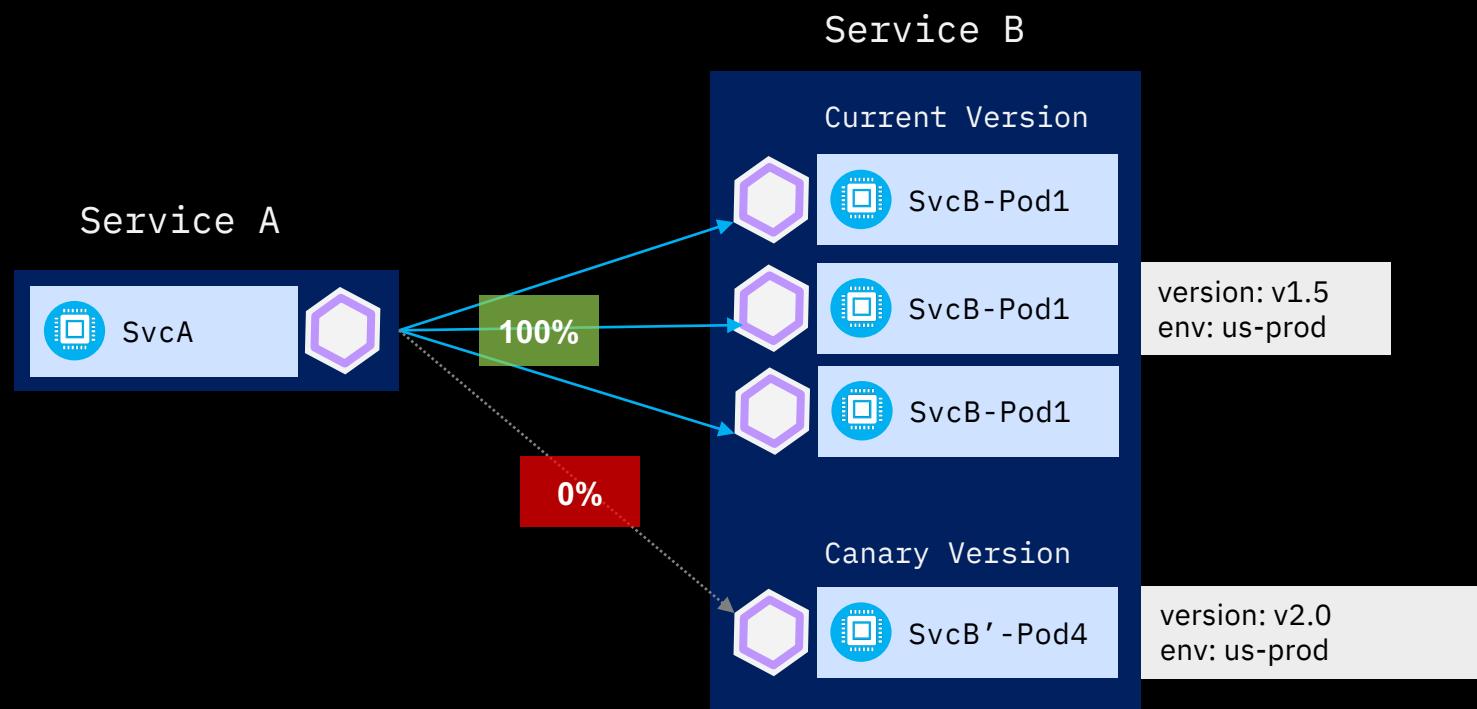


Addressing DevOps Challenges



#	CHALLENGE	ISTIO SOLUTION
CHALLENGE 1	ROLL OUT NEW VERSION WITHOUT DOWNTIME OR CHANGING CODE	TRAFFIC CONTROL
CHALLENGE 2	HOW TO DO CANARY TESTING	TRAFFIC SPLITTING
CHALLENGE 3	HOW TO DO A/B TESTING	TRAFFIC STEERING
CHALLENGE 4	THINGS DON'T ALWAYS GO CORRECTLY IN PRODUCTION...	TRAFFIC MIRRORING RESILIENCY RESILIENCY TESTING
CHALLENGE 5	HOW CAN I LIMIT RATE FOR SOME OF MY SERVICES?	RATE LIMITING
CHALLENGE 6	I NEED TO VIEW AND MONITOR WHAT IS GOING ON WHEN CRISIS ARISES	TELEMETRY
CHALLENGE 7	HOW CAN I SECURE MY SERVICES?	AUTHENTICATION AUTHORIZATION CALICO

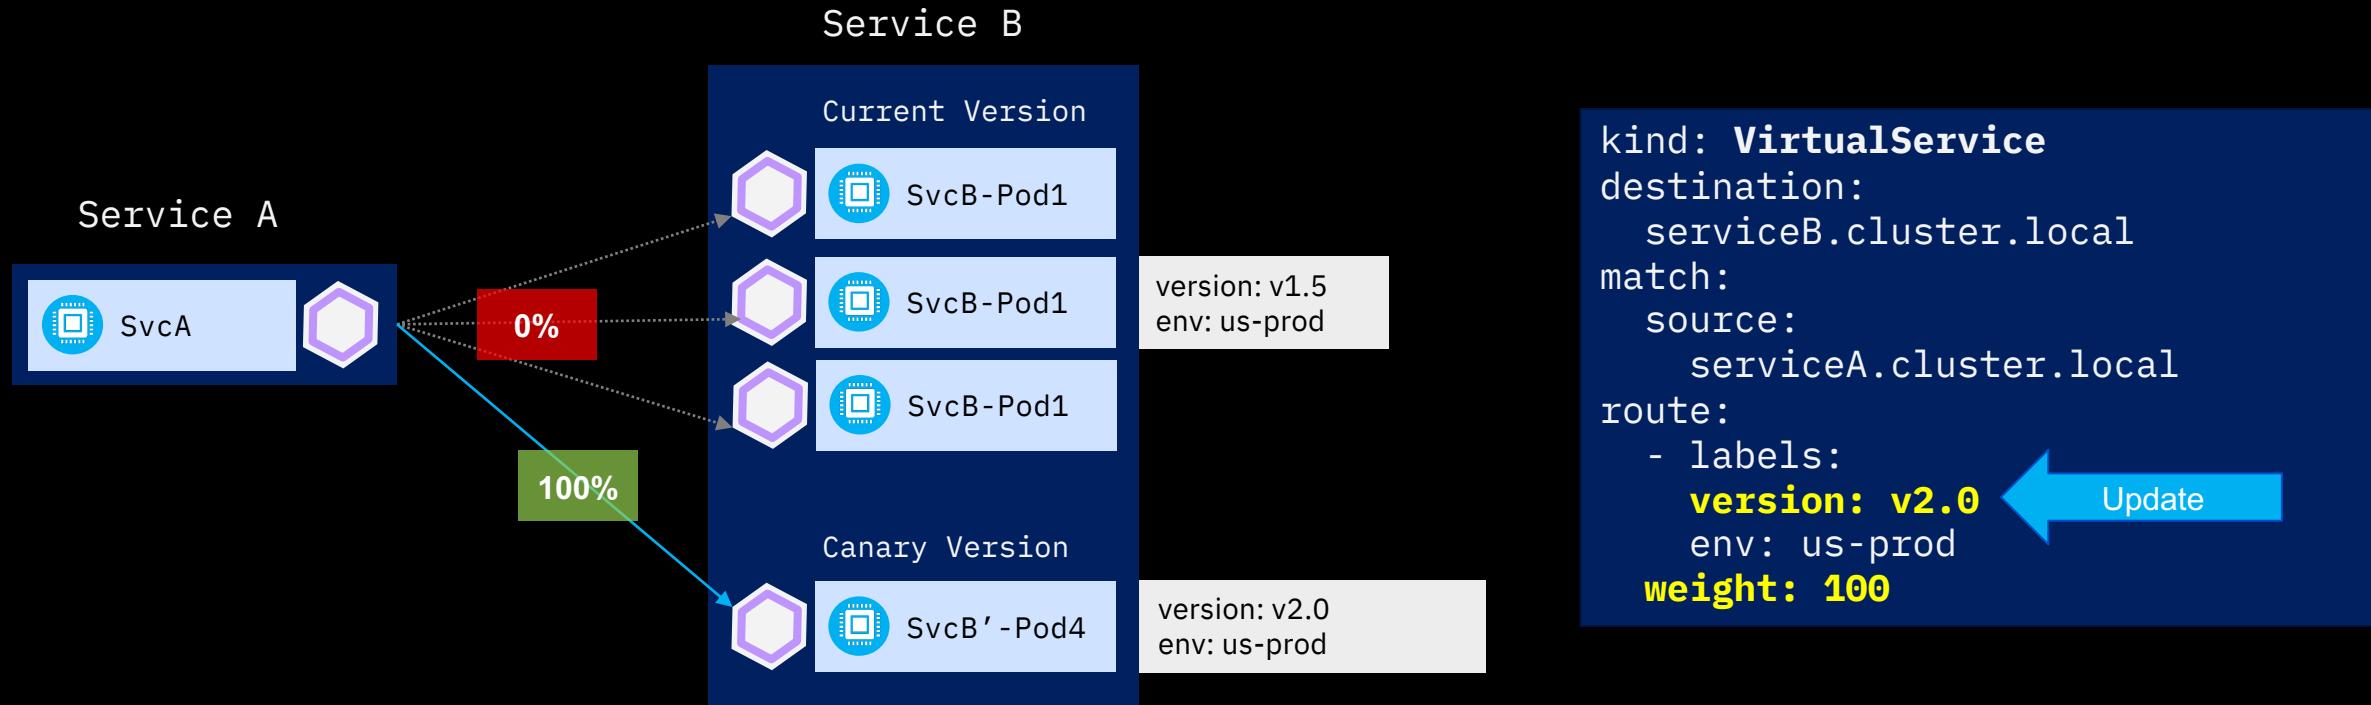
Traffic Control



```
kind: VirtualService
destination:
  serviceB.cluster.local
match:
  source:
    serviceA.cluster.local
route:
  - labels:
      version: v1.5
    env: us-prod
    weight: 100
```

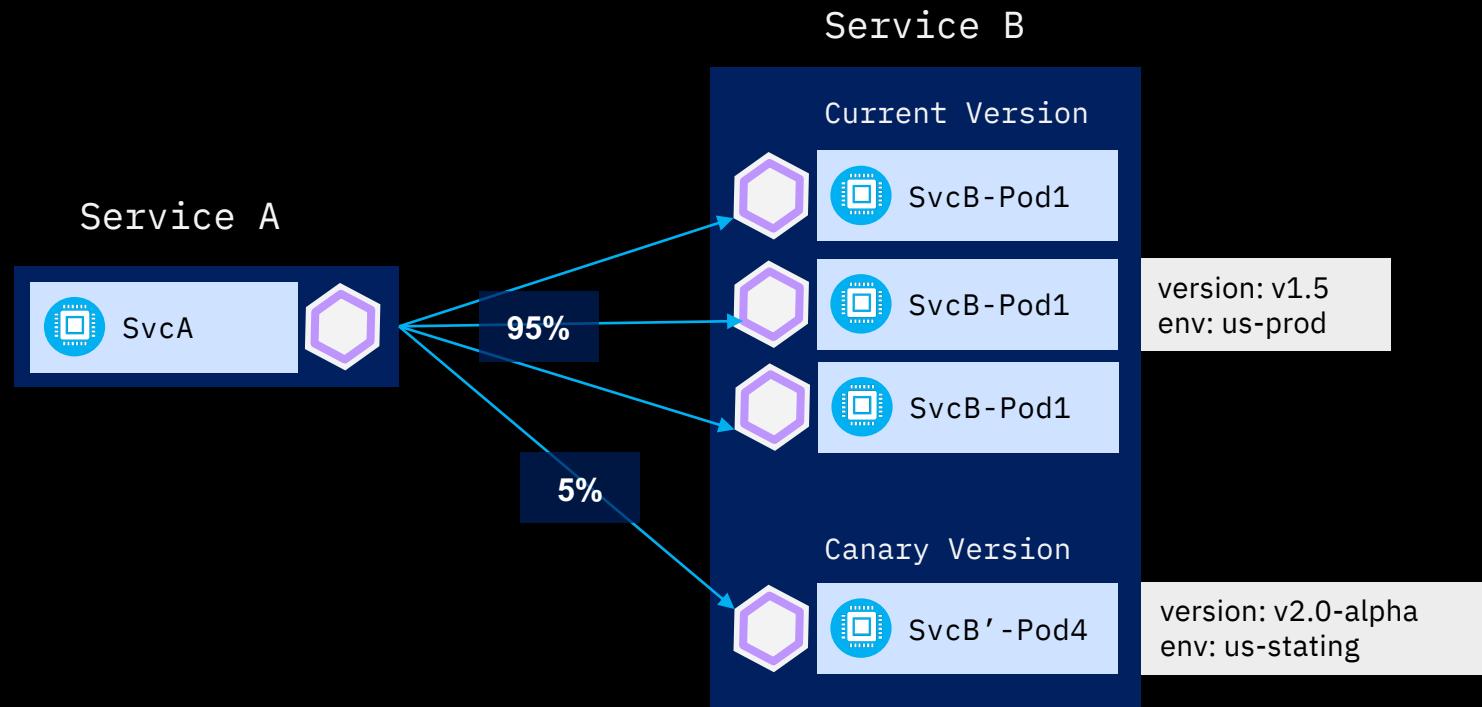
CHALLENGE 1
ROLL OUT NEW VERSION WITHOUT DOWNTIME OR CHANGING CODE

Traffic Control



CHALLENGE 1
ROLL OUT NEW VERSION WITHOUT DOWNTIME OR CHANGING CODE

Traffic Splitting

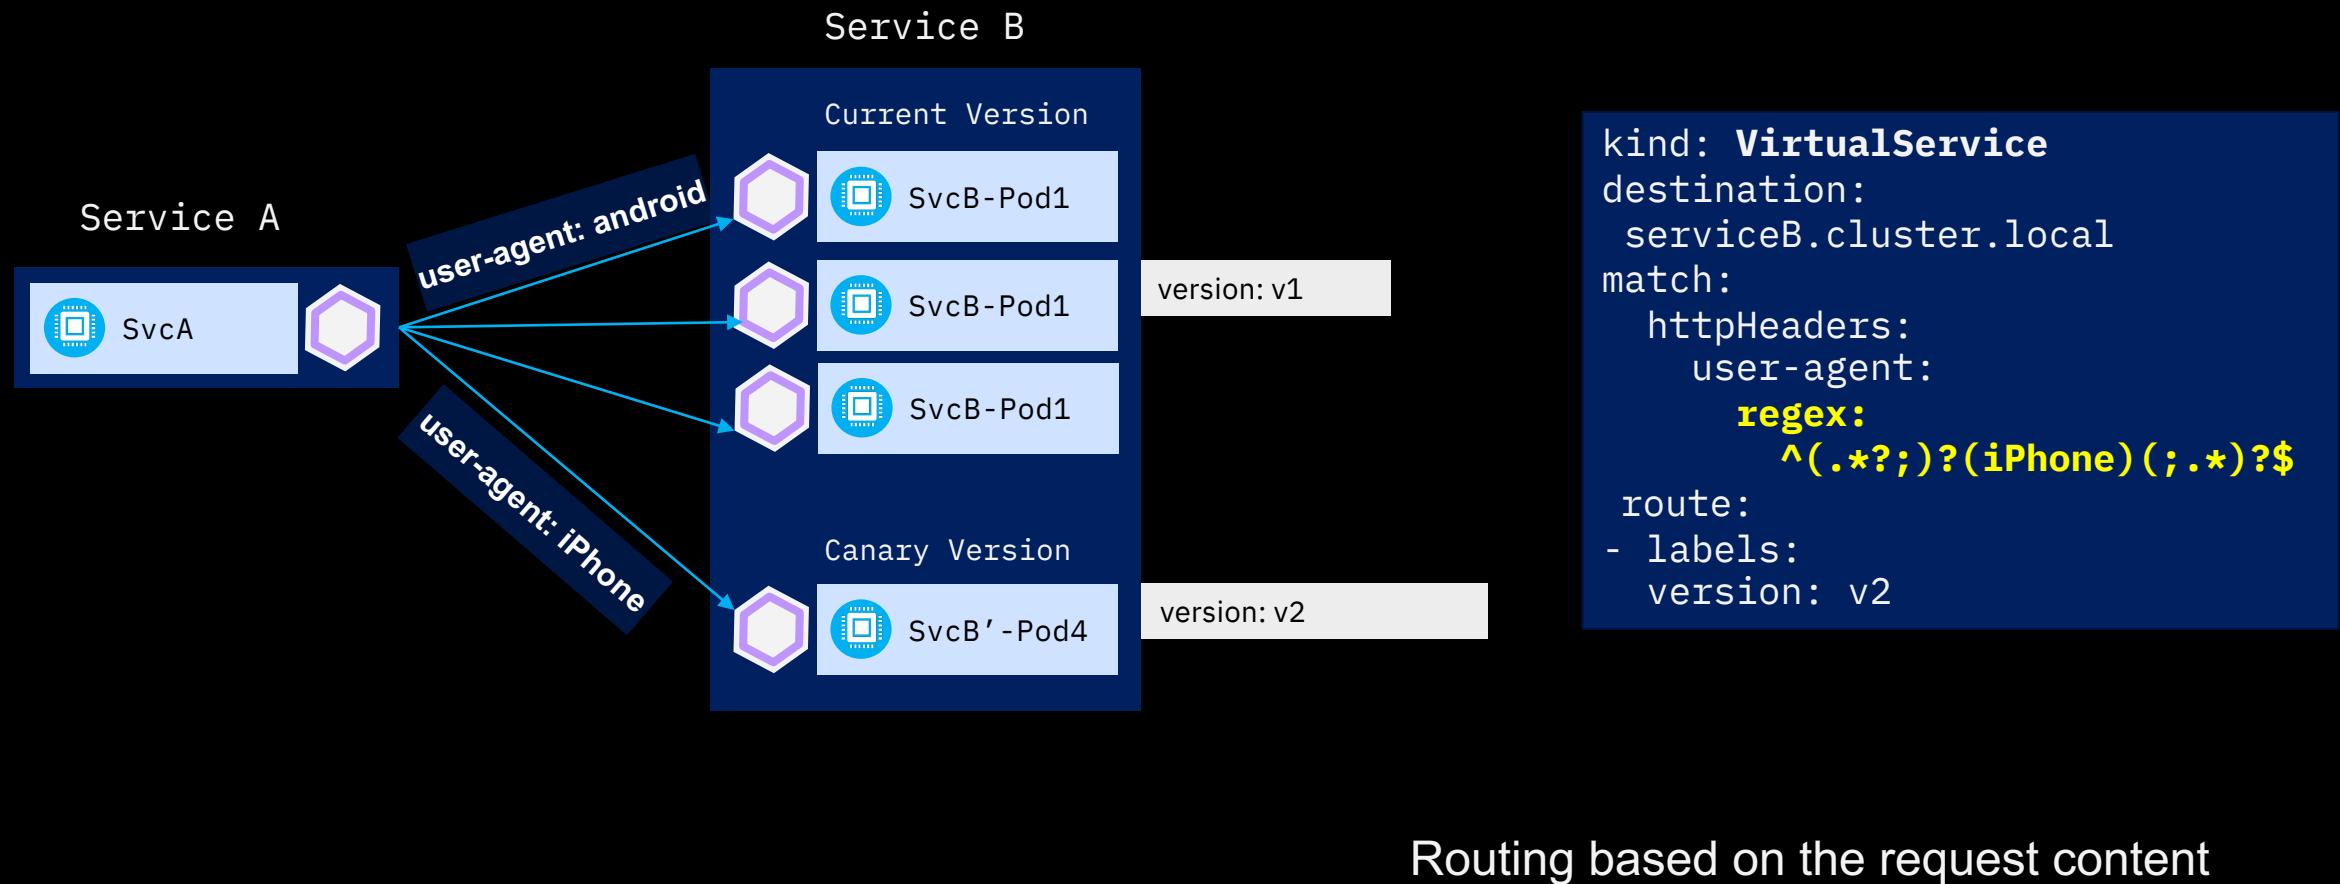


```
kind: VirtualService
destination:
  serviceB.cluster.local
match:
  source:
    serviceA.cluster.local
route:
  - labels:
      version: v1.5
      env: us-prod
  weight: 95
  - labels:
      version: v2.0-alpha
      env: us-staging
  weight: 5
```

CHALLENGE 2 HOW TO DO CANARY TESTING

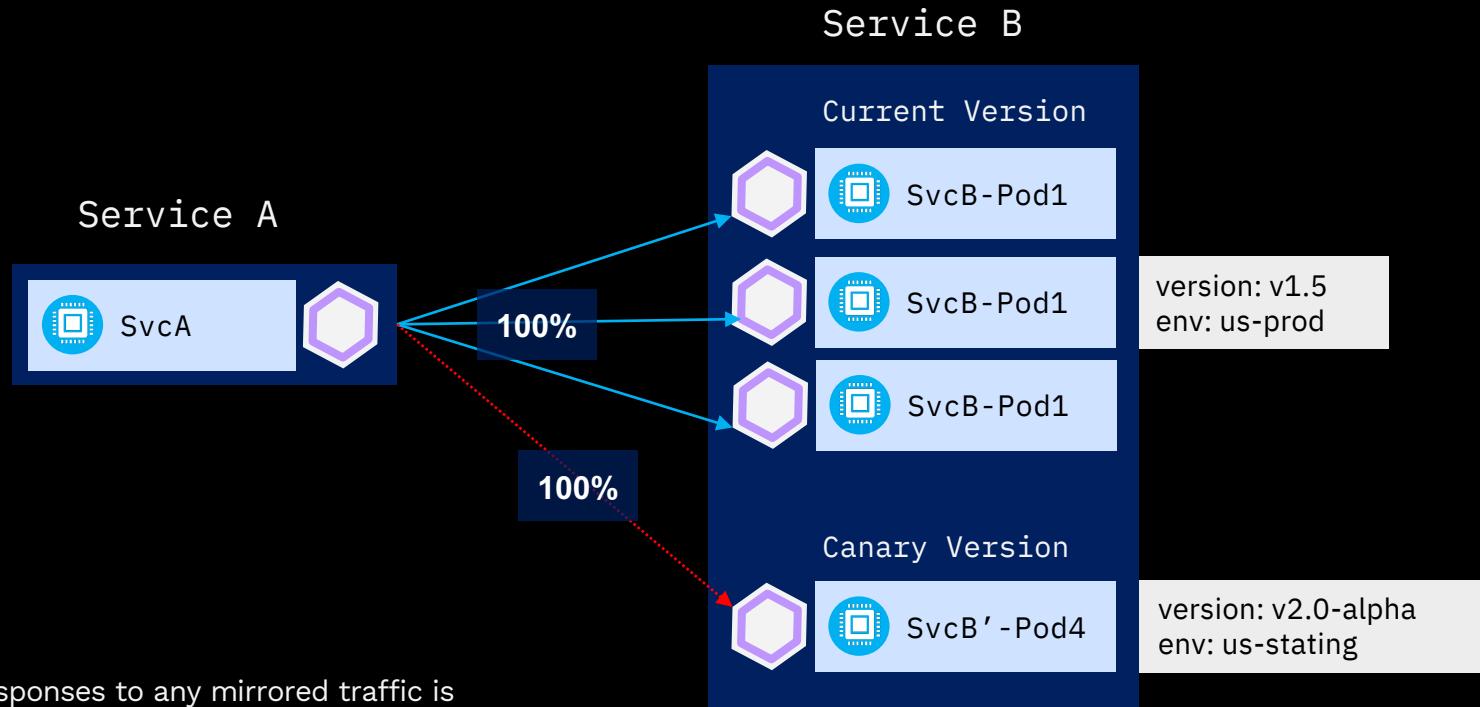
Routing not based on the request content.
Staged rollouts with %-based traffic splits.

Traffic Steering



CHALLENGE 3
HOW TO DO A/B TESTING

Traffic Mirroring



```
kind: VirtualService
destination:
  serviceB.cluster.local
match:
  source:
    serviceA.cluster.local
route:
  - labels:
      version: v1.5
      env: us-prod
    weight: 100
  - labels:
      version: v2.0-alpha
      env: us-staging
    weight: 0
    mirror:
      name: httpbin
      labels:
        version: v2.0-alpha
        env: us-staging
```

CHALLENGE 4
THINGS DON'T ALWAYS GO CORRECTLY IN PRODUCTION...

Resiliency

Istio adds fault tolerance to your application without any changes to code

```
// Circuit breakers
destination: serviceB.example.cluster.local
policy:
- labels:
  version: v1
  circuitBreaker:
    simpleCb:
      maxConnections: 100
      httpMaxRequests: 1000
      httpMaxRequestsPerConnection: 10
      httpConsecutiveErrors: 7
      sleepWindow: 15m
      httpDetectionInterval: 5m
```

Resilience features

- ❖ Timeouts
- ❖ Retries with timeout budget
- ❖ Circuit breakers
- ❖ Health checks
- ❖ AZ-aware load balancing w/ automatic failover
- ❖ Control connection pool size and request load

CHALLENGE 4
THINGS DON'T ALWAYS GO CORRECTLY IN PRODUCTION...

Resiliency

Circuit Breakers

Connection pool

- Limits connections for reviews to invoke ratings
- **Limited to 1 concurrent connection, 1 request per connection (One concurrent request total)**
 - While requests are using all of the connections in a pool, any new requests are rejected

Outlier detection

- **If there are 3 requests in 2 seconds, reviews will be ejected for 3 minutes**
 - Request 1 will take more than 2 seconds, blocking the connection during that time
 - Request 2 won't get a connection, which will generate the first error
 - Request 3 won't get a connection, which will generate the second error, causing ejection

```
kind: DestinationRule
host: reviews
trafficPolicy:
  connectionPool:
    tcp:
      maxConnections: 1
    http:
      http1MaxPendingRequests: 1
      maxRequestsPerConnection: 1
  outlierDetection:
    consecutiveErrors: 2
    interval: 2s
    baseEjectionTime: 3m
  maxEjectionPercent: 100
```

CHALLENGE 4
THINGS DON'T ALWAYS GO CORRECTLY IN PRODUCTION...

Resiliency

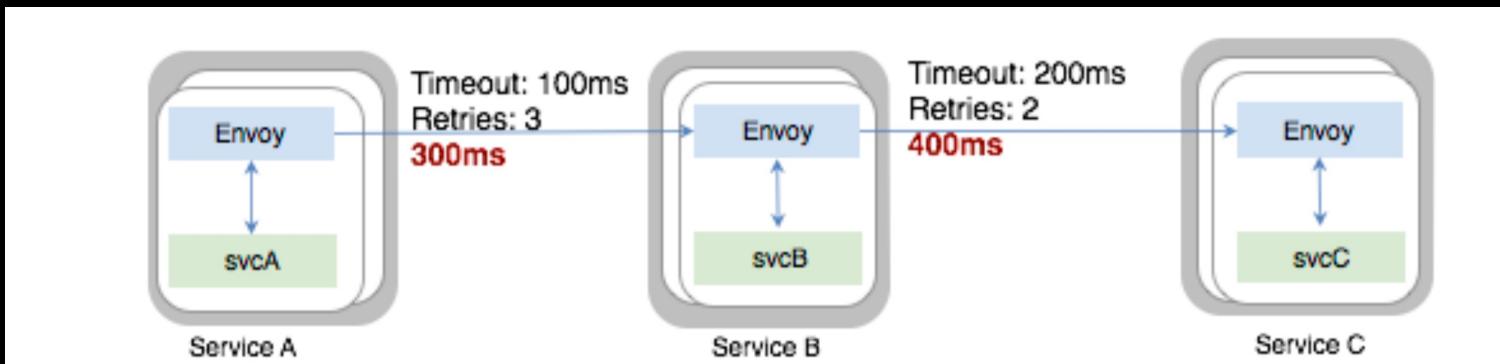
Fault injection

Fault injection can be used for testing

- Faults are caused on a percentage of requests
- Faults can cause a request delay or failure

In this example, ratings is being invoked

- All of the requests from bar have a 2 second timeout
- 40% of the requests from bar also have a 5 second delay



```
kind: VirtualService
hosts:
  - ratings
http:
  - match:
    - headers:
      end-user:
        exact: bar
    fault:
      delay:
        percent: 40
        fixedDelay: 5s
    timeout: 2s
    route:
      - destination:
          host: ratings
```

CHALLENGE 4
HOW DO I INJECT FAULT TO MY
MICROSERVICES TO PREPARE MYSELF?

Resiliency

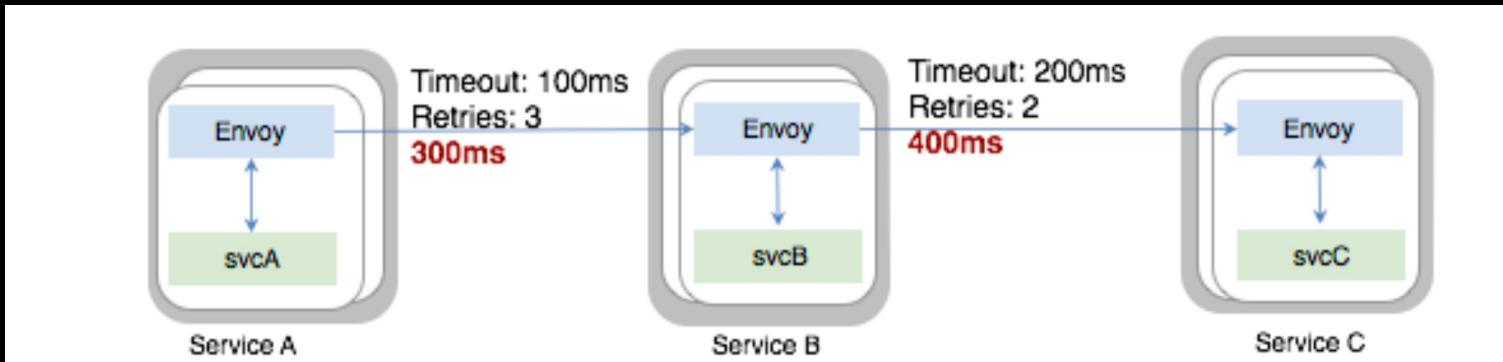
Fault injection

Fault injection can be used for testing

- Faults are caused on a percentage of requests
- Faults can cause a request delay or failure

In this example, ratings is being invoked

- **20% of the requests from foo get an HTTP 500 error**
- **All other requests (not foo or bar) have a 4 second timeout**



CHALLENGE 4
HOW DO I INJECT FAULT TO MY
MICROSERVICES TO PREPARE MYSELF?

```
kind: VirtualService
hosts:
  - ratings
http:
  - match:
    - headers:
      end-user:
        exact: foo
    fault:
      abort:
        percent: 20
        httpStatus: 500
    route:
      - destination:
          host: ratings
  - route:
    - destination:
        host: ratings
    timeout: 4s
```

Timeout is measured after the host is invoked, it is calculated after delay period has passed.

The 40% of requests from bar will time out after 7 seconds (5 sec delay + 2 sec timeout)

Rate limiting

Istio protects your application from rogue actors by imposing ratelimits

Quotas:

```
- name: requestcount.quota.istio-system
  maxAmount: 5000
  validDuration: 1s
  overrides:
    - dimensions:
        destination: ratings
        source: reviews
        sourceVersion: v3
        maxAmount: 1
        validDuration: 1s
    - dimensions:
        destination: ratings
        maxAmount: 100
        validDuration: 1s
```

Rate limit

- ❖ Configurable limits with overrides
- ❖ Multiple rate limiting backends
- ❖ Conditional rate limiting

CHALLENGE 5
HOW CAN I LIMIT RATE FOR SOME OF MY SERVICES?

Rate limiting

Every distinct rate limit configuration represents a counter.

If the number of requests in the last `validDuration` duration exceed `maxAmount`, Mixer returns a `RESOURCE_EXHAUSTED` message to the proxy.

Global rate limit of 500 calls per second.

If “reviews” is called, it’s limited to one call every 5 seconds.

If “reviews” is called from 10.28.11.20, it’s limited to 99 calls per seconds.

```
kind: handler
quotas:
- name: requestcountquota.instance.istio-system
  maxAmount: 500
  validDuration: 1s

overrides:
- dimensions:
    destination: reviews
    maxAmount: 1
    validDuration: 5s

- dimensions:
    destination: productpage
    source: "10.28.11.20"
    maxAmount: 99
    validDuration: 1s
```

CHALLENGE 5
HOW CAN I LIMIT RATE FOR SOME OF MY SERVICES?

Telemetry

- Monitoring & tracing should not be an afterthought in the infrastructure
- Goals
 - Metrics without instrumenting apps
 - Consistent metrics across fleet
 - Trace flow of requests across services
 - Portable across metric backend providers



CHALLENGE 6
I NEED TO VIEW WHAT IS GOING ON WHEN CRISIS ARISES

Kiali

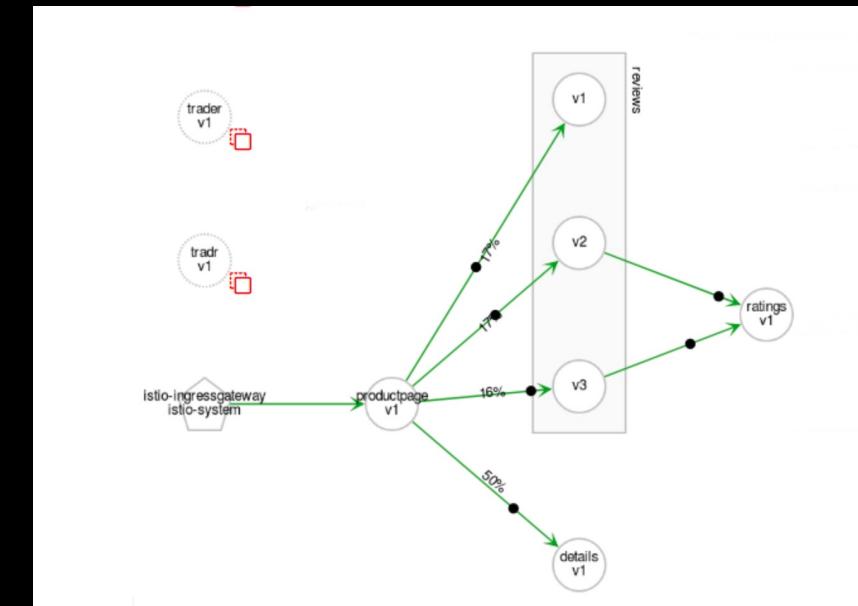
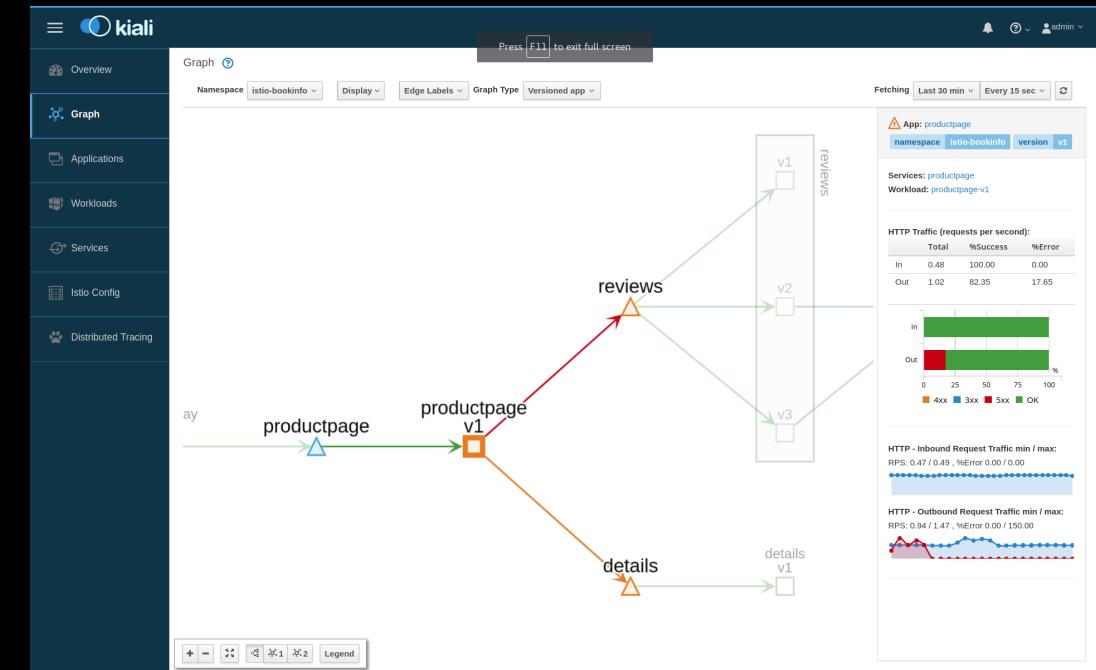
Kiali (greek κιάλι)
monocular or spyglass

Visualise the service mesh topology, features like circuit breakers or request rates

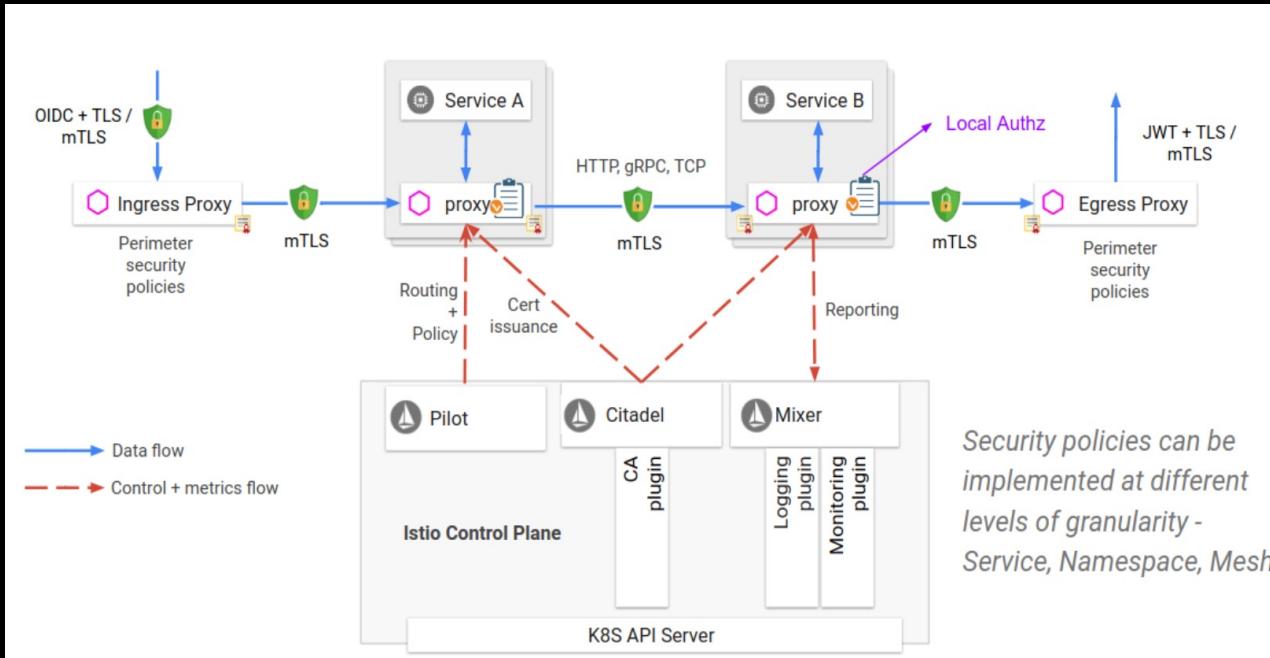
Features

- **Graph**
 - Health
 - Types
 - Side Panel
 - Traffic Animation
- **Applications, Workloads and Services**
 - Detailed Metrics
- **Traffic Routing**
- **Istio compliance**
- **Istio Configuration**

CHALLENGE 6
I NEED TO VIEW WHAT IS GOING ON WHEN CRISIS ARISES



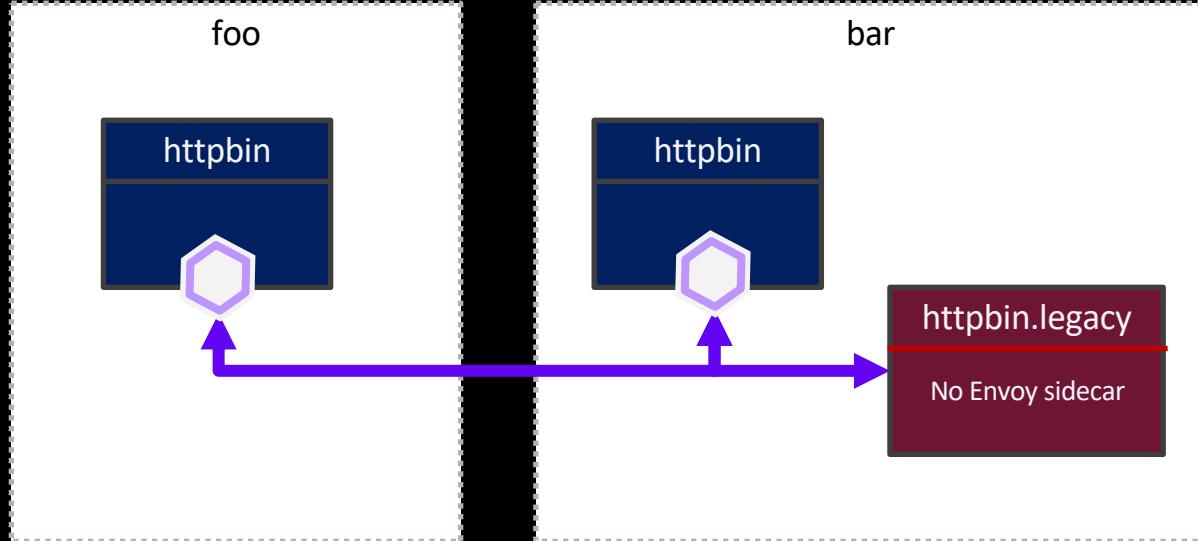
Security



- **Authentication**
 - Transport authentication, also known as service-to-service authentication
 - Origin authentication, also known as end-user authentication
- **Authorization**
 - Based on RBAC
 - Namespace-level, service-level and method-level access control for services

CHALLENGE 7
HOW CAN I SECURE MY SERVICES?

Security - Authentication

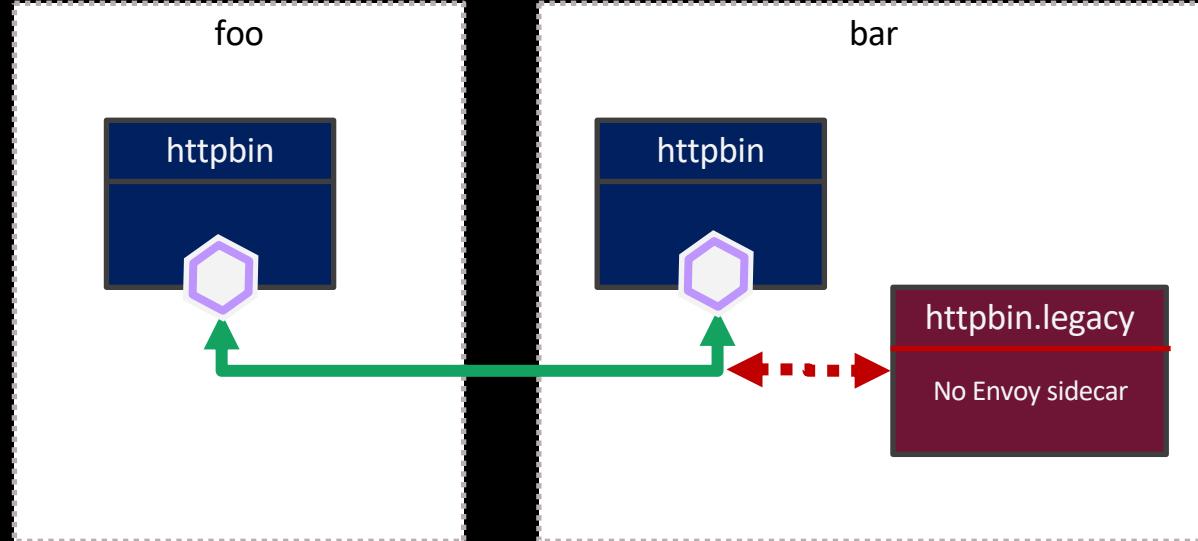


EXAMPLE

CHALLENGE 7
HOW CAN I SECURE MY SERVICES?

- Not encrypted
- Encrypted (TLS)
- No communication

Security - Authentication

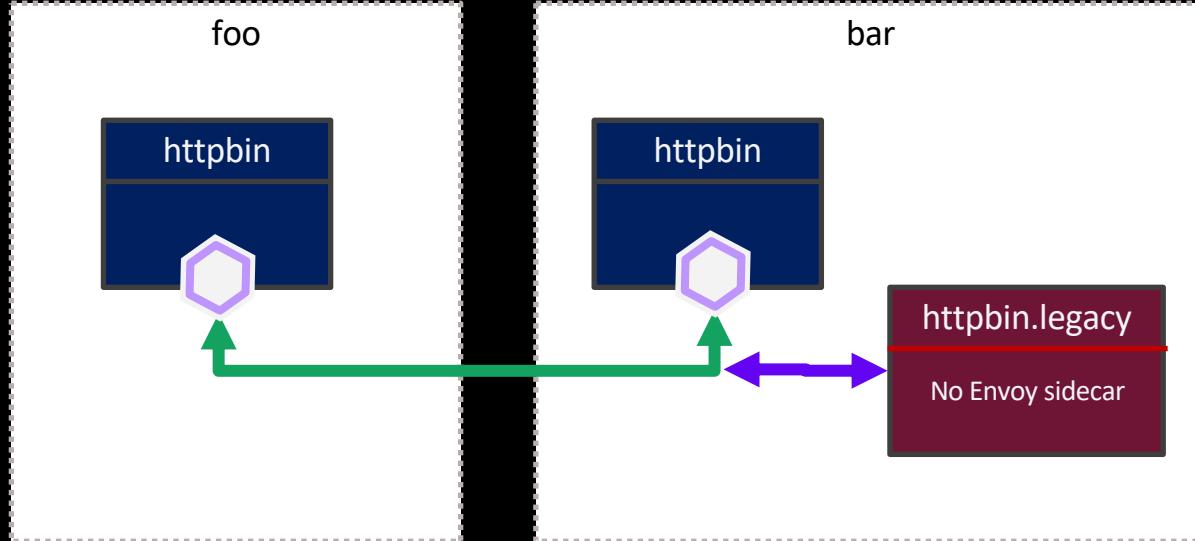


```
kind: DestinationRule
metadata:
  name: "default"
spec:
  host: "*.local"
  trafficPolicy:
    tls:
      mode: ISTIO_MUTUAL
```

CHALLENGE 7
HOW CAN I SECURE MY SERVICES?

- Not encrypted
- Encrypted (TLS)
- No communication

Security - Authentication



```
kind: DestinationRule
metadata:
  name: "httpbin-legacy"
spec:
  host: "httpbin.legacy.svc.cluster.local"
  trafficPolicy:
    tls:
      mode: DISABLE
```

CHALLENGE 7
HOW CAN I SECURE MY SERVICES?

- Not encrypted
- Encrypted (TLS)
- No communication

Security - Authorization

Istio Role Based Access

- **OFF**: Istio authorization is disabled.
- **ON**: enabled for all services in the mesh.
- **ON_WITH_INCLUSION**: enabled for all services specified in the inclusion field.
- **ON_WITH_EXCLUSION**: enabled for all services except the ones in the exclusion field.

```
kind: RbacConfig
metadata:
  name: my-user
spec:
  mode: 'ON_WITH_INCLUSION'
  inclusion:
    services:
      - "webapp.default.svc.cluster.local"
      - "frontend.default.svc.cluster.local"
      - "feedback.default.svc.cluster.local"
```

CHALLENGE 7
HOW CAN I SECURE MY SERVICES?

Security - Authorization

Istio Role Based Access

- **services**: A list of service names.
- **methods**: A list of HTTP method names (GET, POST,...)
- **paths**: HTTP paths (in the form of /packageName.serviceName/methodName)
- **constraints**: additional conditions for your rules.

```
kind: ServiceRole
metadata:
  name: service-viewer
spec:
  rules:
    - services:
        - "webapp.default.svc.cluster.local"
        - "frontend.default.svc.cluster.local"
      methods: ["GET", "HEAD"]
      paths: ["*"]
      constraints:
        - key: request.headers[version]
          values: ["v1", "v2"]
```

CHALLENGE 7
HOW CAN I SECURE MY SERVICES?

Security - Authorization

Access for:

- Service in **Namespace “default” only accessible by authenticated users** and services
- User: "*" assigns the ServiceRole to all (both authenticated and unauthenticated)

```
kind: ServiceRoleBinding
metadata:
  name: binding-products-all-authenticated-users
spec:
  subjects:
    - properties:
        source.principal: "*"
    - properties:
        source.namespace: "default"
  roleRef:
    kind: ServiceRole
    name: "service-viewer"
```

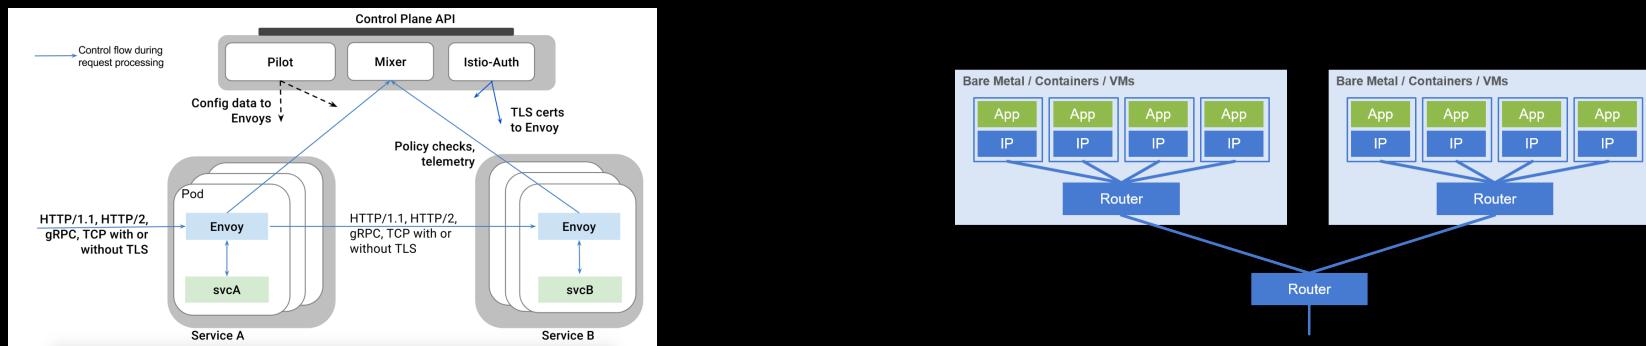
CHALLENGE 7
HOW CAN I SECURE MY SERVICES?

Security

Using Istio in concert with Calico



“RPC” – L7	Layer	“Network” – L3-4
Userspace	Implementation	Kernel
Pod	Enforcement Point	Node

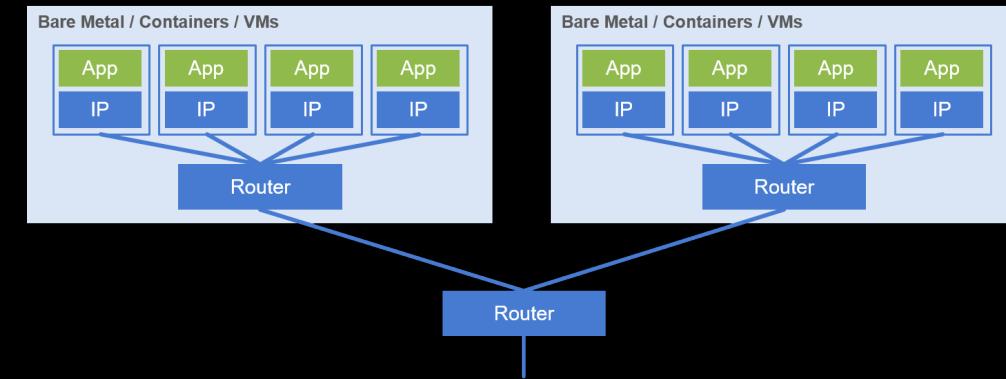


Security

Using Istio in concert with Calico

Operates at **Layer 3**, which is the network layer

- Has the advantage of being **universal** (DNS, SQL, real-time streaming, ...)
- Can extend beyond the service mesh (including to **bare metal or VM** endpoints not under the control of Kubernetes).
- Calico's policy is enforced at the host node, outside the network namespace of the guest pods.
- Based on **iptables**, which are packet filters implemented in the standard Linux kernel, it is extremely fast.



Security

Using Istio in concert with Calico



“RPC” – L7	Layer	“Network” – L3-4
Userspace	Implementation	Kernel
Pod	Enforcement Point	Node
Ideal for applying policy in support of operational goals, like service routing, retries, circuit-breaking, etc	Strengths	Universal, highly efficient, and isolated from the pods, making it ideal for applying policy in support of security goals

Service Mesh - Bad Idea ?

A Service Mesh is not always the right solution...

- ▶ **Service Meshes are Opinionated**

They are a *platform* solution. “Work their way”

- ▶ **Service Meshes are Complex**

Adds considerable complexity with sidecars and control plane

- ▶ **Service Meshes can be Slow**

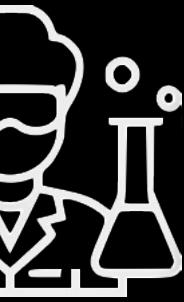
Routing traffic through a series of proxies can get painfully slow (about 700 nodes → reflector)

- ▶ **Service Meshes are for Developers**

Focused primarily on Developer view.

QUESTIONS?





◦◦ Kubernetes Workshop Series
Labs – Self Study

labs



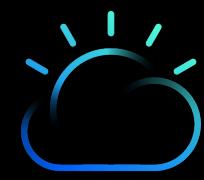
Self Study

Following the lectures there are **hands-on** labs that each participant can complete on their own time.



A screenshot of a website titled "Kubernetes Training". The left sidebar contains a navigation menu with sections like Introduction, Prerequisites, Containers, Kubernetes, Kubernetes Basics, Mesh Networking (Istio), and Kubernetes Security. Below the menu is a GitHub link. The main content area has a dark background with a grid pattern. At the top, it says "Welcome to my Kubernetes Training". It states: "I have developed them over the last year and have held several online and in-person meetups to teach fellow geeks about the awesome world of containers and orchestration. ©2023 Niklaus Hirt". Below this, there's a section titled "Training Modules" with five cards: "Introduction" (with icons for VM Isolation (OS) and Container Isolation (Kernel)), "Prerequisites" (with icons for Shared Kernel, Dynamic Kernel, and Burstable Memory), "Containers Basics" (with a diagram of a container stack), "Kubernetes Basics" (with a dashboard screenshot), "Mesh Networking (Istio)" (with a network diagram), and "Kubernetes Security" (with a screenshot of a security interface).

https://niklaushirt.github.io/k8s_training_web/

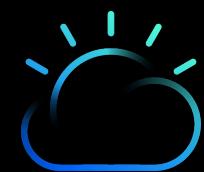


Using the online course

Select the course

The screenshot shows the 'Kubernetes Training' website. On the left is a sidebar with a dark background and white text, listing categories: Introduction, Prerequisites, Containers, Kubernetes Basics, Mesh Networking (Istio), and Kubernetes Security. Below the sidebar is a 'Github' link. The main content area has a dark background with white text. At the top, it says 'Kubernetes Training' and 'Welcome to my Kubernetes Training'. It includes a note from the developer about developing the course over the last year and holding online and in-person meetups. A copyright notice for 2023 Niklaus Hirt is also present. Below this, there's a section titled 'Training Modules' with six cards arranged in two rows of three. The first row contains 'Get Started Here / Introduction', 'Prerequisites', and 'Containers Basics'. The second row contains 'Kubernetes Basics', 'Mesh Networking (Istio)', and 'Kubernetes Security'. Each card features a small thumbnail image and a right-pointing arrow indicating they lead to more detailed pages.

Current course catalog



Using the online course

Labs for the Course



Kubernetes Training

Introduction

Prerequisites

Containers

Containers Basics

Create your first Image

Run your first Image

Creating and running the front...

Container Registry

Run from the registry

Use Portainer

Container Internals

Cleanup

Kubernetes

Kubernetes Basics

Mesh Networking (Istio)

Kubernetes Security

Github

Create your first Image

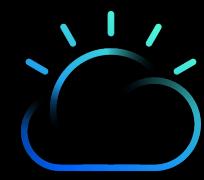
🚀 TASK: Install on Mac or PC Create your first Image

```
graph LR; Dockerfile[Dockerfile] --> build[build]; Code[Code] --> build; Config[Config] --> build; build --> Registry[Local Registry]
```

1. Let's create our first image (the `k8sdemo-backend` image) from this `Dockerfile`:

```
FROM node:8-stretch
```

Course content - Tasks



Using the online course

Commands
to be
executed



```
kubectl delete -f ./training/istio/createTraffic.yaml
kubectl delete -f ./training/istio/samples/bookinfo/platform/k
kubectl delete -f ./training/istio/samples/bookinfo/networkin
kubectl delete -f ./training/istio/samples/bookinfo/networking
kubectl delete -f ./training/istio/samples/bookinfo/networking

istioctl manifest generate --set profile=demo | kubectl delete

kubectl delete ns istio-system
```

Sample
output



```
> STEP 1/11: FROM node:8-stretch
> Resolved "node" as an alias (/etc/containers/registries.conf.d/000-shortnames.conf)
> Trying to pull docker.io/library/node:8-stretch...
> Getting image source signatures
```

Sample
Code
(json, yaml, ...)



```
FROM node:8-stretch

# Change working directory
WORKDIR "/app"

# Update packages and install dependency packages for services
RUN apt-get update \
    && apt-get dist-upgrade -y \
    && apt-get clean \
```

Show more ▾

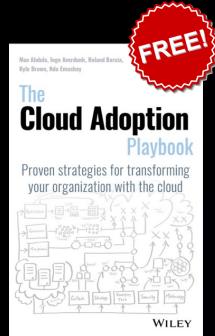
QUESTIONS?



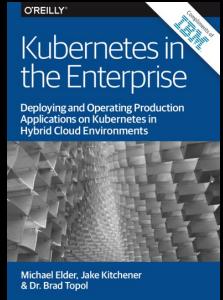
Kubernetes Workshop Series Wrap-up



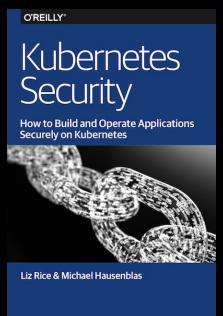
Kubernetes – Some Reading Tips



The de facto guide to improving your enterprise with the cloud, created by distinguished members of our Solution Engineering team
<http://ibm.biz/playbook>



Deploying and Operating Production Applications on Kubernetes in Hybrid Cloud Environments
<https://ibm.co/2LQketN> (excerpt)



<https://kubernetes-security.info/>

Kubernetes – Some Reading Tips

Deployment of Kubernetes Clusters

- Kubespray - <https://github.com/kubernetes-sigs/kubespray>
- Minikube - <https://github.com/kubernetes/minikube>
- Kubeadm - <https://github.com/kubernetes/kubeadm>
- Kops - <https://github.com/kubernetes/kops>
- Red Hat OpenShift Local - <https://developers.redhat.com/products/openshift-local/overview>

Management Tools

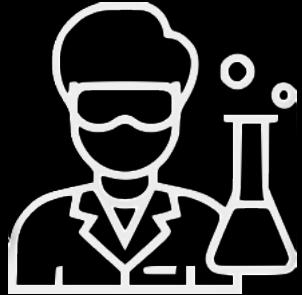
- k9s - <https://github.com/derailed/k9s>
- Lens - <https://k8slens.dev/>

Security

- Kubesec - <https://kubesec.io/>
- Falco - <https://falco.org/>
- kubehunter - <https://github.com/aquasecurity/kube-hunter>
- conftest - <https://github.com/open-policy-agent/conftest>

Deployment

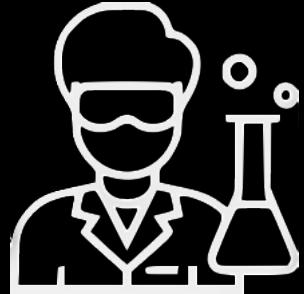
- helm - <https://helm.sh/>
- argocd - <https://argo-cd.readthedocs.io/en/stable/>



Sources and documentation are available here:

Course Documents

https://github.com/niklaushirt/k8s_training_public



Do you want to go further?

Online Course

https://niklaushirt.github.io/k8s_training_web

Course Videos – Kubernetes: from zero to hero

<https://www.youtube.com/channel/UCIS0jmGOQrG2AKKPkTJYj9w/videos>

Course Files

<https://github.com/niklaushirt/training>



Niklaus Hirt



nikh@ch.ibm.com



@nhirt