

JTC01 Docker Labs

Journey to Cloud Training



©2020 Niklaus Hirt / IBM

Task Lab0_LabInformation

Lab0 - Lab information

Docker is an open platform for developing, shipping, and running applications.

Docker enables you to separate your applications from your infrastructure so you can deliver software quickly. With Docker, you can manage your infrastructure in the same ways you manage your applications. By taking advantage of Docker's methodologies for shipping, testing, and deploying code quickly, you can significantly reduce the delay between writing code and running it in production.

In this Lab you will learn the basic operations for building and running Docker containers.

Lab sources

All the source code for the lab is available here:

<https://github.com/niklaushirt/training>

Lab overview

- In this Lab you will learn the basics of Docker:
 1. Docker basics
 2. Build a Docker image
 3. Run a Docker image
 4. Use the Portainer tool
 5. Deploy a more complex Docker application
 6. Push a Docker image into a registry

Lab0 - Lab semantics

Nomenclatures

Shell Commands

The commands that you are going to execute to progress the Labs will look like this:

```
kubectl create -f redis-slave-service.yaml  
> Output Line 1  
> Output Line 2  
> Output Line 3  
...
```

IMPORTANT NOTE: The example output of a command is prefixed by ">" in order to make it more distinguishable.

So in the above example you would only enter/copy-paste

```
kubectl create -f redis-slave-service.yaml
```

 and the output from the command is "Output Line 1" to "Output Line 3"

Code Examples

Code examples are presented like this:

```
apiVersion: lab.ibm.com/v1beta1  
kind: MyResource  
metadata:  
  name: example  
spec:  
  size: 3  
  image: busybox
```

This is only for illustration and is not being actively used in the Labs.

Lab 0 - Prepare the Lab environment

Before starting the Labs, let's make sure that we have the latest source code from the GitHub repository:

<https://github.com/niklaushirt/training>

1. Open a Terminal window by clicking on the Termnial icon in the left sidebar - we will use this extensively later as well
2. Execute the following commands to pull the latest example code from my GitHub repository

```
cd training/  
gitrefresh  
  
~/training/tools/own.sh
```

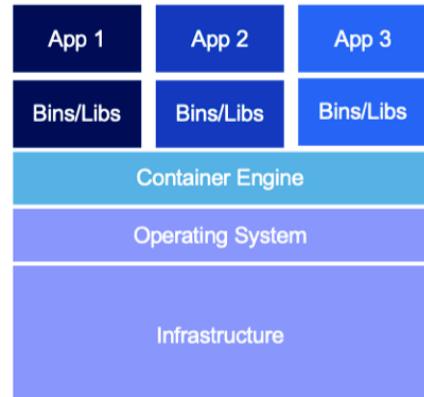
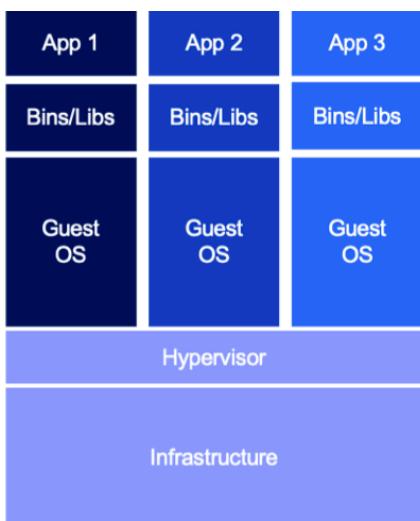
Lab 1 - Get to know Docker

Docker is an open platform for developing, shipping, and running applications.

Docker enables you to separate your applications from your infrastructure so you can deliver software quickly. With Docker, you can manage your infrastructure in the same ways you manage your applications. By taking advantage of Docker's methodologies for shipping, testing, and deploying code quickly, you can significantly reduce the delay between writing code and running it in production.

Docker provides the ability to package and run an application in a loosely isolated environment called a container. The isolation and security allow you to run many containers simultaneously on a given host. Containers are lightweight because they don't need the extra load of a hypervisor, but run directly within the host machine's kernel. This means you can run more containers on a given hardware combination than if you were using virtual machines. You can even run Docker containers within host machines that are actually virtual machines!

More details on `Docker` can be found [here](#).



- + VM Isolation
- Complete OS
- Static Compute
- Static Memory

- + Container Isolation
- + Shared Kernel
- + Burstable Compute
- + Burstable Memory

Dockerfile

A `Dockerfile` is a text document that contains all the commands a user could call on the command line

to assemble an image.

```
FROM node:8-stretch

# Change working directory
WORKDIR "/app"

# Update packages and install dependency packages for services
RUN apt-get update \
&& apt-get dist-upgrade -y \
&& apt-get clean \
&& echo 'Finished installing dependencies'

# Install npm production packages
COPY package.json /app/
RUN cd /app; npm install --production

COPY . /app

ENV NODE_ENV production
ENV BACKEND_URL https://api.nasa.gov/planetary/apod?api_key=DEMO_KEY
ENV PORT 3000

EXPOSE 3000

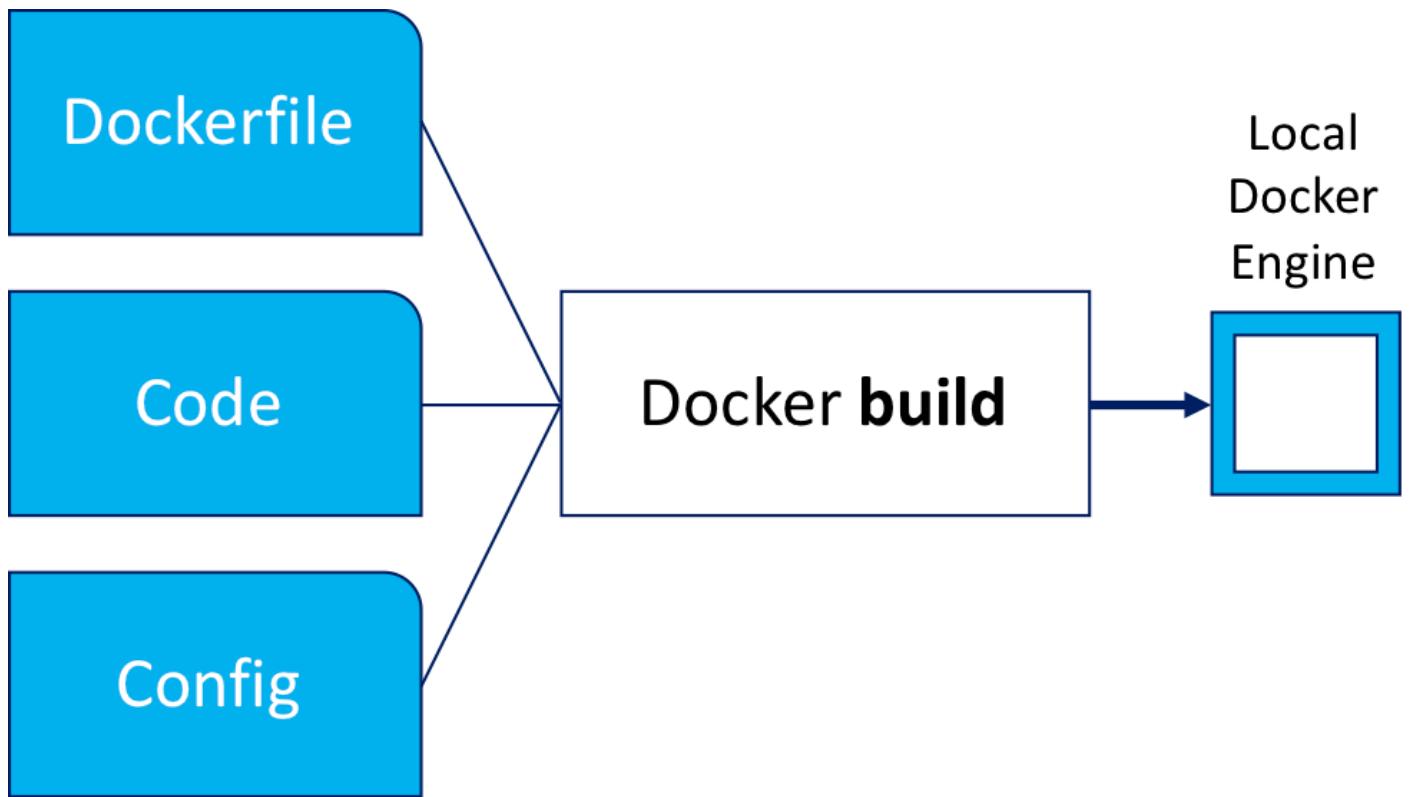
CMD ["npm", "start"]
```

The `docker build` command creates the container from the `Dockerfile`.

More details on Dockerfiles can be found [here](#).

Lab 2 - Docker

Lab 2 - Create your first Image



Let's create our first image (the `k8sdemo-backend` image) from this `Dockerfile` :

```
FROM node:8-stretch

# Change working directory
WORKDIR "/app"

# Update packages and install dependency packages for services
RUN apt-get update \
&& apt-get dist-upgrade -y \
&& apt-get clean \
&& echo 'Finished installing dependencies'

# Install npm production packages
COPY package.json /app/
RUN cd /app; npm install --production

COPY . /app

ENV NODE_ENV production
ENV BACKEND_MESSAGE HelloWorld

ENV PORT 3000

EXPOSE 3000

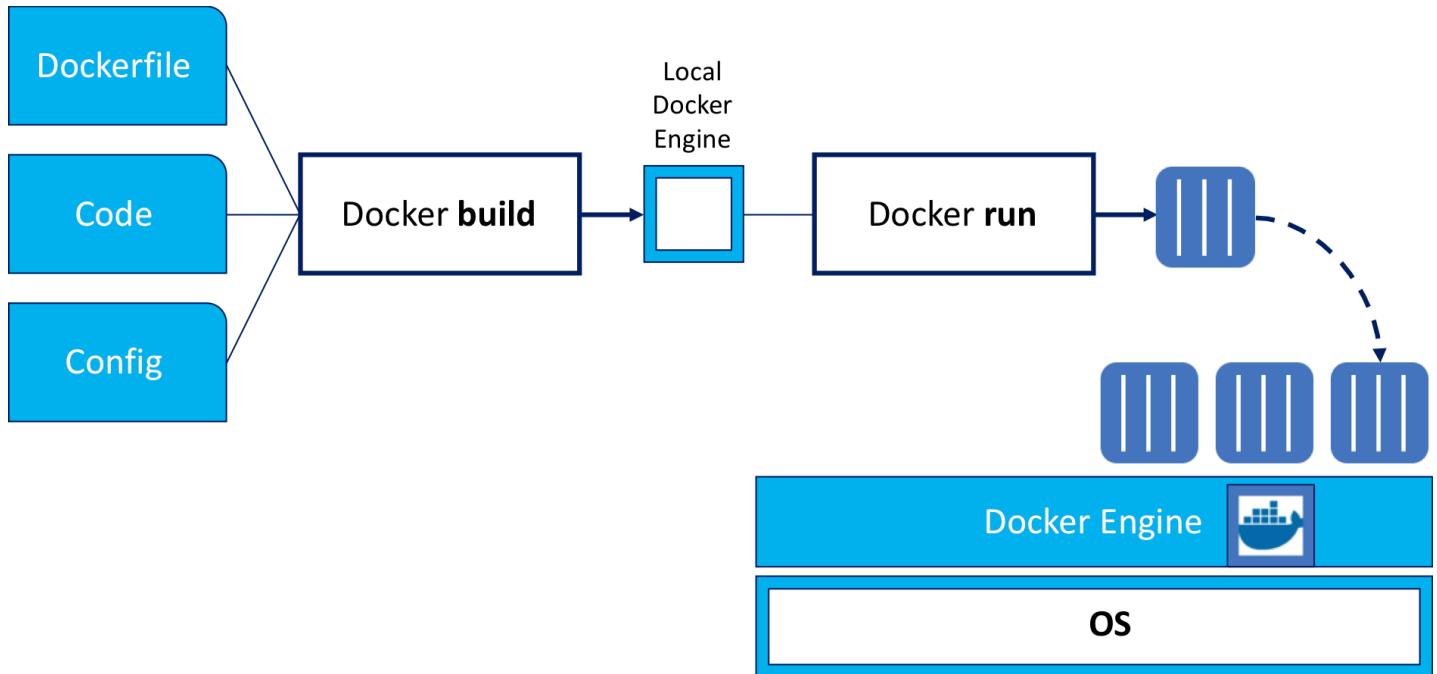
CMD ["npm", "start"]
```

```
cd ~/training/demo-app/k8sdemo_backend

docker build -t k8sdemo-backend:lab .

> Sending build context to Docker daemon 6.975MB
> Step 1/11 : FROM node:8-stretch
>    ---> 7a9afc16a57f
> Step 2/11 : WORKDIR "/app"
>    ---> Using cache
>    ---> a2515f8a3ec5
...
> Step 11/11 : CMD ["npm", "start"]
>    ---> Using cache
>    ---> b9b0f3fea9f7
> Successfully built b9b0f3fea9f7
> Successfully tagged k8sdemo-backend:lab
```

Lab 2 - Run your first Image



```
docker run --rm --name k8sdemo-backend -p 3001:3000 k8sdemo-backend:lab  
> test@0.0.0 start /app  
> node ./bin/www
```

This command runs the backend server:
* `--rm` makes sure that the container is deleted once it's stopped
* `--name` gives the container a fixed name (otherwise you get some pretty funny, automatically generated names - think drunken-weasel)
* `-p` exposes the container port 3000 to the outside port 3001 (we do this so that it does not conflict with port 3000 of the k8sdemo web application we will start later)
* `k8sdemo-backend:lab` is the image we created before

Lab 2 - Use Portainer

Portainer Community Edition is a powerful, open-source management toolset that allows you to easily build, manage and maintain Docker environments.

1. [Open URL](#) or use the Portainer Bookmark
2. Login in with `admin` / `admin` (already prefilled)
3. Select `local` for our Endpoint

The screenshot shows the Portainer interface for the 'local' endpoint. At the top, there's a navigation bar with 'Home' and 'Endpoints' on the left, and 'Portainer support', 'admin', 'my account', and 'log out' on the right. Below the navigation is a search bar labeled 'Search by name, group, tag, status, URL...'. The main area displays the 'Endpoints' section, which includes a 'Refresh' button and a table with the following data:

local		2019-10-07 09:05:25	Group: Unassigned	Standalone 19.03.2
	up	0 stacks	32 containers - 31 healthy 1 warning	1 volume
		26 images		/var/run/docker.sock
		4 6.5 GB	No tags	

At the bottom right, there are buttons for 'Items per page' (set to 10) and a dropdown menu.

Now you get an overview of your local Docker instance.

4. Select `Containers`

Dashboard
Endpoint summary

Endpoint: local | 4 | 6.5 GB - Standalone 19.03.2
URL: /var/run/docker.sock
Tags: -

0 Stacks	32 Containers <small>31 running 1 stopped</small>
26 Images <small>7.2 GB</small>	1 Volume
3 Networks	

portainer.io 1.22.0

5. You get a list of all running containers and you can see our `k8sdemo-backend` container running.

Container list
Containers

Containers

Name	State	Quick actions	Stack	Image	Created	IP Address	Published Ports	Ownership
k8sdemo-backend	running		-	k8sdemo-backend:lab	2019-10-07 09:07:56	172.17.0.3	3001:3000	administrator
k8s_tiller_tiller-deploy-7f4d76c...	running		-	gcr.io/kubernetes-helm/tiller	2019-10-04 17:45:09	-	-	administrator
k8s POD tiller-deploy-7f4d76c...	running		-	k8s.gcr.io/pause:3.1	2019-10-04 17:44:39	-	-	administrator
k8s_dashboard-metrics-scraper...	running		-	kubernetesui/metrics-scraper	2019-10-04 10:34:55	-	-	administrator
k8s_kubernetes-dashboard_kube...	running		-	6802d83967b9	2019-10-04 10:34:49	-	-	administrator
k8s POD kubernetes-dashboard-...	running		-	k8s.gcr.io/pause:3.1	2019-10-04 10:34:48	-	-	administrator
k8s POD dashboard-metrics-scr...	running		-	k8s.gcr.io/pause:3.1	2019-10-04 10:34:48	-	-	administrator
PORTAINER	running		-	portainer/portainer	2019-10-04 10:29:33	172.17.0.2	9010:9000	administrator
k8s_cilium-operator_cilium-op...	running		-	Bede72a3bc5	2019-10-04 10:27:39	-	-	administrator
k8s POD cilium-operator-76df8...	running		-	k8s.gcr.io/pause:3.1	2019-10-04 10:27:39	-	-	administrator

6. Click on the PublishedPorts 3001:3000 to open the backend web interface.

7. In your terminal you will see that this generated some traffic:

```
docker run --rm --name k8sdemo-backend -p 3001:3000 k8sdemo-backend:lab  
> test@0.0.0 start /app  
> node ./bin/www  
  
> GET / 304 225.805 ms --  
> GET /stylesheets/style.css 304 2.175 ms --
```

8. Stop the container by hitting `CTRL-C` in the terminal

So now we have tested our backend component.

Lab 2 - Create the frontend Image

Let's create our first image (the `k8sdemo` frontend image) from this `Dockerfile`:

```
FROM node:8-stretch

# Change working directory
WORKDIR "/app"

# Update packages and install dependency packages for services
RUN apt-get update \
&& apt-get dist-upgrade -y \
&& apt-get clean \
&& echo 'Finished installing dependencies'

# Install npm production packages
COPY package.json /app/
RUN cd /app; npm install --production

COPY . /app

ENV NODE_ENV production
ENV BACKEND_URL https://api.nasa.gov/planetary/apod\?api_key\=DEMO_KEY
ENV PORT 3000

EXPOSE 3000

CMD ["npm", "start"]
```

```
cd ~/training/demo-app/k8sdemo

docker build -t k8sdemo:lab .

> Sending build context to Docker daemon 13.23MB
> Step 1/11 : FROM node:8-stretch
>  --> 7a9afc16a57f
> Step 2/11 : WORKDIR "/app"
>  --> Using cache
>  --> a2515f8a3ec5
...
> Step 11/11 : CMD ["npm", "start"]
>  --> Using cache
>  --> 5293cb32d1f6
> Successfully built 5293cb32d1f6
> Successfully tagged k8sdemo:lab
```

Lab 2 - Run the frontend Image

1. First let's run the backend container again, but this time in the background

```
docker run --rm -d --name k8sdemo-backend -p 3001:3000 k8sdemo-backend:lab  
> 444b0570058b97f0532ef89c92963bb7da6aa1f2d3e27bf8c989da5fb8277fe0
```

This command runs the backend server:

- -d runs the container in the background (as a daemon)

2. Then we start the new Web Frontend container

```
docker run --rm --name k8sdemo -p 3000:3000 --env BACKEND_URL=http://10.0.2.15:  
> test@0.0.0 start /app  
> node ./bin/www
```

This command runs the frontend server:

- --rm makes sure that the container is deleted once it's stopped
- --name gives the container a fixed name
- --env defines the environment variable that points to the `k8sdemo-backend` server API
- -p exposes the container port 3000 to the outside port 3000
- k8sdemo:lab is the image we created before

3. Go back to Portainer and refresh the browser

Name	State	Quick actions	Stack	Image	Created	IP Address	Published Ports	Ownership
k8sdemo	running		-	k8sdemo:lab	2019-10-07 09:26:41	172.17.0.4	3000:3000	administrat
k8sdemo-backend	running		-	k8sdemo-backend:lab	2019-10-07 09:23:50	172.17.0.3	3001:3000	administrat
k8s_tiller_tiller-deploy-7f4d...	running		-	gcr.io/kubernetes-helm/tiller	2019-10-04 17:45:09	-	-	administrat

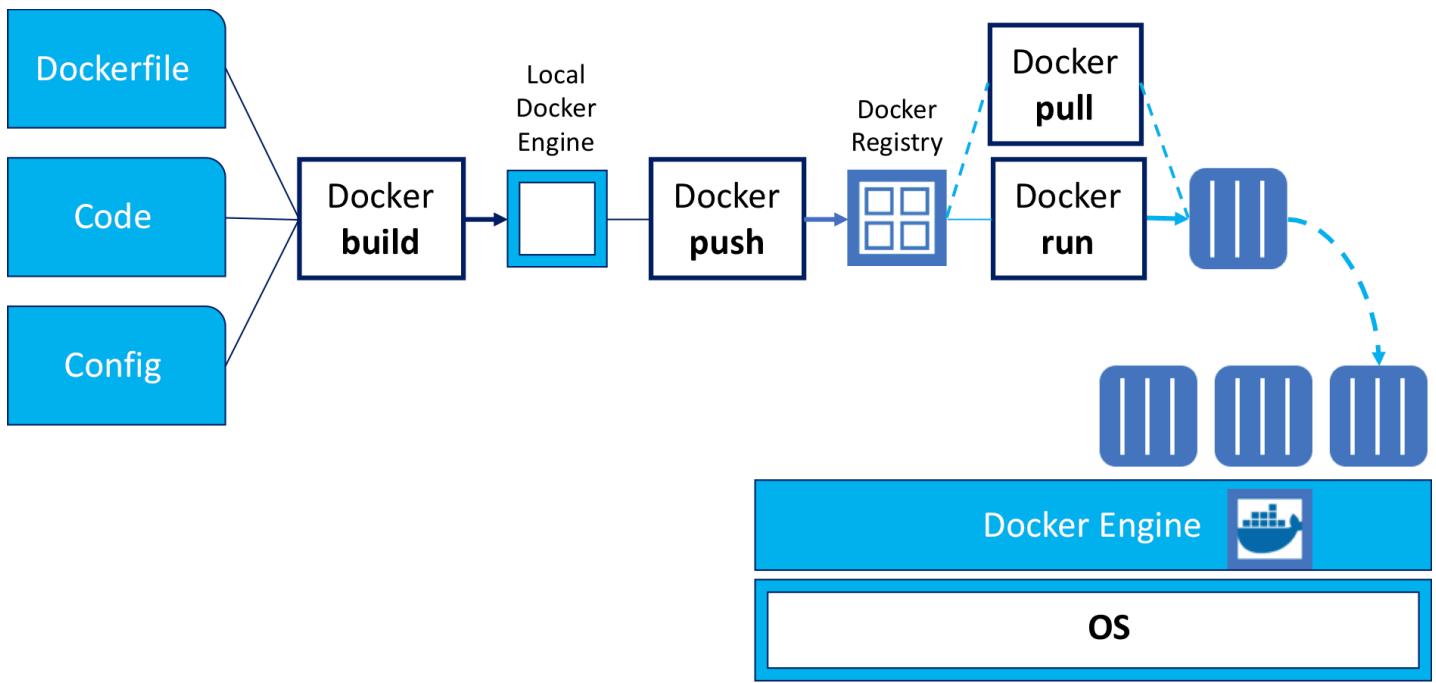
4. Click on the PublishedPorts 3000:3000 for `k8sdemo` to open the web interface.

5. Refresh several times and check in the terminal that some traffic is being generated

```
GET / 304 6.314 ms - -
GET /public/bootstrap.min.css 304 0.975 ms - -
GET /public/bootstrap-theme.min.css 304 0.843 ms - -
GET /public/stylesheets/style.css 304 2.568 ms - -
GET /public/images/ibm_cloud.png 304 0.522 ms - -
GET /public/images/cloud_private.png 304 1.057 ms - -
GET /public/images/back.png 304 0.411 ms - -
The value of BACKEND_URL is: http://k8sdemo-backend
Error: getaddrinfo ENOTFOUND k8sdemo-backend k8sdemo-backend:80
```

6. Stop the container by hitting `CTRL-C` in the terminal

Lab 2 - Push the frontend Image to the registry



1. Let's tag the image with the address of the local Docker registry (localhost:5000).

```
docker tag k8sdemo:lab localhost:5000/k8sdemo:lab
```

2. Expose the local Docker registry.

First execute this in order to be able to access the private registry:

```
kubectl port-forward --namespace kube-system registry-5ng6b 5000:5000 &
```

This exposes the Docker Registry to the Terminal we are using.

If you don't get a command line prompt, press `enter` several times

3. And now push the image to the local registry:

```
docker push localhost:5000/k8sdemo:lab
```

Messages like `Handling connection for 5000` and `Retrying in x seconds` are due to the Lab setup and can be ignored.

Now the image is available to be acquired (pulled) from any machine that has access to the Docker registry.

Lab 2 - Run the frontend Image from the registry

1. Now let's start the Web Frontend container with the image from the registry

```
docker run -d --rm --name k8sdemo -p 3000:3000 --env BACKEND_URL=http://10.0.2.  
> b6e46d8bd60978af7e9e45260111e938da63a64247c9cff3b4e398a6498670a6
```

This command runs the frontend server:

- --rm makes sure that the container is deleted once it's stopped
- --name gives the container a fixed name
- --env defines the environment variable that points to the `k8sdemo-backend` server API
- -p exposes the container port 3000 to the outside port 3000
- -d runs the container in the background (as a daemon)
- `localhost:5000/k8sdemo:lab` is the image we have pushed to the registry before

2. Go back to your browser and refresh the `k8sdemo` web application to make sure that the container has been started.

Lab 2 - Cleanup

To conclude this Lab we have to clean up the containers that we have created

1. Terminate the frontend

```
docker kill k8sdemo  
> k8sdemo
```

2. And Terminate the backend

```
docker kill k8sdemo-backend  
> k8sdemo-backend
```

3. Verify that the two containers have been terminated

```
docker ps | grep k8sdemo
```

This command must return no result.

Congratulations!!! This concludes Lab 2 on Docker