

The Journey to Cloud

DevSecOps
and
From Microservices
to Kubernetes

Niklaus Hirt
DevOps Architect / Cloud Architect
nikh@ch.ibm.com



Who am I?

Niklaus Hirt

Passionate about tech for over 35 years

- High-school in Berne
- Degree in Computer Science at EPFL
- ELCA
- CAST
- IBM

✉ nikh@ch.ibm.com

🐦 @nhirt



Agenda – DevSecOps & Kubernetes Basic Concepts

Module 0: Prepare the Labs

Module 1: DevSecOps

BREAK

Module 2: Microservices

Module 3: Containers with Docker

Module 4: Kubernetes

Module 5: Let's get real

Module 6: Kubernetes Security

LUNCH

Agenda – DevSecOps & Kubernetes Basic Concepts

Module 7: Docker Hands-On

Module 8: Kubernetes Hands-On

Module 9: Kubernetes Security Hands-On (Optional)

Wrap-up

END of course – 16:00



Documentation

https://github.com/niklaushirt/k8s_training_public

Sources

<https://github.com/niklaushirt/training>

Always-on training environment

<http://158.177.137.195:31999/>



°° The Journey to Cloud
Prepare the Labs



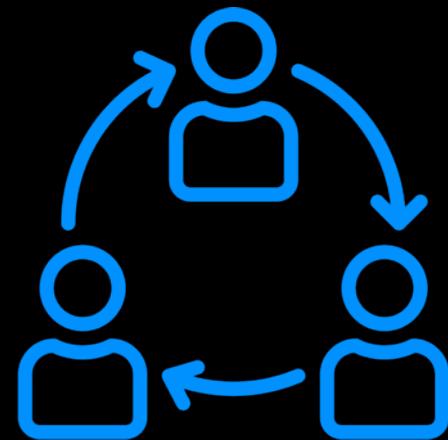
IBM Cloud

Session Objectives

Attendees will be grouped in **teams** wherever it makes sense to facilitate collaborative work.



Following some lectures will be **hands-on** work that each team collaborates to complete.



Teams

black 31701

olive 31711

peru 31715

white 31702

brown 31712

chocolate 31716

red 31703

lightblue 31713

orchid 31717

blue 31704

orange 31708

gold 31718

yellow 31705

purple 31709

pink 31719

lime 31706

maroon 31710

violet 31720

cyan 31707

firebrick 31714



Collector - Accessing team web site

http://158.177.137.195:{port#}

Team name / color will be shown

blue 31704

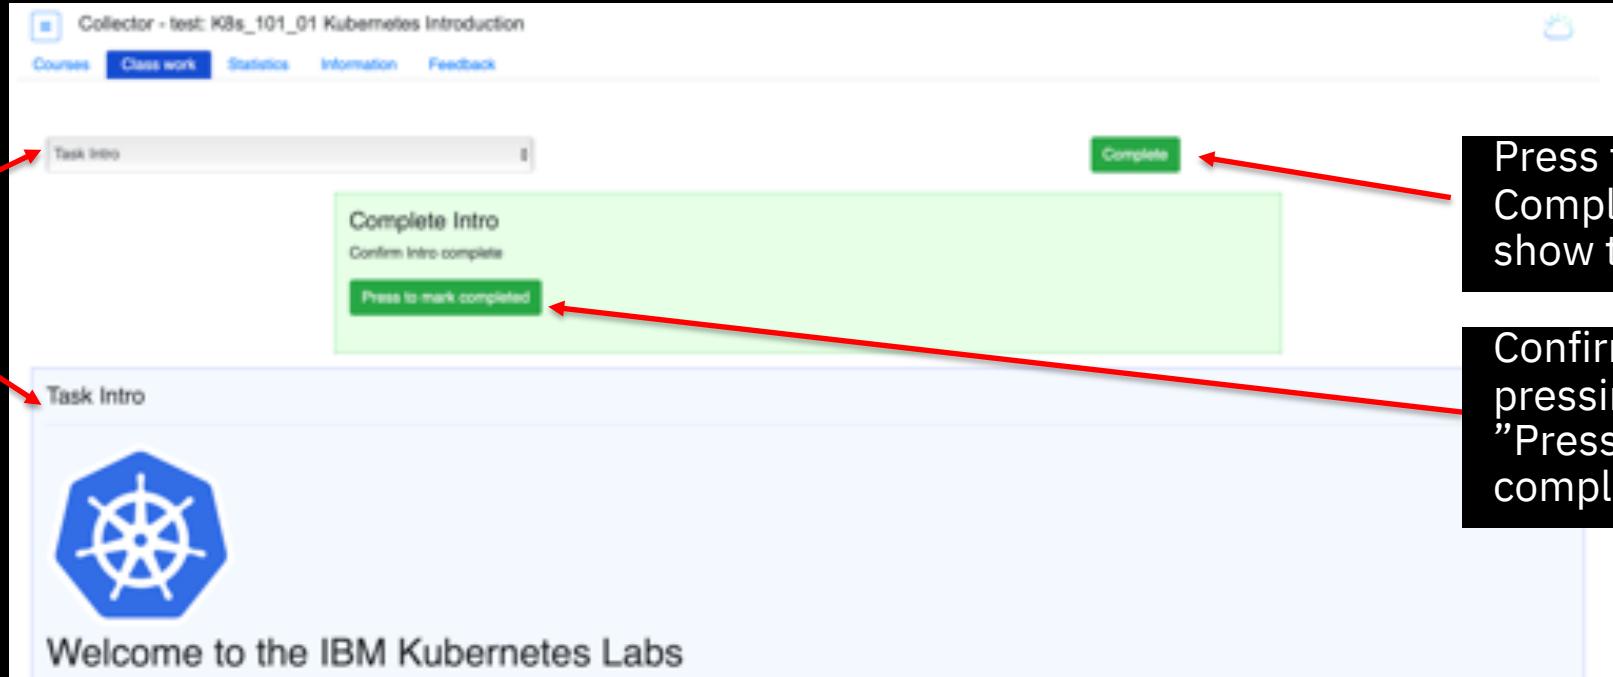
The screenshot shows a web browser window with the title "Collector - blue". The navigation bar includes links for "Courses", "Class work", "Statistics", "Information", and "Feedback". Below the navigation bar, a section titled "Catalog of courses" displays a list of courses under the heading "select course". The listed courses are: KUB01 Lab Setup, KUB02 Kubernetes Introduction, KUB03 Kubernetes Labs, and KUBADV01 Istio. To the right of the course list is a "Begin course" button. A large blue callout box with white text is positioned over the "Begin course" button, containing the instruction "Select course and press button to begin". Red arrows point from the text "Team name / color will be shown" to the "blue" text in the title bar, and from the text "Current course catalog" to the "select course" dropdown.

- KUB01 Lab Setup
- KUB02 Kubernetes Introduction
- KUB03 Kubernetes Labs
- KUBADV01 Istio

Current course catalog

Collector – Class work

Select class work and the blue portion of the screen is shown



Press the green Complete button to show the green portion.

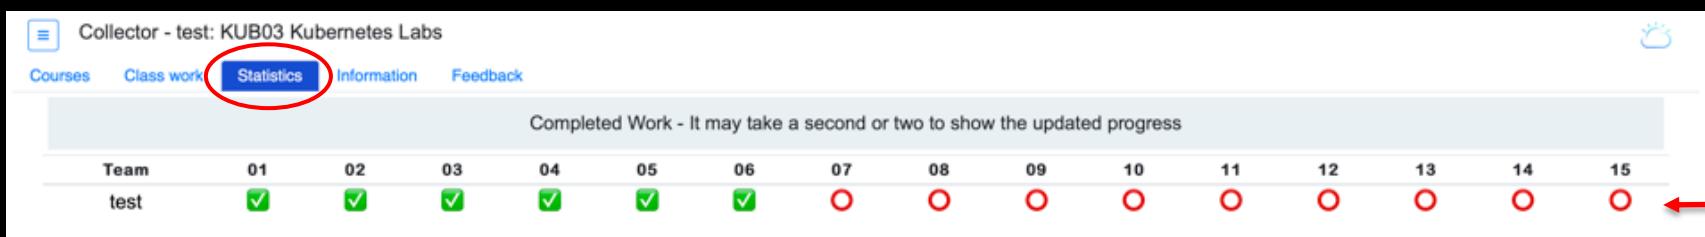
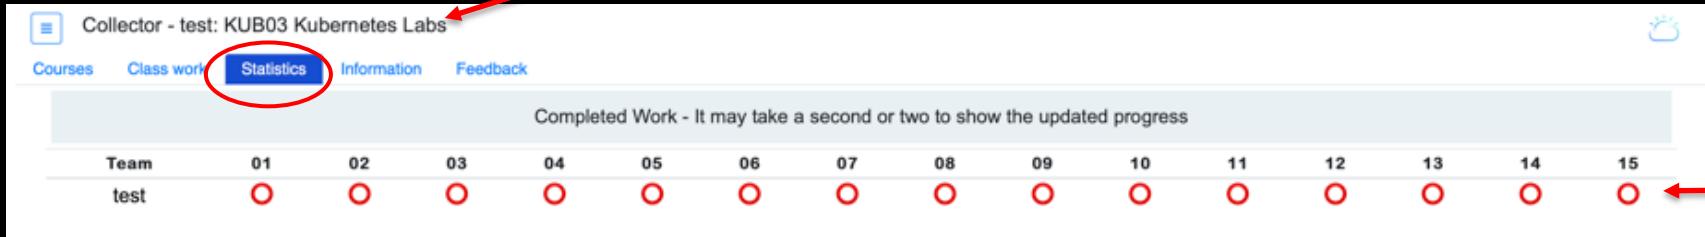
Confirm completion by pressing the green "Press to mark completed" button.



The Complete Button might not show instantly depending on the course settings

Collector – Track course completed work

Course title



Green checkmark - item is completed

Red circle - item is waiting to be completed

The number of items tracked will change based on the current course selected.

Collector – Instructor Dashboard

Remaining Time for the Lab

Collector - instructor: K8s_101_01 Kubernetes Introduction

Remaining time: 0h 29m 50s

Courses Class work Statistics Information Feedback Insight

Completed Work - It may take a second or two to show the updated progress

| Team | 01 | 02 | 03 | 04 | 05 | 06 |
|------------|----|----|----|----|----|----|
| instructor | ✓ | ○ | ○ | ○ | ○ | 0 |
| lightblue | ✓ | ✓ | ✓ | ✓ | ✓ | 1 |
| olive | ✓ | ✓ | ○ | ○ | ○ | 2 |
| peru | ○ | ○ | ○ | ○ | ○ | 3 |
| chocolate | ○ | ○ | ○ | ○ | ○ | 4 |
| pink | ○ | ○ | ○ | ○ | ○ | 5 |
| violet | ○ | ○ | ○ | ○ | ○ | 6 |





JTC90 Lab Setup

EVERYBODY

Task 2: Setup Kubectl

OK

Task 3: Setup git

Task 4: Final Check

?

The Journey to Cloud **DevSecOps**

01



IBM Cloud



disruption

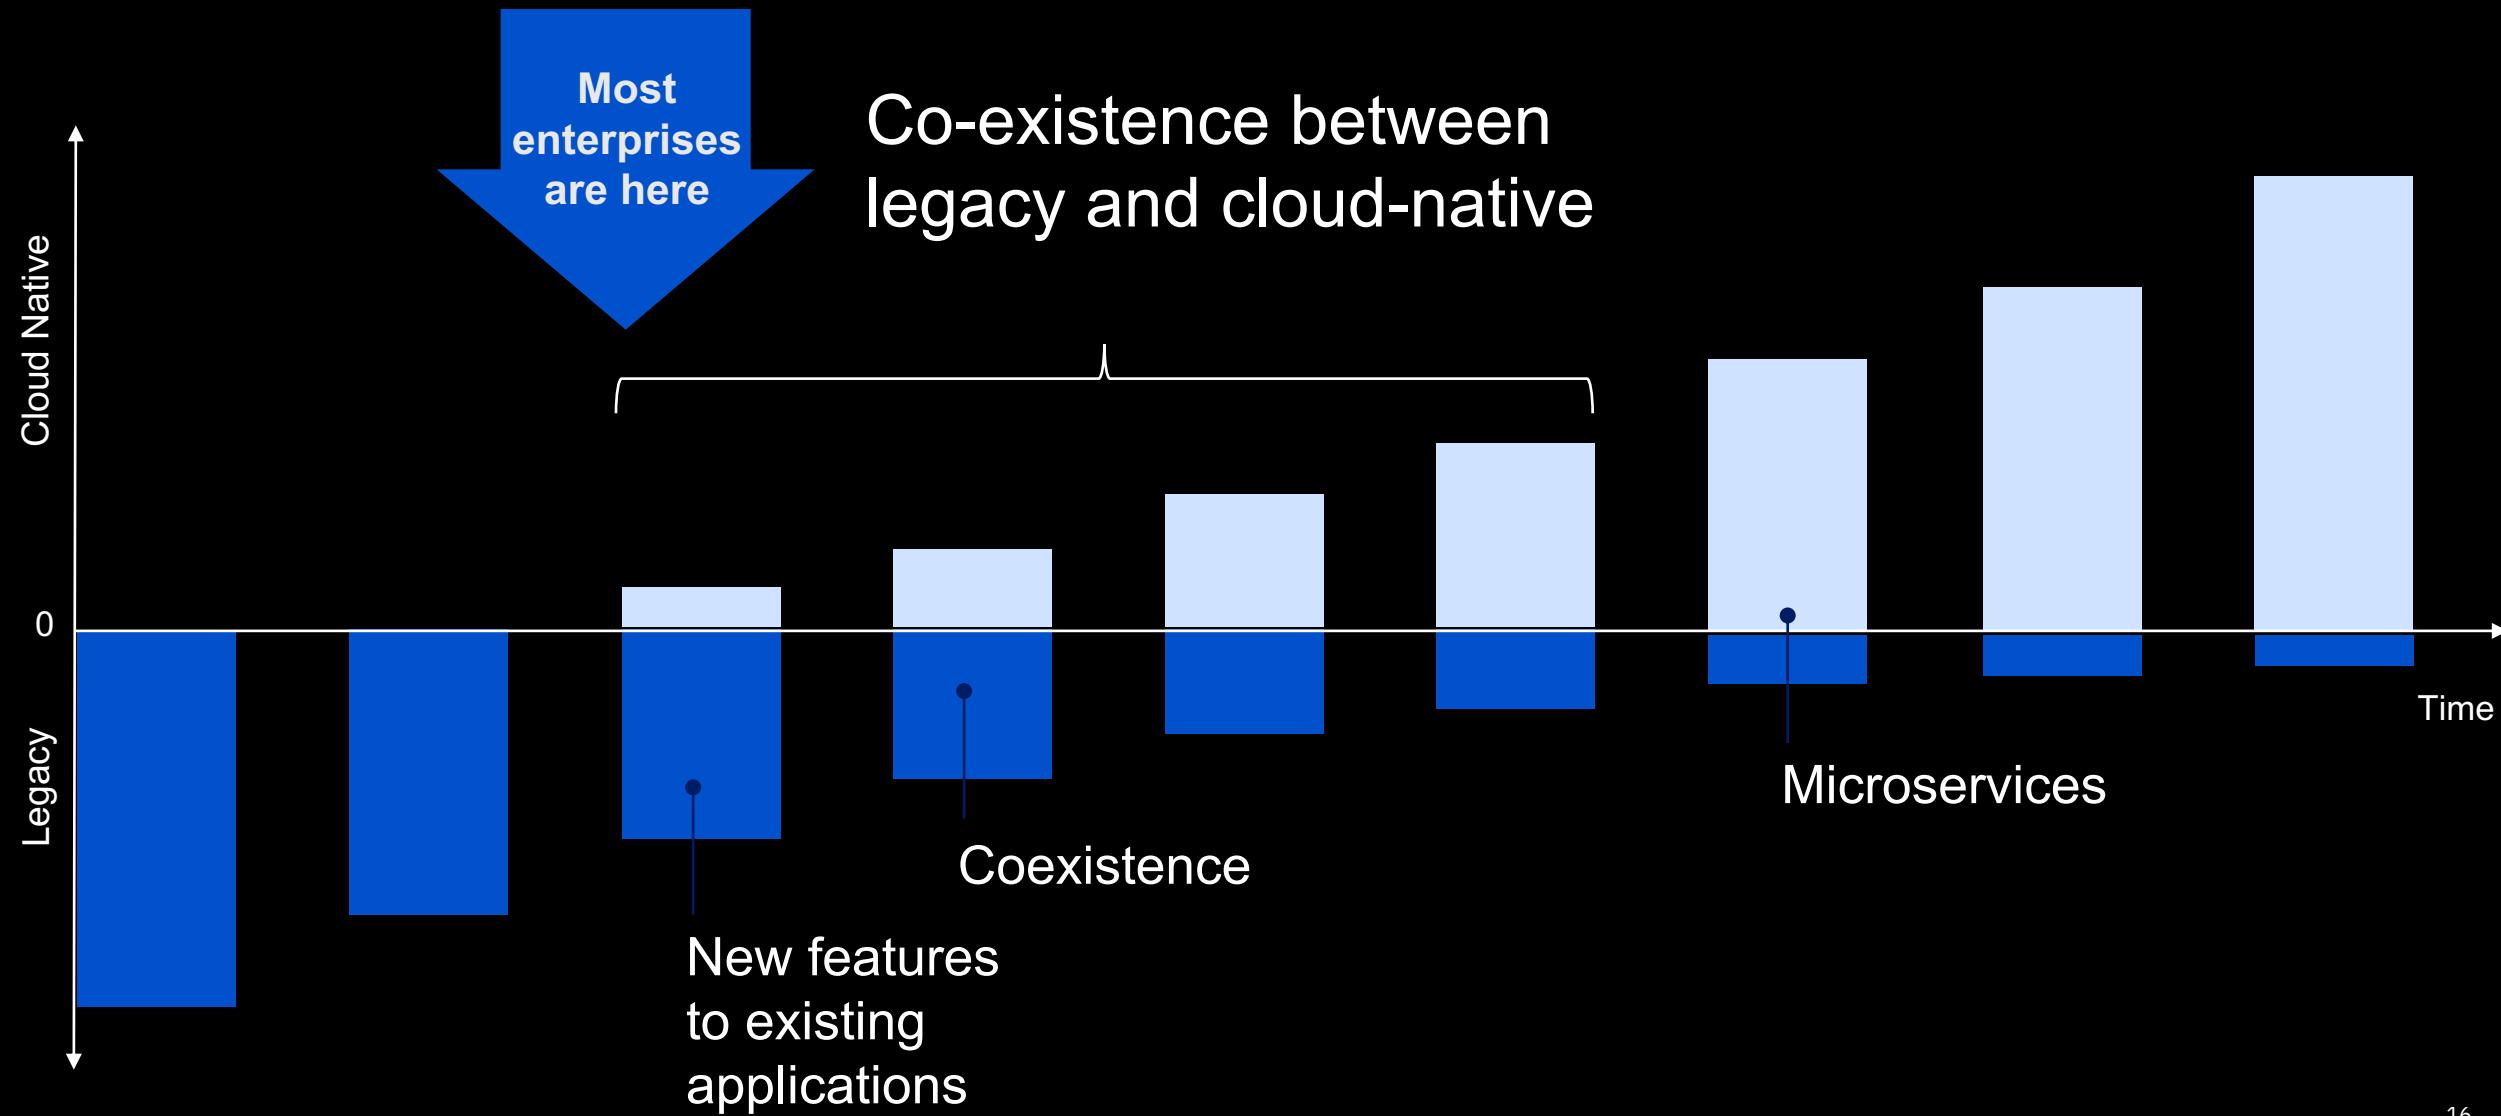
dɪs'ruptʃn/

noun

Business. a **radical change** in an industry, business strategy, etc., especially involving the introduction of a **new product or service** that creates a **new market**

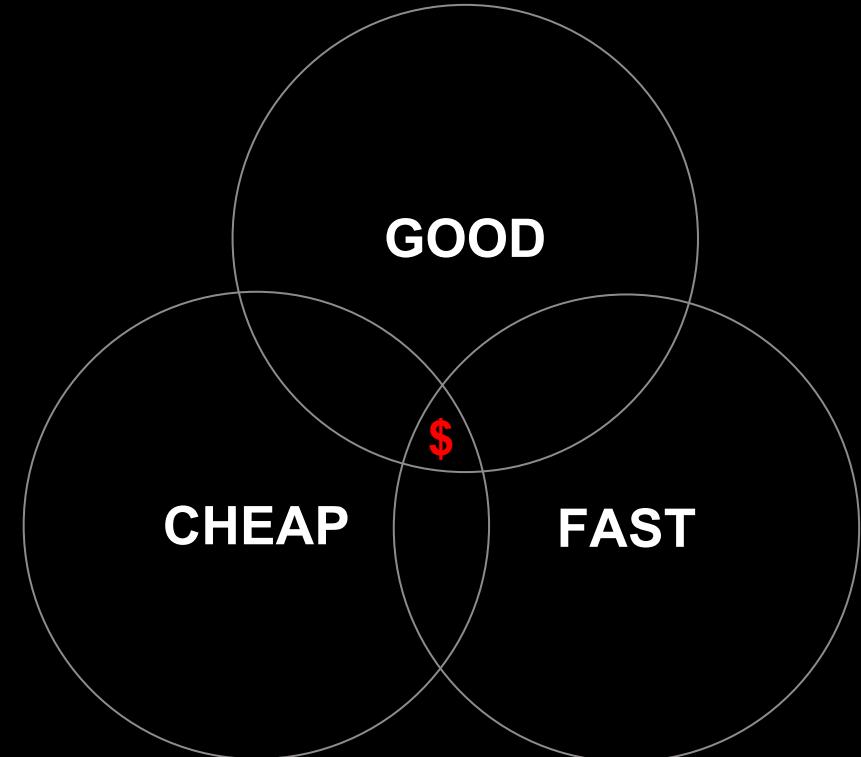
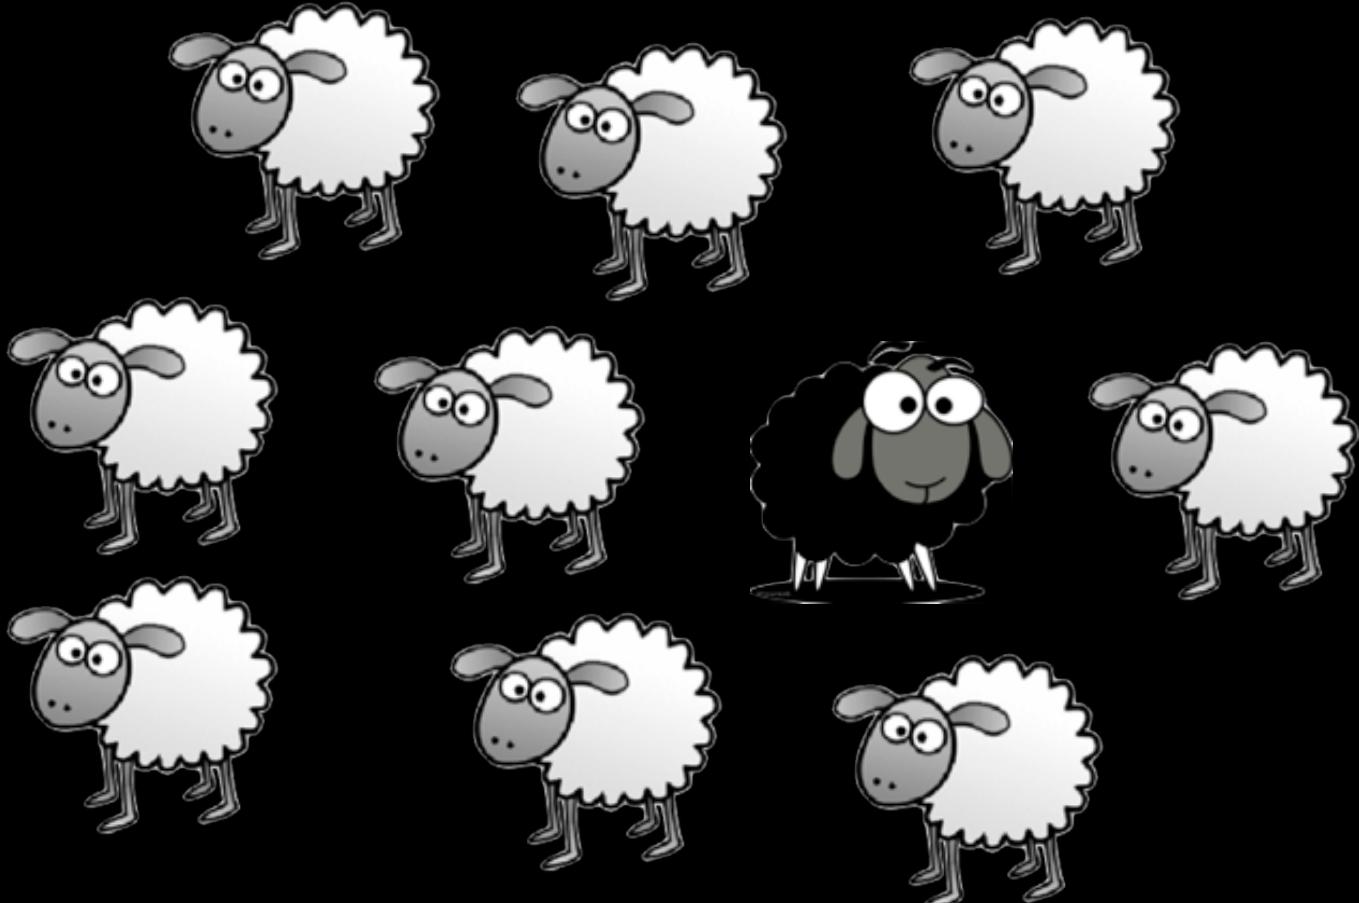
Digital Transformation - Way to go

Cloud native and legacy apps will co-exist for the next 10+ years



Disruption is new reality

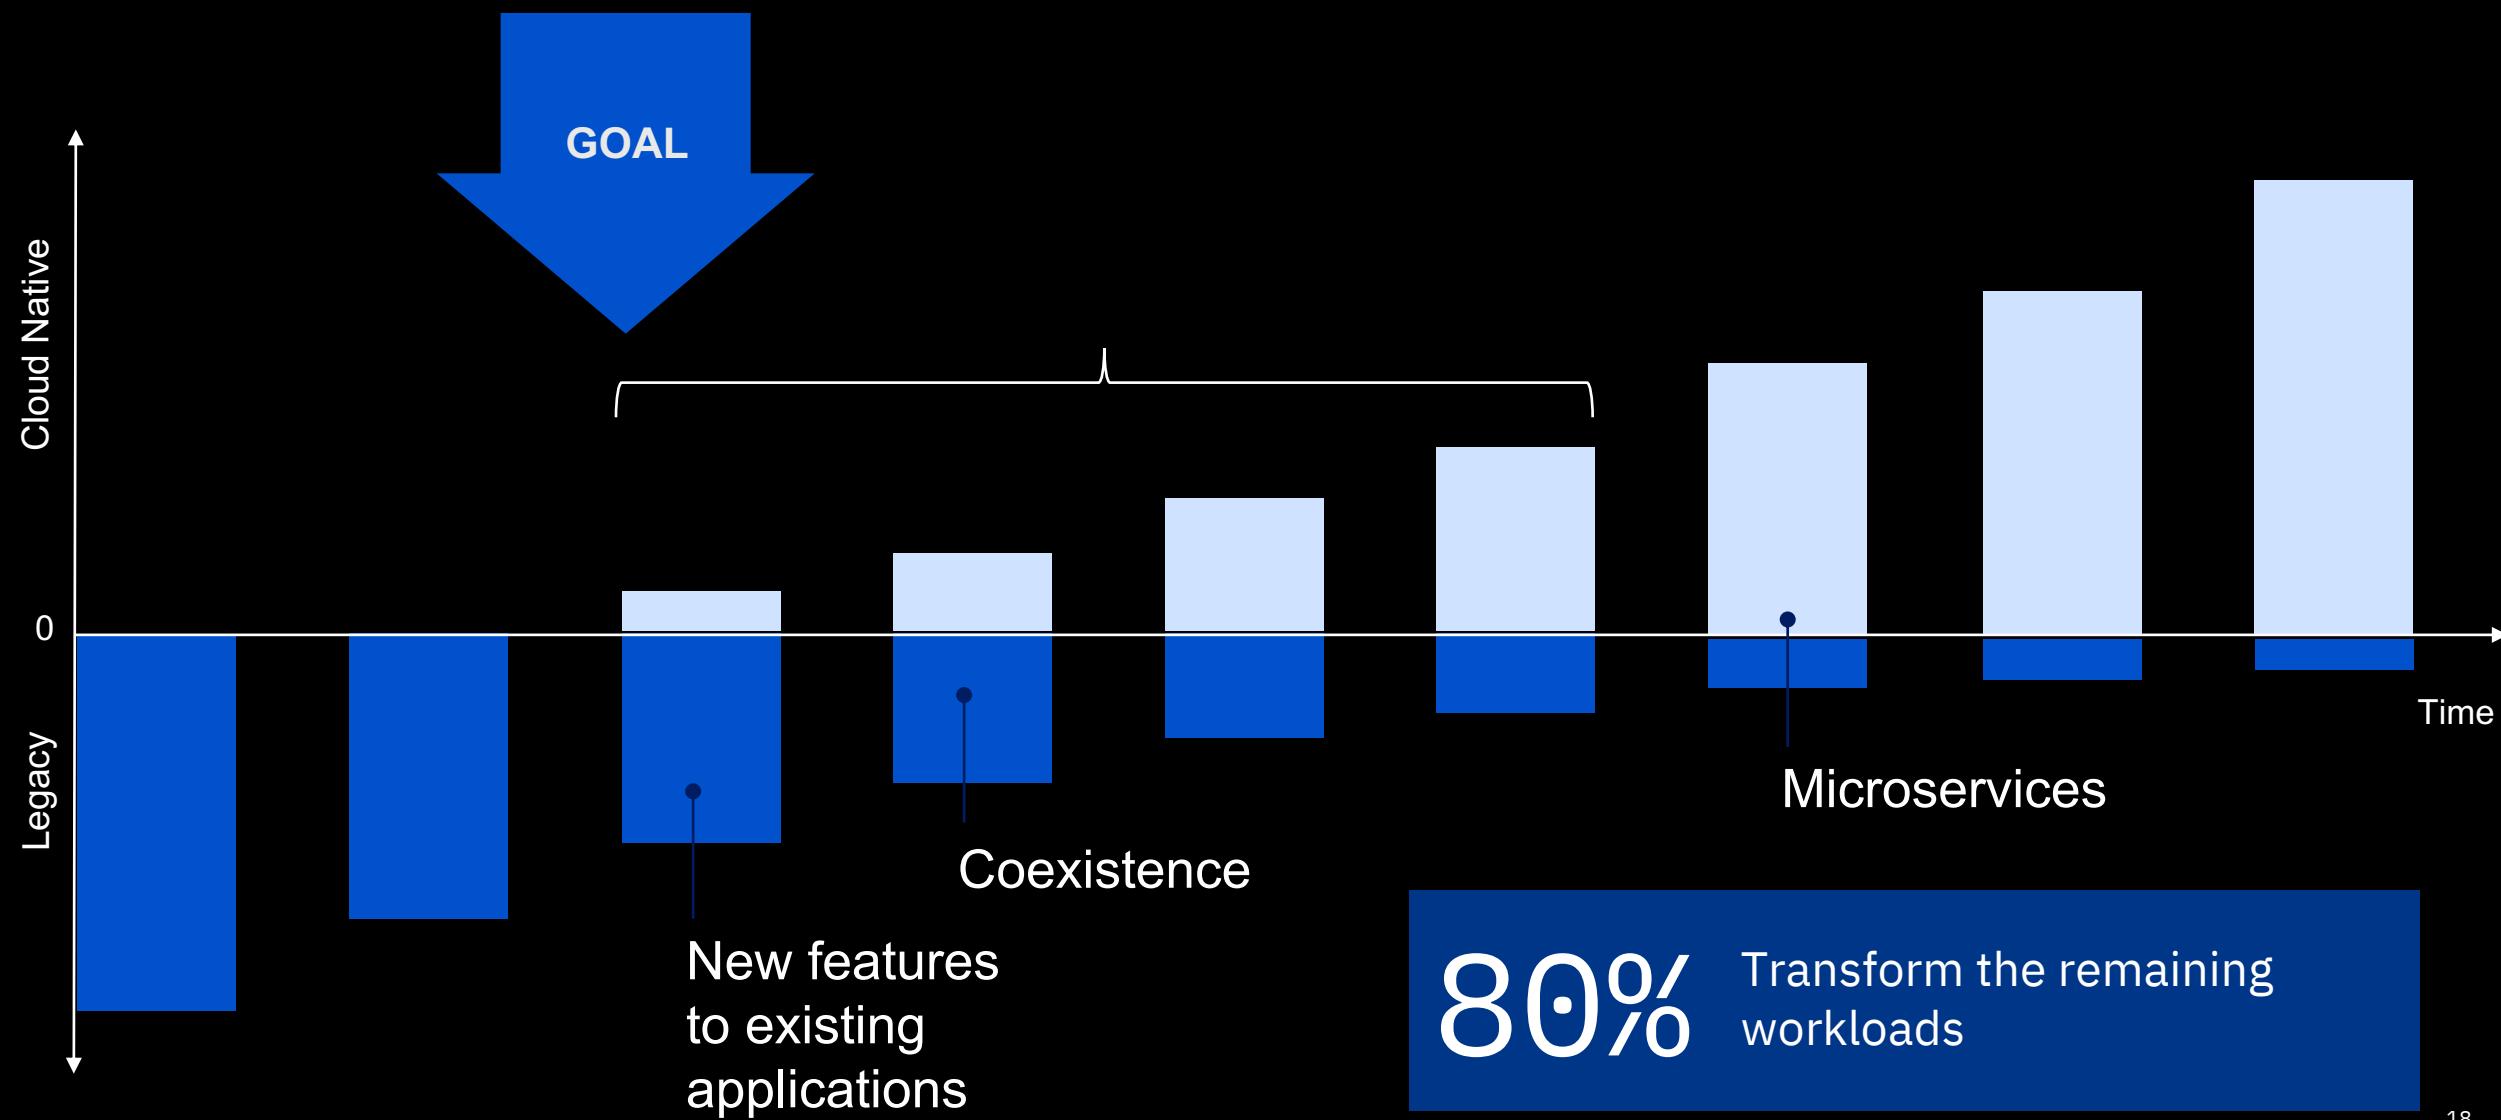
Need to optimize



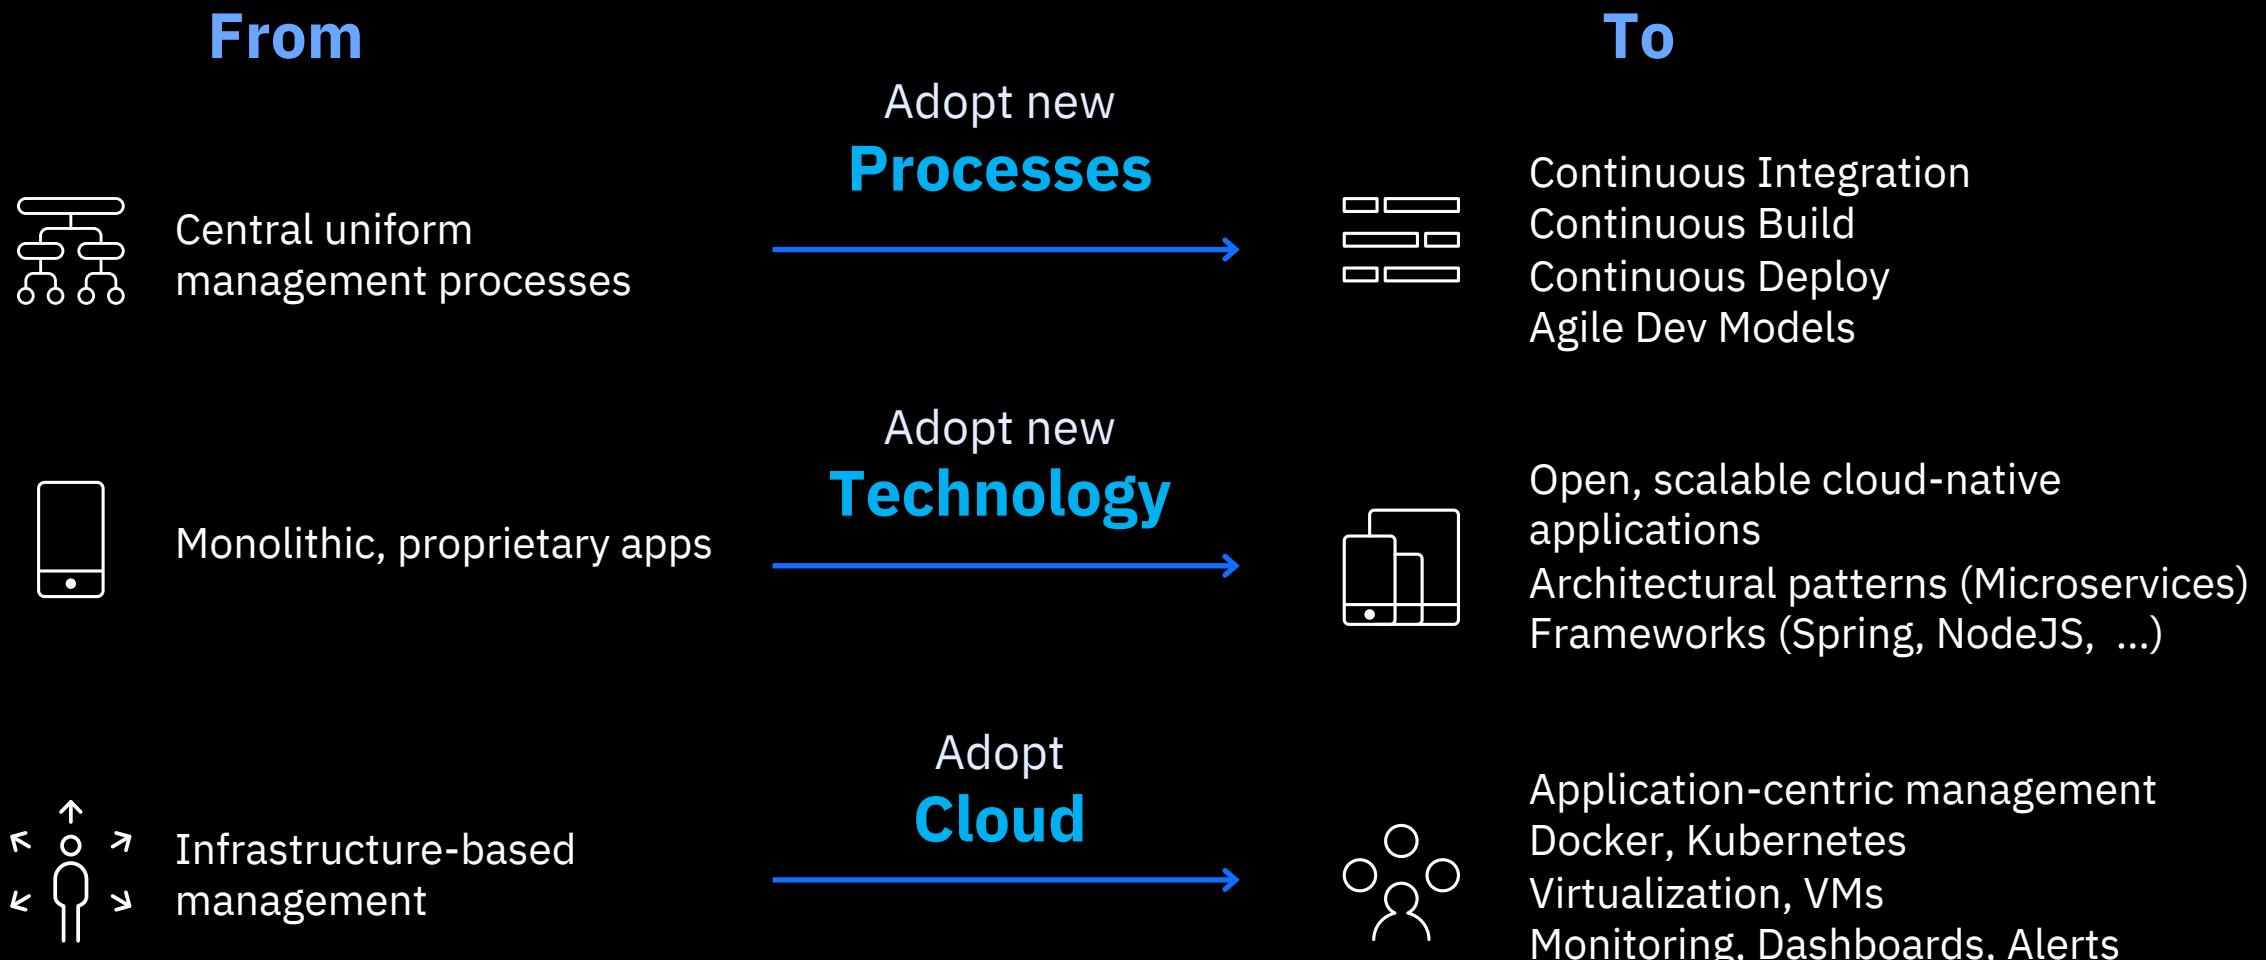
Digital Transformation is the new normal

Digital Transformation

Cloud native and legacy apps will co-exist for the next 10+ years



Digital Transformation – How to get there



DevOps to the rescue



What is DevOps

 just the saddest server
@sadserver

Follow ▾

DevOps is a software engineering culture and practice of putting horrors into containers and then talking about Kubernetes at conferences.

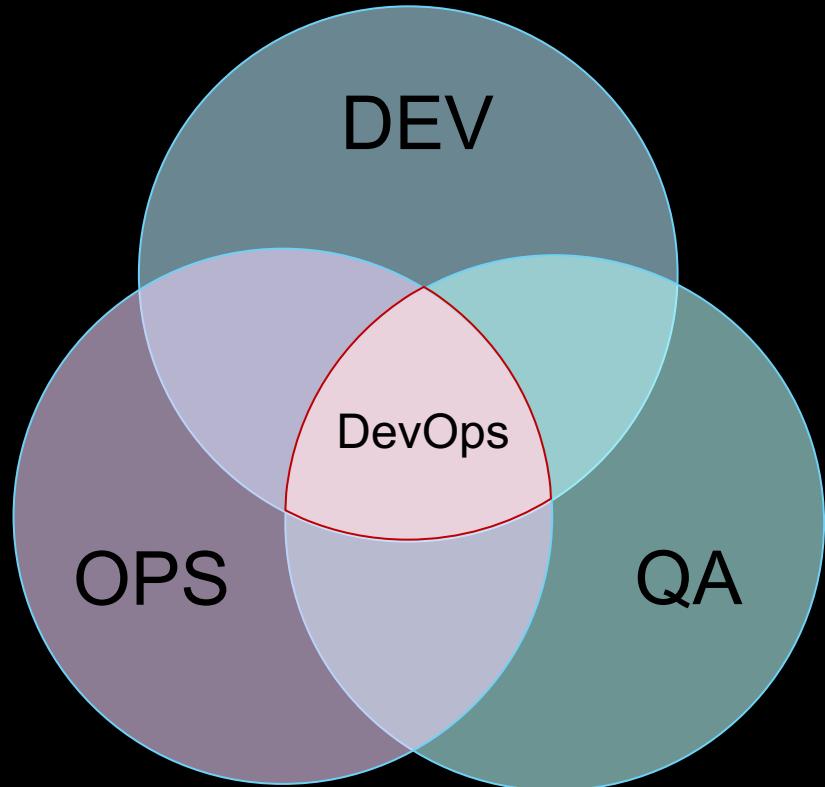
9:48 am - 26 Jun 2018

1,348 Retweets 3,050 Likes



40 1.3K 3.1K

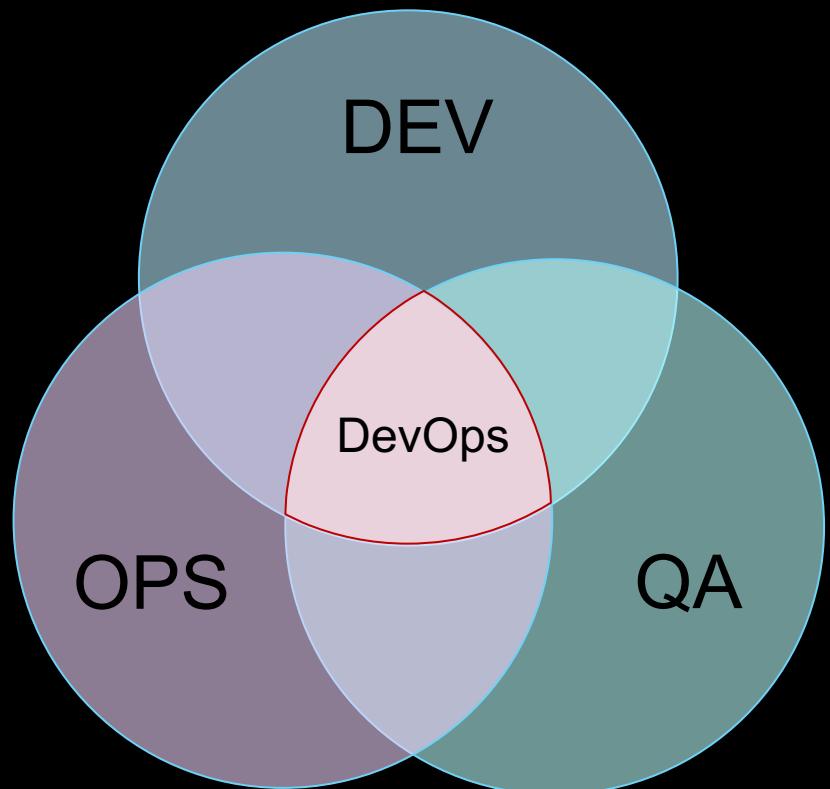
DevOps according to Wikipedia™



DevOps is a software development method that stresses communication, collaboration and integration between software developers and IT professionals.

DevOps is a response to the interdependence of software development and IT operations. It aims to help an organization rapidly produce software products and services.

DevOps according to Wikipedia™



CULTURE

PROCESS

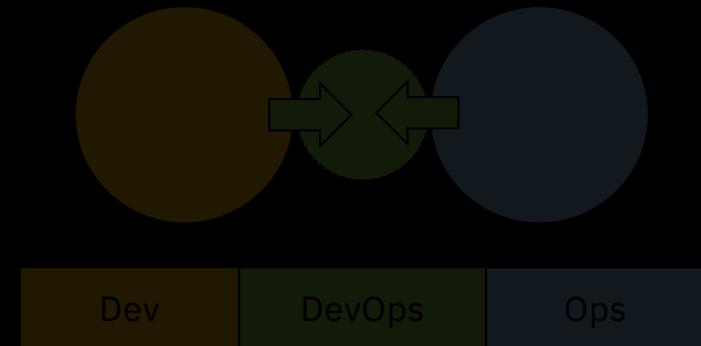
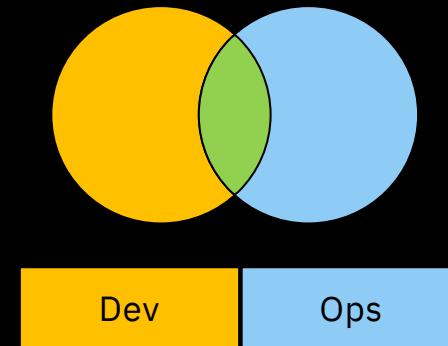
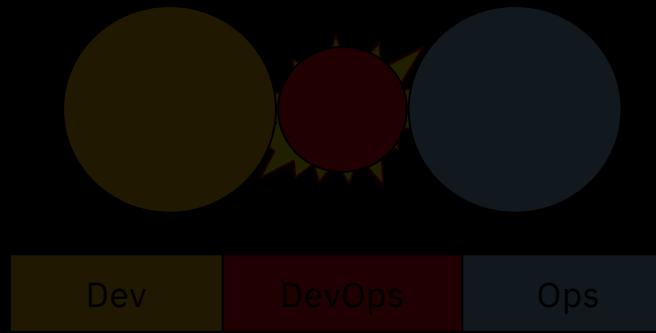
DevOps is a software development **method** that stresses **communication, collaboration** and integration between software **developers** and **IT professionals**

PEOPLE

Tools, Processes
AND **People!**

DevOps - Culture

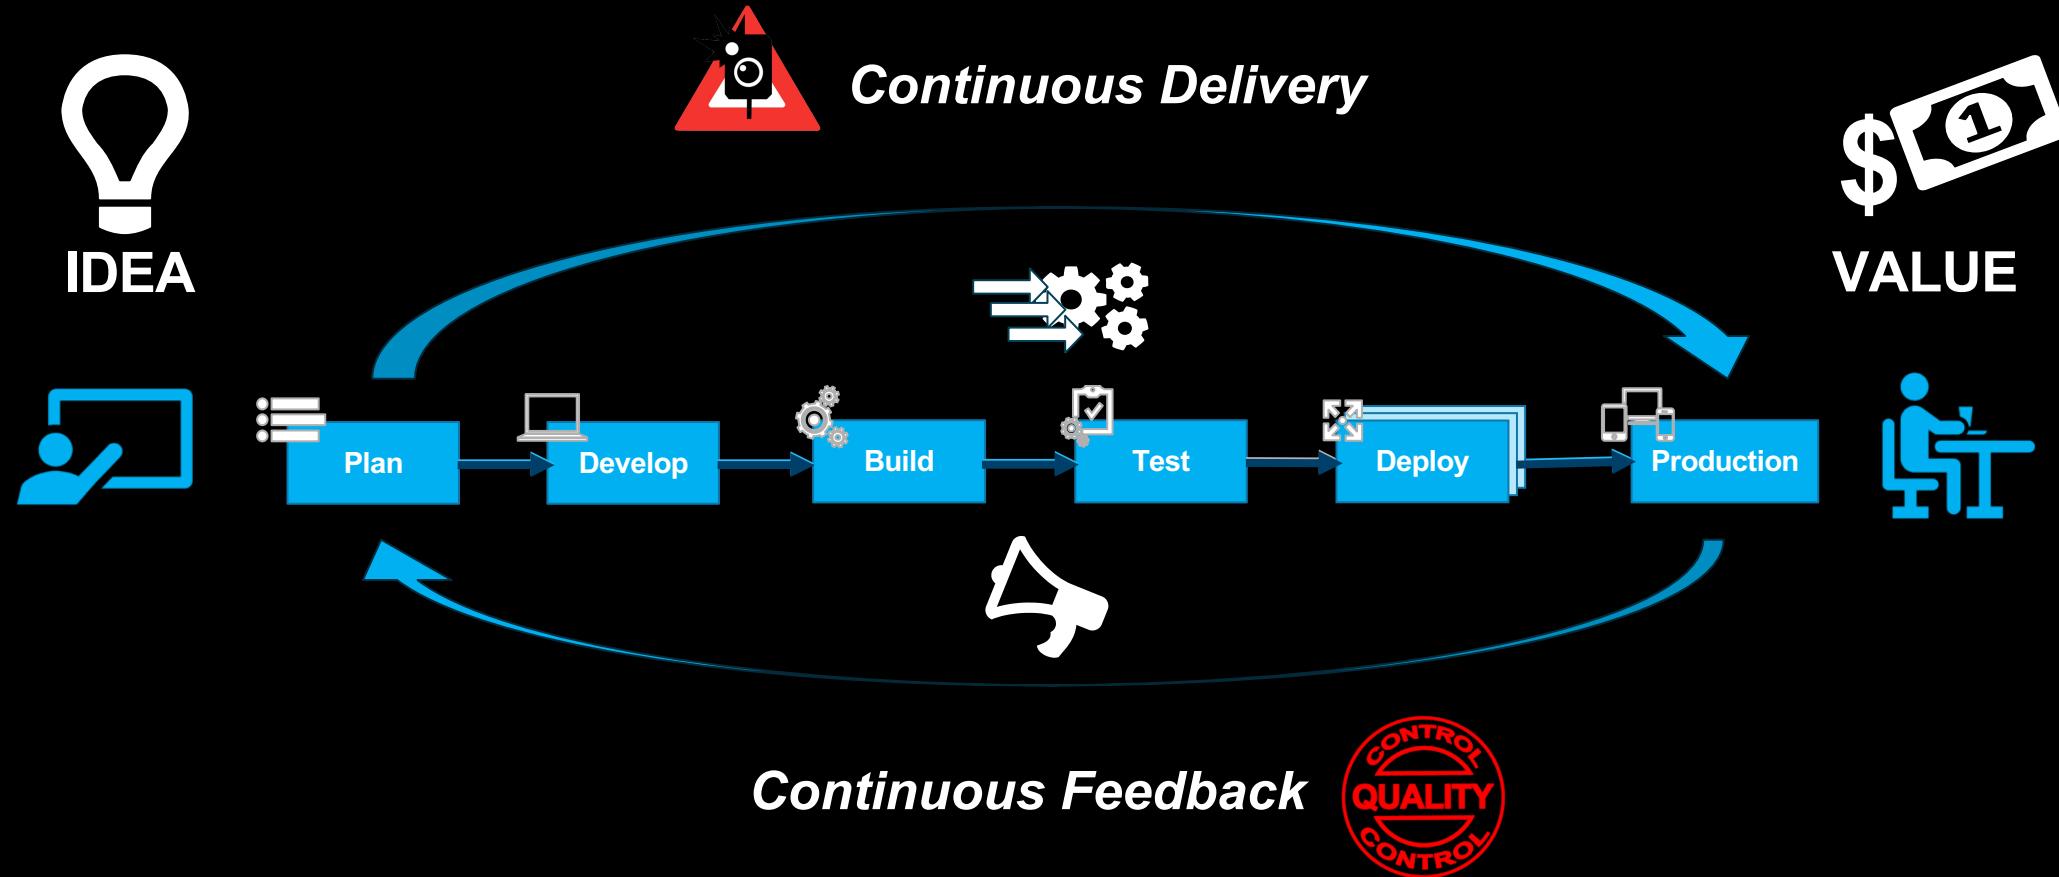
CULTURE 



DevOps – Process

Business Success - get capabilities to the customers

PROCESS ■■■



DevOps – Adoption

What DevOps is NOT

- DevOps is not simply ***combining Development & Operations teams***
- DevOps is not a ***separate team***
- DevOps is not only a ***tool***
- DevOps is not only ***automation***
- DevOps is not a ***one-size-fits-all strategy***

DevOps – Adoption

Getting started on the DevOps journey

- **Be honest**

- **Assess** your own DevOps maturity and aspirations – where are you and where do you want to be?

- **Don't go nuts**

- **Experiment** on a few smaller, low-risk projects – set up a couple of **“two pizza” teams** comprised of Dev/QA + Ops
 - Consider creating a (temporary!) **DevOps "center of excellence"**

- **Cultural change takes time – take reasonable steps**

- Work on **shifting culture** through teambuilding, cross-training, improved communication, perhaps more collaborative tooling

- **Don't try to change the laws of physics**

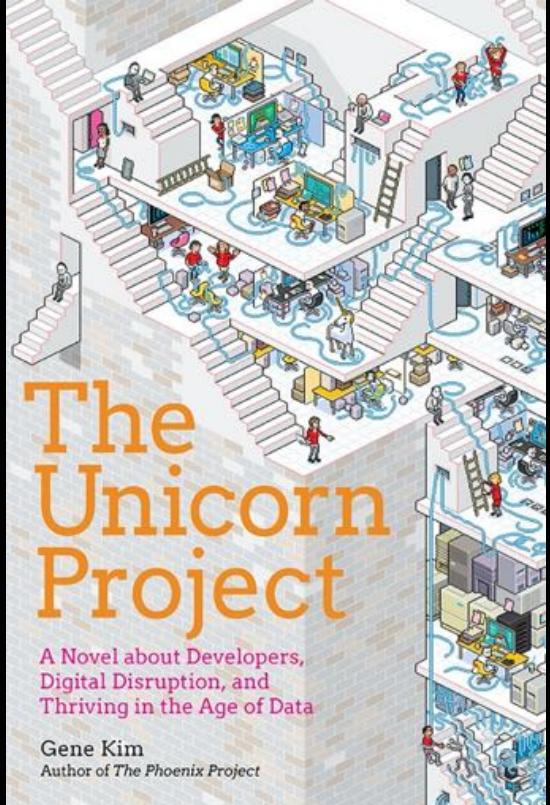
- Reality is that your “Systems of Record” applications move at a **different pace** than your “Systems of Engagement” apps (e.g. mobile)

Remember you must



Start small you should

DevOps – Adoption *Guidance*



A Novel about Developers, Digital Disruption, and Thriving in the Age of Data

DevOps – Adoption

The Five Ideals as of Gene Kim

«We're QA.
We protect the organization from developers.
...
Do not trust them.
Do not get chummy with them.
You give developers an inch, and they'll take a mile.»

from The Unicorn Project by Gene Kim

DevOps – Adoption

The Five Ideals as of Gene Kim

- THE FIRST IDEAL: Locality and Simplicity
- THE SECOND IDEAL: Focus, Flow, and Joy
- THE THIRD IDEAL: Improvement of Daily Work
- THE FOURTH IDEAL: Psychological Safety
- THE FIFTH IDEAL: Customer Focus

«...especially when you have **customers, change is a fact of life.**
A **healthy software** system is one that you can **change at the speed** you need, where people can **contribute easily**, without jumping through hoops.
This is how you make a project that's **fun and worthwhile** contributing to, and where you often find the most vibrant communities.»
from "The Unicorn Project by Gene Kim

DevOps – Adoption

The Five Ideals as of Gene Kim

THE FIRST IDEAL : Locality and Simplicity

Locality in Systems

- Teams need to be able to build and test their changes locally and deploy them without implicating 10 other teams

Organization

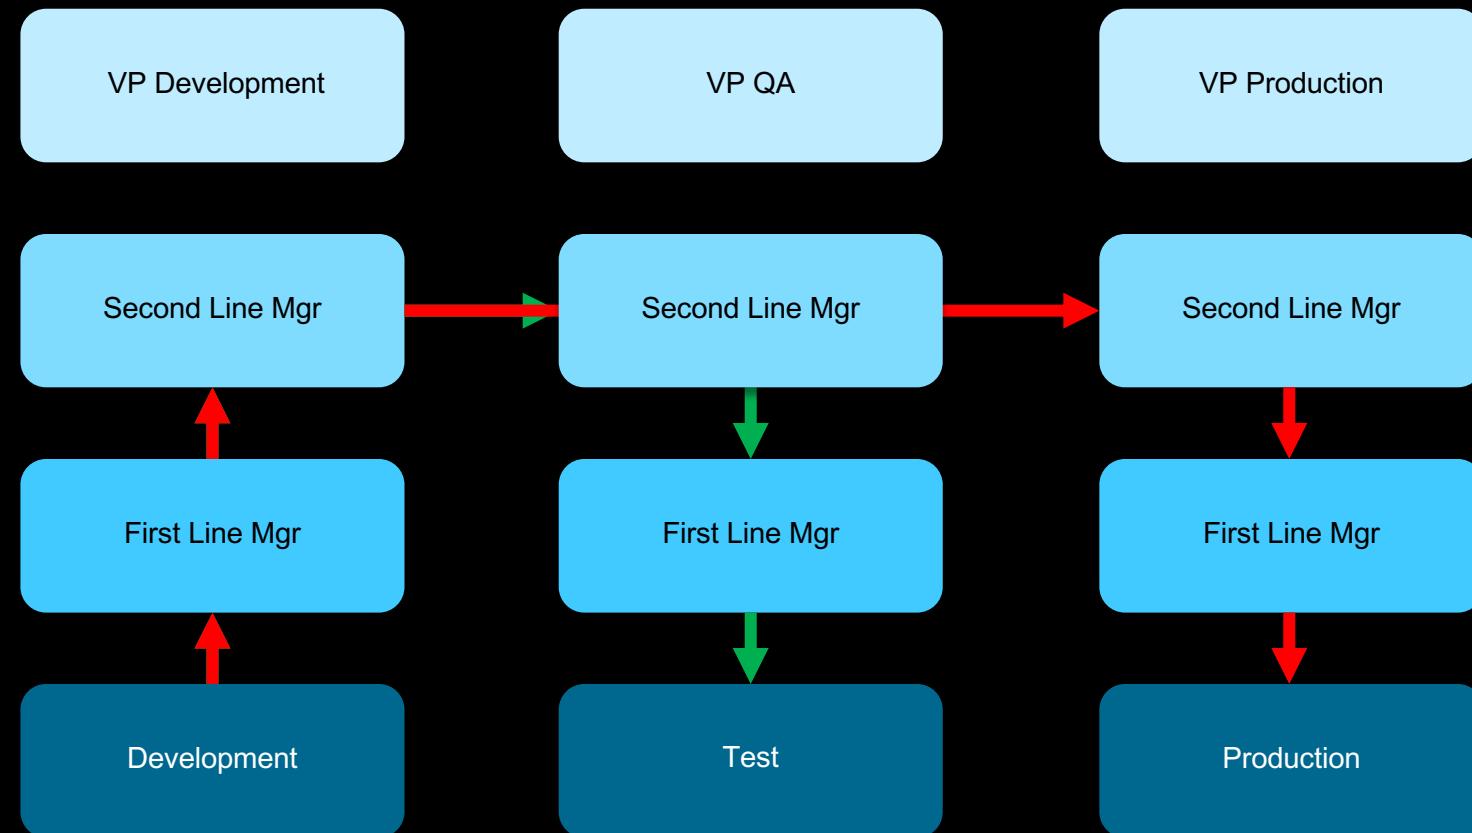
- Every team has the expertise, capability and authority to satisfy customer needs

Simplicity

- Uncouple applications from each other
- Architectural patterns
- Prioritize technical debt reduction
- Two Pizza teams

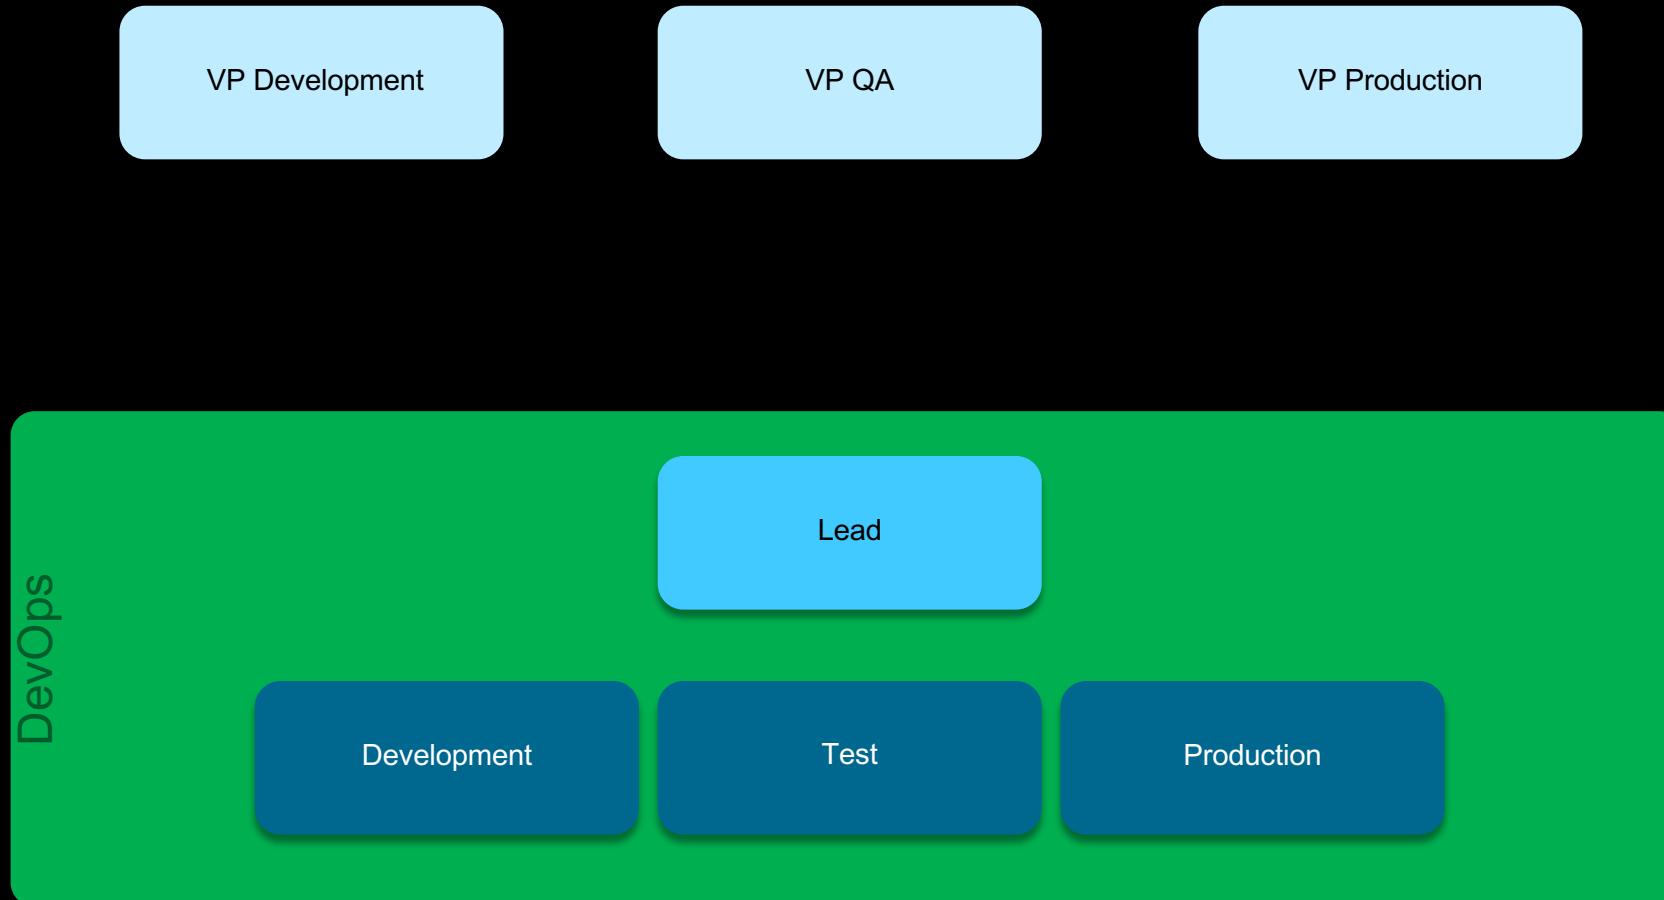
DevOps – Adoption

The Five Ideals as of Gene Kim



DevOps – Adoption

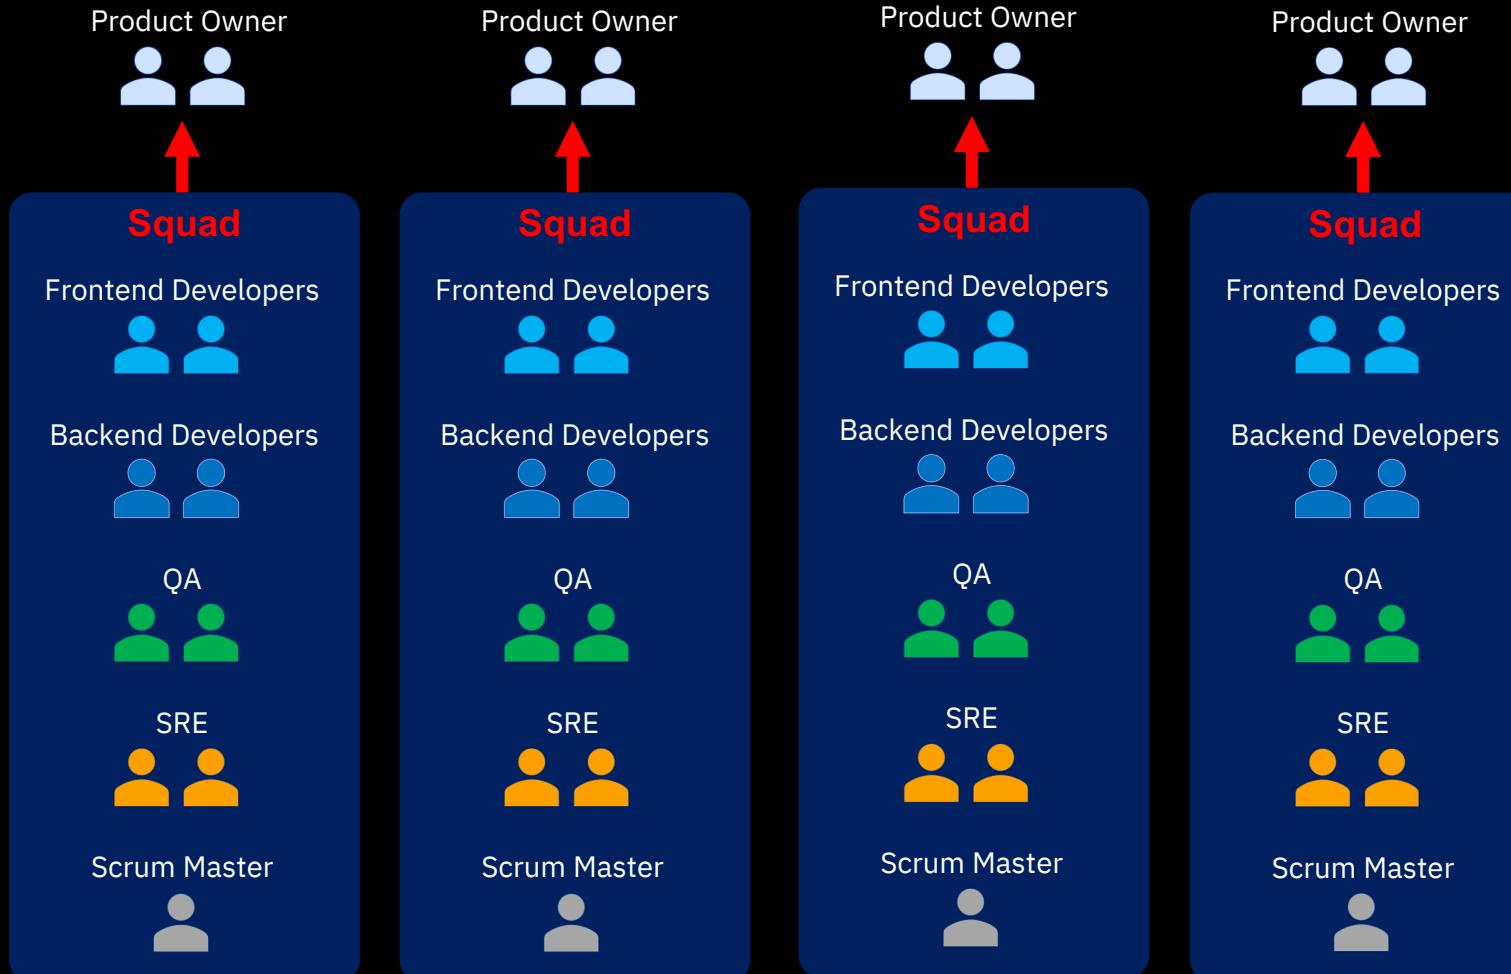
The Five Ideals as of Gene Kim



DevOps – Adoption

The Five Ideals as of Gene Kim

Agile Team Organisation: Squads



Squads contains more than just developers – remove „stigma“

→ rename „Teams“ to „Squads“ 😊 or Crew, Party, ...

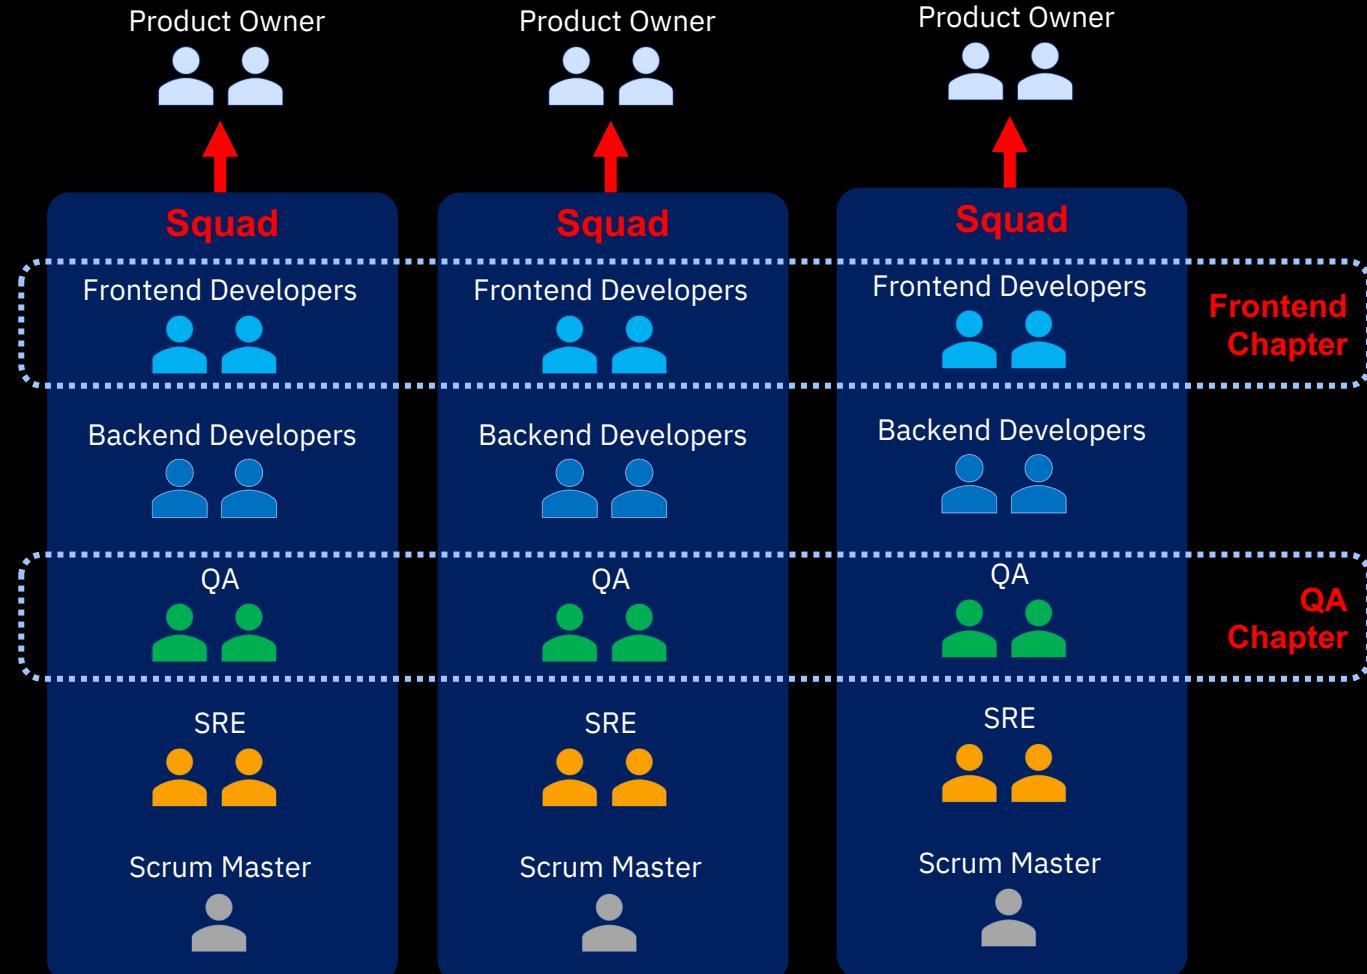
Fully autonomous, cross functional team that has full responsibilities for an application component/functionality/microservice and little to no dependencies on others

5-7 people ideally

DevOps – Adoption

The Five Ideals as of Gene Kim

Agile Team Organisation: Chapters



Chapters are a group or team members working within a special area

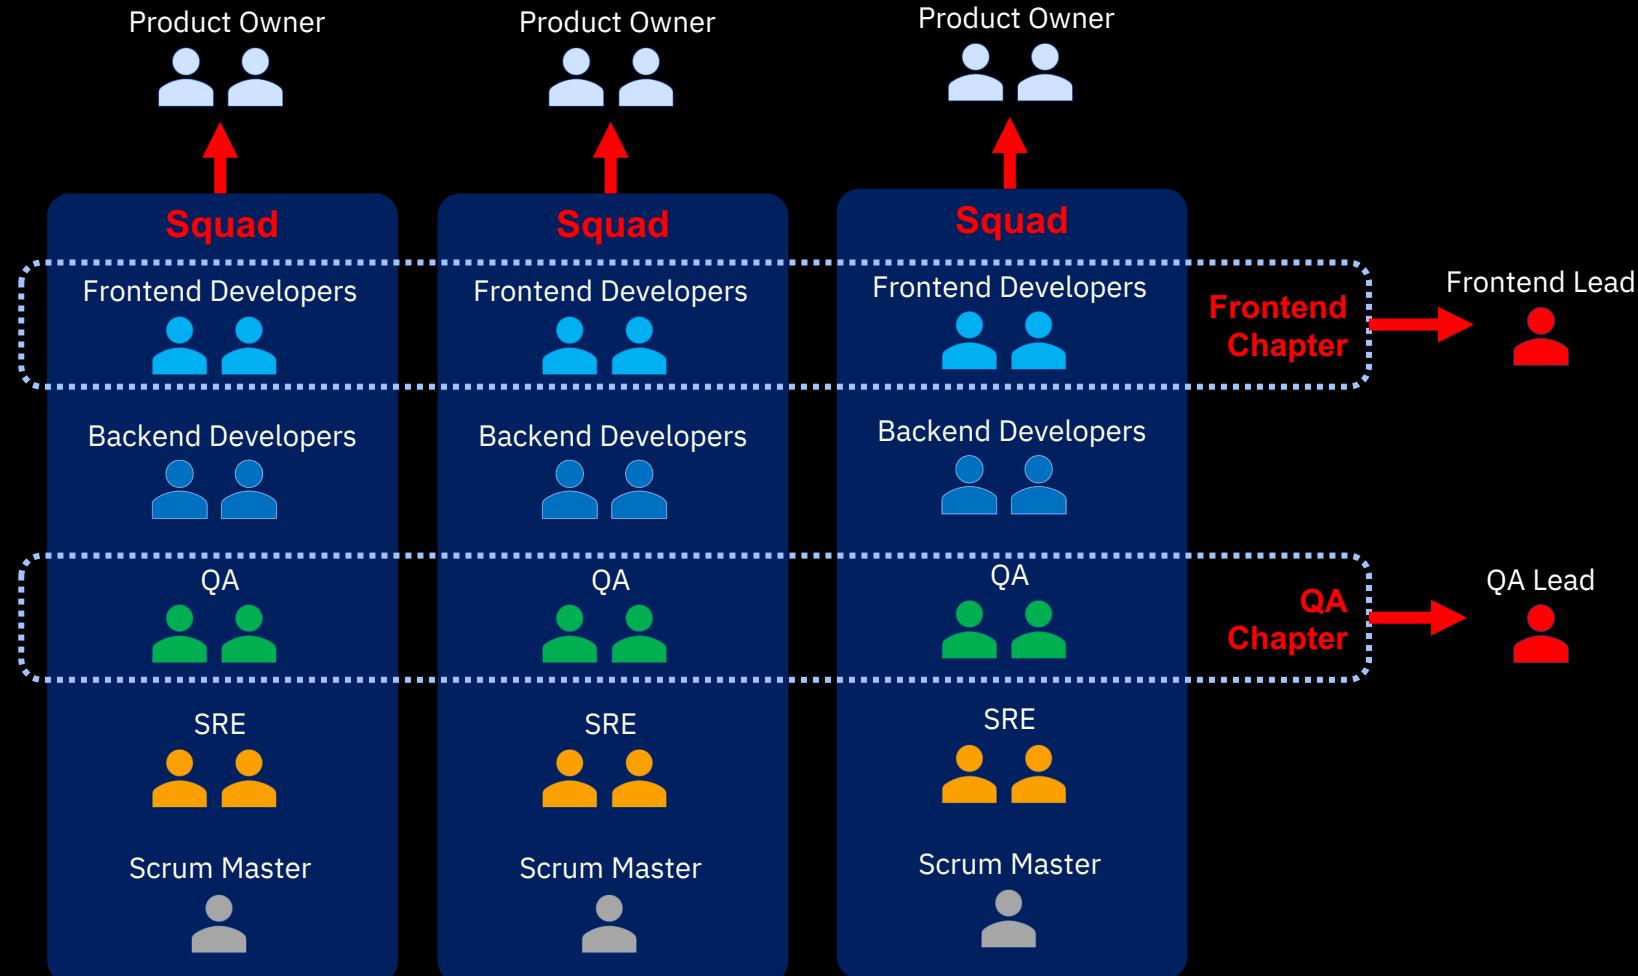
Relay experience to the rest of the chapter and discuss on how other squads can use it.

Promotes team collaboration and innovation.

DevOps – Adoption

The Five Ideals as of Gene Kim

Agile Team Organisation: Chapters



A **Chapter Lead** is the line manager for chapter members.

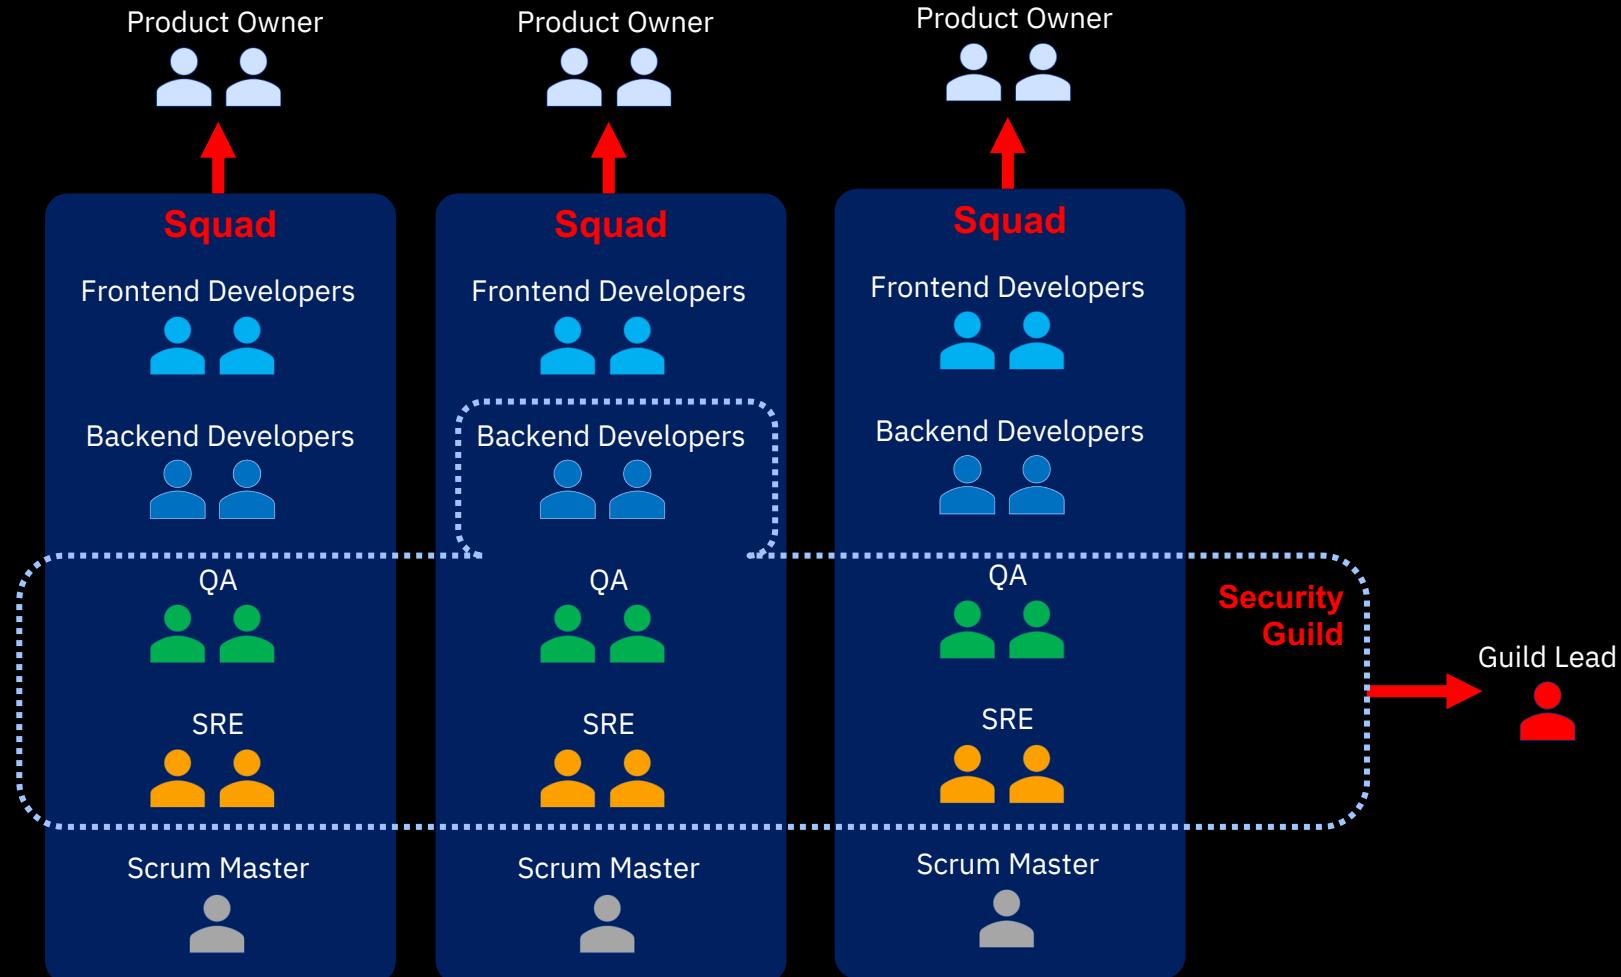
Responsible for developing people .

Remain part of a squad and still do day-to-day work.

DevOps – Adoption

The Five Ideals as of Gene Kim

Agile Team Organisation: Guilds



Guilds cover a problem or something that can be improved and get together a small group of team members together to solve it.

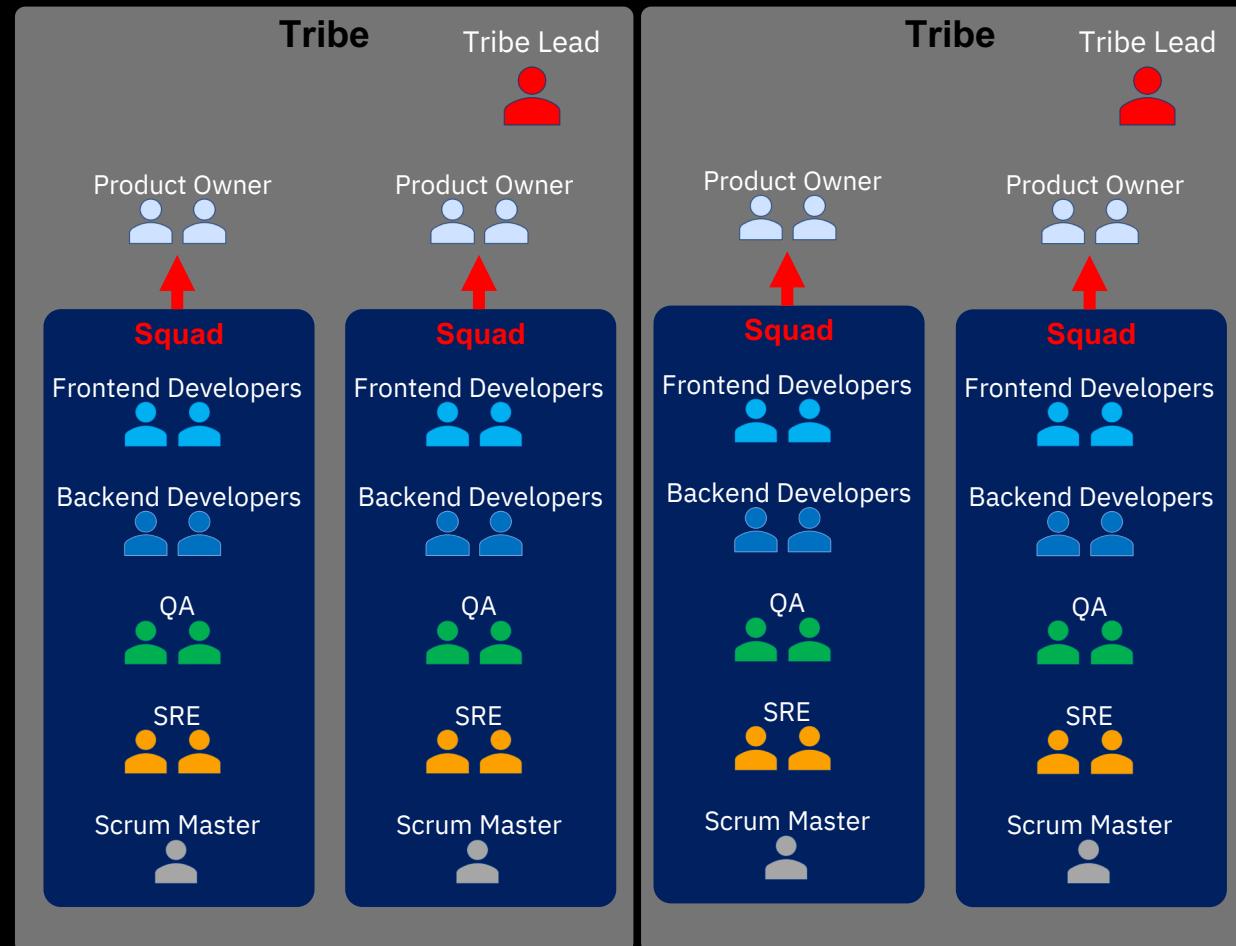
Examples:

- Performance monitoring / optimization
- Testing automation
- Security & vulnerabilities
- Services & Architecture

DevOps – Adoption

The Five Ideals as of Gene Kim

Agile Team Organisation: Tribes



Tribes are a collection of squads that work in related areas like mobile app, backend infrastructure, ...

“Incubator” for the **Squad** mini-startups

Challenging to implement

Only for large complicated infrastructure that needs to be split this way

< 100 people (Dunbar number)

DevOps – Adoption

The Five Ideals as of Gene Kim

THE SECOND IDEAL: Focus, Flow, and Joy

Focus

- Focus on developing their code with minimum dependencies, delays and impediments

Flow

- Focus creates flow of value, therefore joy

Joy

- Energy and time is focused on solving the business problem, and you're having fun

DevOps – Adoption

The Five Ideals as of Gene Kim

Enable developer productivity

- Self-service
- On-demand
- Immediacy and fast feedback
- Focus and flow

Core concepts

- Immutability
- Composability

Look at...

- Microservices
- Containers
- Kubernetes

DevOps – Adoption

The Five Ideals as of Gene Kim

THE THIRD IDEAL: Improvement of Daily Work

Technical Debt

- Pay down technical debt

Architecture

- Enables teams to independently develop, test, and deploy value to customers without being coupled to many other teams

Flow

- Flow with as little interruption and impediments as possible

DevOps – Adoption

The Five Ideals as of Gene Kim

«**Technical Debt** is what you feel the next time
you want to make a change.»

Usually refers to things we need to **clean up**, or where we need to
create or restore simplicity, so that that we can **quickly, confidently,**
and safely make changes to the system

Ward Cunningham in 2003

DevOps – Adoption

The Five Ideals as of Gene Kim

THE FOURTH IDEAL: Psychological Safety

Safety

- Team members feel safe to talk about problems

Blame

- Blameless post-mortems

Feedback

- Create feedback in the system, sooner and faster, with clarity between cause and effect.

DevOps – Culture

Start by understanding your Flow

| Pathological Power-oriented | Bureaucratic Rule-oriented | Generative Performance-oriented |
|---|---|---|
| Low co-operation | Modest co-operation | High co-operation |
| Messengers shot | Messengers neglected | Messengers trained |
| Responsibilities shirked | Narrow responsibilities | Risks are shared |
| Bridging discouraged | Bridging tolerated | Bridging encouraged |
| Failure leads to scapegoating | Failure leads to justice | Failure leads to inquiry |
| Novelty crushed | Novelty leads to problems | Novelty implemented |
| Characterized by large amounts of fear and threat. People often hoard information or withhold it for political reasons, or distort it to make themselves look better. | Organisations protect departments. Those in the department want to maintain their “turf,” insist on their own rules, and generally do things by the book — their book | Organisations focus on the mission. How do we accomplish our goal? Everything is subordinated to good performance, to doing what we are supposed to do. |

DevOps – Adoption

The Five Ideals as of Gene Kim

THE FIFTH IDEAL: Customer Focus

Customer Focus

- Customer focus relates to the difference between **core** and **context**
- Decisions based on **what the customer values**
- Make sure that context doesn't kill core

Core - differentiation

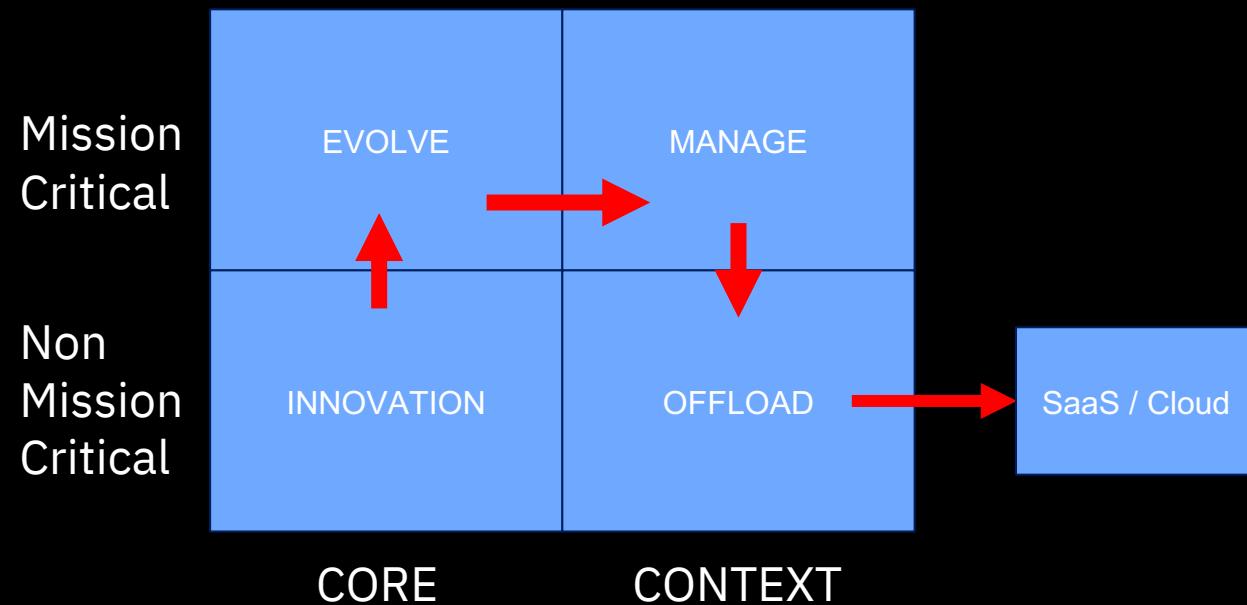
- Lasting durable **business advantage**
- What customers are **willing to pay for**

Context - commodity

- Systems that are mission critical, but **don't create competitive advantage**
- What customers don't want / see, like HR systems, payroll, ...

DevOps – Adoption

The Five Ideals as of Gene Kim



New Products start in the bottom left quadrant to create differentiation for customers but have no customer impact initially.

New products will grow in profitability, becoming mission critical.

Customer needs evolve, products can migrate out of the Core to become commodity because competitors have caught-up.

At the end of life they might become non-critical to the business and can be offloaded.

DevOps – Adoption

Measuring Adoption

| | Elite | High | Medium | Low |
|---|-------------|----------------|------------------|--------------------|
| Lead time for changes | | | | |
| The time it takes to go from code committed to code successfully deployed | 1 hour | 1 day – 1 week | 1 week – 1 month | 1 month – 6 months |
| Deployment frequency | | | | |
| Frequency of <i>production</i> deployments - how often you're delivering something of <i>value</i> to end users | > 1 per day | ~1 per day | 1 week | 1 month |
| Time to restore service | | | | |
| The average time it takes to restore a service (MTTR) | < 1 hour | < 1 day | 1 day – 1 week | > 1 week |
| Change failure rate | | | | |
| How often deployment failures occur in production that require i | 0-15% | 0-15% | 0-15% | 46-60% |

... predictive of software delivery, operational performance, and organizational performance, and they correlate with burnout, employee engagement, and so much more."

DevOps – Adoption

Culture



Culture

- Empower the team - Foster Experimentation
- Generative Actor - "Rogue/TwoPizza Team" - Showcase
- Change takes Time - take reasonable steps
- If something is hard, do it repeatedly



Patterns / Antipatterns

- There should be no (permanent) DevOps Team
- Everyone is DevOps
- Development Executives own their solutions
- Be honest - Assess maturity and aspirations



Management buy-in

- DevOps is about solving a business problem
- DevOps starts at the top – Senior Leadership Team
- Optimize the whole and not the individual 'silos'

Process



Steer

- Continuous Business Planning
- Measure customer value
- Link requirements to releases



Develop/Build

- Continuous Integration
- Agile Methods / SAFe
- Link lifecycle information



Test

- Continuous Testing
- Test Automation
- Test/Service Virtualization
- Make Build Self-Testing



Deploy

- Continuous Deployment
- Automation and Infrastructure Patterns
- Provide Self-Service
- Containers



- Continuous Feedback
- Monitor resources
- Optimize to customer KPIs continuously
- Automate problem isolation and resolution

Tools



Steer

- Rational Doors Next Generation
- cloudMatrix
- Jira
- Rational Team Concert
- Rational Focal Point



Develop/Build

- Rational Team Concert
- IBM MobileFirst Platform
- Git
- Jenkins / Bamboo
- Jira



Test

- Rational Quality Manager
- Rational Test Workbench
- Rational Test Virtualization Server
- IBM Security AppScan
- Selenium



Deploy

- IBM UrbanCode Deploy
- IBM Cloud Orchestrator
- Terraform
- Chef Scripts
- Puppet Scripts

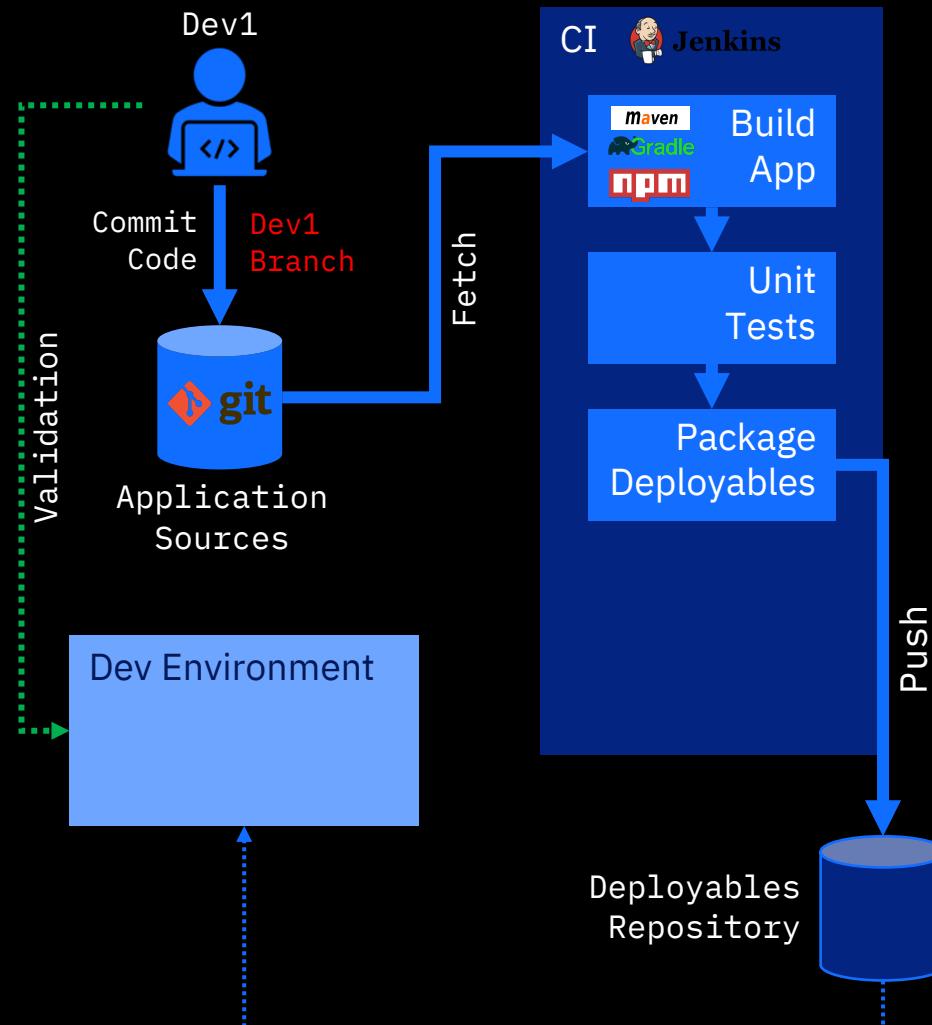


Operate

- Tealeaf Customer Behavior Analysis Suite
- IBM Control Desk
- IBM Application Performance Monitoring

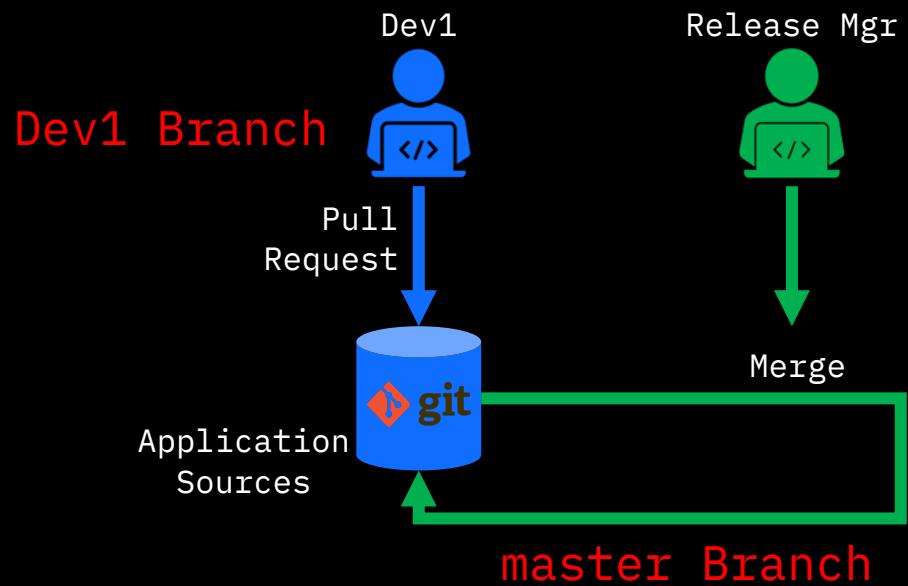
DevOps – CI/CD Pipeline

Development – Dev1 Branch



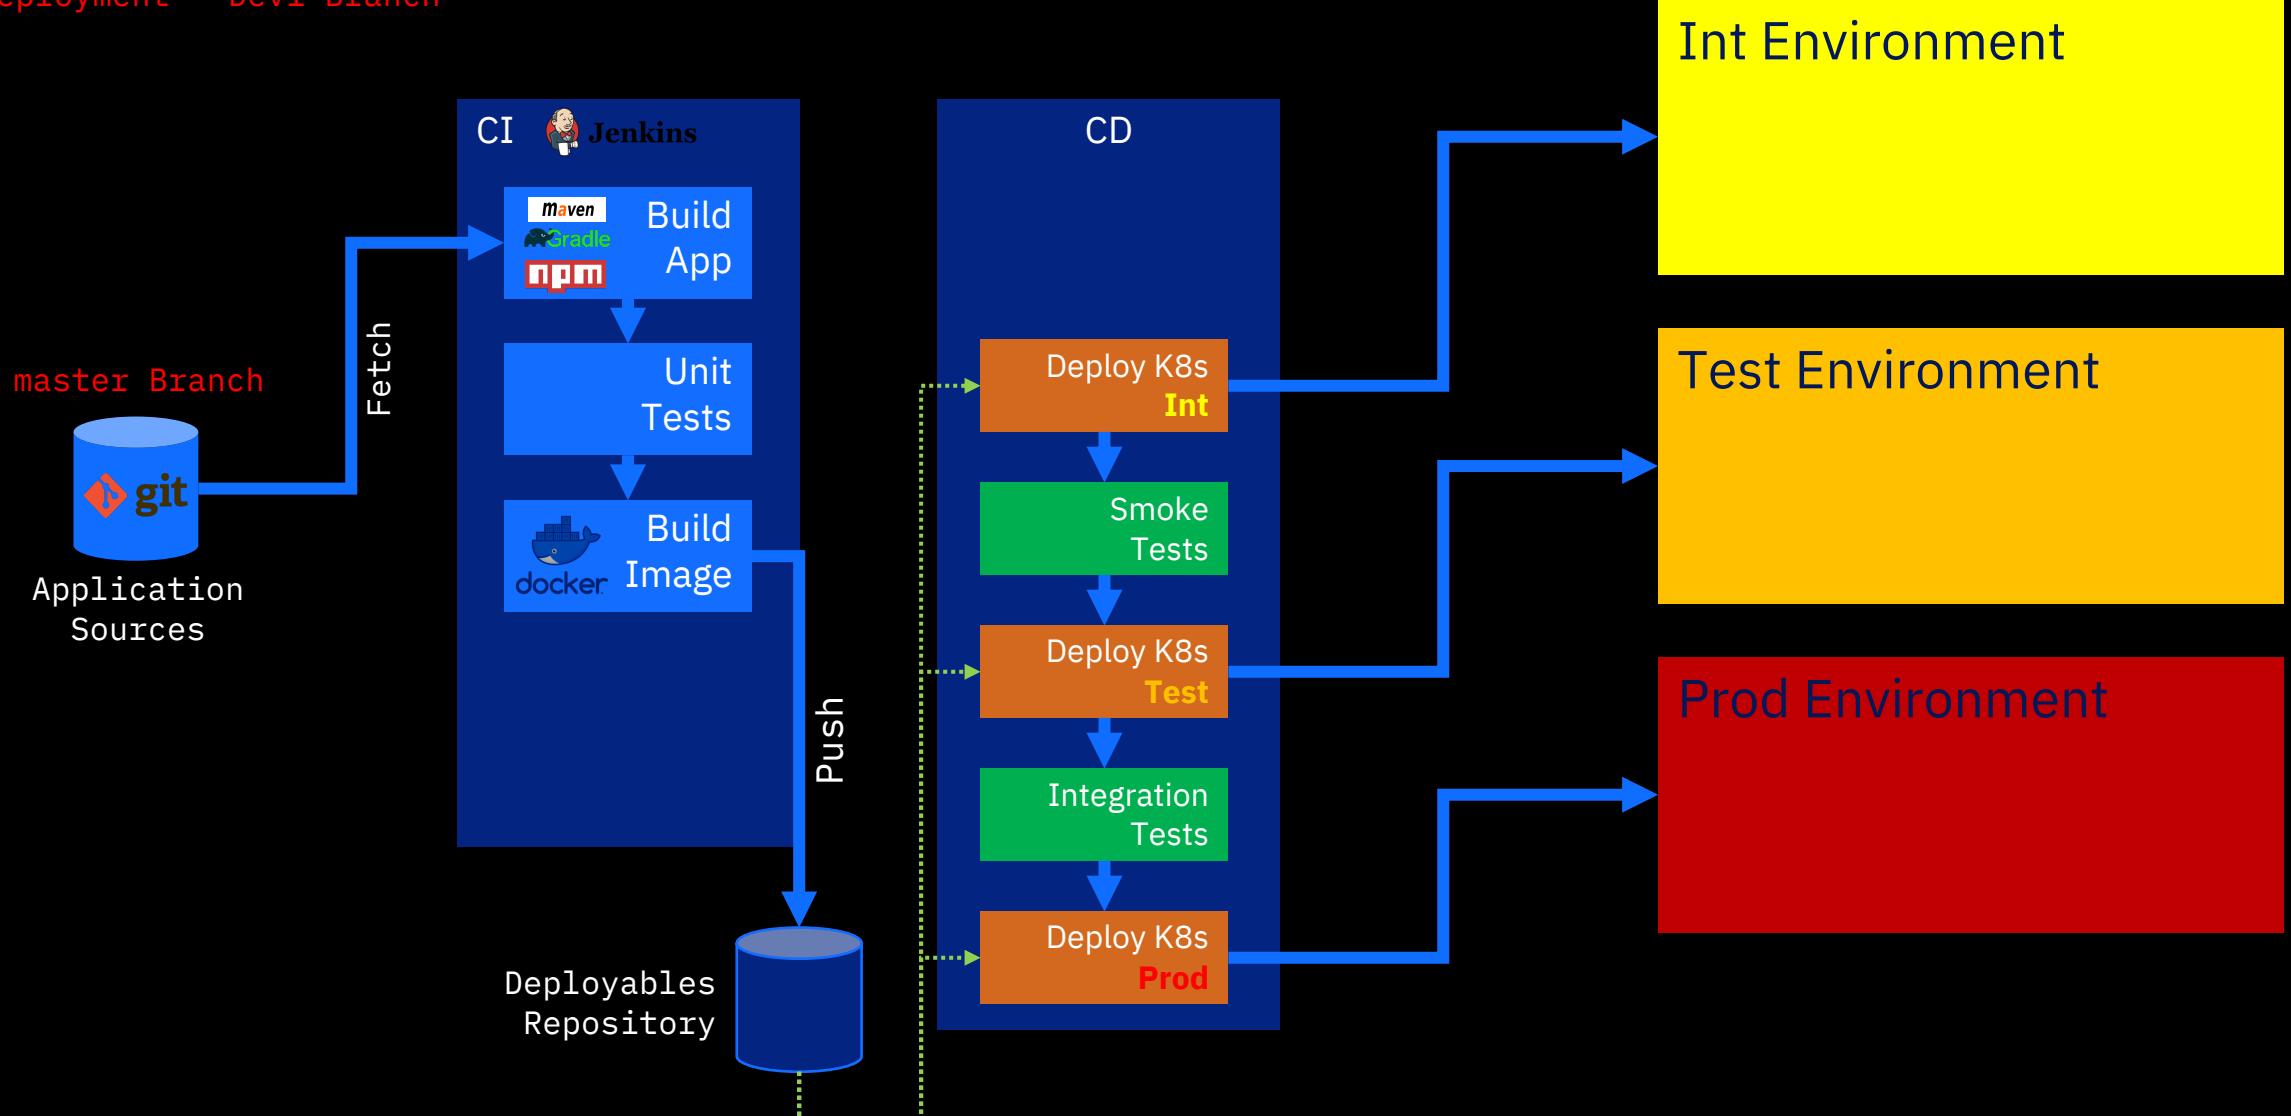
DevOps – CI/CD Pipeline

Merge – Dev1 Branch to master Branch



DevOps – CI/CD Pipeline

Deployment – Dev1 Branch



DevOps – **Git**Ops

GitOps builds on DevOps with **Git as a single source of truth** for declarative infrastructure and applications - the whole system.

1. Having a **completely automated delivery pipeline** that can roll out changes to your infrastructure when changes are made to Git.
2. Operating a fast paced business 24/7 requires **monitoring and observability** baked into the beginning. **Security** is of critical importance.
3. Everything has to be **version controlled** and stored in a single source of truth from which you can recover.

DevOps – **Git**Ops

Full Audit Trail

GitOps is using git, which is an excellent database to persist the changes made to the system.

Everything as a code

In the cloud, both CI and CD should be defined as code first. The desired state of environments should also be set in full as a declarative code.

...

Today, application security is mostly an **afterthought**, that is run in pre-production and often perceived as a **roadblock** to staying ahead of the competition.

Examples

- Getting access to a Git repository (IAM)
- Open firewall ports
- Getting credentials for an API
- Getting certificates

Dev**Sec**Ops

Putting Security in your DevOps transformation

DevSecOps means thinking about application and infrastructure security from the start.

1. The idea is to get **security** back in to the **lifecycle** → shifting security left
2. The goal is to make **security** as **silent** and **seamless** as possible
3. «Everyone is responsible for security» **mindset**
 - Developers write secure code
 - SREs secure the environments

But we still need clear **responsibilities!**

Most organizations have clear **governance of risk**, and out of that we derive **security policies**.

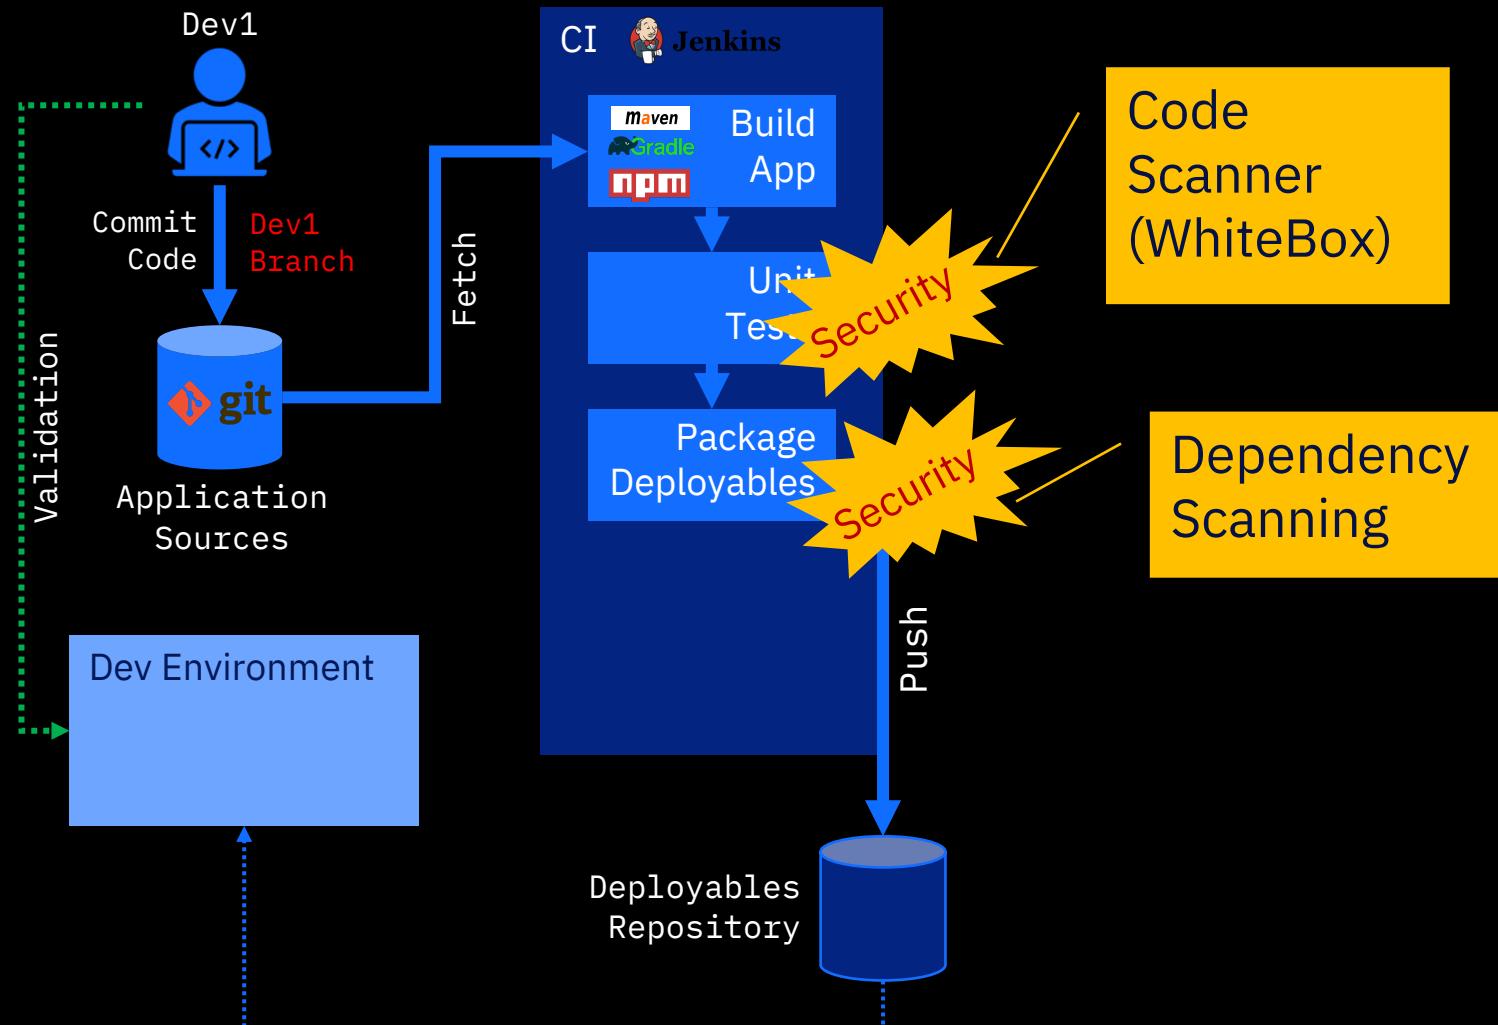
In order to bring Sec into DevOps is to **bake** those **policies** into the processes.

But we still need clear responsibilities!

For example, when you are writing code,
do the developers have a tool that checks for vulnerabilities during the local build
process at **build time**
or do you do this within your **CI process** using Jenkins?
or do you do this within your **CD process** before deployment?

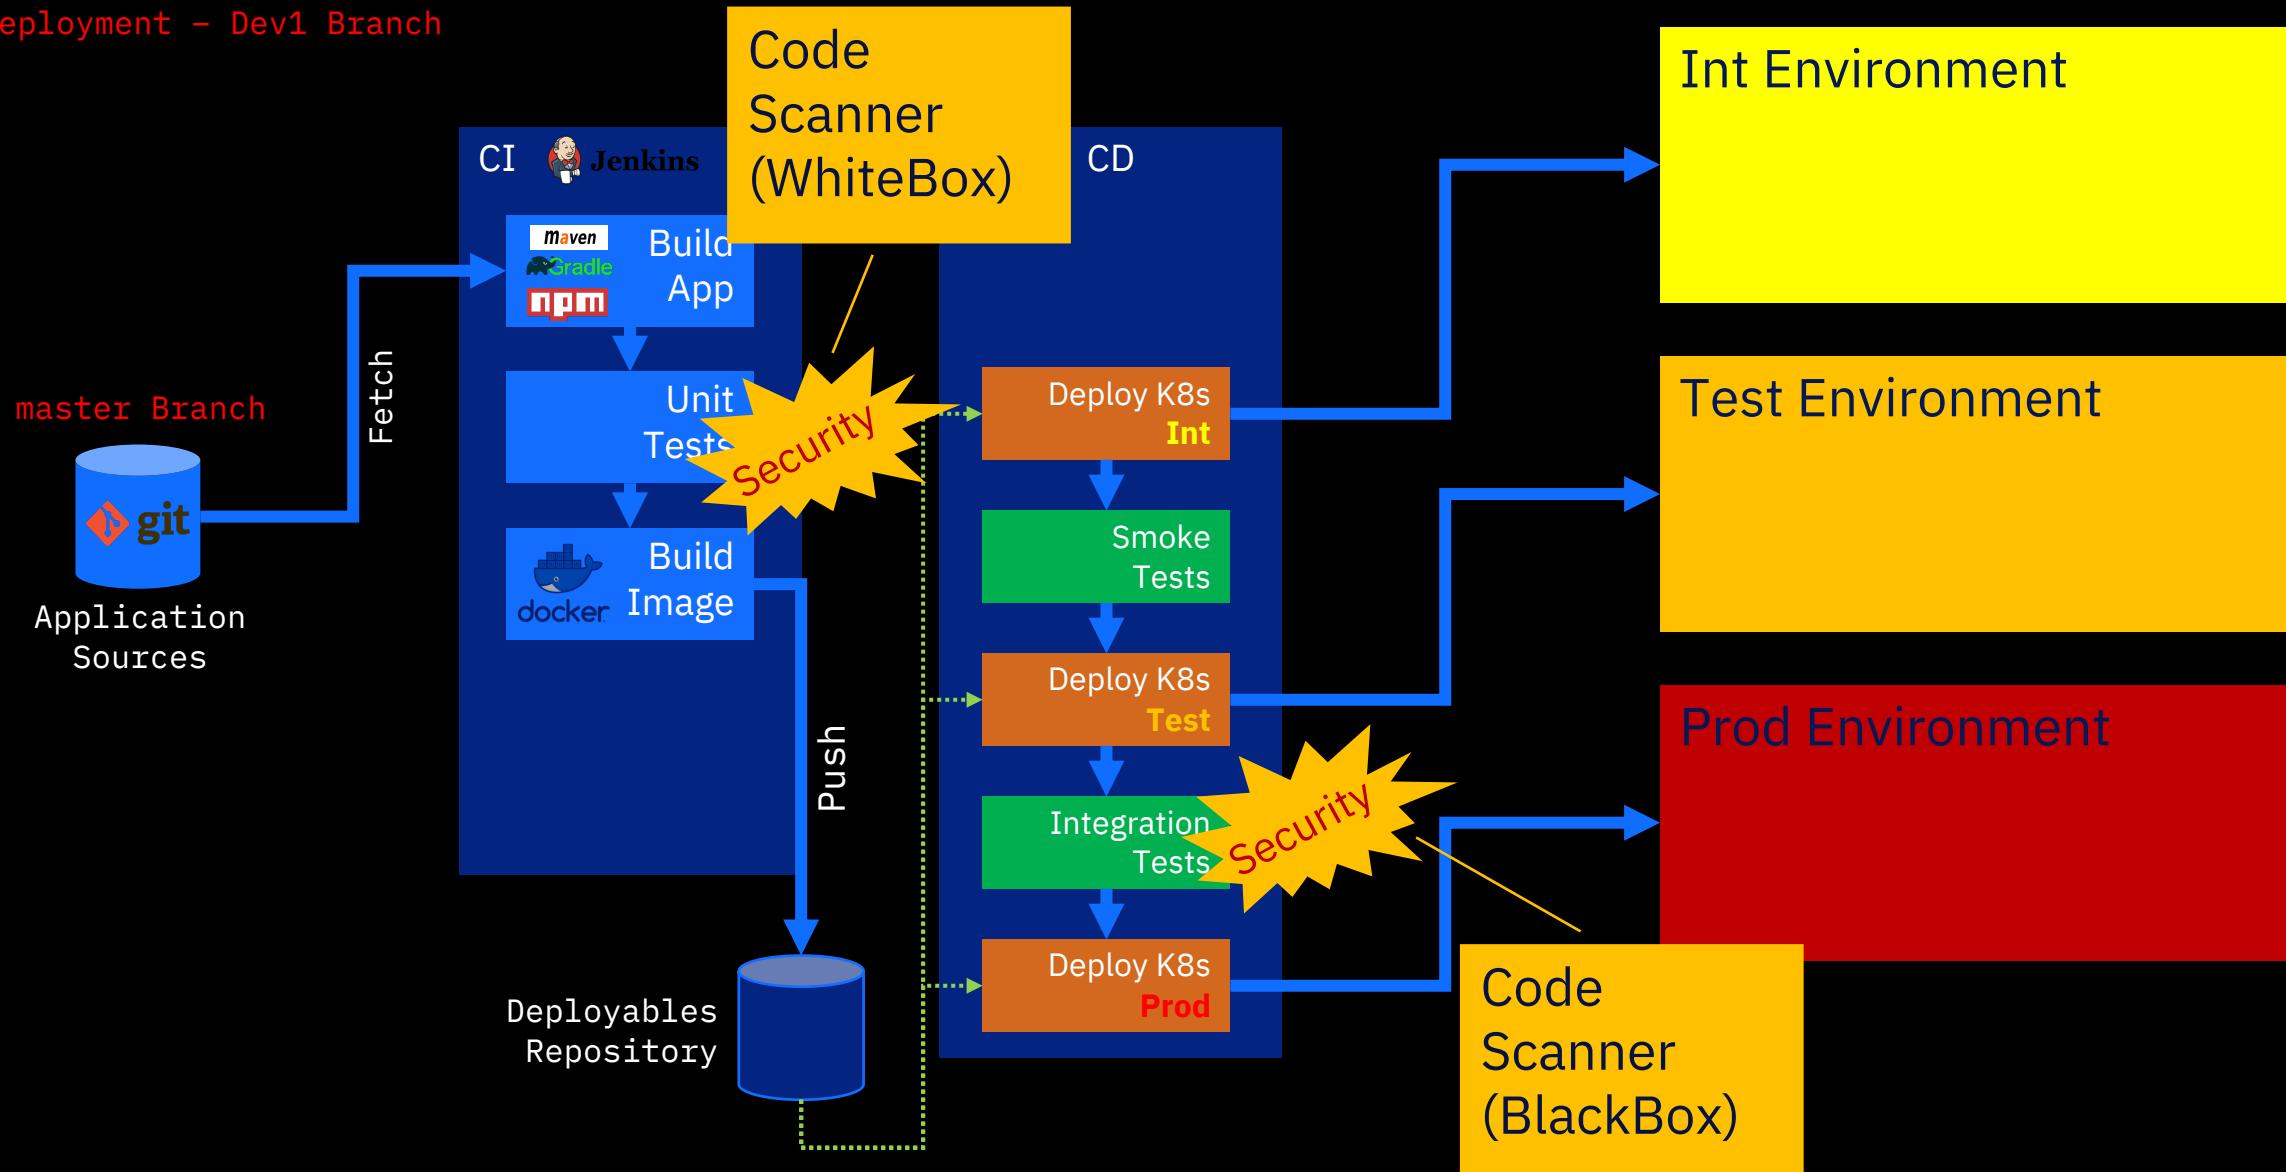
DevSecOps – CI/CD Pipeline

Development – Dev1 Branch

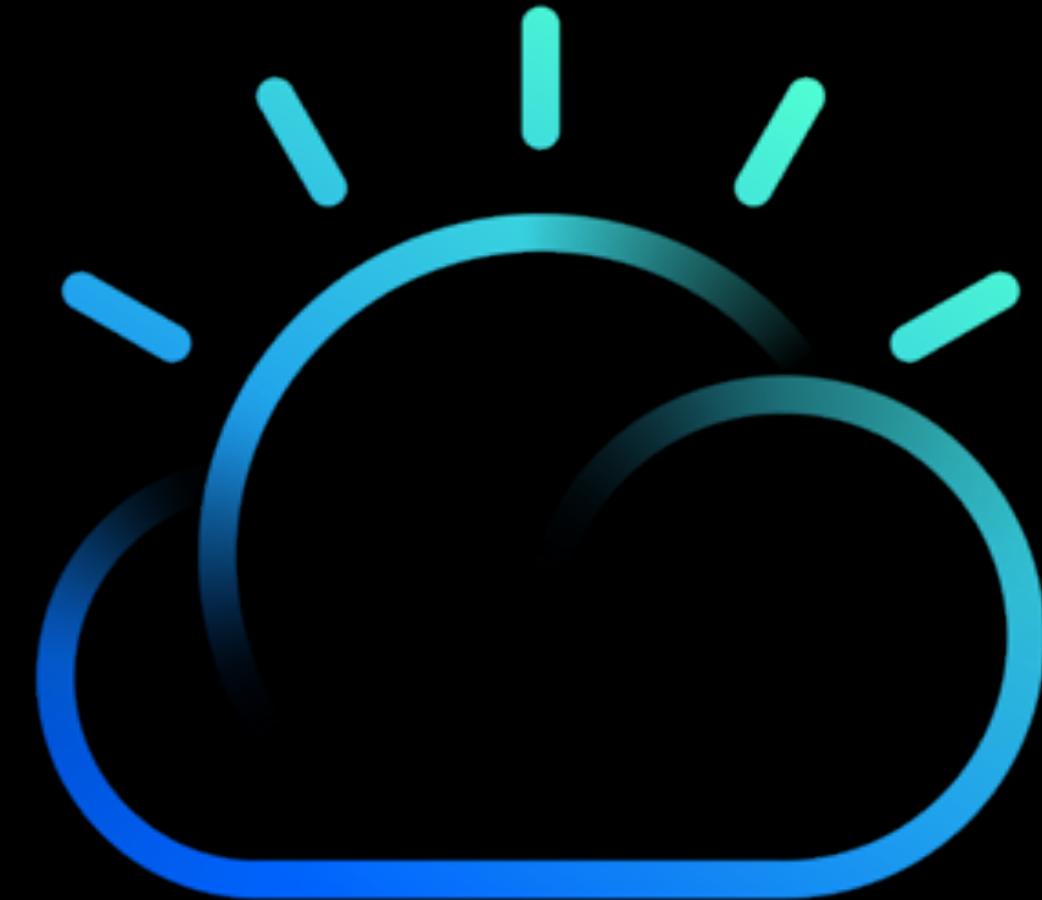


DevSecOps – CI/CD Pipeline

Deployment – Dev1 Branch



QUESTIONS?



The Journey to Cloud **Microservices**

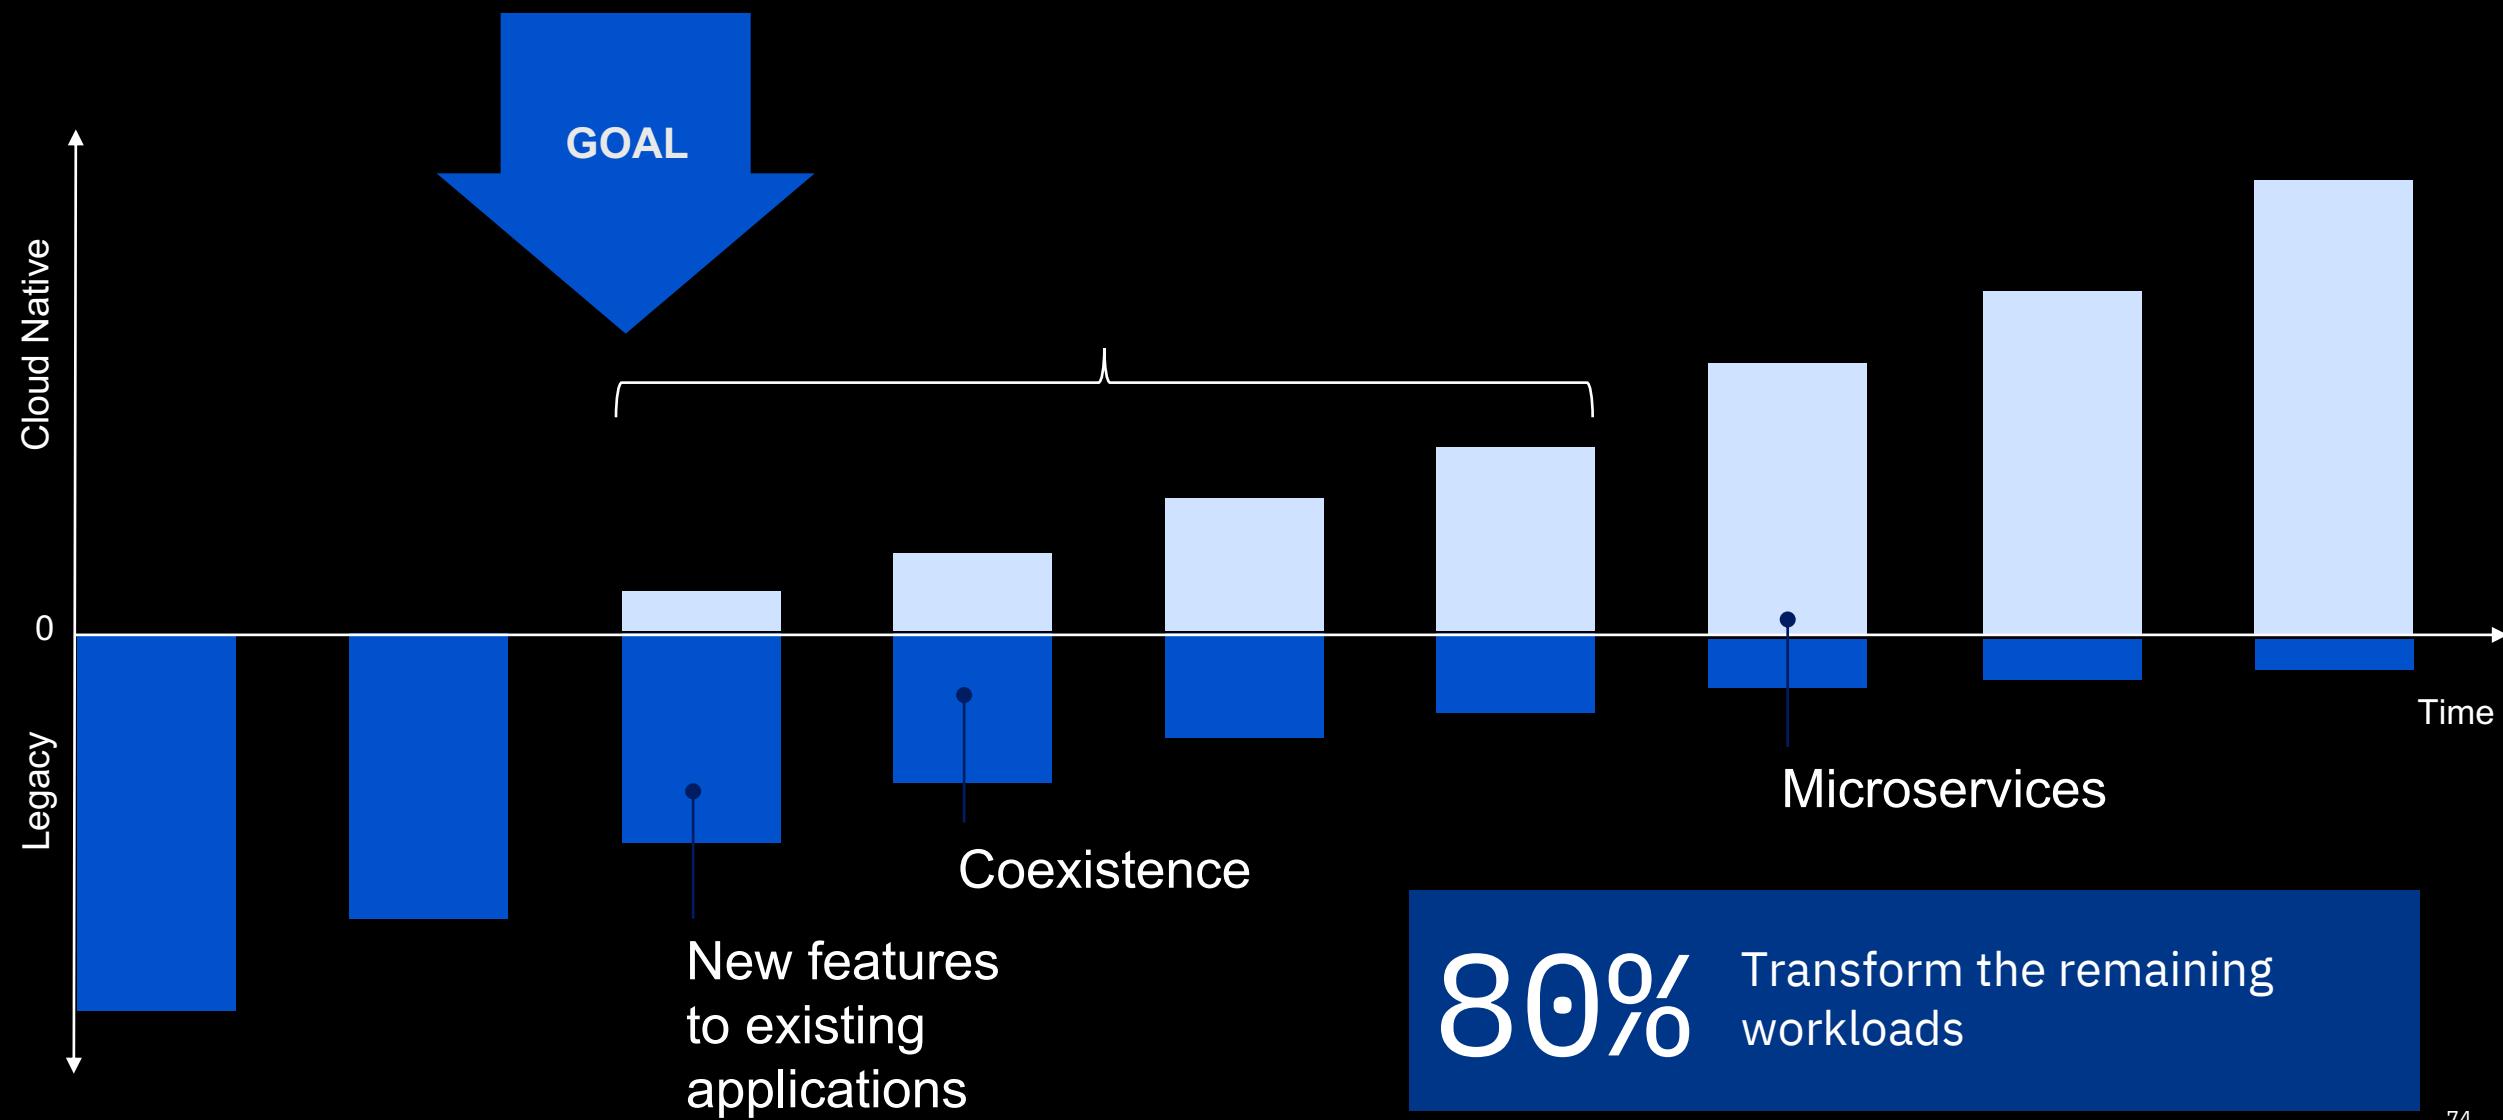
02



IBM Cloud

Digital Transformation

Cloud native and legacy apps will co-exist for the next 10+ years



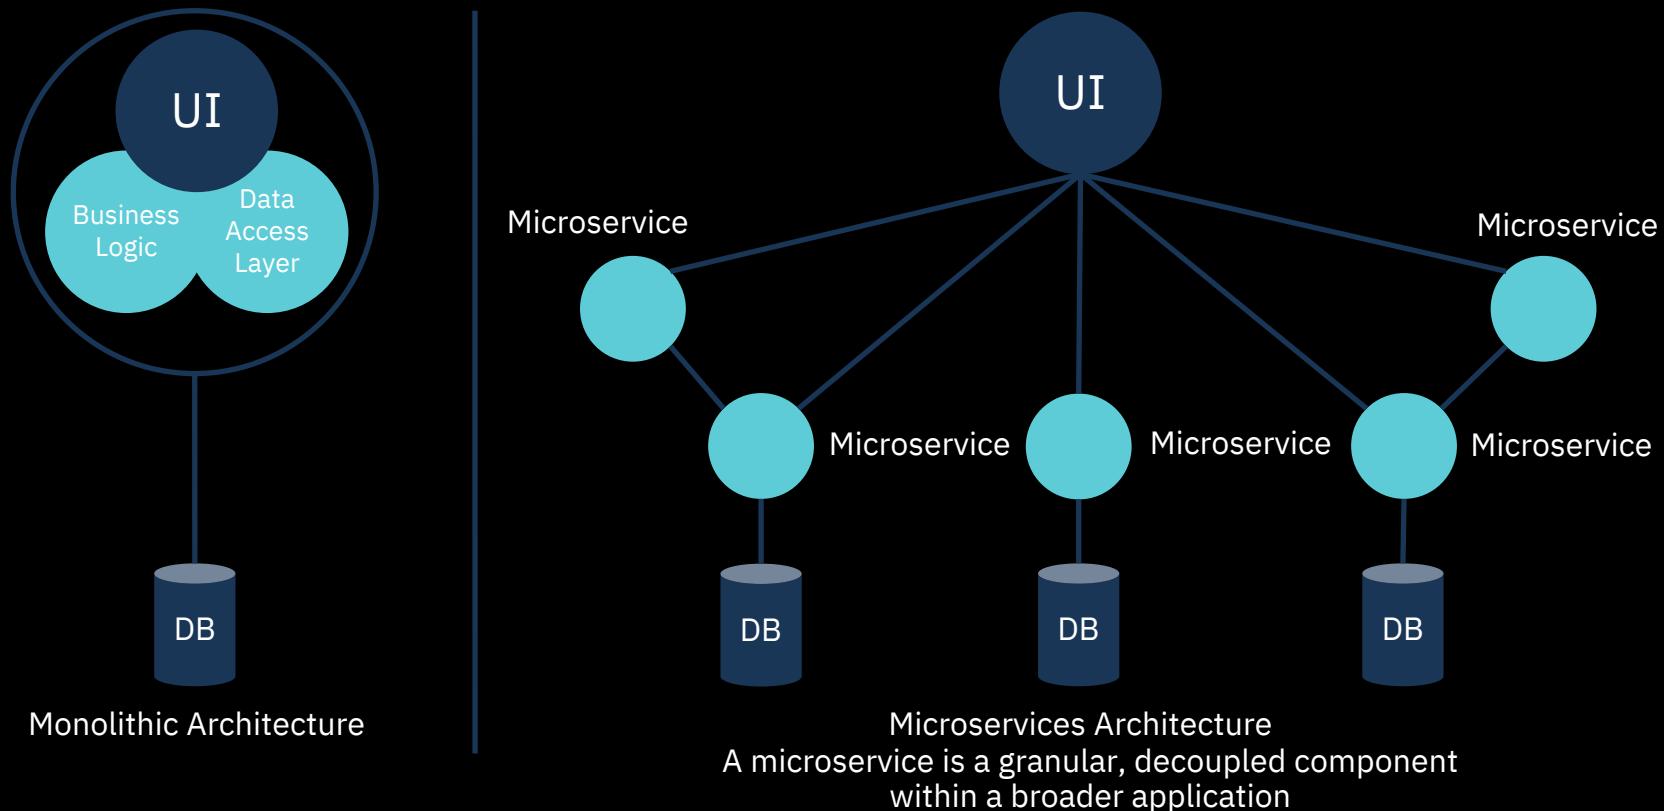
Microservices

Decomposing an application into **single function modules** which are **independently deployed and operated**

Accelerate delivery by minimizing communication and coordination between people

Microservices architecture

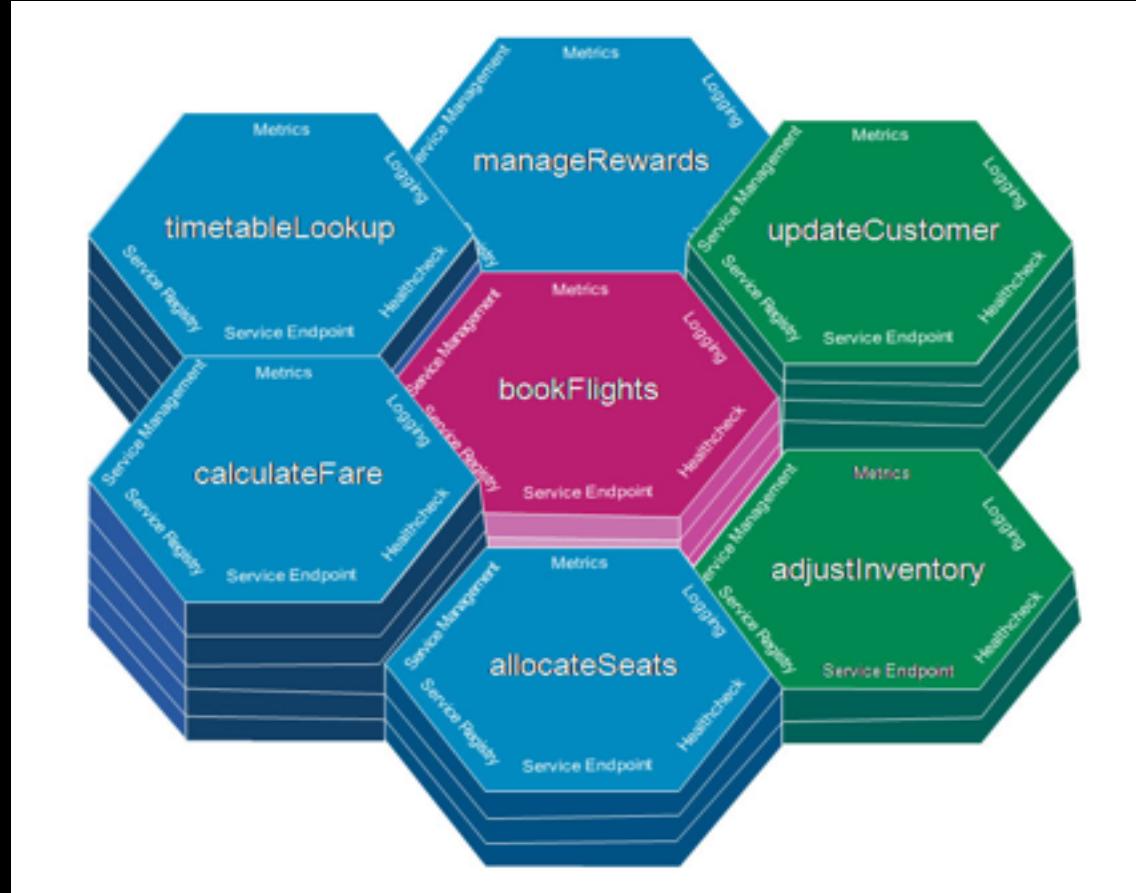
Simplistically, microservices architecture is about breaking down large silo applications into more manageable, fully decoupled pieces



Example application using microservices

Airline reservation application

- Book flights
- Timetable lookup
- Calculate fare
- Allocate seats
- Manage rewards
- Update customer
- Adjust inventory



Microservices – key tenets



Large monoliths are **broken down into many small services**

- Each service runs its own process
- There is one service per container



Services are **optimized for a single function**

- There is only one business function per service
- The Single-responsibility Principle: A microservice should have one, and only one, reason to change



Communication via **REST API** and **message brokers**

- Avoid tight coupling introduced by communication through a database



Per-service continuous delivery (CI/CD)

- Services evolve at different rates
- Let the system evolve, but set architectural principles to guide that evolution



Per-service high availability and clustering decisions

- One size or scaling policy is not appropriate for all
- Not all services need to scale; others require autoscaling up to large numbers

Microservices – advantages

In a microservices architecture each component:

- Is **developed independently** and has **limited, explicit dependencies** on other services
- Is developed by a **single, small team** in which all team members can understand the entire code base
- Is developed on its **own timetable** so new versions are delivered independently of other services
- **Scales and fails independently** which isolates any problems
- Can be developed in a different **language**
- **Manages its own data** to select the best technology and schema

Obvious rules for developing cloud applications

1. Don't code your application directly to a **specific topology**
2. Don't assume the **local file system** is permanent
3. Don't keep **session state** in your application
4. Don't assume any specific **infrastructure dependency**
5. Don't use **infrastructure APIs** from within your application
6. Don't rely on **OS-specific features**
7. Don't **manually install** your application

Read the article : http://www.ibm.com/developerworks/websphere/techjournal/1404_brown/1404_brown.html

That said....

Microservices

are

hard

..but we're
getting ahead
of ourselves

QUESTIONS?



IBM Cloud

The Journey to Cloud **Docker**

03



IBM Cloud

Everybody Loves Containers



A **standard way to package an application and all its dependencies** so that it can be moved between environments and run without changes

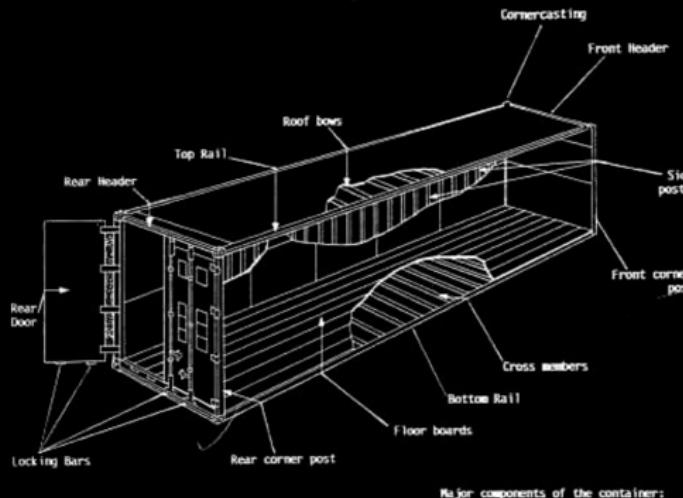
Containers work by **isolating the differences between applications** inside the container so that everything outside the container can be **standardized**

Microservices implementation with Containers

Why it works – separation of concerns

Development

- Worries about what's “**inside**” the container
 - Code
 - Libraries
 - Package Manager
 - Apps
 - Data
- All Linux servers look the same

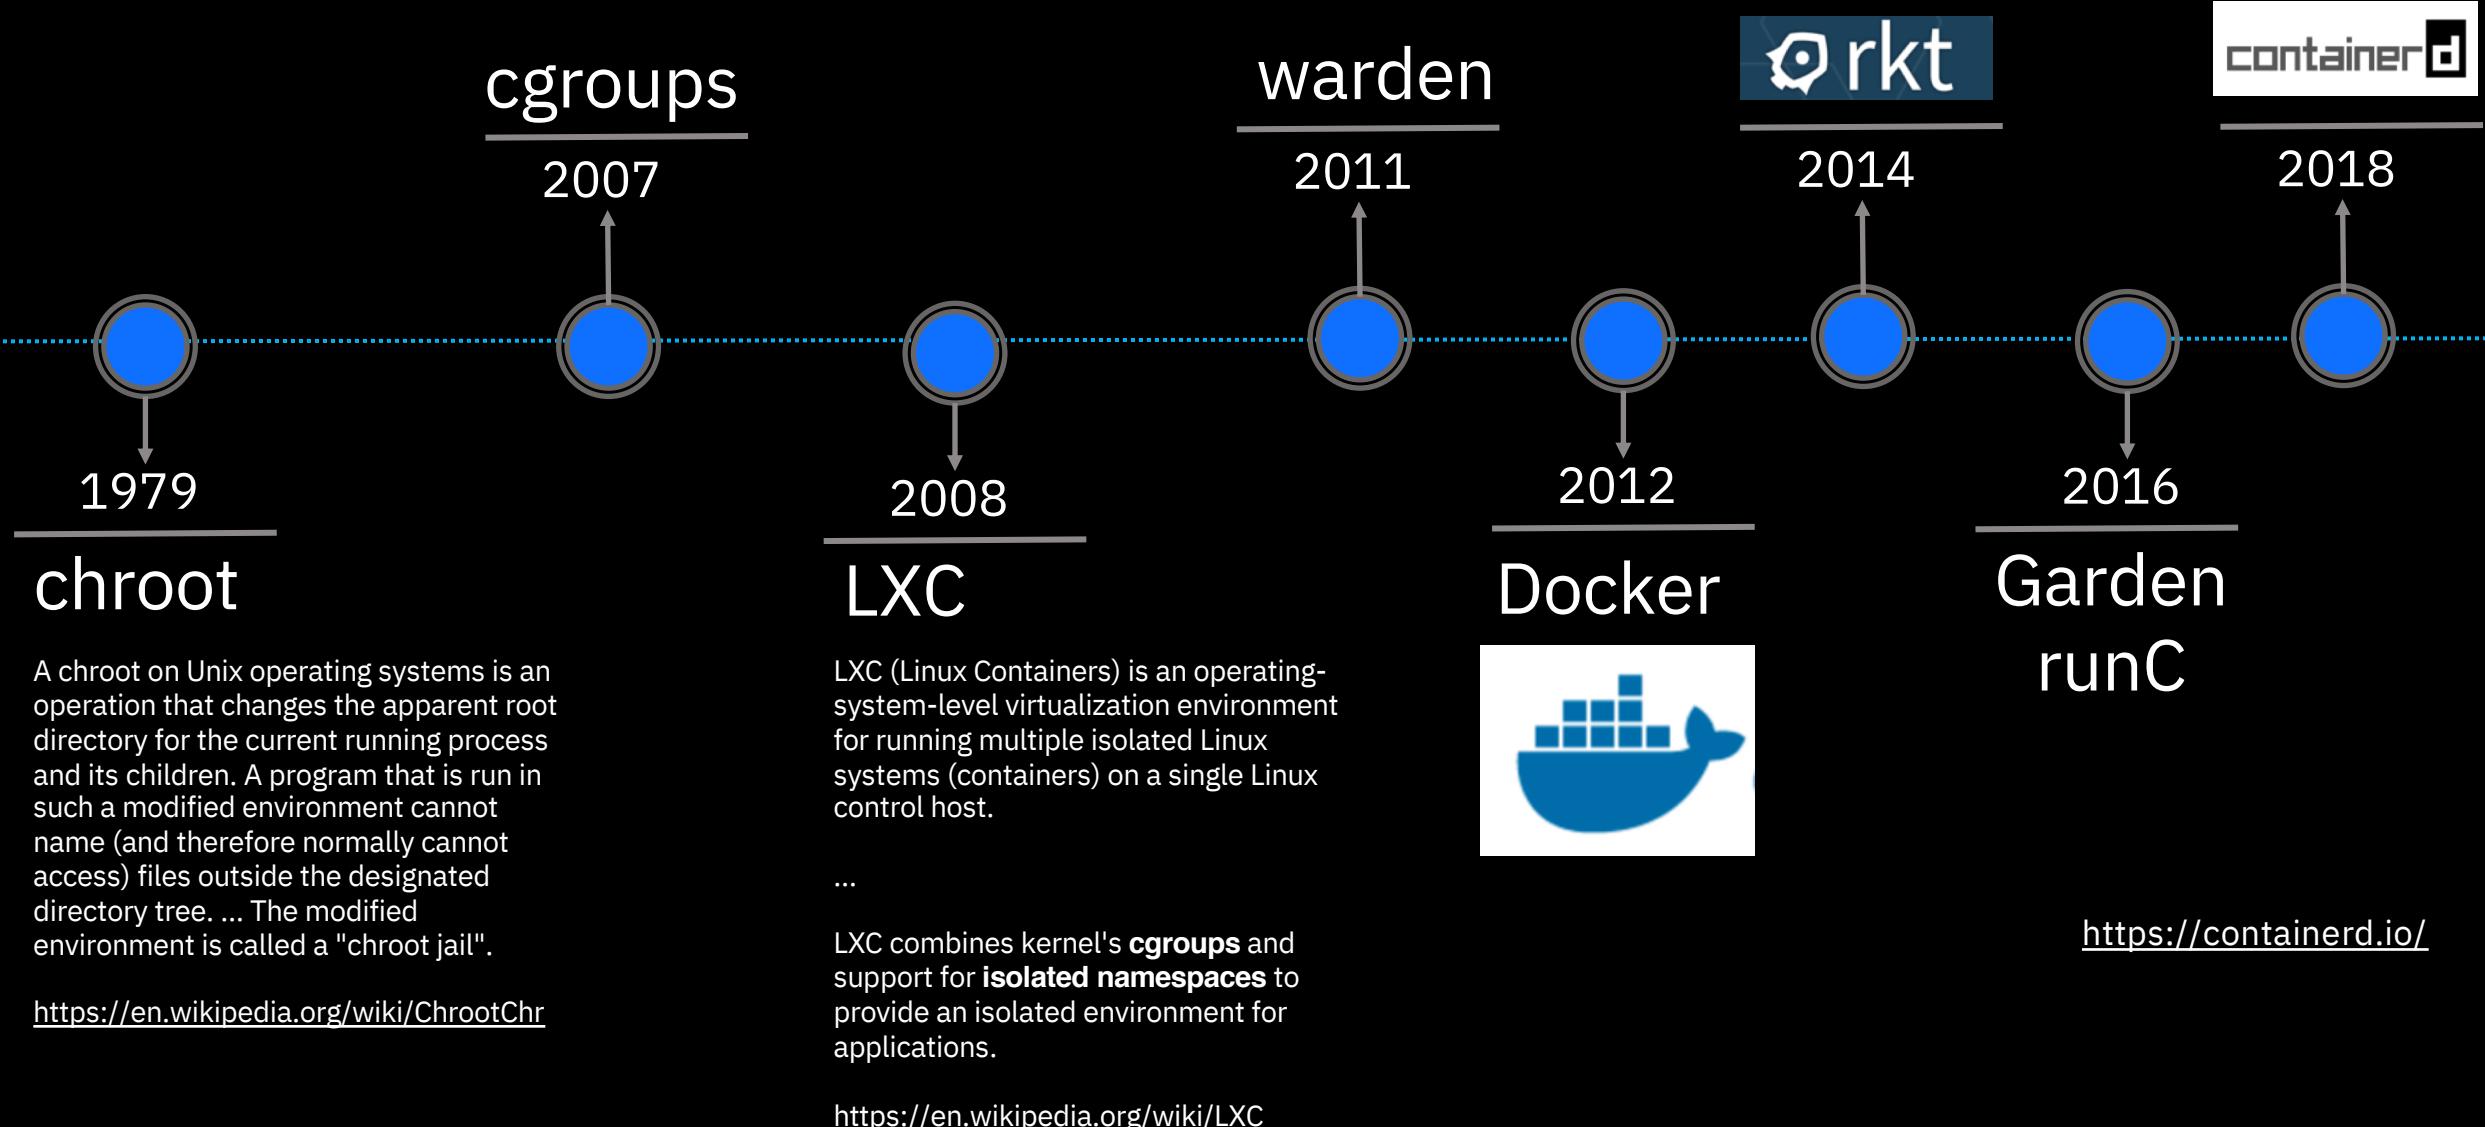


Operations

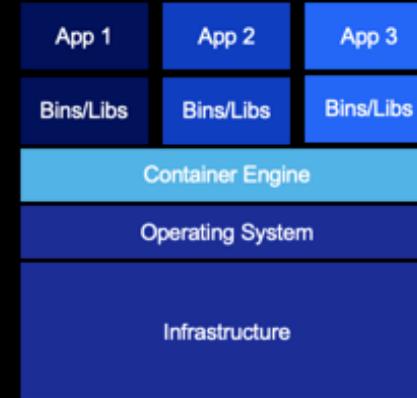
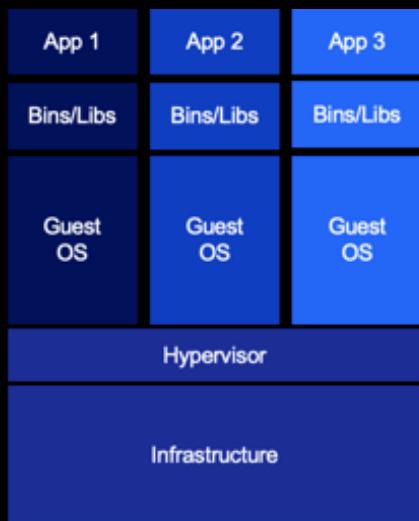
- Worries about what's “**outside**” the container
 - Logging
 - Remote Access
 - Monitoring
 - Network Config
- All containers start, stop, copy, attach, migrate, etc... the same way

Clear ownership boundary between
Dev and IT Ops drives DevOps adoption and fosters agility

Container History



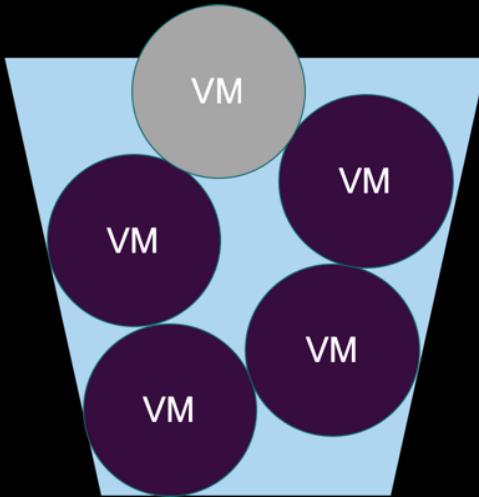
VMs vs. Containers



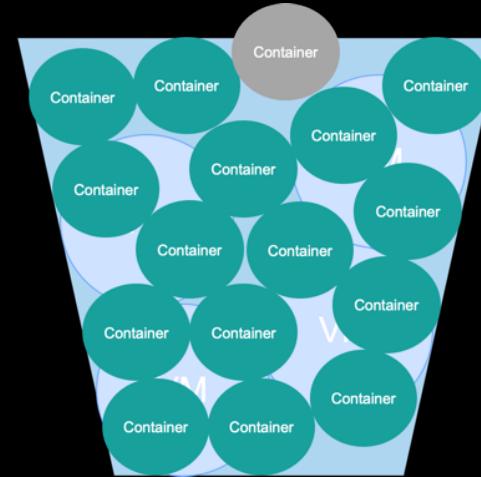
- + VM Isolation
- Complete OS
- Static Compute
- Static Memory

- + Container Isolation
- + Shared Kernel
- + Burstable Compute
- + Burstable Memory

VMs vs. Containers



- + VM Isolation
- Complete OS
- Static Compute
- Static Memory
- Low Resource Utilization



- + Container Isolation
- + Shared Kernel
- + Burstable Compute
- + Burstable Memory
- + High Resource Utilization

Advantages of Containers



Containers are **portable**



Containers are **easy to manage**



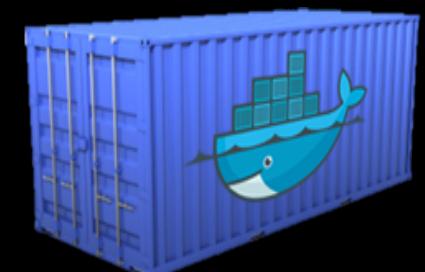
Containers provide “**just enough**” isolation



Containers use hardware **more efficiently**



Containers are **immutable**



Docker Components

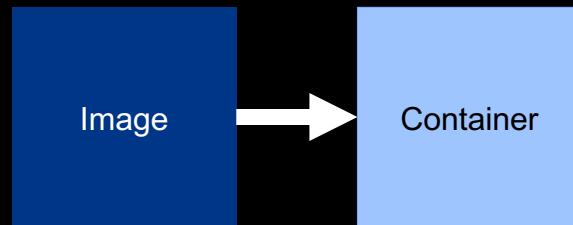
Container

Smallest compute unit



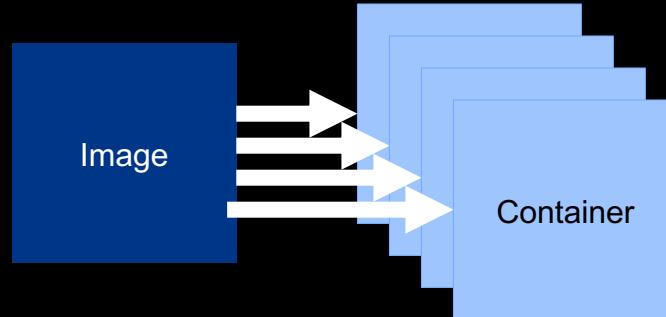
Docker Components

Containers
are created from
Images



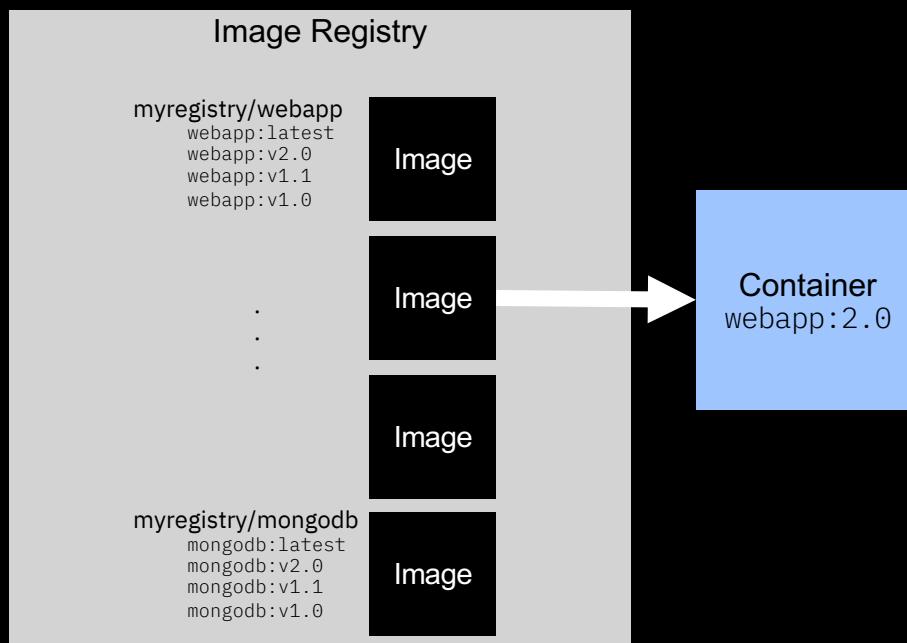
Docker Components

As **many Containers**
as needed can be created from
Images



Docker Components

The **Image Registry**
stores the versioned
Images
to create
Containers



Docker Components



Image

A read-only snapshot of a container stored in Docker Hub to be used as a template for building containers



Container

The standard unit in which the application service resides or transported



Docker Hub/Registry

Available in SaaS or Enterprise to deploy anywhere you choose
Stores, distributes, and shares container images



Docker Engine

A program that creates, ships, and runs application containers
Runs on any physical and virtual machine or server locally, in private or public cloud. Client communicates with Engine to execute commands

Why Containers

Docker Layers

Inheritance

- Build on the work of those who came before you

```
# Pull base image
FROM tomcat:8-jre8

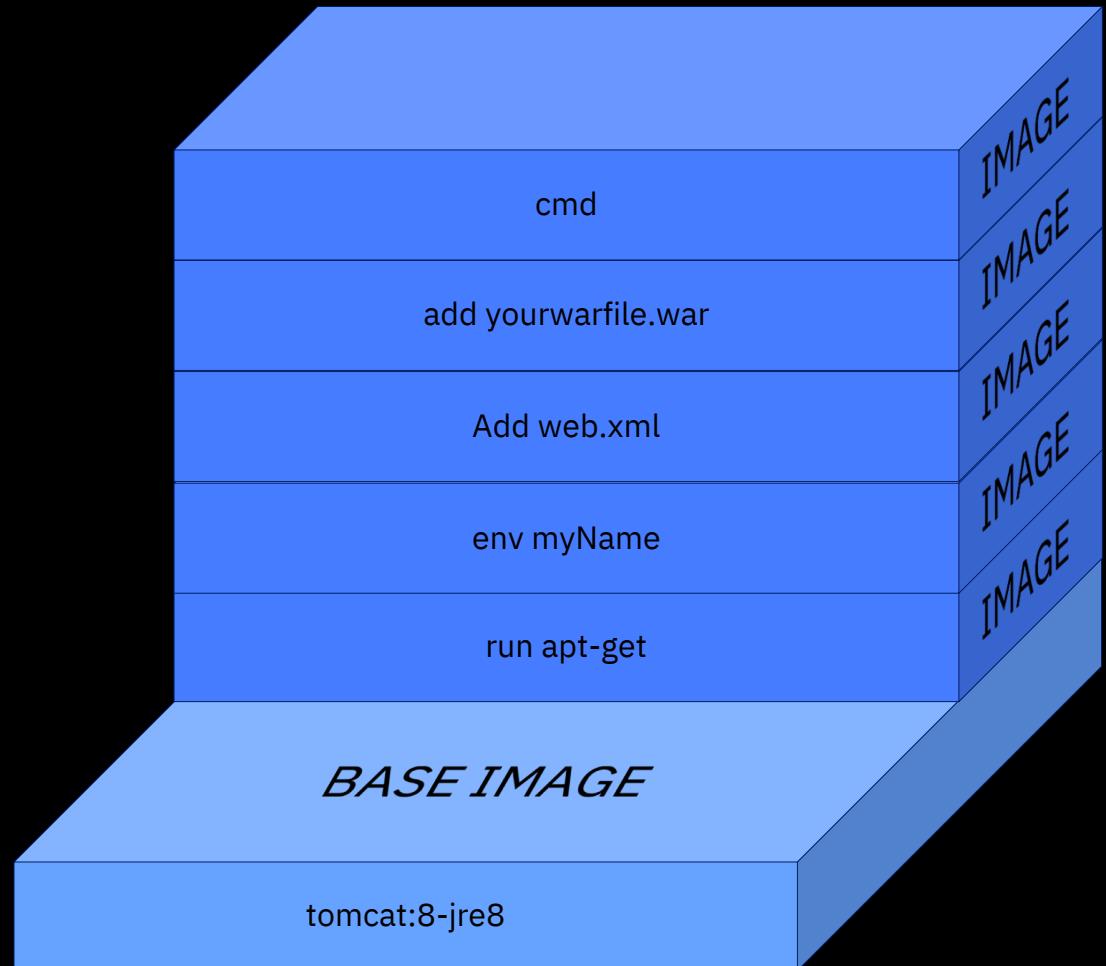
# Maintainer
MAINTAINER youremailaddress"

# Run command
RUN apt-get update && apt-get -y upgrade

# Set variables
ENV myName John Doe

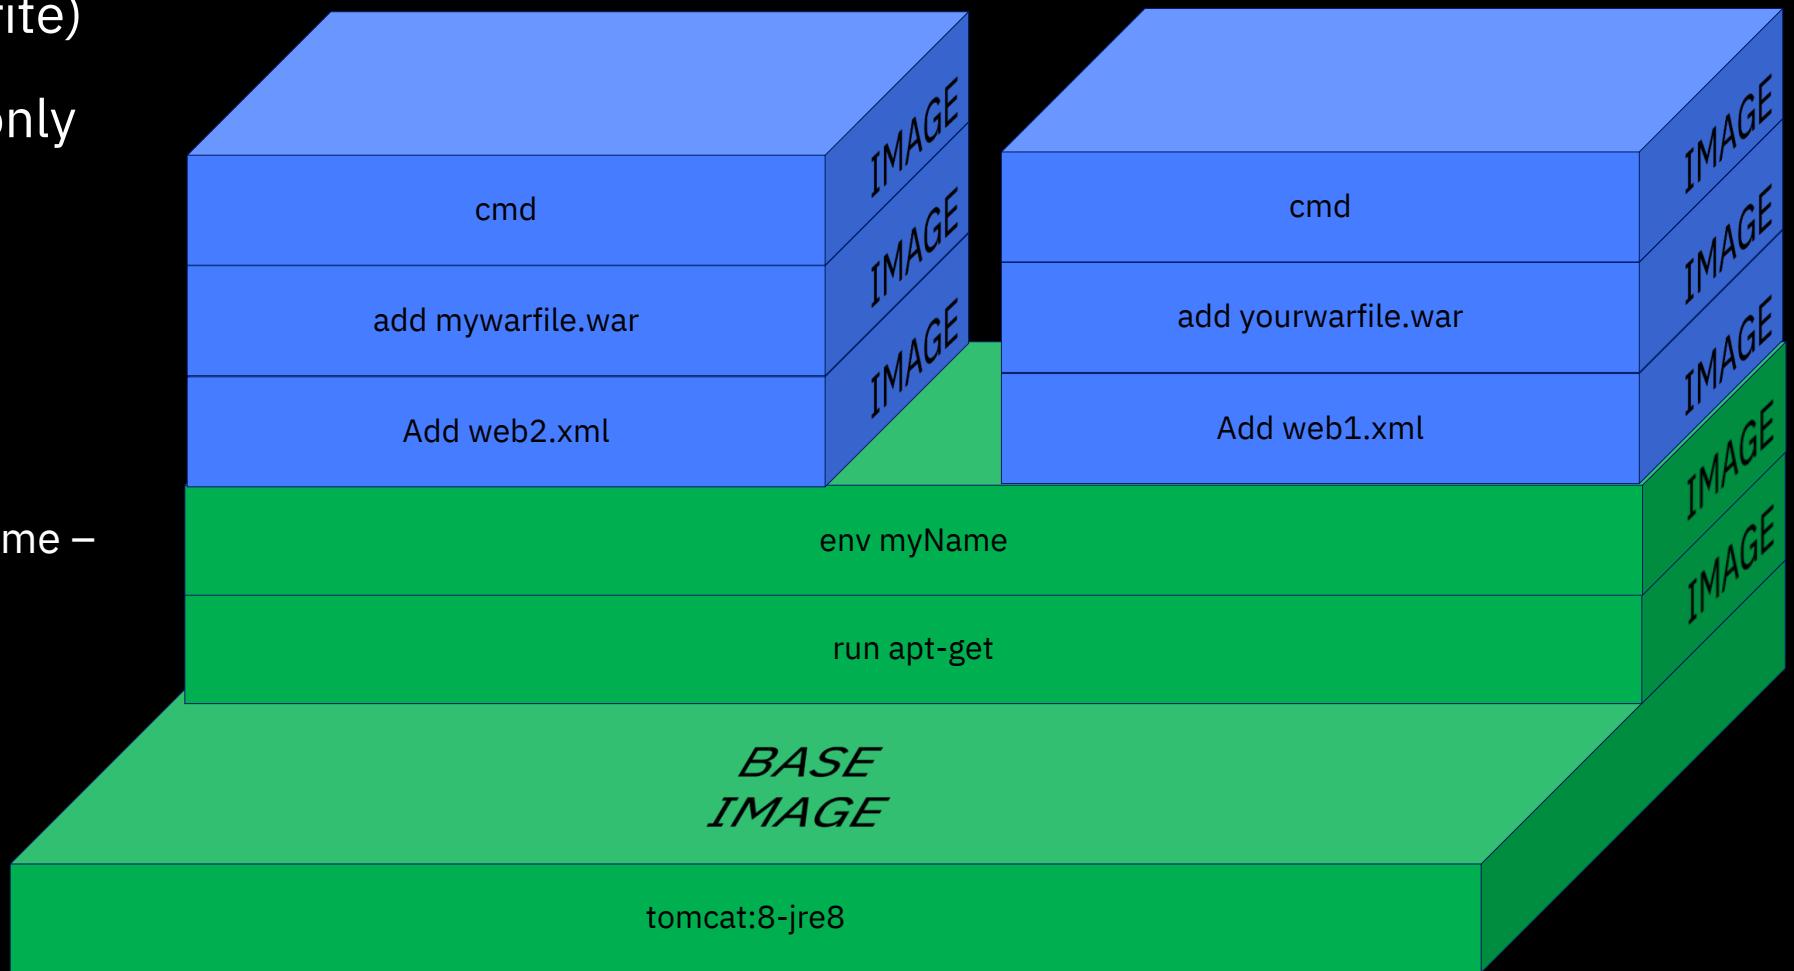
# Copy to images tomcat path
ADD web.xml /usr/local/tomcat/conf/
ADD yourwarfile.war /usr/local/tomcat/webapps/

# Run server
CMD ["catalina.sh", "run"]
```

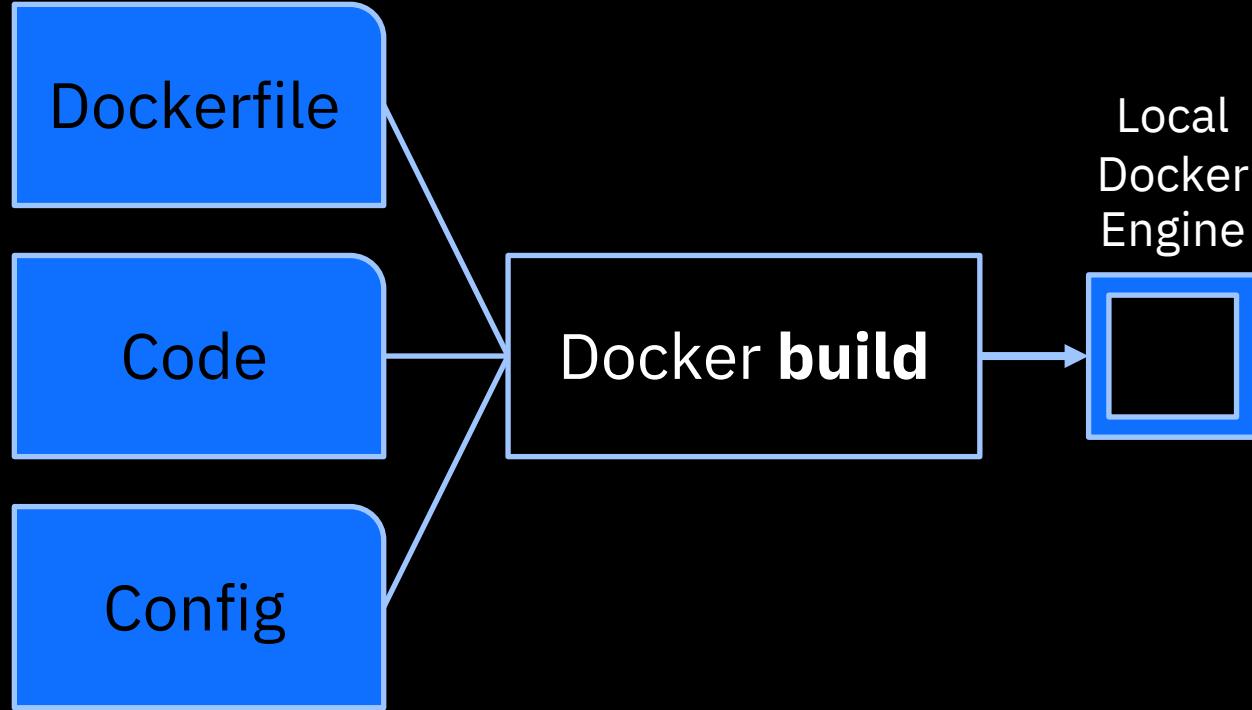


Why Containers

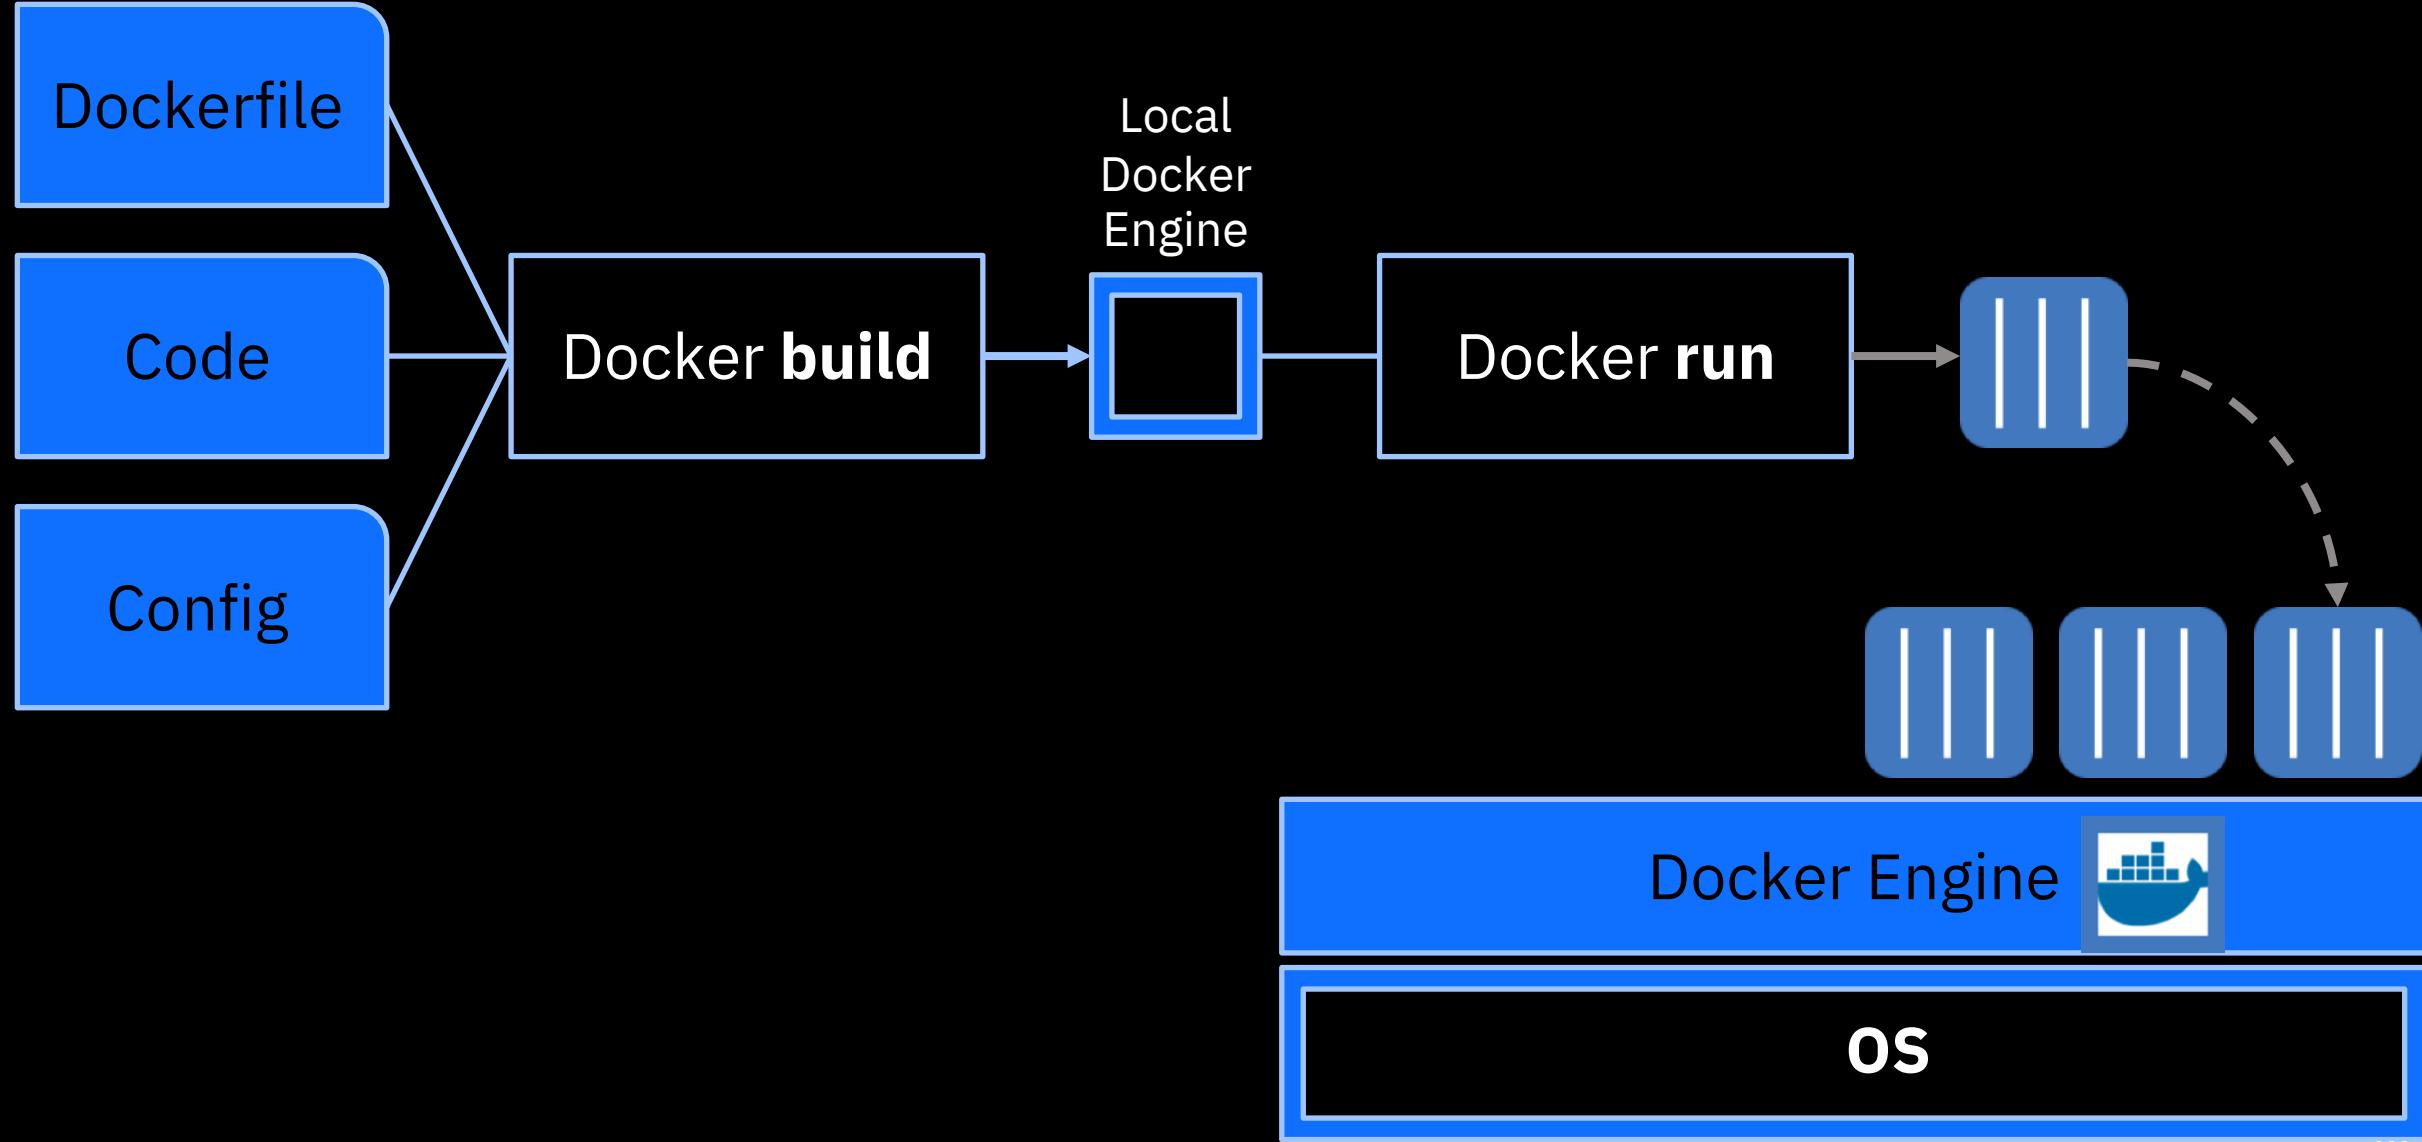
- Docker uses a **layered filesystem** (copy-on-write)
- New files (& edits) are only visible to current/above layers
- Layers allow for **reuse**
 - More containers per host
 - Faster start-up/download time – base layers are "cached"



Docker Basics – Build



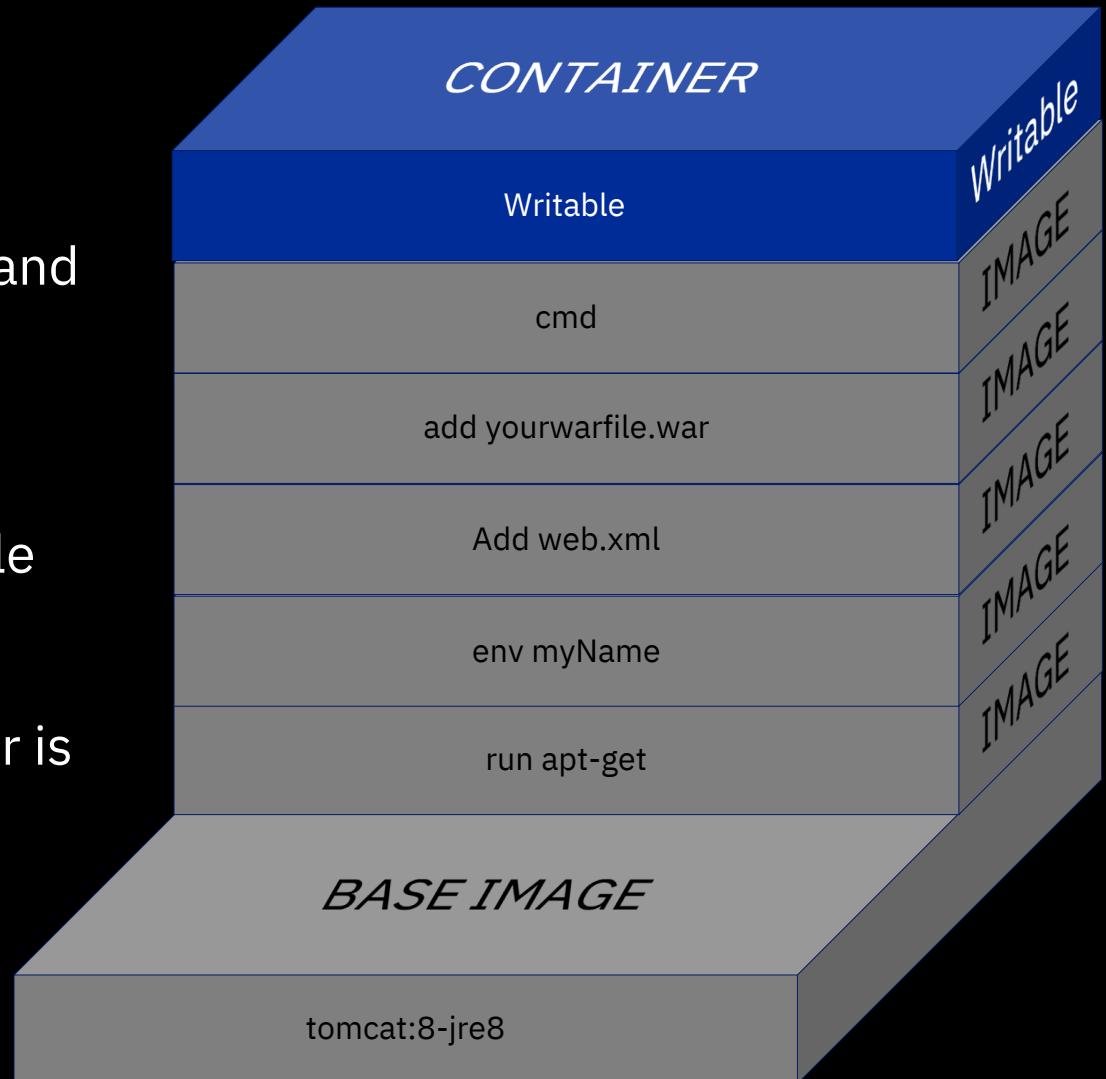
Docker Basics – Run



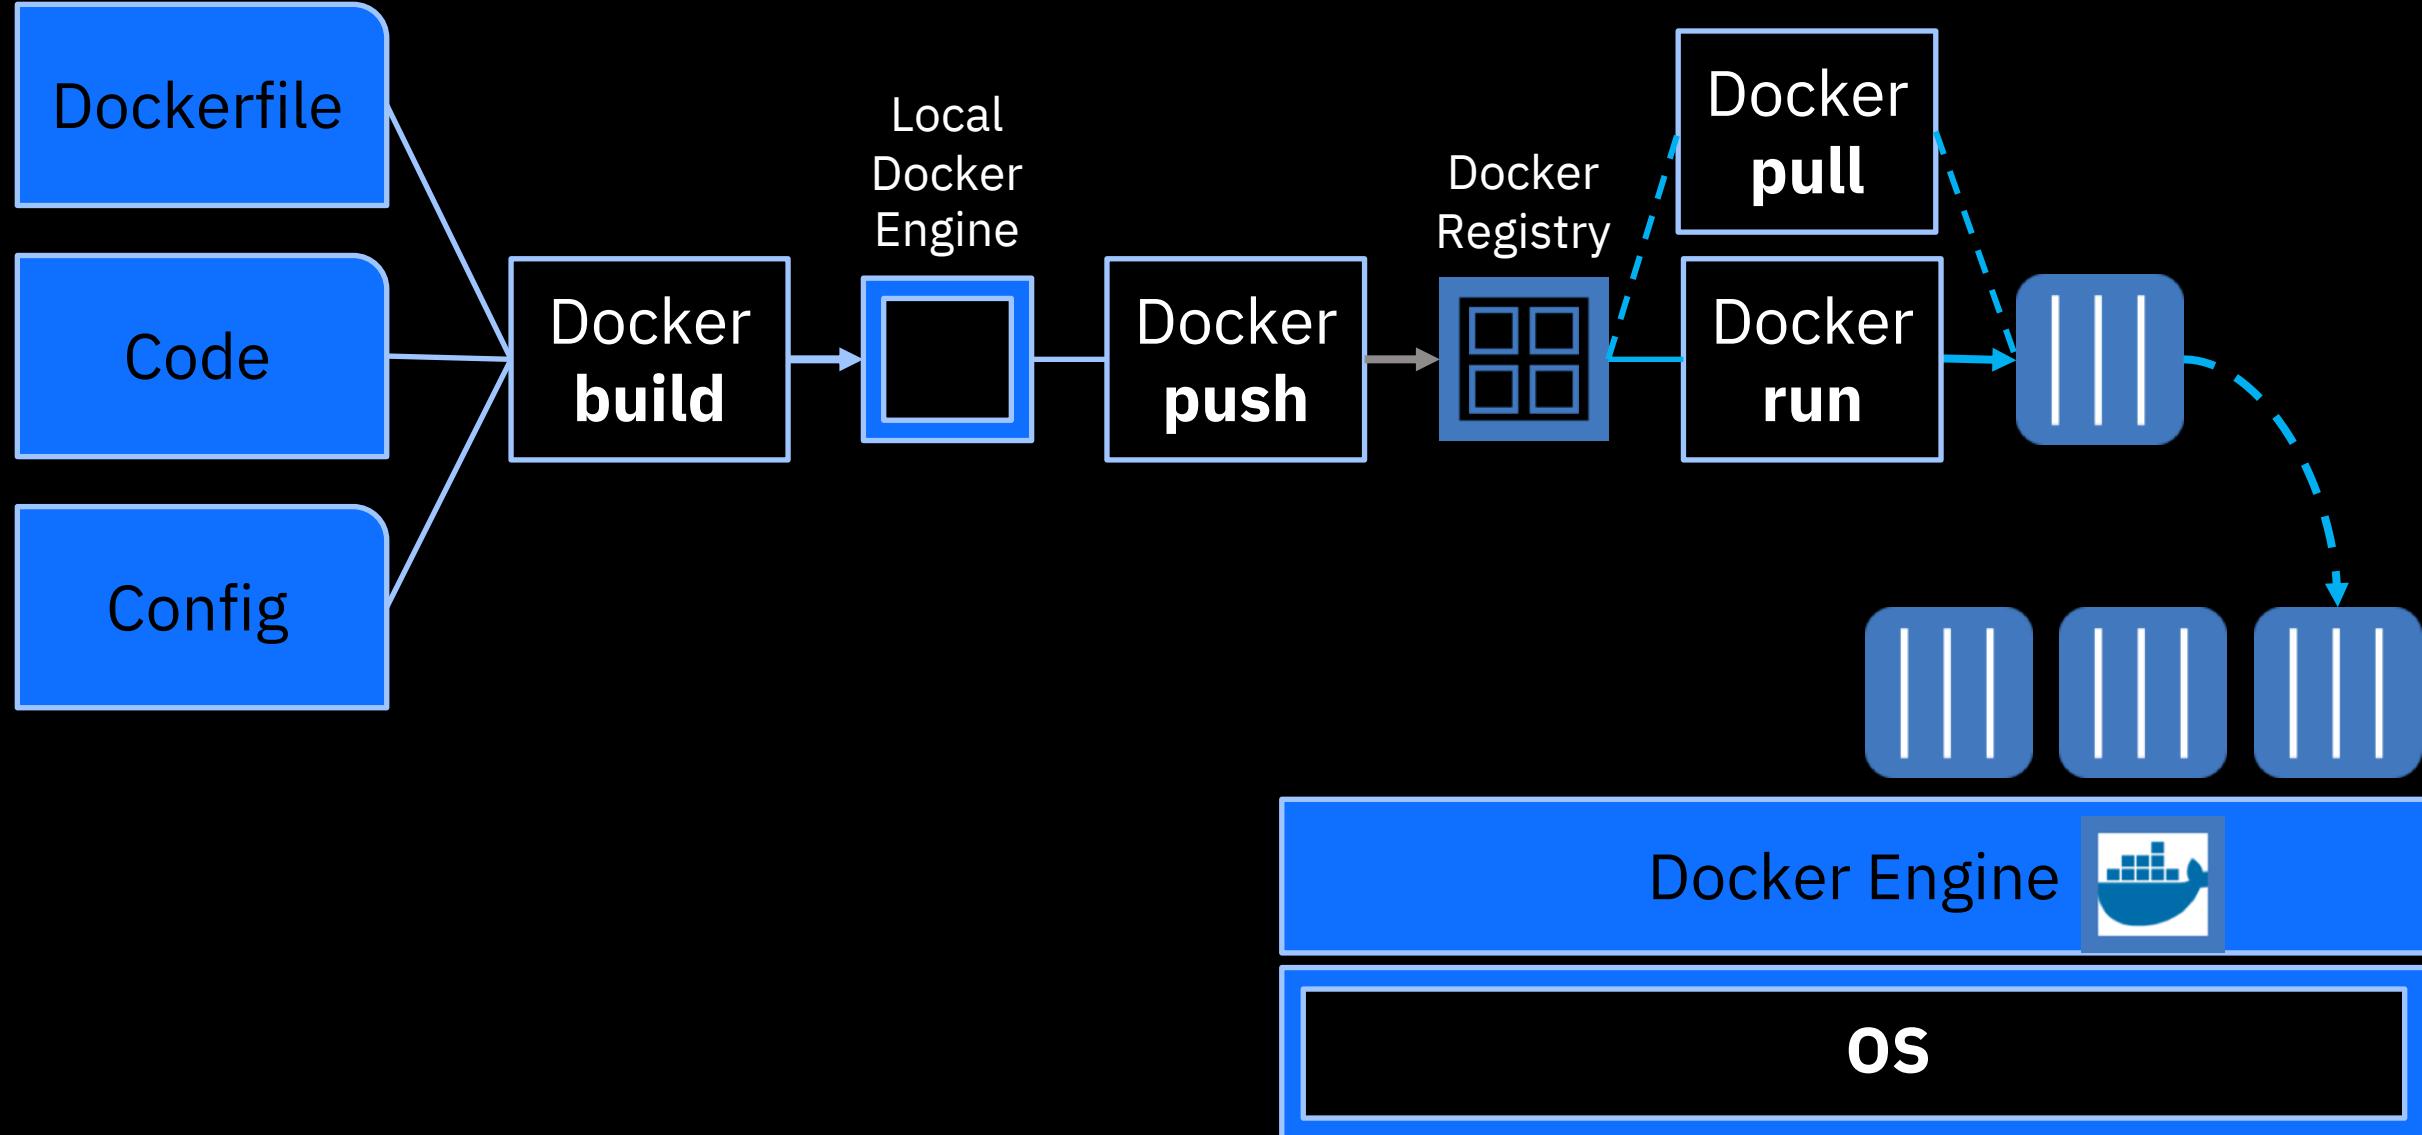
Docker Basics – Run

Docker Layers

- The biggest difference between a **container** and an **image** is the **top writable layer**.
- All writes to the container that add new or modify existing data are stored in this writable layer.
- When a container is deleted, its writable layer is also deleted. The underlying image is unchanged.



Docker Basics – Store, Retrieve & Run with registry



Docker Basics – Registries

Hosting image repositories

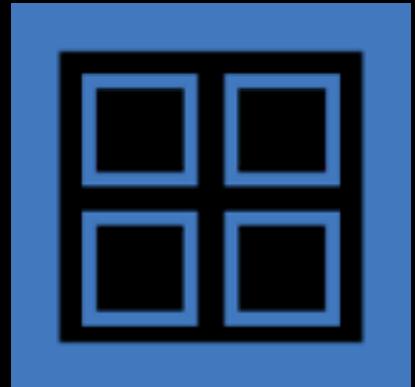
- You can define your own registry
- A registry is managed by a registry container

Public and Private registries

- Public Registry like **Docker Hub**
- <https://hub.docker.com>

Login into the registry

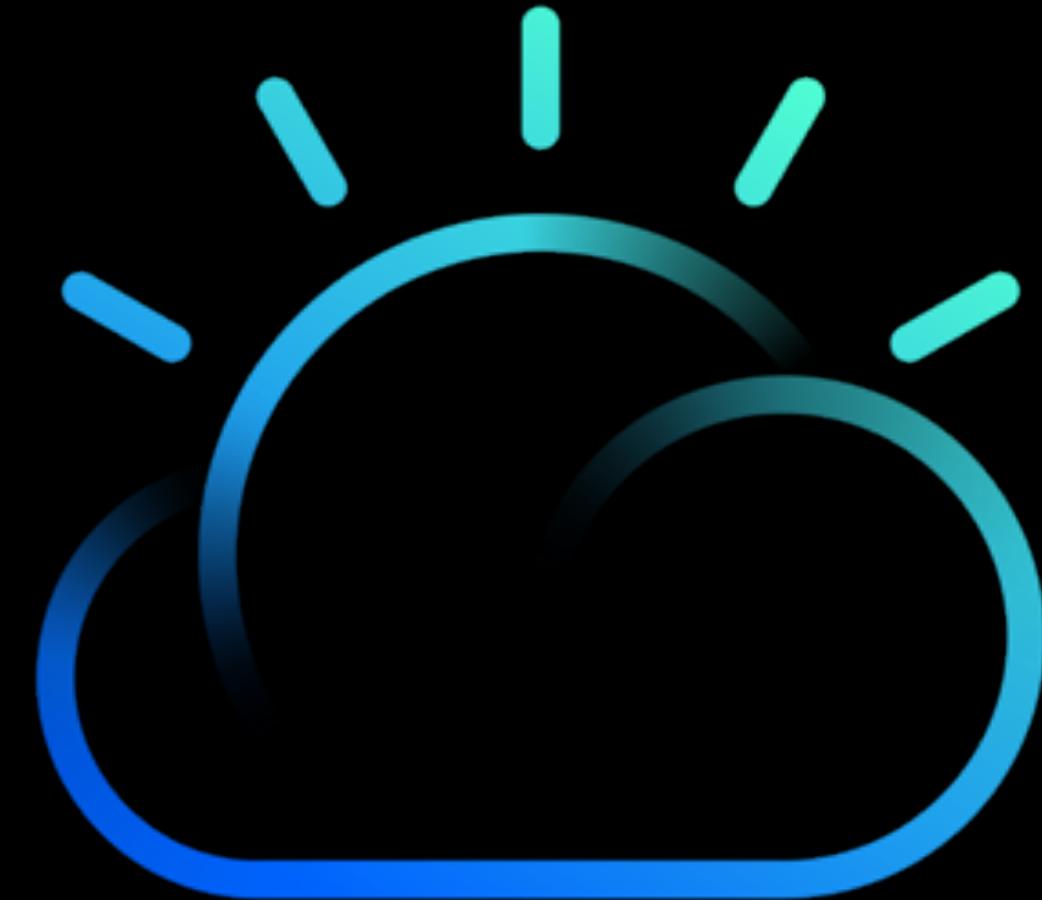
- Docker login domain:port



Docker Recap

- **Containers are not VMs**
- **Containers provide many benefits:**
 - Efficiency
 - Portability
 - Consistency
- **New challenges with containers:**
 - Production apps dependent on open-source projects
 - Existing tools may not be sufficient for container
 - Need to focus on business objectives

QUESTIONS?



The Journey to Cloud **Kubernetes**

04



IBM Cloud

Everybody Loves Containers

But when you go



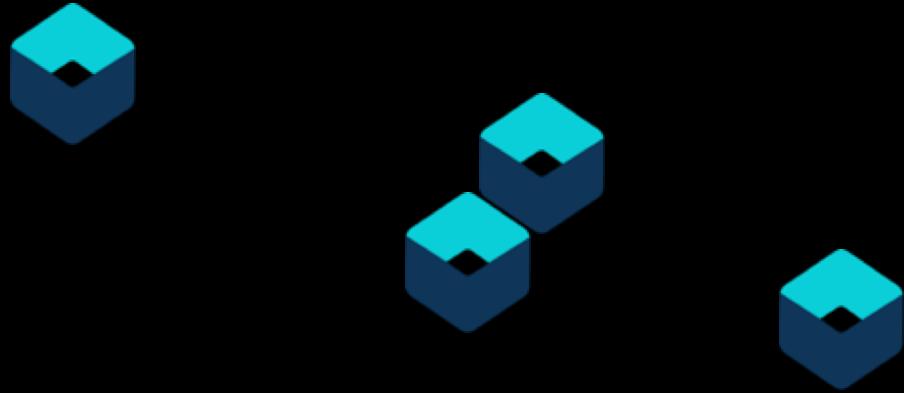
From this....



To this....

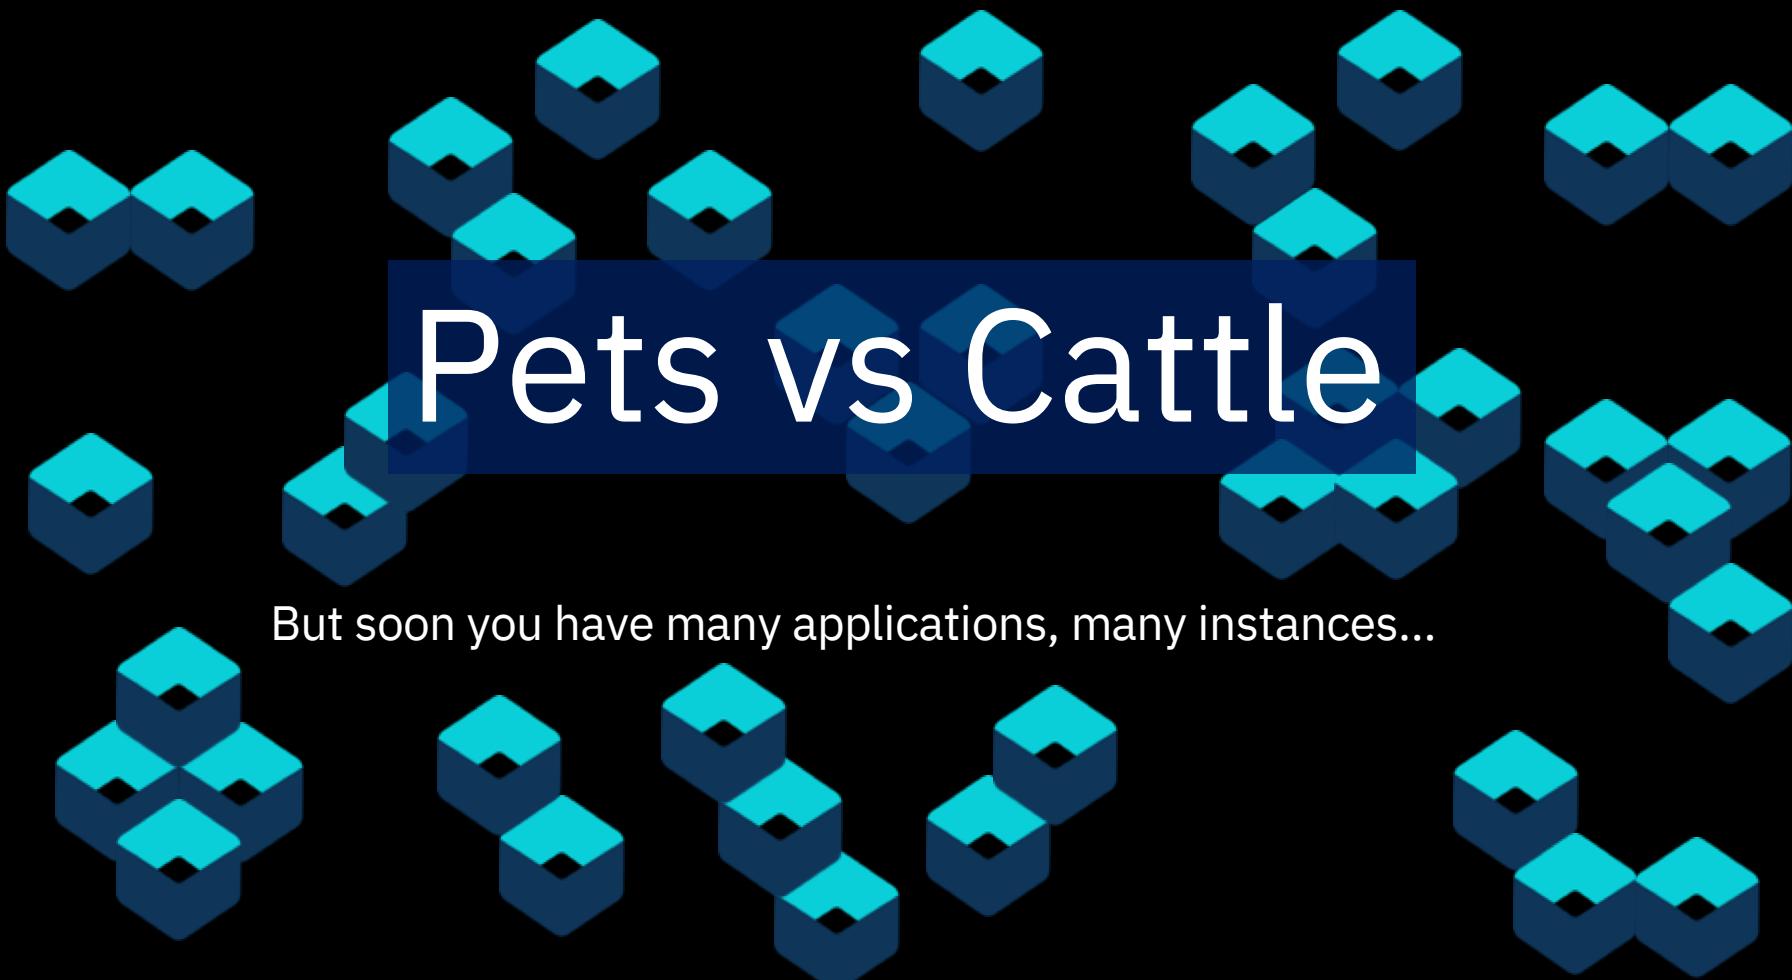


Everyone's container journey starts with one container....



At first the growth is easy to handle....



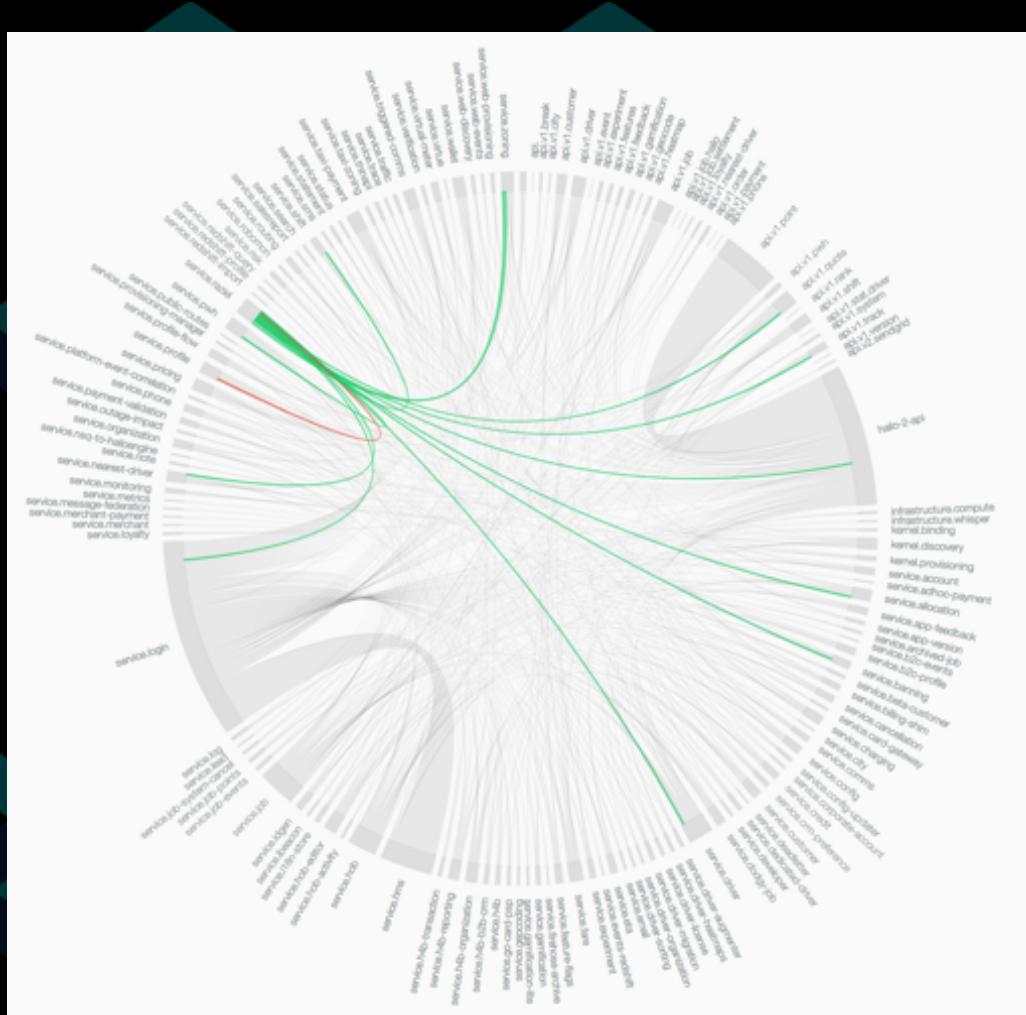


Pets vs Cattle

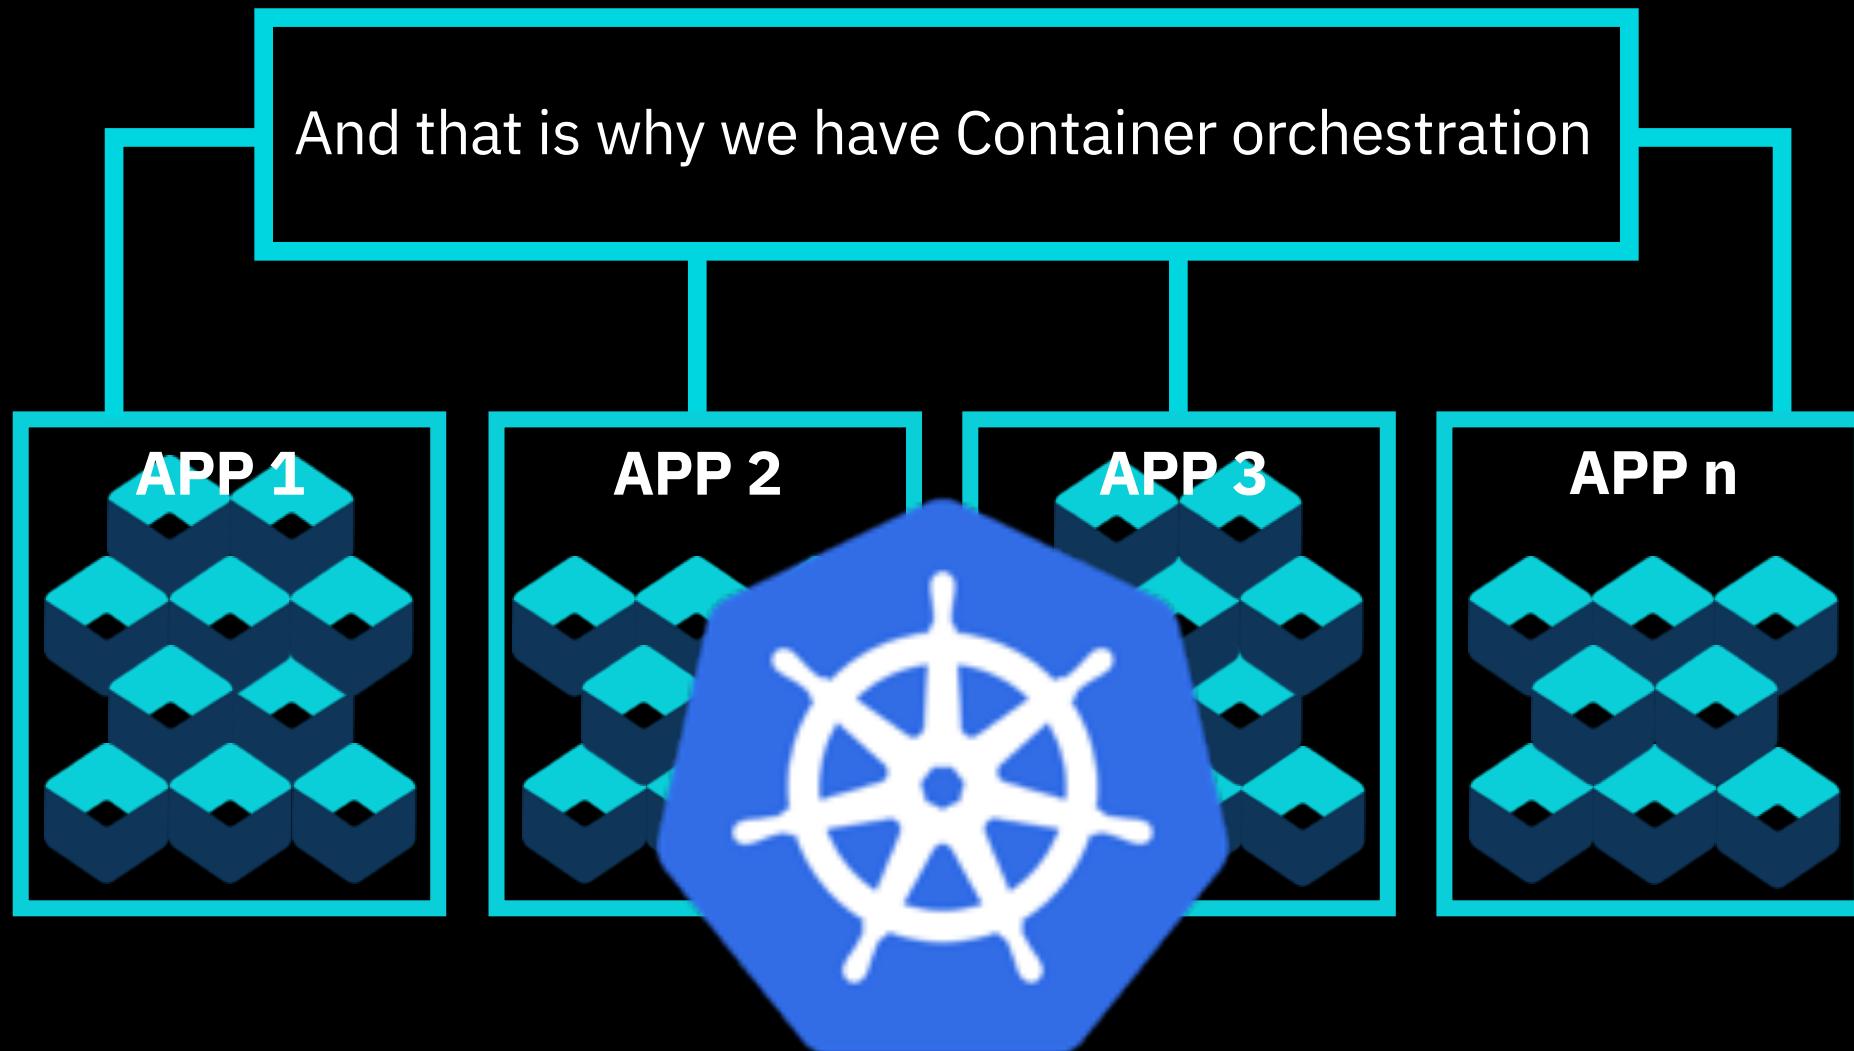
But soon you have many applications, many instances...

The trade off

Improved delivery velocity
in exchange for
increased operational complexity



Enter Kubernetes (K8s)





Imperative Systems

In an imperative system, **the user** knows the desired state and **the user** determines the sequence of commands to transition the system to the desired state.

Declarative Systems

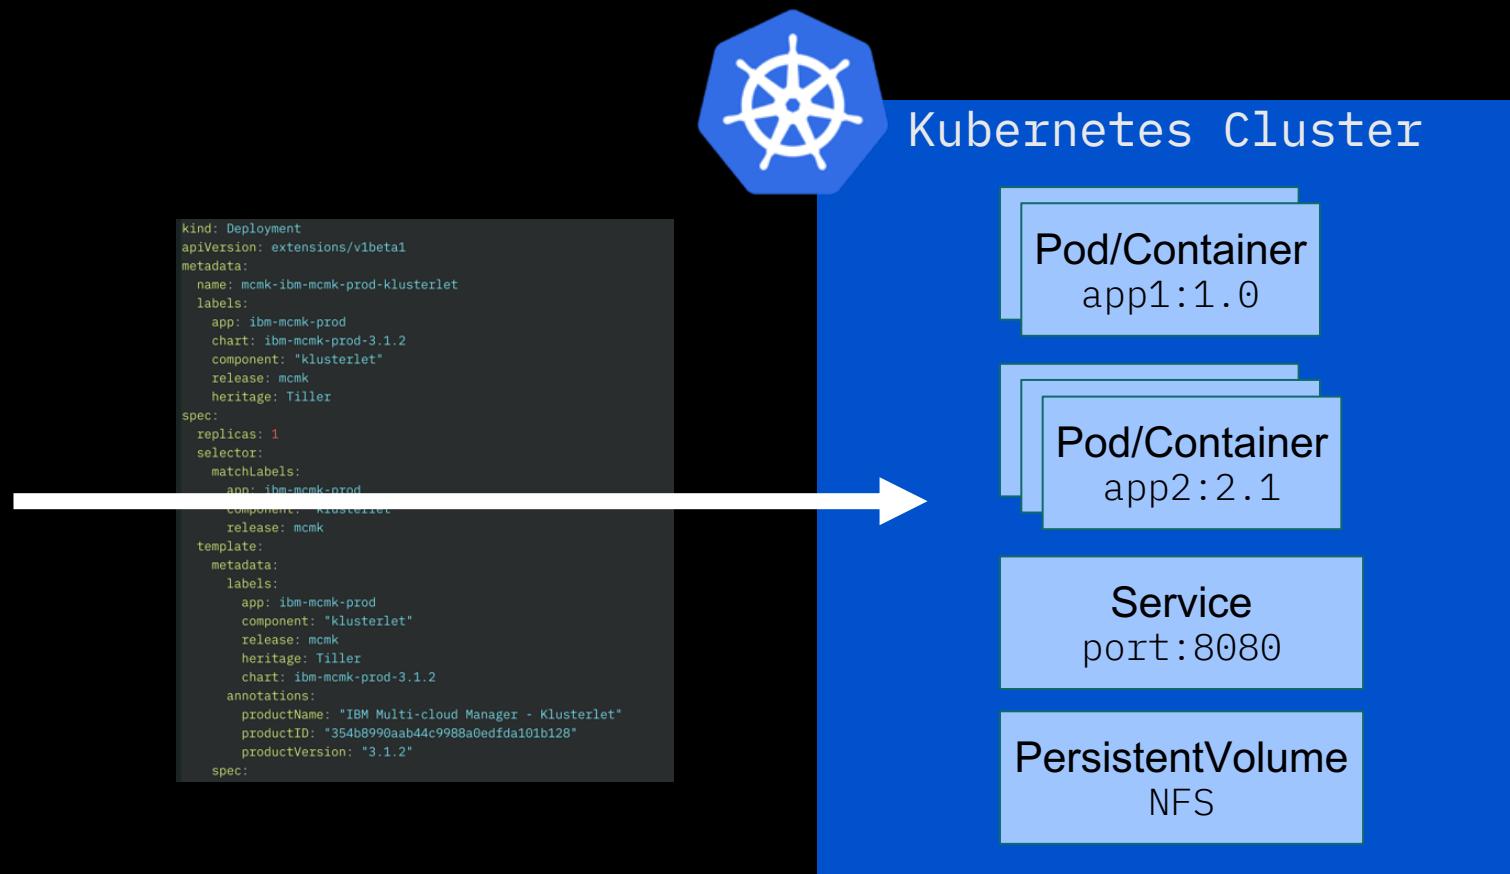
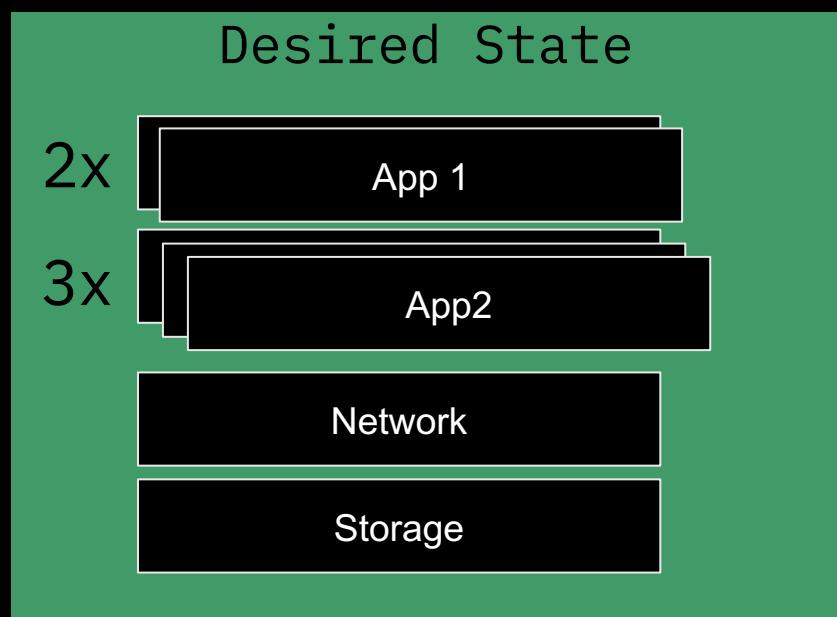
By contrast, in a declarative system, **the user** knows the desired state, supplies a representation of the desired state to the system, then **the system** reads the current state and determines the sequence of commands to transition the system to the desired state.

Kubernetes – Declarative System

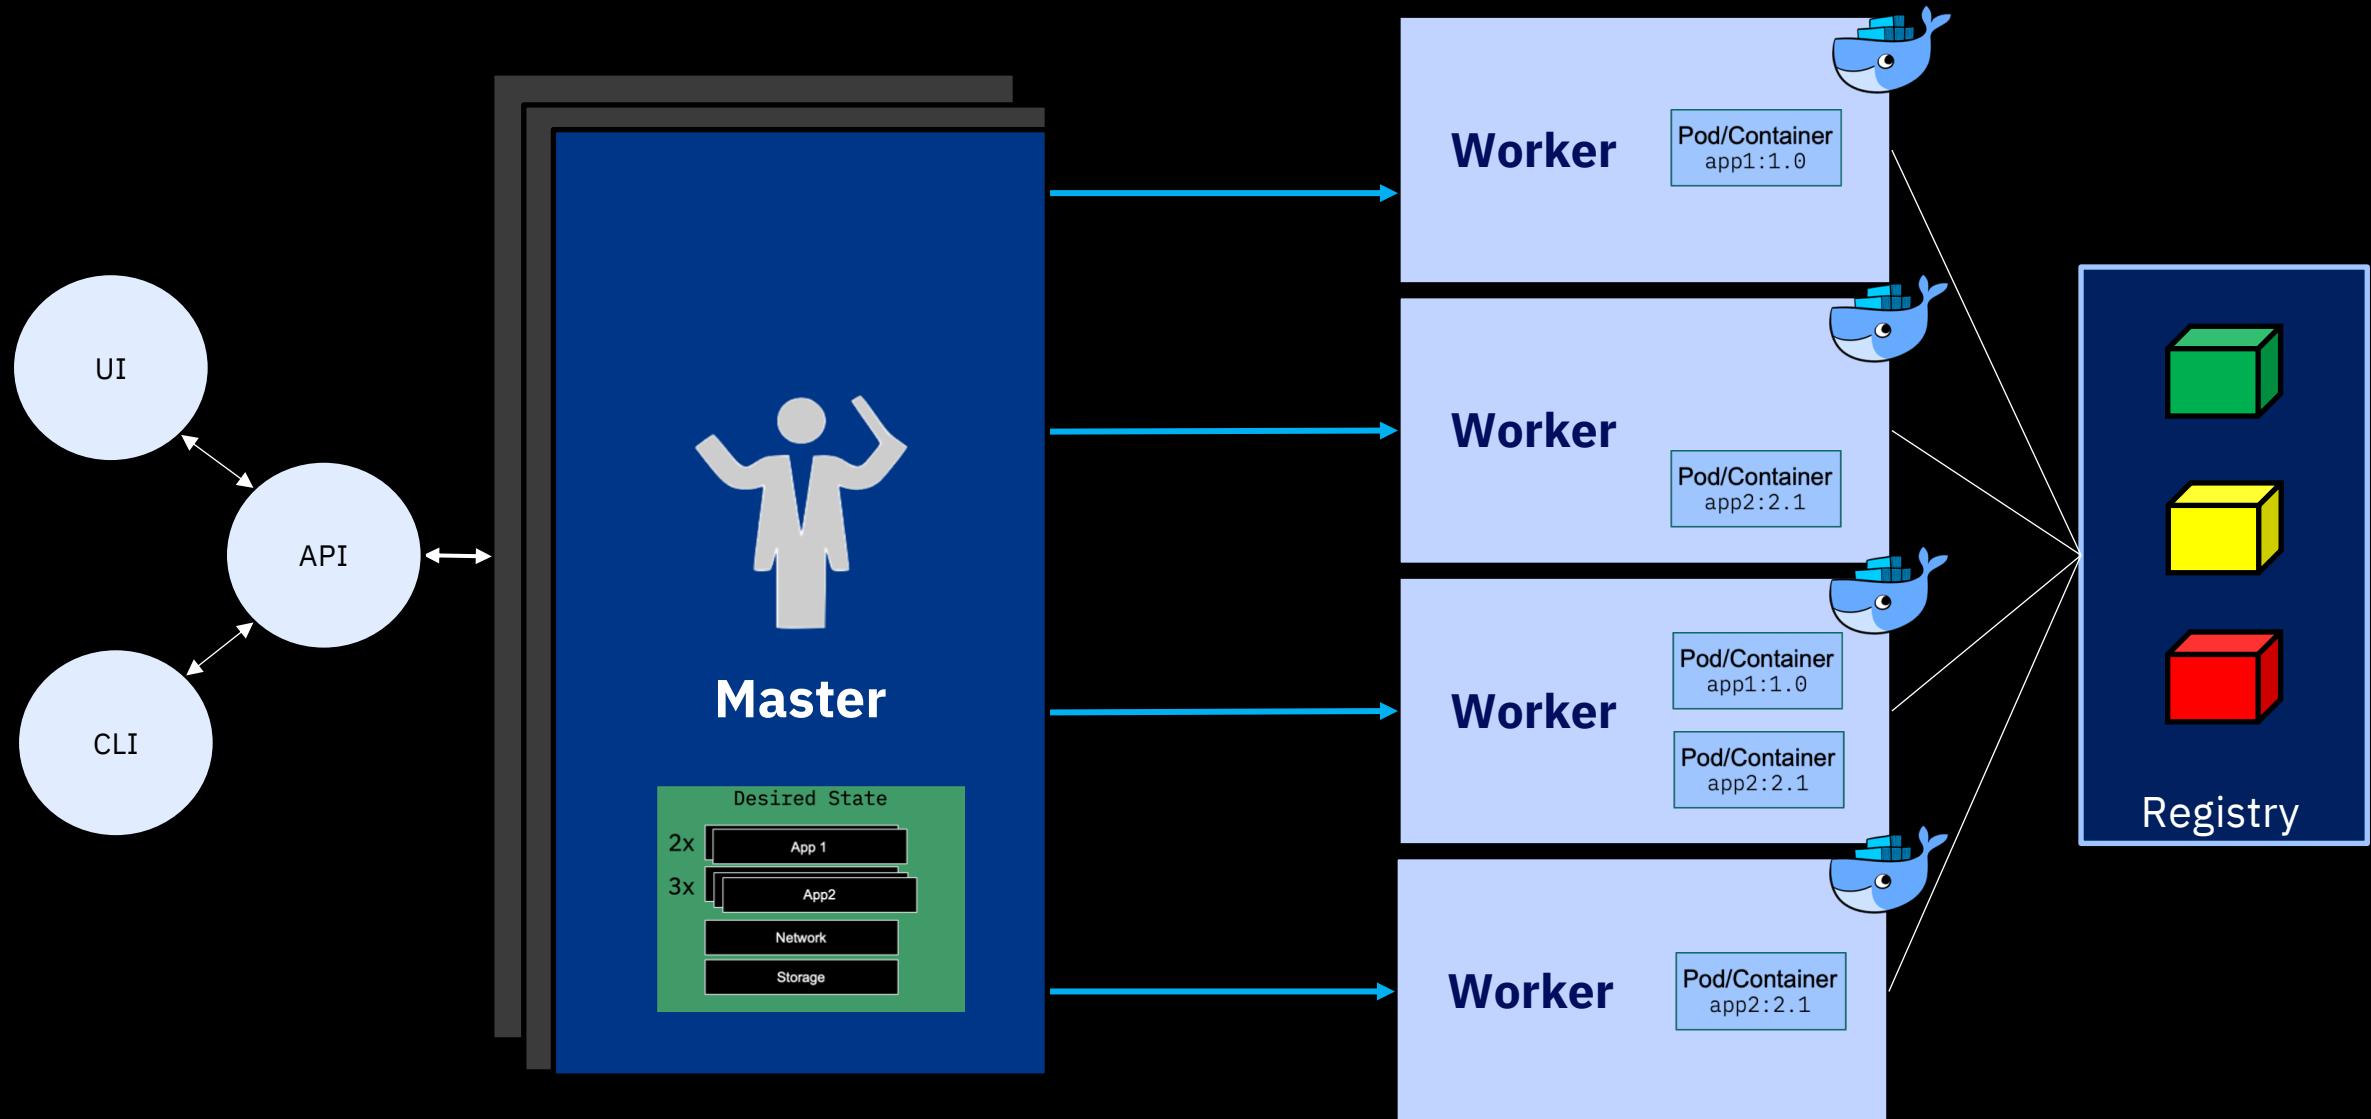


The Desired State

Kubernetes ensures that all the containers running across the cluster are in the desired state at any moment.



Kubernetes Management Architecture





What is Kubernetes?

Container orchestrator

- Runs and manages containers
- Unified API for deploying web applications, batch jobs, and databases
- Maintains and tracks the global view of the cluster
- Supports multiple cloud and bare-metal environments

Manage applications, not machines

- Rolling updates, canary deploys, and blue-green deployments

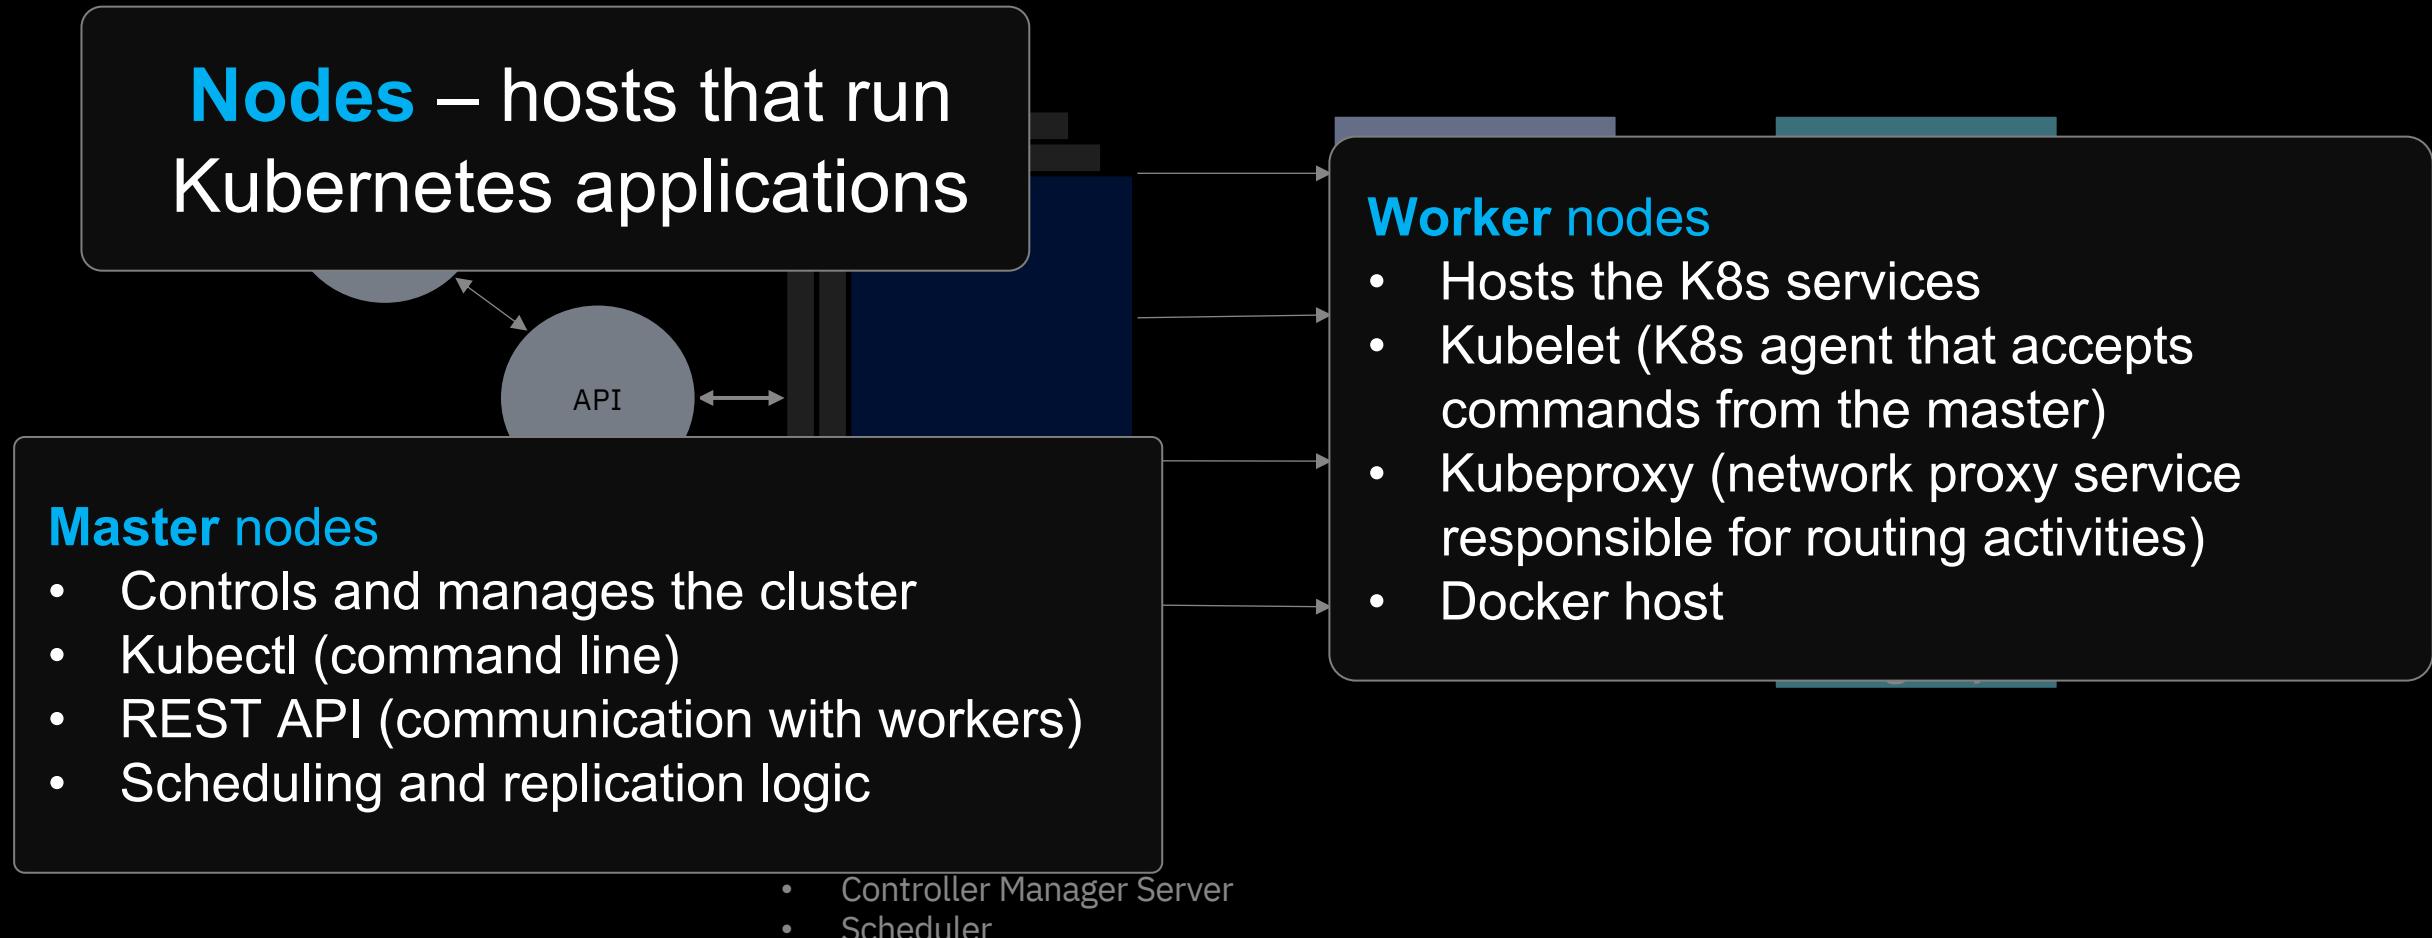
Designed for extensibility

- Rich ecosystem of plug-ins for scheduling, storage, networking

Open source project managed by the Linux Foundation

- Inspired and informed by Google's experiences and internal systems
- 100% open source, written in Go

Kubernetes Management Architecture

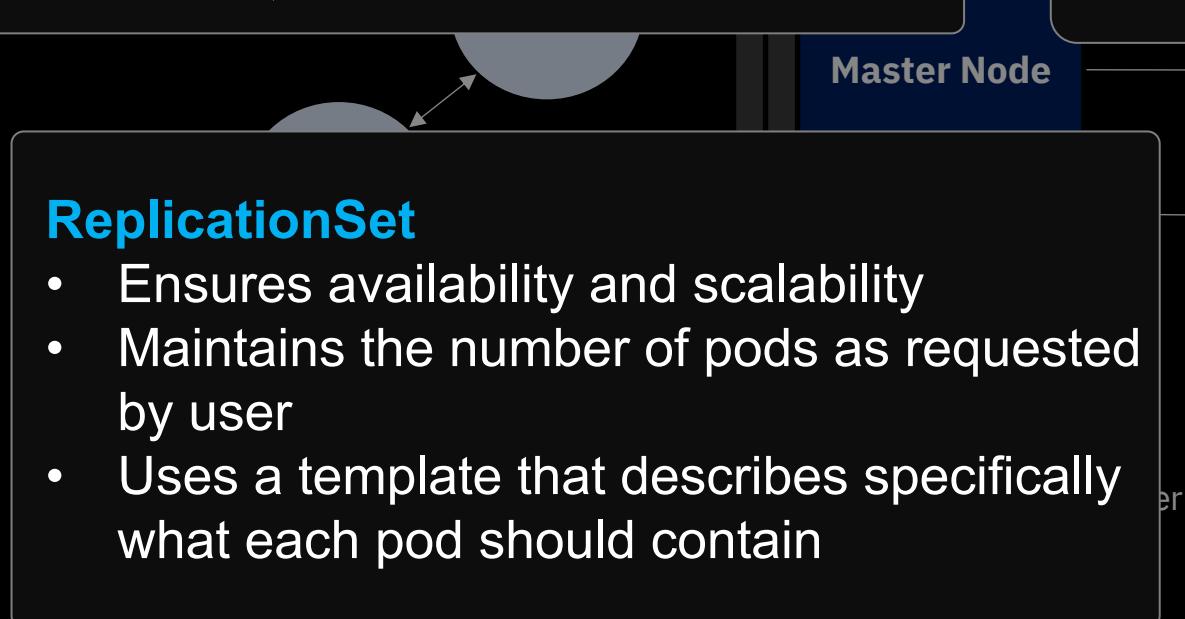


Kubernetes Management Architecture



Pods

- Smallest deployment unit in K8s
- Collection of containers that run on a worker node
- Each has its own IP
- Pod shares a PID namespace, network, and hostname



ReplicationSet

- Ensures availability and scalability
- Maintains the number of pods as requested by user
- Uses a template that describes specifically what each pod should contain

Deployment

- A set of pods to be deployed together
- Declarative - creates a ReplicaSet describing the desired state
- Rollout/ Rollback - Deployment controller changes the actual state to the desired state
- Scale and autoscale: A Deployment can be scaled

Service

- Collections of pods exposed as an endpoint
- A service provides a way to refer to a set of Pods (selected by labels) with a single static IP address

Kubernetes Cluster Architecture

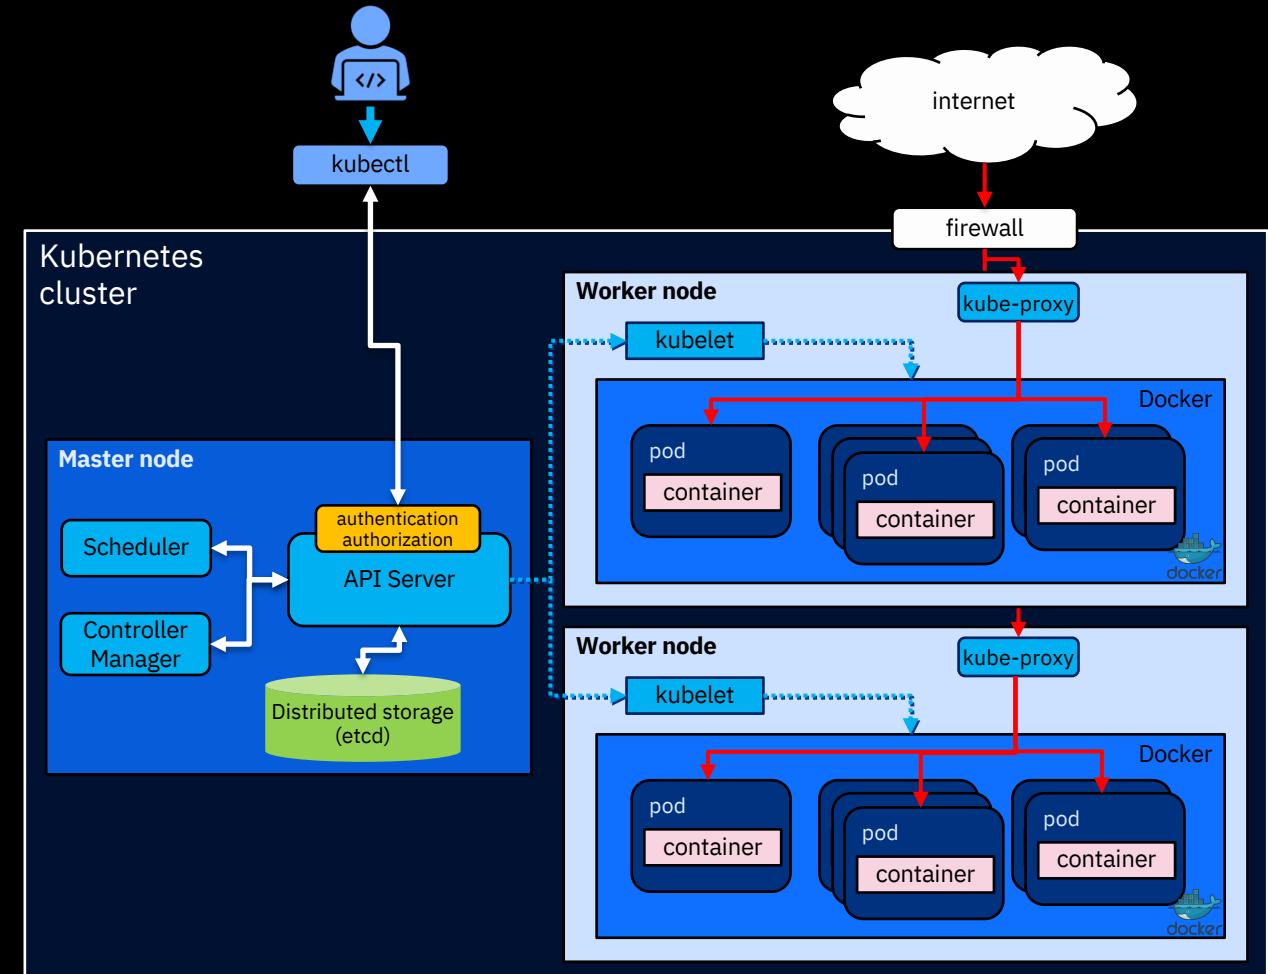


Master node

- Node that manages the cluster
- Scheduling, replication & control
- Multiple nodes for HA

Worker nodes

- Node where pods are run
- Docker engine
- kubelet agent accepts & executes commands from the master to manage pods
- kube-proxy – routes inbound or ingress traffic



kubectl – talking to the Cluster

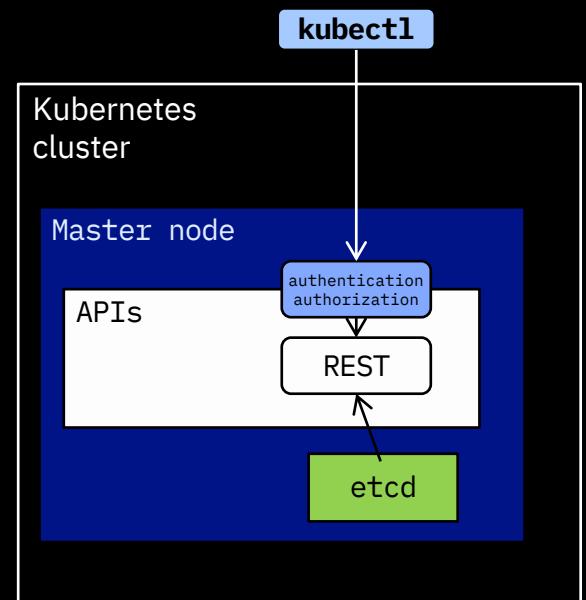


kubectl is a command line interface for running commands against Kubernetes clusters (read state, create objects, ...).

kubectl looks for a file named config in the \$HOME/.kube directory.

Kubernetes uses etcd as a key-value database store.
It stores the configuration of the Kubernetes cluster in etcd.
It also stores the *actual* state of the system and the *desired* state of the system in etcd.

Anything you might read from a `kubectl get xyz` command is stored in etcd.
Any change you make via `kubectl create` will cause an entry in etcd to be updated.



kubectl – talking to the Cluster



kubectl is a command line interface for running commands against Kubernetes clusters.

kubectl looks for a file named config in the \$HOME/.kube directory.

`kubectl [command] [TYPE] [NAME]`

`kubectl create -f example.yaml`

Create objects in yaml file

`kubectl apply -f example.yaml`

Modify objects in yaml file

`kubectl delete -f example.yaml`

Delete objects in yaml file

`kubectl get pods`

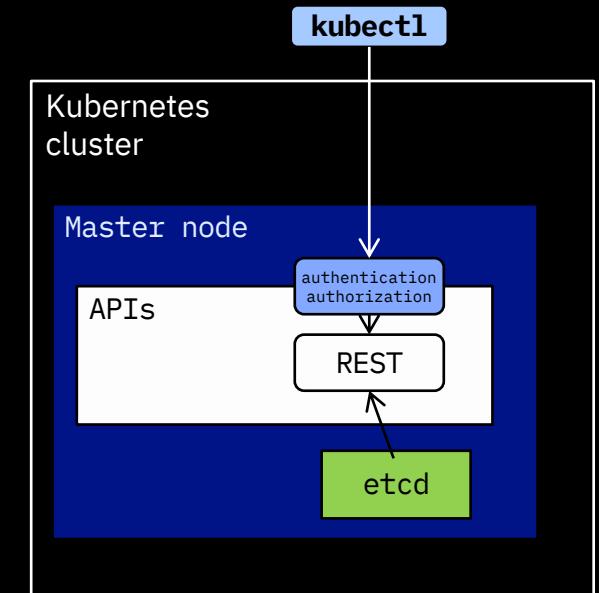
List Pods in Namespace

`kubectl describe nodes <node-name>`

Details about K8s object

`kubectl logs <pod-name>`

Get the logs for a Pod



Pod



- A group of one or more containers is called a pod.
- Containers in a pod are deployed together, and are started, stopped, and replicated as a group.
- **When designing pods ask, “Will these containers work correctly if they land on different machines?”**

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.7.9
      ports:
        - containerPort: 80
```



10.0.0.1

10.0.0.2

Creating a Pod



```
kubectl run nginx --image=nginx:1.7.9  
or  
kubectl create -f example.yaml
```

```
kubectl get pods  
NAME          READY  STATUS    RESTARTS  AGE  
nginx-5bd87f76c-vxc79  1/1    Running   0          29s
```

example.yaml

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: nginx  
spec:  
  containers:  
  - name: nginx  
    image: nginx:1.7.9  
    ports:  
    - containerPort: 80
```



10.0.0.1

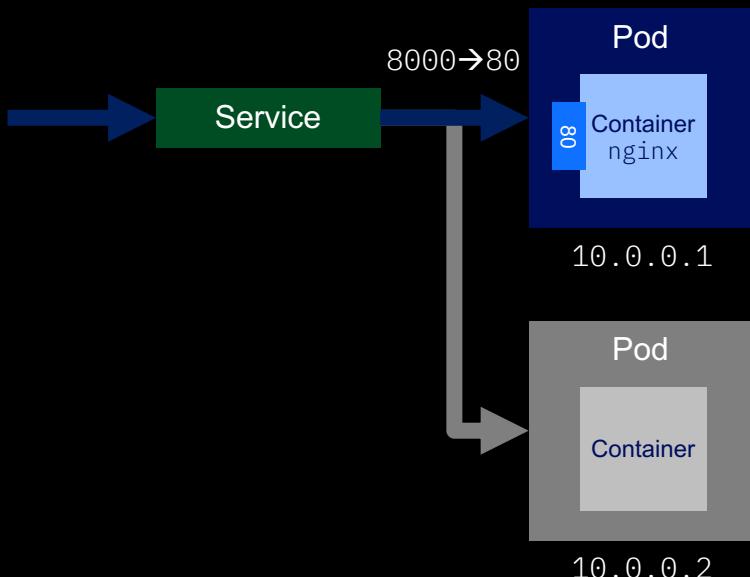


10.0.0.2

Service



- A service provides a way to refer to a set of Pods (selected by labels) with a **single static IP address**.
- Also provide load balancing

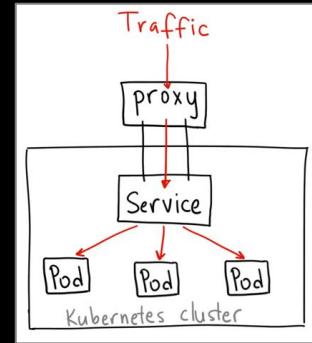


```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  ports:
    - port: 8000
      targetPort: 80
      protocol: TCP
  selector:
    app: nginx
```

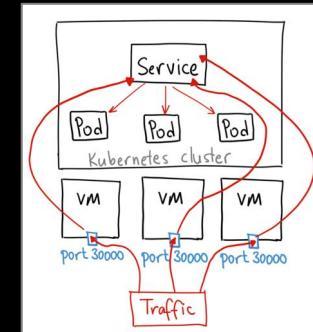
Service



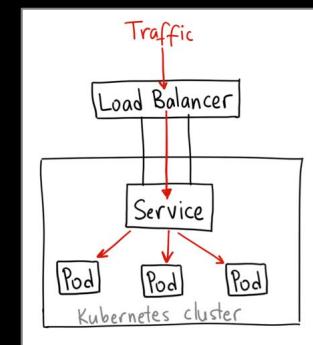
ClusterIP: This type exposes the service on the cluster internal IP. This means that the service is only reachable from within the cluster.



NodePort: This type exposes the service on each Node's static IP address.



LoadBalancer: This service type exposes a service using the cloud provider load balancer.

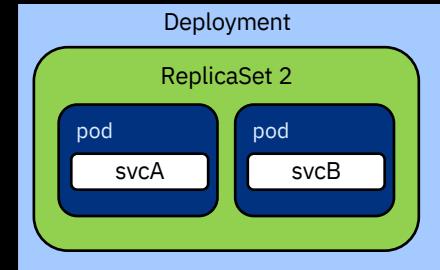


Deployments & ReplicaSets



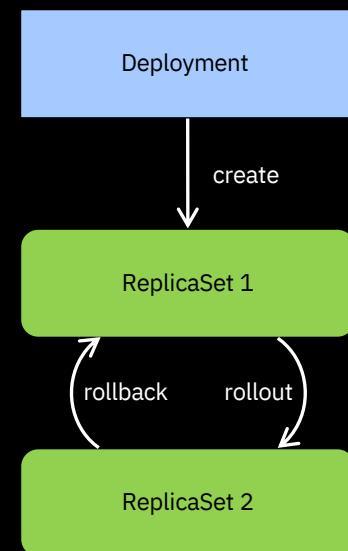
Deployment

- A set of **pods to be deployed together**, such as an application
- **Declarative**: Revising a Deployment **creates a ReplicaSet** describing the desired state
- **Rollout**: Deployment controller changes the actual state to the desired state at a controlled rate
- **Rollback**: Each Deployment revision can be rolled back
- **Scale** and autoscale: A Deployment can be scaled



ReplicaSet

- Cluster-wide pod manager that **ensures the proper number of pods are running** at all times.
- A set of pod templates that describe a set of pod replicas
- Uses a template that describes specifically what each pod should contain
- Ensures that a specified number of pod replicas are running at any given time



Deployment



- A Deployment object defines a Pod creation template and **desired replica count**.
- Create or delete Pods as needed to meet the replica count.
- Manage safely **rolling out changes** to your running Pods.

```
kubectl create -f example.yaml
```

example.yaml

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 2
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.7.9
      ports:
      - containerPort: 80
```

Naming



Name

- Each resource object by type has a unique name

Namespace

- **Resource isolation:** Each namespace is a virtual cluster within the physical cluster
 - Resource objects are scoped within namespaces
 - Low-level resources are not in namespaces: nodes, persistent volumes, and namespaces themselves
 - Names of resources need to be unique within a namespace, but not across namespaces
- **Resource quotas:** Namespaces can divide cluster resources
- Initial namespaces
 - **default** – The default namespace for objects with no other namespace
 - **kube-system** – The namespace for objects created by the Kubernetes system

Resource Quota

- Limits resource consumption per namespace
- Limit can be number of resource objects by type (pods, services, etc.)
- Limit can be total amount of compute resources (CPU, memory, etc.)
- Overcommit is allowed; contention is handled on a first-come, first-served basis

Resources



Linux Foundation – Introduction to Kubernetes

- <https://courses.edx.org/courses/course-v1:LinuxFoundationX+LFS158x+2T2017/course/>

Kubernetes tutorial

- <https://kubernetes.io/docs/tutorials/kubernetes-basics/>

Introduction to container orchestration

- <https://www.exoscale.ch/syslog/2016/07/26/container-orch/>

TNS Research: The Present State of Container Orchestration

- <https://thenewstack.io/tns-research-present-state-container-orchestration/>

Large-scale cluster management at Google with Borg

- <https://research.google.com/pubs/pub43438.html>

Benefits of using a Kubernetes Based Platform

Speed

- Fast boot-time
- Easy scalability
- Rapid deployment

Portability

- Between different environments
- Between private and public cloud

Efficiency

- Better usage of computer resources

Automation

- Next level standardization and automation

13x

More software releases

Eliminate

“works on my machine” syndrome

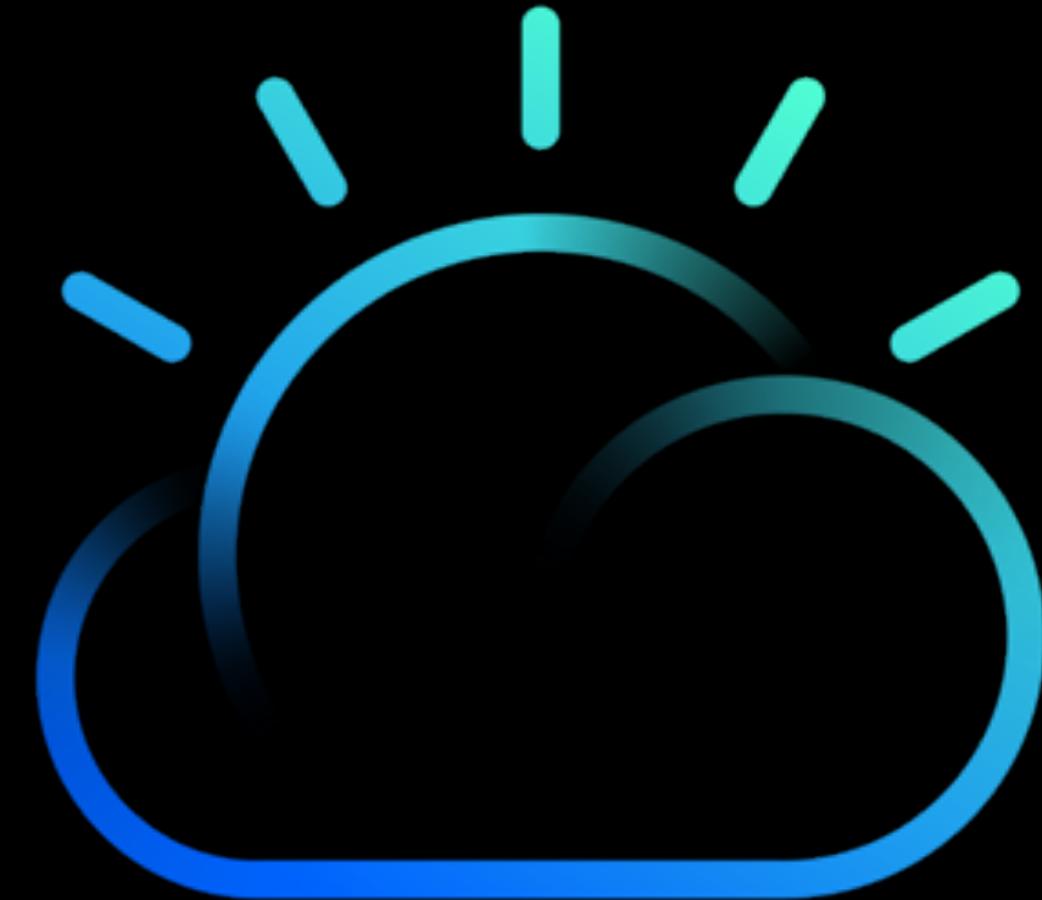
~47%

Reduction of VMs,
OS licensing and
server cost

Reduce

operation cost

QUESTIONS?



The Journey to Cloud
Let's get real

05



IBM Cloud

Microservice transformations – The good, the bad, the ugly

The good

IF implemented correctly Microservices improve scalability and reliability.

- Scale easily and very efficiently
- No single point of failure by isolation
- Be modified without compromising the whole system
- Use of languages and frameworks that best suit the purpose of each service
- Natively ready for CaaS and PaaS

Microservice transformations – The good, the bad, the ugly

The bad

- Added management **complexity**
- **Expertise** required to maintain a microservice-based application
- No or **reduced benefits** if application doesn't need scaling or cloud transformation
- Needs more **robust testing** in order to cover all APIs

Microservice transformations – The good, the bad, the ugly

... and **the downright ugly**

- May require **high initial investment** to run for Dev and Ops overhead
- Higher **initial operational and monitoring costs** through shift in paradigms
- **Non-uniform** application design and architecture through free choice of technologies
- Overload of documentation for every individual component app

Microservices – breaking down the monolith beast

How to identify candidates

- Is there anything that is **scaling differently** than the rest of the system?
- Is there anything that feels “**tacked-on**”?
- Is there anything **changing much faster** than the rest of the system?
- Is there anything requiring more **frequent deployments** than the rest of the system?
- Is there a part of the system that a small team, **operates independently**?
- Is there a **subset of tables** in your datastore that isn’t connected to the rest of the system?

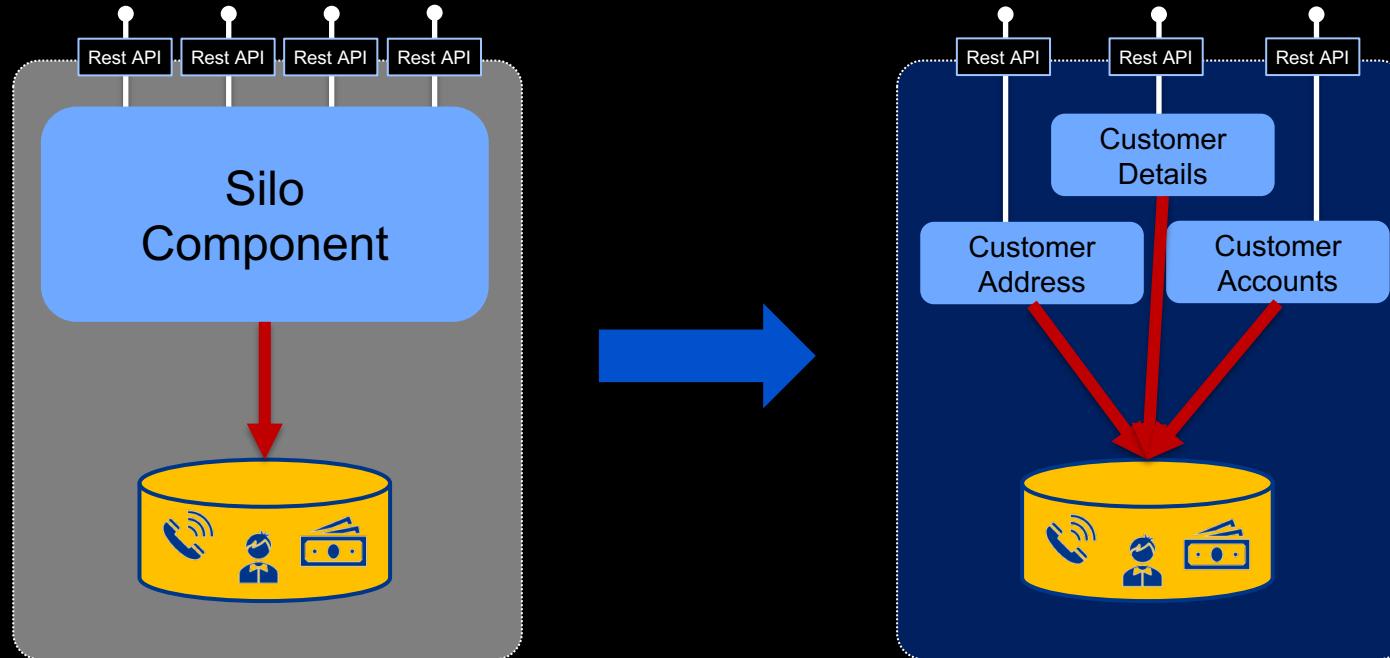
Containerizing your WAR file
doesn't mean you're doing
microservices.

- What is the domain? What is reality?
- Where are the transactional boundaries?
- How should microservices communicate across boundaries?
- What if we just turn the database inside out?

Microservices – Database refactoring

Data Isolation

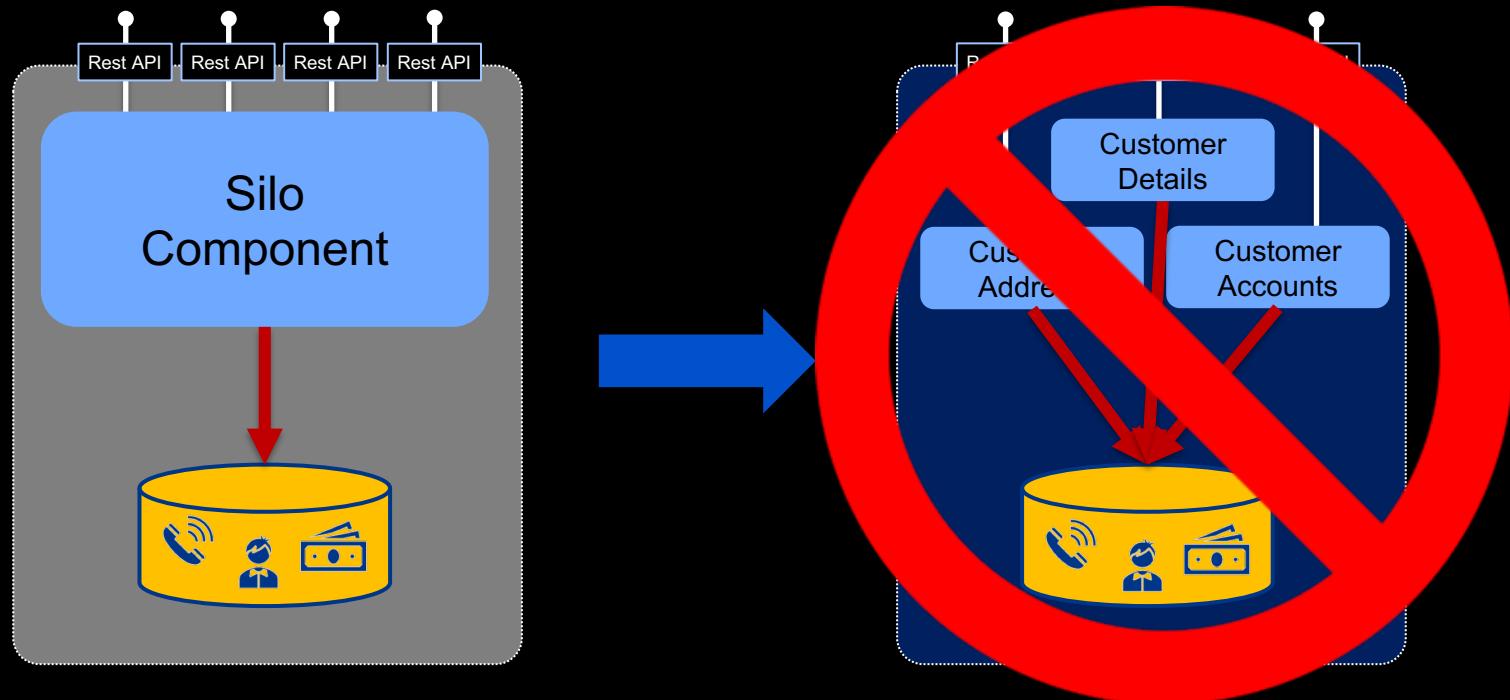
Each microservice must keep its own data. One typical error that happens in the beginning of a microservice transformation is the use of a centralized database, the same way it was used in the monolithic application. This creates a single point of failure and dependency between the microservices.



Microservices – Database refactoring

Data Isolation

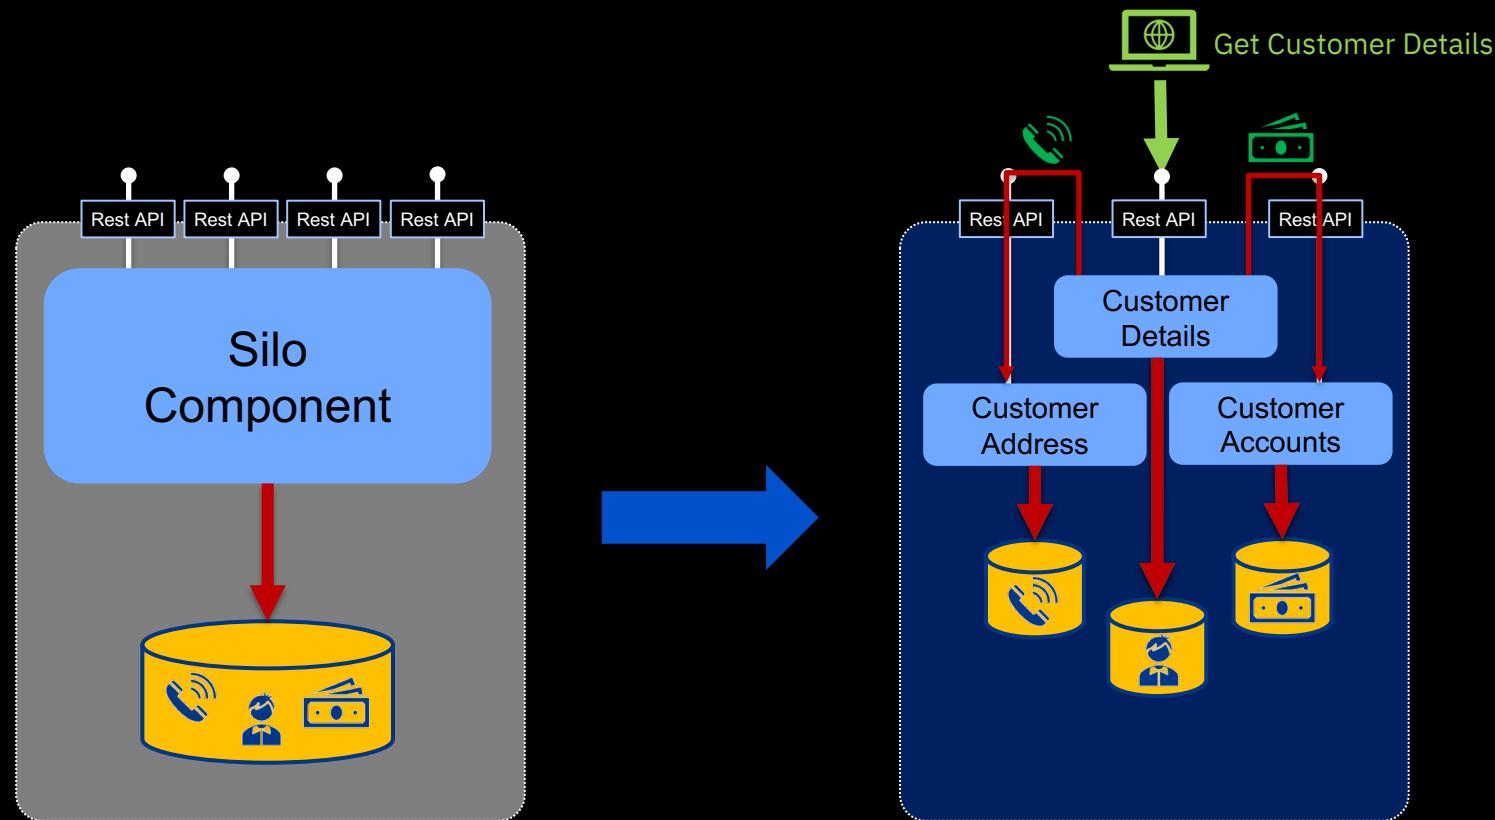
Each microservice must keep its own data. One typical error that happens in the beginning of a microservice transformation is the use of a centralized database, the same way it was used in the monolithic application. This creates a single point of failure and dependency between the microservices.



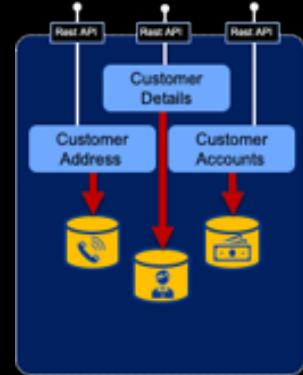
Microservices – Database refactoring

Data Isolation

Each microservice must keep its own data. One typical error that happens in the beginning of a microservice transformation is the use of a centralized database, the same way it was used in the monolithic application. This creates a single point of failure and dependency between the microservices.



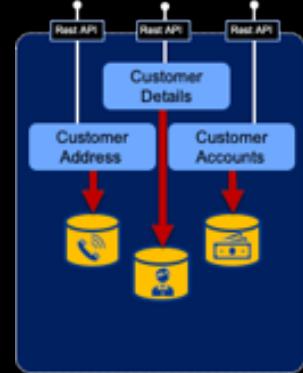
Microservices – per-service database



Benefits

- Avoid unexpected data modification – the API enforces consistency
- Make changes to our system without blocking progress of other parts of the system
- Store the data in a project-owned databases, using the appropriate technology
- Make changes to the schema/databases at our leisure
- Become much more scalable, fault tolerant, and flexible

Microservices – per-service database



Drawbacks

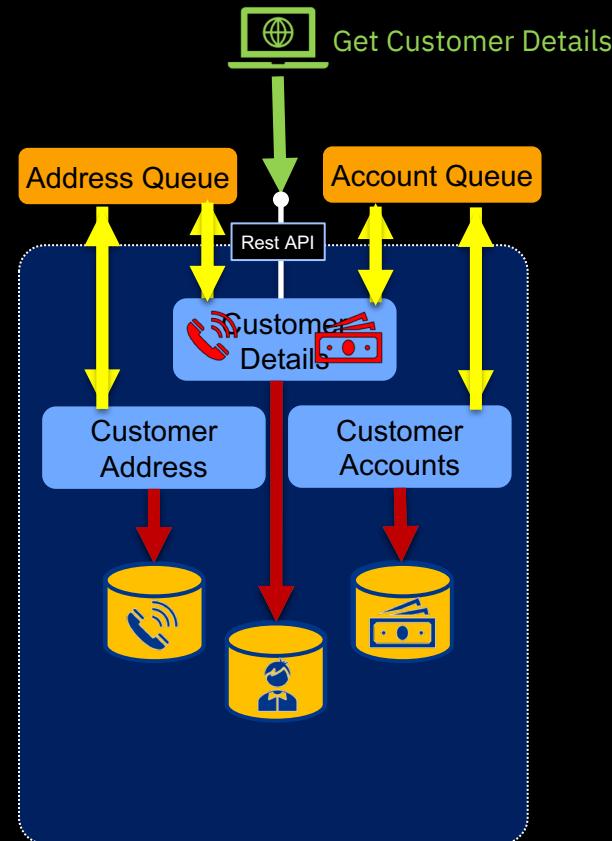
- Costly to refactor
- Needs robust transaction handling by the service (cost!)
- More difficult to debug
- More difficult to operationalize
- More difficult to test (needs well designed API tests)

Event Driven Architecture

Microservices – one step further

Event-Driven Architecture

- Common pattern to maintain data consistency across different services.
- Avoid waiting for ACID transactions to complete
- Makes your application more available and performant
- Provides loose coupling between services



Microservices Good or bad idea?

Microservices – reasons not to use them

Some thoughts

- If **speed** is your goal, microservices aren't the solution (MVP, doesn't have to scale)
- Not all applications are **large enough** to break down into microservices
- Applications that **require tight integration** between individual components and services (real-time processing, ...)
- A moderately large, moderately complex application being maintained by a relatively small development and operations team
- At what point is it in its **lifecycle**? Mission-critical?

Microservices – reasons not to use them

How to detect antipatterns

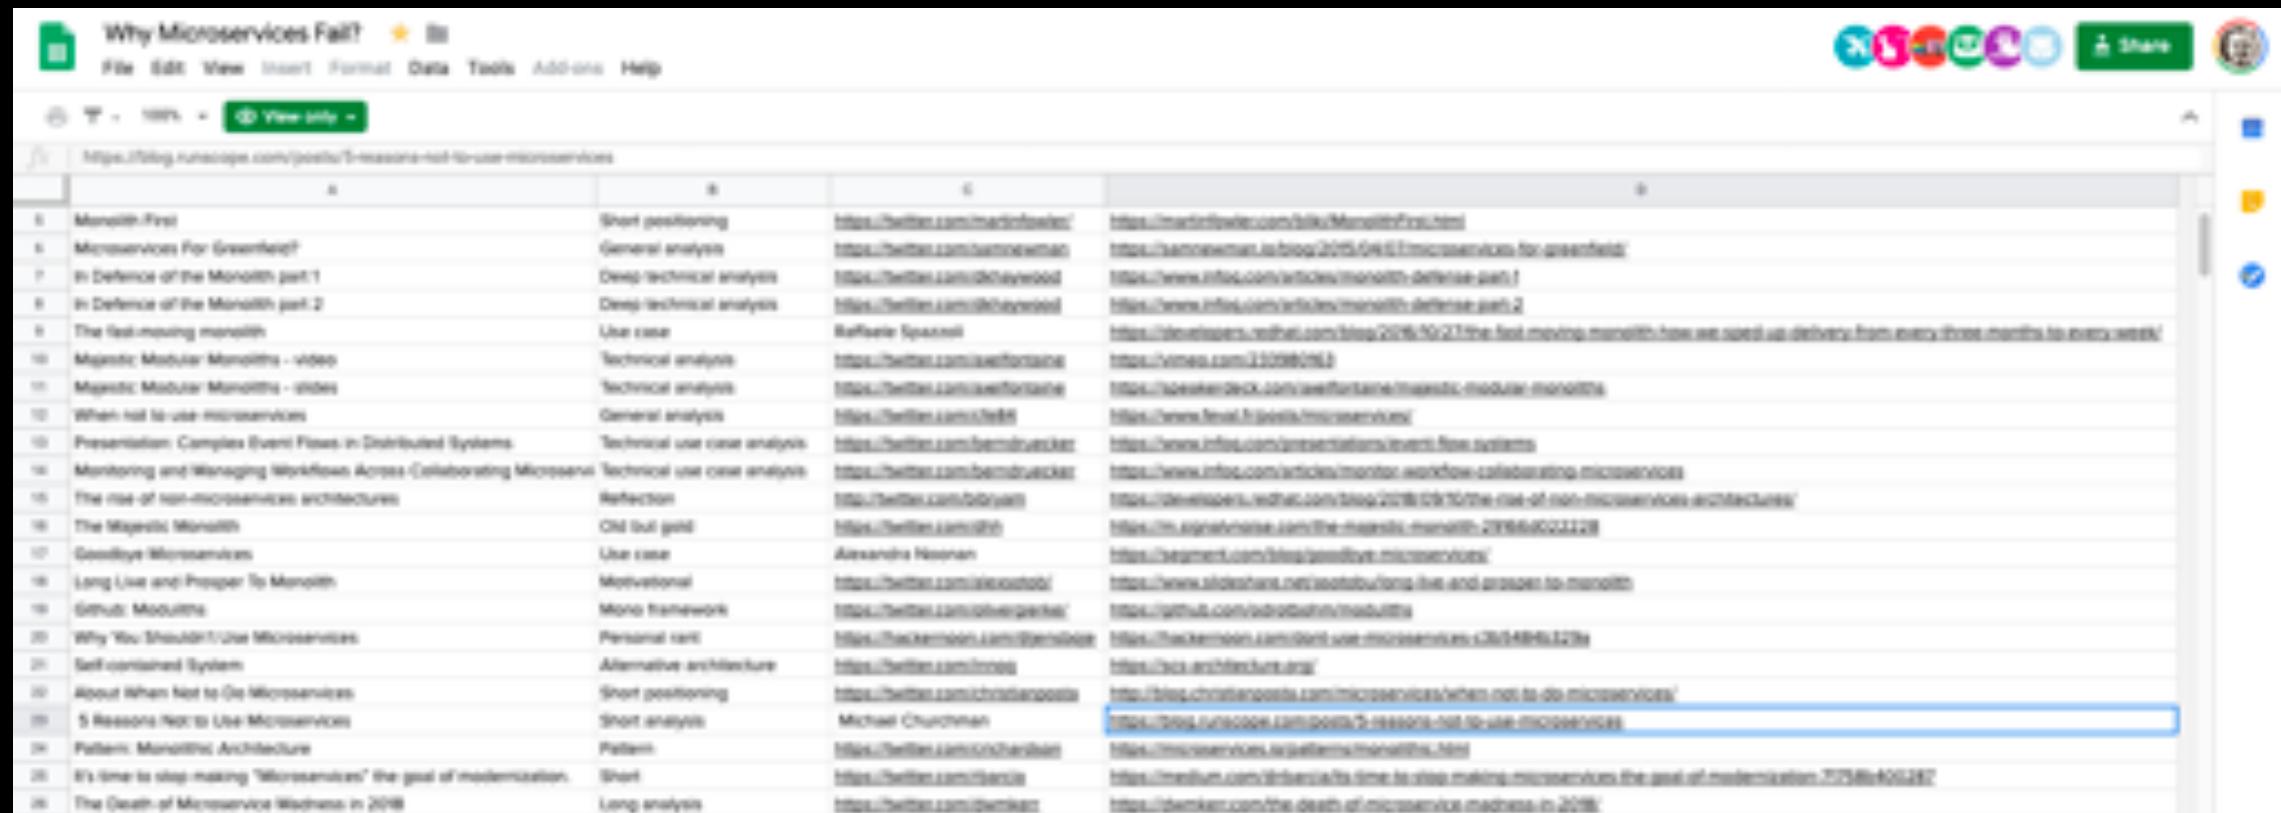
- A change to one microservice often requires **changes to other microservices**
- Many of your microservices **share a datastore**
- Deploying one microservice requires other microservices to be **deployed at the same time**
- Your microservices are **overly chatty**
- The same developers **work across a large number of microservices**
- Your microservices **share a lot of the same code or models**

Microservices – reasons not to use them

What works for **large and complex** applications
does not always work at a **smaller scale**

What makes sense for a **new application**
does not always make sense when maintaining or
updating **existing applications**

Microservices – reasons not to use them



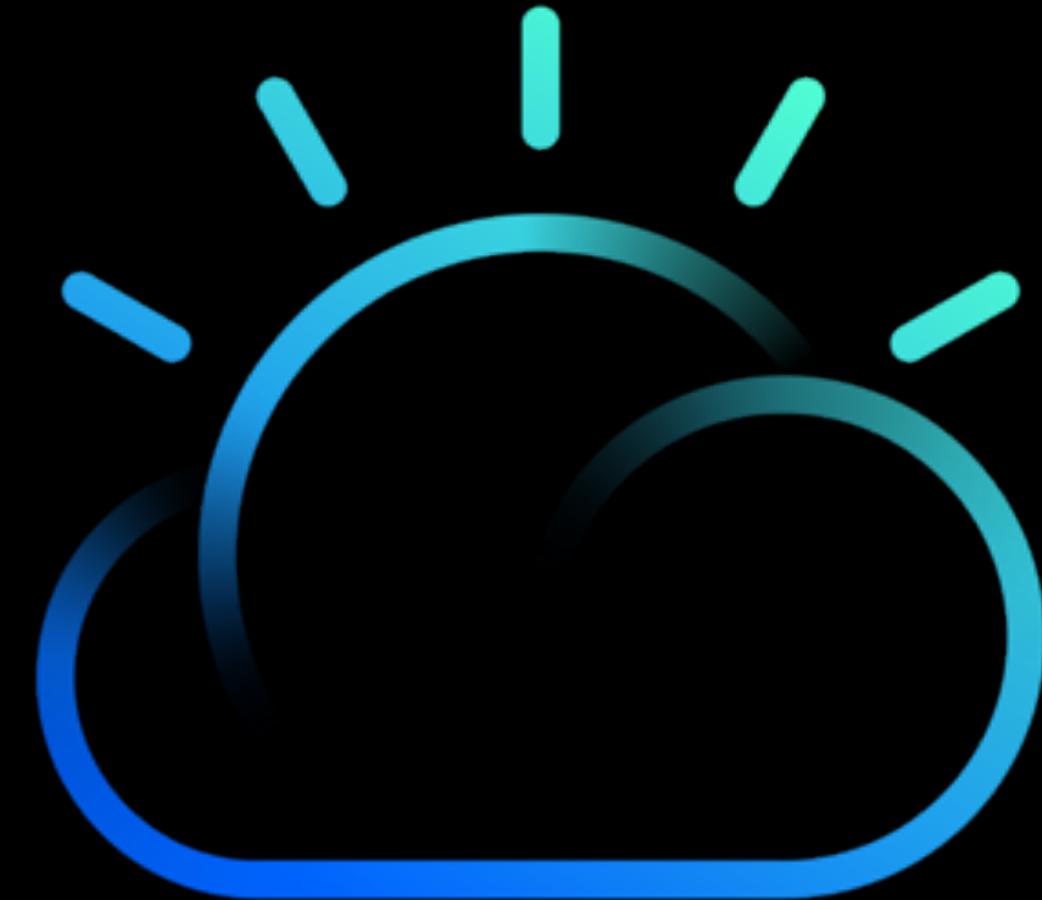
The screenshot shows a Google Sheets document with a title bar 'Why Microservices Fail?' and a toolbar with various icons. The main content is a table with four columns: 'A' (Index), 'B' (Category), 'C' (Type), and 'D' (URL). The table lists 26 items, each with a URL starting with <https://twitter.com>. The items are:

| A | B | C | D |
|----|--|-----------------------------|--|
| 1 | Monolith First | Short positioning | @michaelfowler |
| 2 | Microservices For Greenfield? | General analysis | @bscooperman |
| 3 | In Defence of the Monolith part 1 | Deep technical analysis | @dchaywood |
| 4 | In Defence of the Monolith part 2 | Deep technical analysis | @dchaywood |
| 5 | The fast-moving monolith | User case | @brianf@spacemil |
| 6 | Magnetic Modular Monoliths - video | Technical analysis | @awertson |
| 7 | Magnetic Modular Monoliths - slides | Technical analysis | @awertson |
| 8 | When not to use microservices | General analysis | @steve |
| 9 | Presentation: Complex Event Flows in Distributed Systems | Technical use case analysis | @brettwuecker |
| 10 | Monitoring and Managing Workflows Across Collaborating Microservices | Technical use case analysis | @brettwuecker |
| 11 | The rise of non-microservices architectures | Reflection | @brianw |
| 12 | The Magnetic Monolith | Old school | @09 |
| 13 | Goodbye Microservices | User case | @alexander |
| 14 | Long Live and Prosper To Monolith | Motivational | @alexisotoma |
| 15 | GitHub: Modular | Mono framework | @ulrichgasper |
| 16 | Why You Shouldn't Use Microservices | Personal rant | @john |
| 17 | Self-contained System | Alternative architecture | @lnn |
| 18 | About When Not to Use Microservices | Short positioning | @christianposta |
| 19 | 5 Reasons Not to Use Microservices | Short analysis | @michael |
| 20 | Pattern: Monolithic Architecture | Pattern | @michaelfowler |
| 21 | It's time to stop making "Microservices" the goal of modernization | Short | @bscooper |
| 22 | The Death of Microservice Madness in 2018 | Long analysis | @dchaywood |

https://docs.google.com/spreadsheets/d/1vjnjAI_8TZBv2XhFHra7kEQzQpOHSZpFIWDjynYYf0/edit?pli=1#gid=0

Microservices, when
implemented incorrectly,
can make poorly written
applications even more
dysfunctional

QUESTIONS?



IBM Cloud

The Journey to Cloud **Kubernetes Security**



IBM Cloud

Kubernetes – Security – Back to basics

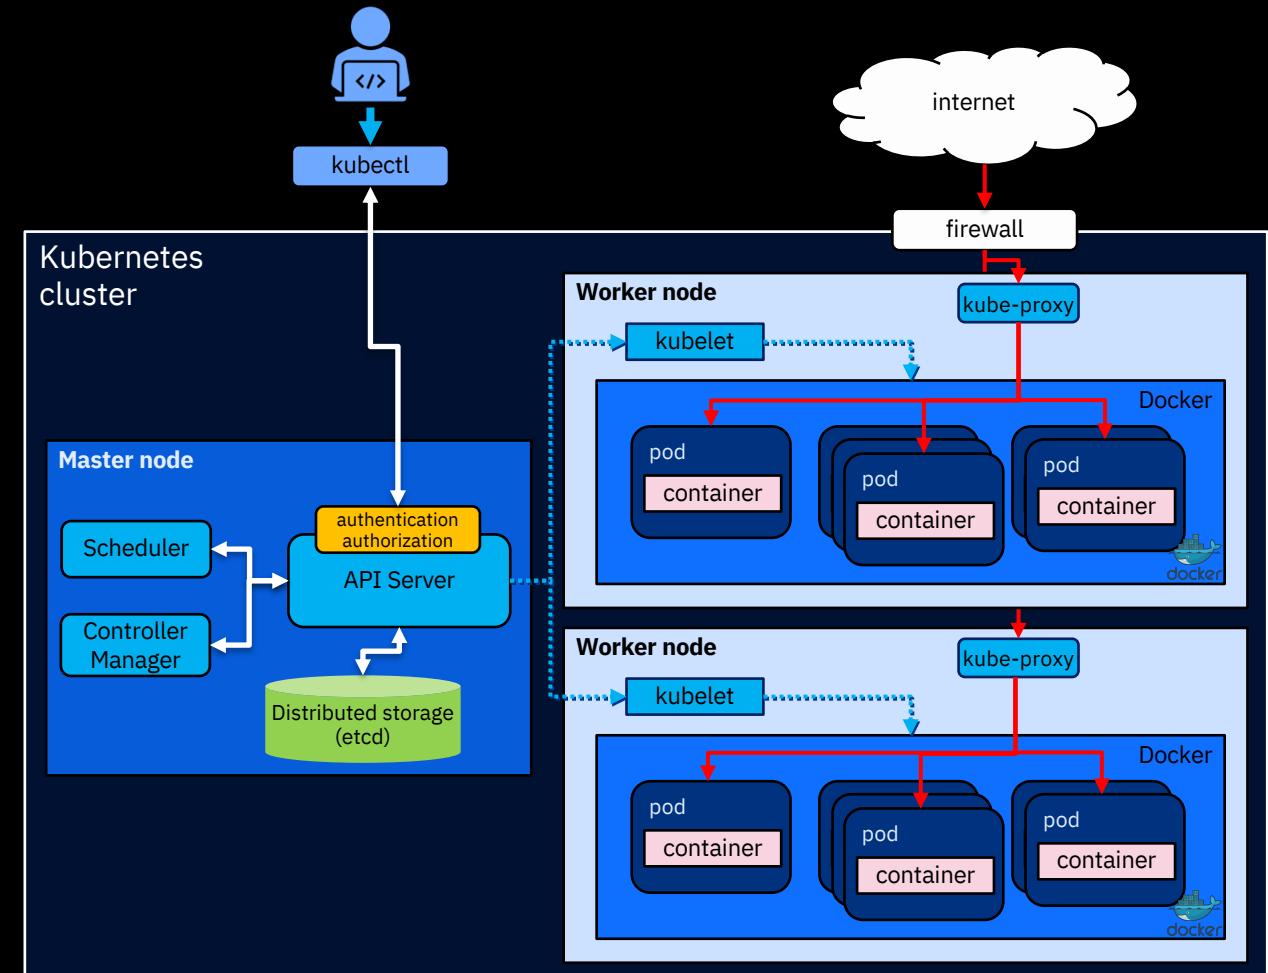


Master node

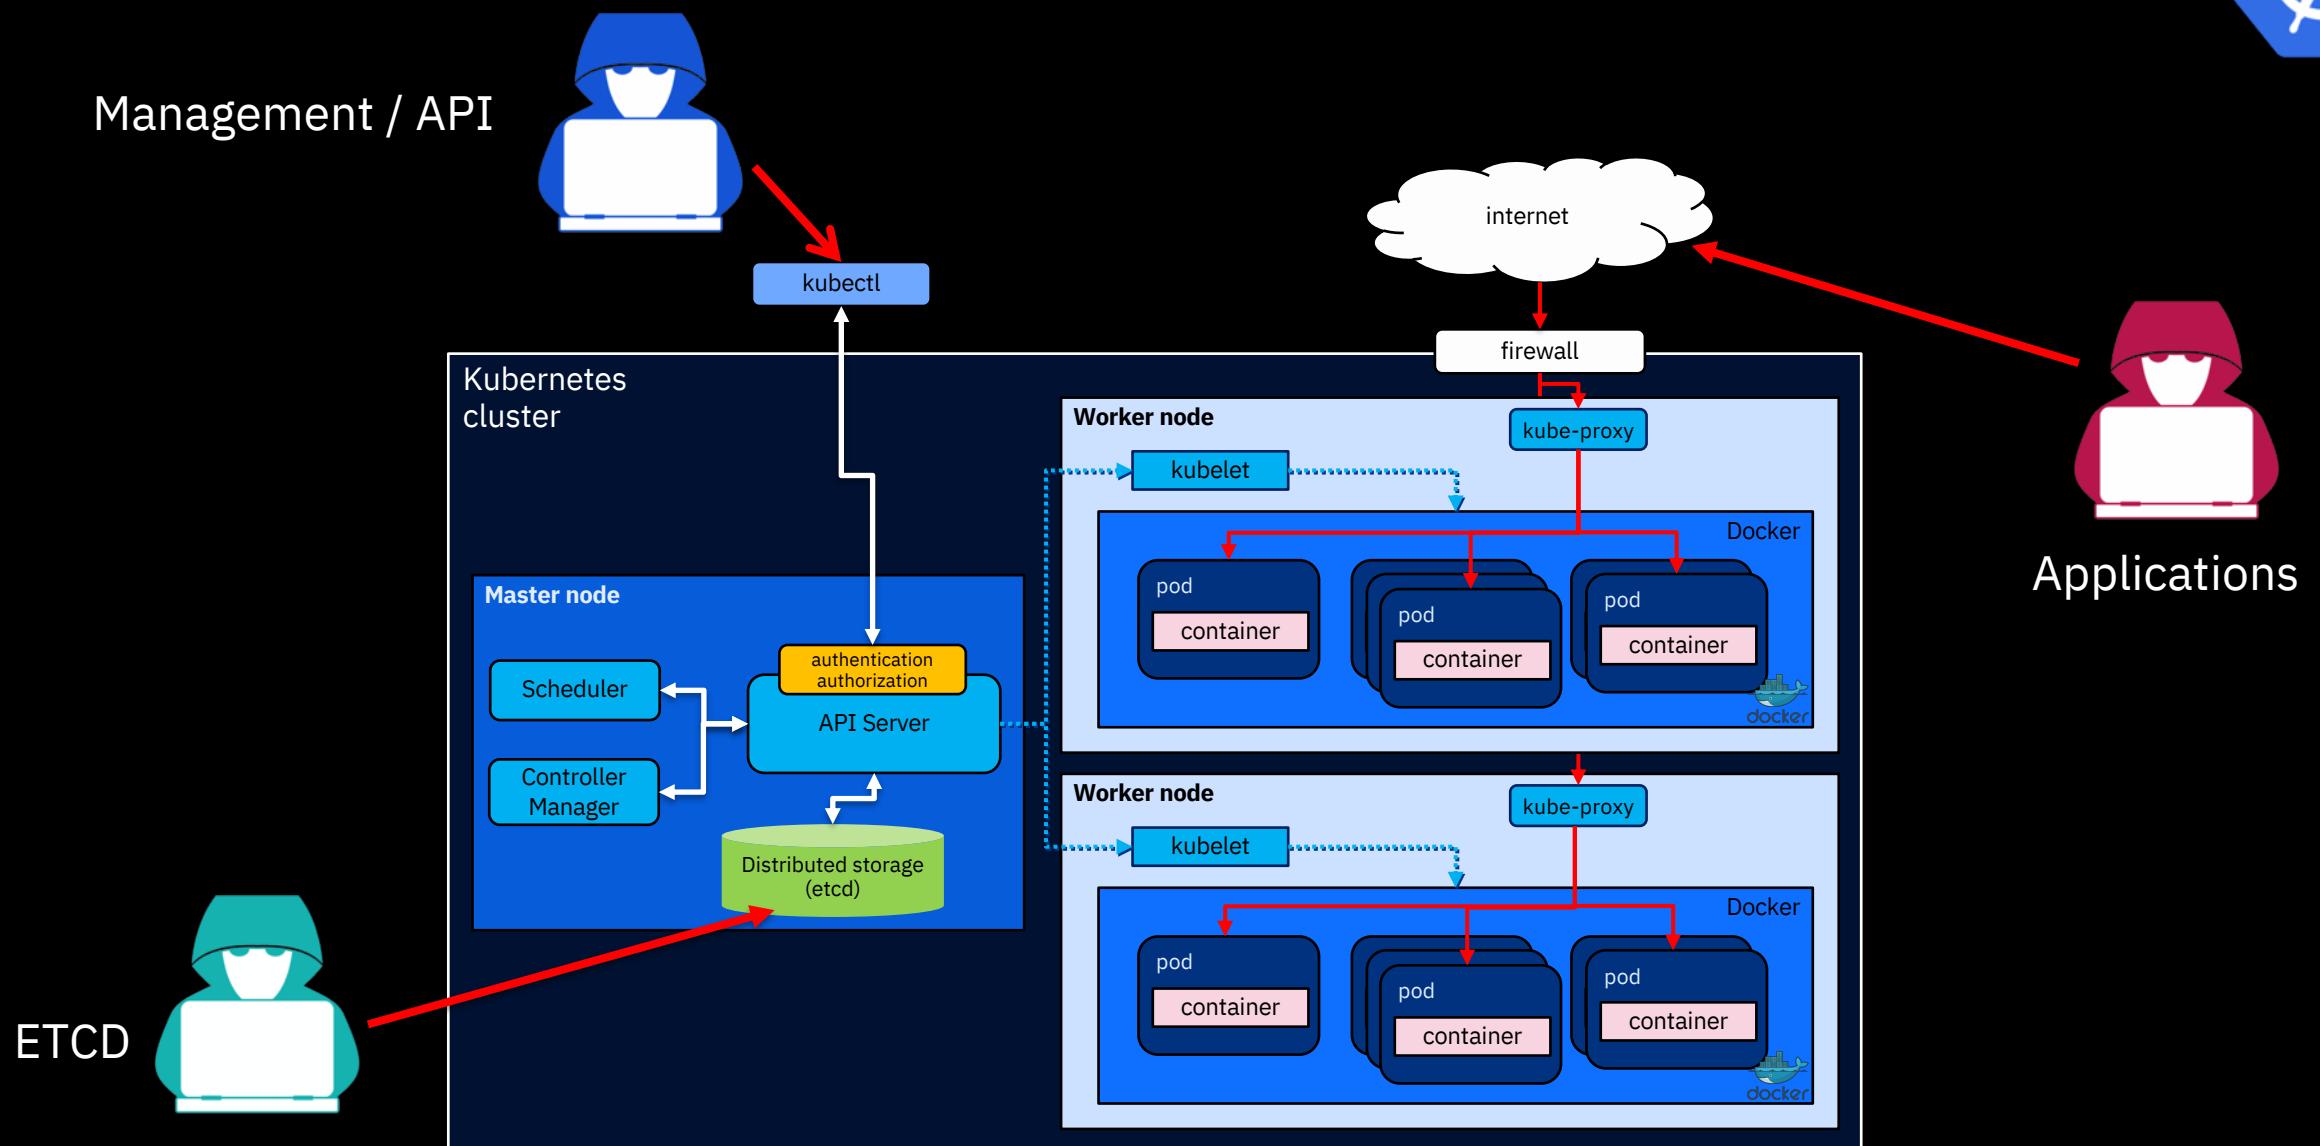
- Node that manages the cluster
- Scheduling, replication & control
- Multiple nodes for HA

Worker nodes

- Node where pods are run
- Docker engine
- kubelet agent accepts & executes commands from the master to manage pods
- kube-proxy – routes inbound or ingress traffic



Kubernetes – Security – Attack surface



Kubernetes – Security Basic Topics

The **Billion Laughs** attack is a particular type of denial of service (DoS) attack which is aimed specifically at XML document parsers. This attack is also referred to as an XML bomb or an exponential entity expansion attack.

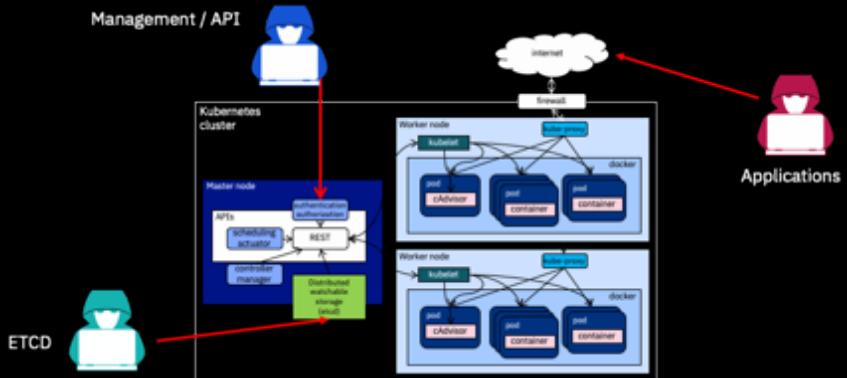
It exploits the fact that nested references to nodes can grow very large when expanded. Because the **kube-apiserver** doesn't perform validation on the manifest, it doesn't detect if those nested references will cause a problem. If the nesting references grow too large, excessive CPU and RAM usage can render the apiserver unresponsive to connections ... hence the Denial of Service.



Kubernetes – Security Basic Topics

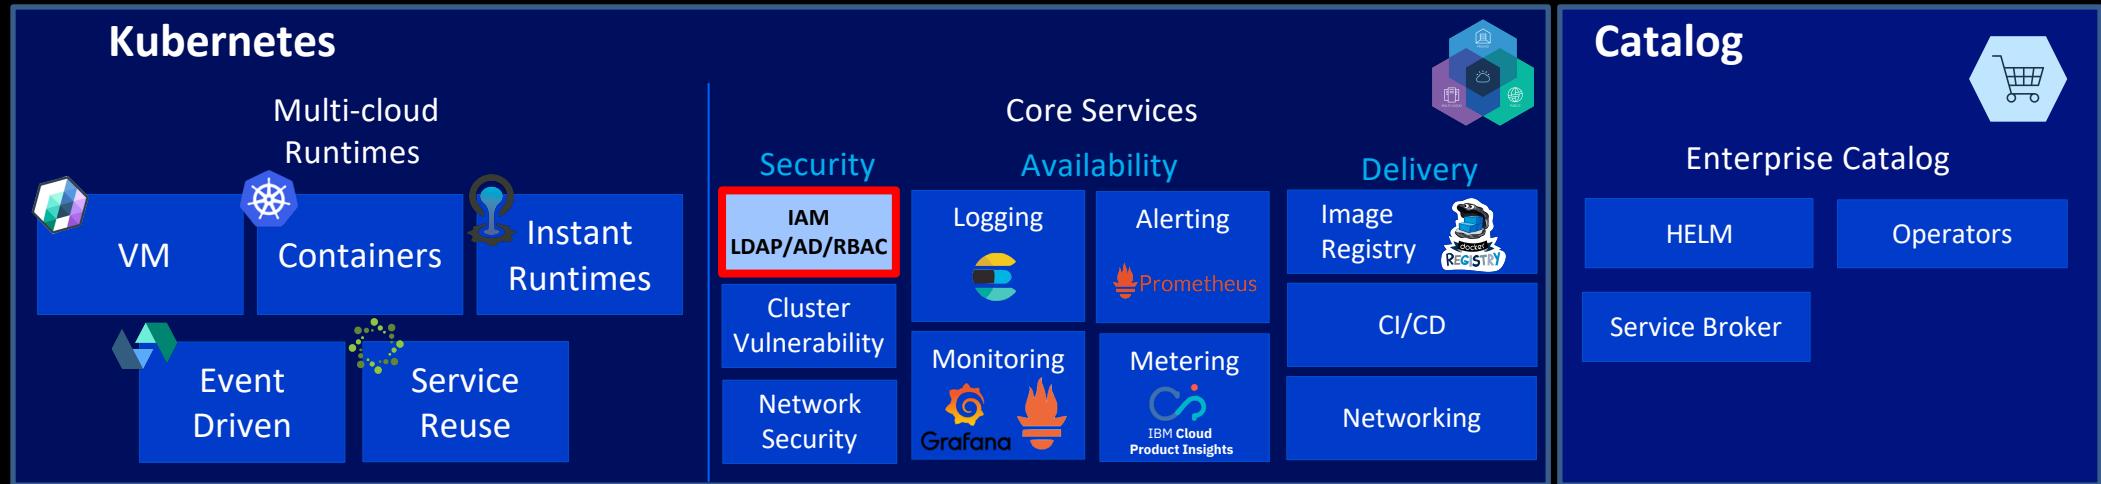
Reduce Kubernetes Attack Surfaces

- Secure access to etcd
- Controlling access to the Kubernetes API
- Controlling access to the Kubelet
- Enforce resource usage limits for workloads
- Rotate infrastructure credentials frequently
- Enable audit logging
- Use Linux security features
- Controlling what privileges containers run with
- Enforcing Network Policies
- Image Scanning of your containers
- Use Kubernetes secrets



Source: <https://kubernetes.io/docs/tasks/administer-cluster/securing-a-cluster>
<https://kubernetes.io/blog/2018/07/18/11-ways-not-to-get-hacked/>

Kubernetes – Core Services - Controlling K8s Access



Authentication – WHO am I - (token, certs, OpenID Connect Provider (OIDC)...)

Developer – create, view, get permission.

Tester – create, delete, update, get permission.

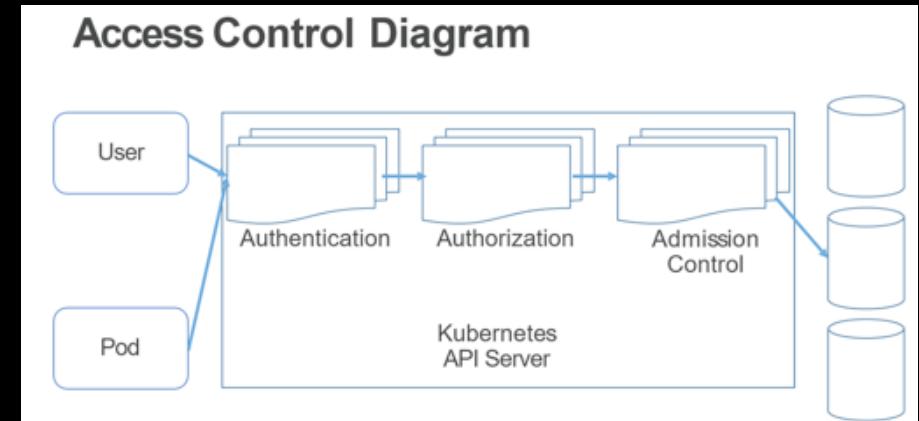
Administrator – All permission

Authorization – CAN I - (RBAC)

Is the user or application authorized or have permissions to access a K8s object?

Admission Control

Intercepts requests to the Kubernetes API server prior to persistence of the object, but after the request is authenticated and authorized



RBAC Basic Elements

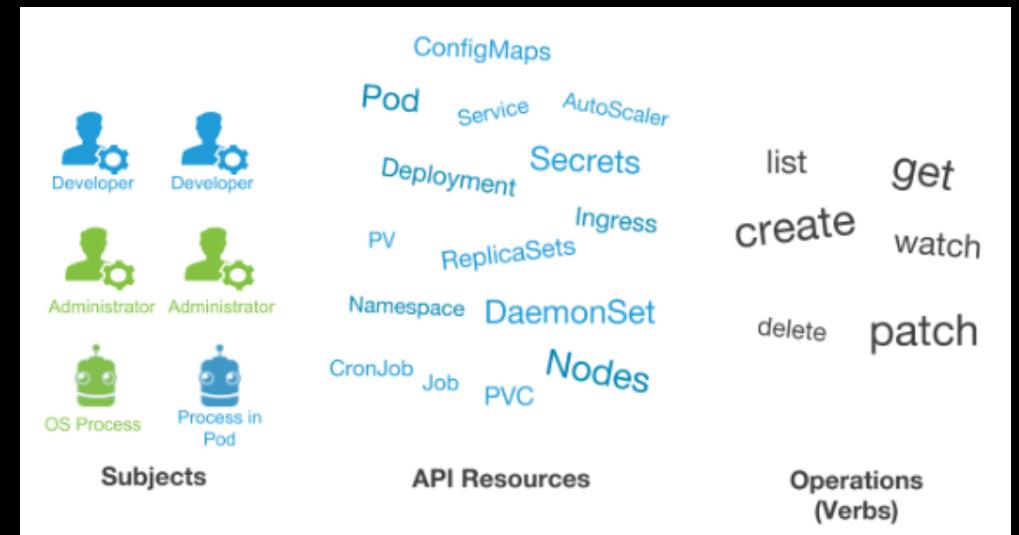
RBAC Building Blocks

Objects

- Pods
- PersistentVolumes
- ConfigMaps
- Deployments
- Nodes
- Secrets
- Namespaces

Verbs

- create
- get
- delete
- list
- update
- edit
- watch
- exec



RBAC Basic Elements

Rules

Operations (verbs) that can be carried out on a group of resources which belong to different API Groups.

Roles and ClusterRoles

Both consist of rules.

- Role: applicable to a single namespace
- ClusterRole: is cluster-wide

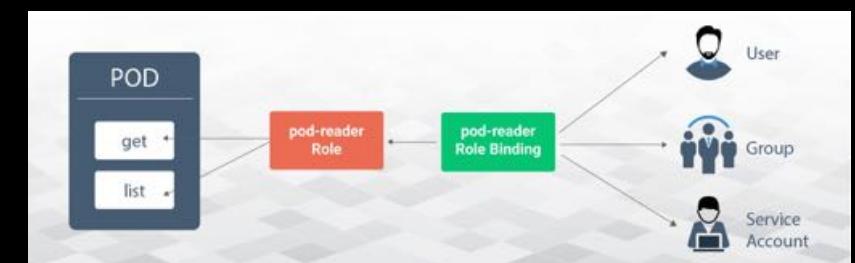
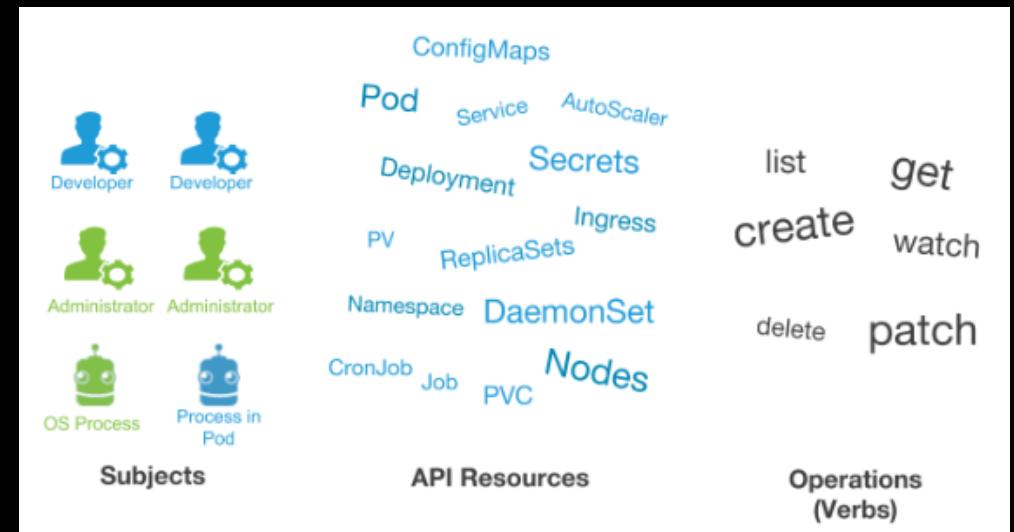
Subjects

Entity that attempts an operation in the cluster

- User Accounts: Humans or processes living outside the cluster.
- Service Accounts: Namespaced account.
- Groups: Reference multiple accounts. (default groups like cluster-admin)

RoleBindings and ClusterRoleBindings

Bind subjects to roles



RBAC Example



```
kind: Role
metadata:
  name: pod-reader
rules:
- apiGroups: [""]
  resources: ["pods"]
  verbs: ["get", "list"]
```

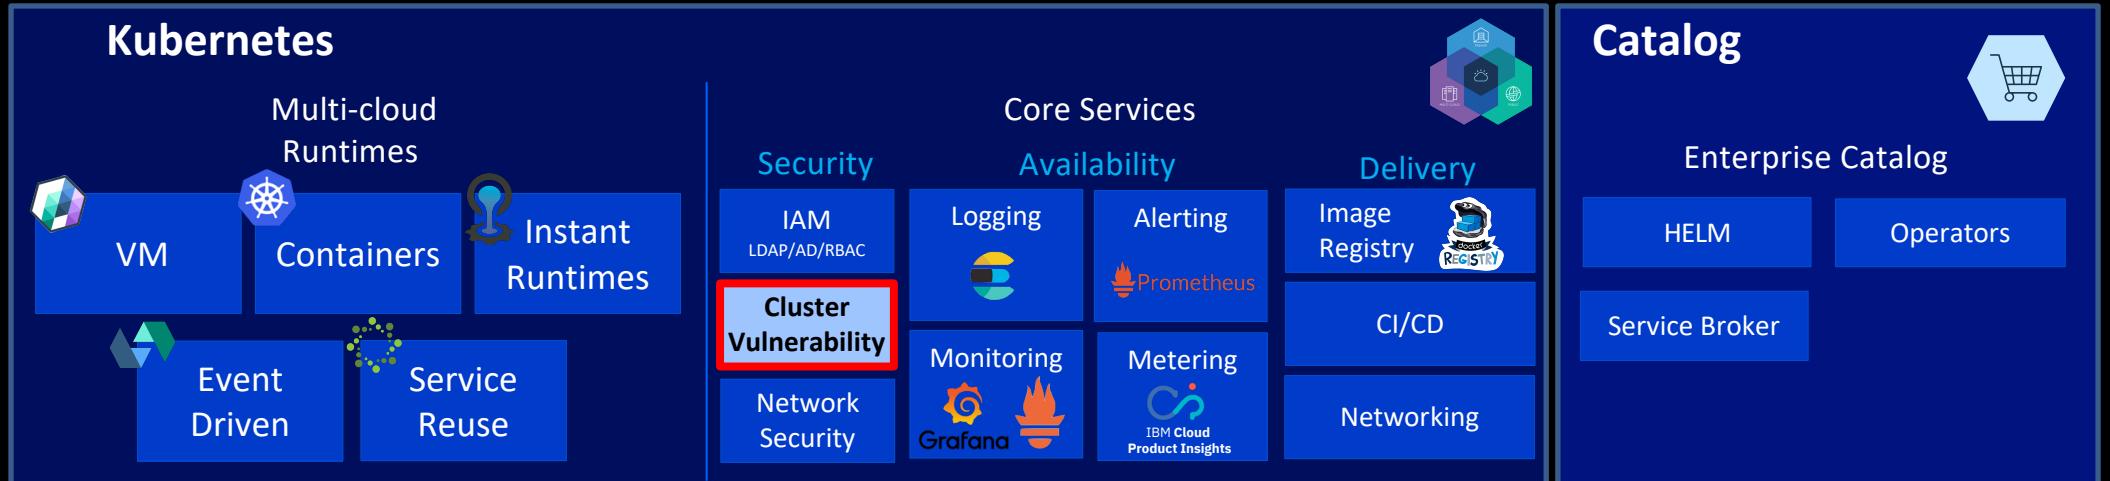
```
kind: RoleBinding
metadata:
  name: pod-reader
subjects:
- kind: User
  name: John
roleRef:
  kind: Role
  name: pod-reader
```

Common RBAC Pitfalls

- **Cluster Administrator Role Granted Unnecessarily**
The built-in **cluster-admin** role grants effectively unlimited access to the cluster.
Should be granted only to the specific users that need it.
- **Duplicated Role Grant**
Role definitions may overlap with each other, making access revocation more difficult.
- **Unused Role**
Roles that are created but not granted to any subject can increase the complexity of RBAC management.
- **Grant of Missing Roles**
Role bindings can reference roles that do not exist.
If the role name is reused those role bindings can unexpectedly become active.

Always adapt
least privileged access
practices

Kubernetes – Core Services – Cluster Security

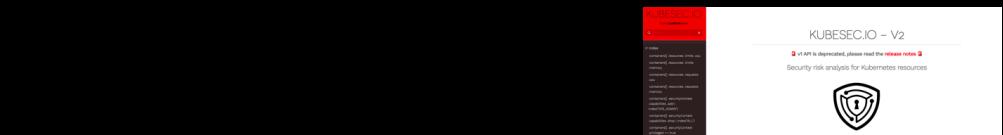


Kubernetes – Cluster Security – Secure Deployments

Scan for K8s Best Practices

kubesc – Validation of security best practices

Performs security risk analysis for Kubernetes resources and tells you what you should change in order to improve the security of those pods. It also gives you a score that you can use to create a minimum standard. The score incorporates a great number of Kubernetes best practices.



```
1  {
2    "object": "Deployment/k8sdemo.default",
3    "valid": true,
4    "message": "Passed with a score of 4 points",
5    "score": 4,
6    "scoring": {
7      "advise": [
8        {
9          "points": 1
10         },
11         {
12           "selector": ".spec .serviceAccountName",
13           "reason": "Service accounts restrict Kubernetes API access and should be configured with least privilege",
14           "points": 3
15         },
16         {
17           "selector": ".metadata .annotations .\\\"container.seccomp.security.alpha.kubernetes.io/pod\\\"",
18           "reason": "Seccomp profiles set minimum privilege and secure against unknown threats",
19           "points": 1
20         }
21       ],
22       "warning": [
23         {
24           "points": 1
25         }
26       ]
27     }
28   }
```

```
{  
  "selector": ".spec .serviceAccountName",  
  "reason": "Service accounts restrict Kubernetes API access and should be configured with least privilege",  
  "points": 3  
},
```

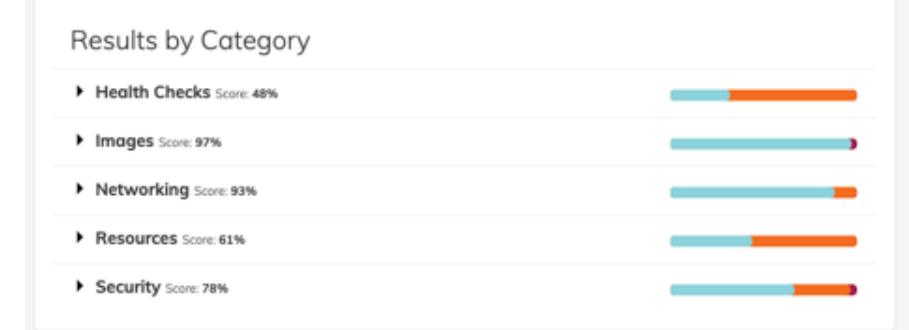
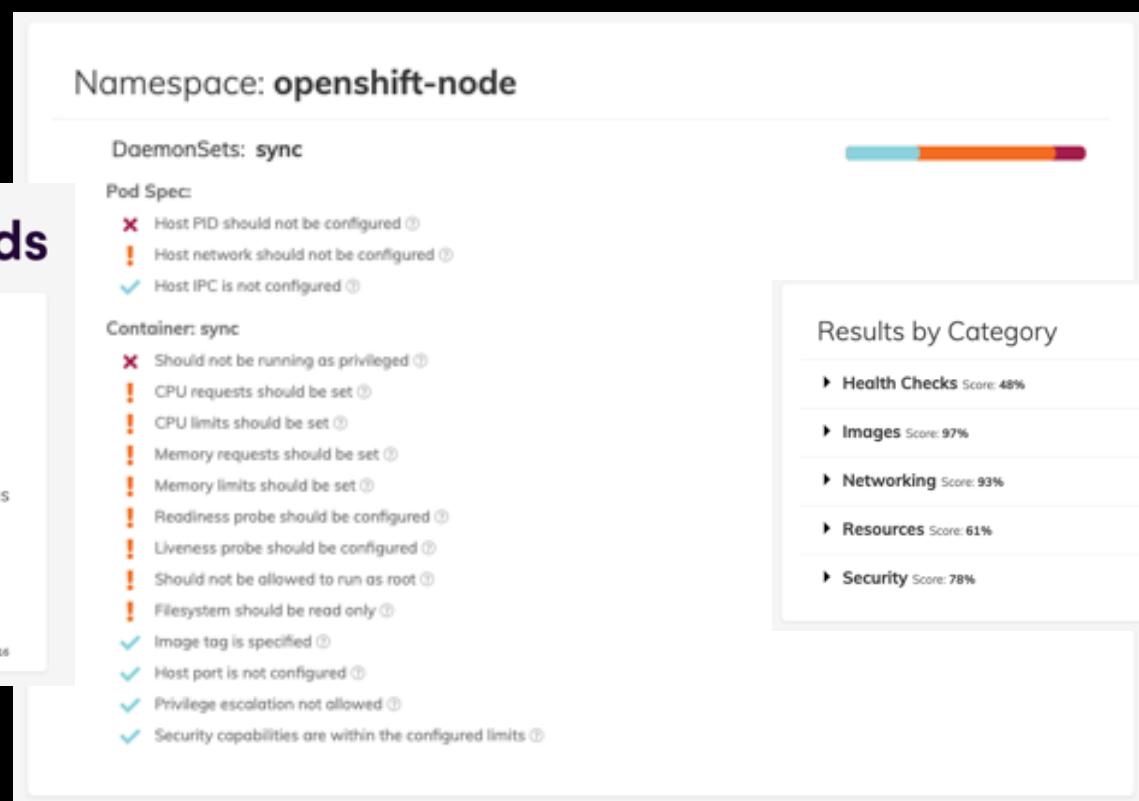
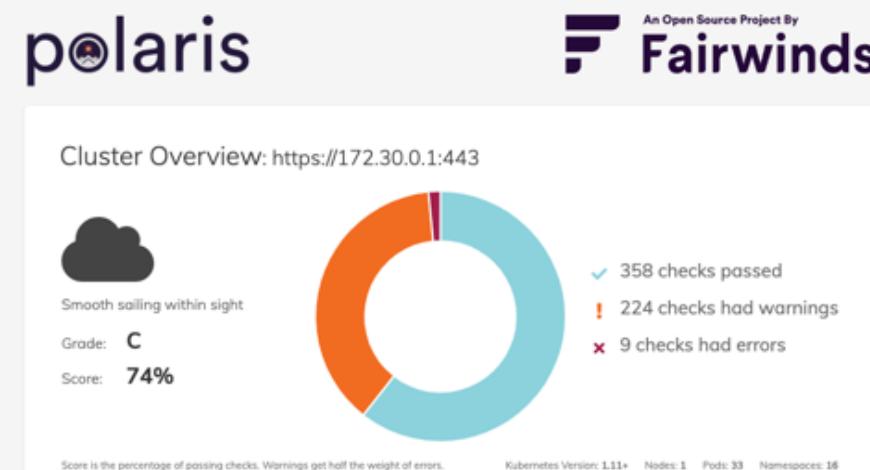
```
42   },
43   {
44     "points": 1
45     },
46     {
47       "selector": "containers[] .securityContext .runAsUser > 10000",
48       "reason": "Run as a high-UID user to avoid conflicts with the host's user table",
49       "points": 1
50     }
51   ]  
}
```

Kubernetes – Cluster Security – Secure Deployments

Scan for K8s Best Practices



Polaris – Validation of best practices



Kubernetes – Cluster Security – Secure Images

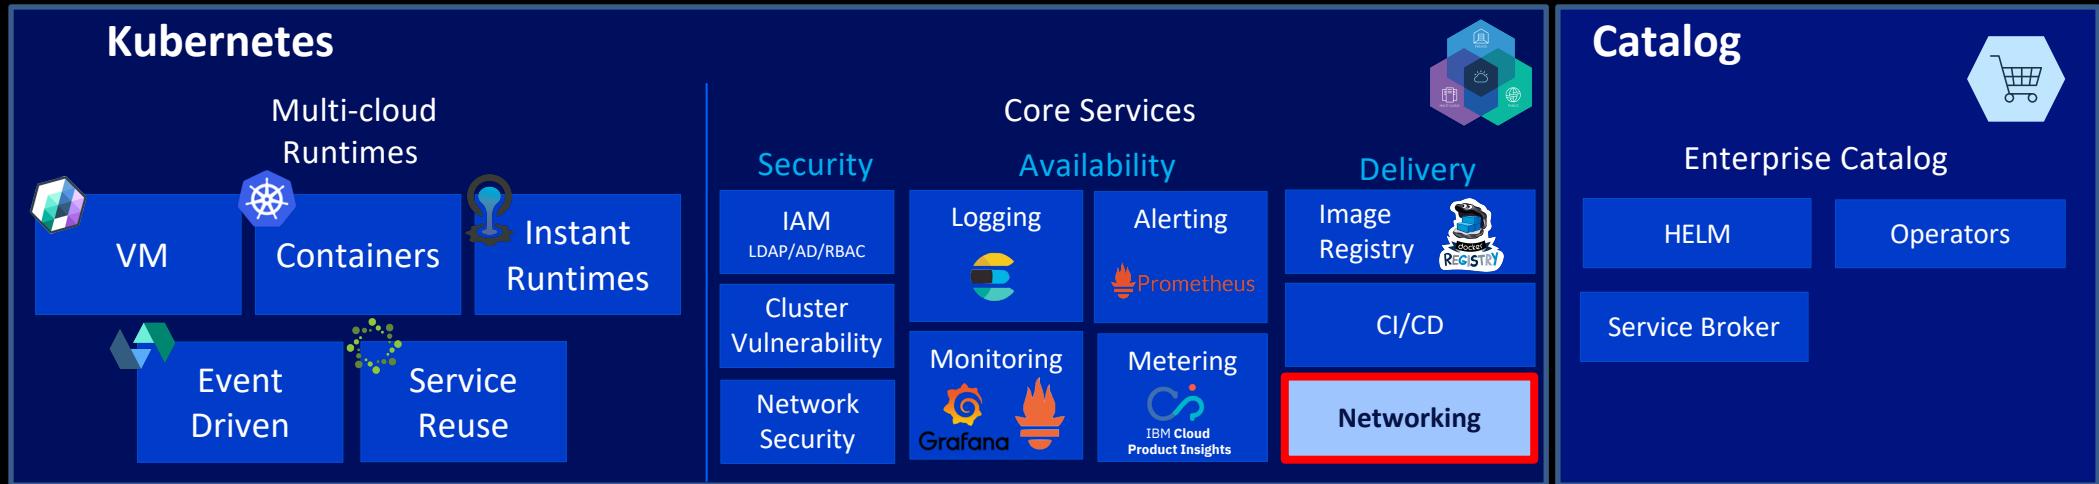
Scan images for vulnerabilities



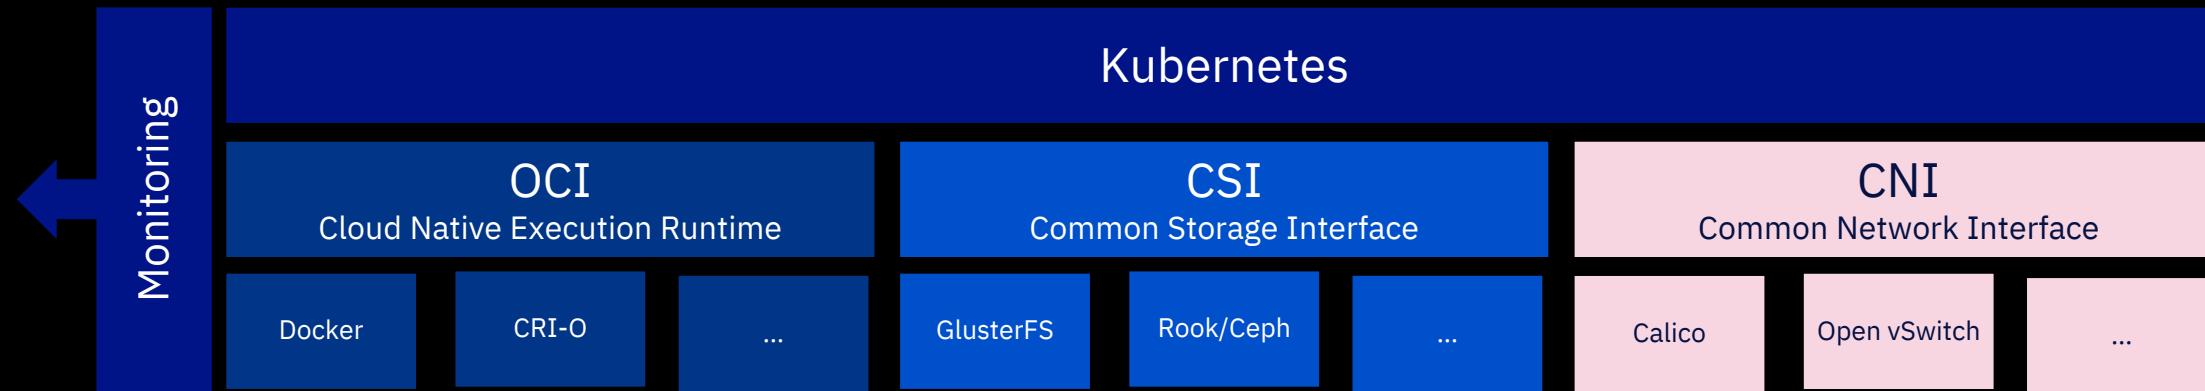
One of the best-known Image scanning tools is Clair. It is an open source project for the static analysis of vulnerabilities in application containers (currently including appc and docker).

| clair timeout: 10ms docker timeout: 10ms no whitelist file Analysising 14 layers GET results from Clair API v1 Found 734 vulnerabilities Unknown: 223 Negligible: 345 Low: 184 Medium: 2 | | | | | | |
|---|---------------|-------------|----------------|---------|--|---|
| Severity | Name | FeatureName | FeatureVersion | FixedBy | Description | Link |
| Medium | CVE-2009-3546 | libwnf | 0.2.6.4-18.6 | | The <code>_gdGetColors</code> function in <code>gd_gdvc</code> in PHP 5.2.11 and 5.3.x before 5.3.1, and the GD Graphics Library 2.x, does not properly verify a certain <code>colorstotal</code> structure member, which might allow remote attackers to conduct buffer overflow or buffer over-read attacks via a crafted GD file, a different vulnerability than CVE-2009-3293. NOTE: some of these details are obtained from third party information. | https://security-tracker.debian.org/tracker/CVE-2009-3546 |
| Medium | CVE-2007-3996 | libwnf | 0.2.6.4-18.6 | | Multiple integer overflows in <code>libgd</code> in PHP before 5.2.4 allow remote attackers to cause a denial of service (application crash) and possibly execute arbitrary code via a large (1) <code>srcW</code> or (2) <code>srcH</code> value to the (a) <code>gdImageCopyResized</code> function, or a large (3) <code>sy</code> (<code>height</code>) or (4) <code>sx</code> (<code>width</code>) value to the (b) <code>gdImageCreate</code> or the (c) <code>gdImageCreateTrueColor</code> function. | https://security-tracker.debian.org/tracker/CVE-2007-3996 |

Kubernetes – Core Services – Mesh Networking



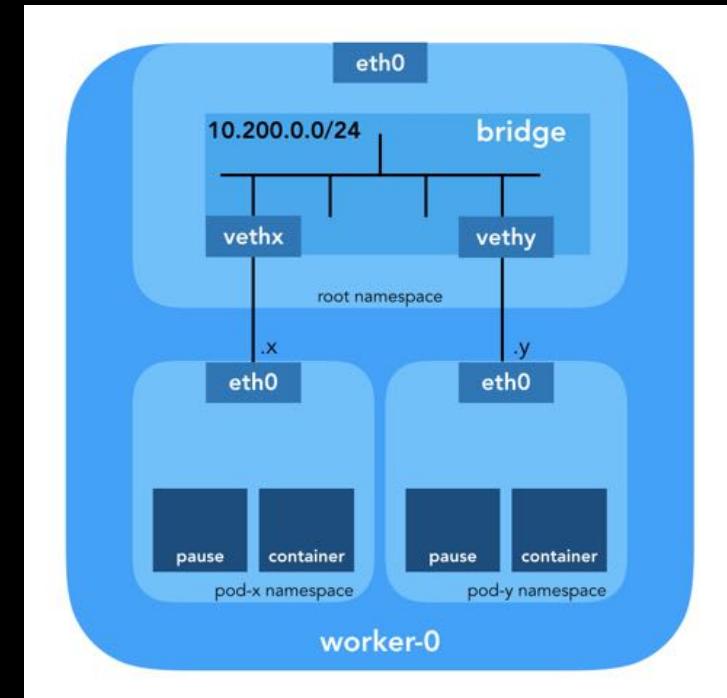
Kubernetes – Core Services – Network Security



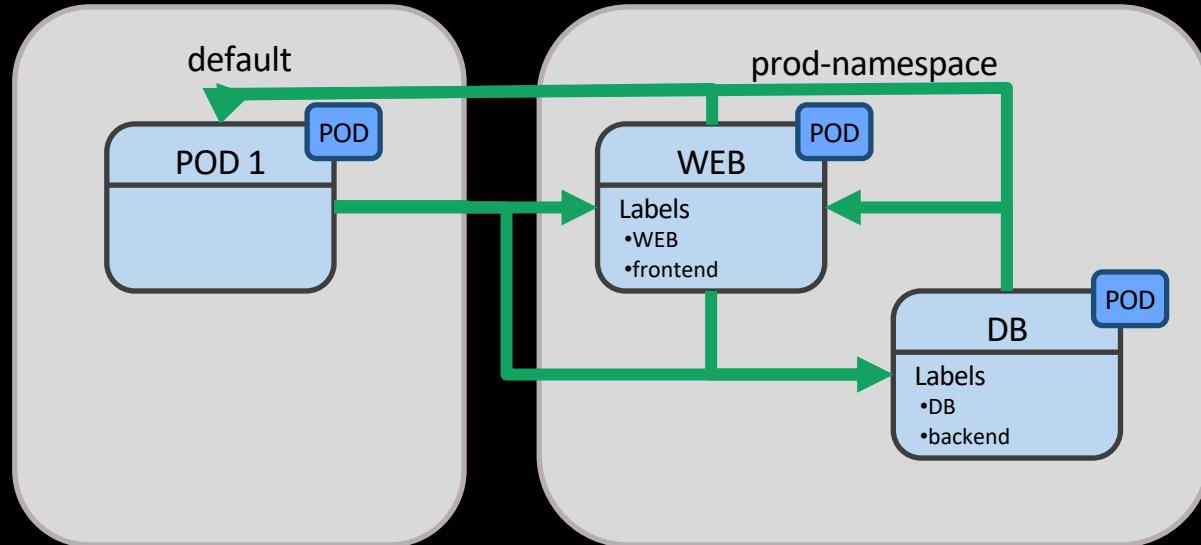
CNI – Common Network Interface

Provides a rich set of security enforcement capabilities running on top of a highly scalable and efficient virtual network

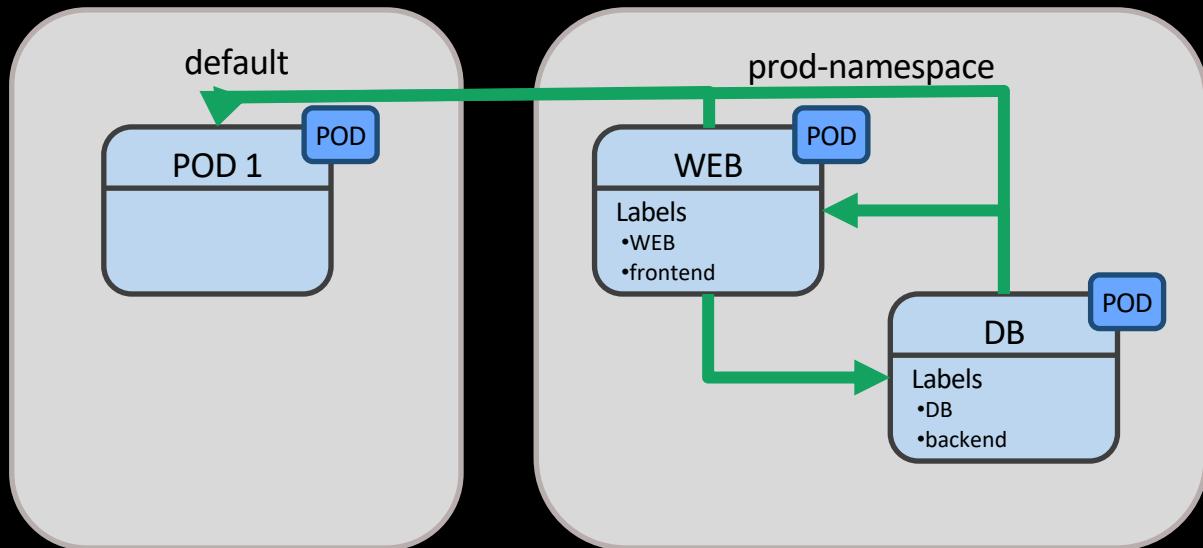
- **Calico**
Virtual networking and network security for containers, VMs, and bare metal services.
- **Open vSwitch**
Production quality, multilayer virtual switch
- ...



Kubernetes – Core Services – Network Security

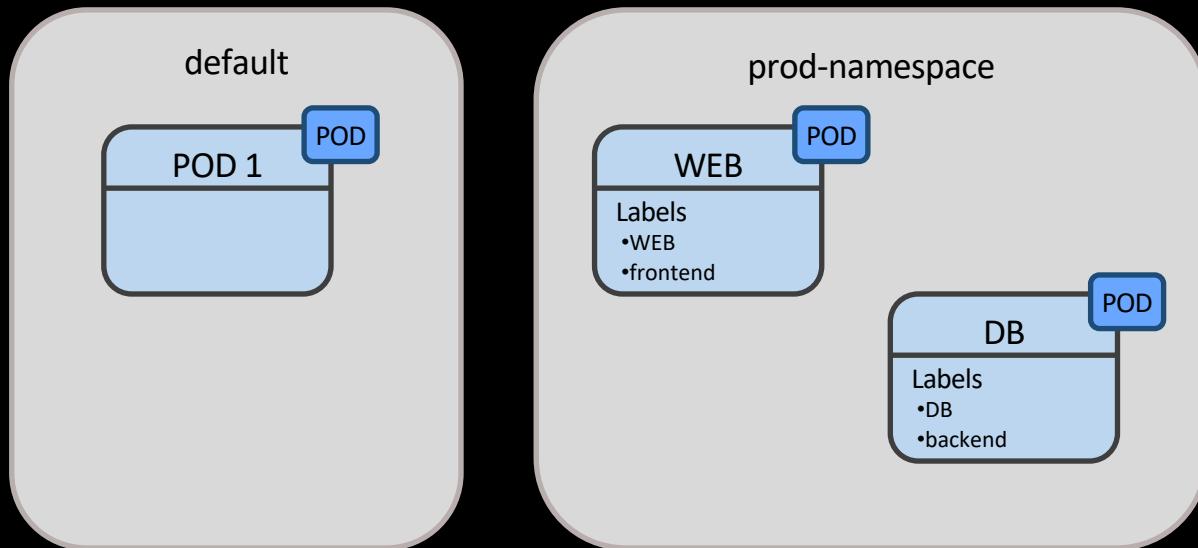


Kubernetes – Core Services – Network Security



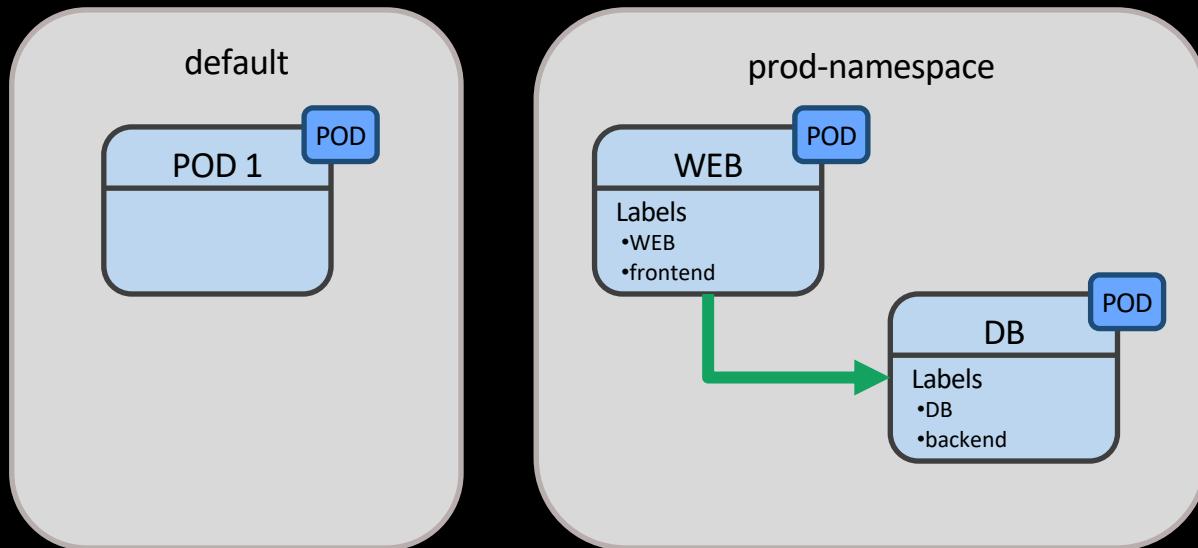
```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: default-deny
  namespace: demo-namespace
spec:
  podSelector:
    matchLabels: {}
```

Kubernetes – Core Services – Network Security



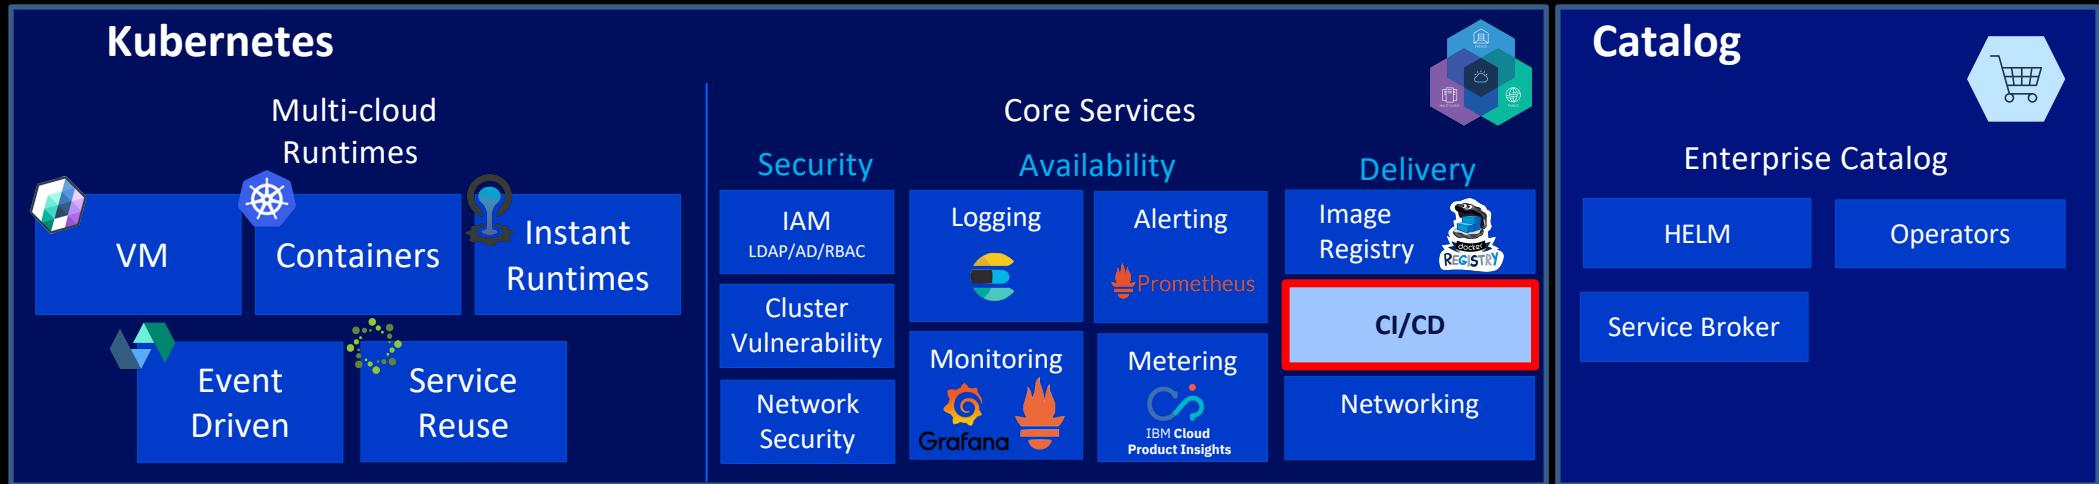
```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: default-deny-all
spec:
  podSelector:
    matchLabels: {}
```

Kubernetes – Core Services – Network Security



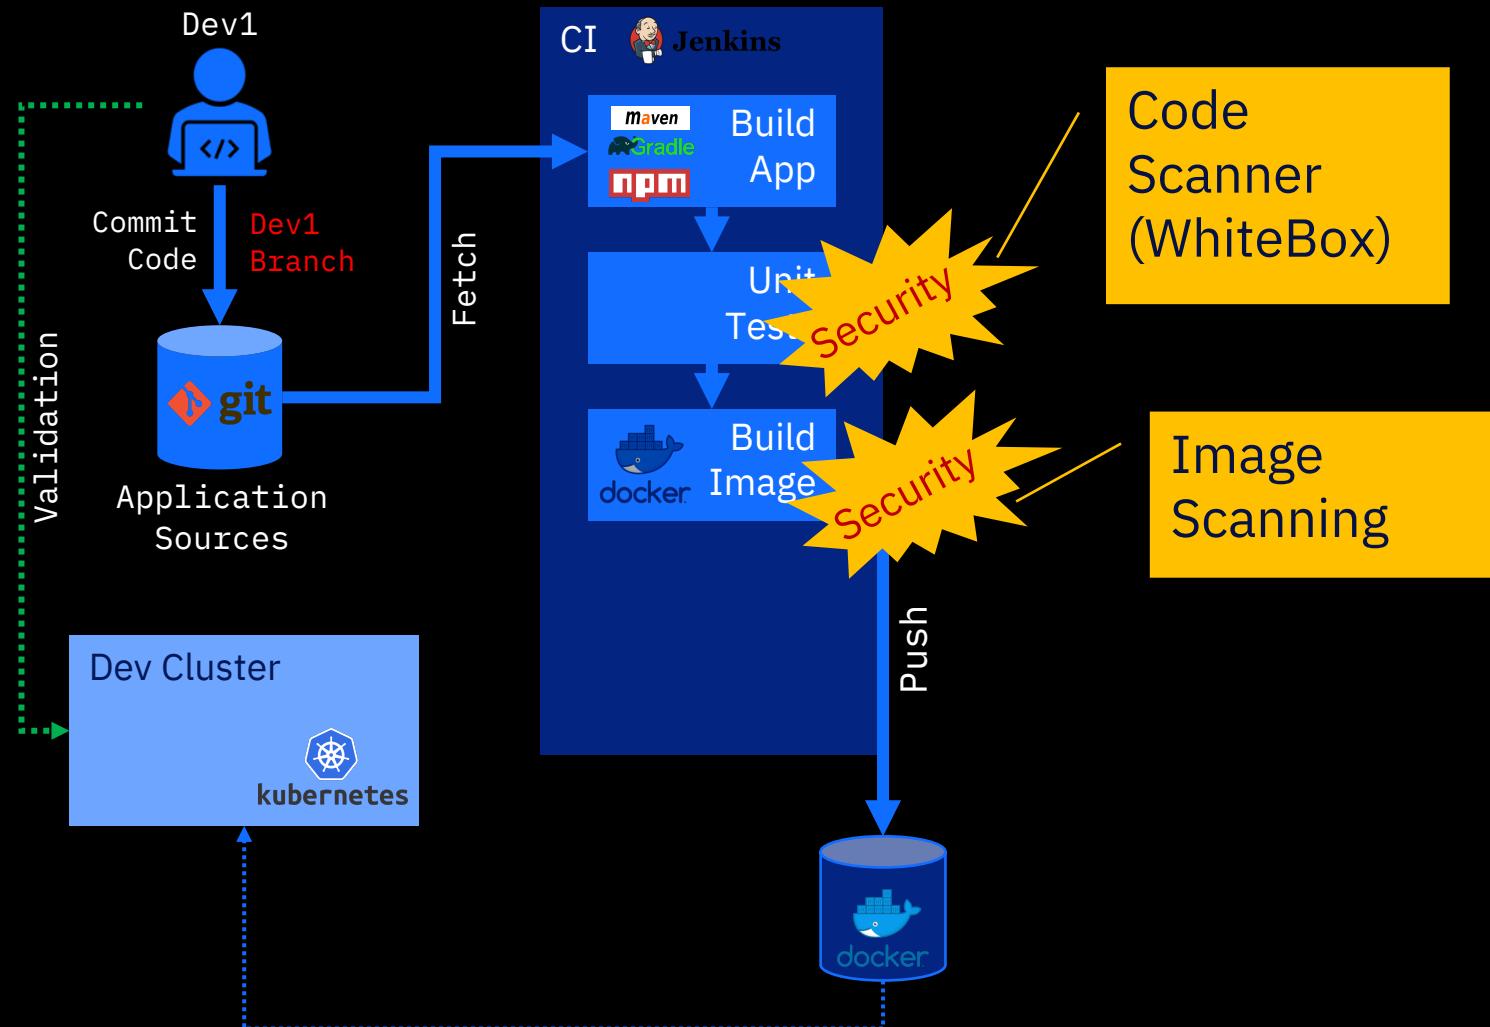
```
kind: NetworkPolicy
metadata:
  name: access-frontend-backend
  namespace: prod-namespace
spec:
  podSelector:
    matchLabels:
      run: DB
  ingress:
  - from:
    - podSelector:
        matchLabels:
          run: WEB
```

Kubernetes – Core Services – Mesh Networking



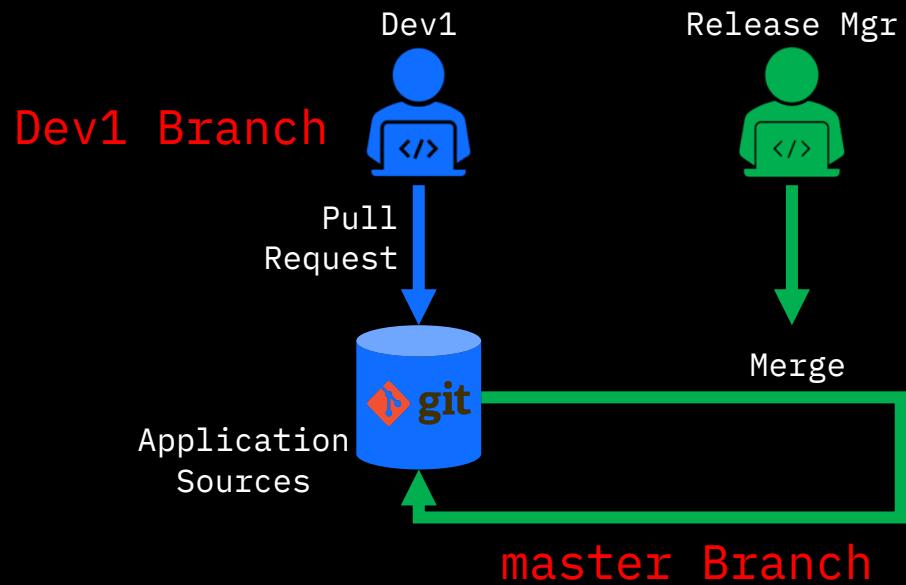
Kubernetes – CI/CD Pipeline

Development – Dev1 Branch



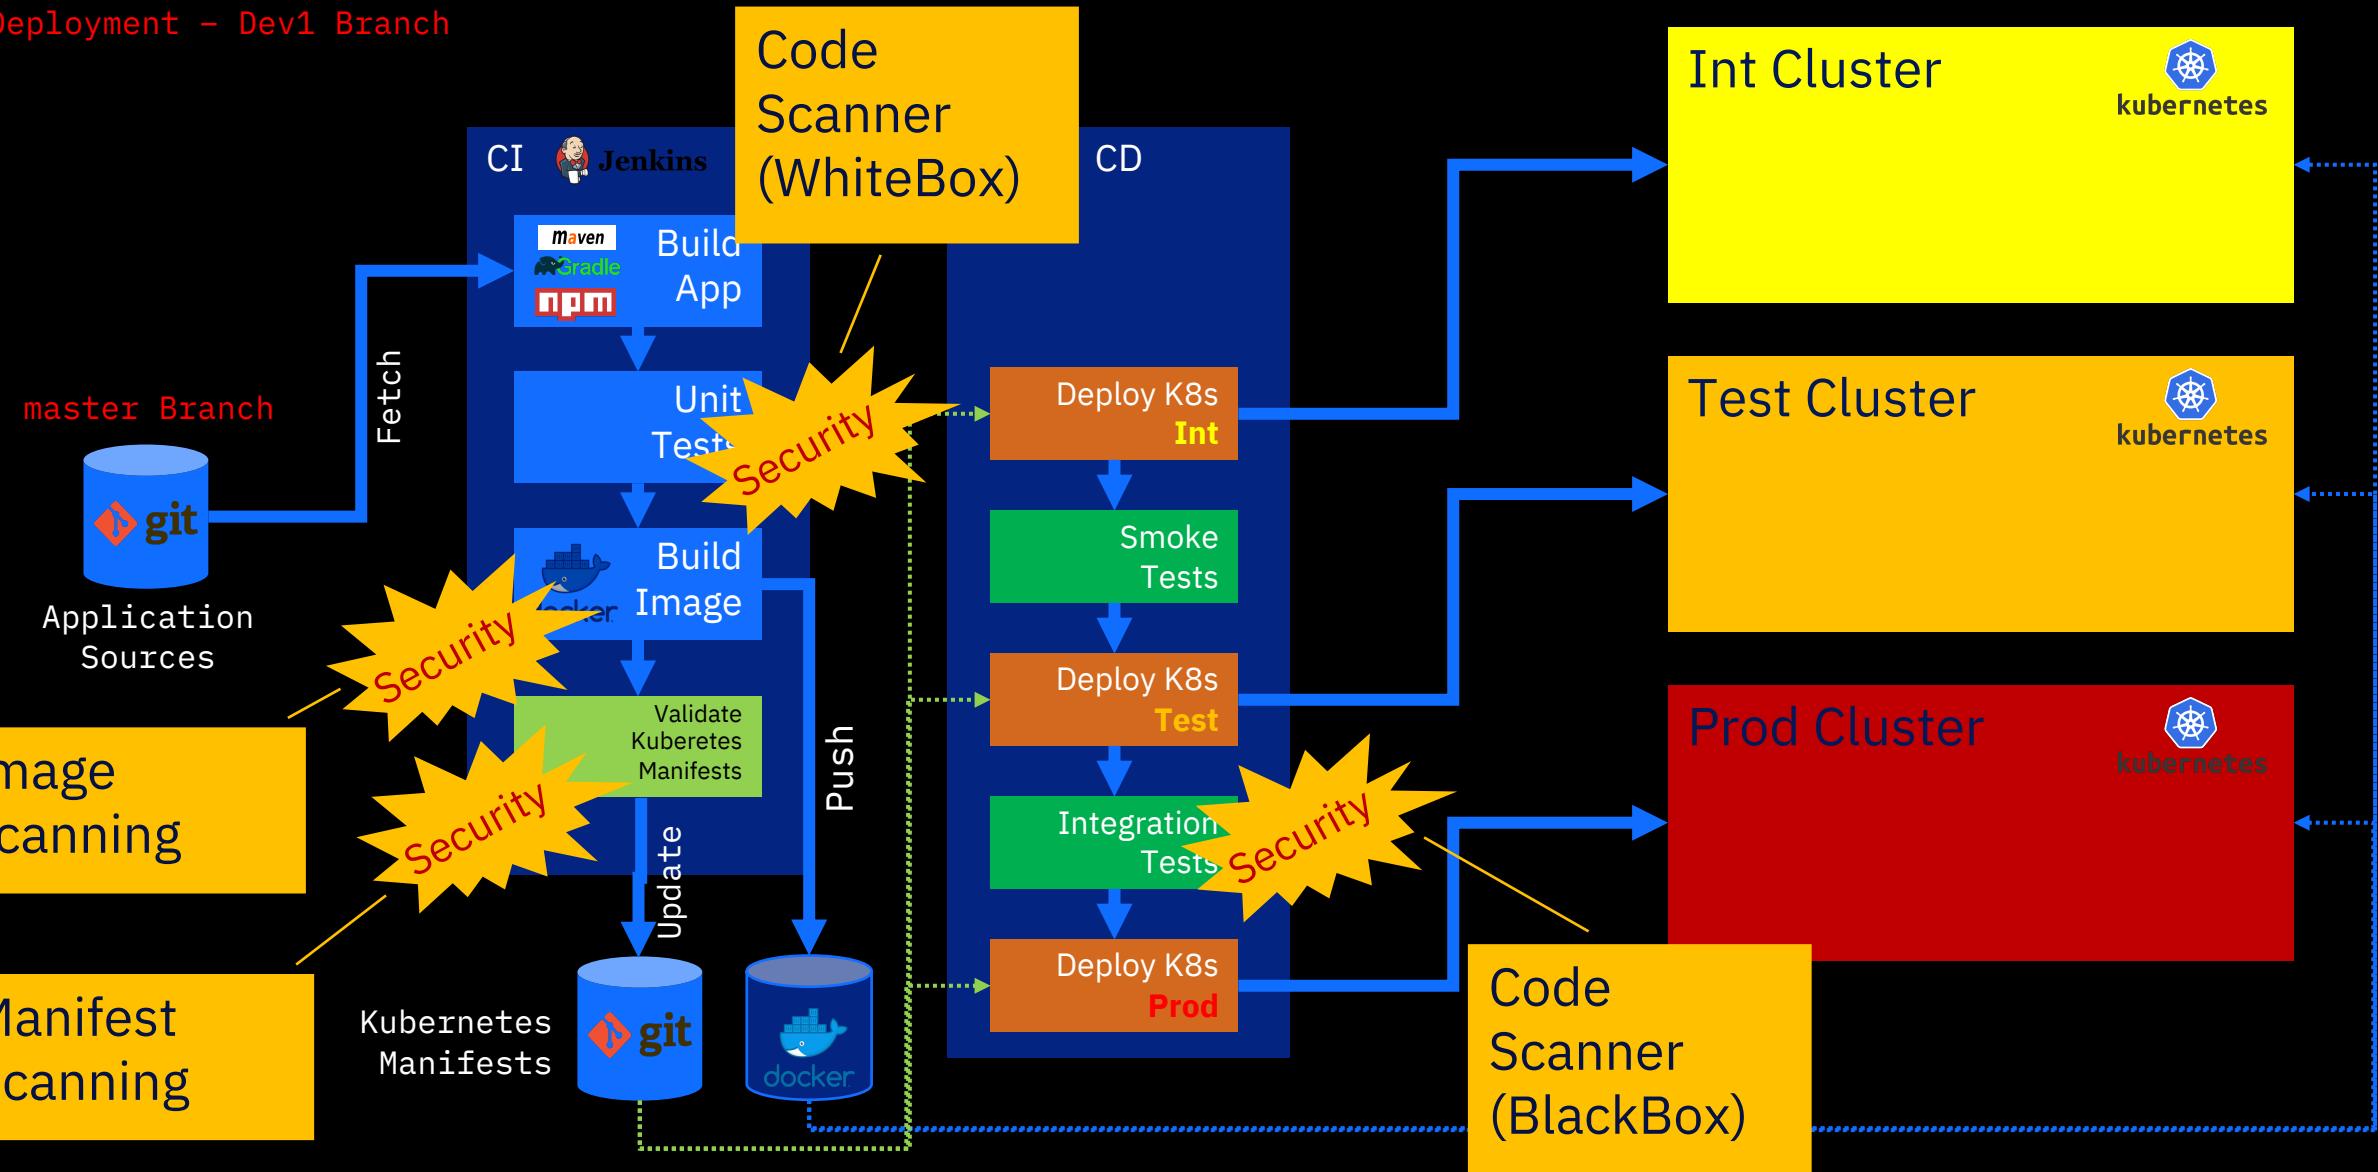
Kubernetes – CI/CD Pipeline

Merge – Dev1 Branch to master Branch

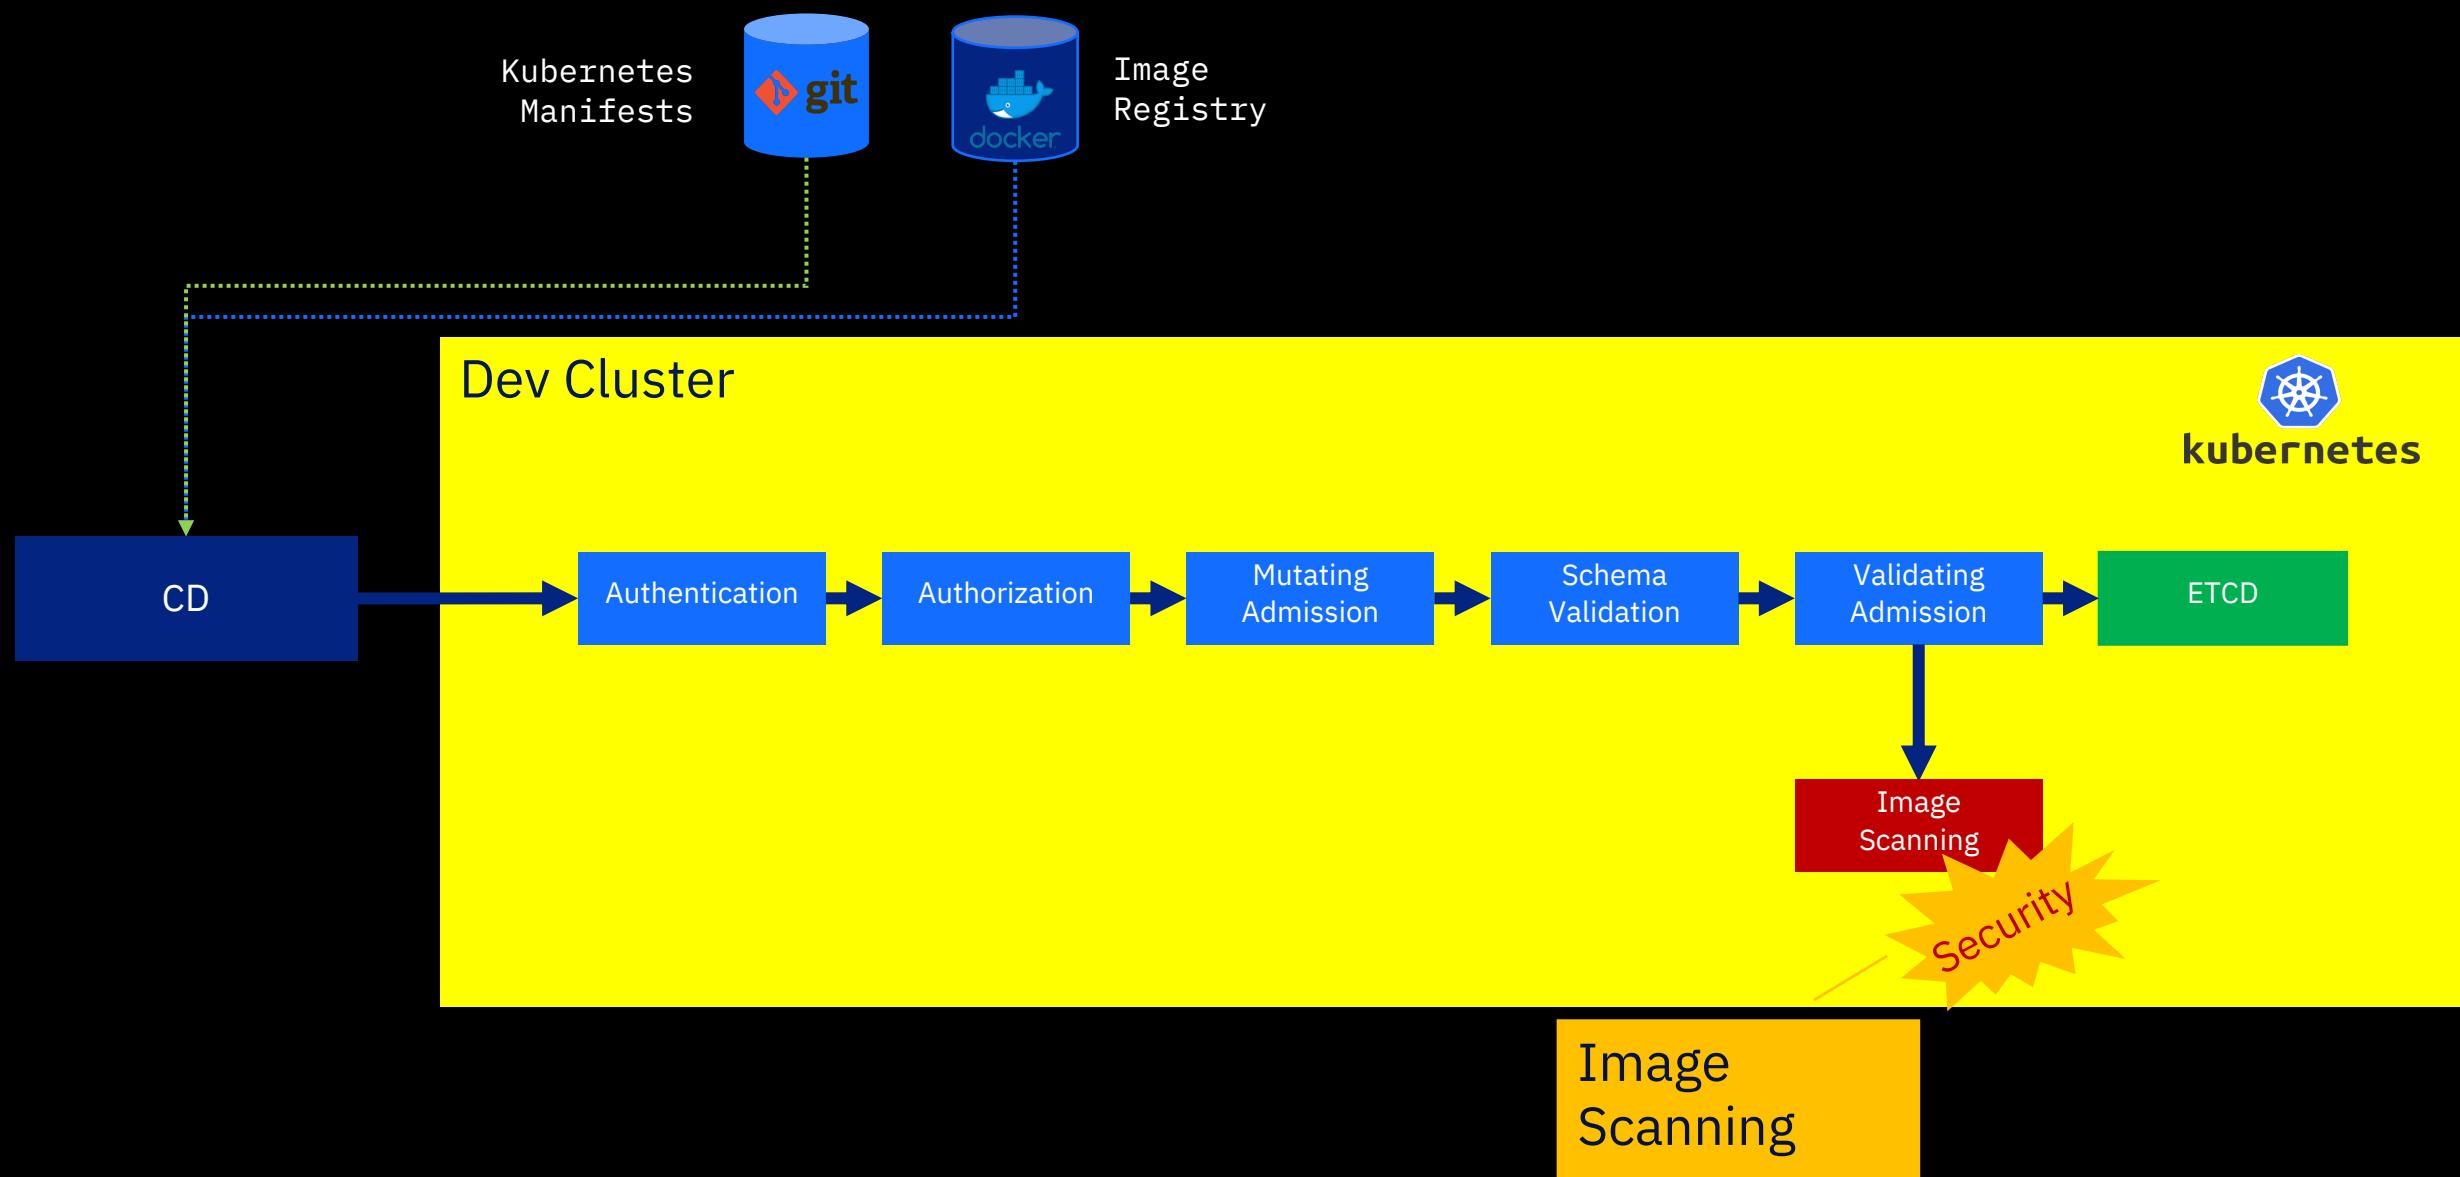


Kubernetes – CI/CD Pipeline

Deployment – Dev1 Branch



Kubernetes – CD Pipeline



DevOps – GitOps

Full Audit Trail

GitOps is using git, which is an excellent database to persist the changes made to the system.

Everything as a code

In the cloud, both CI and CD should be defined as code first. The desired state of environments should also be set in full as a declarative code.

Test Your Manifests Before You Commit

A lot of issues that escape into pre-production environments can be prevented by testing the changes before committing and pushing them.

Two Repos: One For App Source Code, Another For Manifests

Clean separation of application code vs. application config.

Separation of access

The developers who are developing the application, may not necessarily be the same people who can/should push to production environments, either intentionally or unintentionally.

Kubernetes Security – Some tools I use

RBAC

RBAC-lookup - <https://github.com/FairwindsOps/rbac-lookup>

rakkess - <https://github.com/corneliusweig/rakkess>

rbac-view – <https://github.com/jasonrichardsmith/rbac-view>

rback - <https://github.com/team-soteria/rback>

Scanning

Polaris - <https://github.com/FairwindsOps/polaris>

Clair - <https://github.com/coreos/clair>

Management

K9s - <https://github.com/derailed/k9s>

Kubernetes Security – Some Reading Tips

<https://kubernetespodcast.com/episode/065-attacking-and-defending-kubernetes/>

https://en.wikipedia.org/wiki/Red_team

<https://blog.ropnop.com/attacking-default-installs-of-helm-on-kubernetes/>

<https://kubernetes.io/docs/tasks/administer-cluster/encrypt-data/>

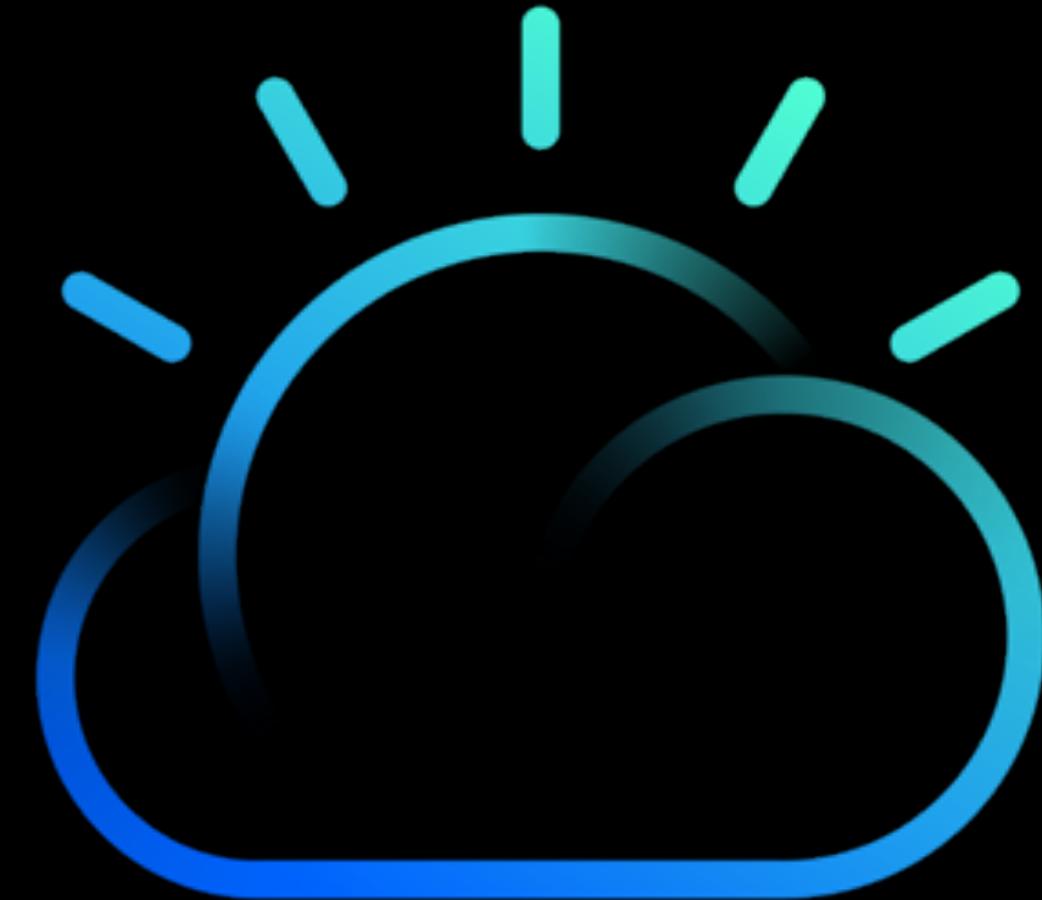
<https://kubernetes.io/blog/2019/03/21/a-guide-to-kubernetes-admission-controllers/>

<https://kubernetes.io/docs/tutorials/clusters/apparmor/>

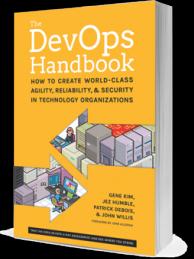
<https://kubernetes.io/docs/concepts/policy/pod-security-policy/>

<https://kubernetes.io/docs/concepts/storage/volumes/#hostpath>

QUESTIONS?

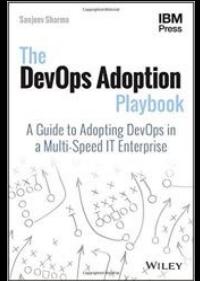


DevOps - – Some Reading Tips



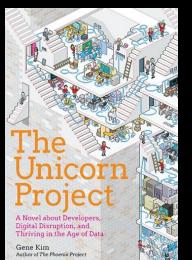
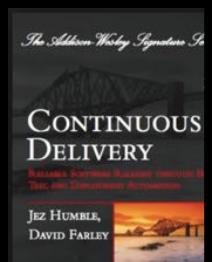
How To Create World-Class
Agility, Reliability, &
Security in Technology
Organizations

Achieve business goals,
maximize IT productivity, and
deliver innovation with
enterprise-scale DevOps.



A view into the cultural
challenges of adopting DevOps
and best practices

Automate deployments using
production-like environments and
accelerate delivery cycles



A Novel about Developers, Digital
Disruption, and Thriving in the Age
of Data

Kubernetes – Some Reading Tips

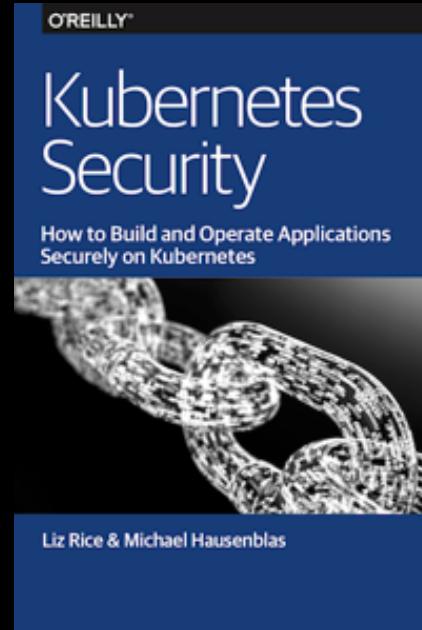


We wrote the books on
cloud adoption...

ibm.com/cloud/garage/adoption

...and Kubernetes in
the enterprise

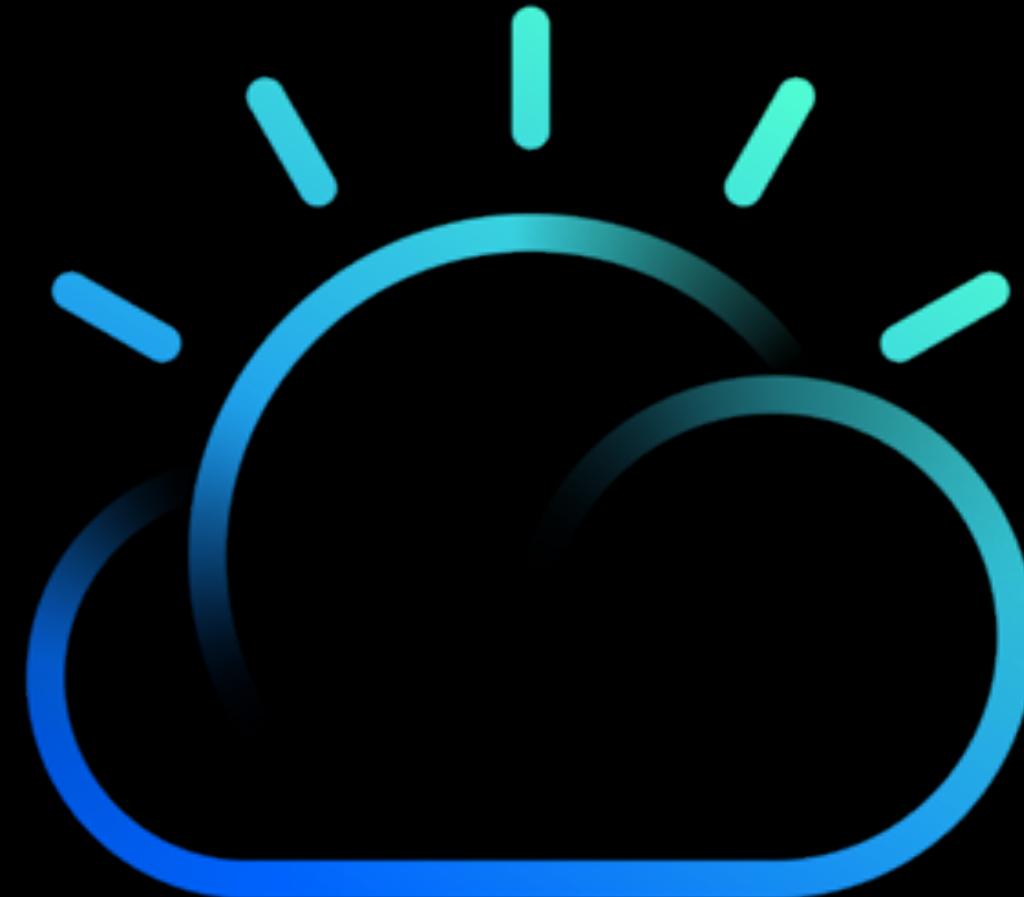
<https://ibm.co/2LQketN>



<https://kubernetes-security.info/>



°° The Journey to Cloud
Docker - Hands-On



IBM Cloud

Remember your Team Color

black 31701

olive 31711

peru 31715

white 31702

brown 31712

chocolate 31716

red 31703

lightblue 31713

orchid 31717

blue 31704

orange 31708

gold 31718

yellow 31705

purple 31709

pink 31719

lime 31706

maroon 31710

violet 31720

cyan 31707

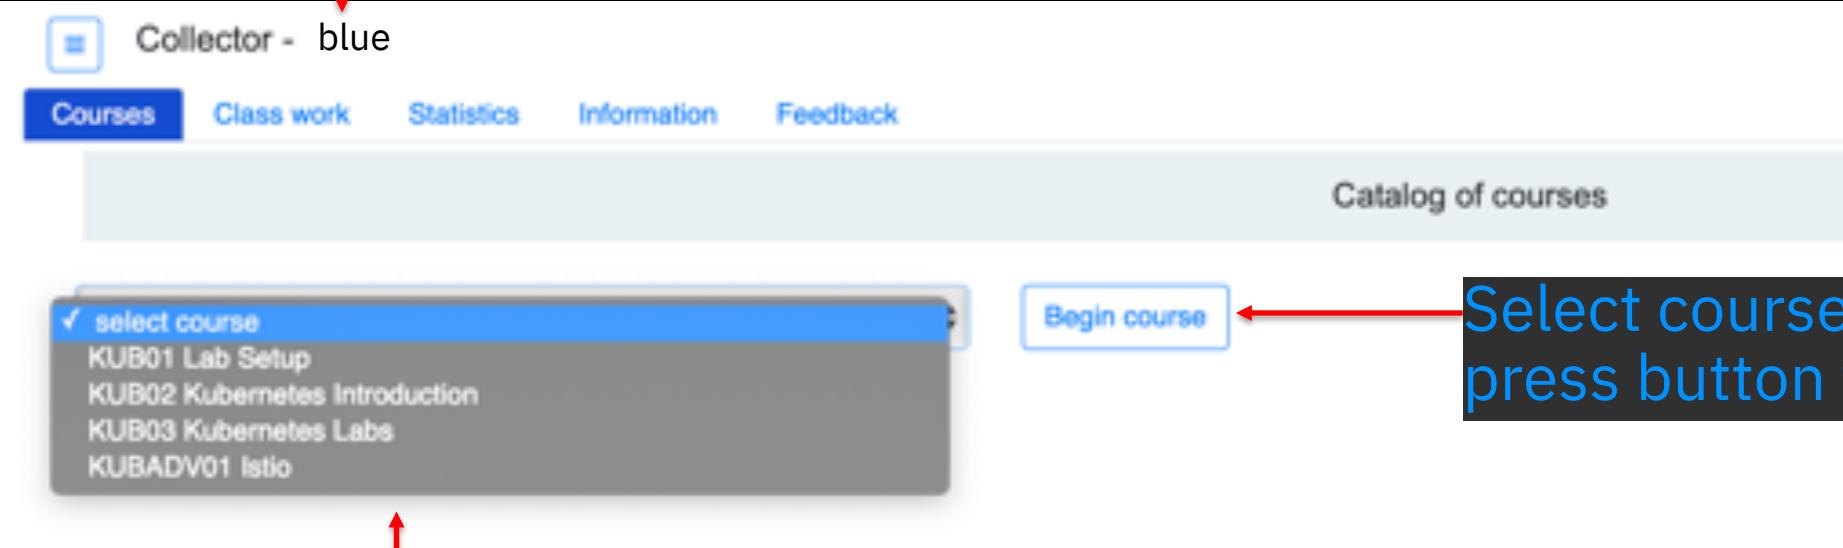
firebrick 31714

Collector - Accessing team web site

`http://158.177.137.195:{port#}`

Team name / color will be shown

blue **31704**



The screenshot shows a web browser displaying the 'Collector - blue' website. At the top, there is a navigation bar with tabs: Courses (which is selected), Class work, Statistics, Information, and Feedback. Below the navigation bar, the page title is 'Catalog of courses'. A dropdown menu titled 'select course' contains four items: KUB01 Lab Setup, KUB02 Kubernetes Introduction, KUB03 Kubernetes Labs, and KUBADV01 Istio. To the right of the dropdown is a blue button labeled 'Begin course'. A large callout box with a red border and white text points to the 'Begin course' button, containing the instruction 'Select course and press button to begin'. A red arrow also points from the text 'Current course catalog' at the bottom left to the 'select course' dropdown.

Current course catalog

Select course and press button to begin



JTC01 – Docker Labs

Lab 0 : Docker basics

Lab 1 : Build a Docker image

Lab 2 : Run a Docker image

Lab 3 : Use the Portainer tool

Lab 4 : Deploy a more complex Docker application

Lab 5 : Push a Docker image into a registry

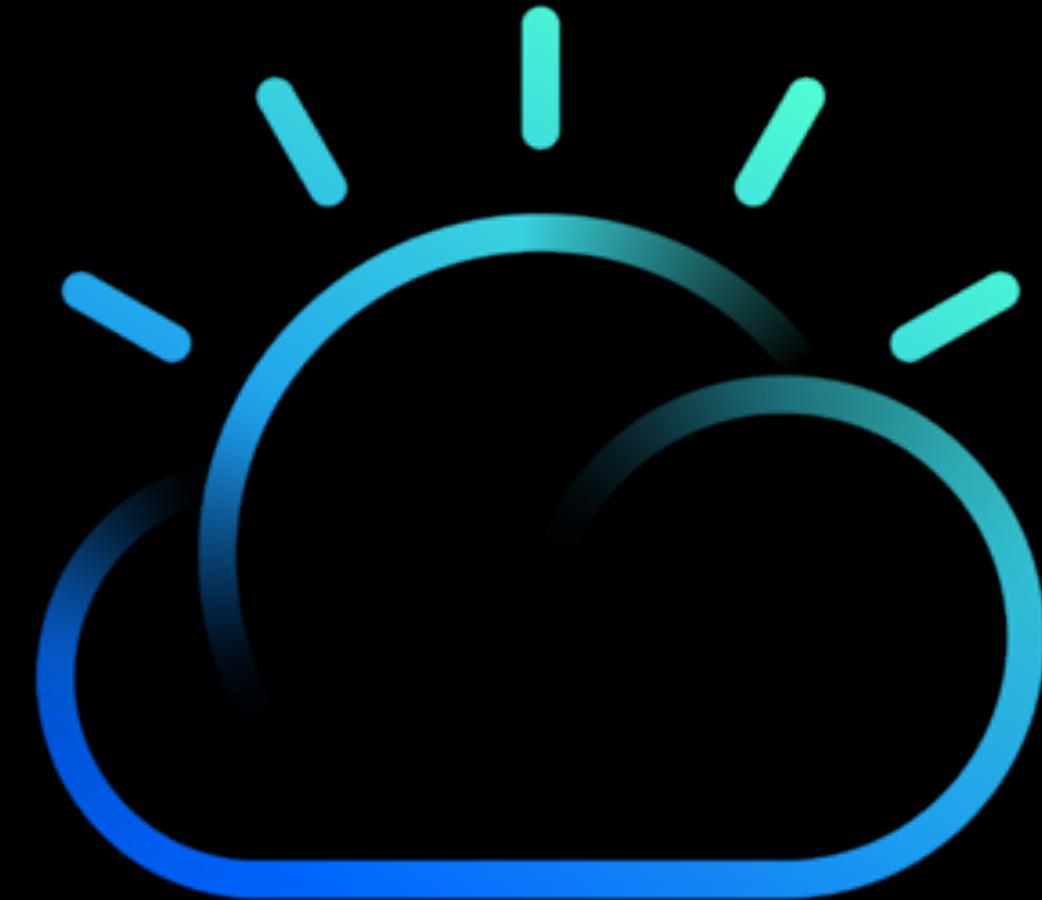


READY
SET
GO!!!!

<http://158.177.137.195:{port#}>

Duration: 60 mins

QUESTIONS?





°° The Journey to Cloud
Kubernetes - Hands-On

08



IBM Cloud



JTC02 – Kubernetes Labs

Lab 1: Refresher/overview over Kubernetes.

Lab 2: Running your first Pod on Kubernetes.

Lab 3: Deploy a Web Application on Kubernetes

Lab 4: Scale an application on Kubernetes

Lab 5: Persisting data in Kubernetes with Volumes



READY
SET
GO!!!!

<http://158.177.137.195:{port#}>

Duration: 120 mins

QUESTIONS?





°° The Journey to Cloud
Kubernetes Security -
Hands-On





JTC16 – Kubernetes Labs

~~Lab 1: Network Policies~~

~~Lab 2: Role Based Access Control (RBAC)~~

~~Lab 3: Service Accounts~~

~~Lab 4: Security Tooling~~

Lab 5: Image scanning



READY
SET
GO!!!!

<http://158.177.137.195:{port#}>

Duration: 30 mins

QUESTIONS?





Documentation

https://github.com/niklaushirt/k8s_training_public

Sources

<https://github.com/niklaushirt/training>

Always-on training environment

<http://158.177.137.195:31999/>

IBM