

## CAS Cloud and Platform Management (CAS CPM 1)

**Niklaus Hirt**  
Cloud and DevOps Architect

[nikh@ch.ibm.com](mailto:nikh@ch.ibm.com)

Rotkreuz, 06.12.2019

# Who am I?

## Niklaus Hirt

Passionate about tech for over 35 years

- High-school in Berne
- Degree in Computer Science at EPFL
- ELCA
- CAST
- IBM

✉ nikh@ch.ibm.com

🐦 @nhirt



# Agenda - Kubernetes Basic Concepts

## Modules

Module 0: Prepare the Labs

Module 1: Microservices

Module 2: Containers with Docker

Module 3: Kubernetes

Module 4: Let's get real

BREAK

Module 5: Docker Hands-On

LUNCH – 12:30

# Agenda - Kubernetes Advanced Concepts

## Modules

Module 6: Kubernetes Hands-On

BREAK

Module 7: Kubernetes Security

oder

Module 8: Mesh Networking with ISTIO

END of course – ca 16:15



Sources and documentation available here:

[https://github.com/niklaushirt/k8s\\_training\\_public](https://github.com/niklaushirt/k8s_training_public)

<https://github.com/niklaushirt/training>



°° The Journey to Cloud  
**Prepare the Labs**



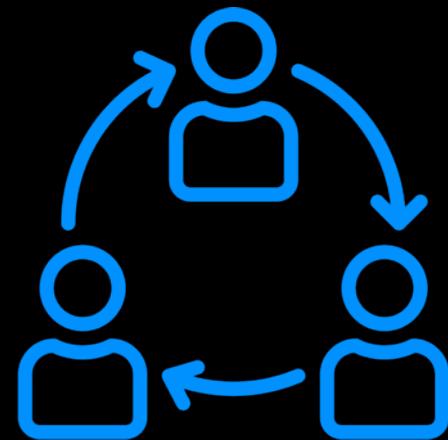
IBM Cloud

# Session Objectives

Attendees will be grouped in **teams** wherever it makes sense to facilitate collaborative work.



Following some lectures will be **hands-on** work that each team collaborates to complete.



# Teams

**black 31701**

**olive 31711**

**peru 31715**

**white 31702**

**brown 31712**

**chocolate 31716**

**red 31703**

**lightblue 31713**

**orchid 31717**

**blue 31704**

**orange 31708**

**gold 31718**

**yellow 31705**

**purple 31709**

**pink 31719**

**lime 31706**

**maroon 31710**

**violet 31720**

**cyan 31707**

**firebrick 31714**



# Collector - Accessing team web site

http://158.177.137.195:{port#}

Team name / color will be shown

blue 31704

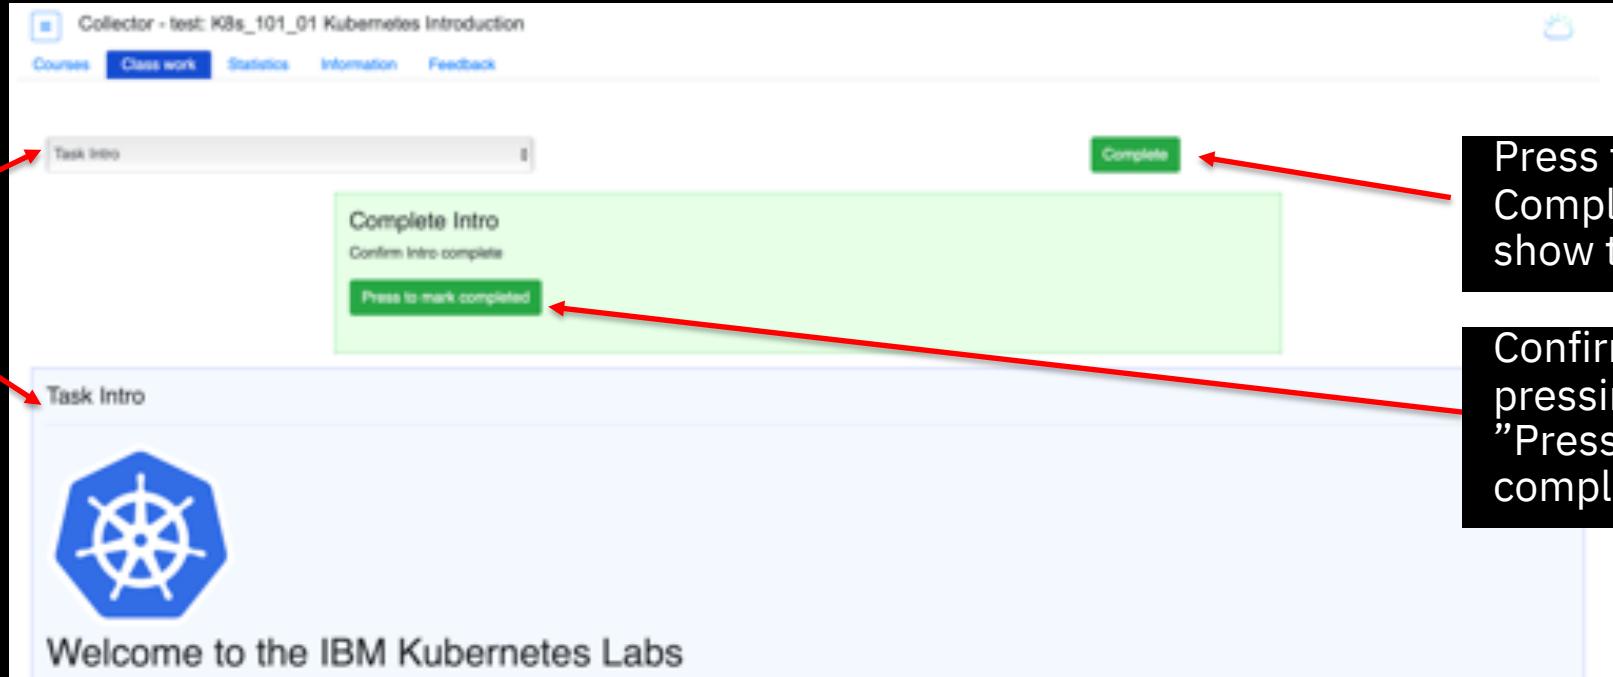
The screenshot shows a web browser window with the title "Collector - blue". The navigation bar includes links for "Courses", "Class work", "Statistics", "Information", and "Feedback". Below the navigation bar, a section titled "Catalog of courses" displays a list of courses under the heading "select course". The listed courses are: KUB01 Lab Setup, KUB02 Kubernetes Introduction, KUB03 Kubernetes Labs, and KUBADV01 Istio. To the right of the course list is a button labeled "Begin course". A red arrow points from the text "Select course and press button to begin" to the "Begin course" button.

- KUB01 Lab Setup
- KUB02 Kubernetes Introduction
- KUB03 Kubernetes Labs
- KUBADV01 Istio

Current course catalog

# Collector – Class work

Select class work and the blue portion of the screen is shown



Press the green Complete button to show the green portion.

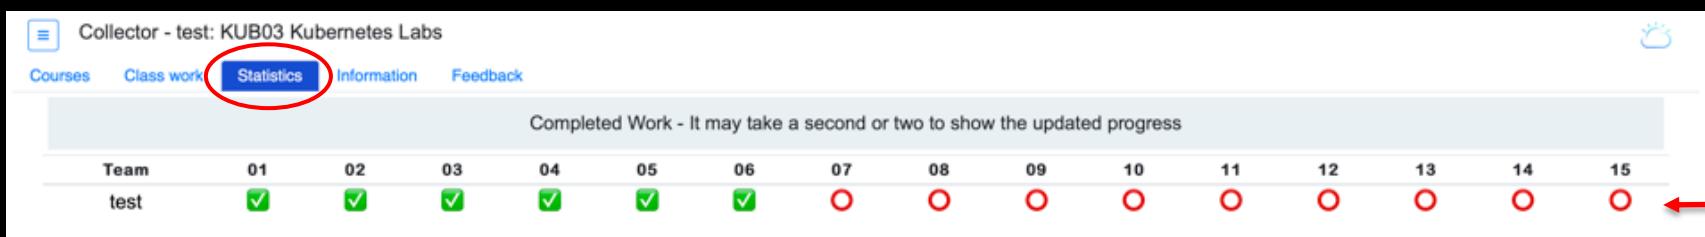
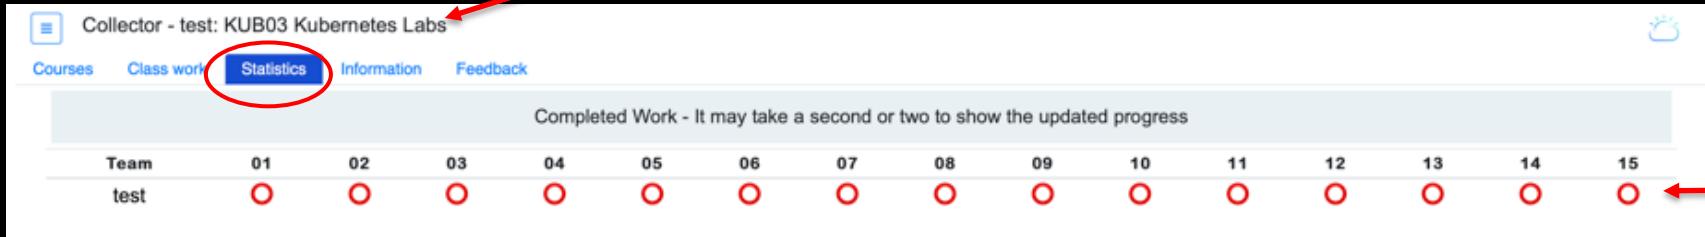
Confirm completion by pressing the green "Press to mark completed" button.



The Complete Button might not show instantly depending on the course settings

# Collector – Track course completed work

Course title



**Green checkmark** - item is completed

**Red circle** - item is waiting to be completed

The number of items tracked will change based on the current course selected.

# Collector – Instructor Dashboard

Remaining Time for the Lab

Collector - instructor: K8s\_101\_01 Kubernetes Introduction

Remaining time: 0h 29m 50s

Courses Class work Statistics Information Feedback Insight

Completed Work - It may take a second or two to show the updated progress

Team	01	02	03	04	05	06
instructor	✓	○	○	○	○	0
lightblue	✓	✓	✓	✓	✓	1
olive	✓	✓	○	○	○	2
peru	○	○	○	○	○	3
chocolate	○	○	○	○	○	4
pink	○	○	○	○	○	5
violet	○	○	○	○	○	6

Collector - instructor: K8s\_101\_01 Kubernetes Introduction





## JTC90 Lab Setup

Task 1: Download Training VM

Task 2: Setup VirtualBox >6.14

Task 3: Start Training VM

Task 4: Login / Check



JTC90 Lab Setup

# EVERYBODY

Task 2: Setup VirtualBox >6.14

# OK

Task 3: Start Training VM

Task 4: Login ? Check

# ?

# The Journey to Cloud **Microservices**

01



IBM Cloud

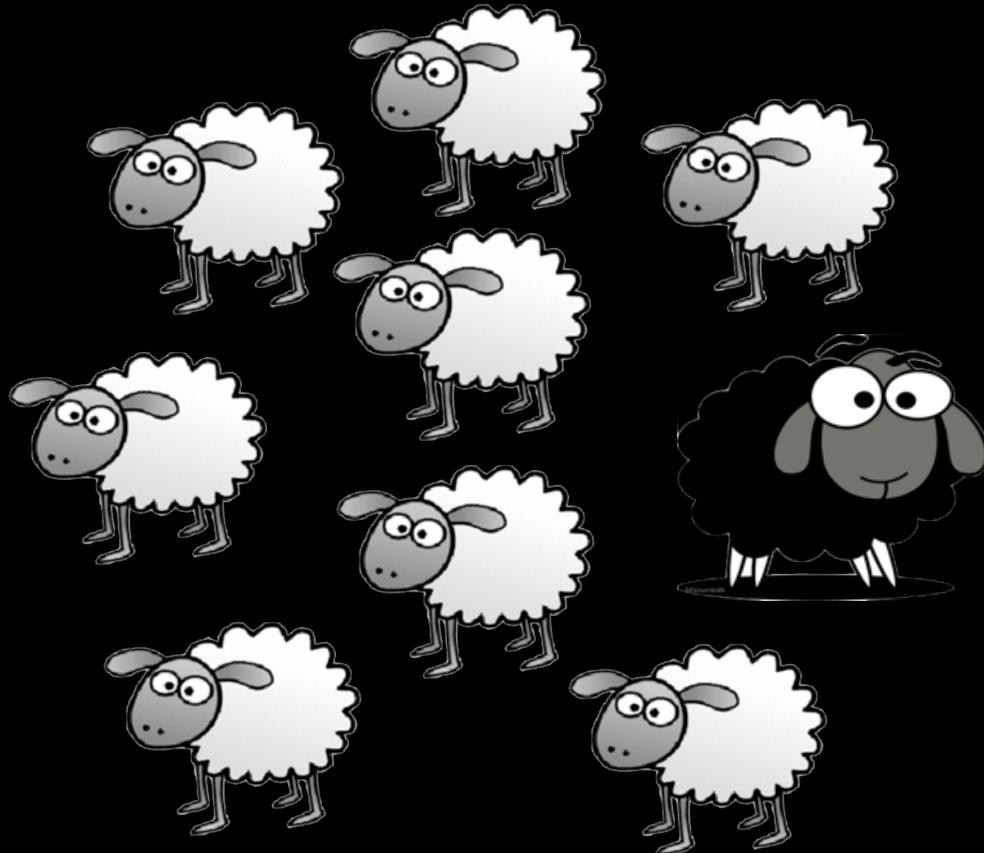


# disruption

dɪs'ruptʃn/  
*noun*

Business. a **radical change** in an industry, business strategy, etc., especially involving the introduction of a **new product or service** that creates a **new market**

# Disruption is new reality



## disruption

dɪs'ruptʃn/

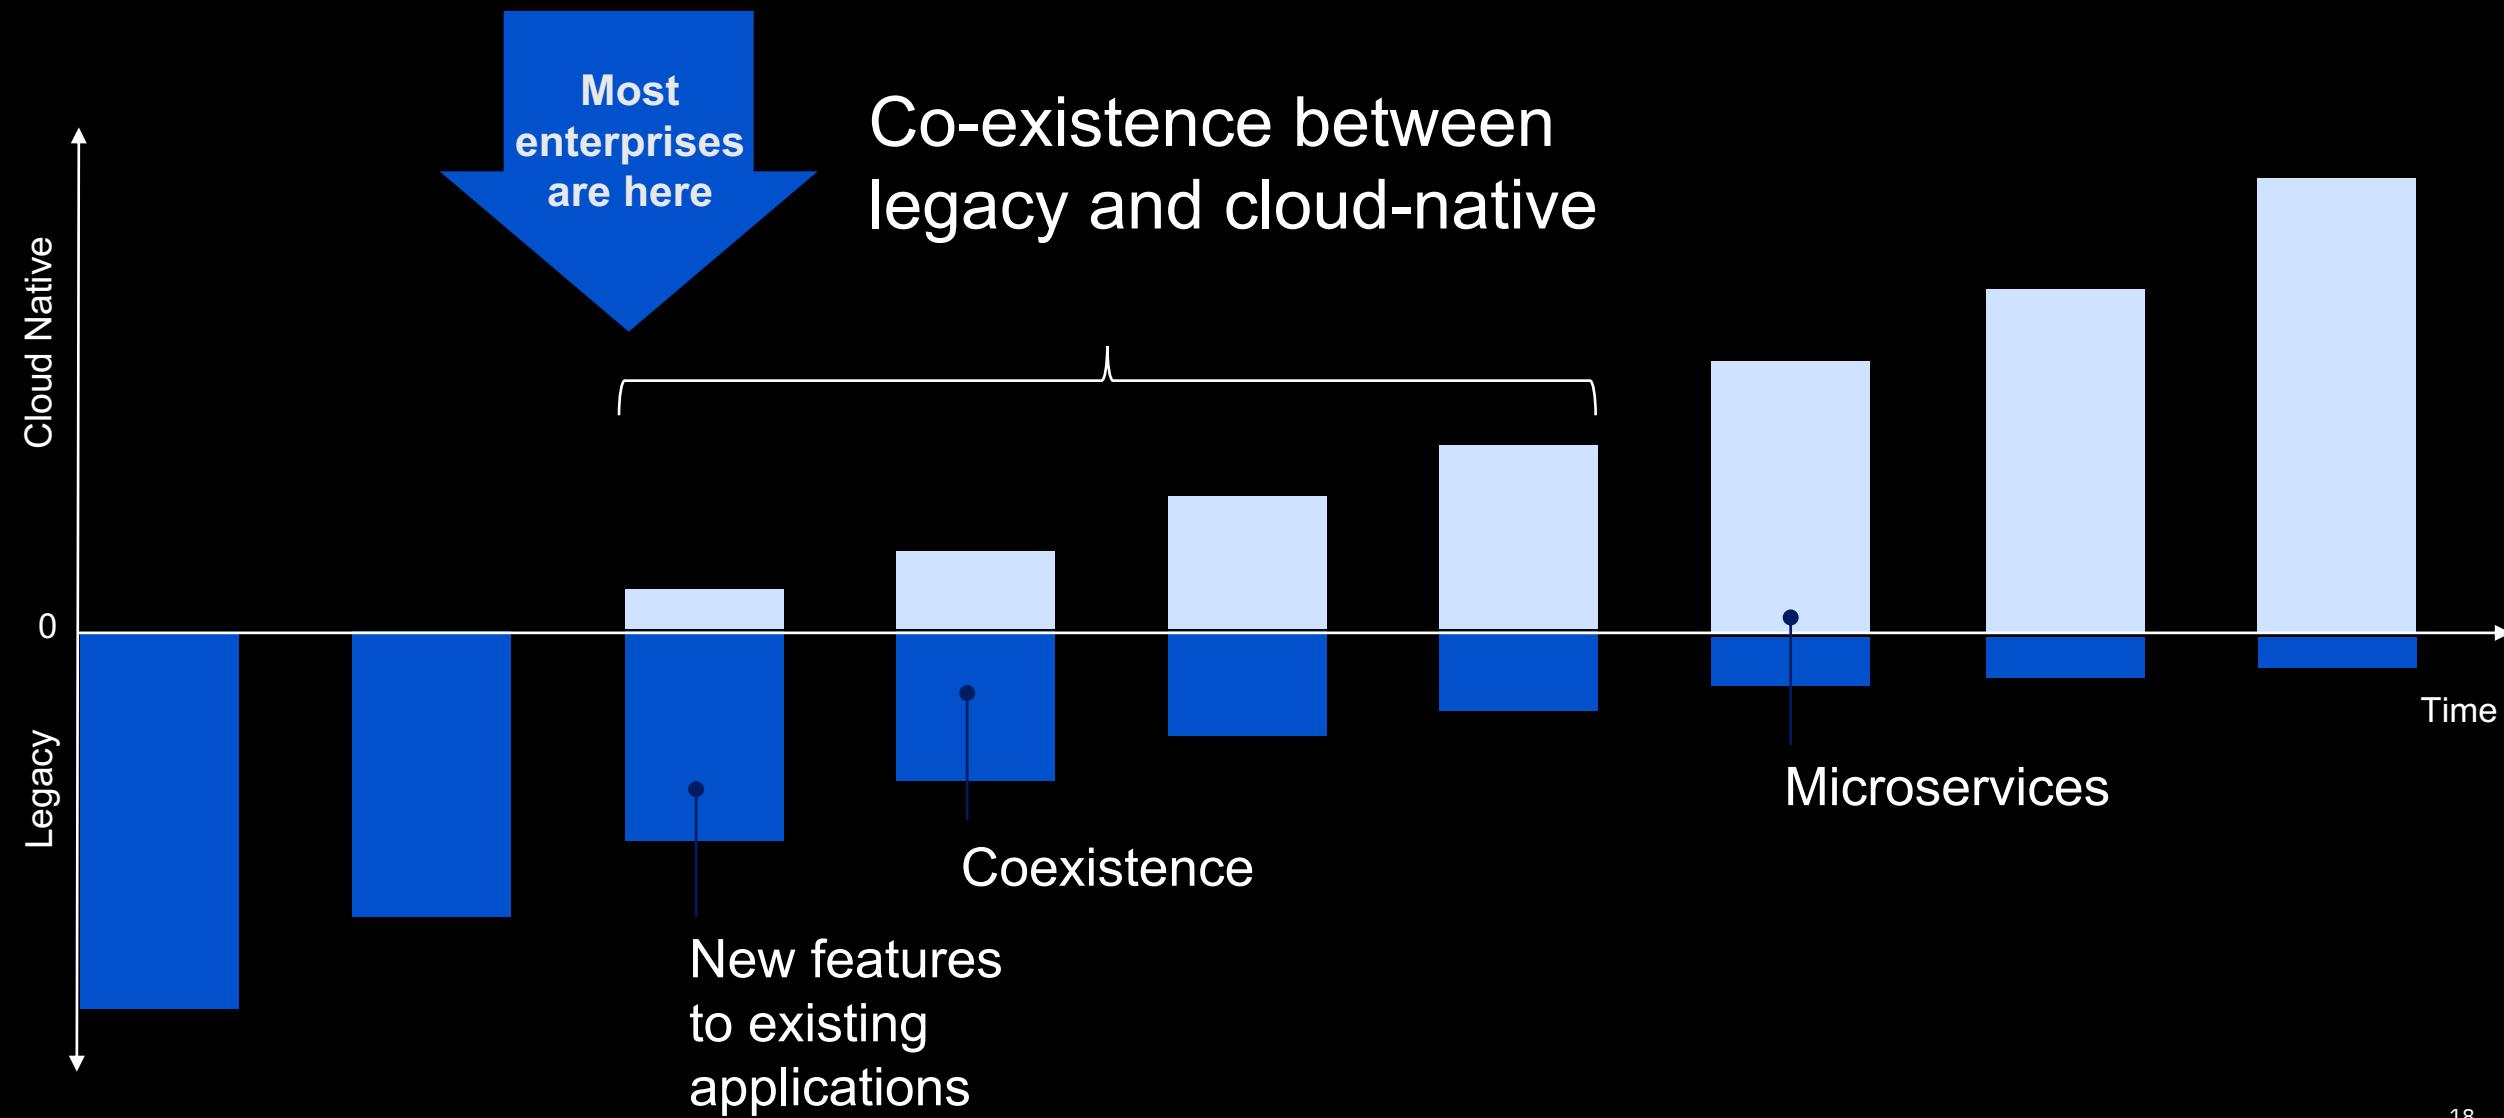
*noun*

Business. a **radical change** in an industry, business strategy, etc., especially involving the introduction of a **new product or service** that creates a **new market**

# Digital Transformation is the new normal

# Digital Transformation - Way to go

*Cloud native and legacy apps will co-exist for the next 10+ years*



# How do you achieve digital transformation?

Adopt new <b>Processes</b>	Adopt new <b>Technology</b>	Adopt new <b>Tools</b>	Adopt <b>Cloud</b>
Continuous Integration	Architectural patterns (Microservices)	Git, Github, Gitlab	Docker, Kubernetes
Continuous Build	Frameworks (Java MicroProfile, Spring ...)	Jenkins	Virtualization, VMs
Continuous Deploy	Node.js , Swift	UrbanCode	Monitoring, Dashboards, Alerts
Agile Dev Models		Docker	
		Kubernetes	

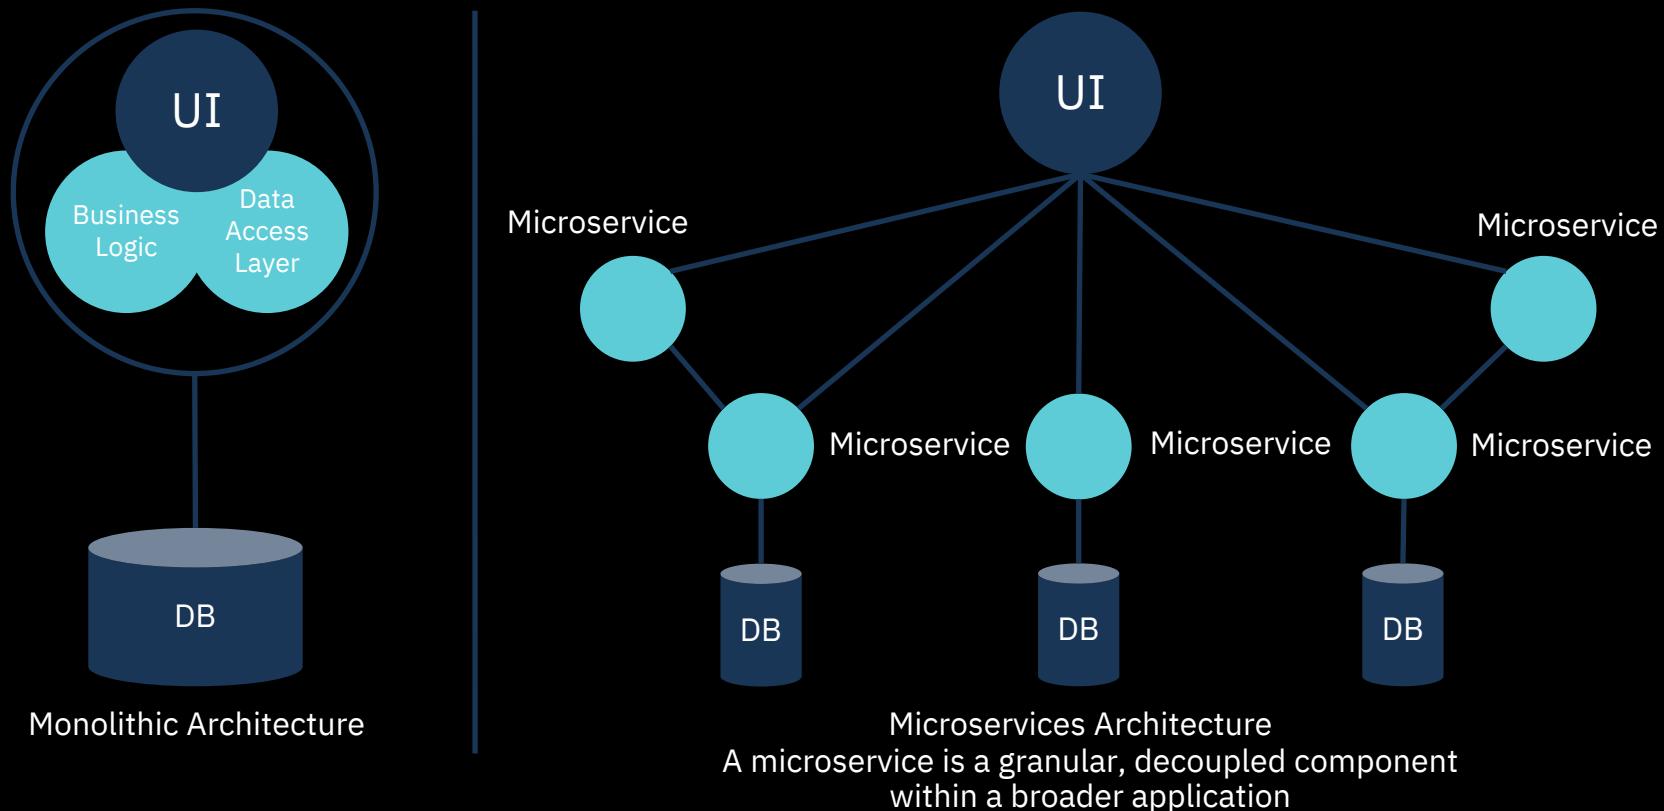
# Microservices

Decomposing an application into **single function modules** which are **independently deployed and operated**

**Accelerate delivery** by minimizing communication and coordination between people

# Microservices architecture

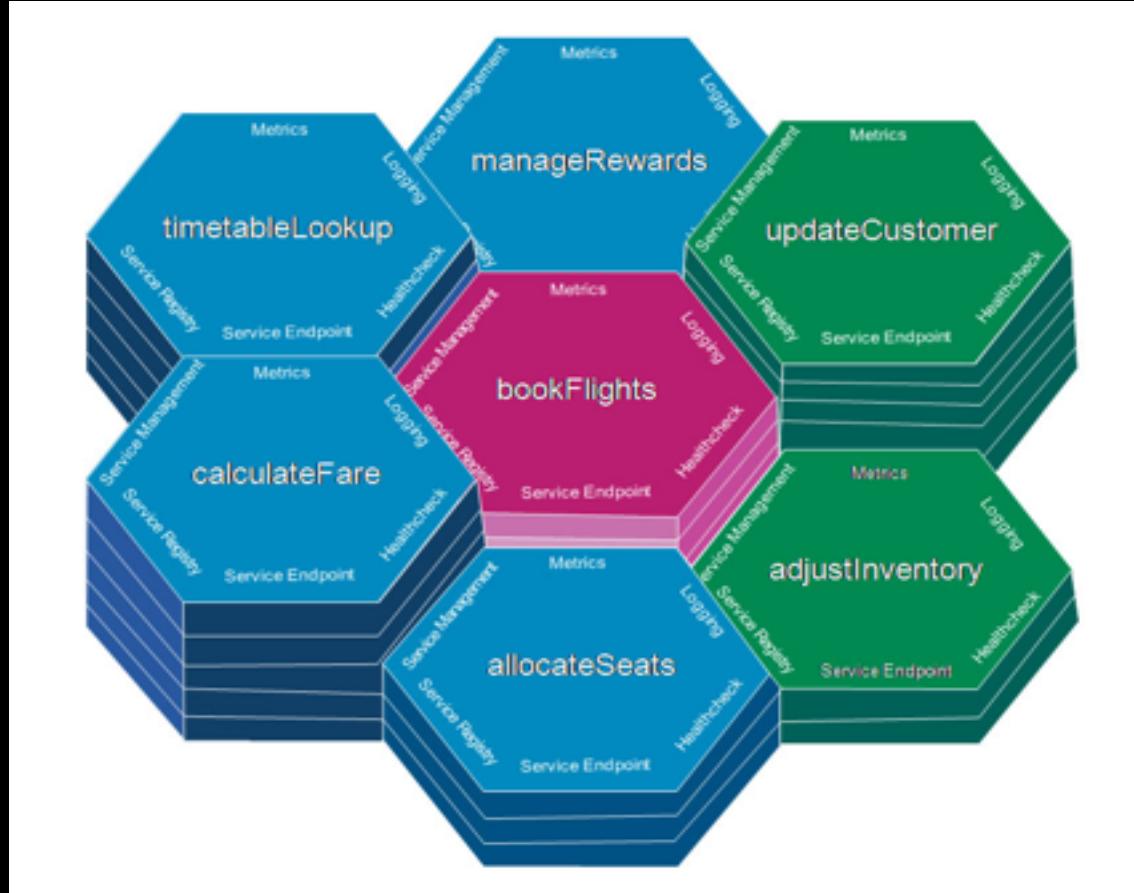
Simplistically, microservices architecture is about breaking down large silo applications into more manageable, fully decoupled pieces



# Example application using microservices

## Airline reservation application

- Book flights
- Timetable lookup
- Calculate fare
- Allocate seats
- Manage rewards
- Update customer
- Adjust inventory



# Microservices – key tenets



Large monoliths are **broken down into many small services**

- Each service runs its own process
- There is one service per container



Services are **optimized for a single function**

- There is only one business function per service
- The Single-responsibility Principle: A microservice should have one, and only one, reason to change



Communication via **REST API** and **message brokers**

- Avoid tight coupling introduced by communication through a database



**Per-service continuous delivery** (CI/CD)

- Services evolve at different rates
- Let the system evolve, but set architectural principles to guide that evolution



**Per-service high availability** and clustering decisions

- One size or scaling policy is not appropriate for all
- Not all services need to scale; others require autoscaling up to large numbers

# Cloud-Native Application Goals

## **Horizontal scaling**

- Application runs in multiple runtimes spread across multiple hosts (VIII)

## **Immutable deployment**

- A runtime is not patched, it's replaced (IX)
- A runtime is stateless (VI)
- Shared functionality in backing services ( IV)

## **Elasticity**

- Automatic scale-out and scale-in to maintain performance
- Achieved via containerization

## **Pay-as-you-go charging model**

- Pay for what you use

# The 12 Factor App

1. Codebase - One codebase tracked in revision control, many deploys
2. Dependencies - Explicitly declare and isolate dependencies
3. Config - Store config in the environment
4. Backing services - Treat backing services as attached resources
5. Build, release, run - Strictly separate build and run stages
6. Processes - Execute the app as one or more stateless processes
7. Port binding - Export services via port binding
8. Concurrency - Scale out via the process model
9. Disposability - Maximize robustness with fast startup and graceful shutdown
10. Dev/prod parity - Keep development, staging, and production as similar as possible
11. Logs - Treat logs as event streams
12. Admin processes - Run admin/management tasks as one-off processes

Not a methodology to do cloud-native design!

<https://12factor.net/>

# The 12 Factor App

## 12 factors for the Impatient

- I. Codebase - use version control (e.g. git)
- II. Dependencies - use a dependency manager (e.g. gradle/maven/sbt)
- III. Config - separate configuration from code (use the OS environment)
- IV. Backing Services - reference resources such as DBs by URLs in the config
- V. Build.release.run - separate build from run. Use versions.
- VI. Processes - run the app as one or more stateless processes.
- VII. Port binding - app should be self-contained. No app server.
- VIII. Concurrency - scale horizontally
- IX. Disposable - fast startup, graceful shutdown
- X. Dev/Prod parity - keep environments similar
- XI. Logs - treat logs as event streams (no FileAppenders!)
- XII. Admin Processes - treat admin processes as one-off events

Not a methodology to do cloud-native design!

<https://12factor.net/>

# Microservices - principles

## One job

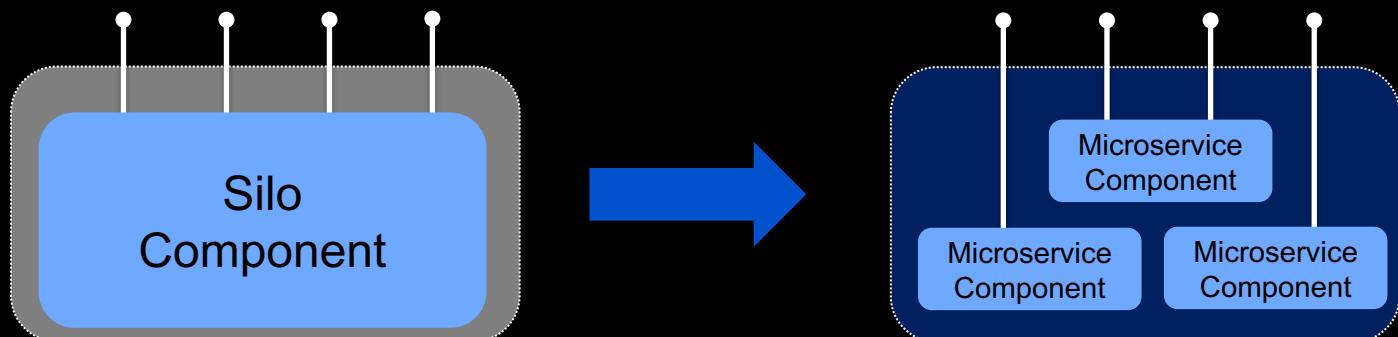
Each microservice must be **optimized for a single function**. Each service is smaller and simpler to write, maintain, and manage. Robert Martin calls this principle the "single responsibility principle."

## Separate processes

Communication between microservices must be conducted through **REST API and message brokers**. All communication from service to service must be through the service API or must use an explicit communication pattern, such as the Claim Check Pattern from Hohpe.

## Execution scope

Although microservices can expose themselves through APIs, the focus is not on interfaces, but on the running components. The granularity of a microservices application is highlighted in this figure:



# Microservices - principles

## CI/CD

Each microservice can be continuously integrated (CI) and continuously delivered (CD). When you build a large application that is composed of many services, you soon realize that different services evolve at different rates. If each service has a unique continuous integration or continuous delivery pipeline, the service can proceed at its own pace. In the monolithic approach, different aspects of the system are all released at the speed of the slowest moving part of the system.

## Resiliency

You can apply high availability and clustering decisions to each microservice. When you build large systems, another realization that you have is that when it comes to clustering, one size does not fit all. The monolithic approach of scaling all the services in the monolith at the same level can lead to the overuse or underuse of services. Even worse, when shared resources are monopolized, services might be neglected. In a large system, you can deploy services that do not need to scale to a minimum number of servers to conserve resources. Other services require scaling up to large numbers.

# Microservices - rules for developing cloud applications

1. Don't code your application directly to a **specific topology**
2. Don't assume the local **file system** is **permanent**
3. Don't keep **session state** in your application
4. Don't **log** to the file system
5. Don't assume any specific **infrastructure dependency**
6. Don't use **infrastructure APIs** from within your application
7. Don't use **obscure protocols**
8. Don't rely on **OS-specific features**
9. Don't **manually install** your application

Read the article : [http://www.ibm.com/developerworks/websphere/techjournal/1404\\_brown/1404\\_brown.html](http://www.ibm.com/developerworks/websphere/techjournal/1404_brown/1404_brown.html)

# Microservices - Advantages

In a microservices architecture each component:

- Is **developed independently** and has **limited, explicit dependencies** on other services
- Is developed by a **single, small team** in which all team members can understand the entire code base
- Is developed on its **own timetable** so new versions are delivered independently of other services
- **Scales and fails independently** which isolates any problems
- Can be developed in a different **language**
- **Manages its own data** to select the best technology and schema

# Microservices - Benefits

- **Efficient teams**
- **Simplified deployment**
- **Right tools for the job**
- **Improved application quality**
- **Scalability**

That said....

# Microservices

are

hard

..but we're  
getting ahead  
of ourselves

QUESTIONS?



# The Journey to Cloud **Docker**

02



IBM Cloud

# Everybody Loves Containers



A **standard way to package an application and all its dependencies** so that it can be moved between environments and run without changes

Containers work by **isolating the differences between applications** inside the container so that everything outside the container can be standardized

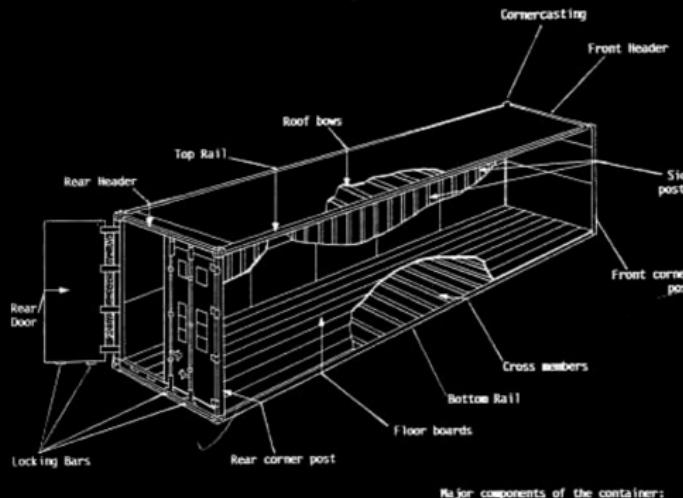
# Microservices implementation with Containers

---

## Why it works – separation of concerns

### Development

- Worries about what's “**inside**” the container
  - Code
  - Libraries
  - Package Manager
  - Apps
  - Data
- All Linux servers look the same

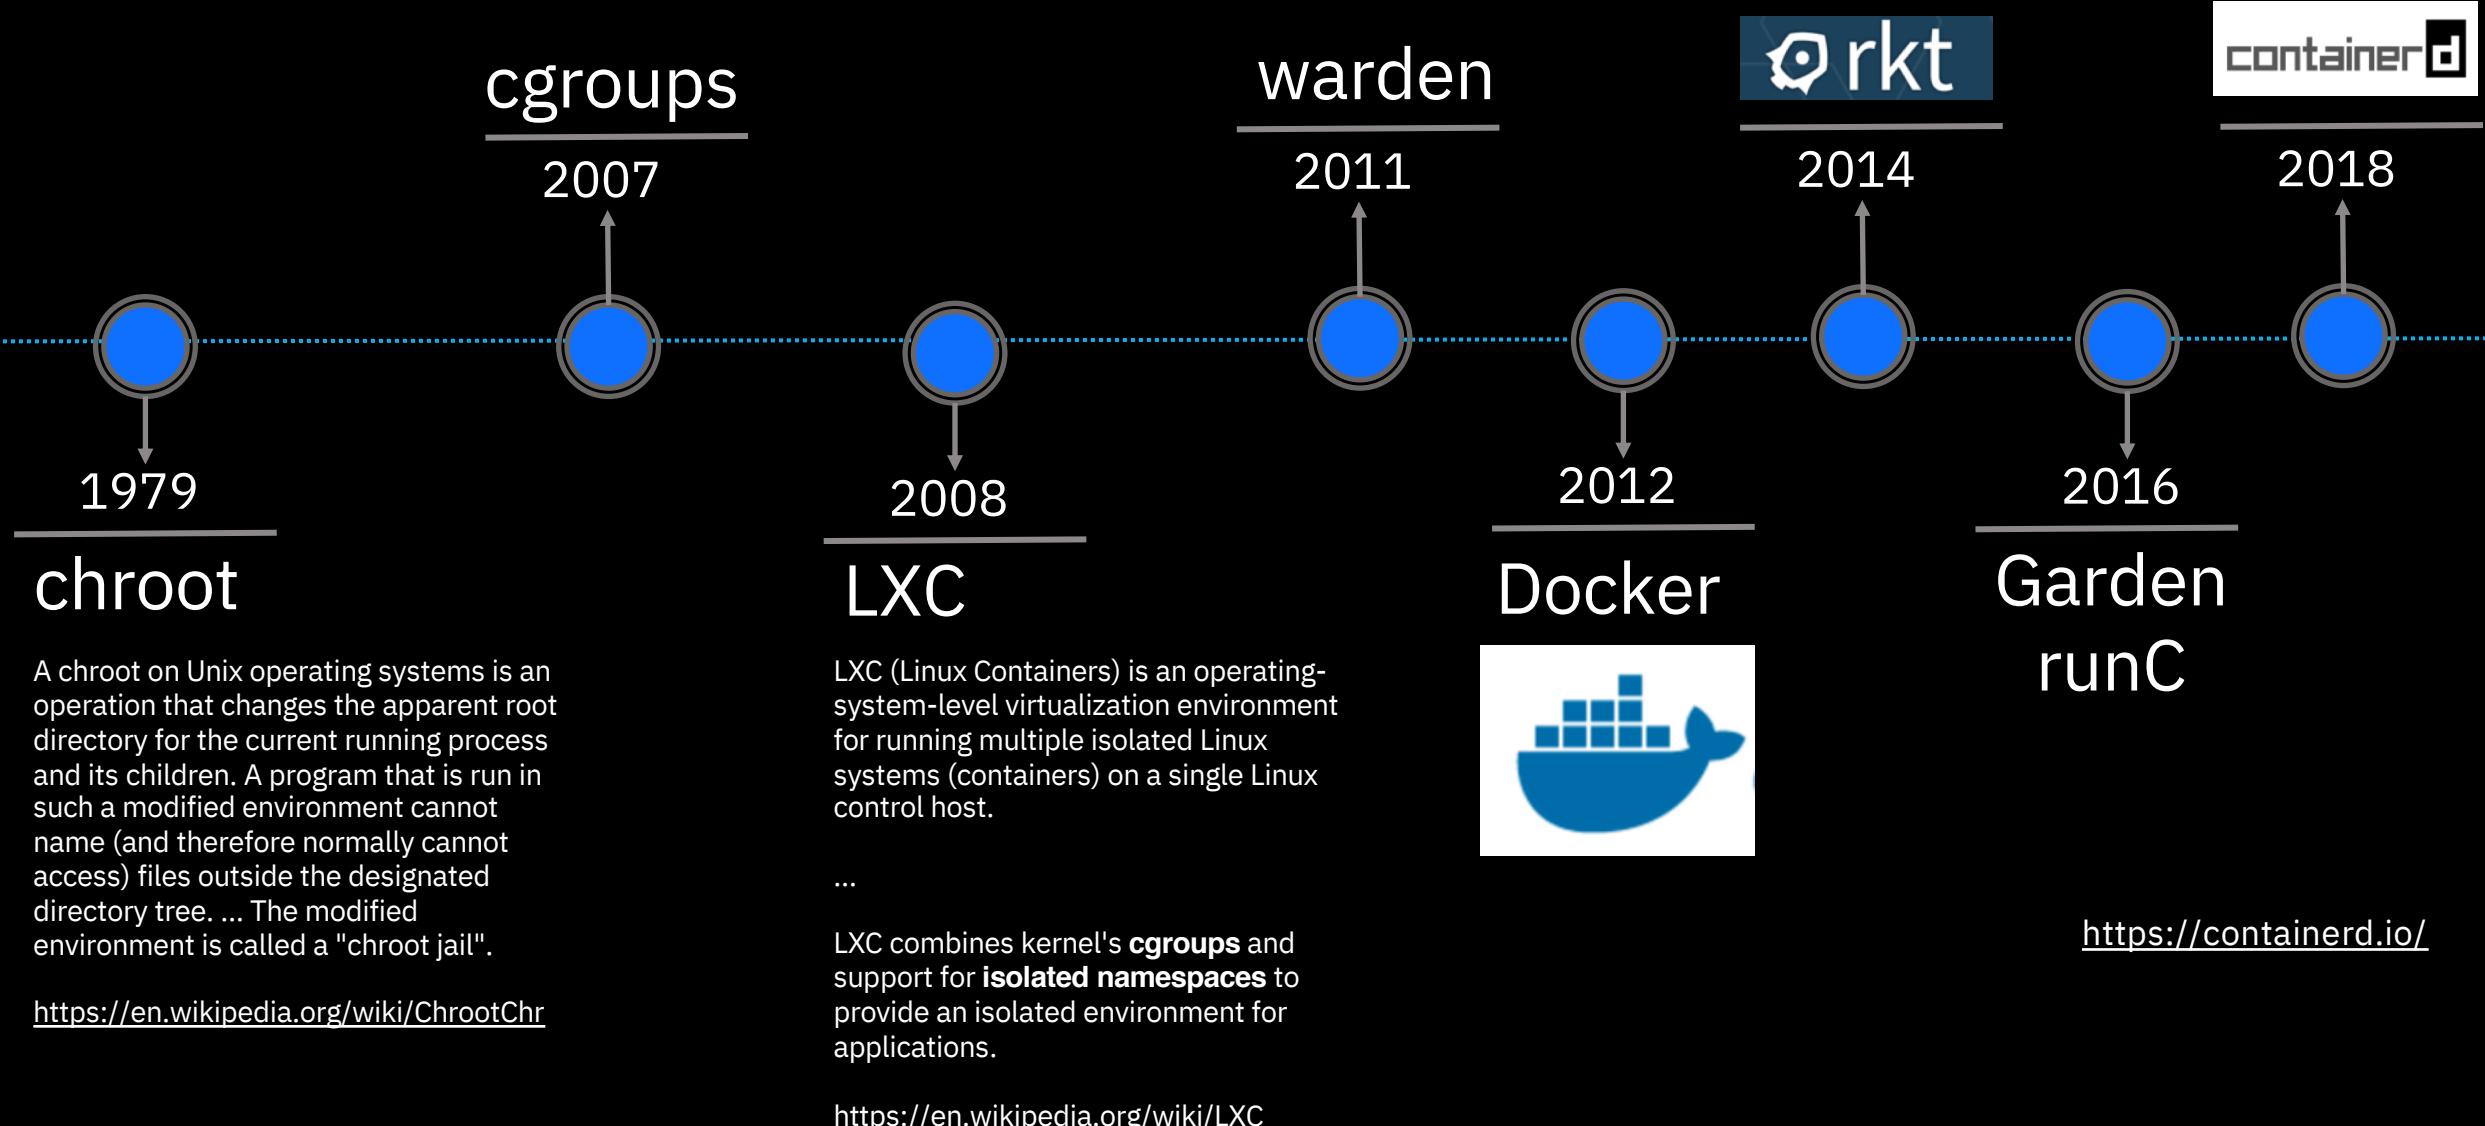


### Operations

- Worries about what's “**outside**” the container
  - Logging
  - Remote Access
  - Monitoring
  - Network Config
- All containers start, stop, copy, attach, migrate, etc... the same way

Clear ownership boundary between  
Dev and IT Ops drives DevOps adoption and fosters agility

# Container History



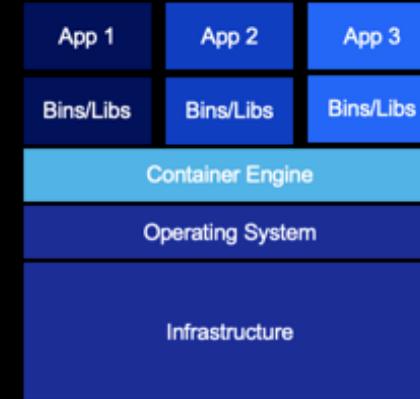
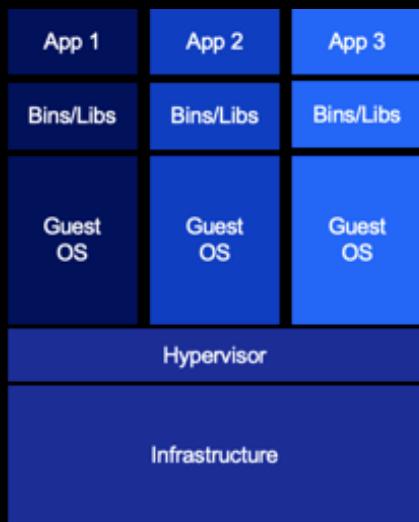
# The life of Containers

---

The process inside the container does not know that it is inside a container

It might feel a little bit lonely but the outside world looks perfectly normal to it

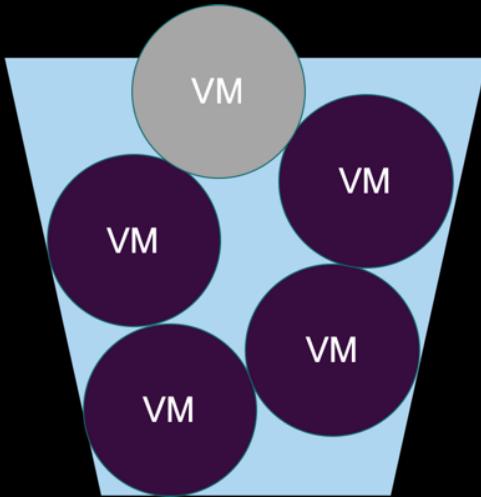
# VMs vs. Containers



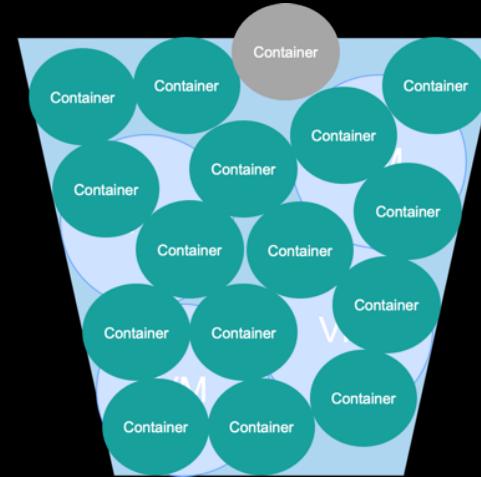
- + VM Isolation
- Complete OS
- Static Compute
- Static Memory

- + Container Isolation
- + Shared Kernel
- + Burstable Compute
- + Burstable Memory

# VMs vs. Containers



- + VM Isolation
- Complete OS
- Static Compute
- Static Memory
- Low Resource Utilization



- + Container Isolation
- + Shared Kernel
- + Burstable Compute
- + Burstable Memory
- + High Resource Utilization

# VM vs. Containers

Features	Virtual Machines (VMs)	Application Containers	Application Container Benefits
<b>Virtualization</b>	In VMs, the virtualization occurs in server hardware.	In containers, the virtualization occurs at the OS level.	Faster time-to-market
<b>Standardization</b>	VMs are standardized systems that have capabilities similar to that of bare metal computers.	Containers are not standardized, as their kernel OS has varying degrees of complexity.	Makes applications more portable by isolating them from the underlying infrastructure
<b>Resource management</b>	Each VM has its own kernel OS, binary, and library.	Containers share the same host OS, and if needed, both binary and library.	Lower overhead costs
<b>Density</b>	VMs require a few gigabytes (GB) of space that limits them into a single server.	Containers require a few megabytes (MB) of space that enables many containers to fit into a single server.	Lighter weight than VMs
<b>Boot time</b>	Minutes	Seconds	Easy to update and release newer versions of applications

# Advantages of Containers



Containers are **portable**



Containers are **easy to manage**



Containers provide “**just enough**” isolation



Containers use hardware **more efficiently**



Containers are **immutable**



# Docker Components

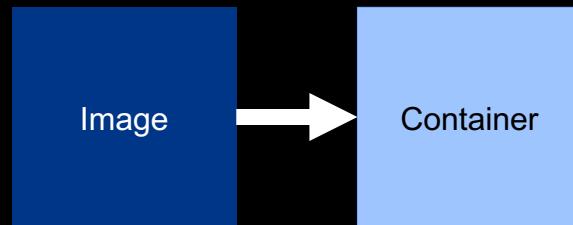
## Container

Smallest compute unit



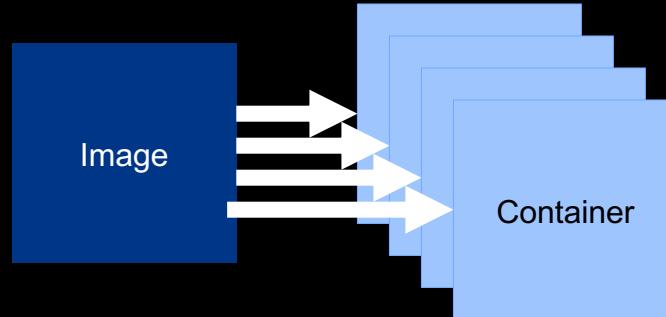
# Docker Components

**Containers**  
are created from  
**Images**



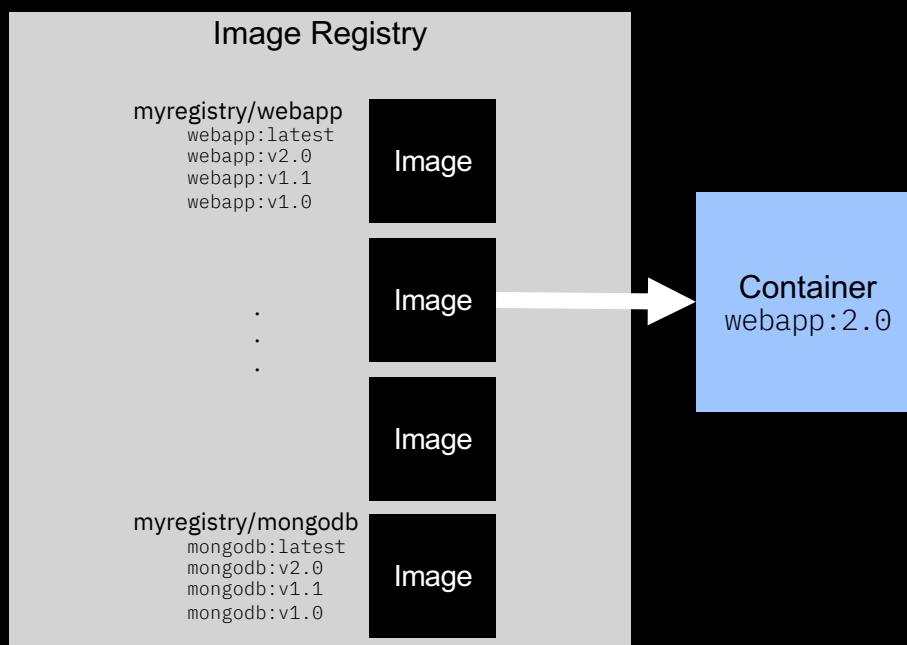
# Docker Components

As **many Containers**  
as needed can be created from  
**Images**



# Docker Components

The **Image Registry**  
stores the versioned  
**Images**  
to create  
**Containers**



# Docker Components



## Image

A read-only snapshot of a container stored in Docker Hub to be used as a template for building containers



## Container

The standard unit in which the application service resides or transported



## Docker Hub/Registry

Available in SaaS or Enterprise to deploy anywhere you choose  
Stores, distributes, and shares container images



## Docker Engine

A program that creates, ships, and runs application containers  
Runs on any physical and virtual machine or server locally, in private or public cloud. Client communicates with Engine to execute commands

# Docker Images

## Docker Layers

### Inheritance

- Build on the work of those who came before you

```
# Pull base image
FROM tomcat:8-jre8

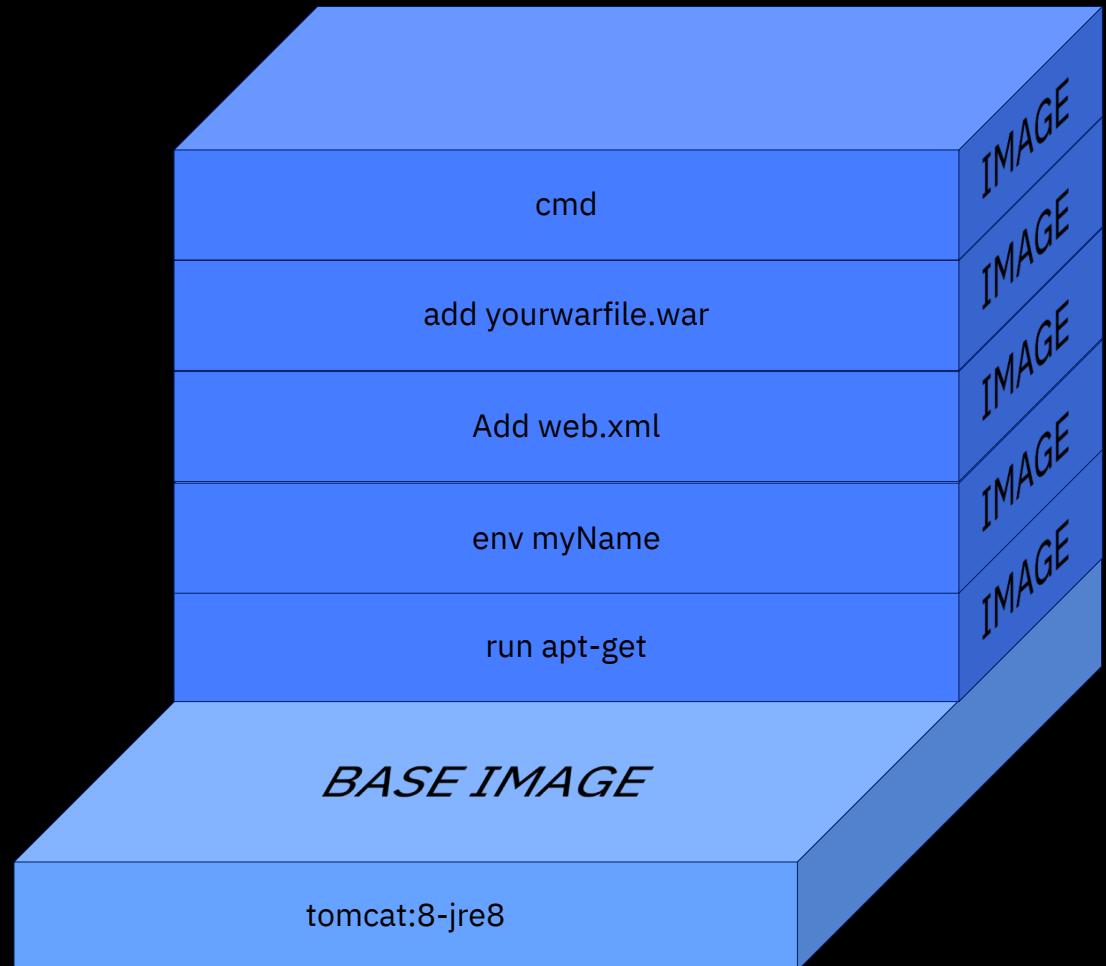
# Maintainer
MAINTAINER youremailaddress"

# Run command
RUN apt-get update && apt-get -y upgrade

# Set variables
ENV myName John Doe

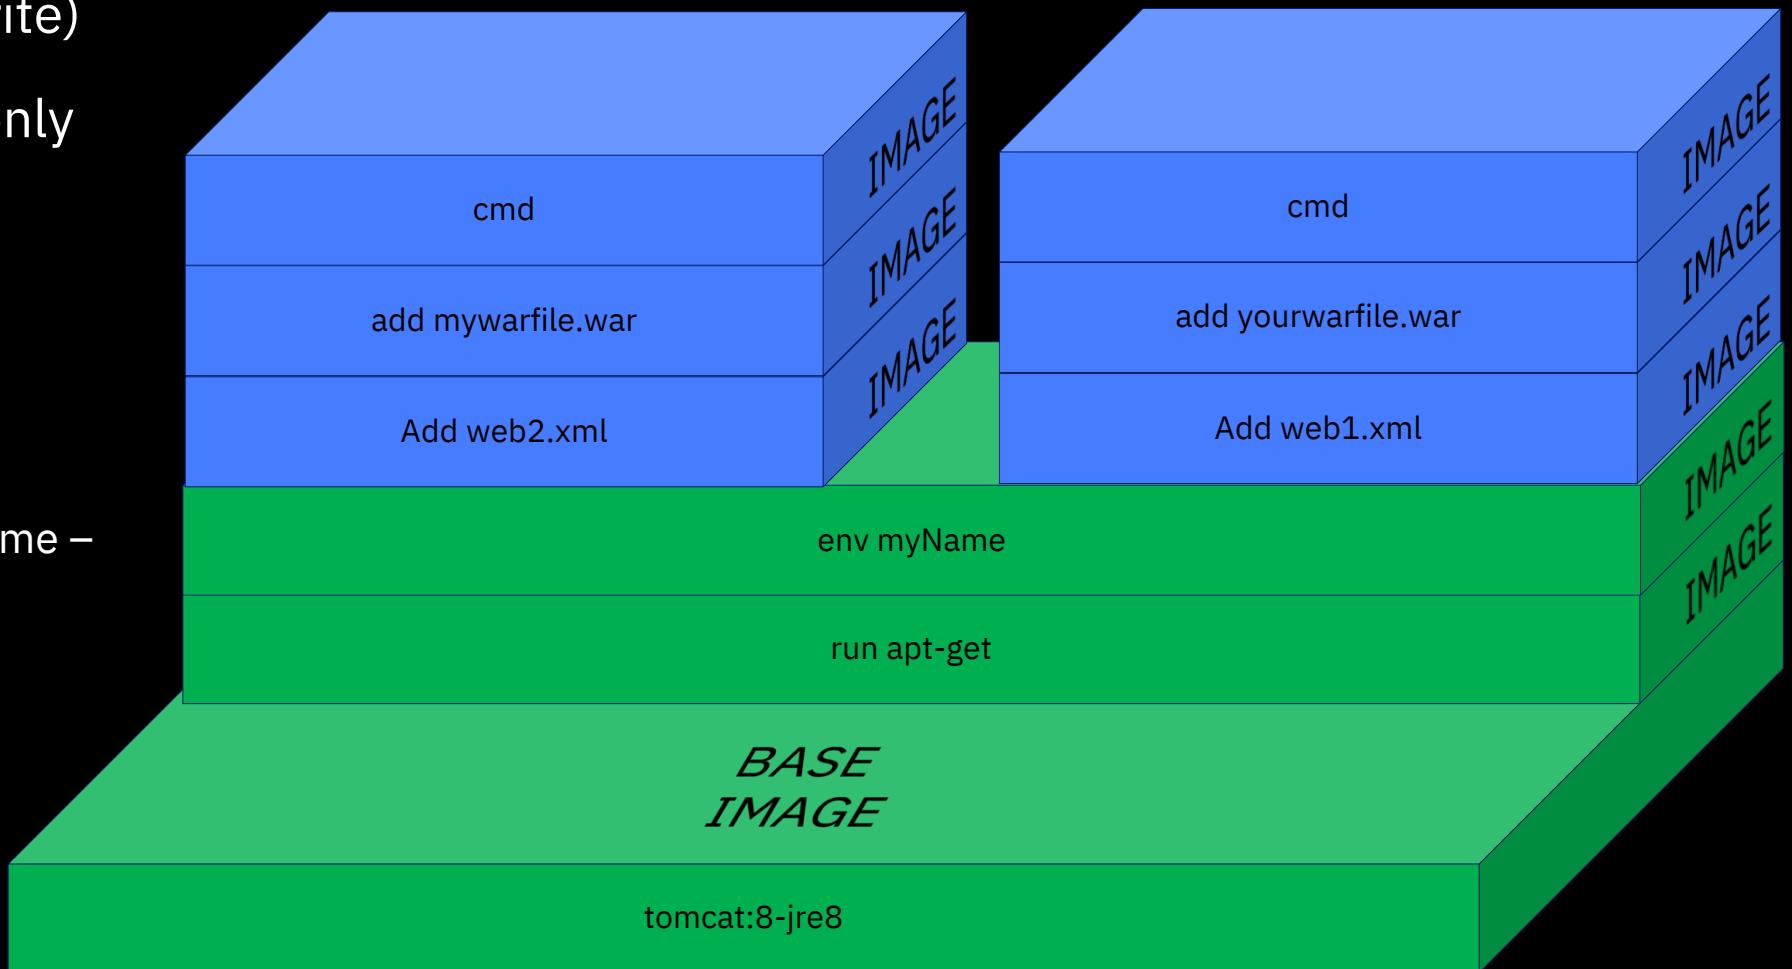
# Copy to images tomcat path
ADD web.xml /usr/local/tomcat/conf/
ADD yourwarfile.war /usr/local/tomcat/webapps/

# Run server
CMD ["catalina.sh", "run"]
```

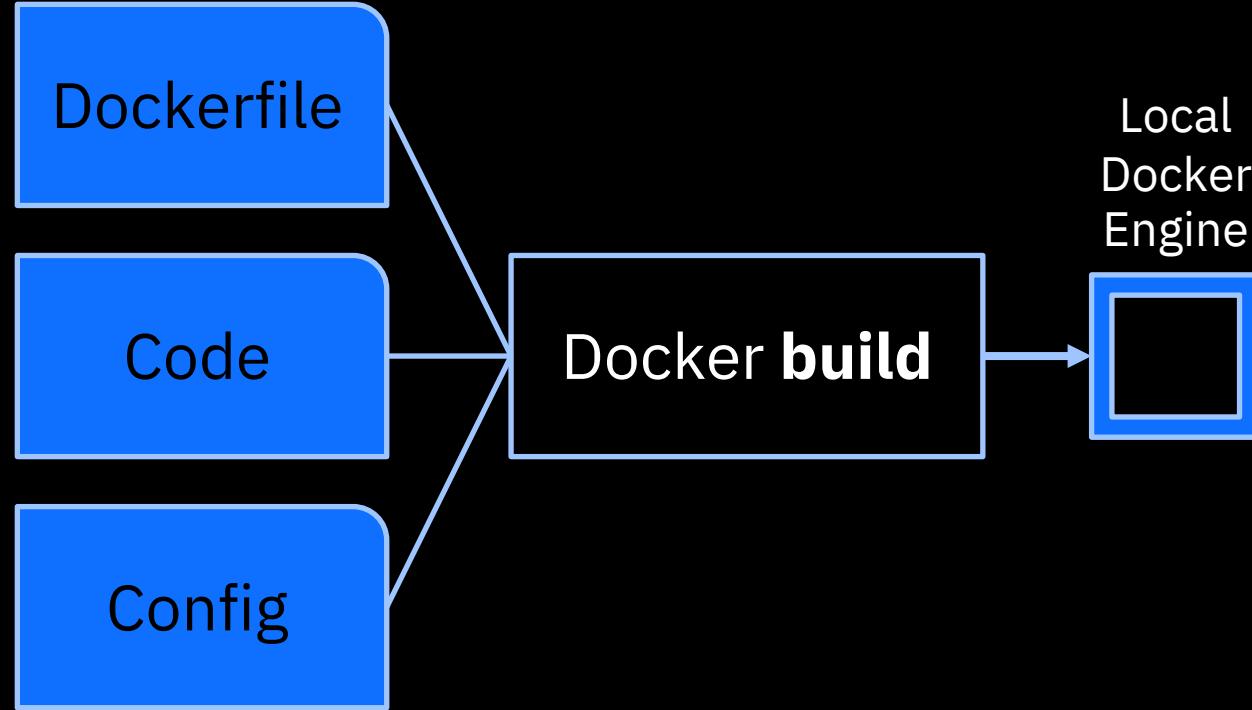


# Docker Images

- Docker uses a **layered filesystem** (copy-on-write)
- New files (& edits) are only visible to current/above layers
- Layers allow for **reuse**
  - More containers per host
  - Faster start-up/download time – base layers are "cached"



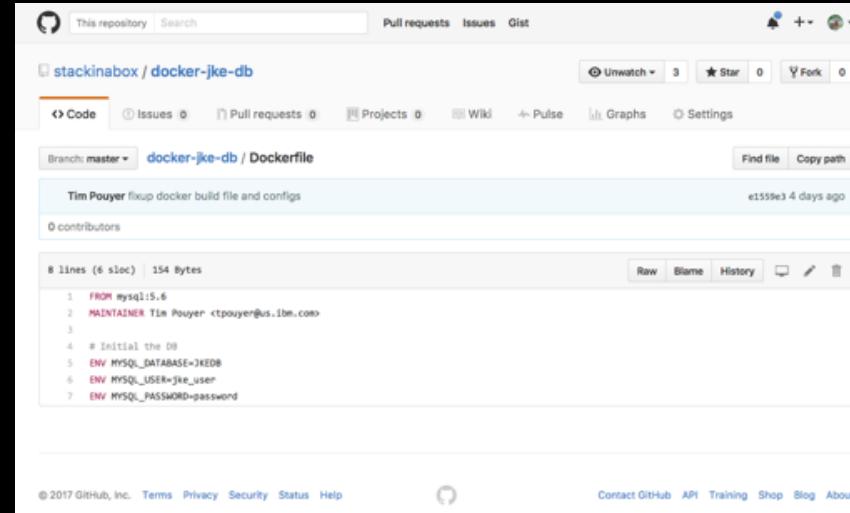
# Docker Basics – Build



# Docker Basics - Build

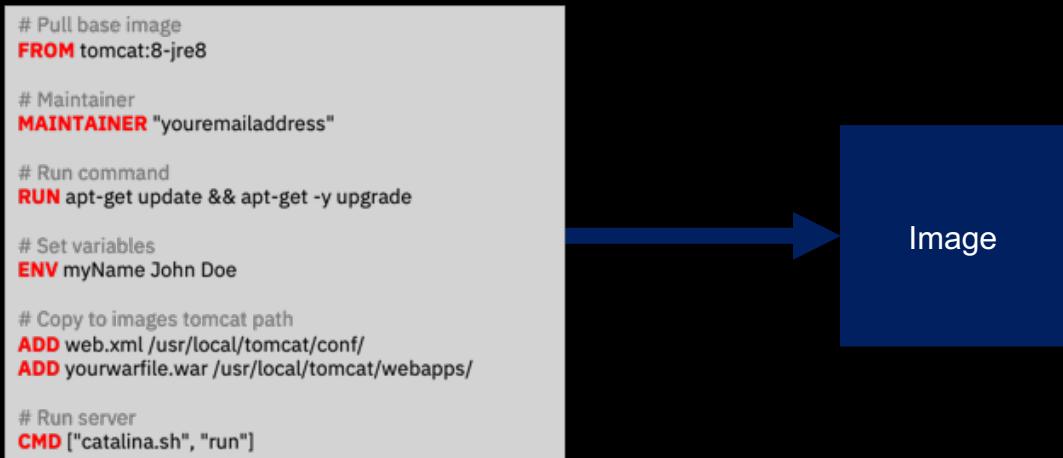
## Dockerfile

- Text based file describing:
  - Previous Layer
  - Environment Variables
  - Commands used to populate data/software/frameworks/etc...
  - Command to run when executed



A screenshot of a GitHub repository page for 'stackinabox / docker-jke-db'. The 'Dockerfile' tab is selected, showing the following code:

```
FROM mysql:5.6
MAINTAINER Tim Poyer <tpoyer@us.ibm.com>
# Initial the db
ENV MYSQL_DATABASE=jKEdb
ENV MYSQL_USER=jke_user
ENV MYSQL_PASSWORD=password
```



A diagram illustrating the build process. On the left, a white rectangular box contains a Dockerfile with various commands. A large blue arrow points from this box to a solid blue square labeled 'Image' on its right side.

```
# Pull base image
FROM tomcat:8-jre8

# Maintainer
MAINTAINER "youremailaddress"

# Run command
RUN apt-get update && apt-get -y upgrade

# Set variables
ENV myName John Doe

# Copy to images tomcat path
ADD web.xml /usr/local/tomcat/conf/
ADD yourwarfile.war /usr/local/tomcat/webapps/

# Run server
CMD ["catalina.sh", "run"]
```

# Docker Basics – Build -Dockerfile

one process per container

- A text file that builds an image using Docker directives

```
# Pull base image
FROM tomcat:8-jre8

# Maintainer
MAINTAINER "youremailaddress"

# Run command
RUN apt-get update && apt-get -y upgrade

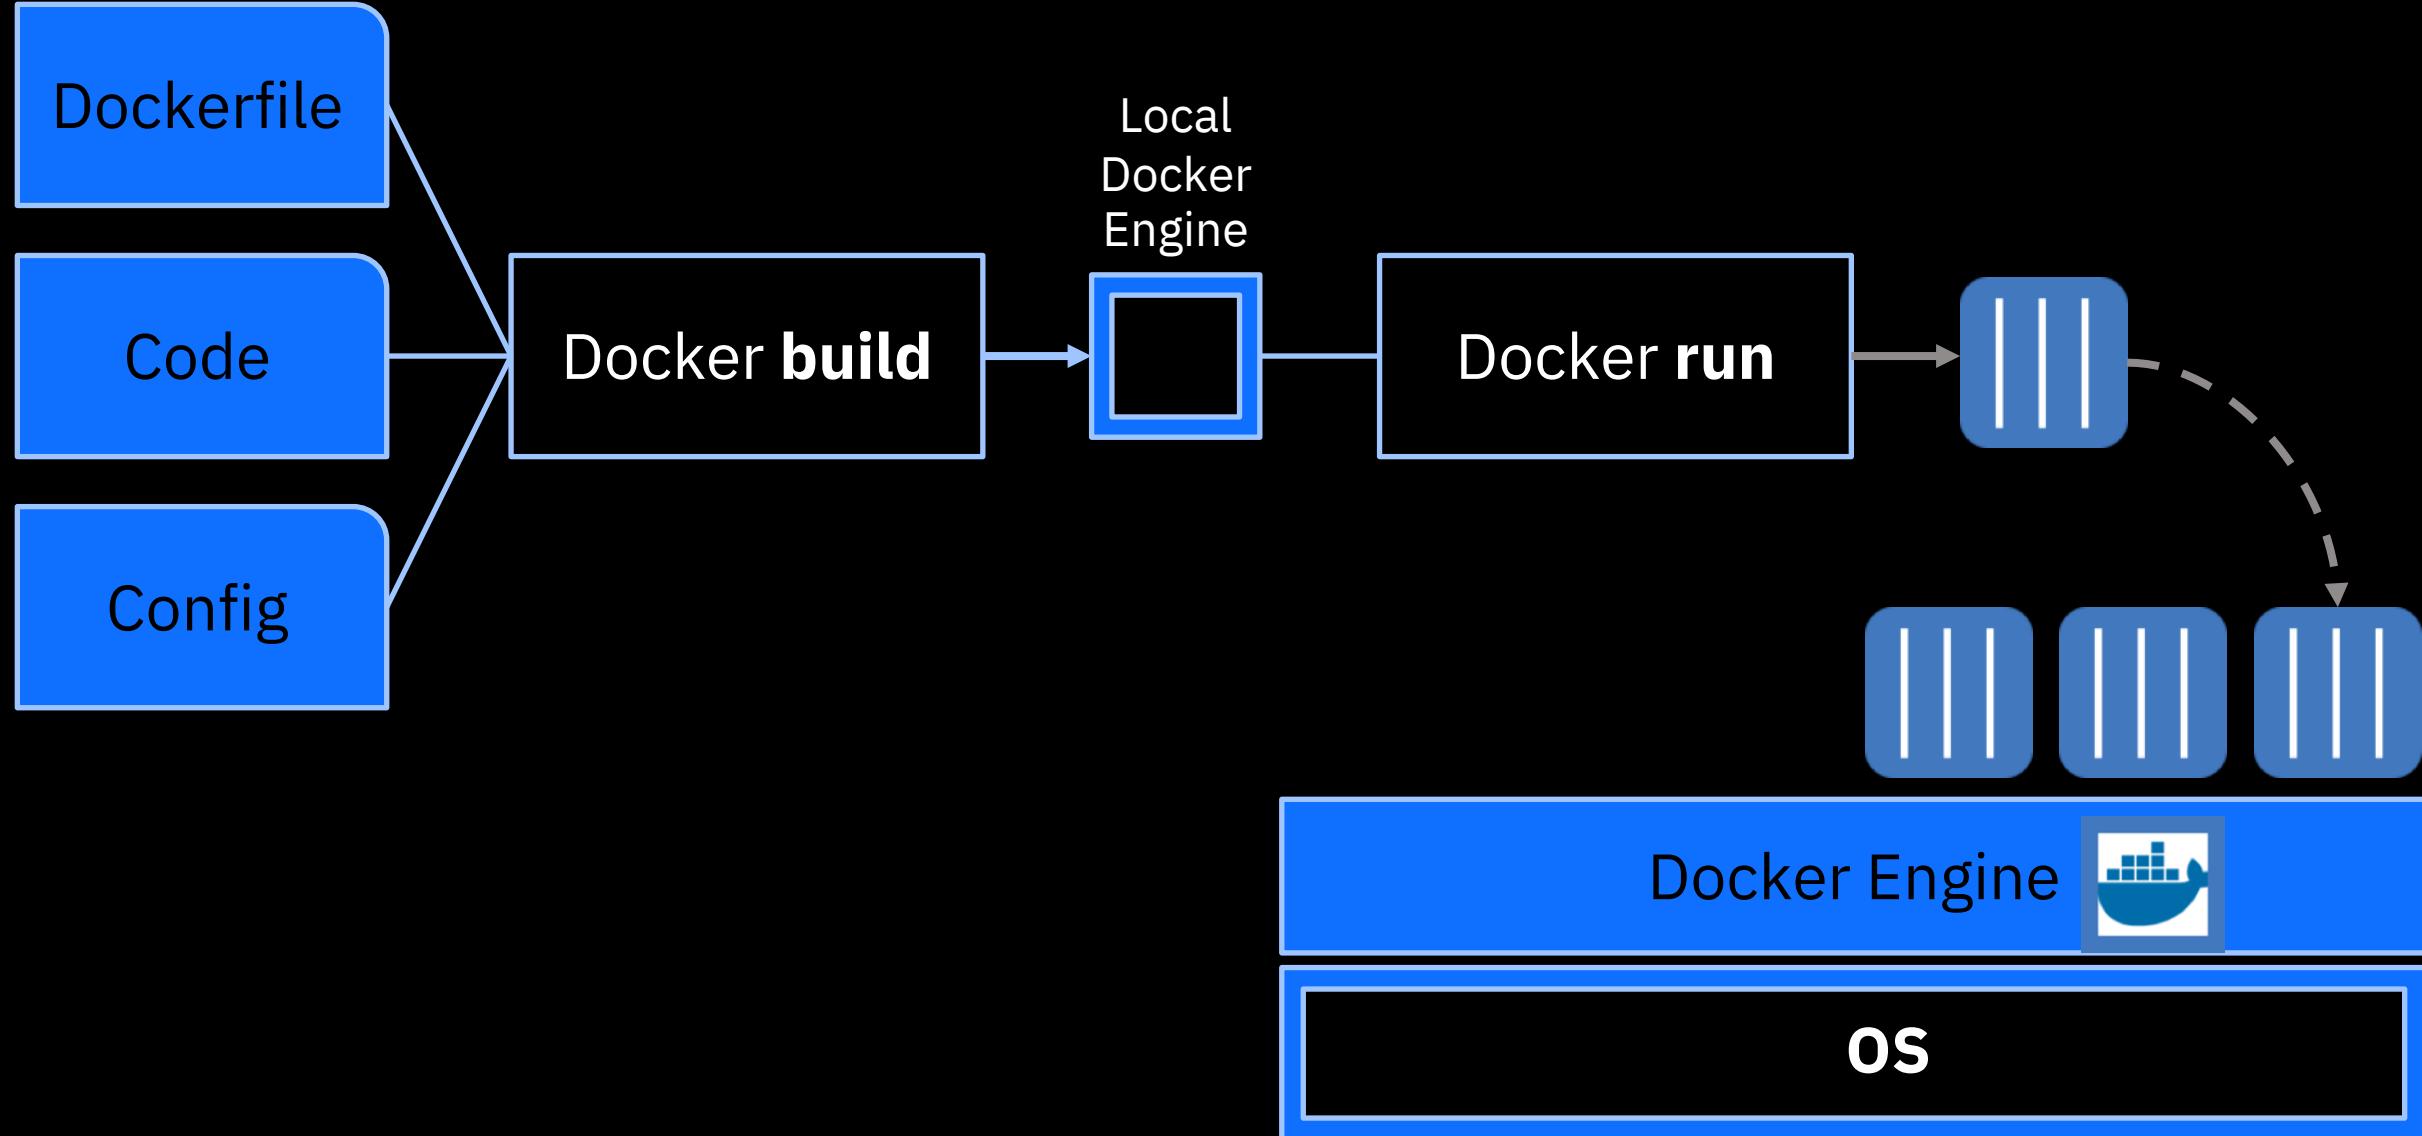
# Set variables
ENV myName John Doe

# Copy to images tomcat path
ADD web.xml /usr/local/tomcat/conf/
ADD yourwarfile.war /usr/local/tomcat/webapps/

# Run server
CMD ["catalina.sh", "run"]
```

```
docker build -t myimage ./Dockerfile
```

# Docker Basics – Run



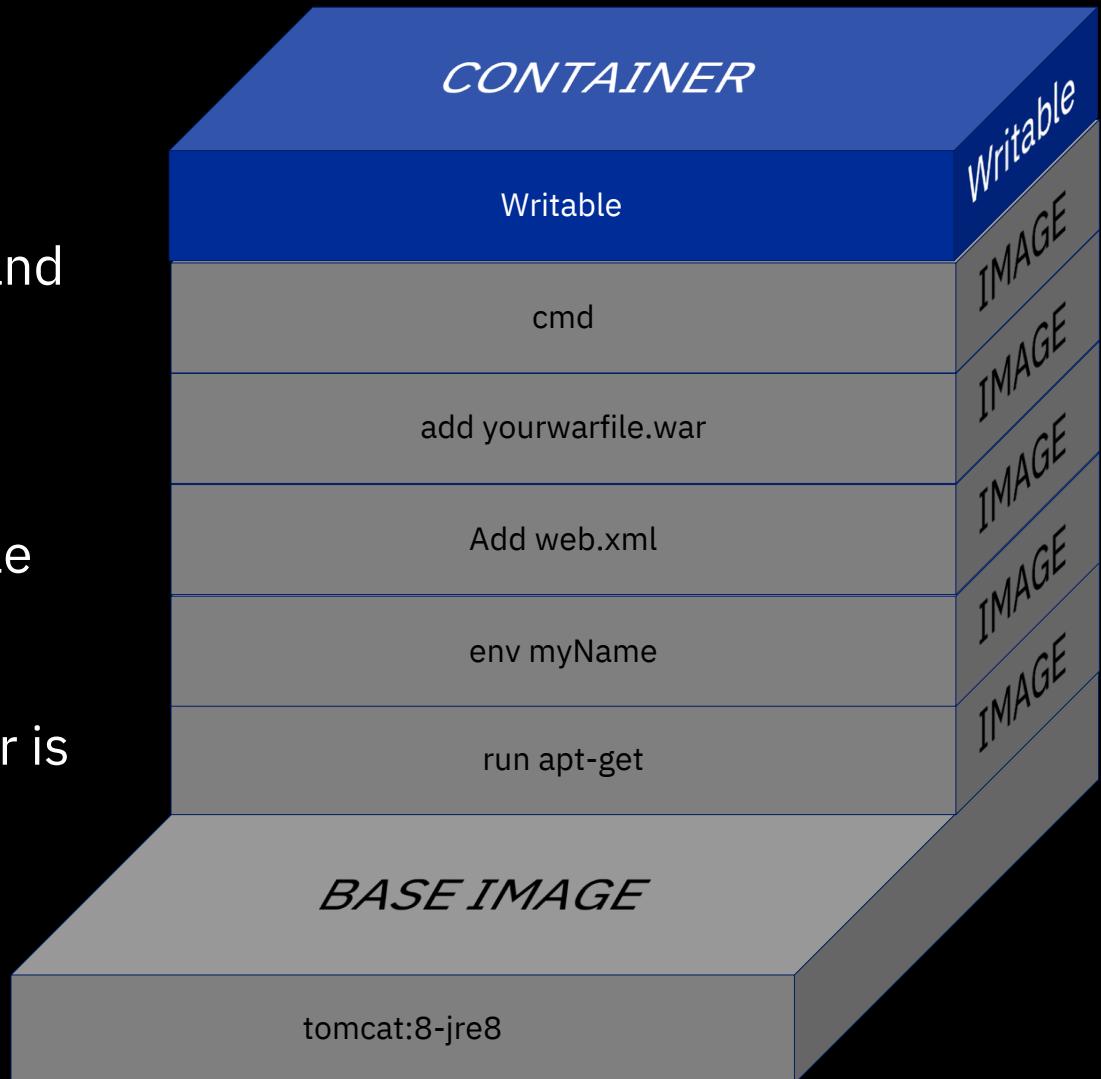
# Docker Basics – Run

---

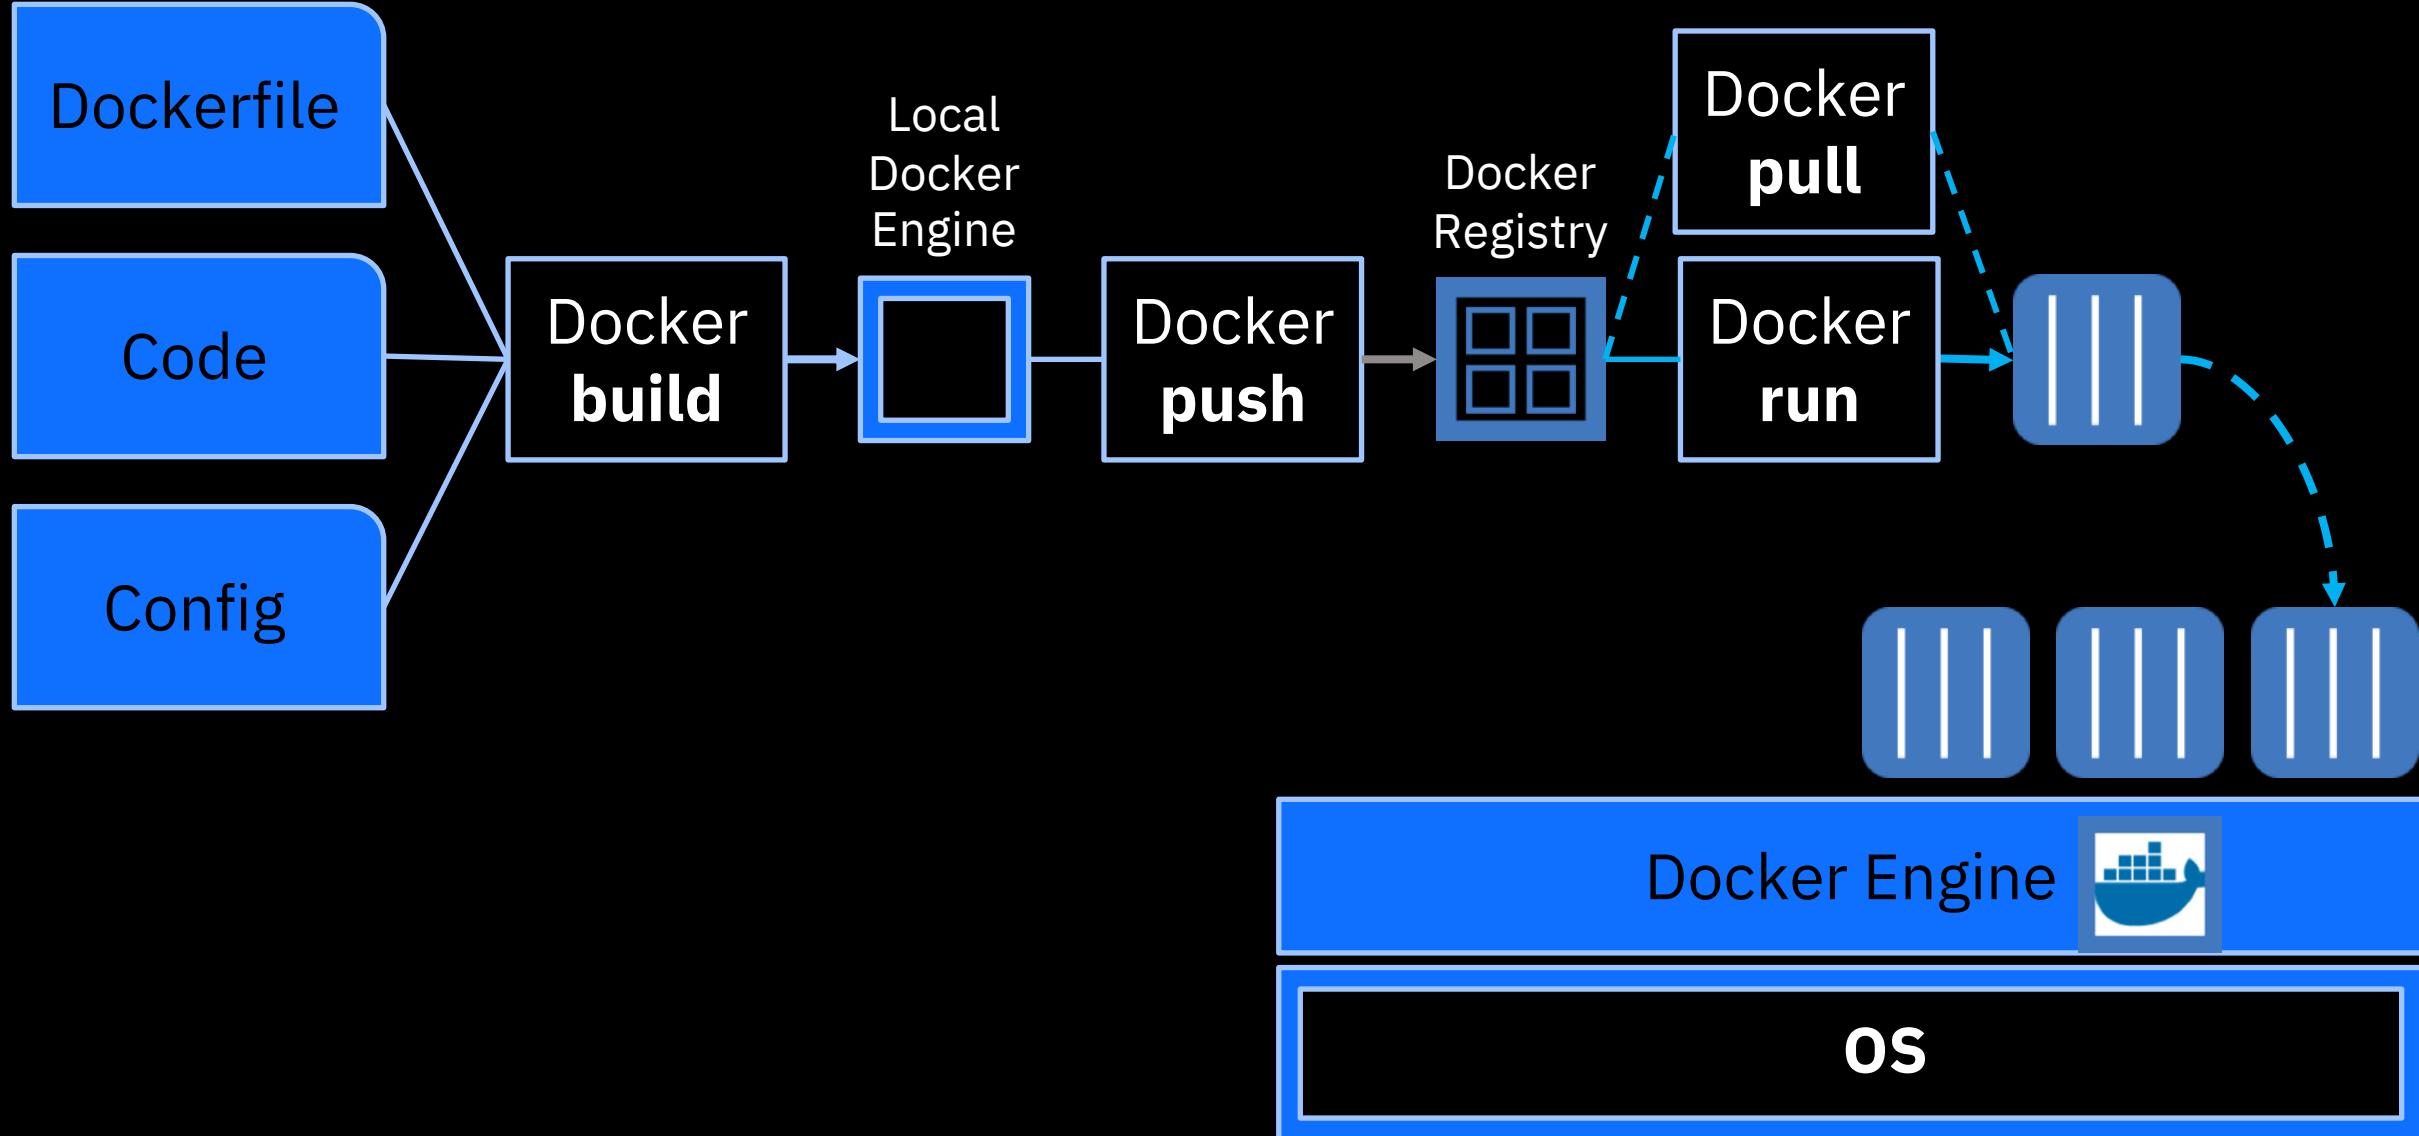
## Docker Layers

### Inheritance

- The biggest difference between a **container** and an **image** is the **top writable layer**.
- All writes to the container that add new or modify existing data are stored in this writable layer.
- When a container is deleted, its writable layer is also deleted. The underlying image is unchanged.



# Docker Basics – Store, Retrieve & Run with registry



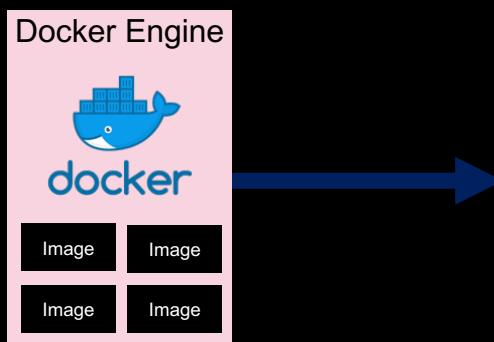
# Docker Basics – Store & Retrieve

## Docker Registry

- Private Local
- Public Docker Hub
- Private Shared

docker images

```
tpouyer at Laptop in ~/Development/workspace/docker/docker-jke-db on master
$ docker images
REPOSITORY          TAG      IMAGE ID      CREATED       SIZE
stackinabox/demo-docker    dev      4e2a1f6cc1a4   25 hours ago  528.3 MB
stackinabox/demo-jke        1.1      a596fbf2c3b7   3 days ago   672 MB
stackinabox/demo-jke        dev      a596fbf2c3b7   3 days ago   672 MB
stackinabox/demo-jke        latest   a596fbf2c3b7   3 days ago   672 MB
stackinabox/demo-jke        1.1      a596fbf2c3b7   3 days ago   672 MB
stackinabox/demo-jke        latest   a596fbf2c3b7   3 days ago   672 MB
stackinabox/jke-web         latest   a4be0cdad8bc   4 days ago   462.3 MB
stackinabox/demo-jke-db     dev      54bdbf42b999   4 days ago   327.5 MB
stackinabox/jke-db          dev      54bdbf42b999   4 days ago   327.5 MB
stackinabox/jke-db          latest   53c878fbf034   7 days ago   477 MB
websphere-liberty           6.2.2.0  7e27c3fb9c96   10 days ago  445.7 MB
mysql                       5.6      a896fd82dcfd5  13 days ago  327.5 MB
mysql                       latest   f3694c67abdb   13 days ago  400.1 MB
$
```



myregistry/webapp

webapp:latest  
webapp:v2.0  
webapp:v1.1  
webapp:v1.0

myregistry/mongodb

mongodb:latest  
mongodb:v2.0  
mongodb:v1.1  
mongodb:v1.0

# Docker Basics – Run

## Docker run

- Local (containerd)

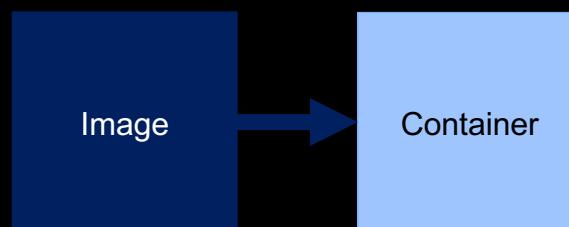
```
tpouyer at Laptop in ~/Development/workspace/docker/docker-jke-db on master
$ docker run -d postgres:latest
4eec29ae5eaa20798b1af8fa37873f7fc28cbb7cb789986b44f66cd94a784e0a

tpouyer at Laptop in ~/Development/workspace/docker/docker-jke-db on master
$ docker ps
CONTAINER ID        IMAGE               COMMAND      CREATED          STATUS          PORTS
ORTS
4eec29ae5eaa        postgres:latest   "/docker-entrypoint.s"   5 seconds ago   Up 4 seconds   5
432/tcp

tpouyer at Laptop in ~/Development/workspace/docker/docker-jke-db on master
$ 
```

```
docker run -d -e MYVAR=foo myimage:1.0.0
docker run -ti myimage:1.0.0 /bin/bash
```

...



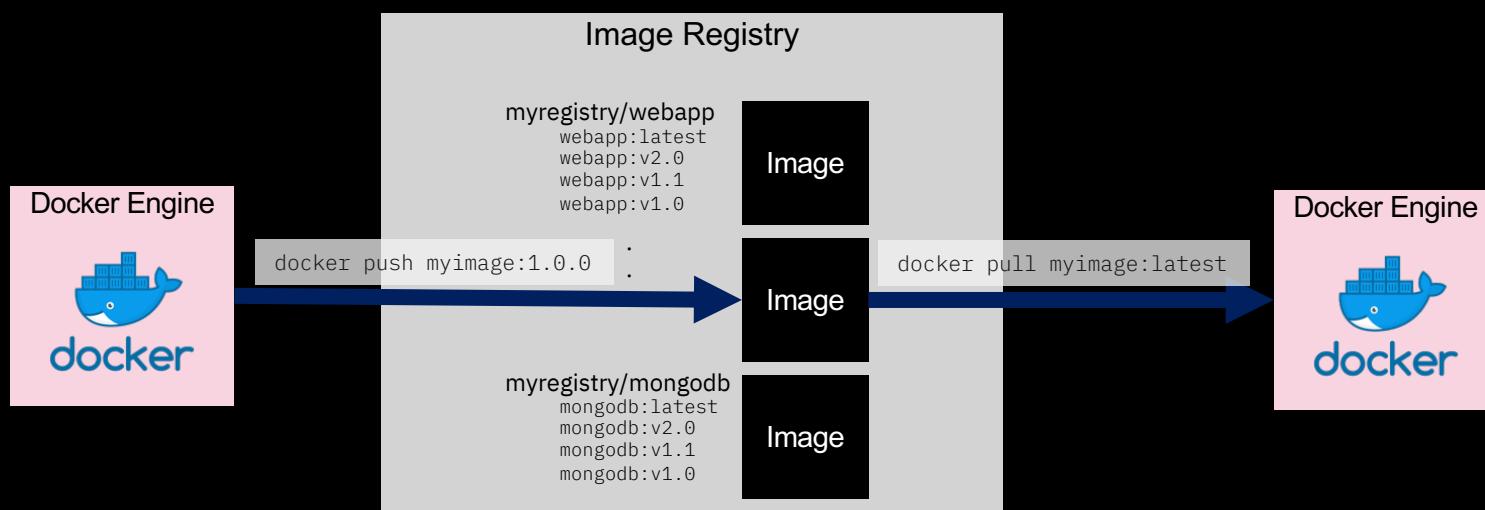
# Docker Basics – Store & Retrieve with registry

## Docker Registry

- Private Local
- Public Docker Hub
- Private Shared

```
docker pull myimage:latest  
docker push myimage:1.0.0
```

```
touyer at Laptop in ~/Development/workspace/docker/docker-jke-db on master  
$ docker pull postgres:latest  
latest: Pulling from library/postgres  
5040bd298390: Already exists  
f08454c3c700: Pull complete  
4db038cdfe03: Pull complete  
e1d9ba315f03: Pull complete  
25e8ee93170e: Pull complete  
3f28084c3f51: Pull complete  
78c91f0aedcd: Pull complete  
93ab52dbcbb8: Pull complete  
27ec75825613: Pull complete  
28ef691a9920: Pull complete  
0f0dd28755c9: Pull complete  
2a4a824861f7: Pull complete  
Digest: sha256:0842a7ef786aa2658623885160cb38451eb3d40856e7d222ae8069b6e6296877  
Status: Downloaded newer image for postgres:latest  
  
touyer at Laptop in ~/Development/workspace/docker/docker-jke-db on master  
$
```



# Docker Basics – Registries

## Hosting image repositories

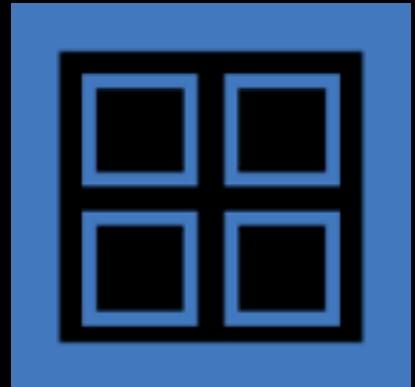
- You can define your own registry
- A registry is managed by a registry container

## Public and Private registries

- Public Registry like **Docker Hub**
- <https://hub.docker.com>

## Login into the registry

- Docker login domain:port



# Docker Recap

- **Containers are not VMs**
- **Containers provide many benefits:**
  - Efficiency
  - Portability
  - Consistency
- **New challenges with containers:**
  - Production apps dependent on open-source projects
  - Existing tools may not be sufficient for container
  - Need to focus on business objectives

QUESTIONS?



# The Journey to Cloud **Kubernetes**

03



IBM Cloud

# Everybody Loves Containers

But when you go



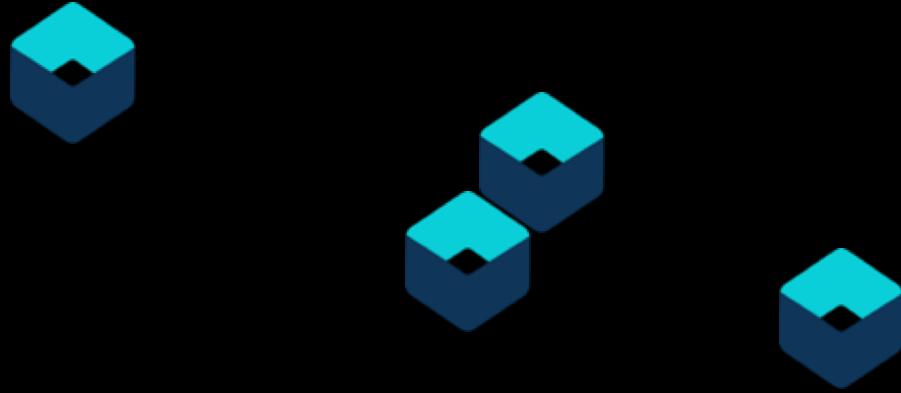
From this....



To this....



Everyone's container journey starts with one container....



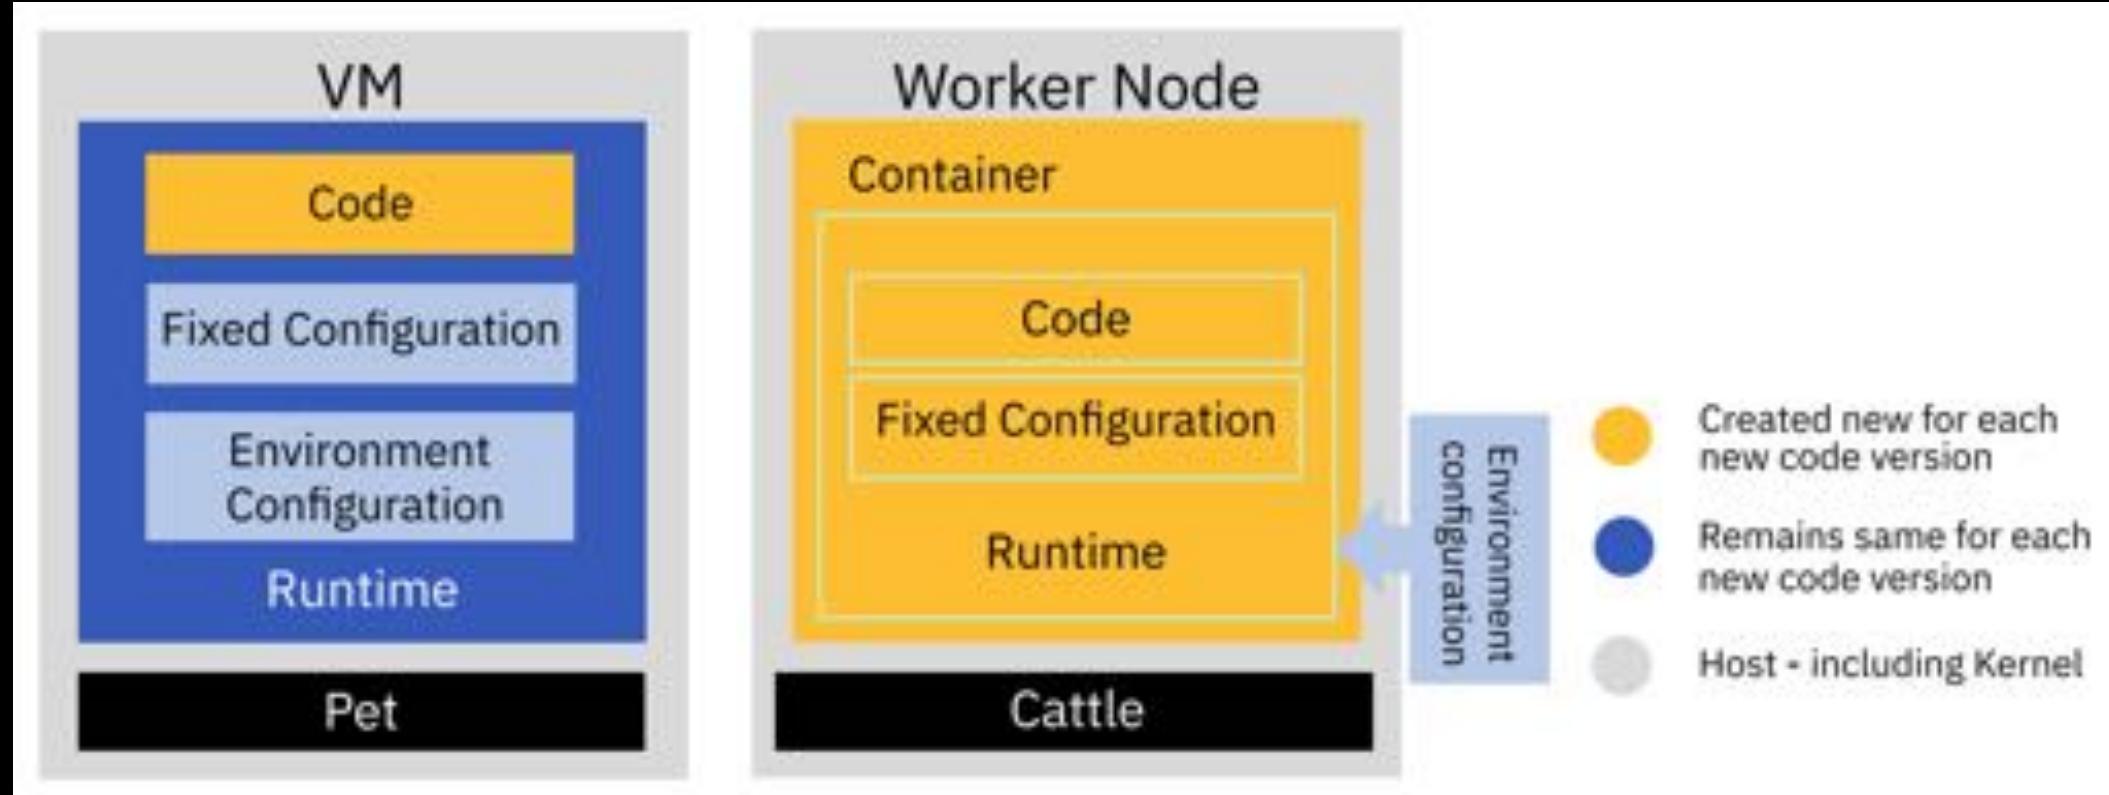
At first the growth is easy to handle....



# Pets vs Cattle

But soon you have many applications, many instances...

# Pets vs Cattle

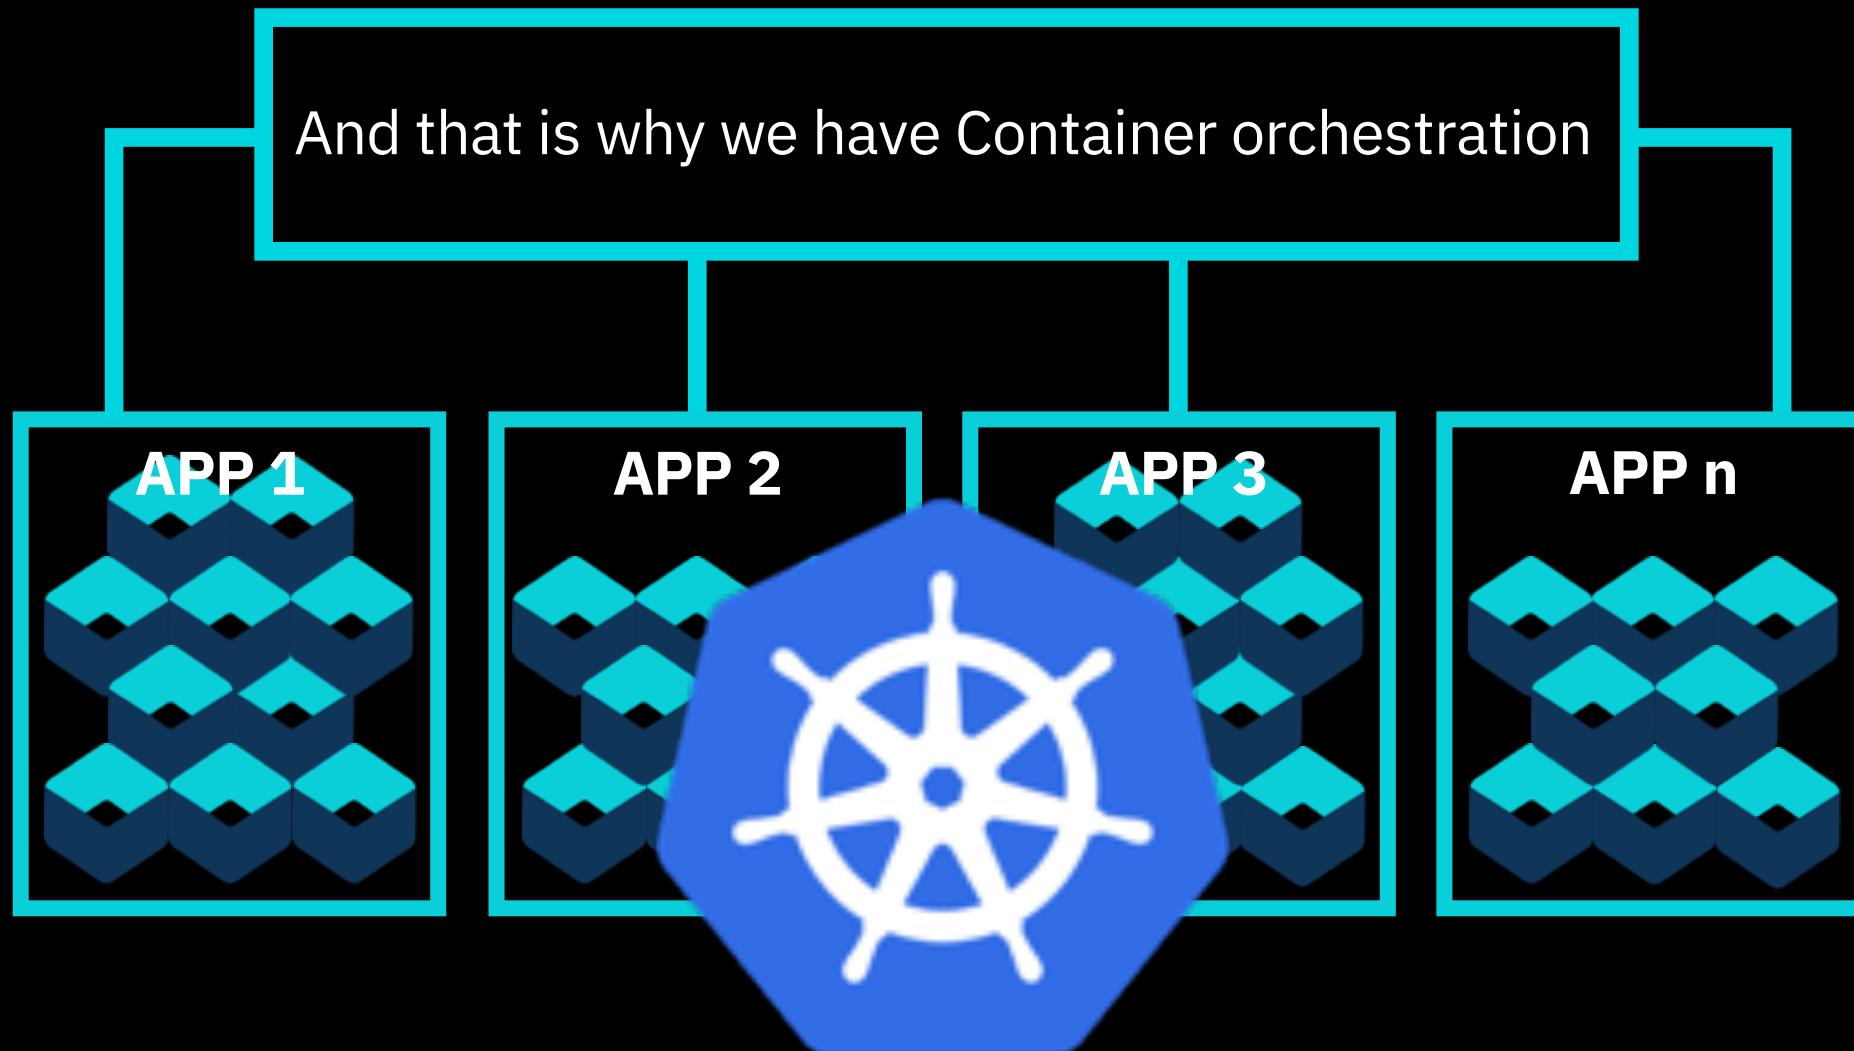


# The trade off

**Improved delivery velocity**  
in exchange for  
**increased operational complexity**



# Enter Kubernetes (K8s)



# Kubernetes – Declarative System



## Imperative Systems

In an imperative system, **the user** knows the desired state and **the user** determines the sequence of commands to transition the system to the desired state.

## Declarative Systems

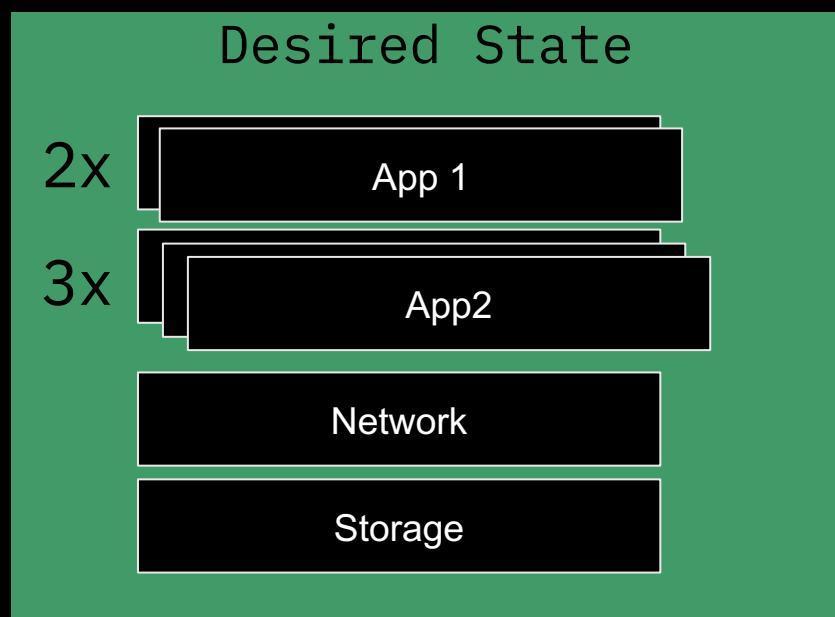
By contrast, in a declarative system, **the user** knows the desired state, supplies a representation of the desired state to the system, then **the system** reads the current state and determines the sequence of commands to transition the system to the desired state.

# Kubernetes – Declarative System

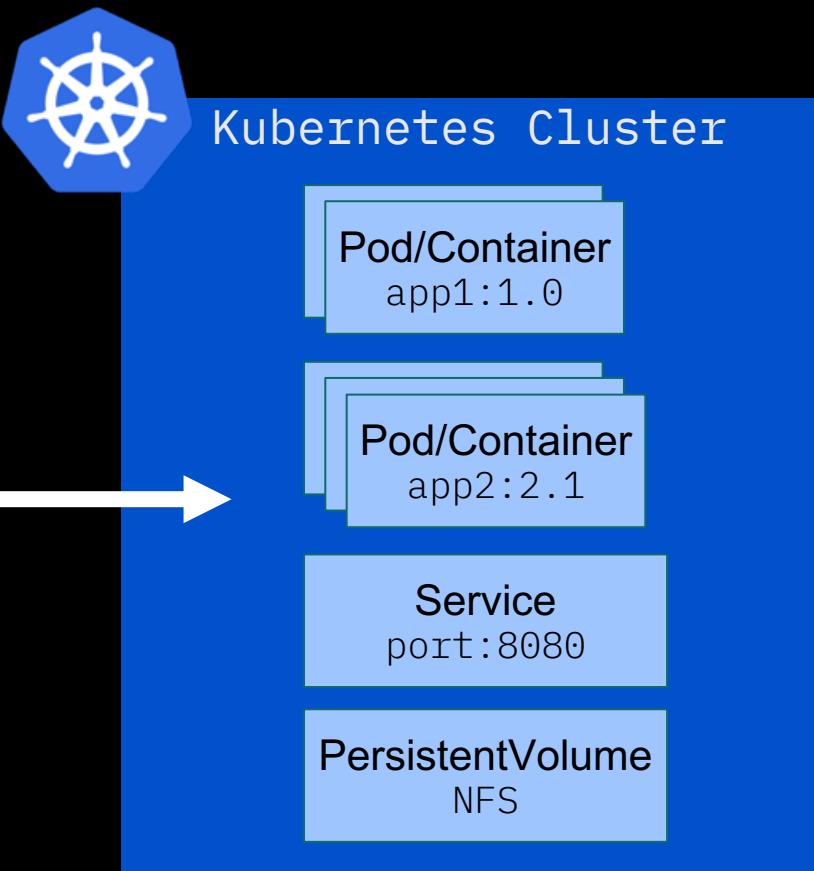


## The Desired State

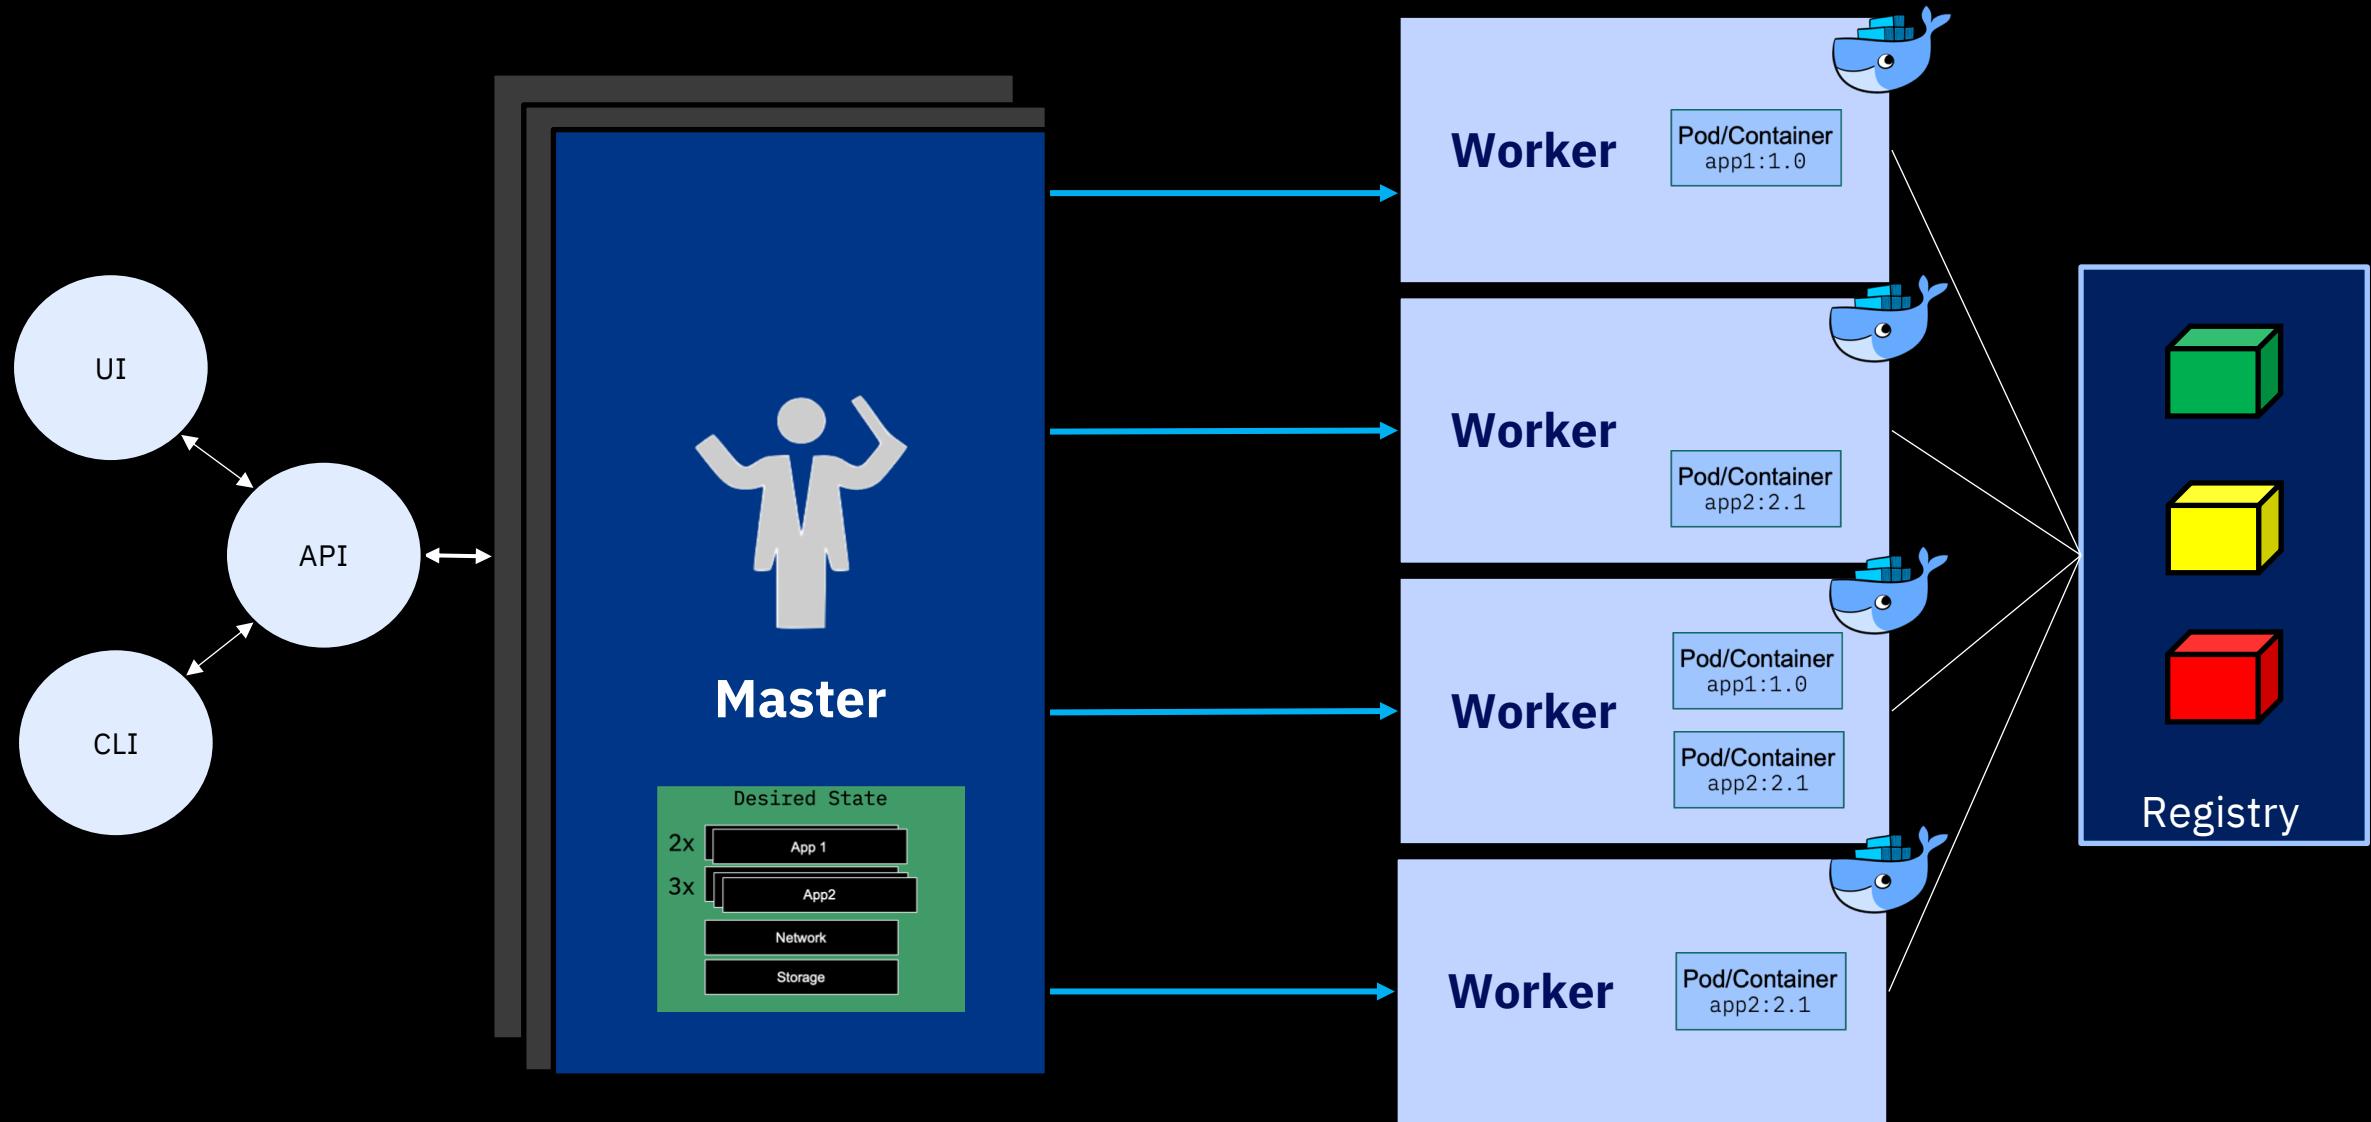
Kubernetes ensures that all the containers running across the cluster are in the desired state at any moment.



```
kind: Deployment
apiVersion: extensions/v1beta1
metadata:
  name: mcmk-ibm-mcmk-prod-klusterlet
  labels:
    app: ibm-mcmk-prod
    chart: ibm-mcmk-prod-3.1.2
    component: "klusterlet"
    release: mcmk
    heritage: Tiller
spec:
  replicas: 1
  selector:
    matchLabels:
      app: ibm-mcmk-prod
      component: klusterlet
      release: mcmk
  template:
    metadata:
      labels:
        app: ibm-mcmk-prod
        component: "klusterlet"
        release: mcmk
        heritage: Tiller
        chart: ibm-mcmk-prod-3.1.2
      annotations:
        productName: "IBM Multi-cloud Manager - Klusterlet"
        productID: "354b8990aab44c9988a0edfdal01b128"
        productVersion: "3.1.2"
    spec:
```



# Kubernetes Management Architecture





# What is Kubernetes?

## Container orchestrator

- Runs and manages containers
- Unified API for deploying web applications, batch jobs, and databases
- Maintains and tracks the global view of the cluster
- Supports multiple cloud and bare-metal environments

## Manage applications, not machines

- Rolling updates, canary deploys, and blue-green deployments

## Designed for extensibility

- Rich ecosystem of plug-ins for scheduling, storage, networking

## Open source project managed by the Linux Foundation

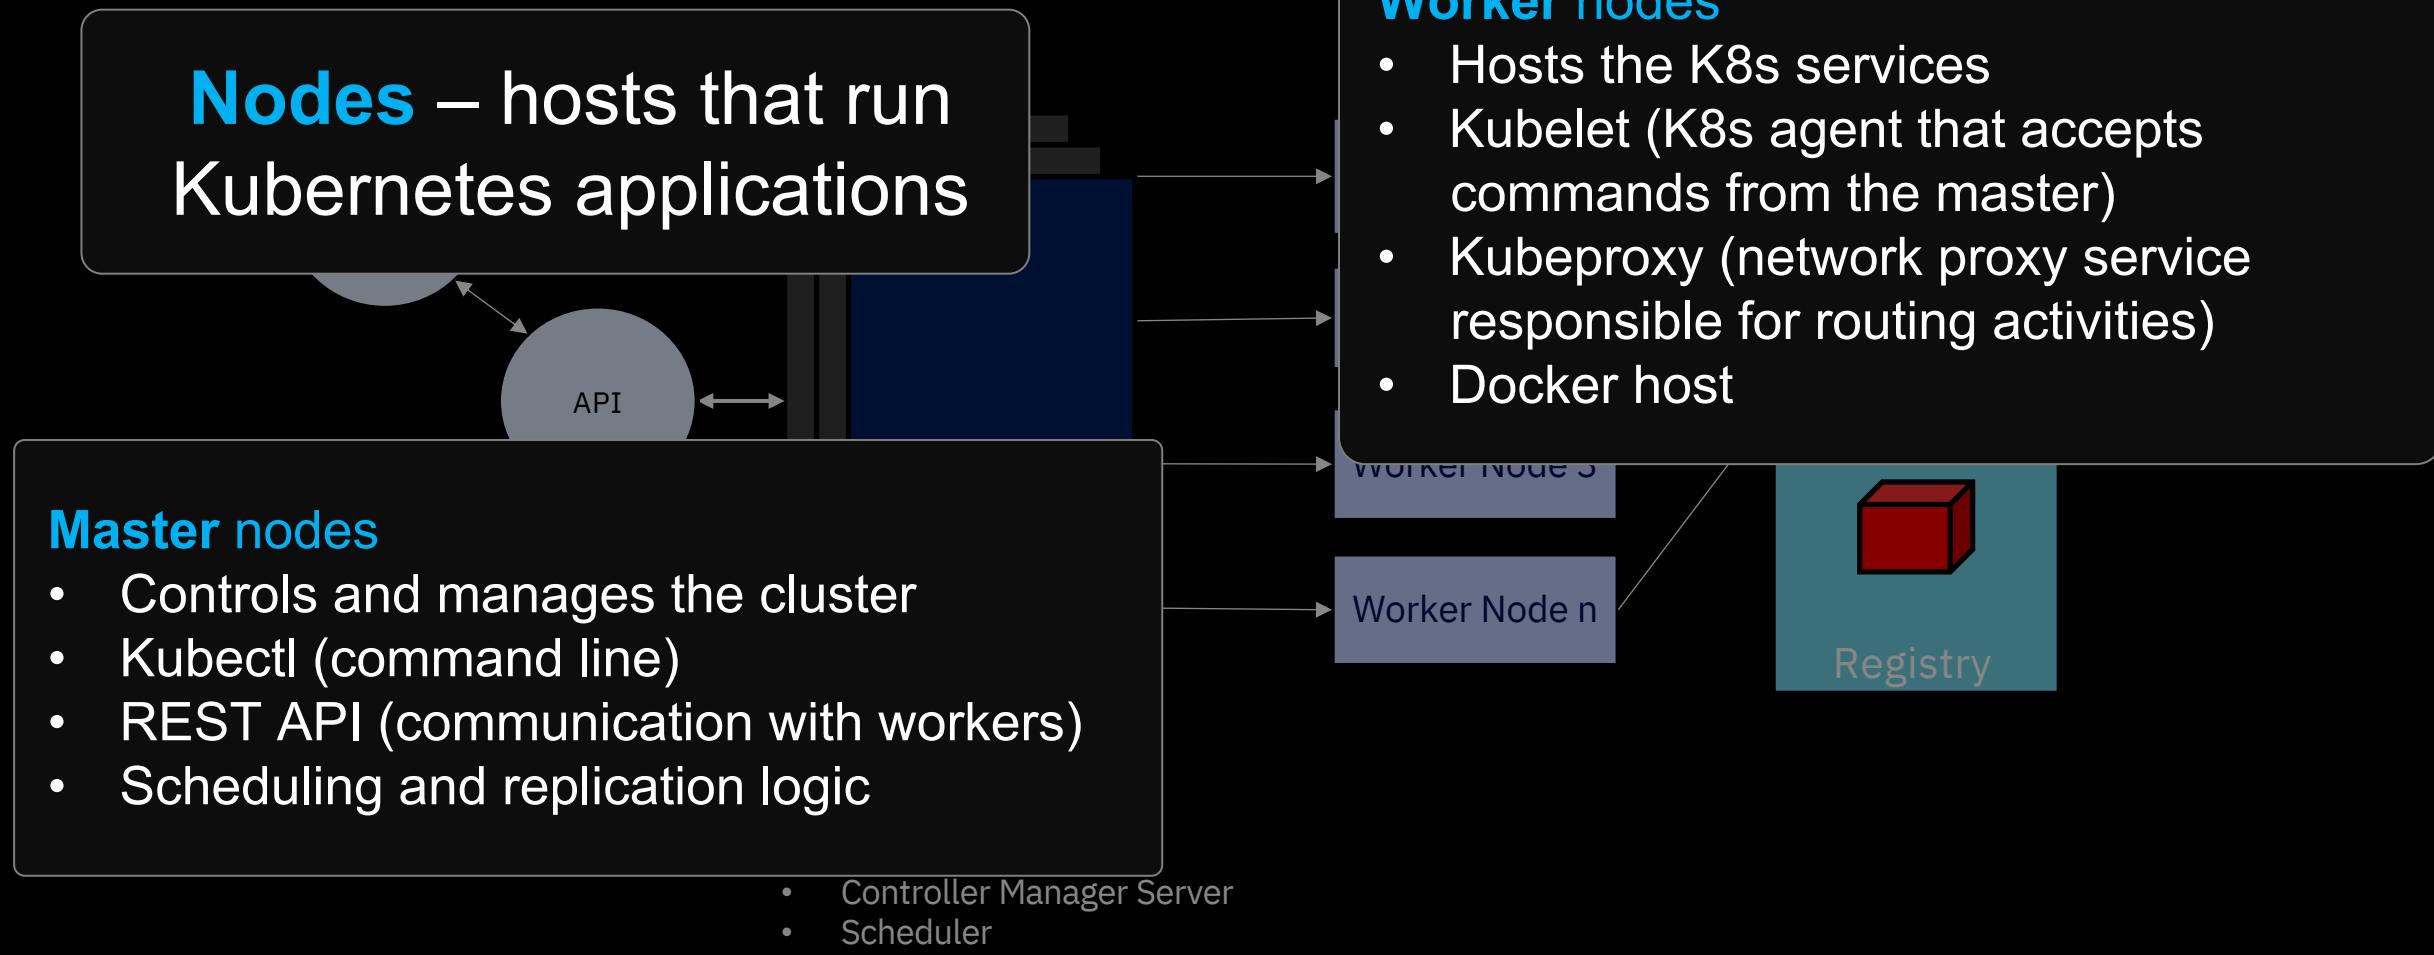
- Inspired and informed by Google's experiences and internal systems
- 100% open source, written in Go

# Kubernetes Strengths



- Kubernetes has a **clear governance model** managed by the Linux Foundation
- A growing and **vibrant** Kubernetes **ecosystem**
- Kubernetes **avoids dependency and vendor lock-in**
- Kubernetes supports a **wide range of deployment options**

# Kubernetes Management Architecture

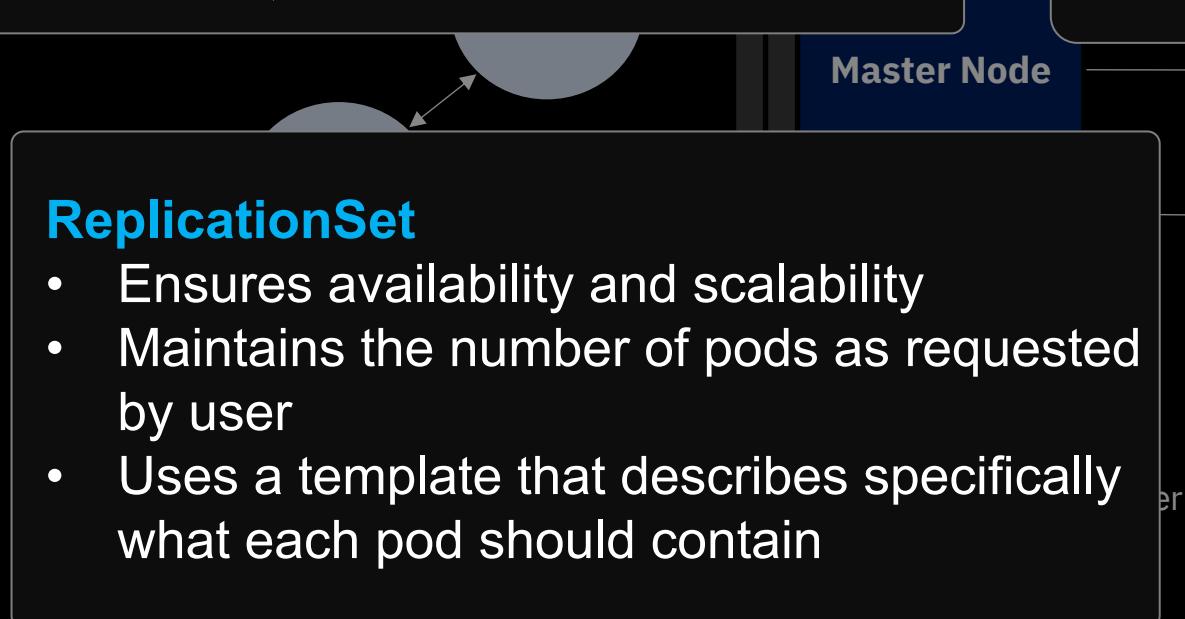


# Kubernetes Management Architecture



## Pods

- Smallest deployment unit in K8s
- Collection of containers that run on a worker node
- Each has its own IP
- Pod shares a PID namespace, network, and hostname



## Deployment

- A set of pods to be deployed together
- Declarative - creates a ReplicaSet describing the desired state
- Rollout/ Rollback - Deployment controller changes the actual state to the desired state
- Scale and autoscale: A Deployment can be scaled

## Service

- Collections of pods exposed as an endpoint
- A service provides a way to refer to a set of Pods (selected by labels) with a single static IP address

# Kubernetes Management Architecture



## ConfigMap

- Configuration values to be used by containers in a pod
- Stores configuration outside of the container image, making containers more reusable

## Labels

- Metadata assigned to K8s resources
- Key-value pairs for identification
- Critical to K8s as it relies on querying the cluster for resources that have certain labels

## Secrets

- Sensitive info that containers need to consume
- Encrypted in special volumes mounted automatically

- Etcd
- API Server
- Controller Manager Server

**API Server** – Kubernetes API server

Worker Node 3

**etcd** - a highly-available key value store which K8s uses for persistent storage of all of its REST API objects

**Scheduler** – schedules pods in worker nodes

# What is container orchestration?



## Container orchestration

- Manages the deployment, placement, and lifecycle of workload containers

## Cluster management

- Federates multiple hosts into one target

## Scheduling

- Distributes containers across nodes

## Service discovery

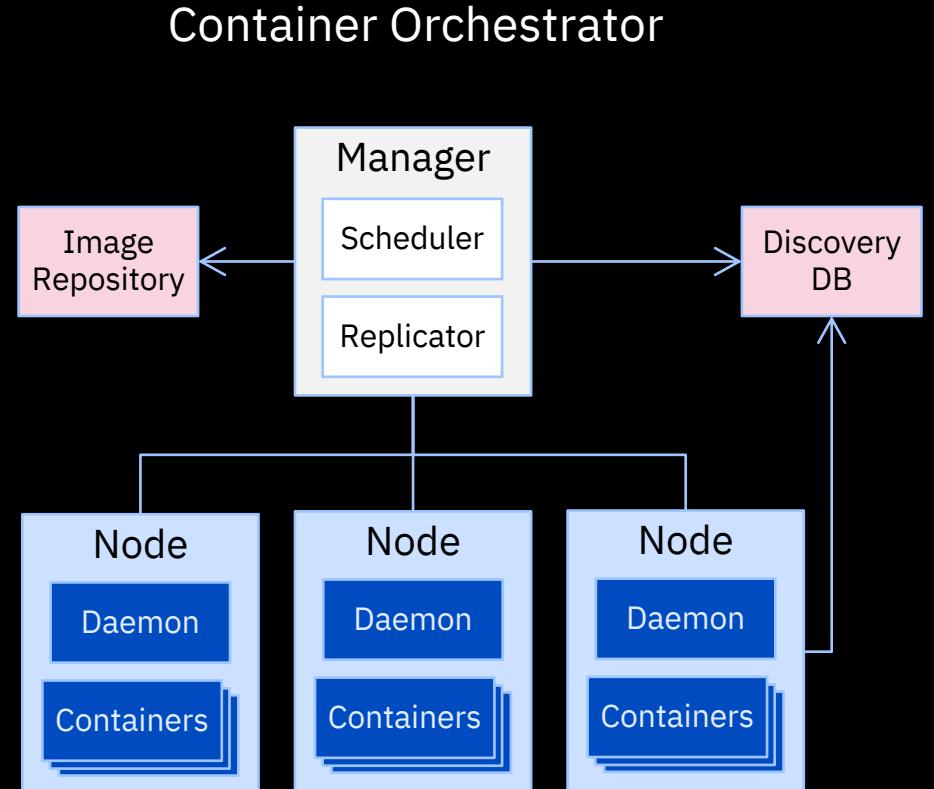
- Knows where the containers are located
- Distributes client requests across the containers

## Replication

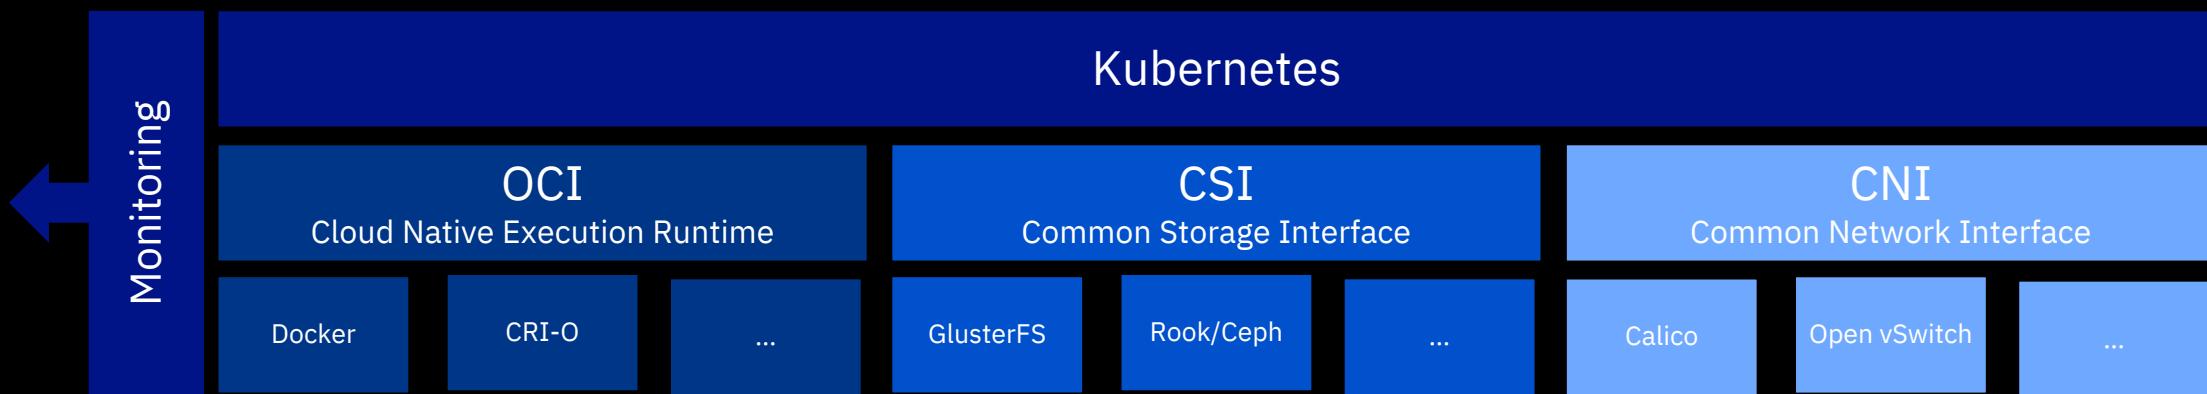
- Ensures the right number of nodes and containers

## Health management

- Replaces unhealthy containers and nodes



# Kubernetes Cluster Architecture



# Kubernetes Cluster Architecture

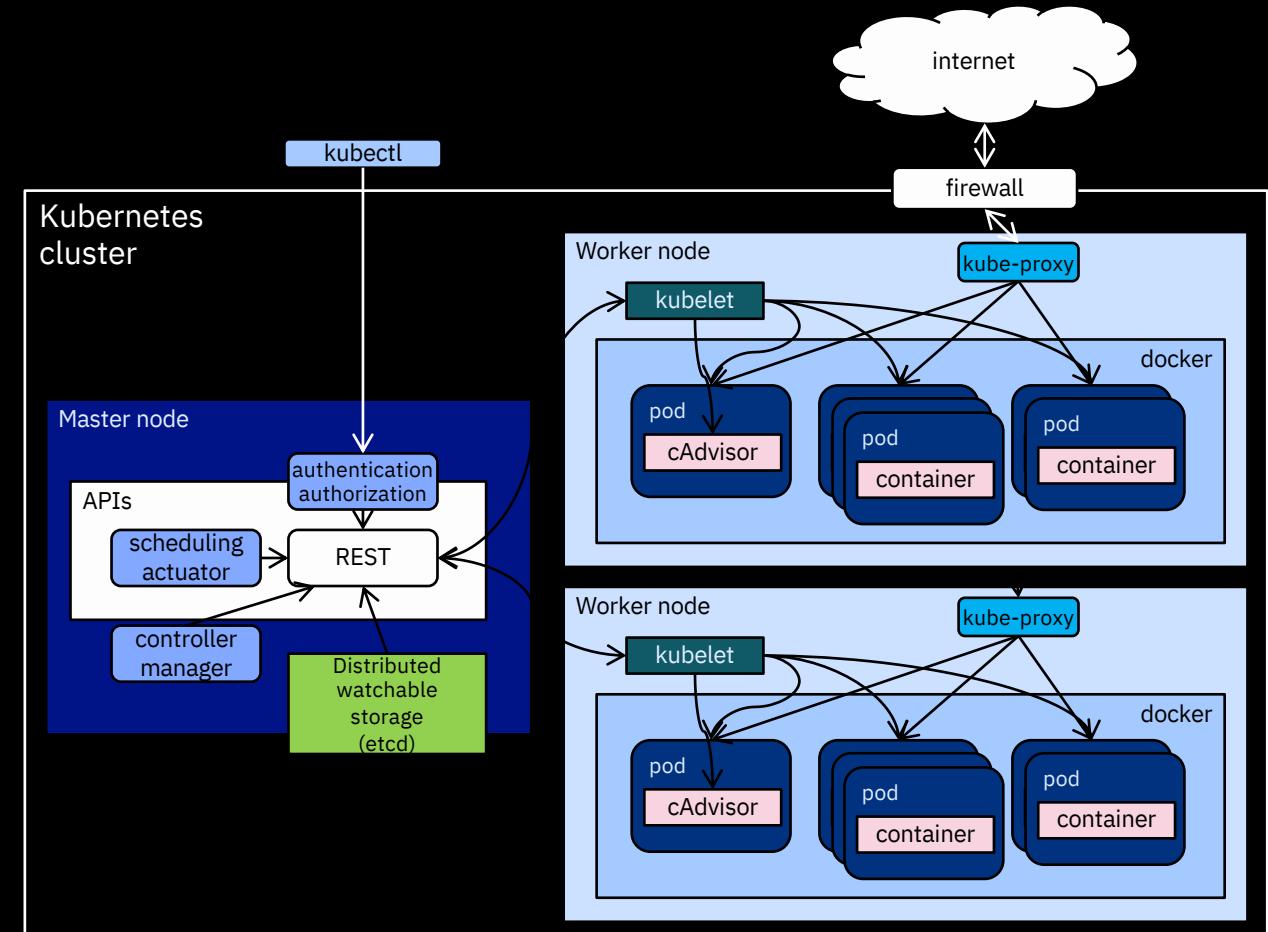


## Master node

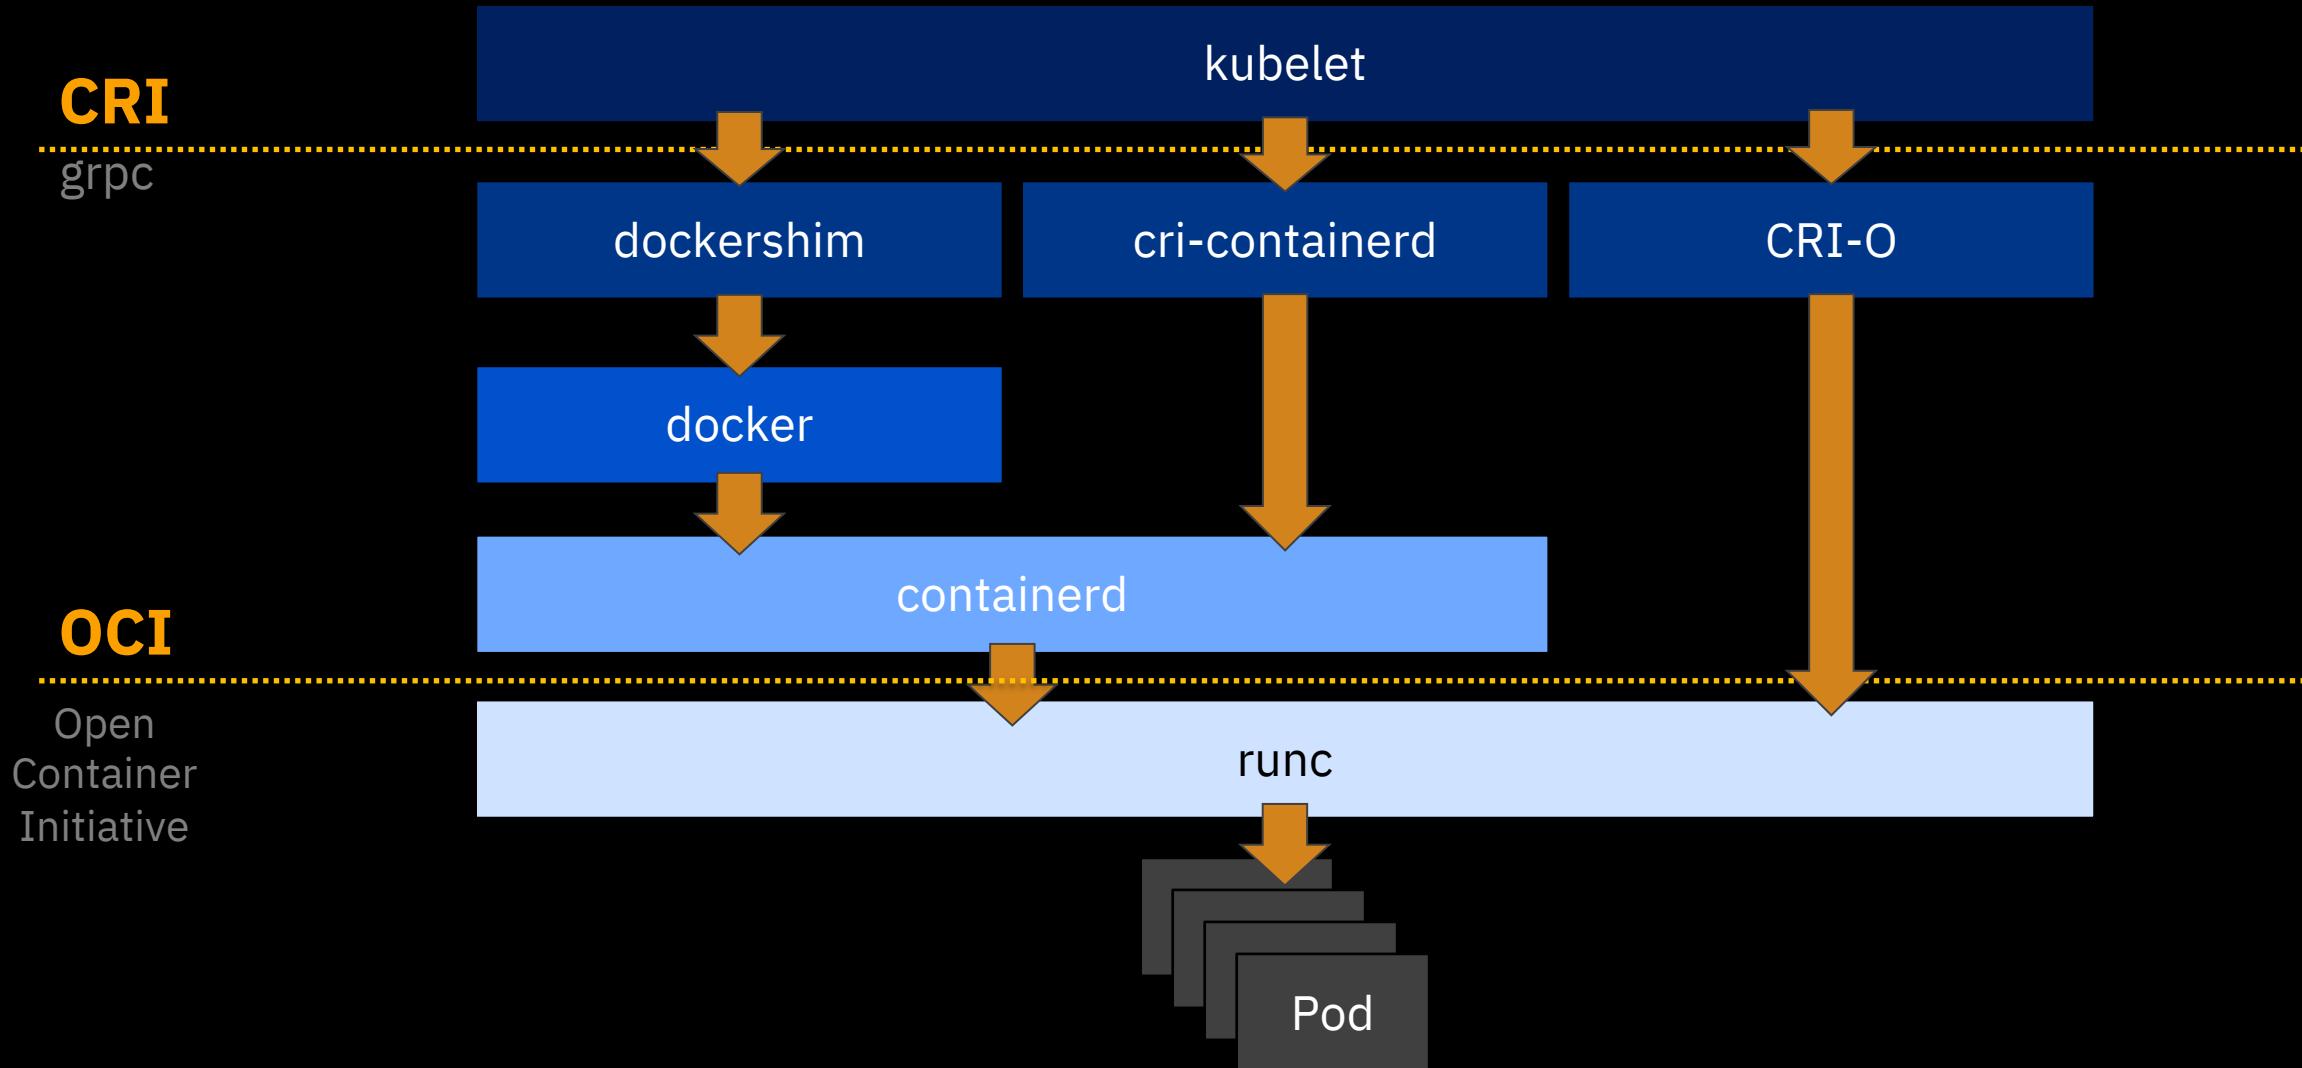
- Node that manages the cluster
- Scheduling, replication & control
- Multiple nodes for HA

## Worker nodes

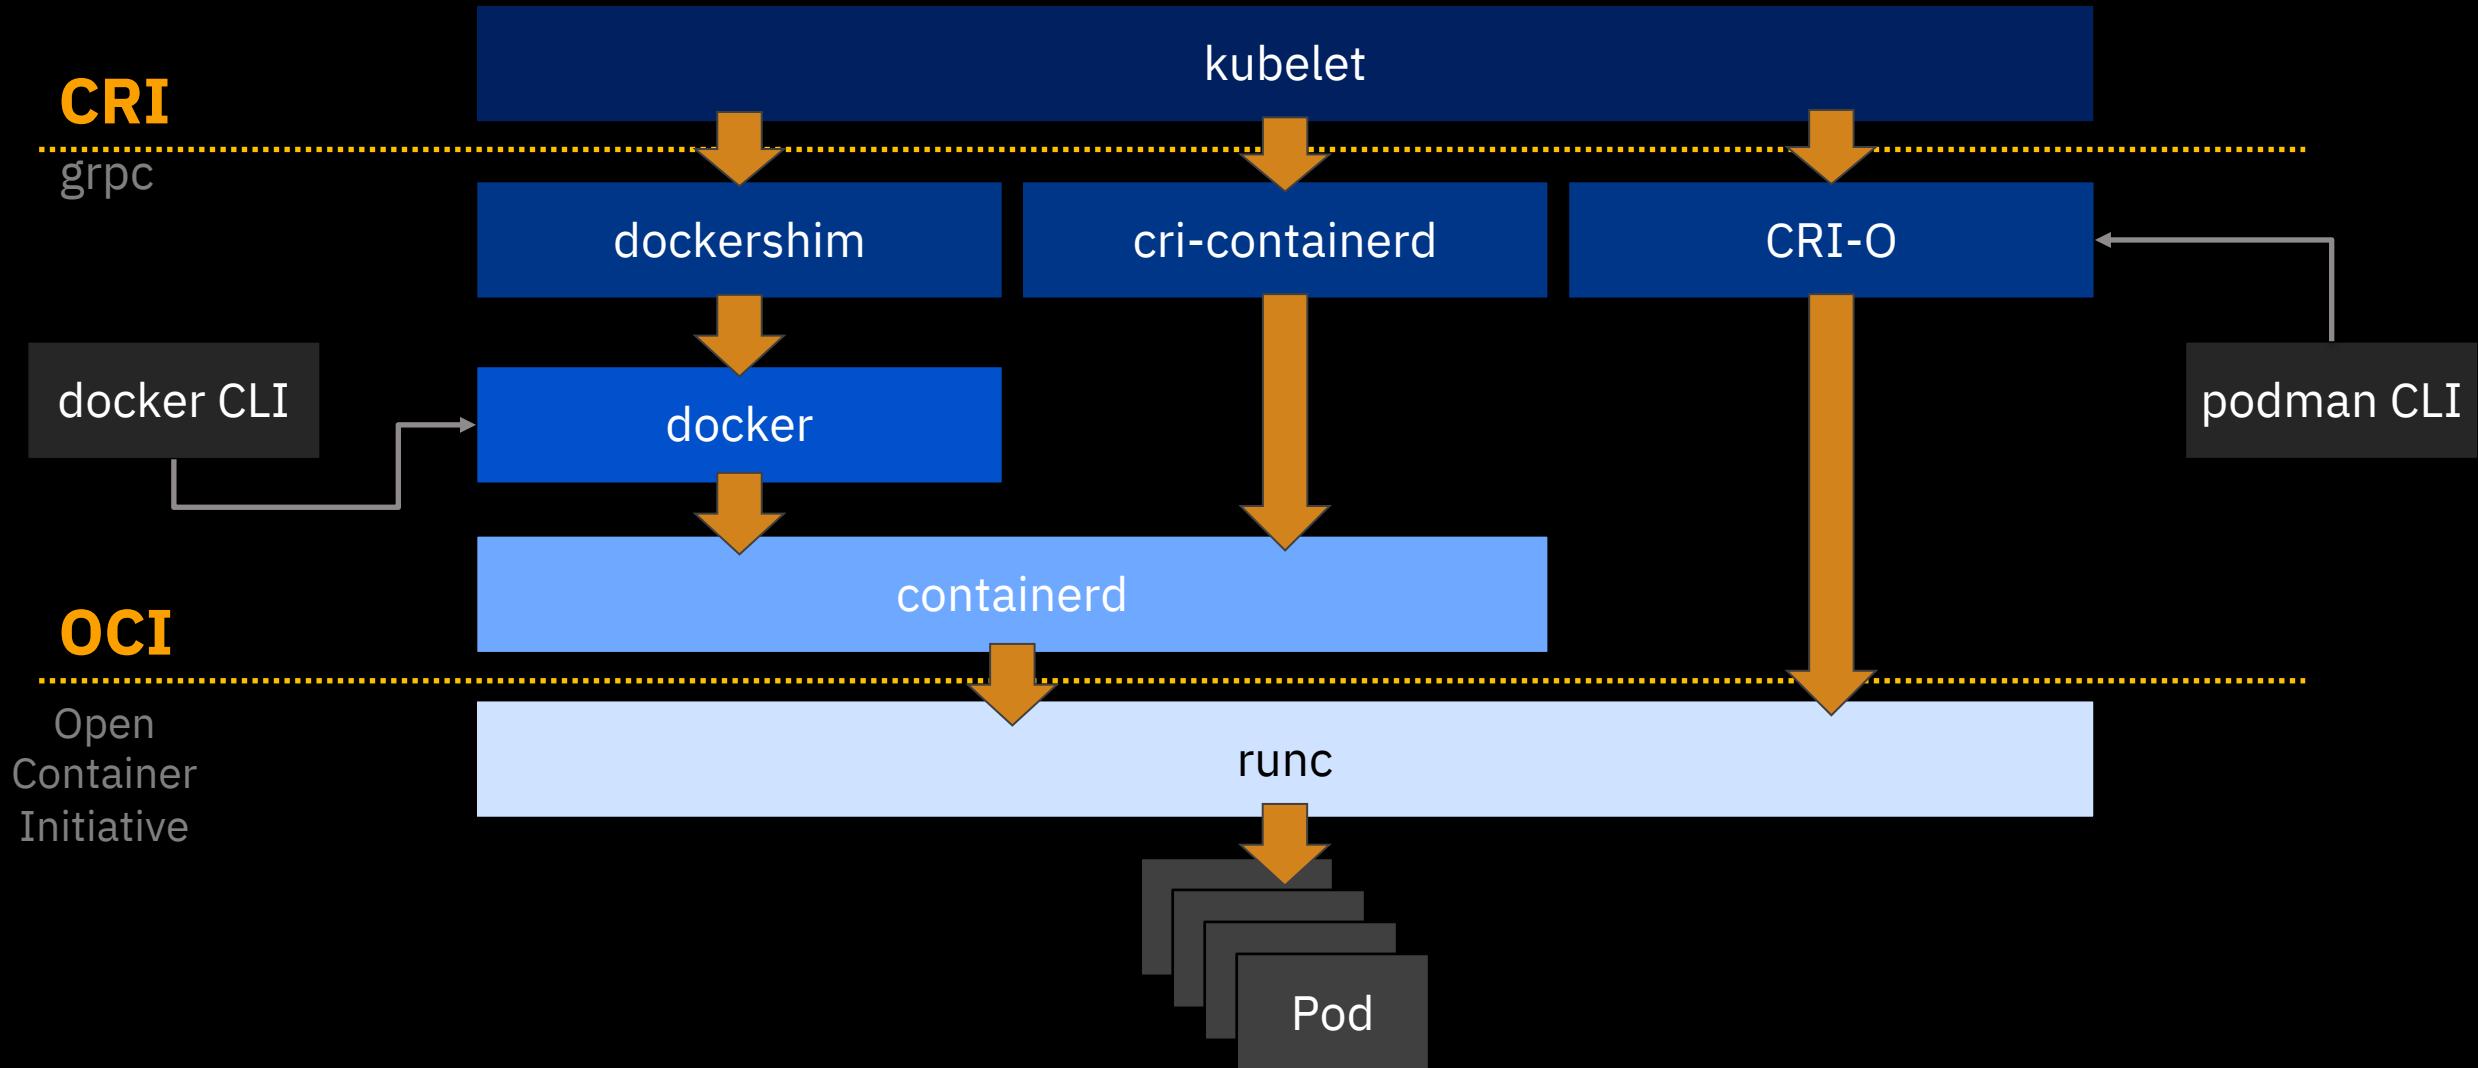
- Node where pods are run
- Docker engine
- kubelet agent accepts & executes commands from the master to manage pods
- cAdvisor – Container Advisor provides resource usage and performance statistics
- kube-proxy – routes inbound or ingress traffic



# Kubernetes – Common Runtime Interface



# Kubernetes – Common Runtime Interface



# Master Node Components



## Etcd

- A highly-available key value store
- All cluster data is stored here

## API Server

- Exposes API for managing Kubernetes
- Used by kubectl CLI

## Controller manager

- Daemon that runs controllers, which are the background threads that handle routine tasks in the cluster
- Node Controller – Responsible for noticing and responding when nodes go down
- Replication Controller – Replaced by ReplicaSet
- Endpoints Controller – Populates the Endpoints object (that is, joins services and pods)
- Service Account & Token Controllers – Create default accounts and API access tokens for new namespaces

## Scheduler

- Selects the worker node each pods runs in

# kubectl – talking to the Cluster



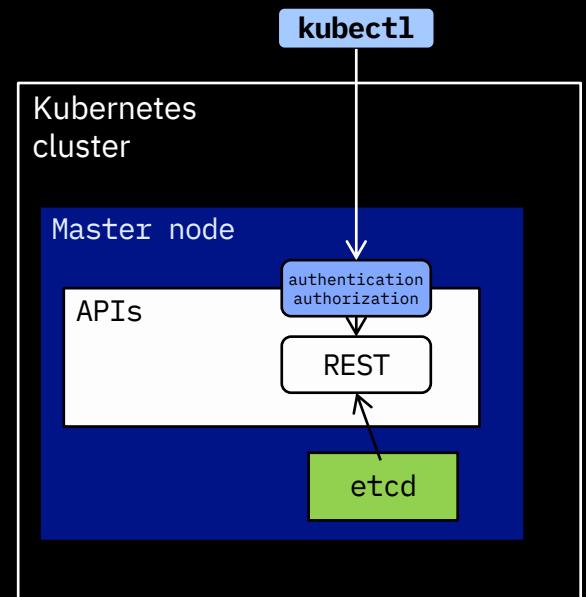
kubectl is a command line interface for running commands against Kubernetes clusters (read state, create objects, ...).

kubectl looks for a file named config in the \$HOME/.kube directory.

Kubernetes uses etcd as a key-value database store.  
It stores the configuration of the Kubernetes cluster in etcd.  
It also stores the *actual* state of the system and the *desired* state of the system in etcd.

Anything you might read from a `kubectl get xyz` command is stored in etcd.

Any change you make via `kubectl create` will cause an entry in etcd to be updated.



# kubectl – talking to the Cluster



kubectl is a command line interface for running commands against Kubernetes clusters.

kubectl looks for a file named config in the \$HOME/.kube directory.

`kubectl [command] [TYPE] [NAME]`

`kubectl create -f example.yaml`

Create objects in yaml file

`kubectl apply -f example.yaml`

Modify objects in yaml file

`kubectl delete -f example.yaml`

Delete objects in yaml file

`kubectl get pods`

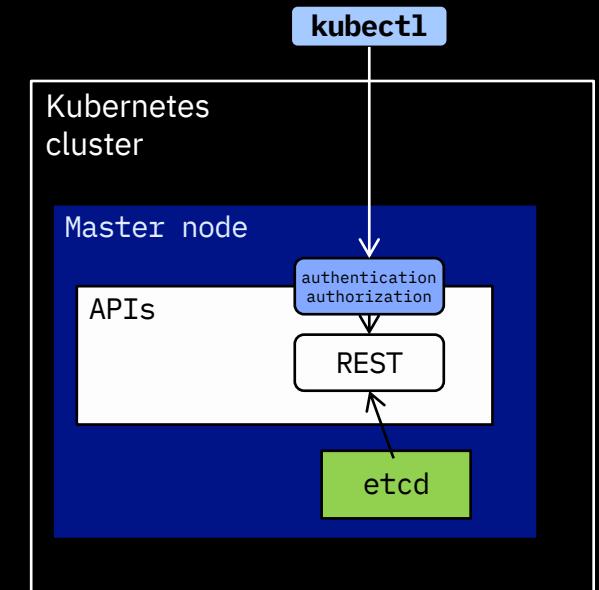
List Pods in Namespace

`kubectl describe nodes <node-name>`

Details about K8s object

`kubectl logs <pod-name>`

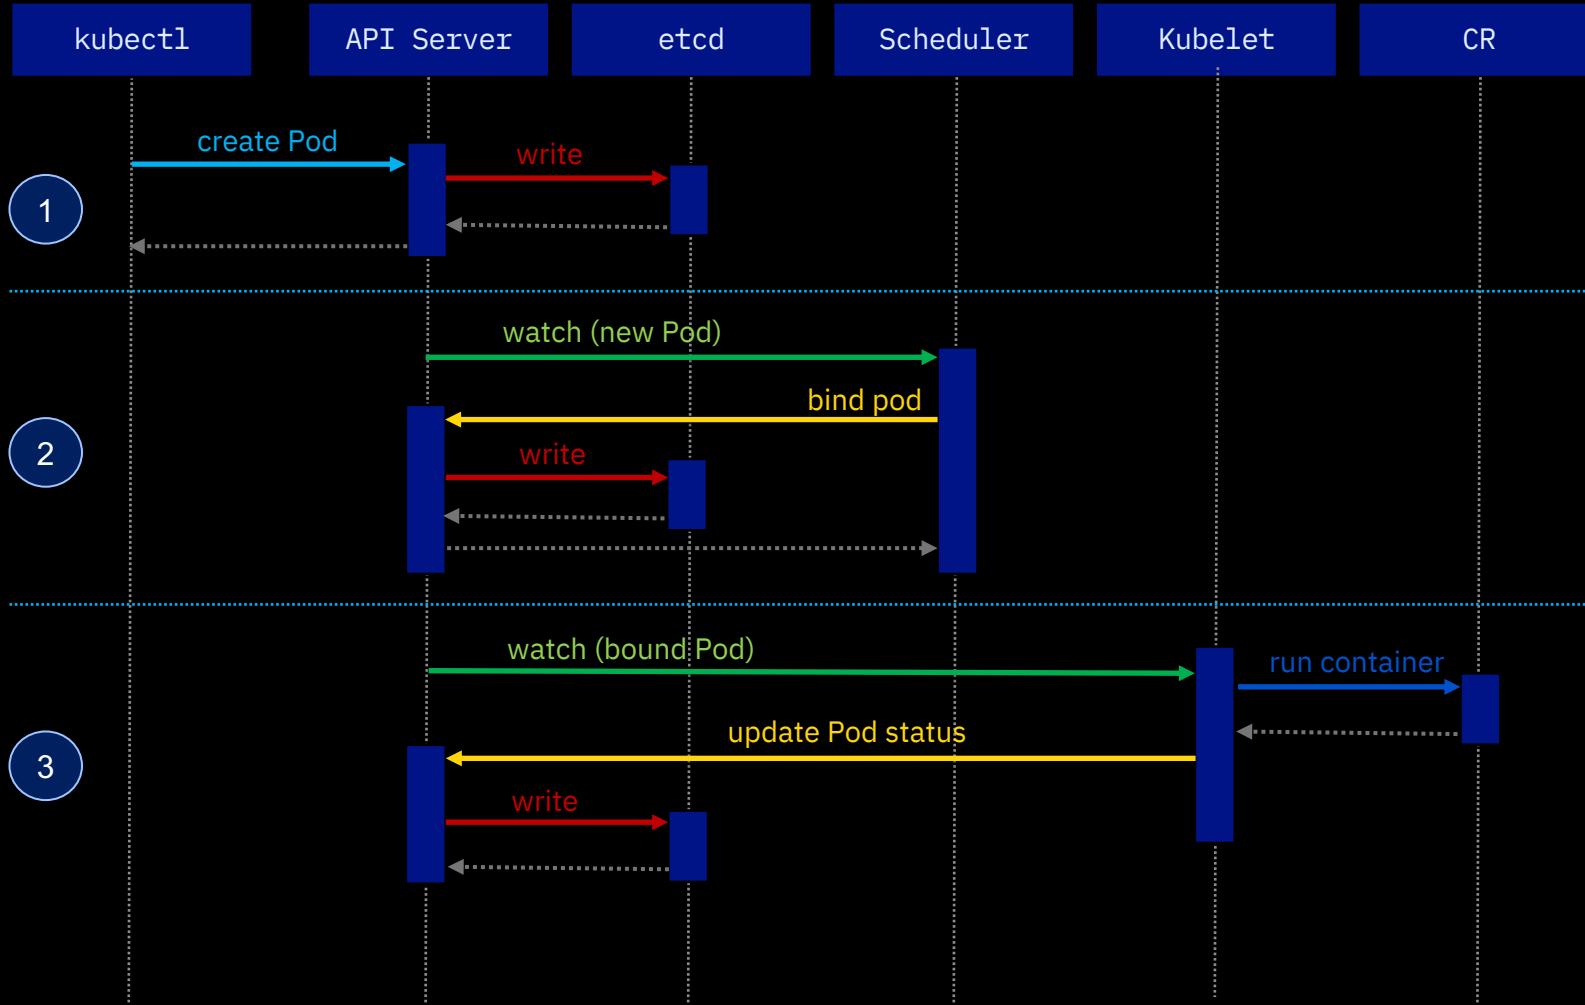
Get the logs for a Pod



# kubectl – talking to the Cluster



What does this actually do: `kubectl run nginx --image= nginx:1.7.9`



1. Create a Pod via the API Server and the API server writes it to etcd
2. The scheduler notices an “unbound” Pod and decides which node to run that Pod on. It writes that binding back to the API Server.
3. The Kubelet notices a change in the set of Pods that are bound to its node. It, in turn, runs the container via the container runtime (i.e. Docker).

The Kubelet monitors the status of the Pod via the container runtime. As things change, the Kubelet will reflect the current status back to the API Server

# Pod

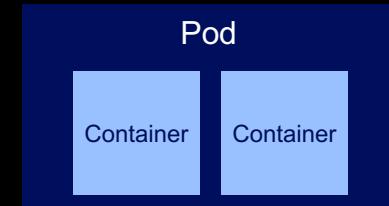


- A group of one or more containers is called a pod.
- Containers in a pod are deployed together, and are started, stopped, and replicated as a group.
- Containers in pod share the same network interface.
- Applications in the same pod
  - Share IP Address and port space
  - Share the same hostname
  - Can communicate using native IPC
  - Can share mounted storage
- Applications in different pods
  - Have different IP Addresses
  - Have different hostnames
  - Pods running on the same node might as well be on different servers
- **When designing pods ask, “Will these containers work correctly if they land on different machines?”**

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.7.9
      ports:
        - containerPort: 80
```



10.0.0.1



10.0.0.2

# Creating a Pod



```
kubectl run nginx --image=nginx:1.7.9  
or  
kubectl create -f example.yaml
```

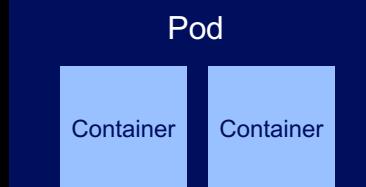
```
kubectl get pods  
NAME          READY  STATUS    RESTARTS  AGE  
nginx-5bd87f76c-vxc79  1/1    Running   0          29s
```

example.yaml

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: nginx  
spec:  
  containers:  
  - name: nginx  
    image: nginx:1.7.9  
  ports:  
  - containerPort: 80
```



10.0.0.1

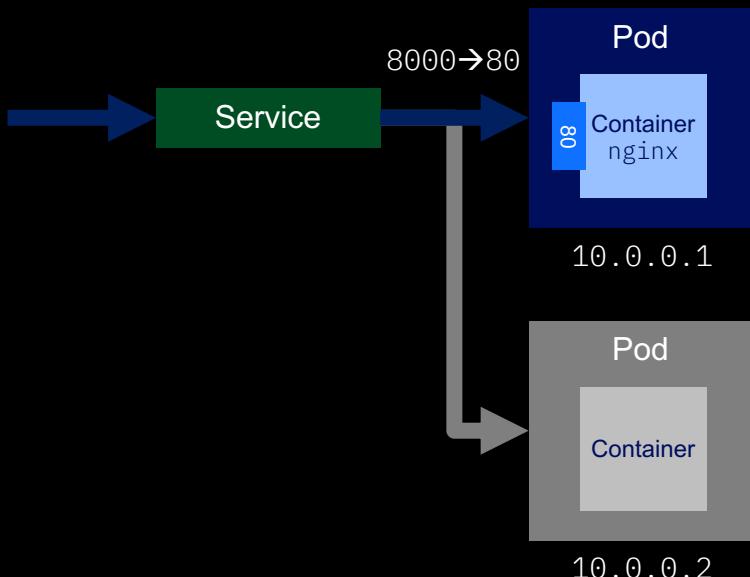


10.0.0.2

# Service



- A service provides a way to refer to a set of Pods (selected by labels) with a single static IP address.
- Also provide load balancing

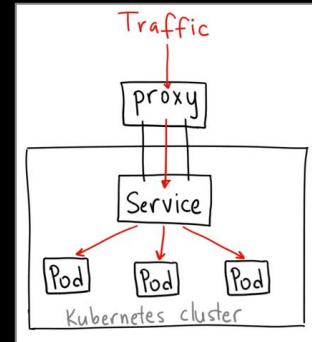


```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  ports:
    - port: 8000
      targetPort: 80
      protocol: TCP
  selector:
    app: nginx
```

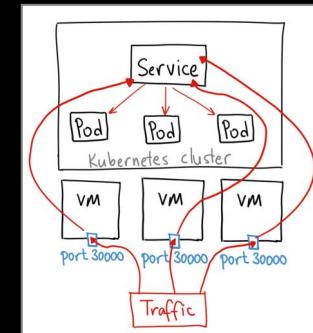
# Service



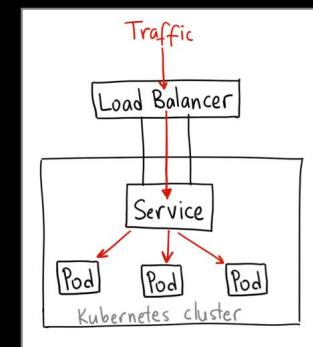
**ClusterIP:** This type exposes the service on the cluster internal IP. This means that the service is only reachable from within the cluster.



**NodePort:** This type exposes the service on each Node's static IP address.



**LoadBalancer:** This service type exposes a service using the cloud provider load balancer.



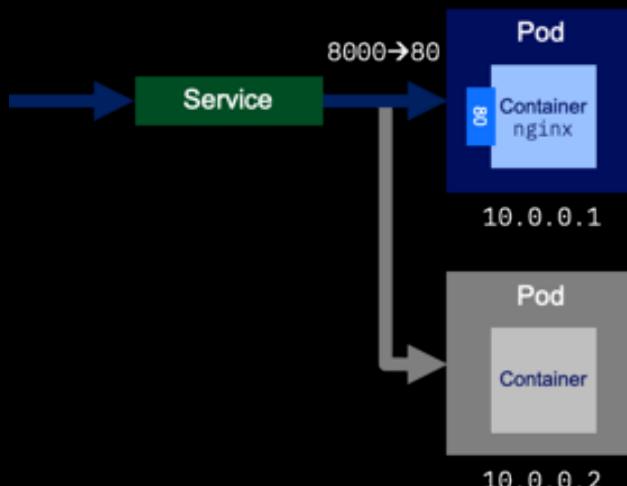
# Creating a Service



```
kubectl expose deployment nginx --type="NodePort" --port=8000 --target-port=80  
or  
kubectl create -f example.yaml
```

```
kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
nginx	NodePort	10.106.115.135	<none>	8000:32499/TCP	6s



example.yaml

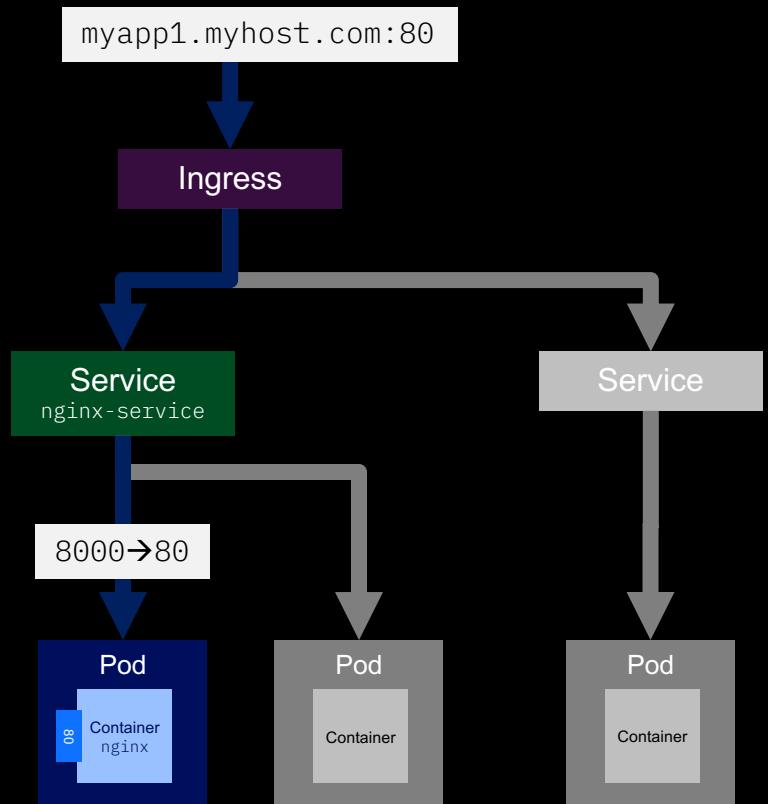
```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  ports:
    - port: 8000
      targetPort: 80
      protocol: TCP
  selector:
    app: nginx
```

# Ingress



An ingress can be configured to give services externally-reachable URLs, load balance traffic, terminate SSL, and offer name based virtual hosting. An ingress controller is responsible for fulfilling the ingress, usually with a loadbalancer.

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: nginx-ingress
  annotations:
    ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - host: myapp1.myhost.com
    http:
      paths:
      - backend:
          serviceName: nginx-service
          servicePort: 80
```

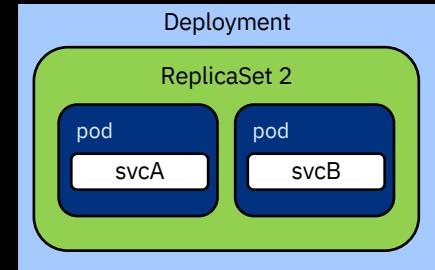


# Deployments & ReplicaSets



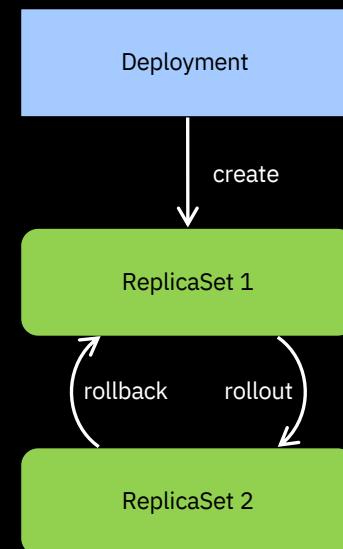
## Deployment

- A set of **pods to be deployed together**, such as an application
- **Declarative**: Revising a Deployment **creates a ReplicaSet** describing the desired state
- **Rollout**: Deployment controller changes the actual state to the desired state at a controlled rate
- **Rollback**: Each Deployment revision can be rolled back
- **Scale** and autoscale: A Deployment can be scaled



## ReplicaSet

- Cluster-wide pod manager that **ensures the proper number of pods are running** at all times.
- A set of pod templates that describe a set of pod replicas
- Uses a template that describes specifically what each pod should contain
- Ensures that a specified number of pod replicas are running at any given time



# Deployment



- A Deployment object defines a Pod creation template and **desired replica count**.
- Create or delete Pods as needed to meet the replica count.
- Manage safely **rolling out changes** to your running Pods.

```
kubectl create -f example.yaml
```

example.yaml

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 2
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.7.9
      ports:
      - containerPort: 80
```

# StatefulSet, DaemonSet, Job...



## **StatefulSet**

- Intended to be used with stateful applications and distributed systems.
  - Pods are created sequentially
  - Ordinal index and stable network identity

## **DaemonSet**

- Ensures that all (or some) nodes run a copy of a pod
  - running a cluster storage daemon
  - running a logs collection daemon
  - running a node monitoring daemon

## **Job**

- Creates one or more pods and ensures that a specified number of them successfully terminate

## **Cron Job**

- Manage time based jobs, once at a specified in time, or repeatedly at a specified time point

# Naming



## Name

- Each resource object by type has a unique name

## Namespace

- **Resource isolation:** Each namespace is a virtual cluster within the physical cluster
  - Resource objects are scoped within namespaces
  - Low-level resources are not in namespaces: nodes, persistent volumes, and namespaces themselves
  - Names of resources need to be unique within a namespace, but not across namespaces
- **Resource quotas:** Namespaces can divide cluster resources
- Initial namespaces
  - **default** – The default namespace for objects with no other namespace
  - **kube-system** – The namespace for objects created by the Kubernetes system

## Resource Quota

- Limits resource consumption per namespace
- Limit can be number of resource objects by type (pods, services, etc.)
- Limit can be total amount of compute resources (CPU, memory, etc.)
- Overcommit is allowed; contention is handled on a first-come, first-served basis

# Kubernetes User Security



- **Authentication**
  - OIDC Tokens
  - ServiceAccount Tokens
- **Authorization**
  - RBAC
  - Role
  - RoleBinding
  - ClusterRole
  - ClusterRoleBinding
- **Namespaces**
  - Help to scope access control to a set of resources in a Kubernetes cluster

```
---  
apiVersion: rbac.authorization.k8s.io/v1beta1  
kind: ClusterRole  
metadata:  
  name: viewer  
rules:  
- apiGroups:  
  - apps  
  - extensions  
resources:  
- deployments  
- replicaset  
verbs:  
- get  
- list  
- watch  
---  
apiVersion: rbac.authorization.k8s.io/v1  
kind: ClusterRoleBinding  
metadata:  
  name: cicd-viewer  
roleRef:  
  apiGroup: rbac.authorization.k8s.io  
  kind: ClusterRole  
  name: viewer  
subjects:  
- kind: User  
  name: IAM#cicd@us.ibm.com  
  apiGroup: rbac.authorization.k8s.io
```

# Configuring Resources and Containers



## Label

- **Metadata** assigned to Kubernetes resources (pods, services, etc.)
- Key-value pairs for identification
- Critical to Kubernetes as it relies on querying the cluster for resources that have certain labels

## Selector

- An expression that **matches labels** to identify related resources

Deployment details	
Type	Detail
Name	cam-ui-basic
Namespace	services
Created	39 minutes ago
Labels	app=cam-ibm-cam,chart=ibm-cam-1.3.0,heritage=Tiller,name=cam-ui-basic,release=cam
Selector	name=cam-ui-basic
Replicas	Desired: 1   Total: 1   Updated: 1   Available: 1
RollingUpdateStrategy	Max unavailable: 1   Max surge: 1
MinReadySeconds	0

Service details	
Type	Detail
Name	cam-ui-basic
Namespace	services
Created	40 minutes ago
Type	ClusterIP
Labels	app=cam-ibm-cam,chart=ibm-cam-1.3.0,heritage=Tiller,name=cam-ui-basic,release=cam
Selector	name=cam-ui-basic
Cluster IP	10.0.0.165
External IP	-
Port	cam-ui-basic 39002/TCP
Node port	None
Session affinity	None

# Configuring Resources and Containers



## **ConfigMap**

- Configuration values to be used by containers in a pod
- Stores configuration outside of the container image, making containers more reusable

## **Secret**

- Sensitive info that containers need to read or consume
- Encrypted in special volumes mounted automatically

# Configuring Resources and Containers



```
apiVersion: extensions/v1beta1
kind: ConfigMap
apiVersion: v1
metadata:
  name: example-config
data:
  allowed: '"true"'
  enemies: aliens
  lives: "3"
```

```
kind: Deployment
spec:
  containers:
    - name: test-container
      image: xxx
      ...
  env:
    - name: SPECIAL_LEVEL_KEY
      valueFrom:
        configMapKeyRef:
          name: example-config
          key: enemies
```

Can be used as normal environment variable

# Kubernetes Autoscaling



## Horizontal Pod Autoscaling (HPA)

- Automatically scales the number of pods in a replication controller, deployment, or replica set
- Matches the observed average CPU utilization to the specified target
- Fetches metrics in two different ways: direct Heapster access and REST client access
- Kubernetes Heapster enables container cluster monitoring and performance analysis
- Default config: query every 30 sec, maintain 10% tolerance, wait 3 min after scale-up, wait 5 min after scale-down

```
$ kubectl autoscale deployment <deployment-name> --cpu-percent=50  
--min=1 --max=10 deployment "<hpa-name>" autoscaled
```

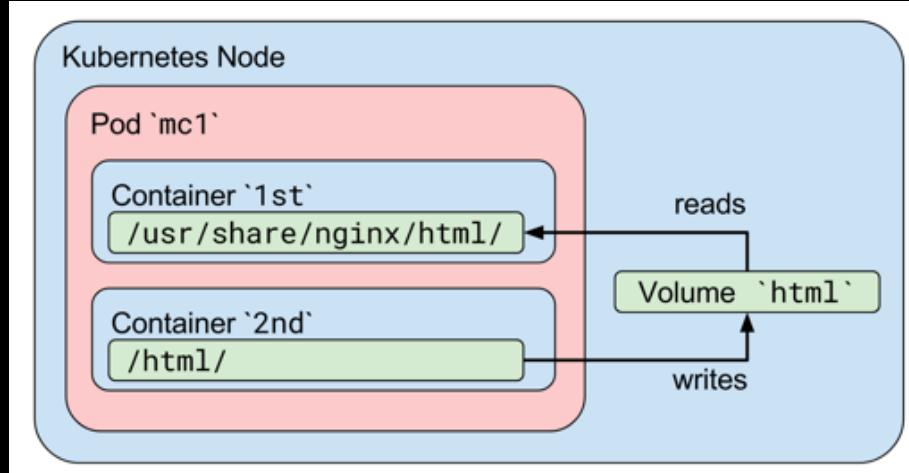
Creates a horizontal pod autoscaler

- An HPA instance
- Maintains between 1 and 10 replicas of the pods controlled by the deployment
- Maintains an average CPU utilization across all pods of 50%

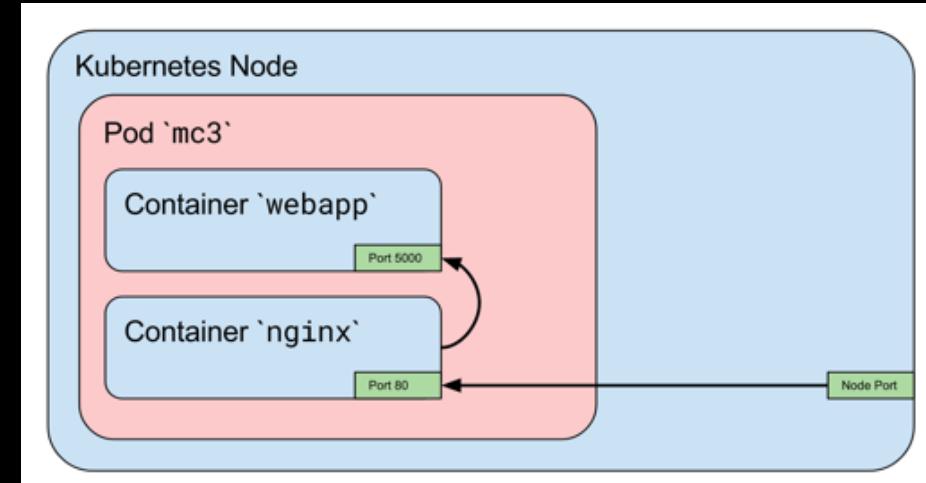
# Multi-Container Pod Design Patterns



## Shared Volumes

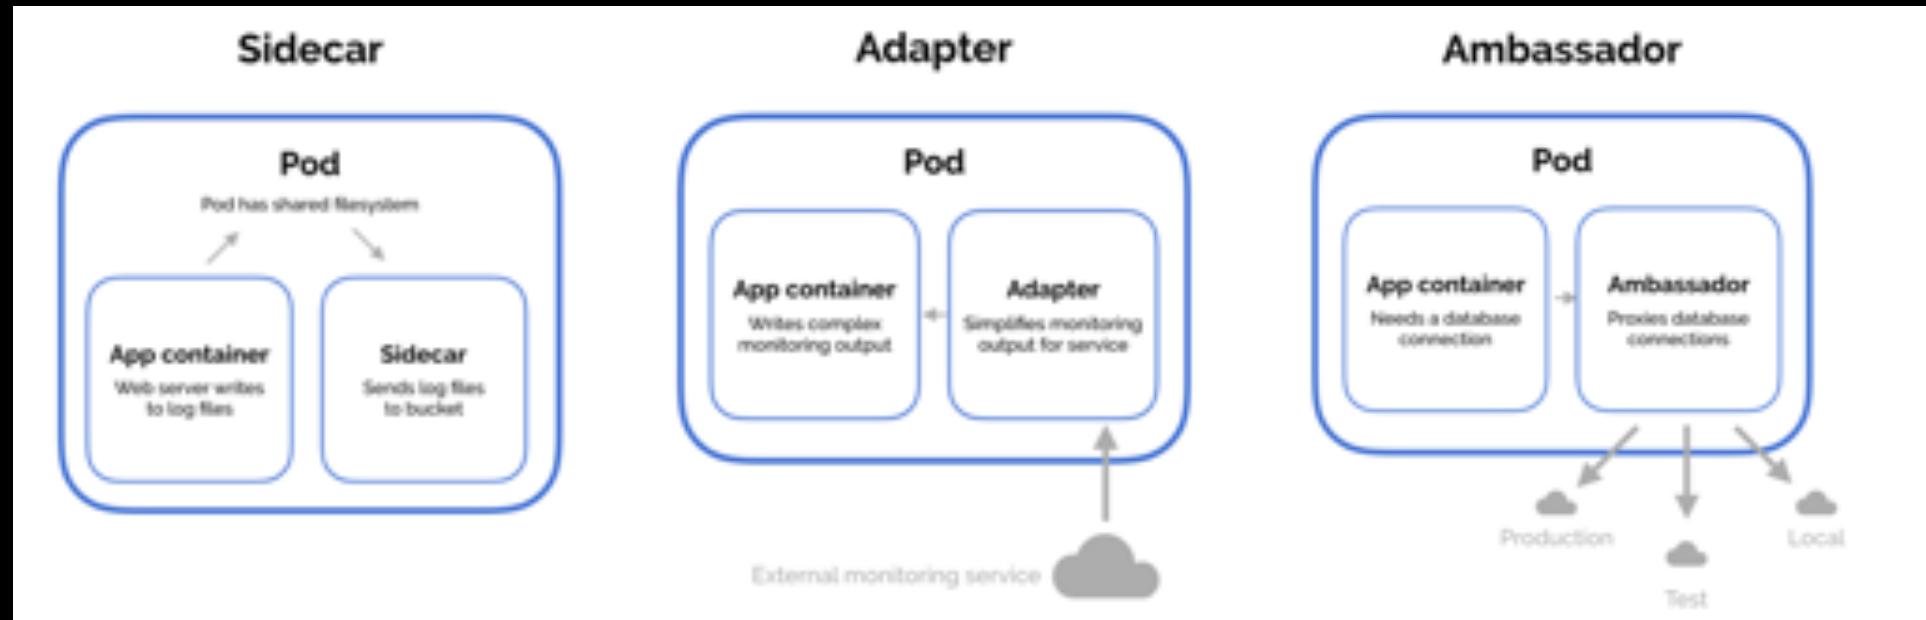


## No Network isolation



**In what order containers are being started in a Pod?**

# Multi-Container Pod Design Patterns



## Sidecar pattern

The sidecar pattern consists of a main application—i.e. your web application—plus a helper container with a **responsibility that is essential to your application**, but is not necessarily part of the application itself.

## Adapter pattern

The adapter pattern is used to standardize and normalize application output or **monitoring data** for aggregation.

## Ambassador pattern

The ambassador pattern is a useful way to connect containers with the outside world (**proxy**).

# HELM Charts



- Helm is the open standard for **Application Packaging and Deployment** for Kubernetes (like npm, yum or apt-get)
- Helm charts **automate the deployment of resources** and prerequisites including locations of Docker images



Catalog

All Categories > Search items Filter

Category	Chart Name	Description	Status
Blockchain	acs-engine-autoscaler	DEPRECATED Scales worker nodes within agent pools	KUBE Deprecated
Business Automation	aerospike	A metric chart for Aerospike in Kubernetes	KUBE
Data	airflow	Airflow is a platform to programmatically author, schedule and monitor workflows	KUBE
Data Science & Analytics	ark	A Helm chart for ark	KUBE
DevOps	artifactory	DEPRECATED Universal Repository Manager supporting all major packaging formats, build tools	KUBE Deprecated
Integration	artifactory-ha	DEPRECATED Universal Repository Manager supporting all major packaging formats, build tools	KUBE Deprecated
IoT	artifactory-ha	Universal Repository Manager supporting all major packaging formats, build tools and CI serv...	KUBE Deprecated
Network	audit-logging	Audit logging storage and search management solution	Ingresscharts
Operations	auditbeat	A lightweight shipper to audit the activities of users and processes on your systems	KUBE
Runtimes & Frameworks	auth-apis	SCP IAM Token Service	Ingresscharts
Security	auth-idp	SCP Security Authentication Provider	Ingresscharts
Storage	auth-pap	SCP IAM Policy Administration	Ingresscharts
Tools	auth-pdp	SCP IAM Policy Decision	Ingresscharts
Other	aws-cluster-autoscaler	Scales worker nodes within auto-scaling groups	KUBE Deprecated
	boltbond	BoltBond is an innovative payment network and a new kind of money	KUBE
	bookstack	BookStack is a simple, self-hosted, easy-to-use platform for organizing and storing informa...	KUBE

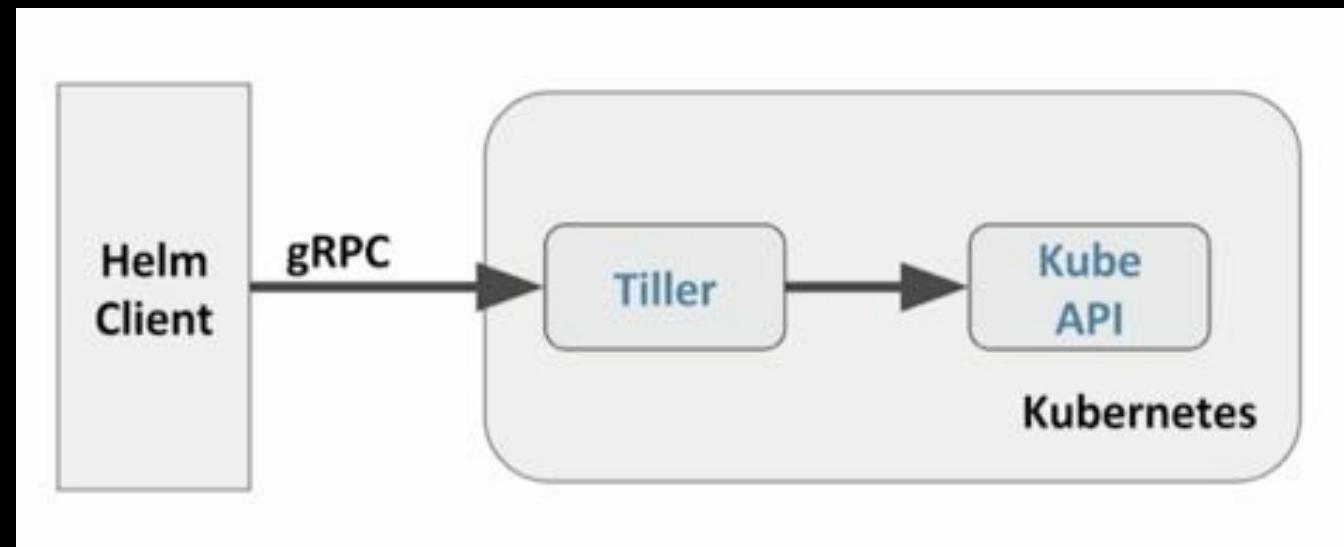
# HELM Example



**Helm has 2 Components :**

**Tiller** : Tiller is the in-cluster component of Helm. It interacts directly with the Kubernetes API server to install, upgrade, query, and remove Kubernetes resources. It also stores the objects that represent releases.

**Client** : The Helm client interacts with the Tiller and stores information in a local directory.

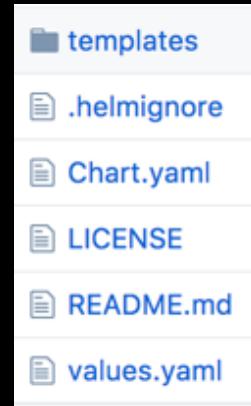


# HELM Example



- **Chart :**

A Helm package that contains information sufficient for installing a set of Kubernetes resources into a Kubernetes cluster.



- **Chart.yaml**

Basic information about the chart (name, versions, links, icon, ...)

- **Templates (yaml)**

Kubernetes deployment manifests (yaml) with placeholders for external parameters.

- **values.yaml**

A set of variables that will show up in the UI (or customized at the helm command line) and be passed on at deploy time to the deployment manifests.

Deployment Chart

```
1 apiVersion: v1
2 kind: PersistentVolumeClaim
3 metadata:
4   name: {{ template "fullname" . }}-pv-claim
5 spec:
6   accessModes:
7     - ReadWriteMany
8   resources:
9     requests:
10       storage: {{ .Values.storage.size }}
```

Values.yaml file

```
storage:
  size: 2Gi
```



## IBM Helm Charts

<https://github.com/ibm/charts>

A screenshot of the GitHub repository page for "IBM / charts". The page shows the repository's landing page with the title "IBM / charts", navigation links for "Code", "Pull requests 0", "Projects 0", "ZenHub", "Wiki", and "Insights", and a brief description: "The IBM/charts repository provides helm charts for IBM and Third Party middleware." Below the description are several tags: "kubernetes", "helm", "helm-charts", "ibm", "ibm-bluemix", and "docker".

The IBM/charts repository provides helm charts for IBM and Third Party middleware.

kubernetes helm helm-charts ibm ibm-bluemix docker

# Kubectl Commands



Get the state of your cluster

```
$ kubectl cluster-info
```

Get all the nodes of your cluster

```
$ kubectl get nodes -o wide
```

Get info about the pods of your cluster

```
$ kubectl get pods -o wide
```

Get info about the replication controllers of your cluster

```
$ kubectl get rc -o wide
```

Get info about the services of your cluster

```
$ kubectl get services
```

Get full config info about a Service

```
$ kubectl get service NAME_OF_SERVICE -o json
```

Get the IP of a Pod

```
$ kubectl get pod NAME_OF_POD -template={{.status.podIP}}
```

Delete a Pod

```
$ kubectl delete pod NAME
```

Delete a Service

```
$ kubectl delete service NAME_OF_SERVICE
```

# Resources



## Linux Foundation – Introduction to Kubernetes

- <https://courses.edx.org/courses/course-v1:LinuxFoundationX+LFS158x+2T2017/course/>

## Kubernetes tutorial

- <https://kubernetes.io/docs/tutorials/kubernetes-basics/>

## Introduction to container orchestration

- <https://www.exoscale.ch/syslog/2016/07/26/container-orch/>

## TNS Research: The Present State of Container Orchestration

- <https://thenewstack.io/tns-research-present-state-container-orchestration/>

## Large-scale cluster management at Google with Borg

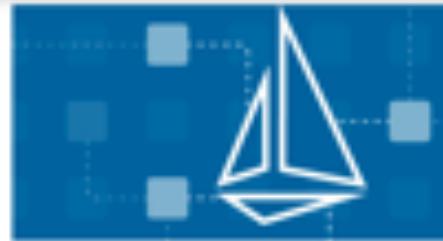
- <https://research.google.com/pubs/pub43438.html>

# [Cognitiveclass.ai](#) – Free basic to advanced classes on AI, Cloud... etc

 Learning Paths Courses Mobile Apps Badges Business Competitions  Explore new learning opportunities Login



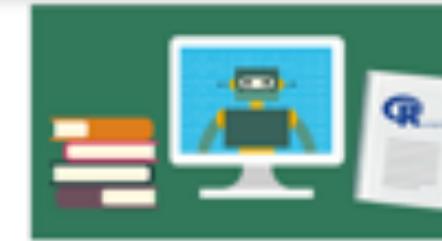
[How to Build Watson AI & Swift APIs and Make Money](#)  
Cognitive Class: SW0201EN  
 Intermediate



[Beyond the Basics: Istio and IBM Cloud Kubernetes Service](#)  
Cognitive Class: CO0401EN  
 Advanced



[Deep Learning with TensorFlow](#)  
Cognitive Class: MLD120ENv2  
 Advanced



[Machine Learning with R](#)  
Cognitive Class: ML0151EN  
 Intermediate



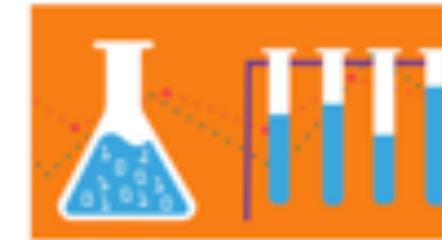
[Data Analysis with Python](#)  
Cognitive Class: DA0101EN  
 Beginner



[Data Visualization with R](#)  
Cognitive Class: DV0151EN  
 Beginner



[Spark Fundamentals I](#)  
Cognitive Class: BD0211EN  
 Beginner



[Data Science Methodology](#)  
Cognitive Class: DS0103EN  
 Beginner

# Benefits of using a Kubernetes Based Platform

## Speed

- Fast boot-time
- Easy scalability
- Rapid deployment

## Portability

- Between different environments
- Between private and public cloud

## Efficiency

- Better usage of computer resources

## Automation

- Next level standardization and automation

**13x**

More software releases

**Eliminate**

“works on my machine” syndrome

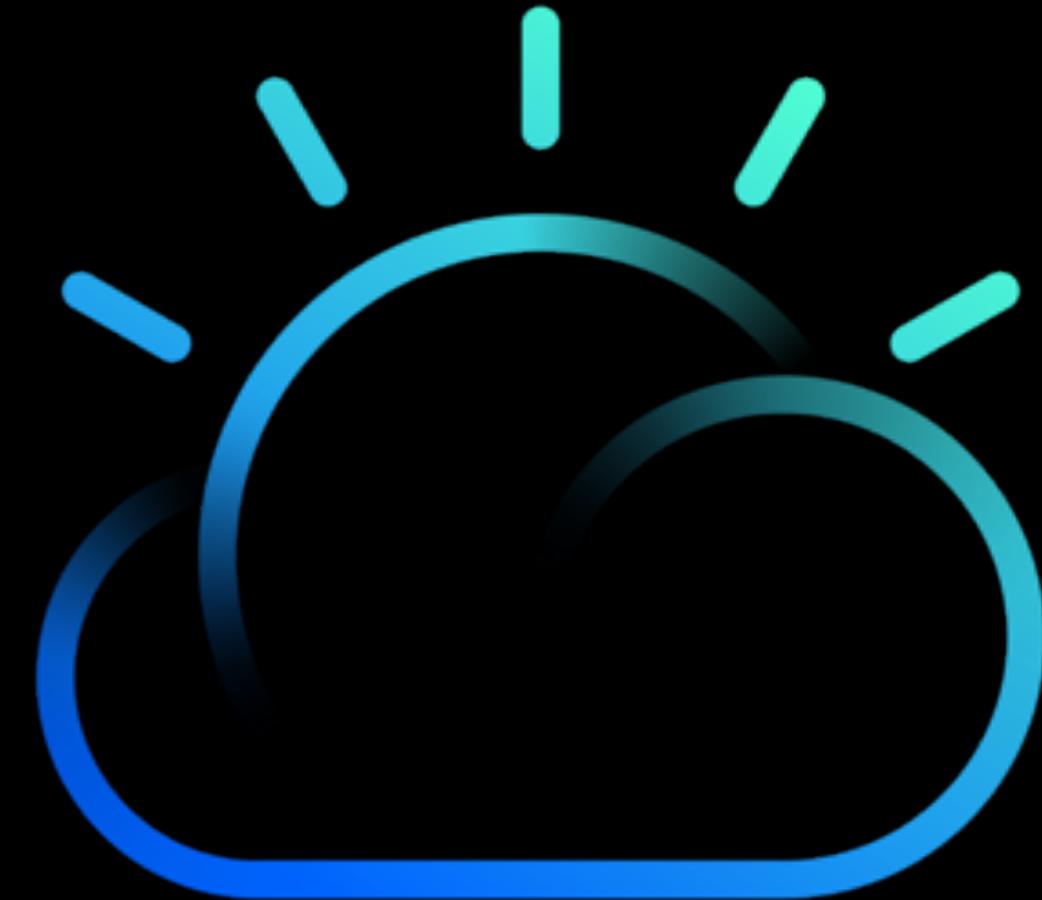
**~47%**

Reduction of VMs,  
OS licensing and  
server cost

**Reduce**

operation cost

QUESTIONS?



The Journey to Cloud  
**Let's get real**

04



IBM Cloud

# Microservice transformations – The good, the bad, the ugly

## The good

**IF** implemented correctly Microservices improve scalability and reliability.

- Scale easily and very efficiently
- No single point of failure by isolation
- Be modified without compromising the whole system
- Use of languages and frameworks that best suit the purpose of each service
- Natively ready for CaaS and PaaS

# Microservice transformations – The good, the bad, the ugly

## The bad

- Added management **complexity**
- **Expertise** required to maintain a microservice-based application
- No or **reduced benefits** if application doesn't need scaling or cloud transformation
- Needs more **robust testing** in order to cover all APIs

# Microservice transformations – The good, the bad, the ugly

## ... and **the downright ugly**

- May require **high initial investment** to run for Dev, deployment and Ops overhead
- Higher **initial operational and monitoring costs** through shift in paradigms
- **Non-uniform** application design and architecture through free choice of technologies
- **Overload of documentation** for every individual component app

Microservices, when  
implemented incorrectly,  
can make poorly written  
applications even  
more dysfunctional

# Microservices – breaking down the monolith beast

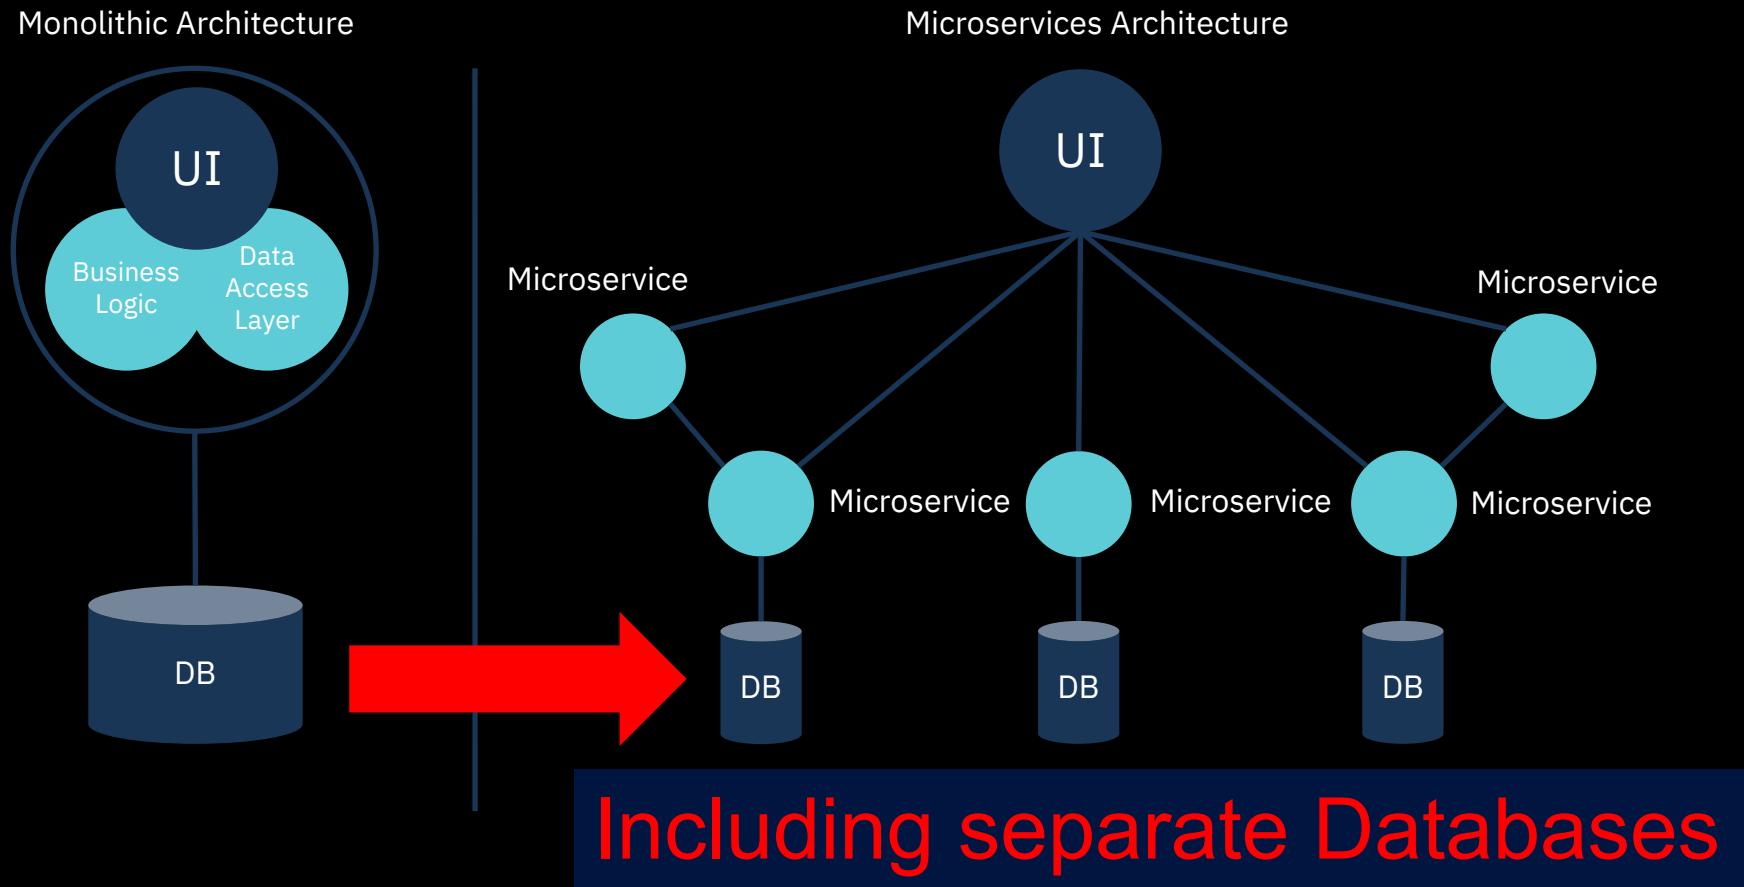
## How to identify candidates

- Is there anything that is scaling at a differently than the rest of the system?
- Is there anything that feels “tacked-on”?
- Is there anything changing much faster than the rest of the system?
- Is there anything requiring more frequent deployments than the rest of the system?
- Is there a part of the system that a small team, operates independently?
- Is there a subset of tables in your datastore that isn’t connected to the rest of the system?

# Databases

# Microservices – Database refactoring

Simplistically, microservices architecture is about breaking down large silo applications into more manageable, fully decoupled pieces



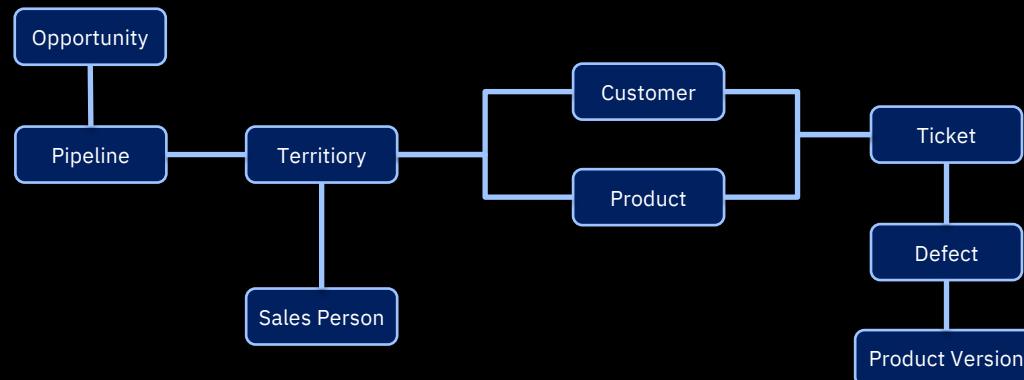
Containerizing your WAR file  
doesn't mean you're doing  
microservices.

- What is the domain? What is reality?
- Where are the transactional boundaries?
- How should microservices communicate across boundaries?
- What if we just turn the database inside out?

# Microservices – Bounded Context

Bounded Context is a central pattern in Domain-Driven Design

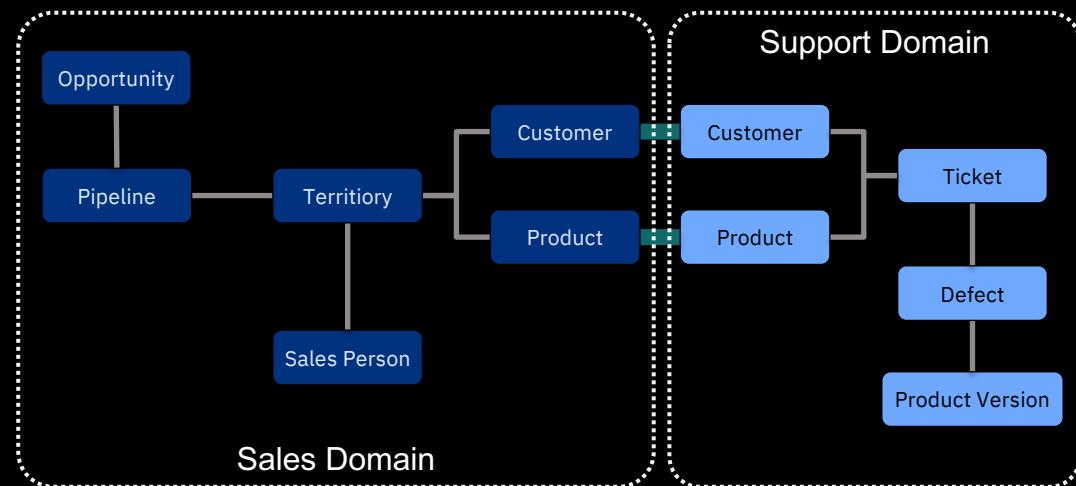
- DDD deals with large models by dividing them into different Bounded Contexts



# Microservices – Bounded Context

Bounded Context is a central pattern in Domain-Driven Design

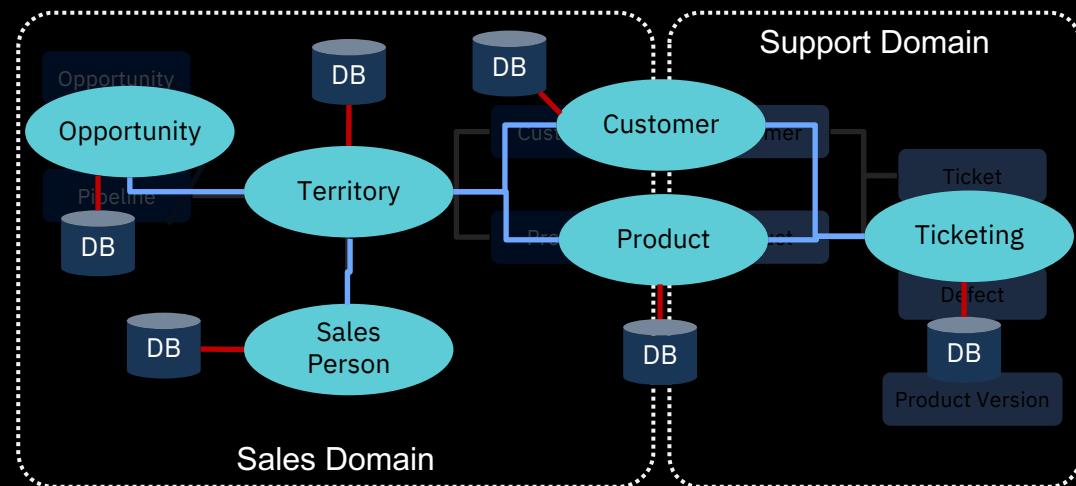
- DDD deals with large models by dividing them into different Bounded Contexts
- Helps to build and refine a model that represents our domains, contained within a boundary that defines our context.



# Microservices – Bounded Context

Bounded Context is a central pattern in Domain-Driven Design

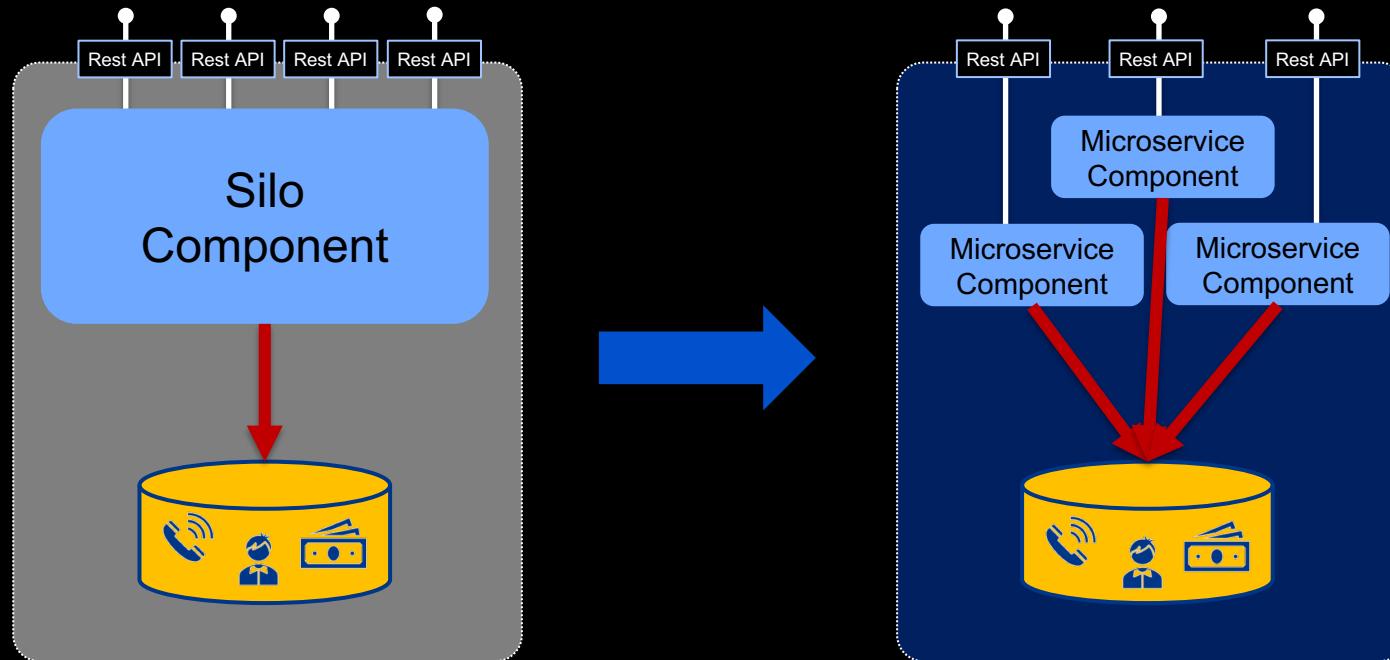
- DDD deals with large models by dividing them into different Bounded Contexts
- Helps to build and refine a model that represents our domains, contained within a boundary that defines our context.
- These boundaries end up being our microservices



# Microservices – Database refactoring

## Data Isolation

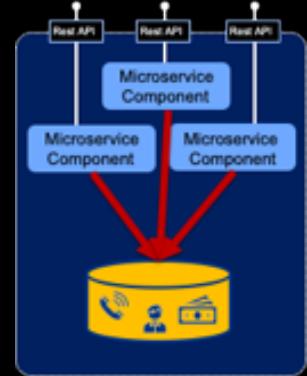
Each microservice must keep its own data. One typical error that happens in the beginning of a microservice transformation is the use of a centralized database, the same way it was used in the monolithic application. This creates a single point of failure and dependency between the microservices.



# Microservices – monolithic database

## Challenges

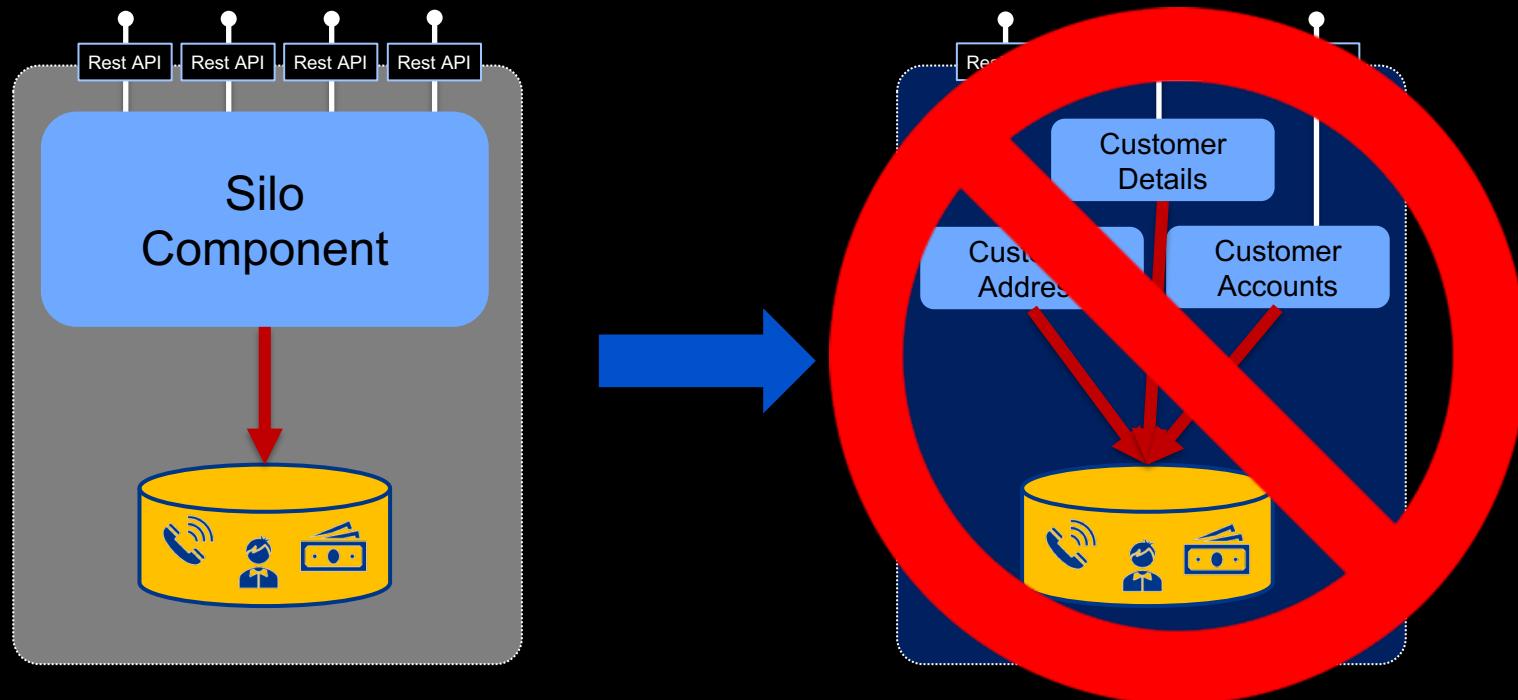
- Inability to deploy your service changes independently through **tight coupling**.
- Schema **changes need to be coordinated** amongst all the services
- **Difficult to scale** individual services
- Difficult to improve application performance (DB/Table size)
- All your services have to use a **relational database** even when a no-SQL datastore would bring benefits



# Microservices – Database refactoring

## Data Isolation

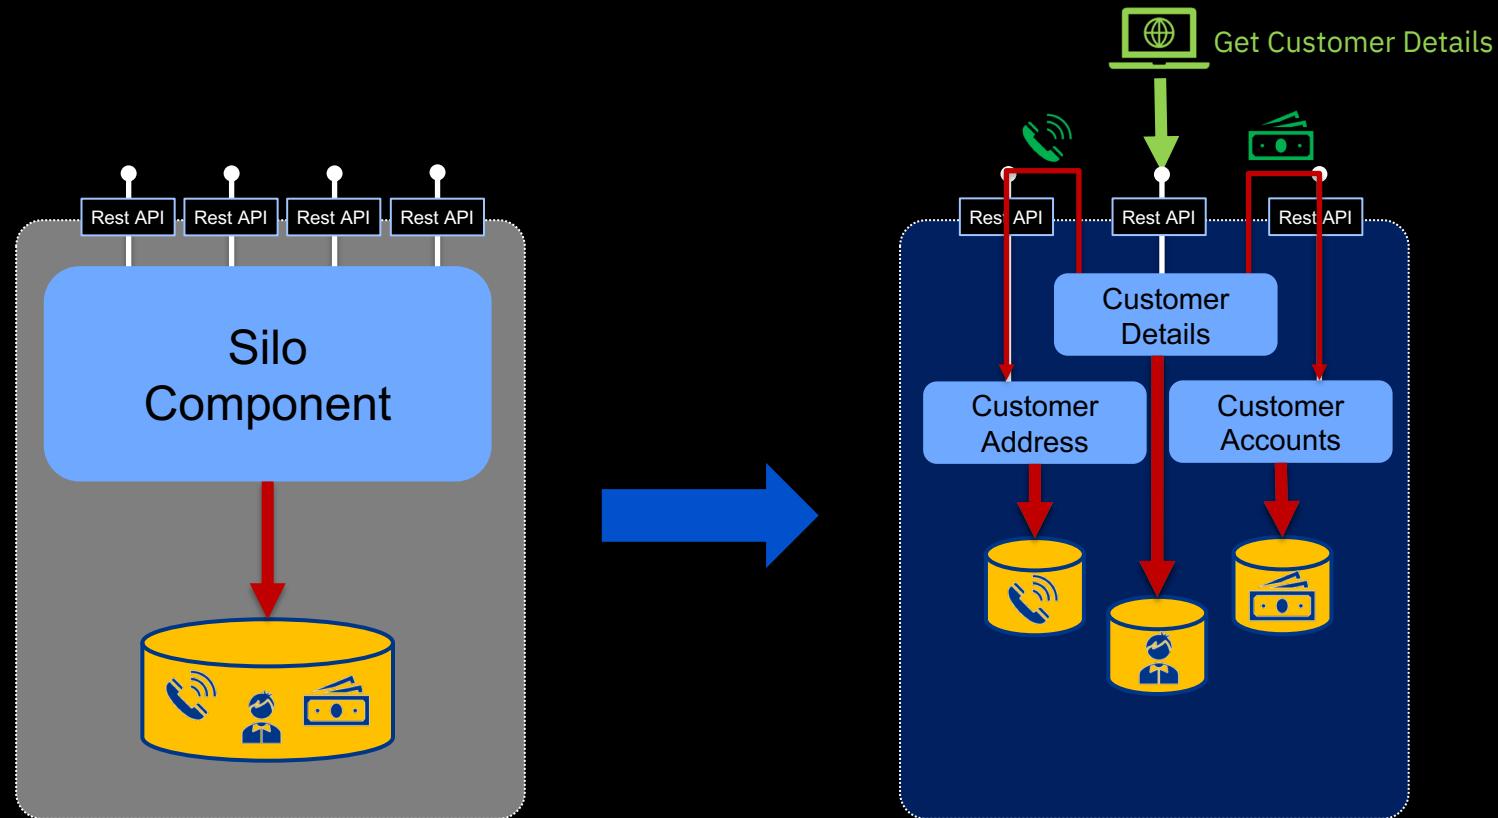
Each microservice must keep its own data. One typical error that happens in the beginning of a microservice transformation is the use of a centralized database, the same way it was used in the monolithic application. This creates a single point of failure and dependency between the microservices.



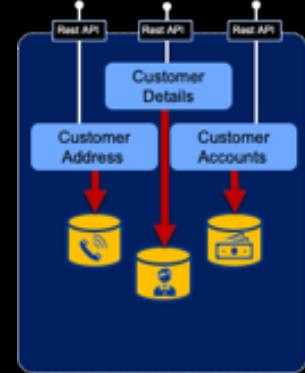
# Microservices – Database refactoring

## Data Isolation

Each microservice must keep its own data. One typical error that happens in the beginning of a microservice transformation is the use of a centralized database, the same way it was used in the monolithic application. This creates a single point of failure and dependency between the microservices.



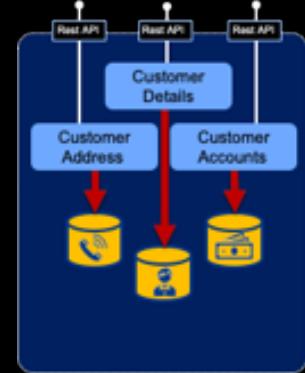
# Microservices – per-service database



## Benefits

- Avoid unexpected data modification – the API enforces consistency
- Make changes to our system without blocking progress of other parts of the system
- Store the data in a project-owned databases, using the appropriate technology
- Make changes to the schema/databases at our leisure
- Become much more scalable, fault tolerant, and flexible

# Microservices – per-service database



## Drawbacks

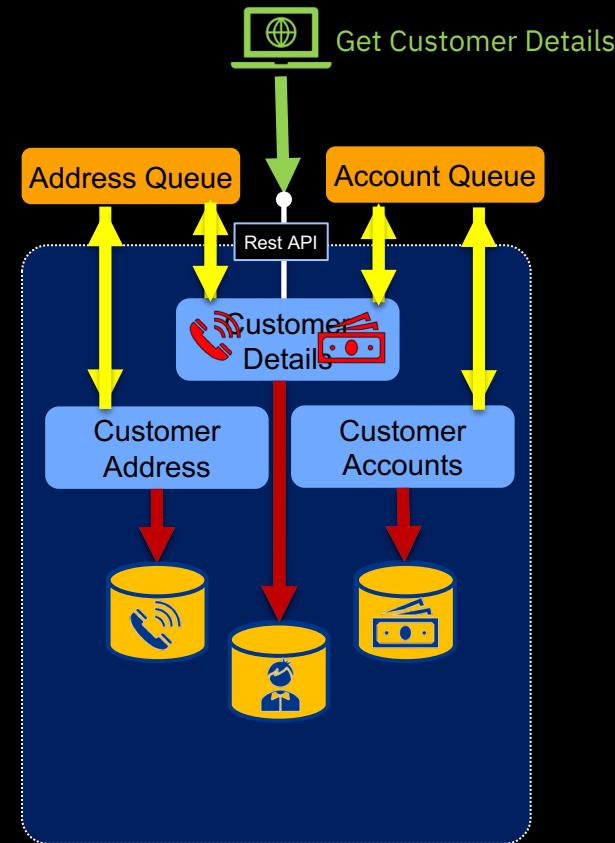
- Costly to refactor
- Needs robust transaction handling by the service (cost!)
- More difficult to debug
- More difficult to operationalize
- More difficult to test (needs well designed API tests)

# Event Driven Architecture

# Microservices – one step further

## Event-Driven Architecture

- Common pattern to maintain data consistency across different services.
- Avoid waiting for ACID transactions to complete
- Makes your application more available and performant
- Provides loose coupling between services



# Microservices Good or bad idea?

# Microservices – reasons not to use them

## Some thoughts

- If speed is your goal, microservices aren't the solution (MVP, doesn't have to scale)
- At what point is it in its lifecycle? Mission-critical?
- A moderately large, moderately complex application being maintained by a relatively small development and operations team
- Not all applications are large enough to break down into microservices
- Applications that require tight integration between individual components and services (real-time processing, ...)

# Microservices – reasons not to use them

## How to detect antipatterns

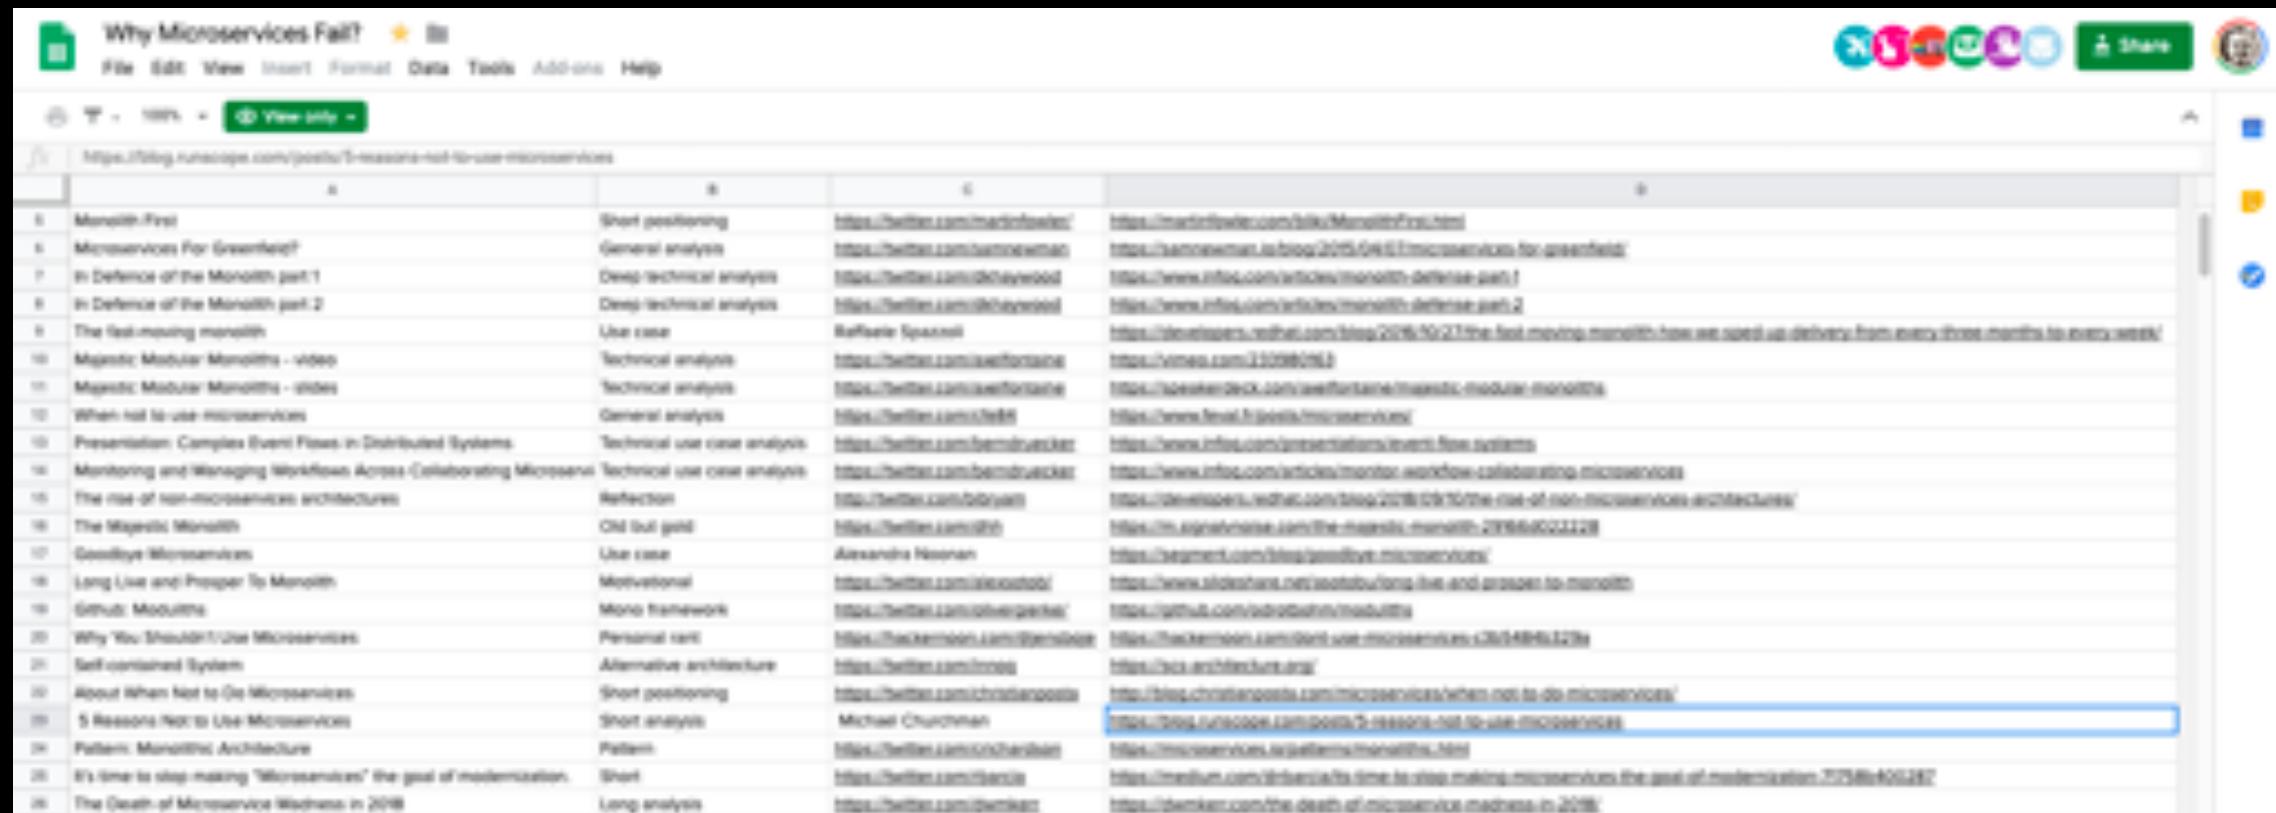
- A change to one microservice often requires changes to other microservices
- Deploying one microservice requires other microservices to be deployed at the same time
- Your microservices are overly chatty
- The same developers work across a large number of microservices
- Many of your microservices share a datastore
- Your microservices share a lot of the same code or models

# Microservices – reasons not to use them

What works for **large and complex** applications  
does not always work at a **smaller scale**

What makes sense for a **new application**  
does not always make sense when maintaining or  
updating **existing applications**

# Microservices – reasons not to use them

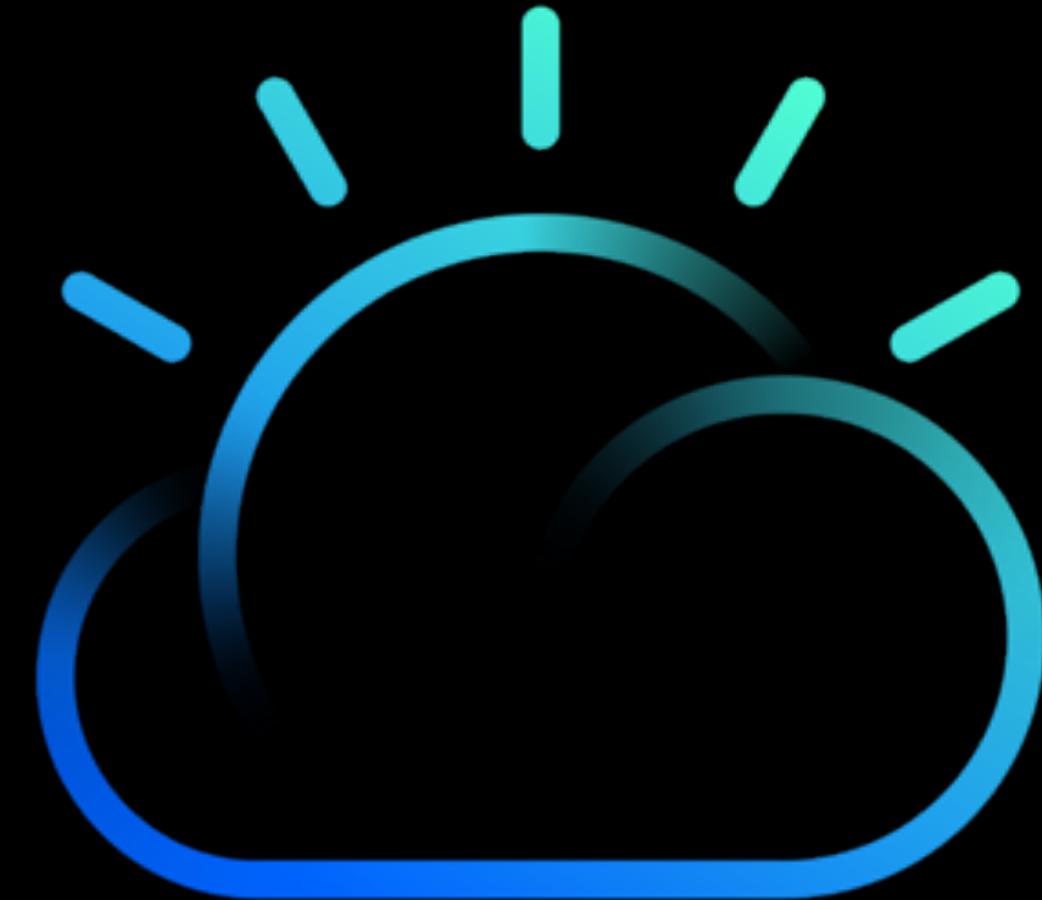


The screenshot shows a Google Sheets document with a title bar 'Why Microservices Fail?' and a toolbar with various icons. The main content is a table with four columns: A, B, C, and D. Column A lists 26 items, column B lists their types, column C lists their URLs, and column D lists their descriptions. The table has a light gray background and white text. The URL for item 25 is highlighted with a blue border.

A	B	C	D
1	Monolith First	Short positioning	<a href="https://twitter.com/michaelfowler/">https://twitter.com/michaelfowler/</a>
2	Microservices For Greenfield?	General analysis	<a href="https://twitter.com/tarceveman">https://twitter.com/tarceveman</a>
3	In Defence of the Monolith part 1	Deep technical analysis	<a href="https://twitter.com/dchaywood">https://twitter.com/dchaywood</a>
4	In Defence of the Monolith part 2	Deep technical analysis	<a href="https://twitter.com/dchaywood">https://twitter.com/dchaywood</a>
5	The fast-moving monolith	User case	<a href="https://medium.com/@michaelfowler/monolith-first-why-we-deliver-every-three-months-in-every-week-102f3a232011">https://medium.com/@michaelfowler/monolith-first-why-we-deliver-every-three-months-in-every-week-102f3a232011</a>
6	Magnetic Modular Monoliths - video	Technical analysis	<a href="https://www.infoq.com/presentations/magnetic-modular-monoliths-video">https://www.infoq.com/presentations/magnetic-modular-monoliths-video</a>
7	Magnetic Modular Monoliths - slides	Technical analysis	<a href="https://www.infoq.com/presentations/magnetic-modular-monoliths-slides">https://www.infoq.com/presentations/magnetic-modular-monoliths-slides</a>
8	When not to use microservices	General analysis	<a href="https://www.infoq.com/presentations/when-not-to-use-microservices">https://www.infoq.com/presentations/when-not-to-use-microservices</a>
9	Presentation: Complex Event Flows in Distributed Systems	Technical use case analysis	<a href="https://www.infoq.com/presentations/complex-event-flows-distributed-systems">https://www.infoq.com/presentations/complex-event-flows-distributed-systems</a>
10	Monitoring and Managing Workflows Across Collaborating Microservices	Technical use case analysis	<a href="https://www.infoq.com/presentations/monitor-and-manage-workflows-across-collaborating-microservices">https://www.infoq.com/presentations/monitor-and-manage-workflows-across-collaborating-microservices</a>
11	The rise of non-microservices architectures	Reflection	<a href="https://www.infoq.com/presentations/the-rise-of-non-microservices-architectures">https://www.infoq.com/presentations/the-rise-of-non-microservices-architectures</a>
12	The Magnetic Monolith	Old school	<a href="https://medium.com/@michaelfowler/the-magnetic-monolith-2098d4023228">https://medium.com/@michaelfowler/the-magnetic-monolith-2098d4023228</a>
13	Goodbye Microservices	User case	<a href="https://medium.com/@alexanderheinecke/goodbye-microservices">https://medium.com/@alexanderheinecke/goodbye-microservices</a>
14	Long Live and Prosper To Monolith	Motivational	<a href="https://www.slideshare.net/alexanderheinecke/long-live-and-prosper-to-monolith">https://www.slideshare.net/alexanderheinecke/long-live-and-prosper-to-monolith</a>
15	GitHub: Modular	Mono framework	<a href="https://github.com/alexanderheinecke/modular">https://github.com/alexanderheinecke/modular</a>
16	Why You Shouldn't Use Microservices	Personal rant	<a href="https://thekennen.com/2016/09/13/why-you-shouldnt-use-microservices/">https://thekennen.com/2016/09/13/why-you-shouldnt-use-microservices/</a>
17	Self-contained System	Alternative architecture	<a href="https://twitter.com/alexanderheinecke">https://twitter.com/alexanderheinecke</a>
18	About When Not to Use Microservices	Short positioning	<a href="http://blog.christianposta.com/microservices/when-not-to-use-microservices/">http://blog.christianposta.com/microservices/when-not-to-use-microservices/</a>
19	5 Reasons Not to Use Microservices	Short analysis	<a href="https://www.infoq.com/presentations/5-reasons-not-to-use-microservices">https://www.infoq.com/presentations/5-reasons-not-to-use-microservices</a>
20	Pattern: Monolithic Architecture	Pattern	<a href="https://twitter.com/michaelfowler">https://twitter.com/michaelfowler</a>
21	It's time to stop making "Microservices" the goal of modernization	Short	<a href="https://medium.com/@michaelfowler/its-time-to-stop-making-microservices-the-goal-of-modernization-7758b401287">https://medium.com/@michaelfowler/its-time-to-stop-making-microservices-the-goal-of-modernization-7758b401287</a>
22	The Death of Microservice Madness in 2018	Long analysis	<a href="https://davidsarno.com/the-death-of-microservice-madness-in-2018/">https://davidsarno.com/the-death-of-microservice-madness-in-2018/</a>

[https://docs.google.com/spreadsheets/d/1vjnjAII\\_8TZBv2XhFHra7kEQzQpOHSZpFIWDjynYYf0/edit?pli=1#gid=0](https://docs.google.com/spreadsheets/d/1vjnjAII_8TZBv2XhFHra7kEQzQpOHSZpFIWDjynYYf0/edit?pli=1#gid=0)

QUESTIONS?





°° The Journey to Cloud  
**Docker - Hands-On**

05



IBM Cloud

# Remember your Team Color

**black 31701**

**olive 31711**

**peru 31715**

**white 31702**

**brown 31712**

**chocolate 31716**

**red 31703**

**lightblue 31713**

**orchid 31717**

**blue 31704**

**orange 31708**

**gold 31718**

**yellow 31705**

**purple 31709**

**pink 31719**

**lime 31706**

**maroon 31710**

**violet 31720**

**cyan 31707**

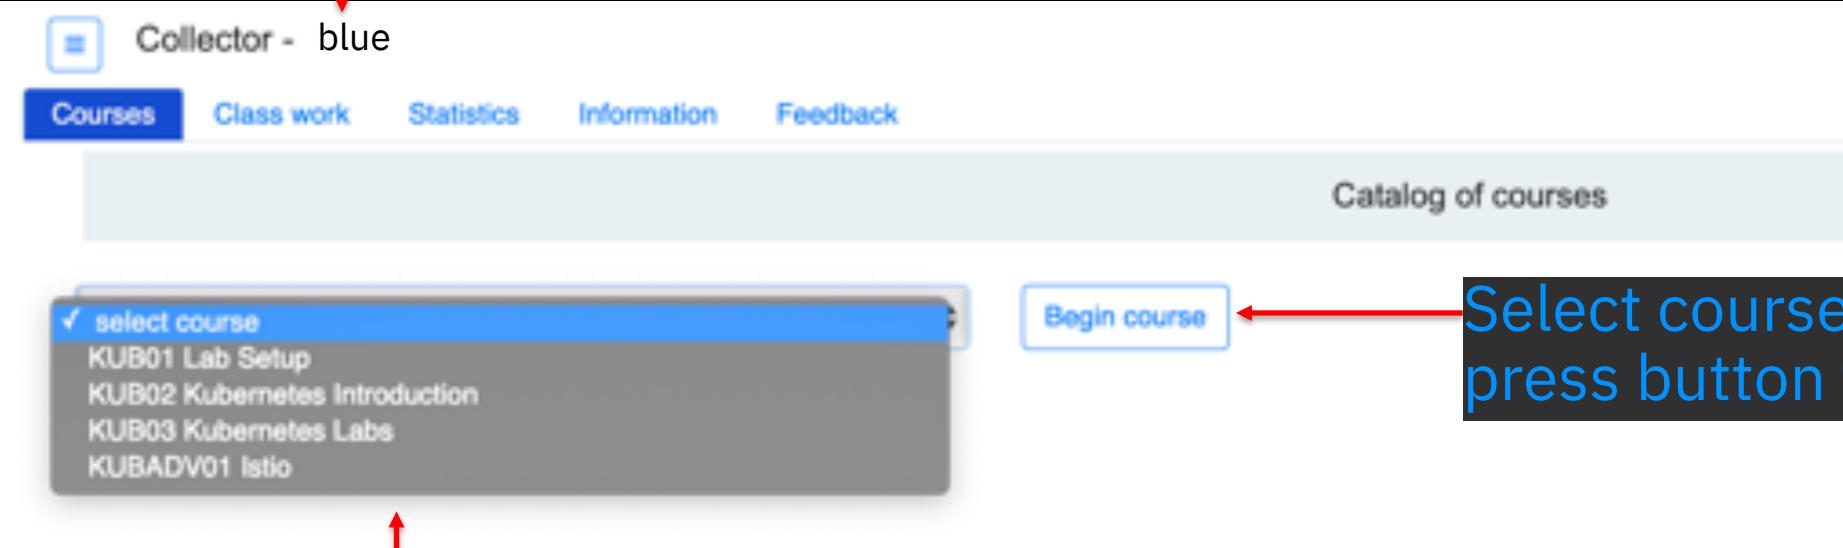
**firebrick 31714**

# Collector - Accessing team web site

`http://158.177.137.195:{port#}`

Team name / color will be shown

blue **31704**



The screenshot shows a web browser displaying the 'Collector - blue' website. At the top, there is a navigation bar with links for 'Courses' (which is highlighted in blue), 'Class work', 'Statistics', 'Information', and 'Feedback'. Below the navigation bar, a section titled 'Catalog of courses' contains a list of courses under the heading 'select course'. The courses listed are: KUB01 Lab Setup, KUB02 Kubernetes Introduction, KUB03 Kubernetes Labs, and KUBADV01 Istio. To the right of this list is a blue button labeled 'Begin course'. A red arrow points from the text 'Select course and press button to begin' to the 'Begin course' button. Another red arrow points from the text 'Current course catalog' to the list of courses.

Team name / color will be shown

blue **31704**

Collector - blue

Courses Class work Statistics Information Feedback

Catalog of courses

select course

- KUB01 Lab Setup
- KUB02 Kubernetes Introduction
- KUB03 Kubernetes Labs
- KUBADV01 Istio

Begin course

Select course and press button to begin

Current course catalog



# JTC01 Docker Labs

## Lab 1:

- Creating images
- Run the Images
- Use Portainer tool
- Use private registries

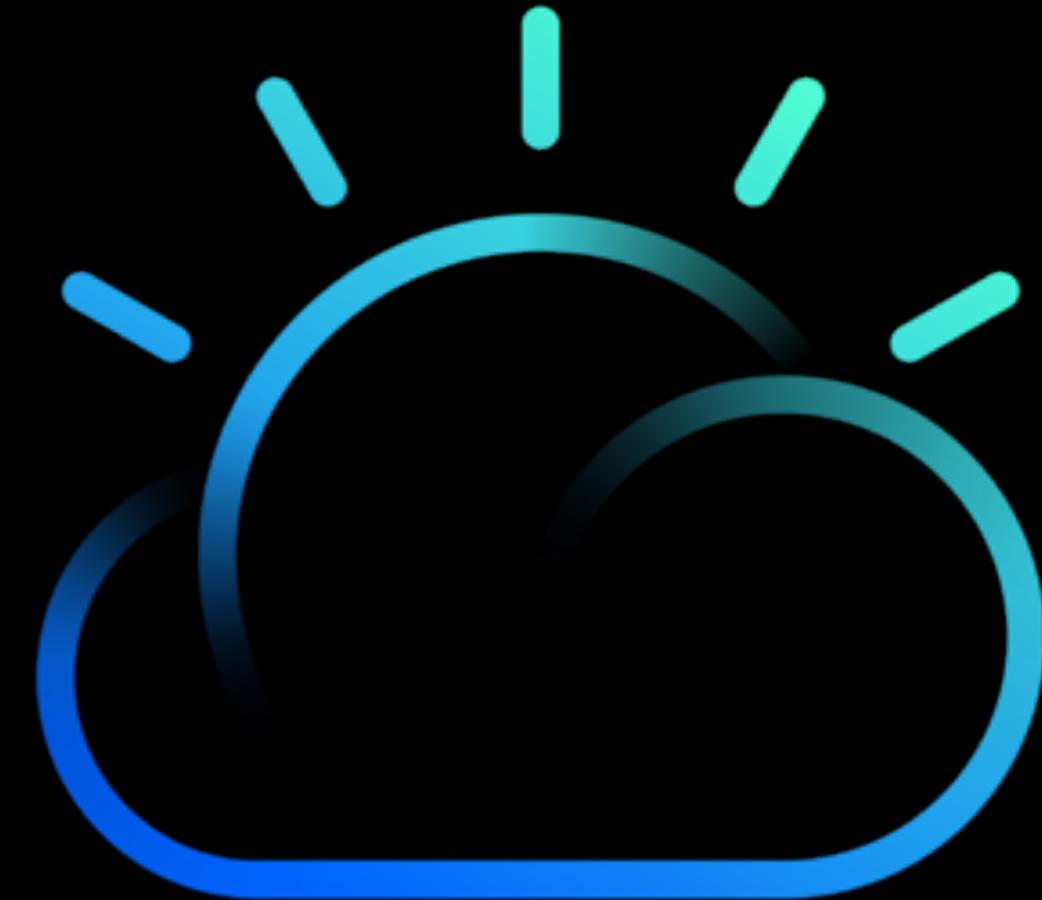


READY  
SET  
GO!!!!

<http://158.177.137.195:{port#}>

Duration: 60 mins

QUESTIONS?





Sources and documentation will be available here:

<https://github.com/niklaushirt/training>



°° The Journey to Cloud  
**Kubernetes - Hands-On**

06



IBM Cloud

# Remember your Team Color

**black 31701**

**olive 31711**

**peru 31715**

**white 31702**

**brown 31712**

**chocolate 31716**

**red 31703**

**lightblue 31713**

**orchid 31717**

**blue 31704**

**orange 31708**

**gold 31718**

**yellow 31705**

**purple 31709**

**pink 31719**

**lime 31706**

**maroon 31710**

**violet 31720**

**cyan 31707**

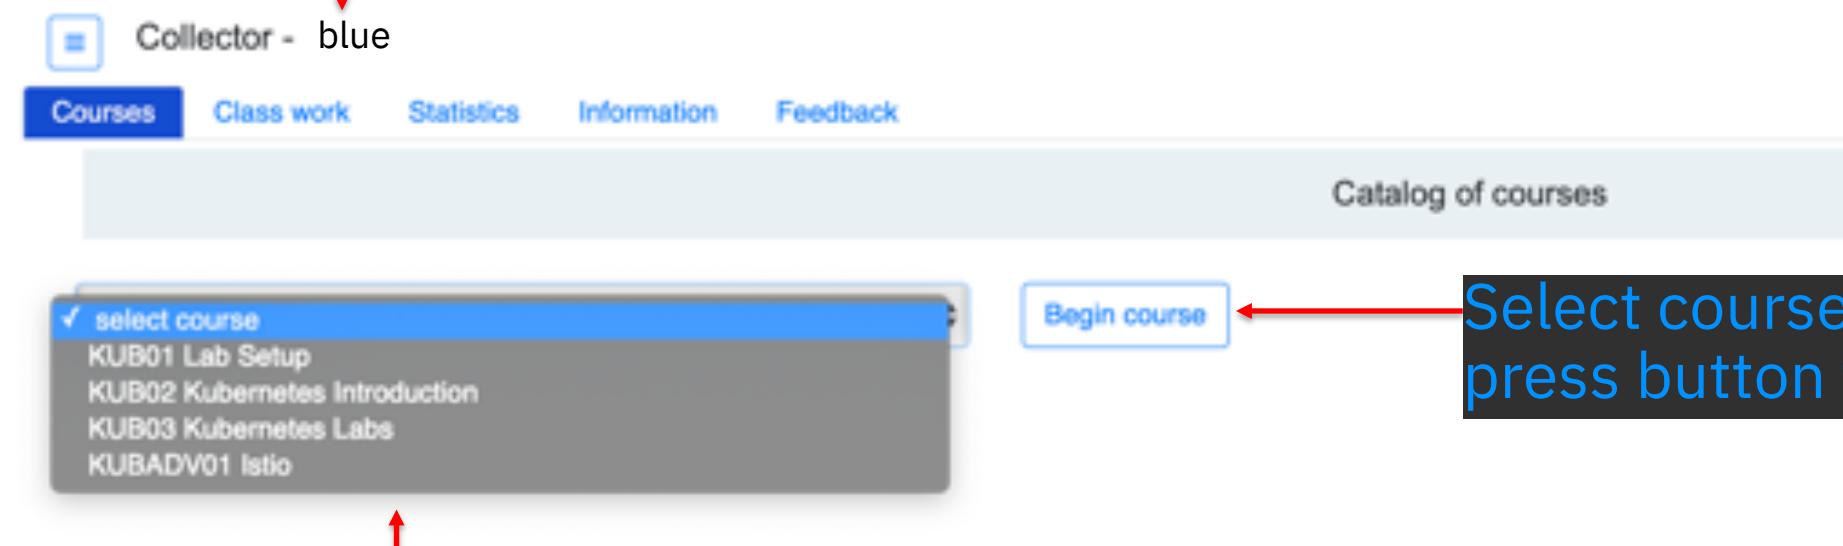
**firebrick 31714**

# Collector - Accessing team web site

`http://158.177.137.195:{port#}`

Team name / color will be shown

blue **31704**



The screenshot shows a web browser displaying the 'Collector - blue' website. At the top, there is a navigation bar with links for 'Courses' (which is highlighted in blue), 'Class work', 'Statistics', 'Information', and 'Feedback'. Below the navigation bar, a section titled 'Catalog of courses' contains a list of courses under the heading 'select course'. The courses listed are: KUB01 Lab Setup, KUB02 Kubernetes Introduction, KUB03 Kubernetes Labs, and KUBADV01 Istio. To the right of this list is a blue button labeled 'Begin course'. A red arrow points from the text 'Select course and press button to begin' to the 'Begin course' button. Another red arrow points from the text 'Current course catalog' to the list of courses.

Team name / color will be shown

blue **31704**

Collector - blue

Courses Class work Statistics Information Feedback

Catalog of courses

select course

- KUB01 Lab Setup
- KUB02 Kubernetes Introduction
- KUB03 Kubernetes Labs
- KUBADV01 Istio

Begin course

Select course and press button to begin

Current course catalog



## JTC02 Kubernetes Labs

Lab 1: Overview Kubernetes

Lab 2: Running your first Pod on Kubernetes

Lab 3: Deploy a Web Application on Kubernetes

Lab 4: Scale an application on Kubernetes.

Lab 5: Persisting data in Kubernetes with Volumes



READY  
SET  
GO!!!!

<http://158.177.137.195:{port#}>

Duration: 90 mins

QUESTIONS?





Sources and documentation will be available here:

<https://github.com/niklaushirt/training>

# The Journey to Cloud **Kubernetes Security**



IBM Cloud

# Kubernetes – Security – Back to basics

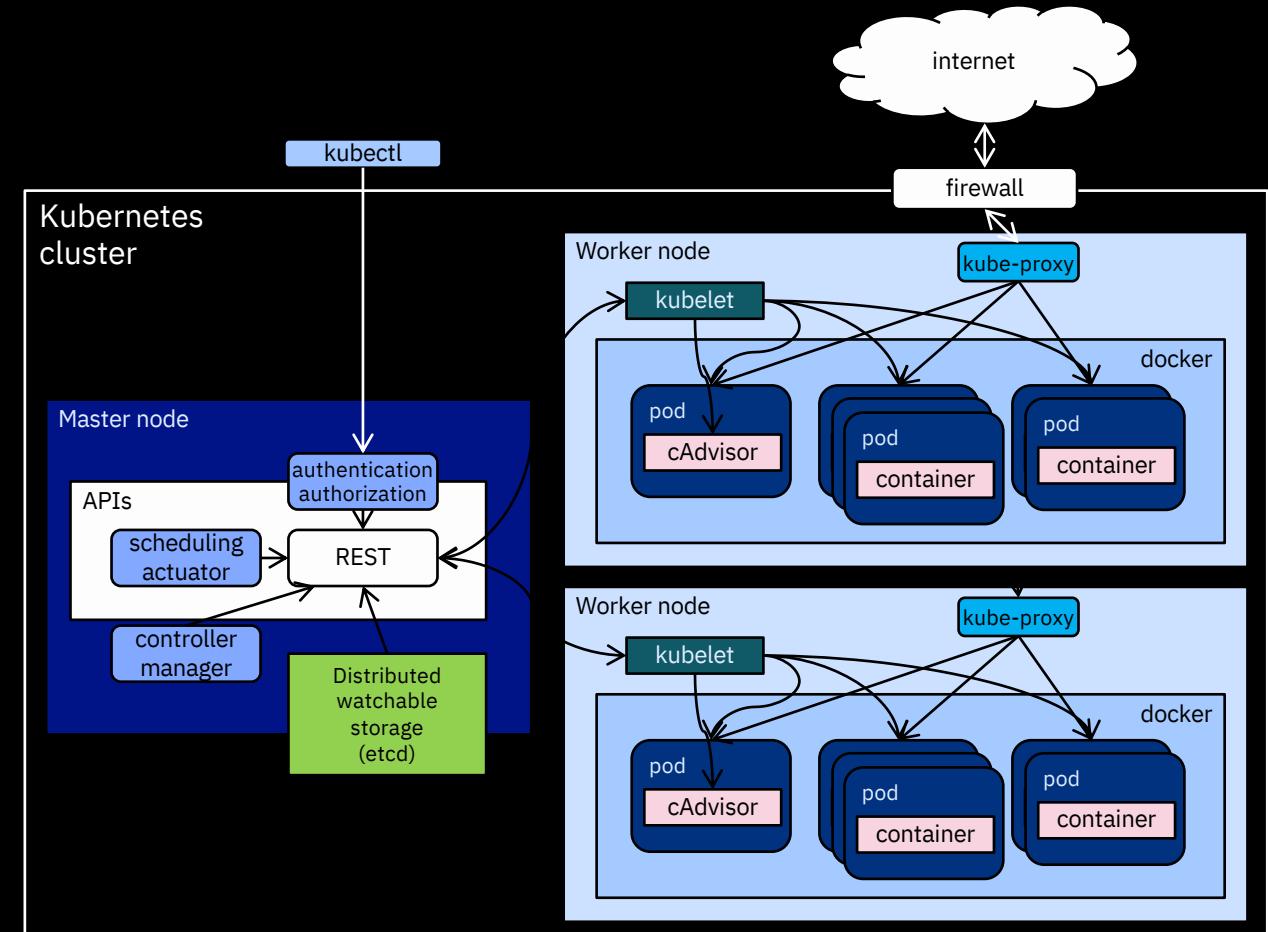


## Master node

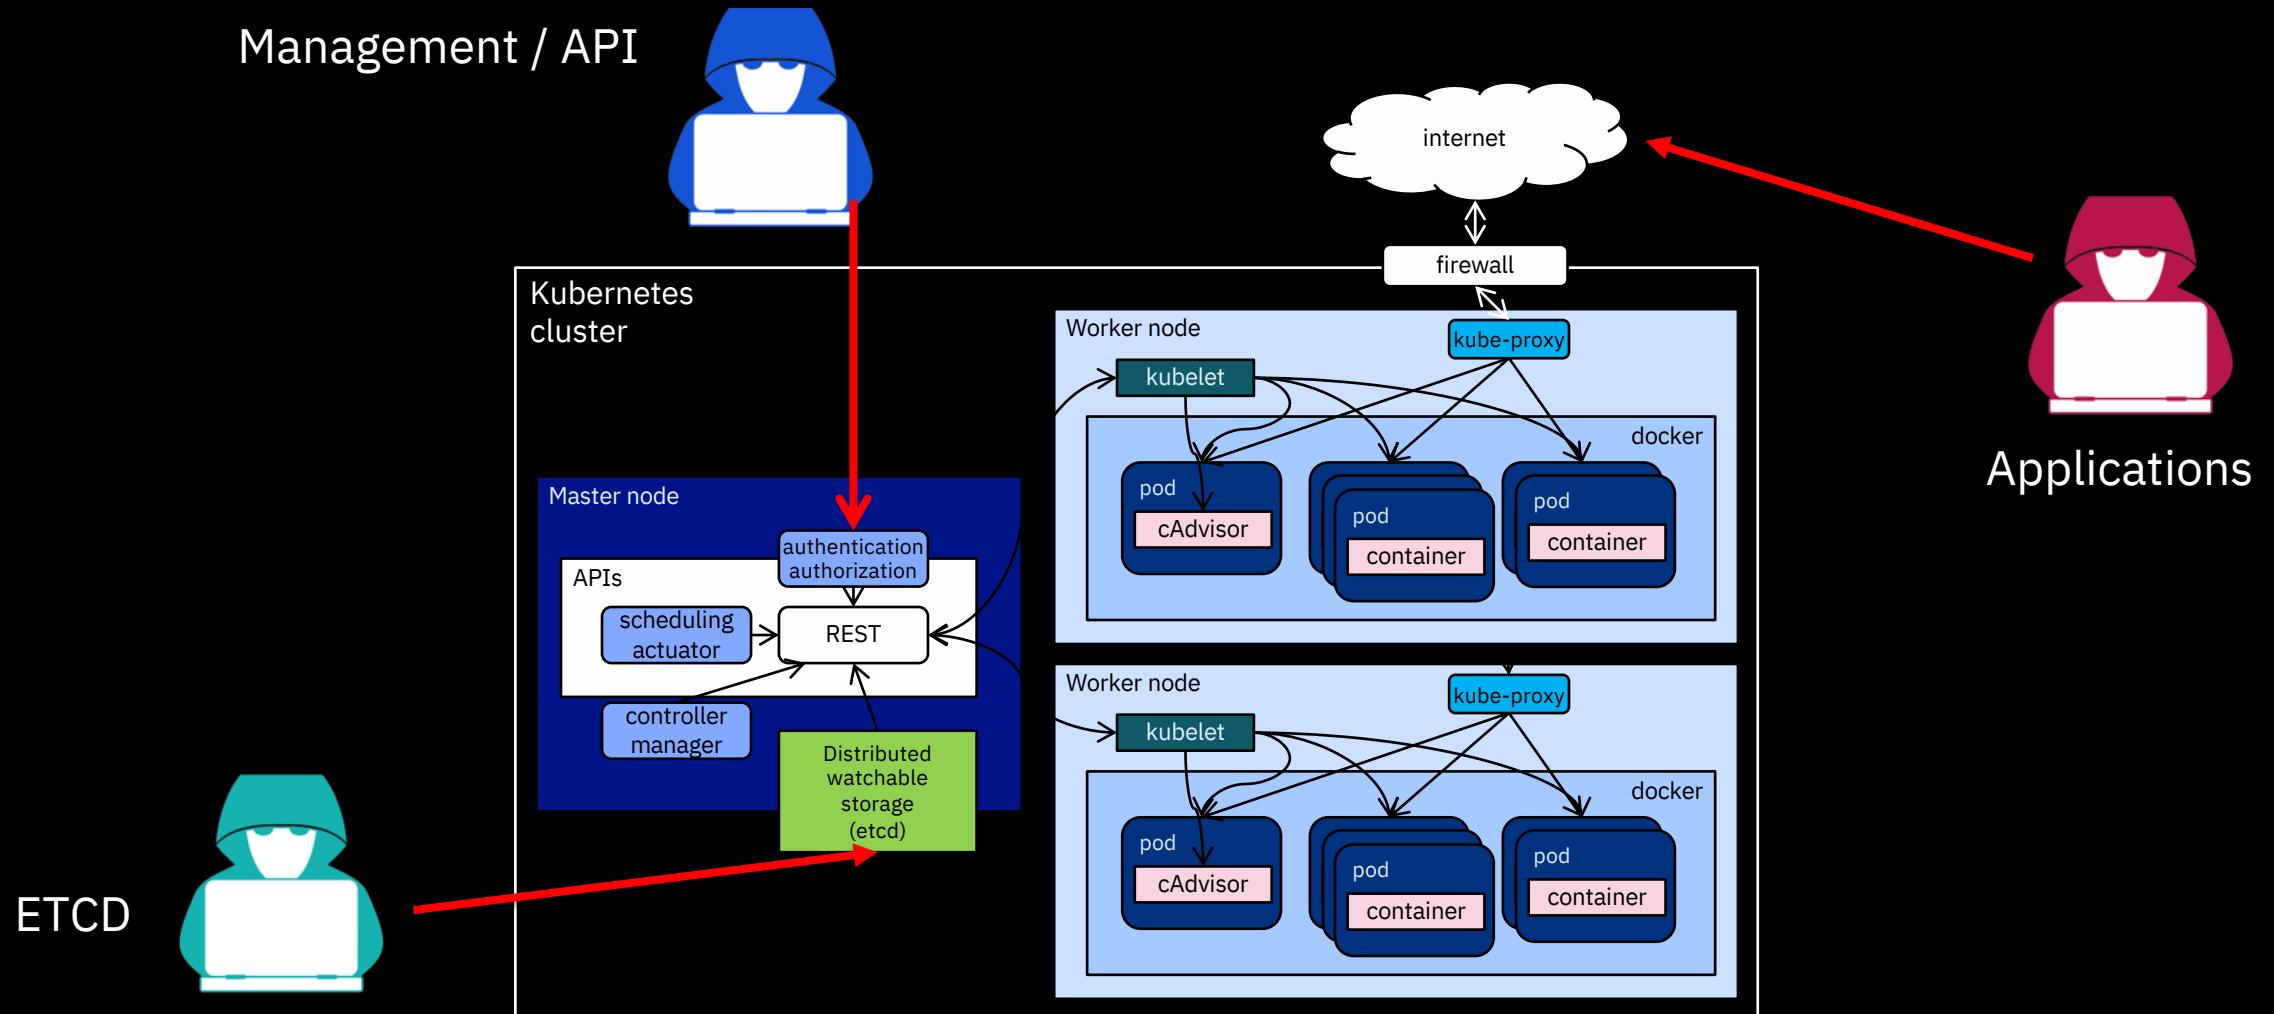
- Node that manages the cluster
- Scheduling, replication & control
- Multiple nodes for HA

## Worker nodes

- Node where pods are run
- Docker engine
- kubelet agent accepts & executes commands from the master to manage pods
- cAdvisor – Container Advisor provides resource usage and performance statistics
- kube-proxy – routes inbound or ingress traffic



# Kubernetes – Security – Attack surface



# Kubernetes – Security Basic Topics

The **Billion Laughs** attack is a particular type of denial of service (DoS) attack which is aimed specifically at XML document parsers. This attack is also referred to as an XML bomb or an exponential entity expansion attack.

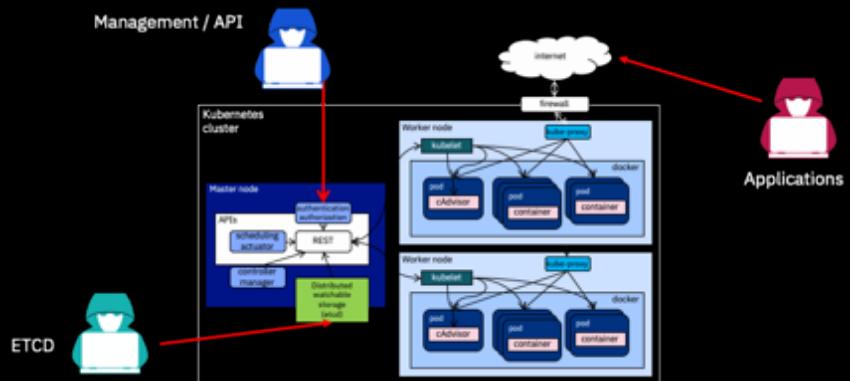
It exploits the fact that nested references to nodes can grow very large when expanded. Because the **kube-apiserver** doesn't perform validation on the manifest, it doesn't detect if those nested references will cause a problem. If the nesting references grow too large, excessive CPU and RAM usage can render the apiserver unresponsive to connections ... hence the Denial of Service.



# Kubernetes – Security Basic Topics

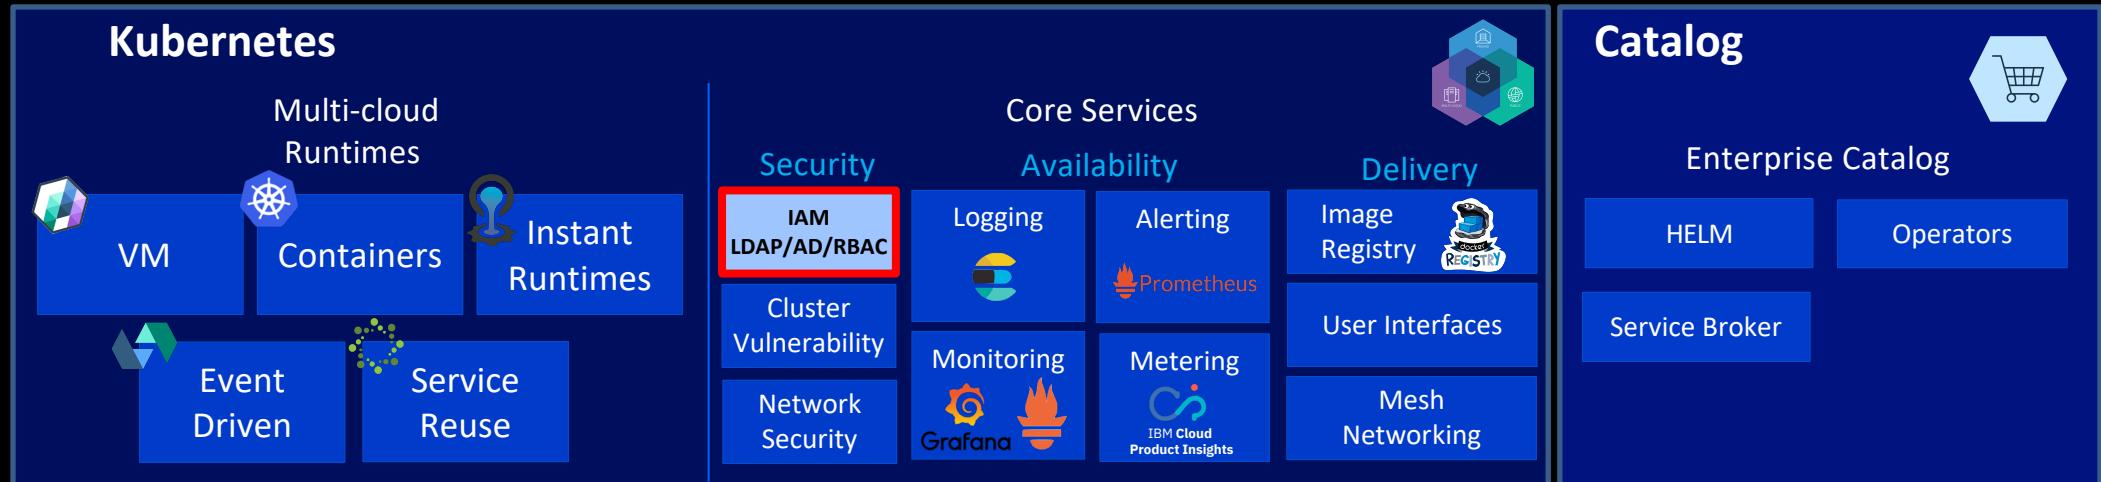
## Reduce Kubernetes Attack Surfaces

- Secure access to etcd
- Controlling access to the Kubernetes API
- Controlling access to the Kubelet
- Enforce resource usage limits for workloads
- Rotate infrastructure credentials frequently
- Enable audit logging
- Use Linux security features
- Controlling what privileges containers run with
- Enforcing Network Policies
- Image Scanning of your containers



Source: <https://kubernetes.io/docs/tasks/administer-cluster/securing-a-cluster>  
<https://kubernetes.io/blog/2018/07/18/11-ways-not-to-get-hacked/>

# Kubernetes – Core Services - Controlling K8s Access



**Authentication** – WHO am I - (token, certs, OpenID Connect Provider (OIDC)...)

*Developer* – create, view, get permission.

*Tester* – create, delete, update, get permission.

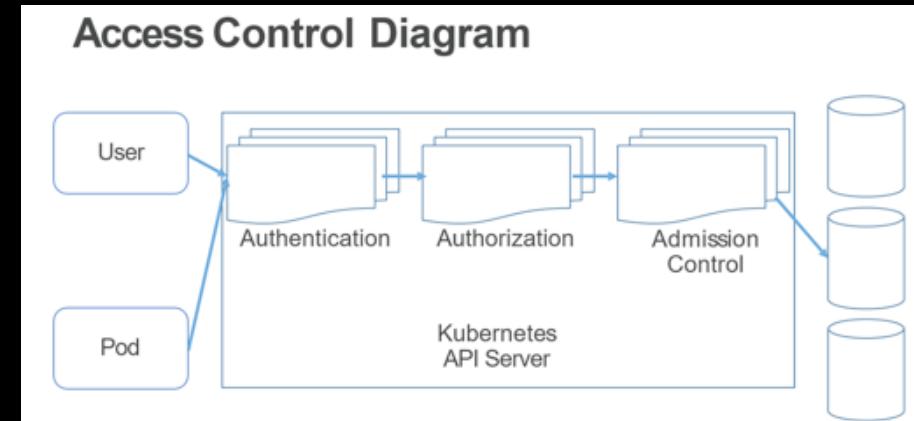
*Administrator* – All permission

**Authorization** – CAN I - (RBAC)

Is the user or application authorized or have permissions to access a K8s object?

**Admission Control**

Intercepts requests to the Kubernetes API server prior to persistence of the object, but after the request is authenticated and authorized





## Integrating with LDAP/AD

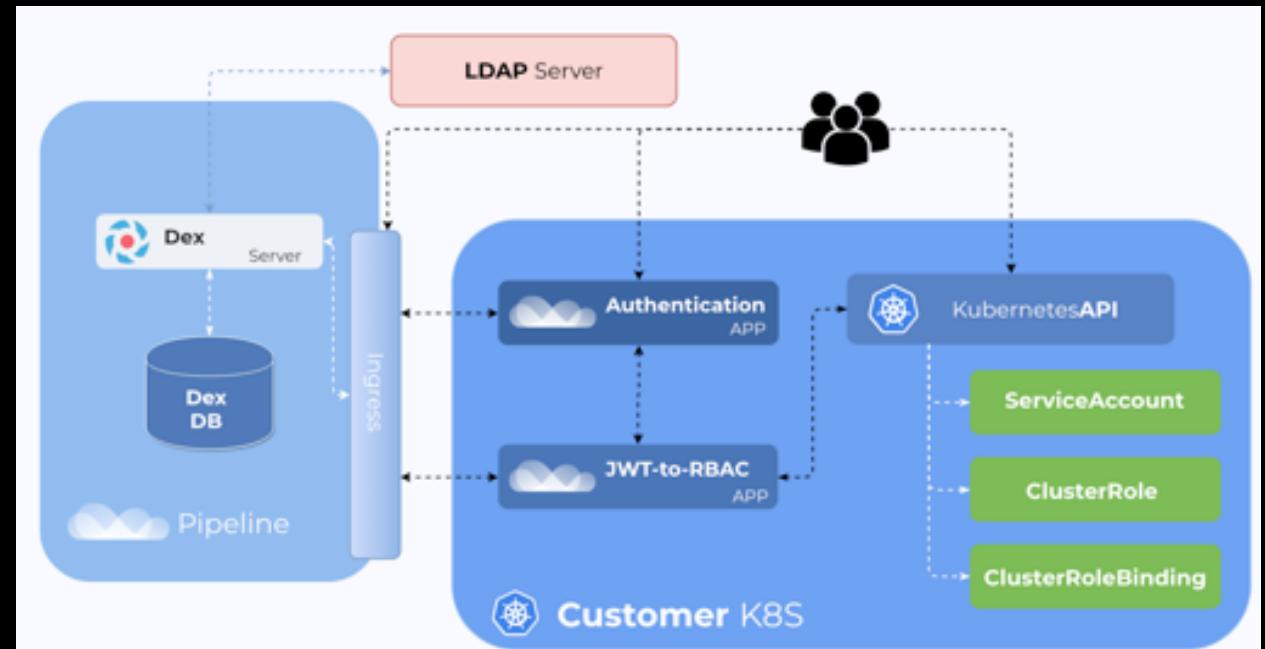
### OpenID Connect

ID Tokens are JSON Web Tokens (JWTs)

### dex

Dex acts as a portal to other identity providers through connectors

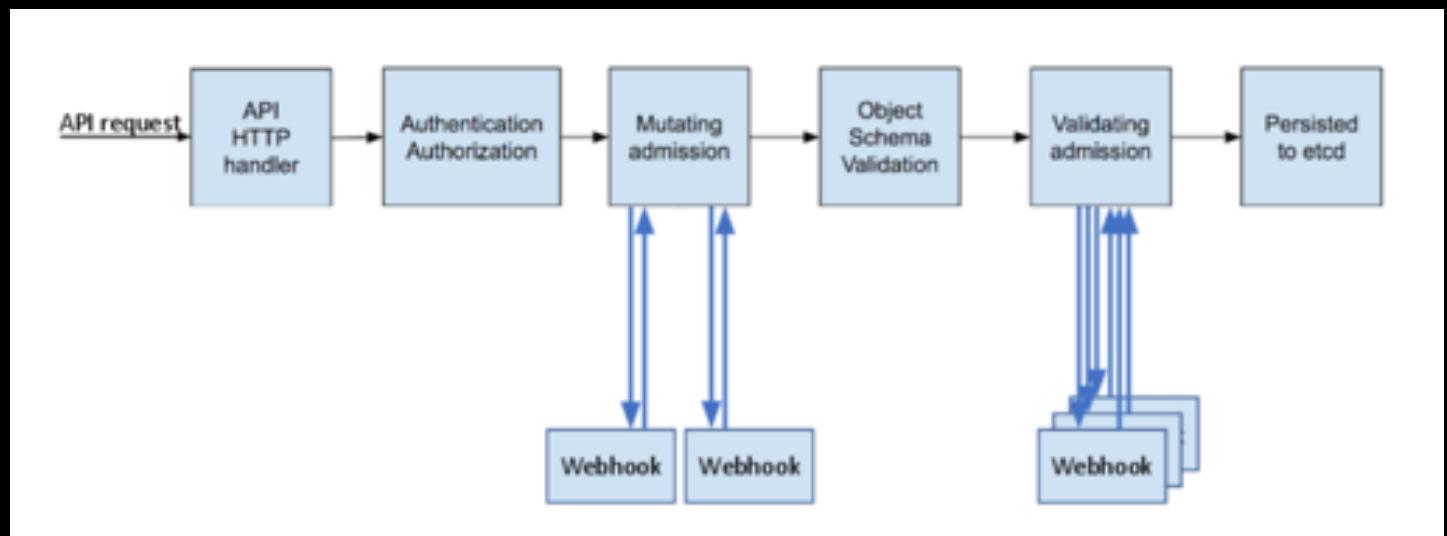
dex defer authentication to LDAP servers, SAML providers, or established identity providers like GitHub, Google, and Active Directory.



# Kubernetes WebHooks

### WebHooks Type

- **mutating**: to dynamically change incoming deployments on the fly (think automatic Istio sidecar injection), and
- **validating**: to accept or reject those same deployments based on the rules in your callback.
- → Image Scanner



## RBAC Basic Elements

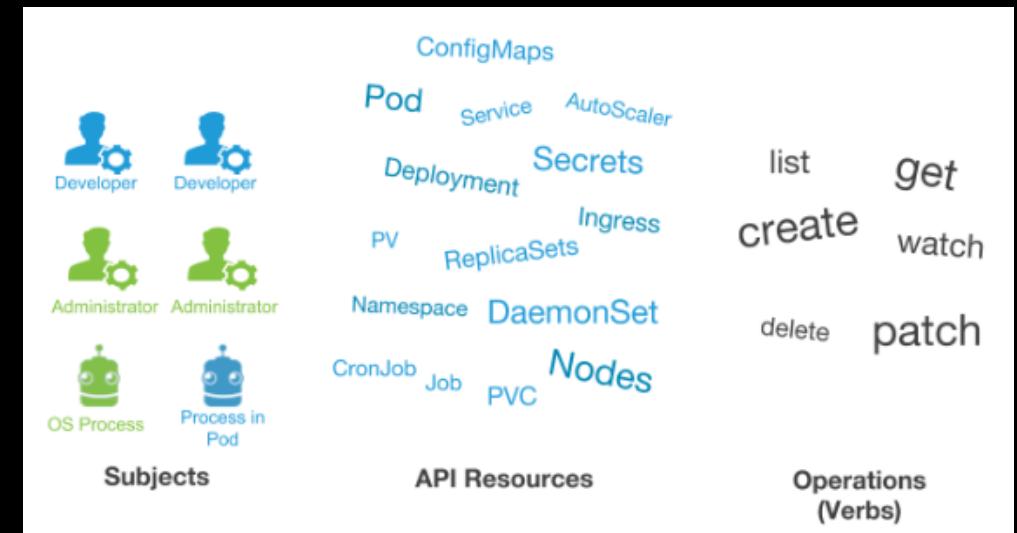
### RBAC Building Blocks

#### Objects

- Pods
- PersistentVolumes
- ConfigMaps
- Deployments
- Nodes
- Secrets
- Namespaces

#### Verbs

- create
- get
- delete
- list
- update
- edit
- watch
- exec



## RBAC Basic Elements

### Rules

Operations (verbs) that can be carried out on a group of resources which belong to different API Groups.

### Roles and ClusterRoles

Both consist of rules.

- Role: applicable to a single namespace
- ClusterRole: is cluster-wide

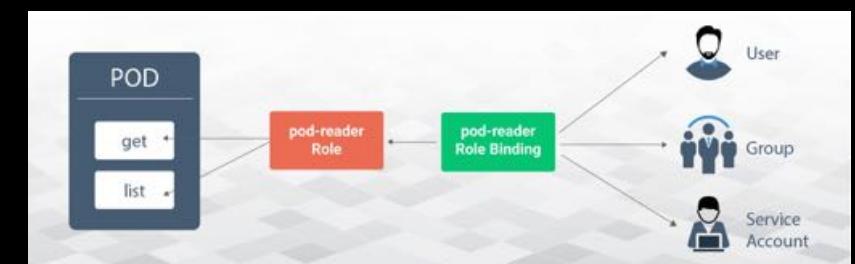
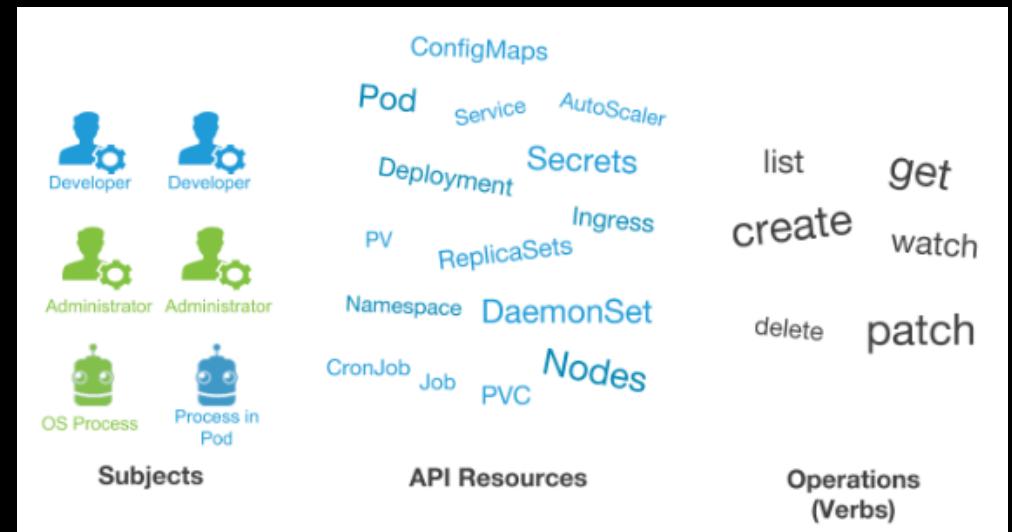
### Subjects

Entity that attempts an operation in the cluster

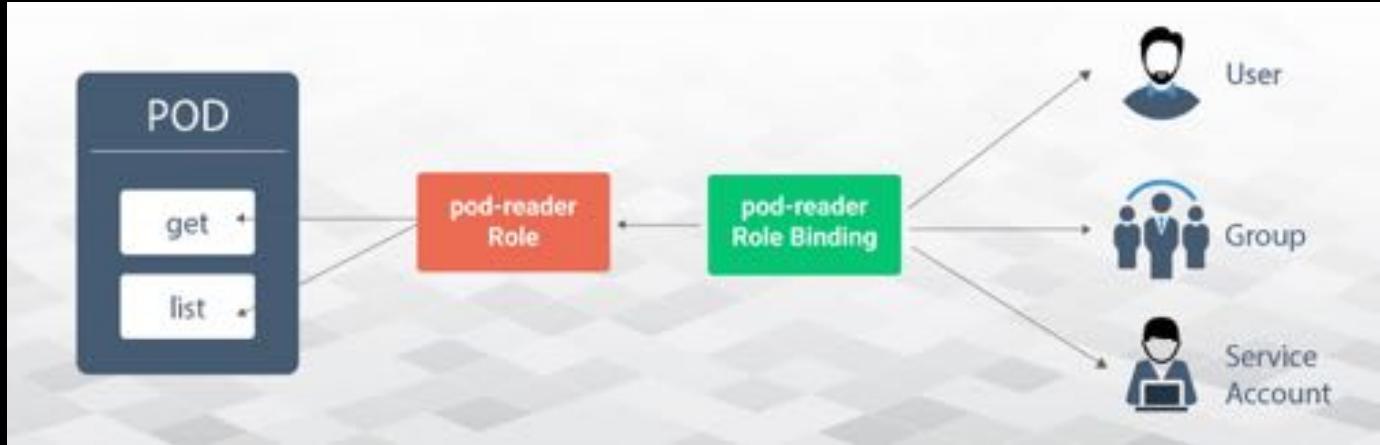
- User Accounts: Humans or processes living outside the cluster.
- Service Accounts: Namespaced account.
- Groups: Reference multiple accounts. (default groups like cluster-admin)

### RoleBindings and ClusterRoleBindings

Bind subjects to roles



## RBAC Example



```
kind: Role
metadata:
  name: pod-reader
rules:
- apiGroups: [""]
  resources: ["pods"]
  verbs: ["get", "list"]
```

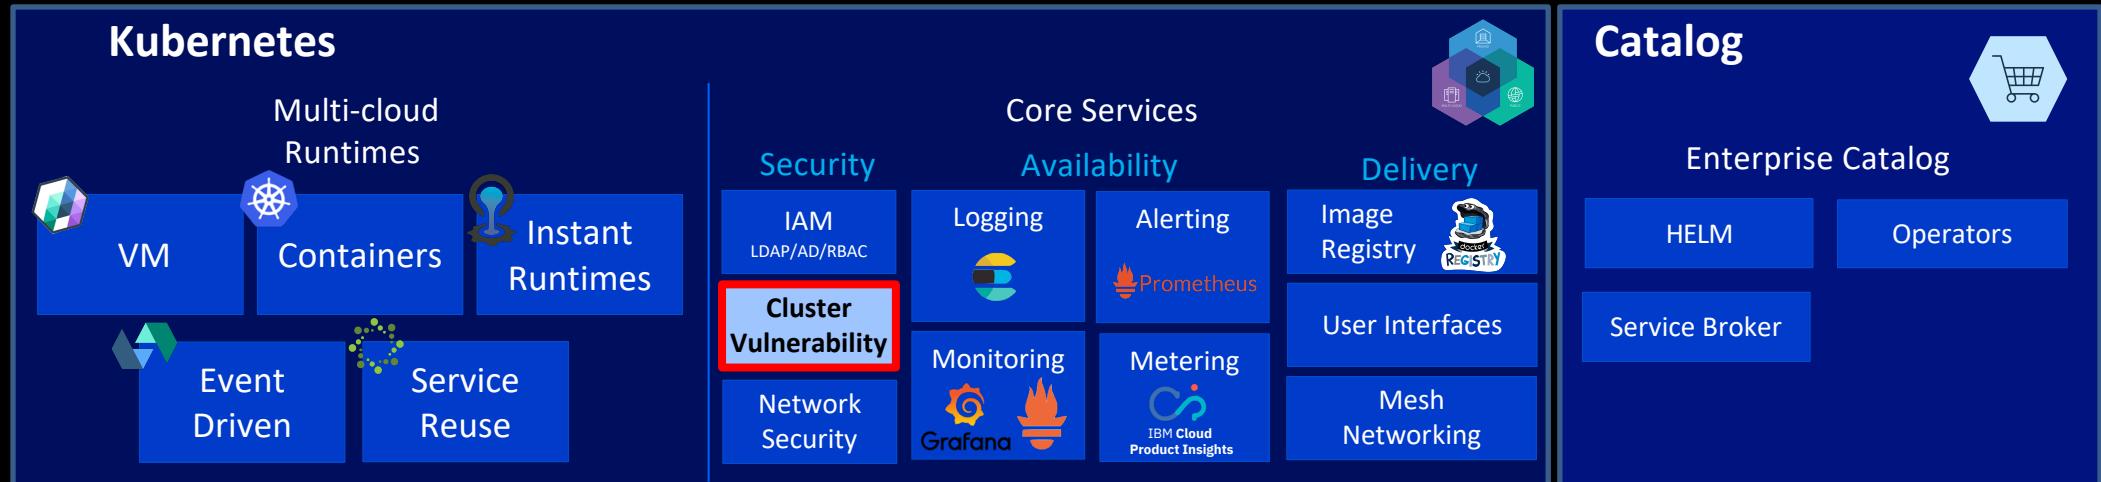
```
kind: RoleBinding
metadata:
  name: pod-reader
subjects:
- kind: User
  name: John
roleRef:
  kind: Role
  name: pod-reader
```

# Common RBAC Pitfalls

- **Cluster Administrator Role Granted Unnecessarily**  
The built-in **cluster-admin** role grants effectively unlimited access to the cluster.  
Should be granted only to the specific users that need it.
- **Duplicated Role Grant**  
Role definitions may overlap with each other, making access revocation more difficult.
- **Unused Role**  
Roles that are created but not granted to any subject can increase the complexity of RBAC management.
- **Grant of Missing Roles**  
Role bindings can reference roles that do not exist.  
If the role name is reused those role bindings can unexpectedly become active.

Always adapt  
least privileged access  
practices

# Kubernetes – Core Services – Cluster Security

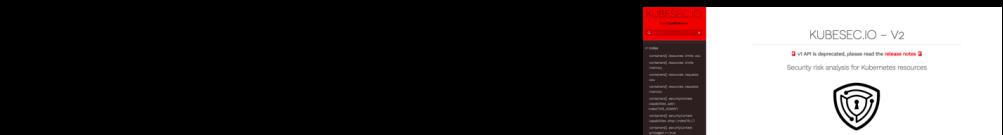


# Kubernetes – Cluster Security – Secure Deployments

## Scan for K8s Best Practices

### kubesc – Validation of security best practices

Performs security risk analysis for Kubernetes resources and tells you what you should change in order to improve the security of those pods. It also gives you a score that you can use to create a minimum standard. The score incorporates a great number of Kubernetes best practices.



```
"object": "Deployment/k8sdemo.default",
"valid": true,
"message": "Passed with a score of 4 points",
"score": 4,
"scoring": {
  "advise": [
    {
      "selector": ".spec .serviceAccountName",
      "reason": "Service accounts restrict Kubernetes API access and should be configured with least privilege",
      "points": 3
    },
    {
      "selector": "metadata .annotations .\\\"container.seccomp.security.alpha.kubernetes.io/pod\\\"",
      "reason": "Seccomp profiles set minimum privilege and secure against unknown threats",
      "points": 1
    }
  ]
},
```

```
{
  "selector": ".spec .serviceAccountName",
  "reason": "Service accounts restrict Kubernetes API access and should be configured with least privilege",
  "points": 3
},
```

```
  "points": 1
},
{
  "selector": "containers[] .securityContext .runAsUser > 10000",
  "reason": "Run as a high-UID user to avoid conflicts with the host's user table",
  "points": 1
}
]
```

# Kubernetes – Cluster Security – Secure Images

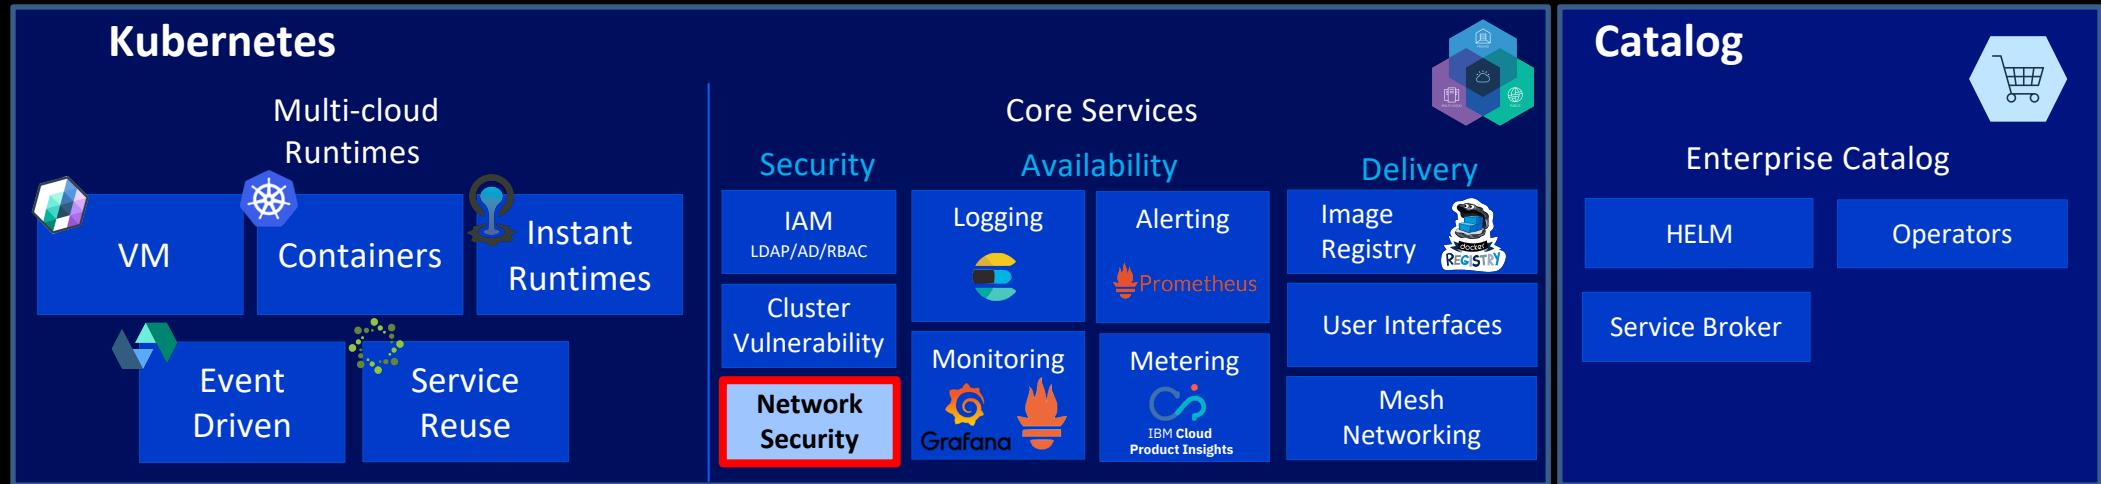
## Scan images for vulnerabilities



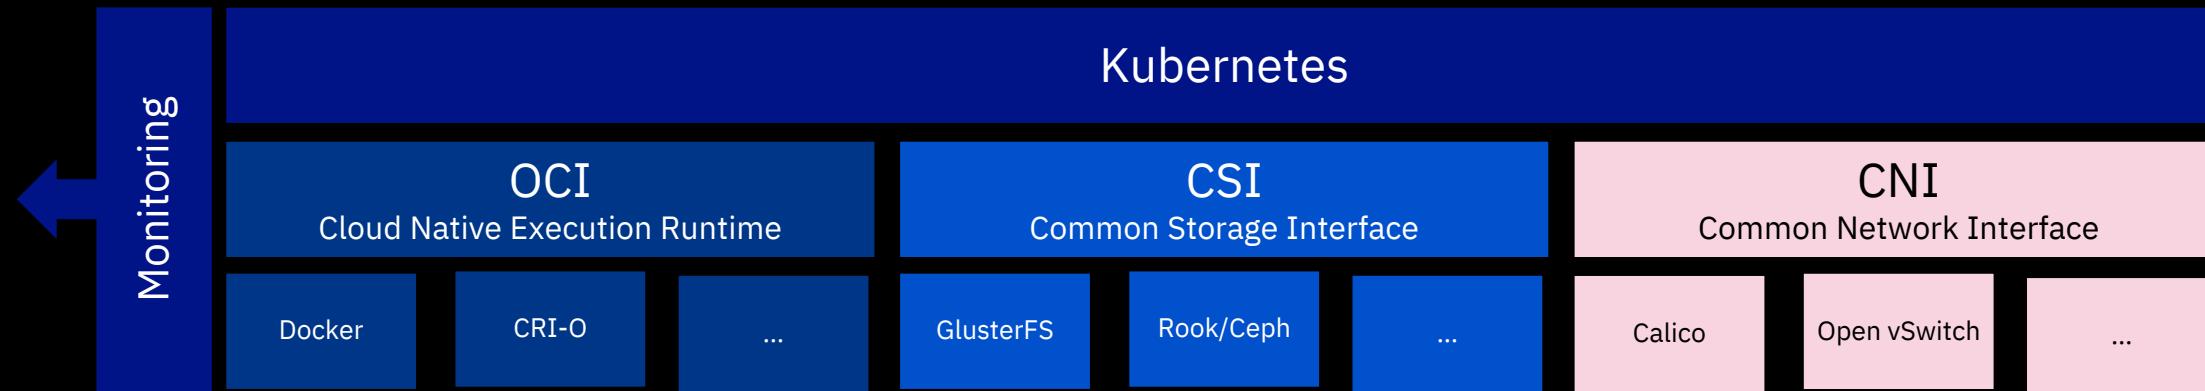
One of the best-known Image scanning tools is Clair. It is an open source project for the static analysis of vulnerabilities in application containers (currently including appc and docker).

clair timeout: 10ms docker timeout: 10ms no whitelist file Analysising 14 layers GET results from Clair API v1 Found 734 vulnerabilities Unknown: 223 Negligible: 345 Low: 184 Medium: 2						
Severity	Name	FeatureName	FeatureVersion	FixedBy	Description	Link
Medium	CVE-2009-3546	libwnf	0.2.6.4-18.6		The <code>_gdGetColors</code> function in <code>gd_gdvc</code> in PHP 5.2.11 and 5.3.x before 5.3.1, and the GD Graphics Library 2.x, does not properly verify a certain <code>colorstotal</code> structure member, which might allow remote attackers to conduct buffer overflow or buffer over-read attacks via a crafted GD file, a different vulnerability than CVE-2009-3293. NOTE: some of these details are obtained from third party information.	<a href="https://security-tracker.debian.org/tracker/CVE-2009-3546">https://security-tracker.debian.org/tracker/CVE-2009-3546</a>
Medium	CVE-2007-3996	libwnf	0.2.6.4-18.6		Multiple integer overflows in <code>libgd</code> in PHP before 5.2.4 allow remote attackers to cause a denial of service (application crash) and possibly execute arbitrary code via a large (1) <code>srcW</code> or (2) <code>srcH</code> value to the (a) <code>gdImageCopyResized</code> function, or a large (3) <code>sy</code> ( <code>height</code> ) or (4) <code>sx</code> ( <code>width</code> ) value to the (b) <code>gdImageCreate</code> or the (c) <code>gdImageCreateTrueColor</code> function.	<a href="https://security-tracker.debian.org/tracker/CVE-2007-3996">https://security-tracker.debian.org/tracker/CVE-2007-3996</a>

# Kubernetes – Core Services – Network Security



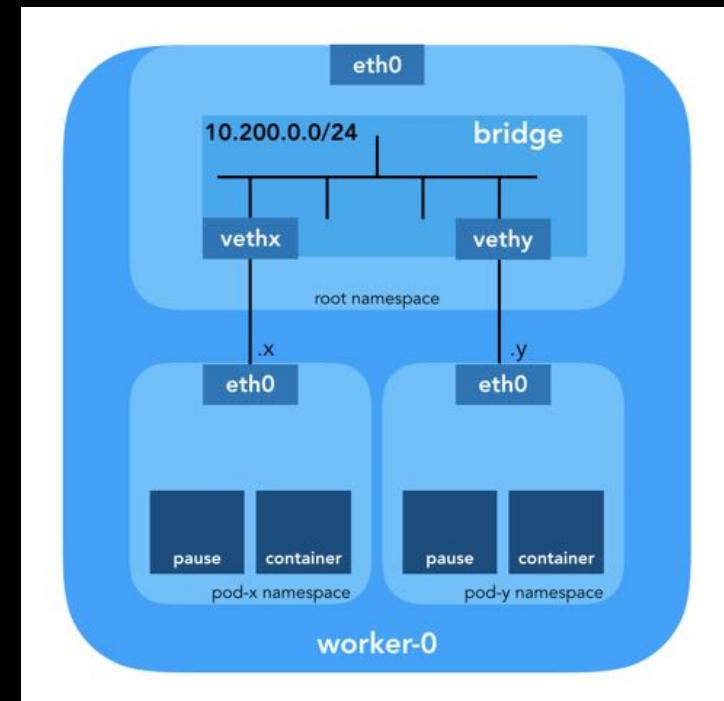
# Kubernetes – Core Services – Network Security



## CNI – Common Network Interface

Provides a rich set of security enforcement capabilities running on top of a highly scalable and efficient virtual network

- **Calico**  
Virtual networking and network security for containers, VMs, and bare metal services.
- **Open vSwitch**  
Production quality, multilayer virtual switch
- ...

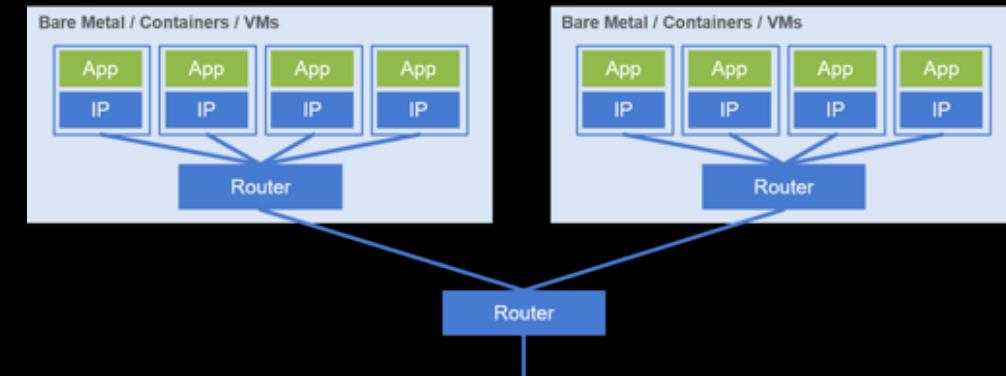


# Kubernetes – Core Services – Network Security

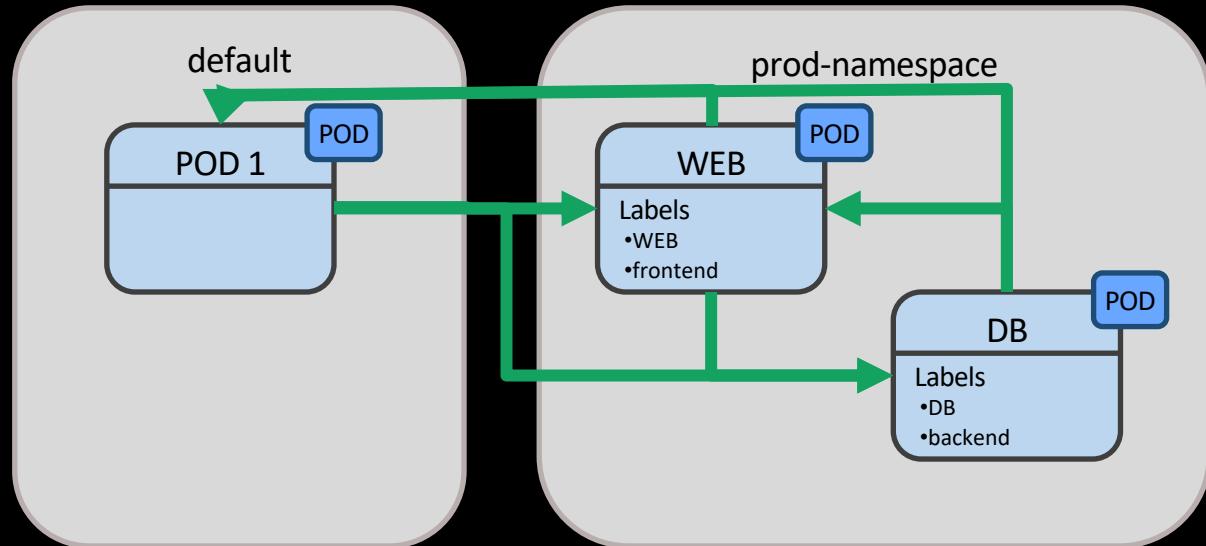


Operates at **Layer 3**, which is the network layer

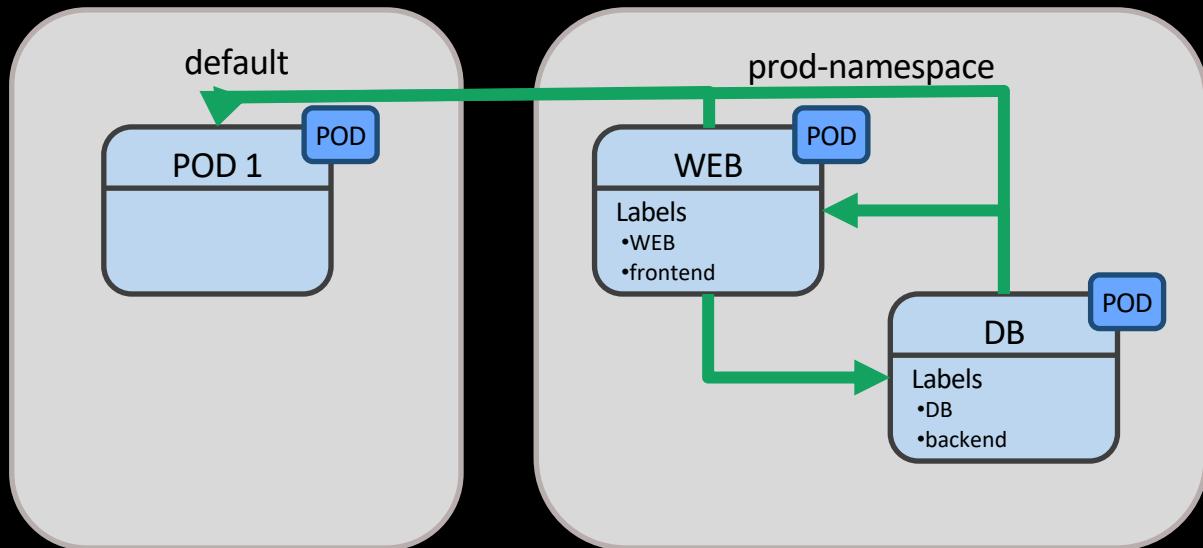
- Has the advantage of being **universal** (DNS, SQL, real-time streaming, ...)
- Can extend beyond the service mesh (including to **bare metal or VM** endpoints not under the control of Kubernetes).
- Calico's policy is enforced at the host node, outside the network namespace of the guest pods.
- Based on **iptables**, which are packet filters implemented in the standard Linux kernel, it is extremely fast.



# Kubernetes – Core Services – Network Security

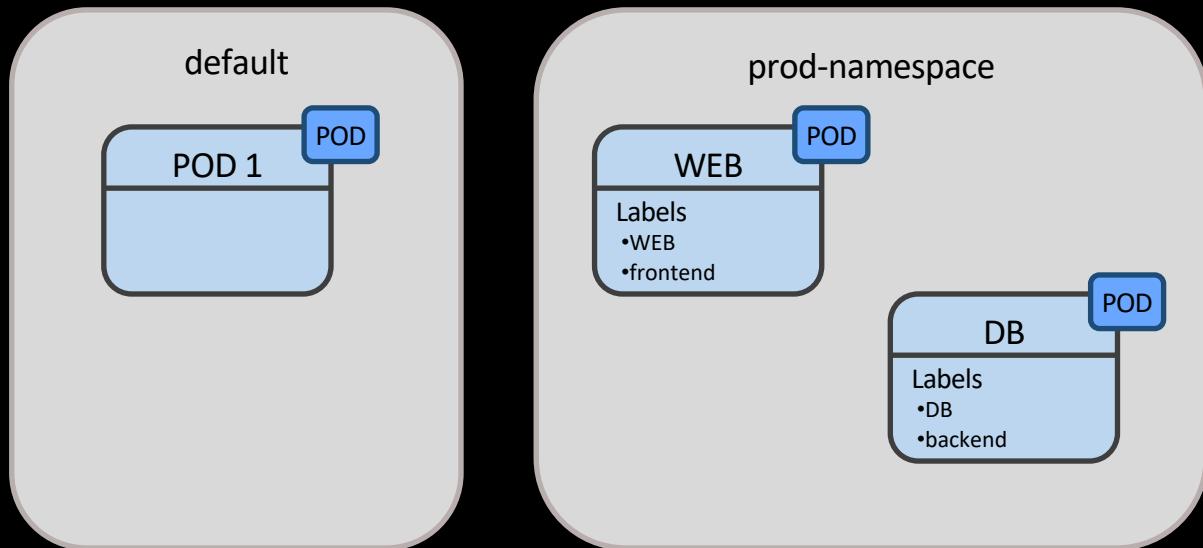


# Kubernetes – Core Services – Network Security



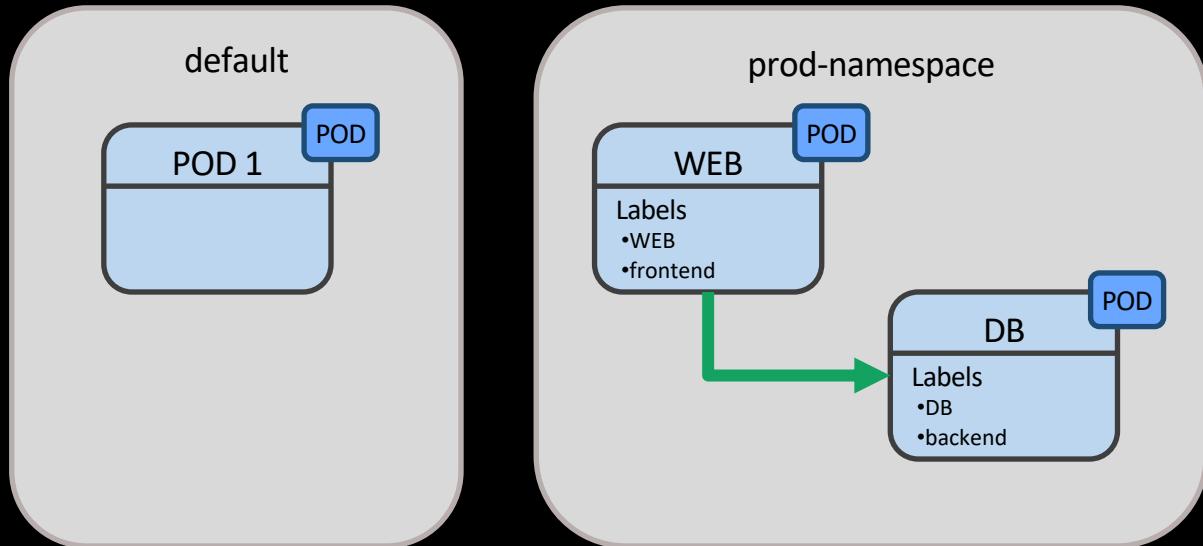
```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: default-deny
  namespace: demo-namespace
spec:
  podSelector:
    matchLabels: {}
```

# Kubernetes – Core Services – Network Security



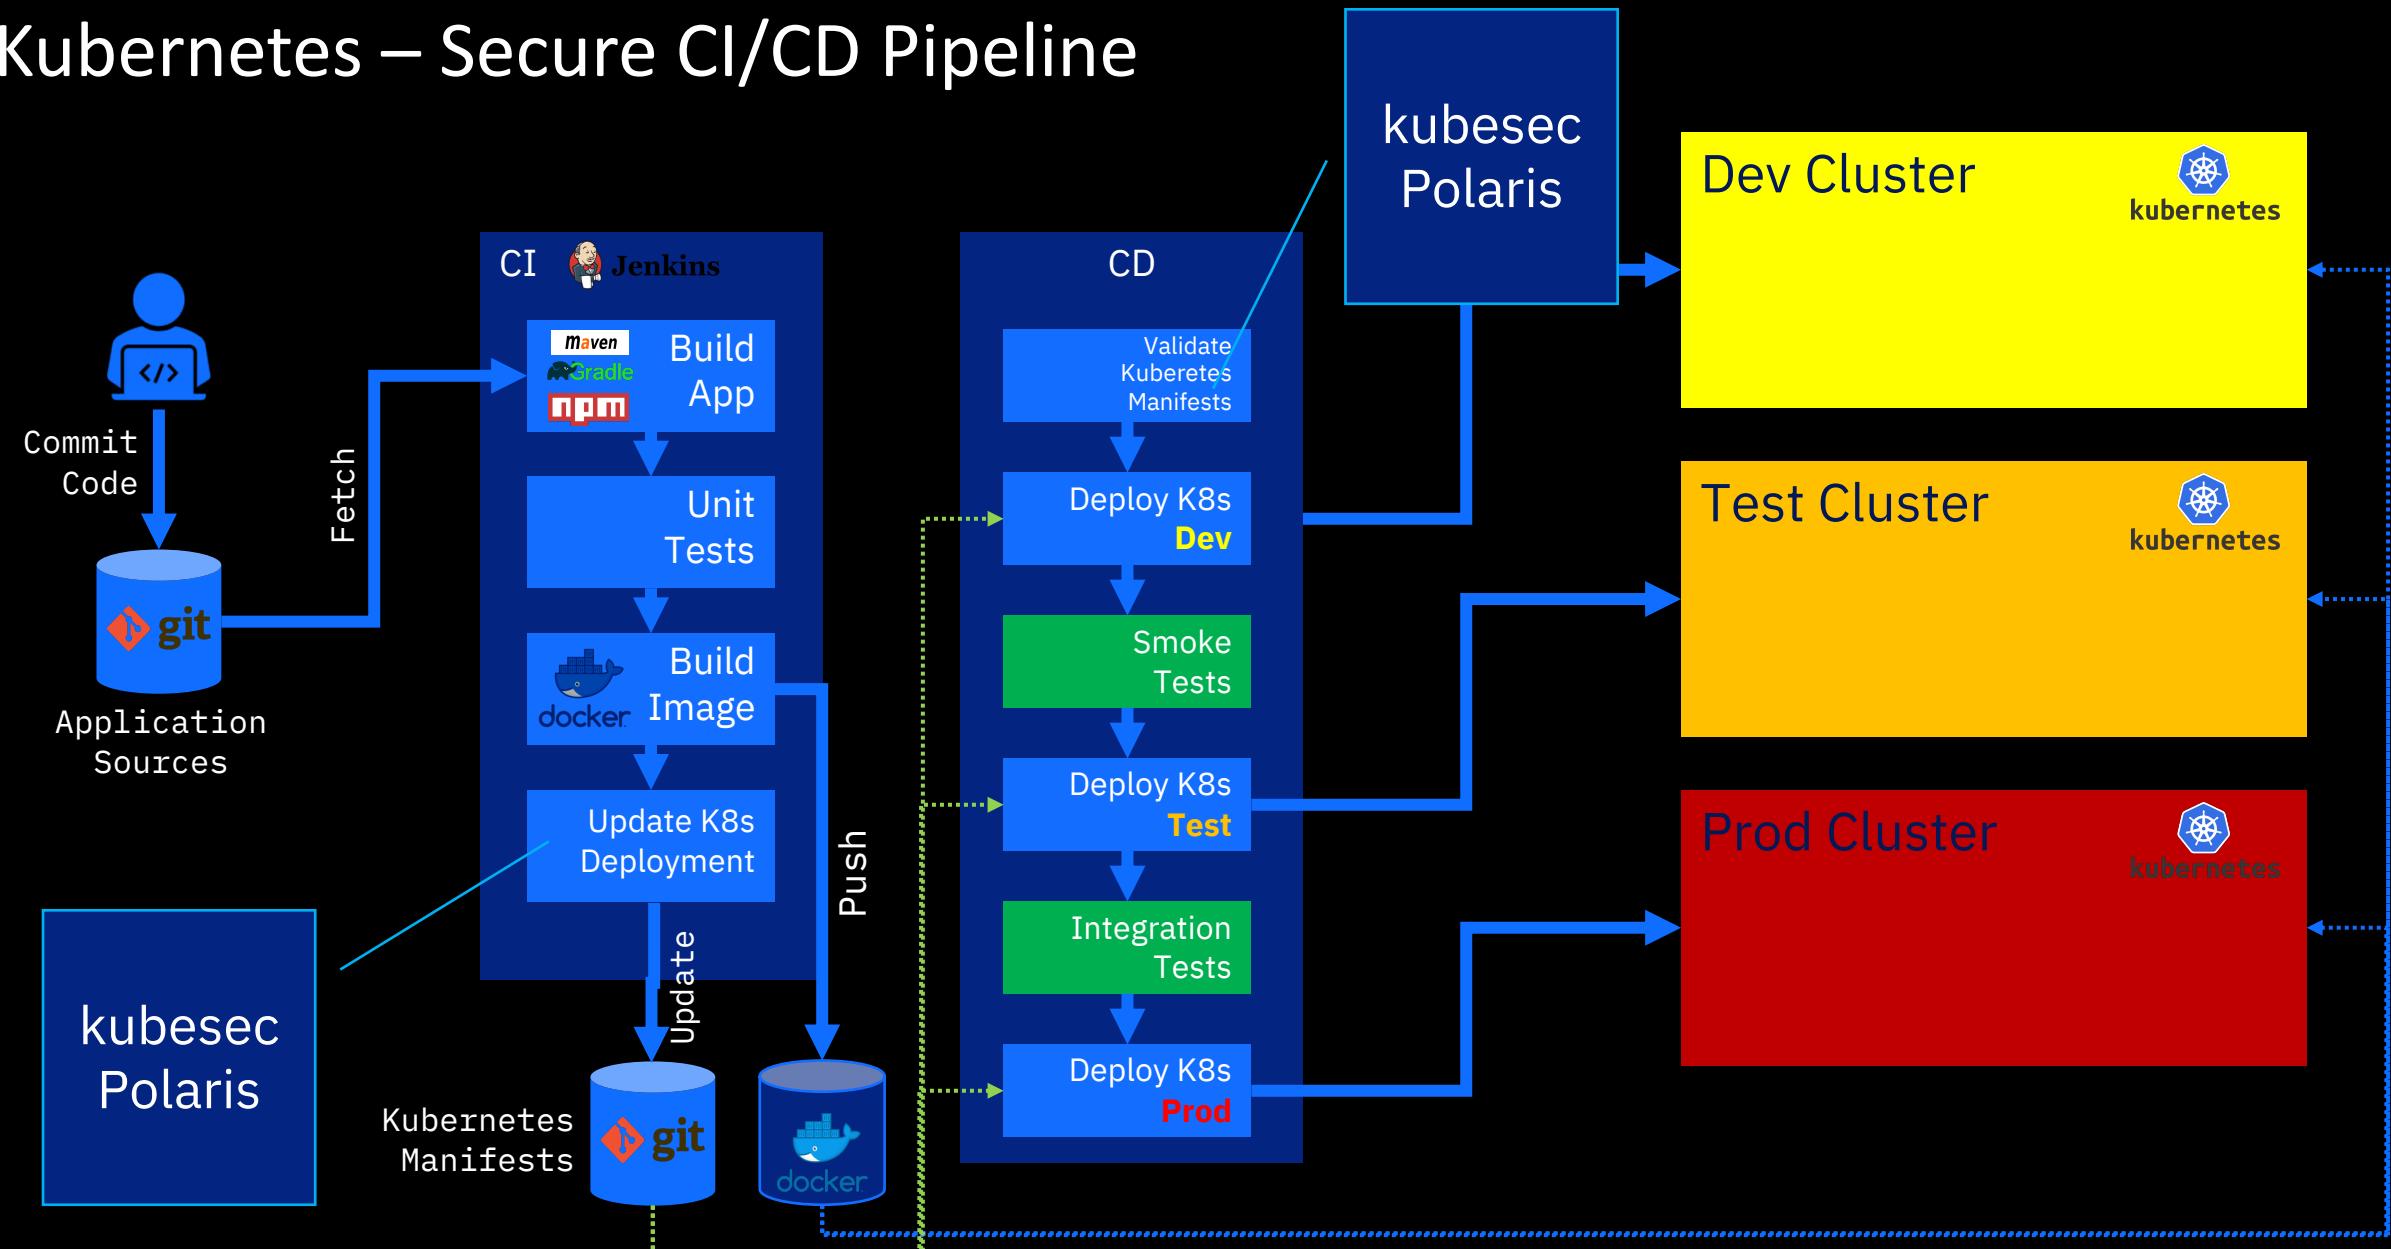
```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: default-deny-all
spec:
  podSelector:
    matchLabels: {}
```

# Kubernetes – Core Services – Network Security

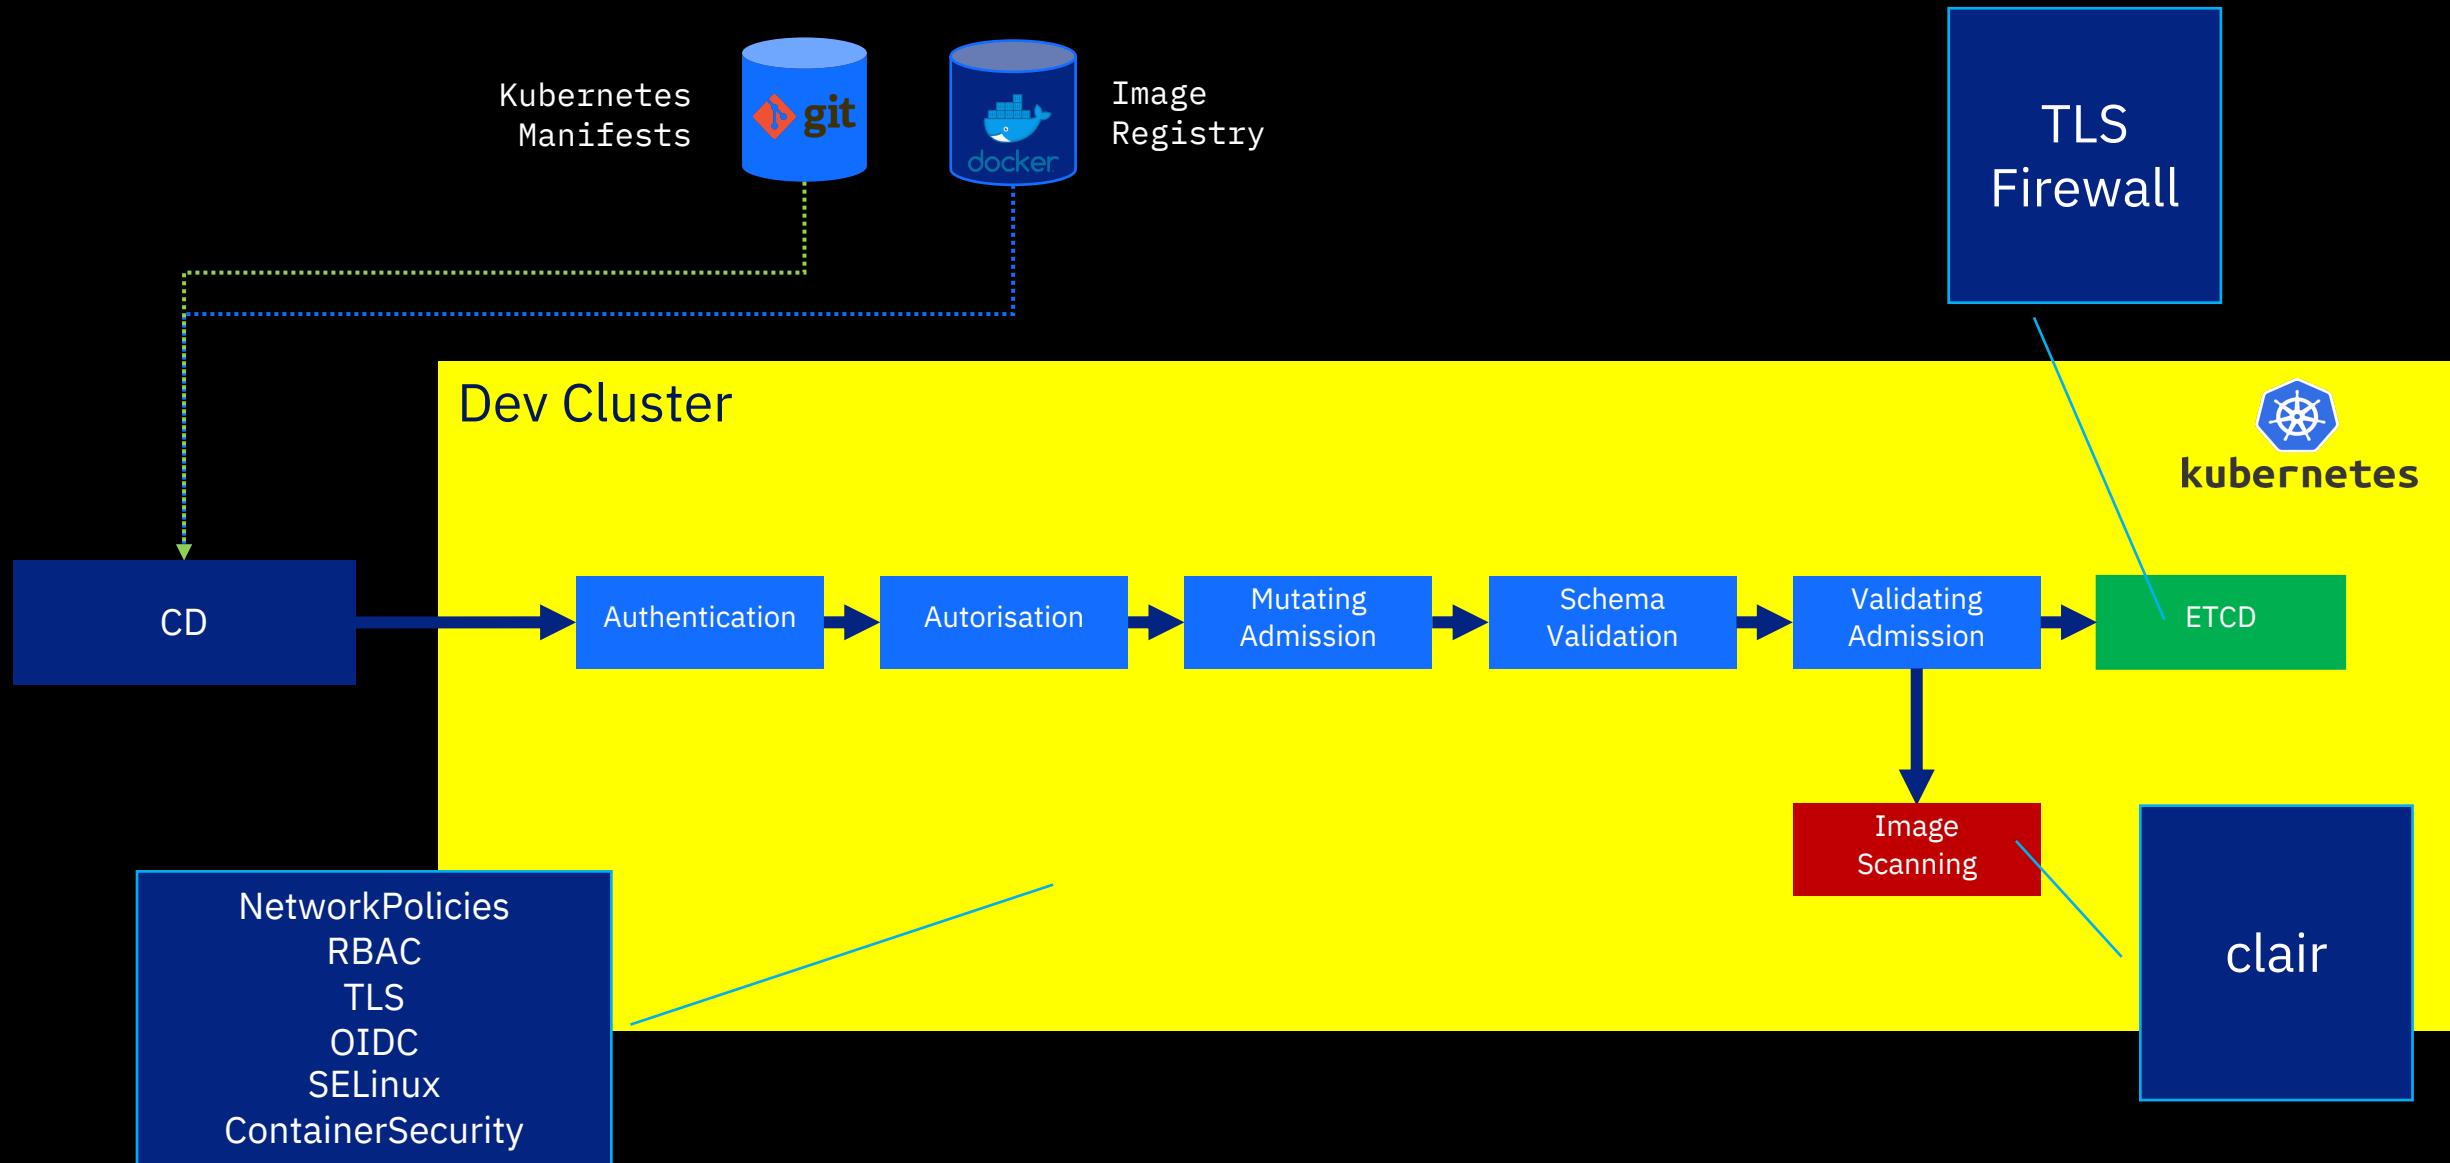


```
kind: NetworkPolicy
metadata:
  name: access-frontend-backend
  namespace: prod-namespace
spec:
  podSelector:
    matchLabels:
      run: DB
  ingress:
    - from:
        - podSelector:
            matchLabels:
              run: WEB
```

# Kubernetes – Secure CI/CD Pipeline



# Kubernetes – Secure CI/CD Pipeline



# Kubernetes – Tip 1

## Use **Namespaces** Liberally

**“Namespaces are cheap.”**

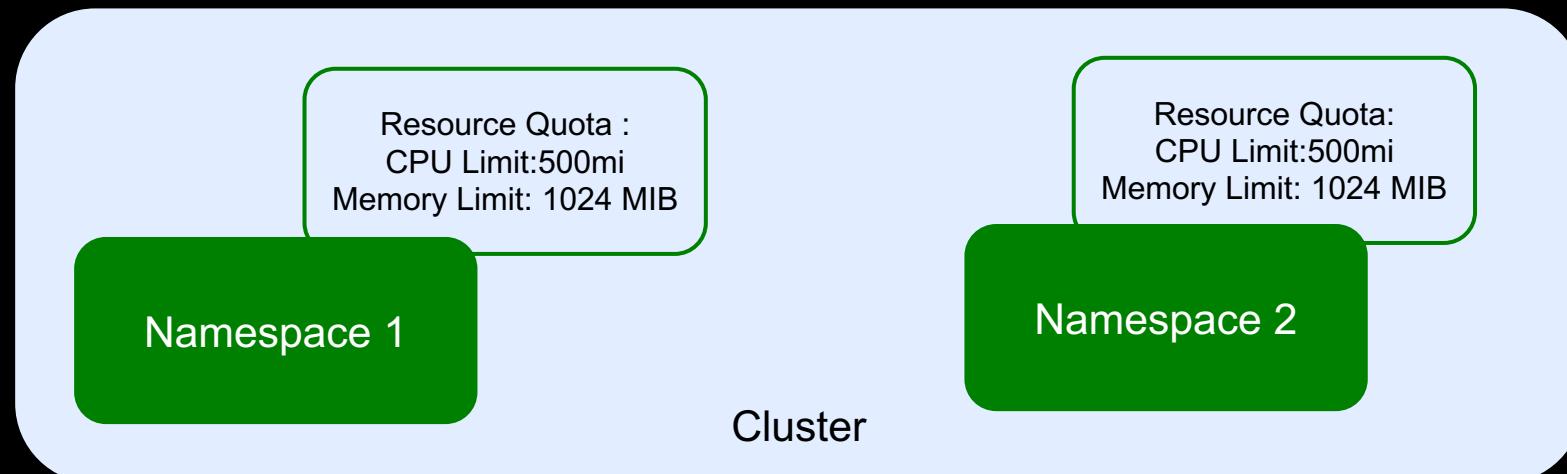
Use them to separate things like infrastructure tooling and applications.

This allows you to restrict access easily using **RBAC** and to limit the scope of applications. It will also make your **network policy** creation easier when you decide to do it.

# Kubernetes – Tip 2

## Set resource quota

Once quota in a namespace for compute resources set, the users are forced to set requests or limits for those values. Avoids starving of Cluster.



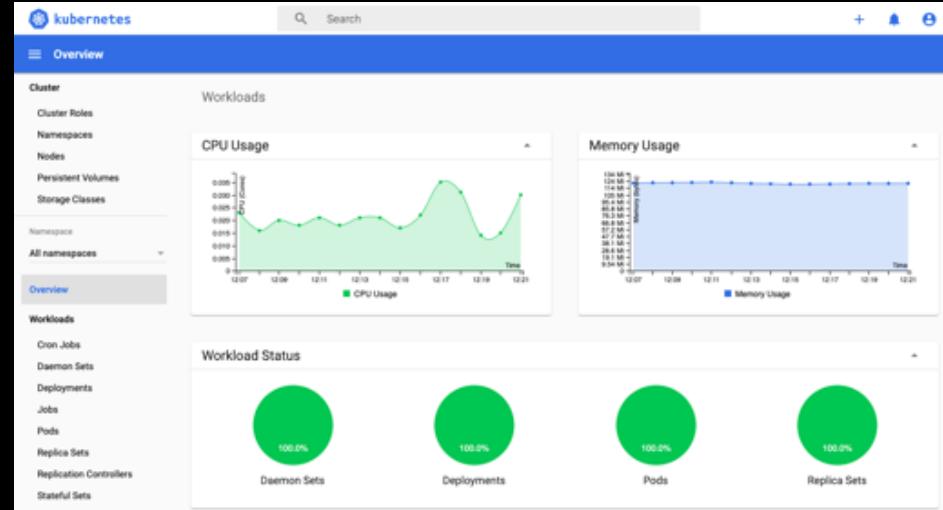
# Kubernetes – Tip 3

## Secure Kubernetes Dashboard

The Kubernetes Dashboard has often been used for attacks.  
Easy to deploy with high privileges.

Best practice is to:

- Allow only authenticated access
- Use RBAC
- Ensure Dashboard service account has limited access
- Don't expose to the internet (use kubectl proxy for local access)



# Kubernetes – Tip 4

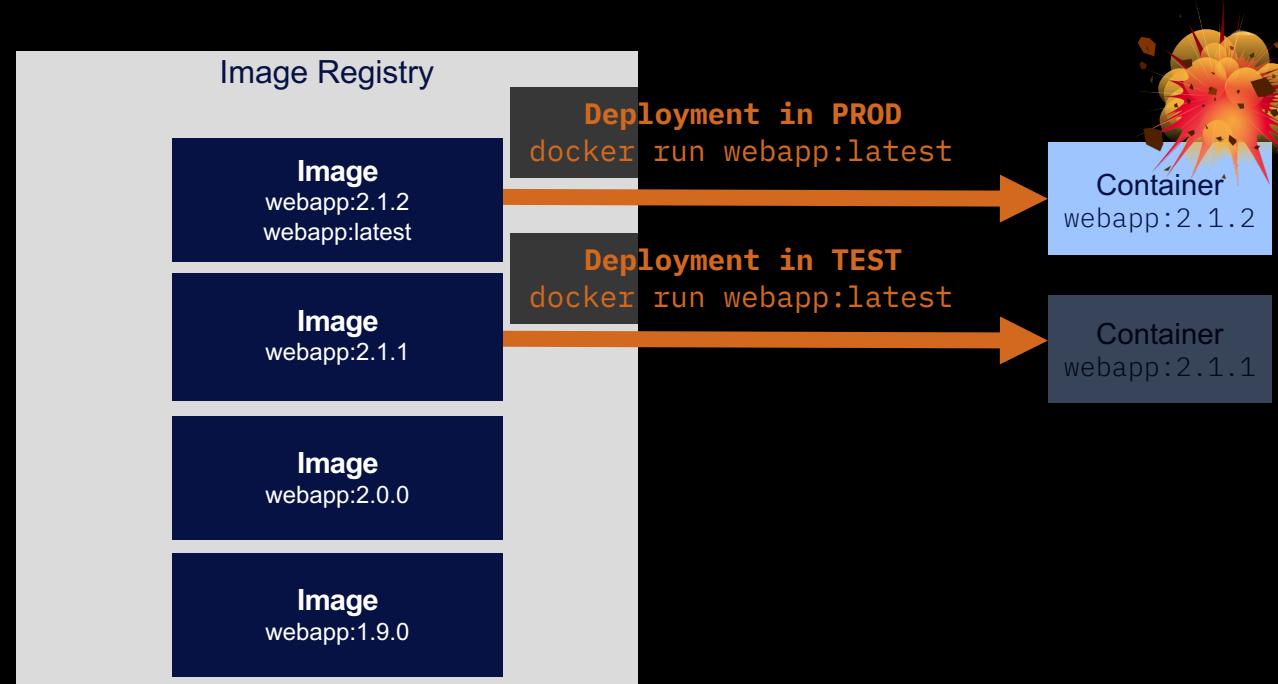
## Avoid Rolling Tags at Any Cost

Avoid commands like `docker pull myregistry/myimage:latest`.

This “latest” is an example of a rolling tag (i.e. a tag that will point to different images over time).

If you want your deployments to be secure and maintainable, make sure that you use immutable images (for example: “`myregistry/myimage:1.1.2`”).

With this approach, every time you deploy or scale, you know what image you are using and you will have the guarantee that the deployed image has been tested and validated.



# Kubernetes – Tip 5

## **Don't run containers as the root user**

By default, containers run as the root user inside the container. This is easy to avoid using the pod specification to set a high-numbered UID.

```
spec:  
  securityContext:  
    runAsUser: 10324
```

# Kubernetes – Tip 6

## **Don't run containers as privileged**

Running a container or pod as privileged gives it the ability to make modifications to the host. This is a large security issue that is easy to mitigate by just not doing it.

# Kubernetes – Tip 7

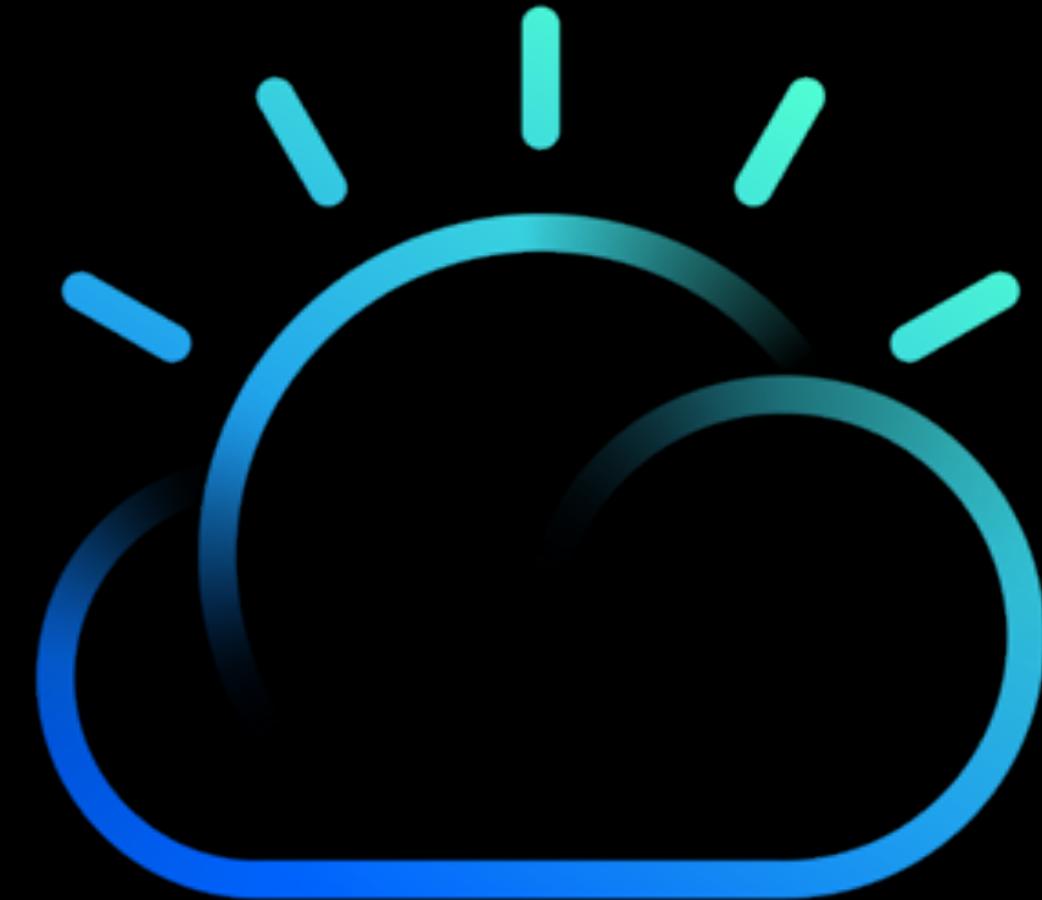
## **Don't use the default list of capabilities**

Docker runs containers with a significant set of Linux capabilities by default, many of which your app might not require.

The following config will drop all Linux capabilities by default, allowing you to add only the specific capabilities your app actually needs:

```
spec:  
  containers:  
    - name: foo  
      securityContext:  
        capabilities:  
          drop:  
            - ALL
```

QUESTIONS?



# The Journey to Cloud **Mesh Networking**



IBM Cloud

# The trade off

**Improved delivery velocity**  
in exchange for  
**increased operational complexity**



# Common DevOps Challenge 1

How do I **roll out** a newer version of my microservice  
**without down time**?

How do I **ensure traffic** continue to go to the current version  
before the newer version is tested and ready?

# Common DevOps Challenge 2

How do **canary testing**?

How do I proceed to a **full rollout** after satisfactory testing of the new version?

# Common DevOps Challenge 3

How do I do **A/B testing**?

- Release a new version to a subset of users in a precise way

I want to leverage crowdsourced testing. How do I **test** the new version **with a subset of users**?

I have **launched B in the dark**, but how can I keep B to myself or a small testing group?

# Other common DevOps Challenges

4. Things don't always go correctly in production... How do I **inject fault** to my microservices to prepare myself?
5. My services can only **handle certain rate**, how can I limit rate for some of my services?

# Other common DevOps Challenges

6. I need to **view and monitor** what is going on with each of my services when crisis arises.
7. How can I **secure my services**.

# Service Mesh

**Dedicated infrastructure layer  
to make  
service-to-service communication  
fast, safe and reliable**

# Istio



A service mesh designed to connect, manage and secure micro services

The image shows two news snippets side-by-side. The left snippet is from Forbes, dated May 25, 2017, with 3,859 views. It's titled "Google, IBM And Lyft Want To Simplify Microservices Management With Istio". The right snippet is from ZDNet, dated May 24, 2017, with 1,100 views. It's titled "Google, IBM, and Lyft launch open source project Istio". Both snippets mention Istio as a vendor-neutral way to connect, secure, manage, and monitor microservices.

The image shows a Google Cloud blog post titled "Istio: a modern approach to developing and managing microservices" by Varun Talwar, Product Manager, Cloud Service Platform, published on May 24, 2017. The post discusses the alpha release of Istio and its purpose of providing a uniform way to help connect, secure, manage and monitor microservices.

Launched in May 2017 by Google, Lyft and IBM

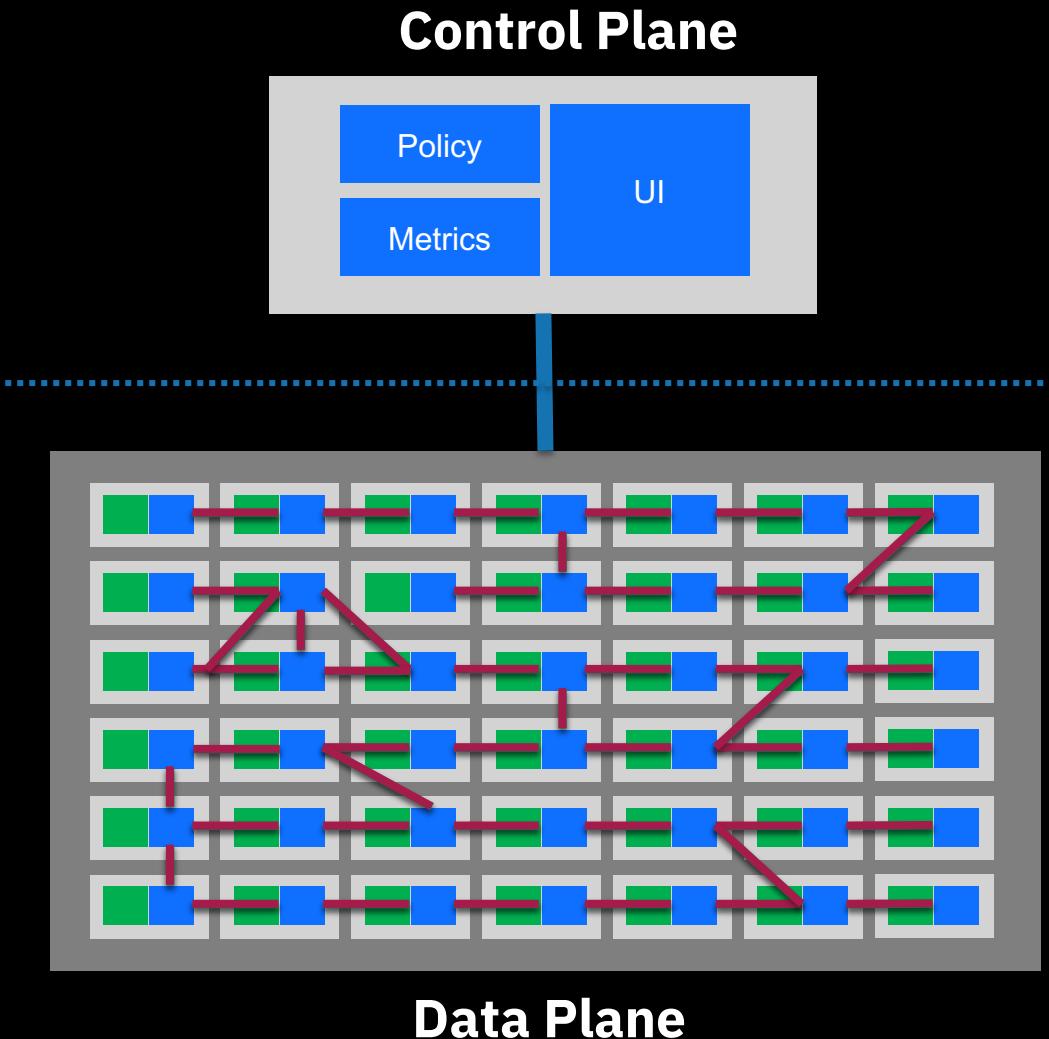
Open Source

Zero Code Changes

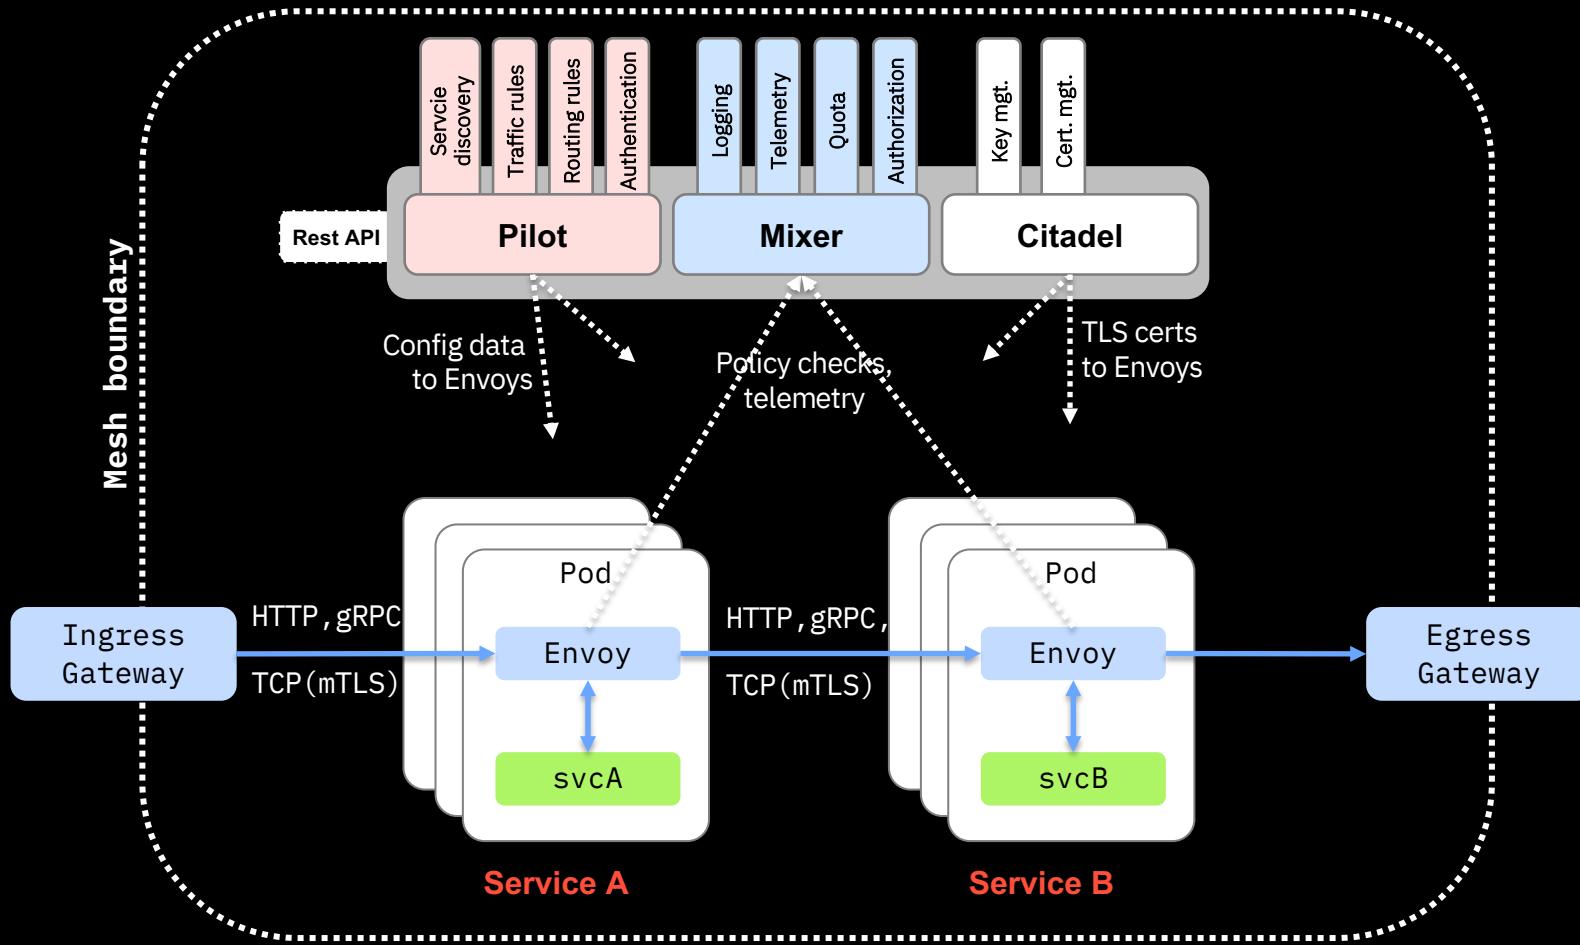
# Service Mesh

## Tent poles:

- Control over request routing (CI/CD release patterns)
- Resiliency (retries, timeouts, etc)
- Cascading failure prevention (circuit breaking)
- Load balancing
- Provide and manage TLS termination between endpoints
- Metrics to provide instrumentation at the service-to-service layer



# ISTIO - Architecture

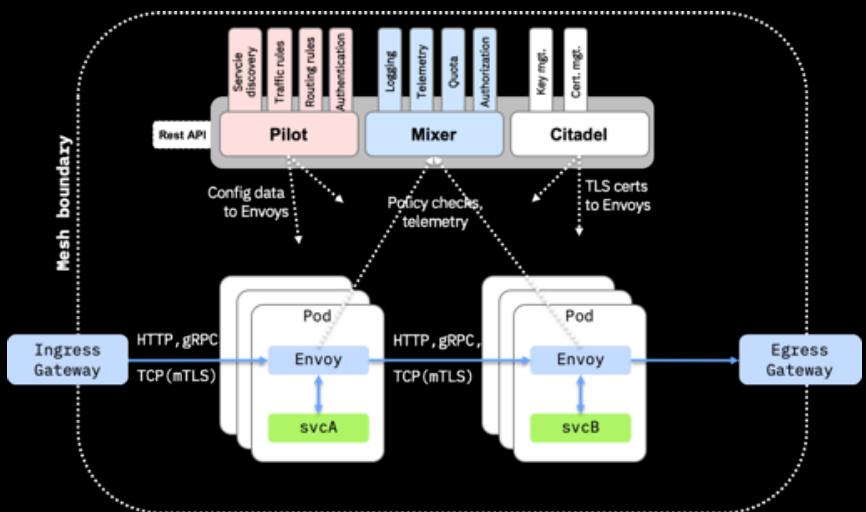


●—● Istio control plane traffic. Request routing rules, resilience configuration (circuit breakers, timeouts, retries), policies (ACLs, rate limits, auth), and metrics/reports from proxies.

●—● User/application traffic. HTTP/1.1, HTTP/2, gRPC, TCP with or without TLS

## Operates at **Layer 7**

- policies can be applied based on virtual host, URL, or other HTTP headers.
- Flexibility in processing.
- Allows it to be distributed



## Custom resource definitions

## Ingress Configuration

```
kind: Gateway
metadata:
  name: helloworld-gateway
spec:
  selector:
    istio: ingressgateway
  servers:
    - hosts:
        - myapp.demo.com
      port:
        name: http
        number: 80
      protocol: HTTP
```



<https://myapp.demo.com/demo>

## URL Routing

```
kind: VirtualService
metadata:
  name: helloworld
spec:
  hosts:
    - myapp.demo.com
  gateways:
    - helloworld-gateway
  http:
    - match:
        - uri:
            exact: /demo
      route:
        - destination:
            host: helloworld
            subset: v1
            weight: 90
        - destination:
            host: helloworld
            subset: v2
            weight: 10
```

```
kind: DestinationRule
metadata:
  name: helloworld-destination
spec:
  host: helloworld
  subsets:
    - name: v1
      labels:
        version: v1
    - name: v2
      labels:
        version: v2
```

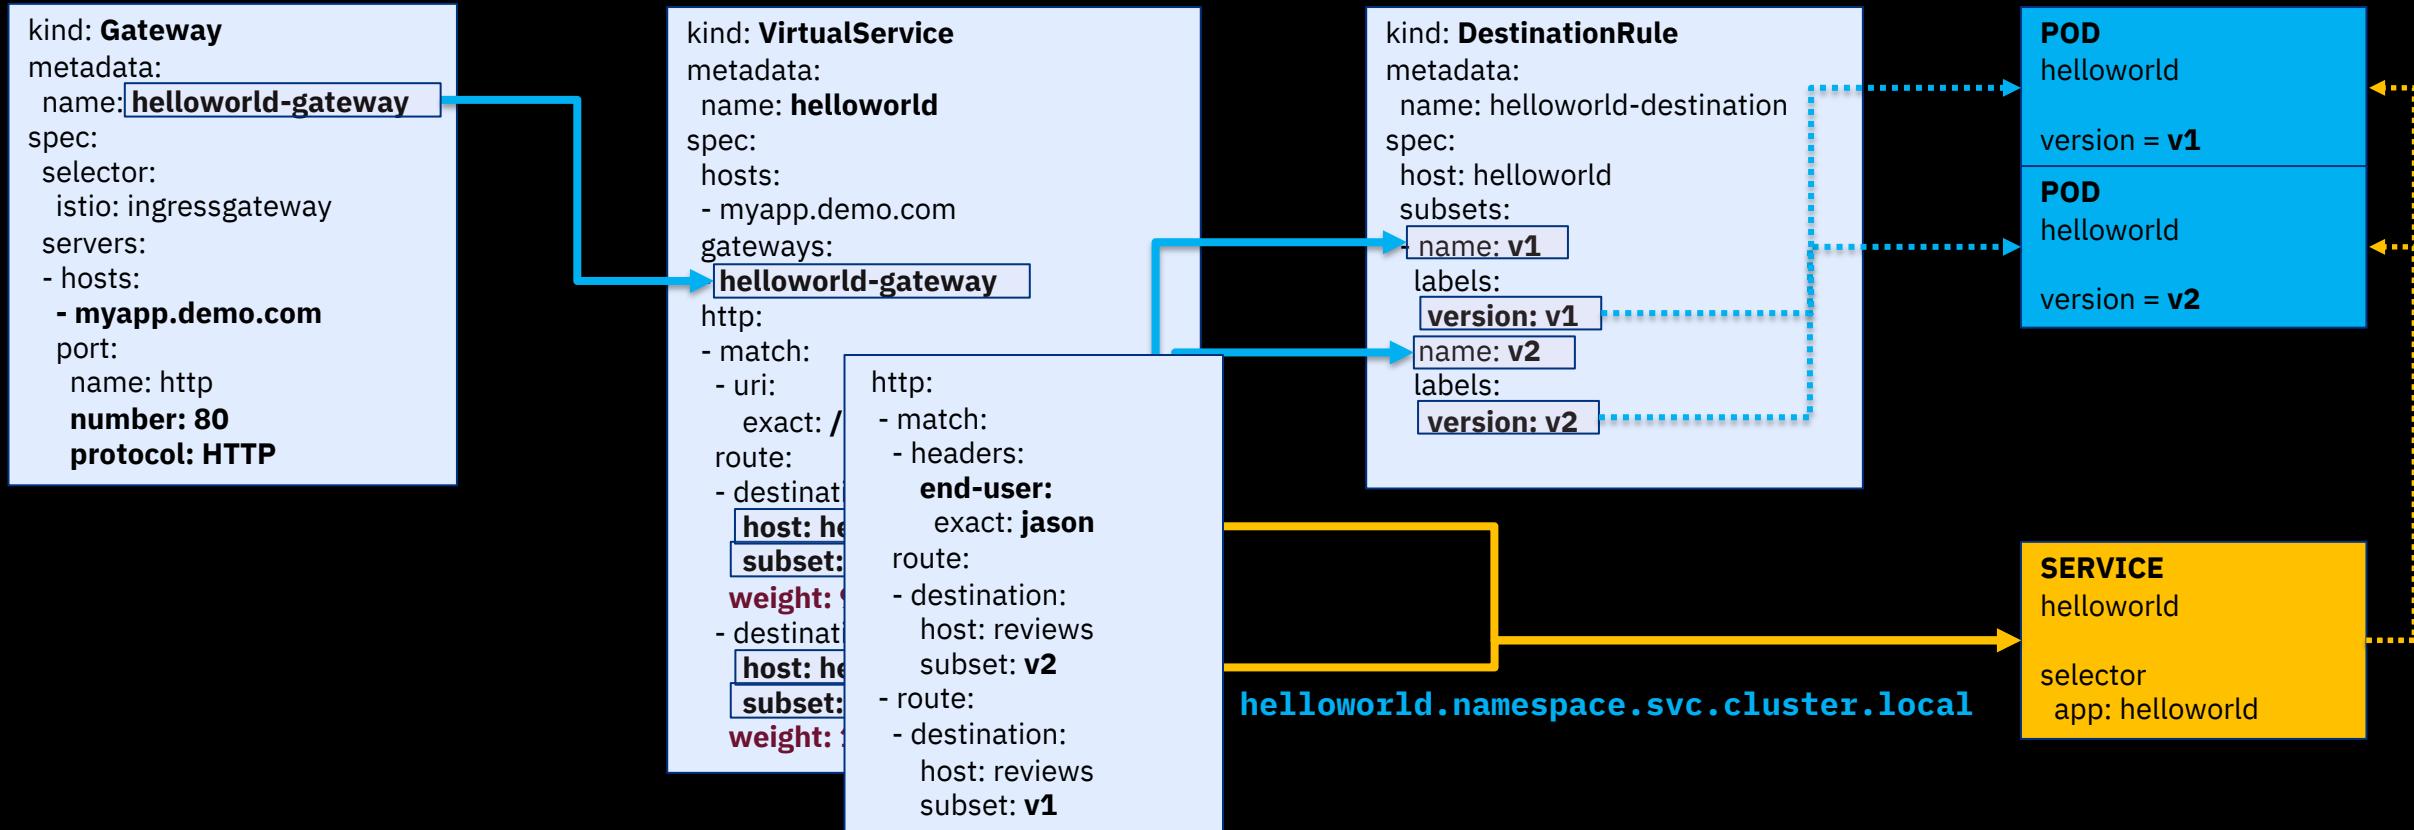
POD  
helloworld  
version = v1

POD  
helloworld  
version = v2

helloworld.namespace.svc.cluster.local

SERVICE  
helloworld  
selector  
app: helloworld

## Custom resource definitions

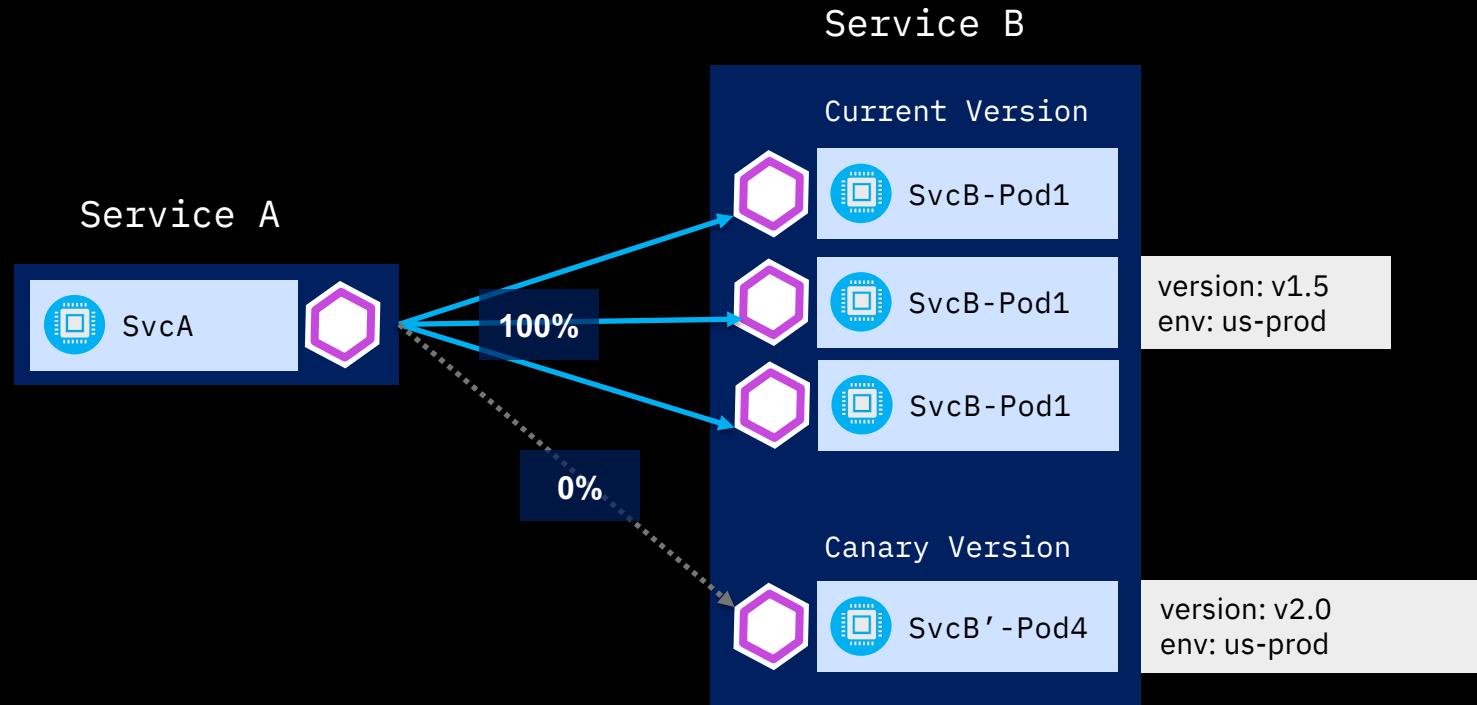


# Addressing DevOps Challenges



#	CHALLENGE	ISTIO SOLUTION
<b>CHALLENGE 1</b>	ROLL OUT NEW VERSION WITHOUT DOWNTIME OR CHANGING CODE	TRAFFIC CONTROL
<b>CHALLENGE 2</b>	HOW TO DO CANARY TESTING	TRAFFIC SPLITTING
<b>CHALLENGE 3</b>	HOW TO DO A/B TESTING	TRAFFIC STEERING
<b>CHALLENGE 4</b>	THINGS DON'T ALWAYS GO CORRECTLY IN PRODUCTION...	TRAFFIC MIRRORING RESILIENCY RESILIENCY TESTING
<b>CHALLENGE 5</b>	HOW CAN I LIMIT RATE FOR SOME OF MY SERVICES?	RATE LIMITING
<b>CHALLENGE 6</b>	I NEED TO VIEW AND MONITOR WHAT IS GOING ON WHEN CRISIS ARISES	TELEMETRY
<b>CHALLENGE 7</b>	HOW CAN I SECURE MY SERVICES?	AUTHENTICATION AUTHORIZATION CALICO

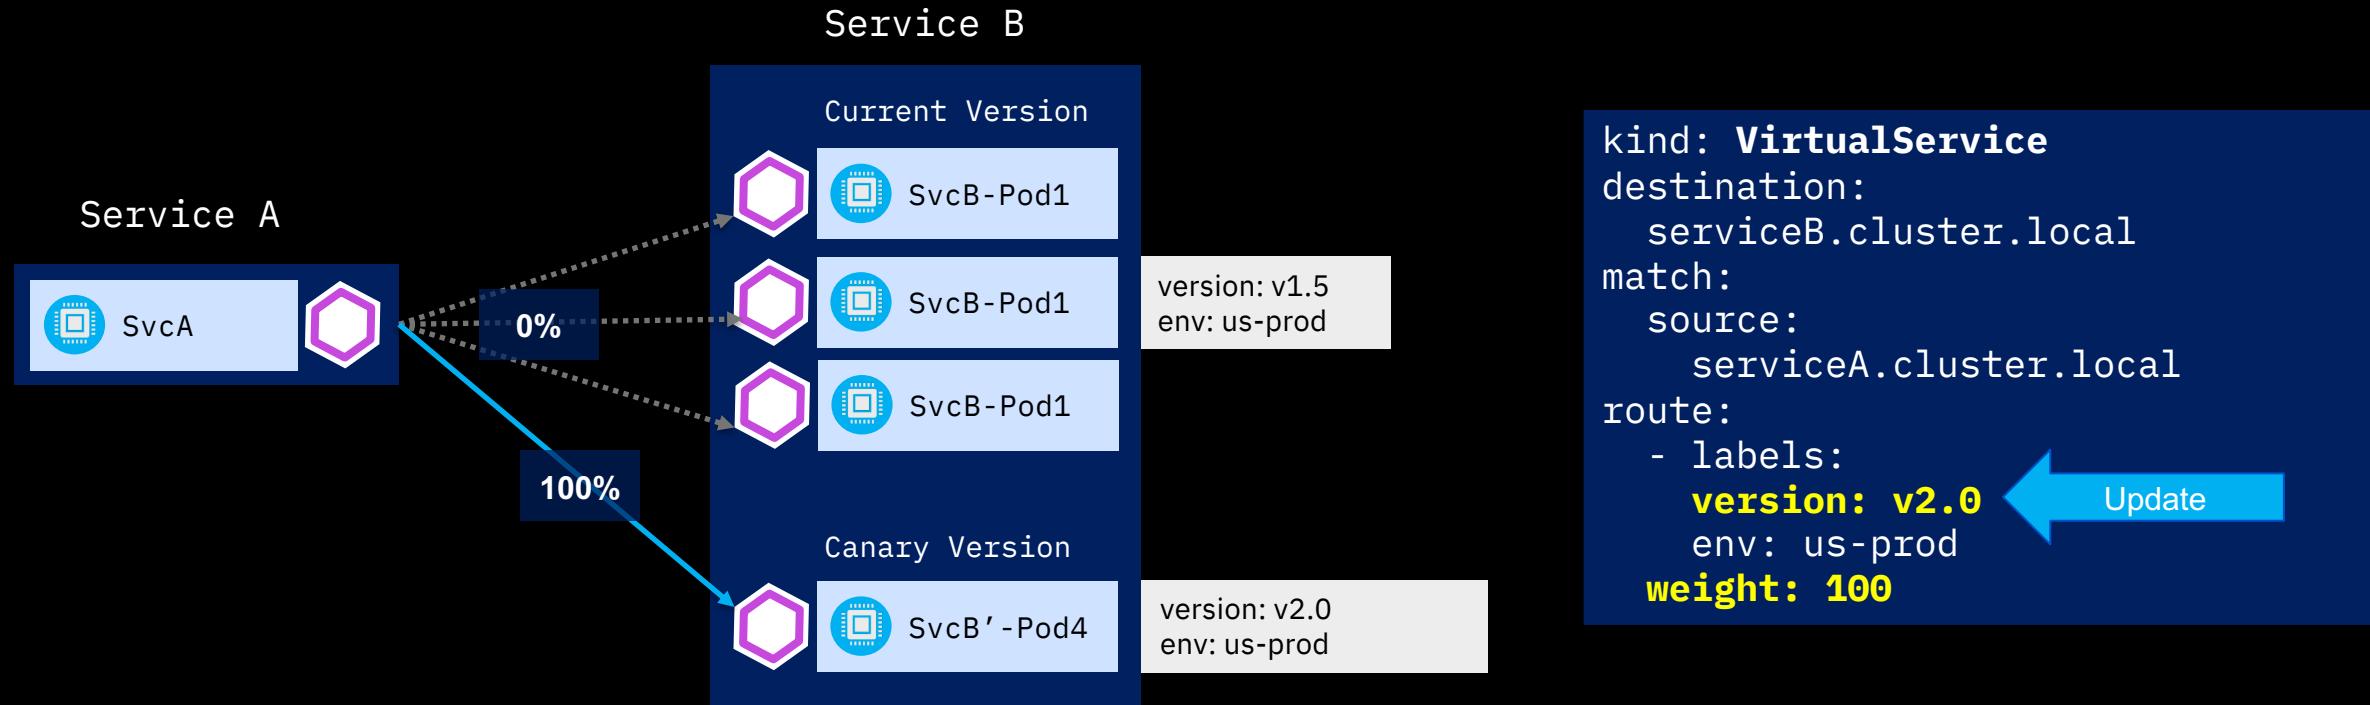
# Traffic Control



```
kind: VirtualService
destination:
  serviceB.cluster.local
match:
  source:
    serviceA.cluster.local
route:
  - labels:
      version: v1.5
    env: us-prod
    weight: 100
```

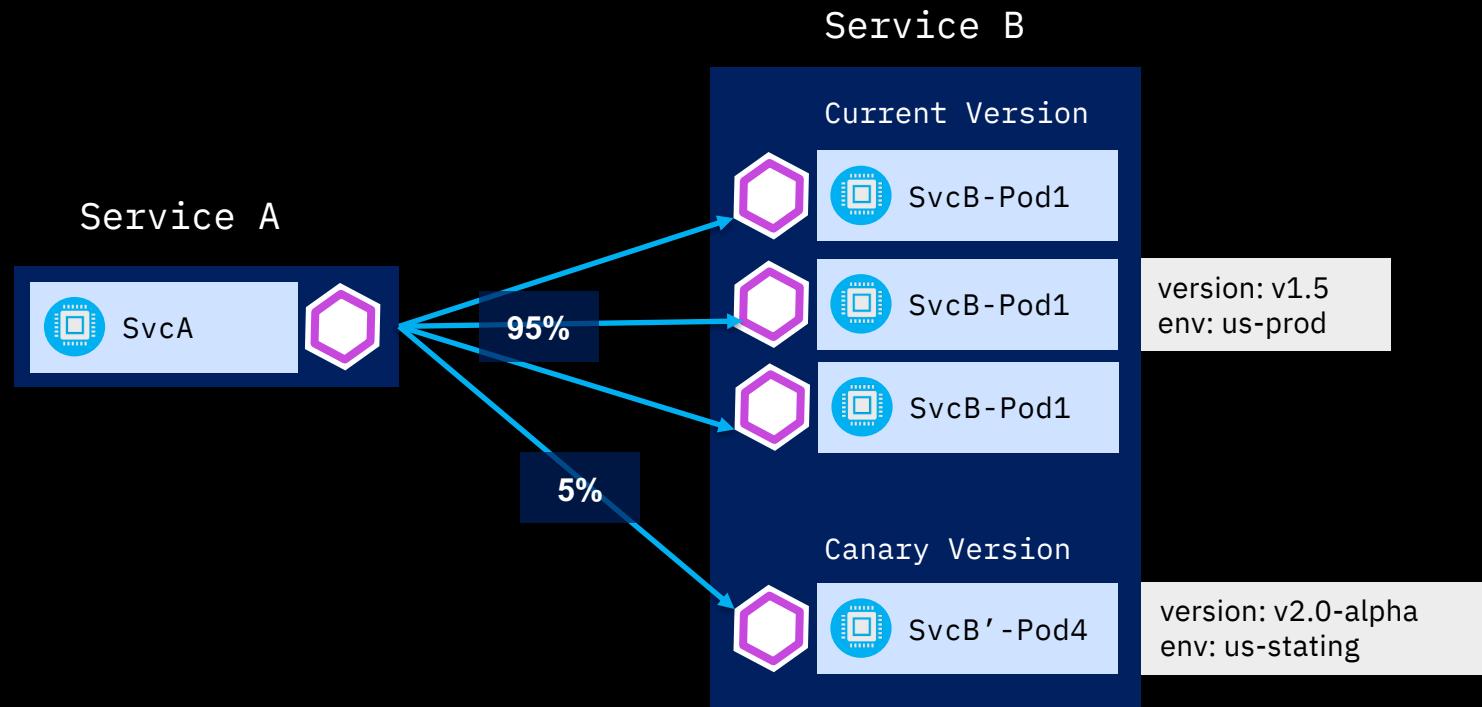
**CHALLENGE 1**  
ROLL OUT NEW VERSION WITHOUT DOWNTIME OR CHANGING CODE

# Traffic Control



**CHALLENGE 1**  
ROLL OUT NEW VERSION WITHOUT DOWNTIME OR CHANGING CODE

# Traffic Splitting

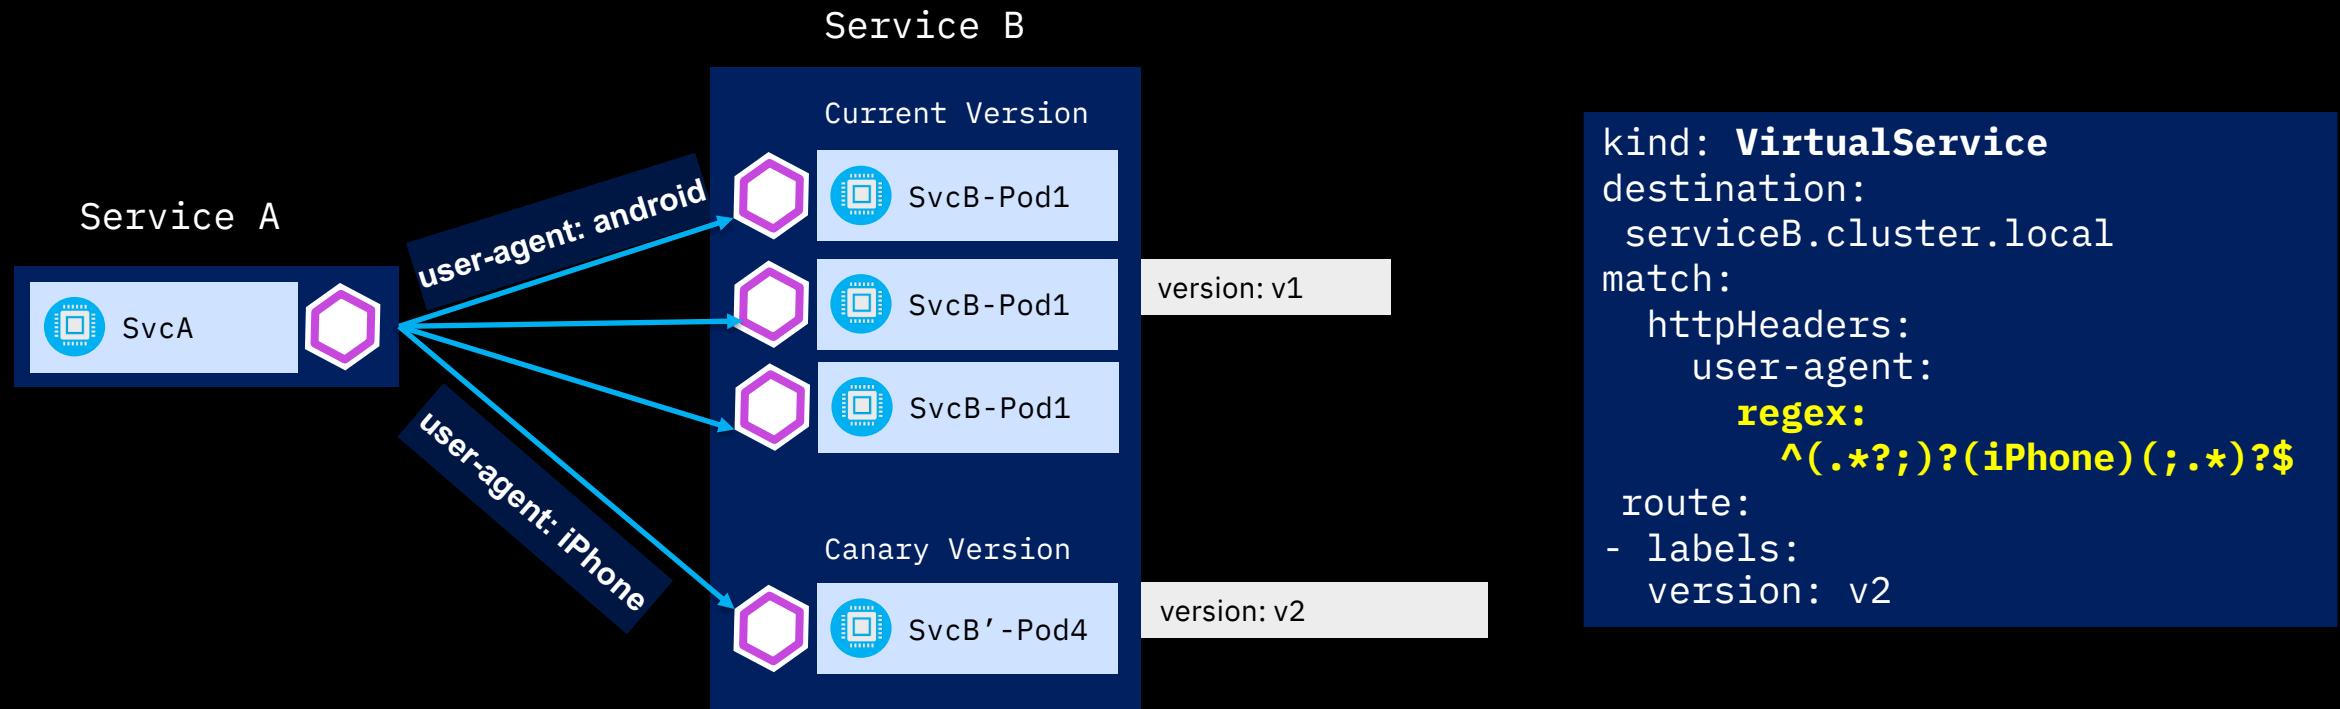


```
kind: VirtualService
destination:
  serviceB.cluster.local
match:
  source:
    serviceA.cluster.local
route:
  - labels:
      version: v1.5
      env: us-prod
weight: 95
  - labels:
      version: v2.0-alpha
      env: us-staging
weight: 5
```

## CHALLENGE 2 HOW TO DO CANARY TESTING

Routing not based on the request content.  
Staged rollouts with %-based traffic splits.

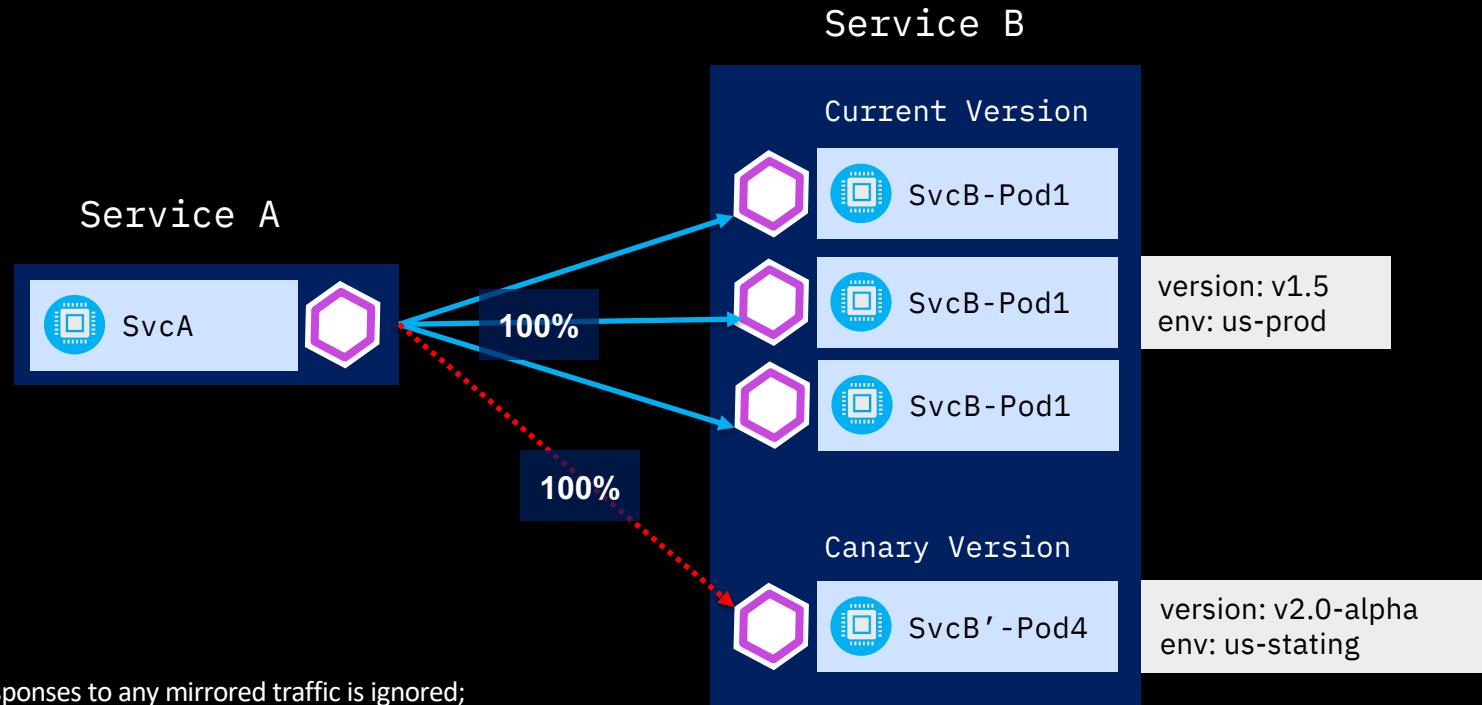
# Traffic Steering



Routing based on the request content

**CHALLENGE 3**  
**HOW TO DO A/B TESTING**

# Traffic Mirroring



```
kind: VirtualService
destination:
  serviceB.cluster.local
match:
  source:
    serviceA.cluster.local
route:
  - labels:
      version: v1.5
      env: us-prod
    weight: 100
  - labels:
      version: v2.0-alpha
      env: us-staging
    weight: 0
    mirror:
      name: httpbin
      labels:
        version: v2.0-alpha
        env: us-staging
```

**CHALLENGE 4**  
THINGS DON'T ALWAYS GO CORRECTLY IN PRODUCTION...

# Resiliency

Istio adds fault tolerance to your application without any changes to code

```
// Circuit breakers
destination: serviceB.example.cluster.local
policy:
- labels:
  version: v1
circuitBreaker:
  simpleCb:
    maxConnections: 100
    httpMaxRequests: 1000
    httpMaxRequestsPerConnection: 10
    httpConsecutiveErrors: 7
    sleepWindow: 15m
    httpDetectionInterval: 5m
```

## Resilience features

- ❖ Timeouts
- ❖ Retries with timeout budget
- ❖ Circuit breakers
- ❖ Health checks
- ❖ AZ-aware load balancing w/ automatic failover
- ❖ Control connection pool size and request load

**CHALLENGE 4**  
**THINGS DON'T ALWAYS GO CORRECTLY IN PRODUCTION...**

# Rate limiting

Istio protects your application from rogue actors by imposing ratelimits

## Quotas:

```
- name: requestcount.quota.istio-system
  maxAmount: 5000
  validDuration: 1s
  overrides:
    - dimensions:
        destination: ratings
        source: reviews
        sourceVersion: v3
        maxAmount: 1
        validDuration: 1s
    - dimensions:
        destination: ratings
        maxAmount: 100
        validDuration: 1s
```

## Rate limit

- ❖ Configurable limits with overrides
- ❖ Multiple rate limiting backends
- ❖ Conditional rate limiting

**CHALLENGE 5**  
**HOW CAN I LIMIT RATE FOR SOME OF MY SERVICES?**

# Telemetry

Monitoring & tracing should not be an afterthought in the infrastructure

## Goals

- Metrics without instrumenting apps
- Consistent metrics across fleet
- Trace flow of requests across services
- Portable across metric backend providers



**CHALLENGE 6**  
**I NEED TO VIEW WHAT IS GOING ON WHEN CRISIS ARISES**

# Kiali

**Kiali** (greek κιάλι)  
*monocular or spyglass*

Visualise the service mesh topology, features like circuit breakers or request rates

## Features

### Graph

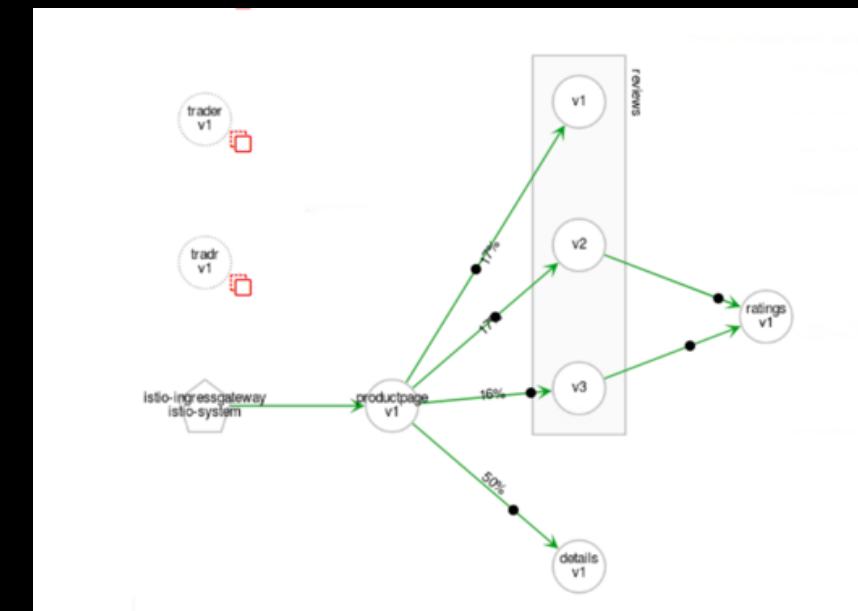
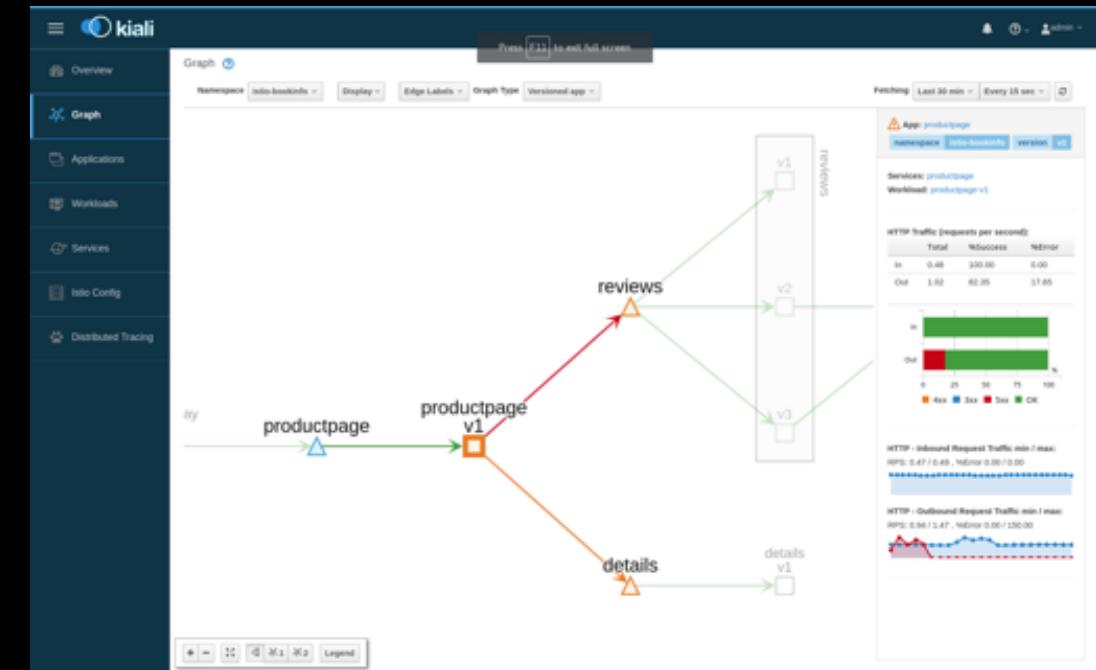
- Health
- Types
- Side Panel
- Traffic Animation

### Applications, Workloads and Services

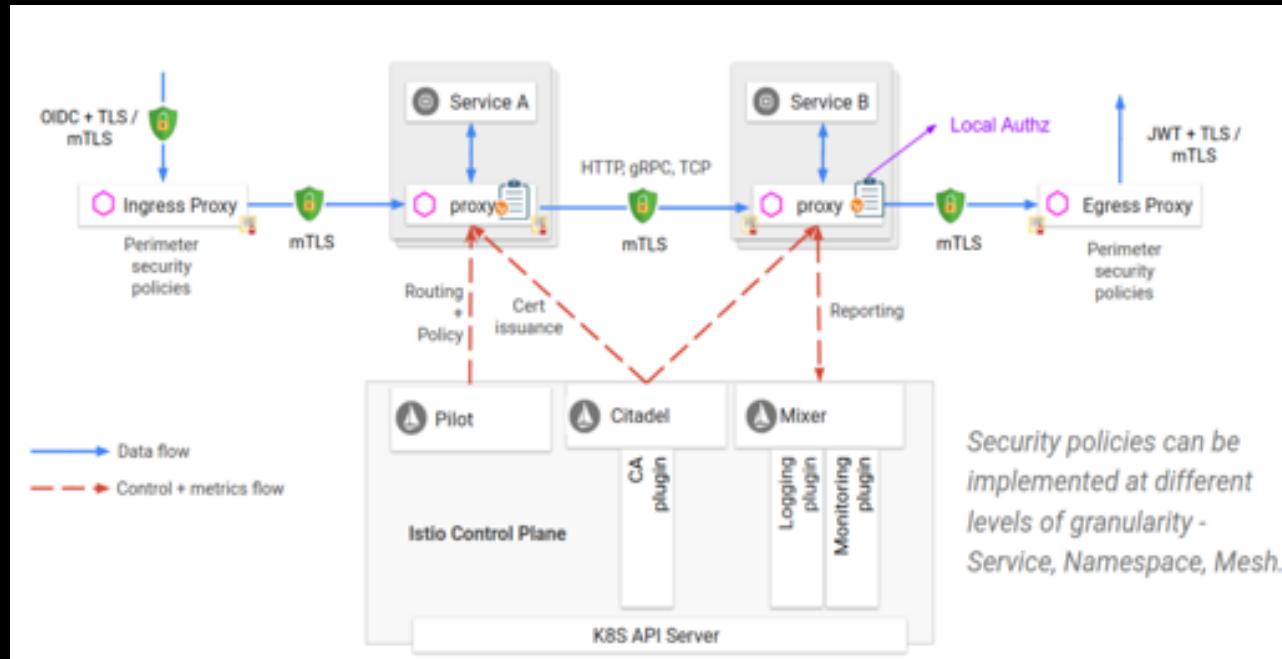
- Detailed Metrics
- Traffic Routing
- Istio compliance
- Istio Configuration

### CHALLENGE 6

I NEED TO VIEW WHAT IS GOING ON WHEN CRISIS ARISES



# Security



## Authentication

Transport authentication, also known as service-to-service authentication  
Origin authentication, also known as end-user authentication

## Authorization

Based on RBAC  
Namespace-level, service-level and method-level access control for services

**CHALLENGE 7**  
**HOW CAN I SECURE MY SERVICES?**

# Service Mesh - Bad Idea ?

A Service Mesh is not always the right solution...

- ▶ **Service Meshes are Opinionated**

They are a *platform* solution. “Work their way”

- ▶ **Service Meshes are Complex**

Adds considerable complexity with sidecars and control plane

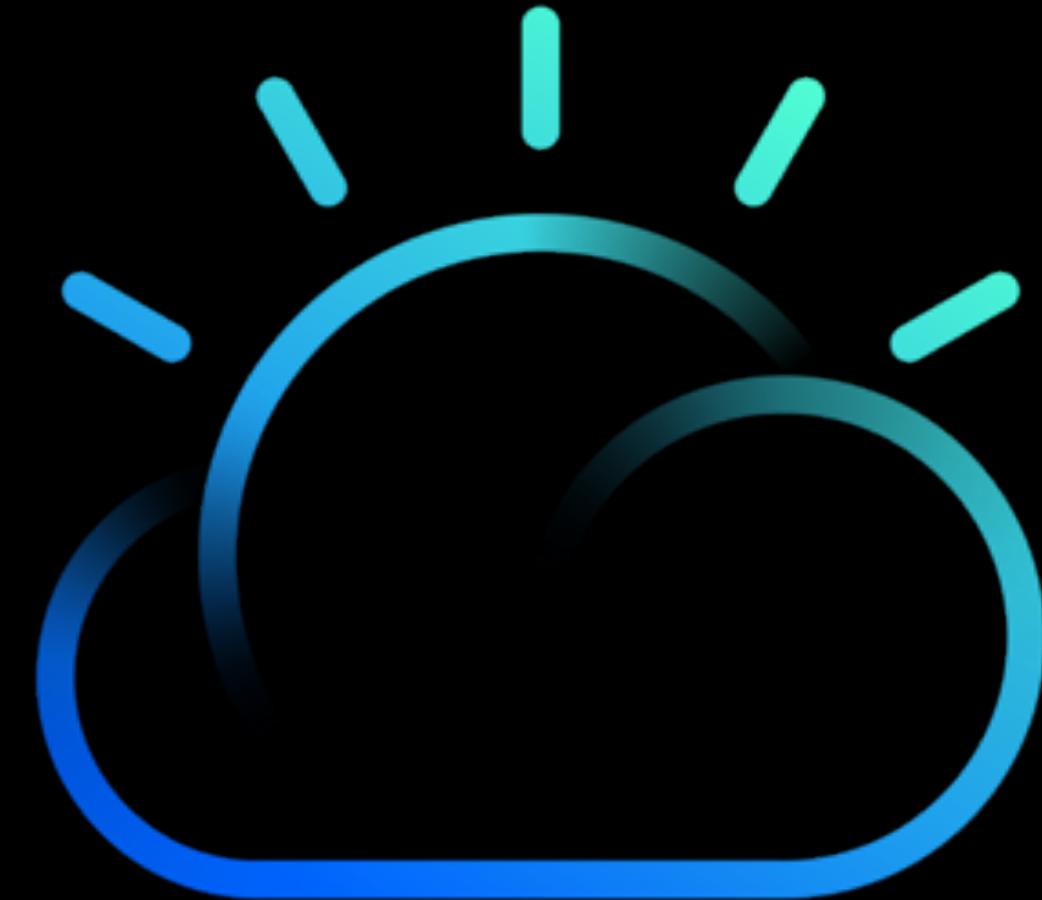
- ▶ **Service Meshes can be Slow**

Routing traffic through a series of proxies can get painfully slow (about 700 nodes → reflector)

- ▶ **Service Meshes are for Developers**

Focused primarily on Developer view.

QUESTIONS?



# Thank you

Niklaus Hirt  
Cloud / DevOps Architect

—

nikh@ch.ibm.com  
+41 79 948 72 46  
ibm.com

