

The Journey to Cloud

JTC16

Kubernetes Security

Niklaus Hirt

DevOps Architect / Cloud Architect

nikh@ch.ibm.com



Who am I?

Niklaus Hirt

Passionate about tech for over 35 years

- High-school in Berne
- Degree in Computer Science at EPFL
- ELCA
- CAST
- IBM

nikh@ch.ibm.com
@nhirt



Module – ADVANCED II

Module 0: Prepare the Labs

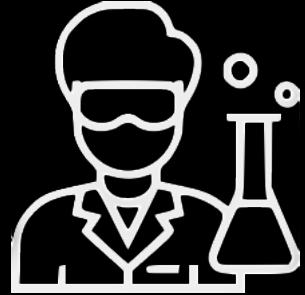
Module 1: Security Introduction

Module 2: Security Elements

BREAK

Module 3: Wrap-up and Tips

Module 4: Hands-On Lab



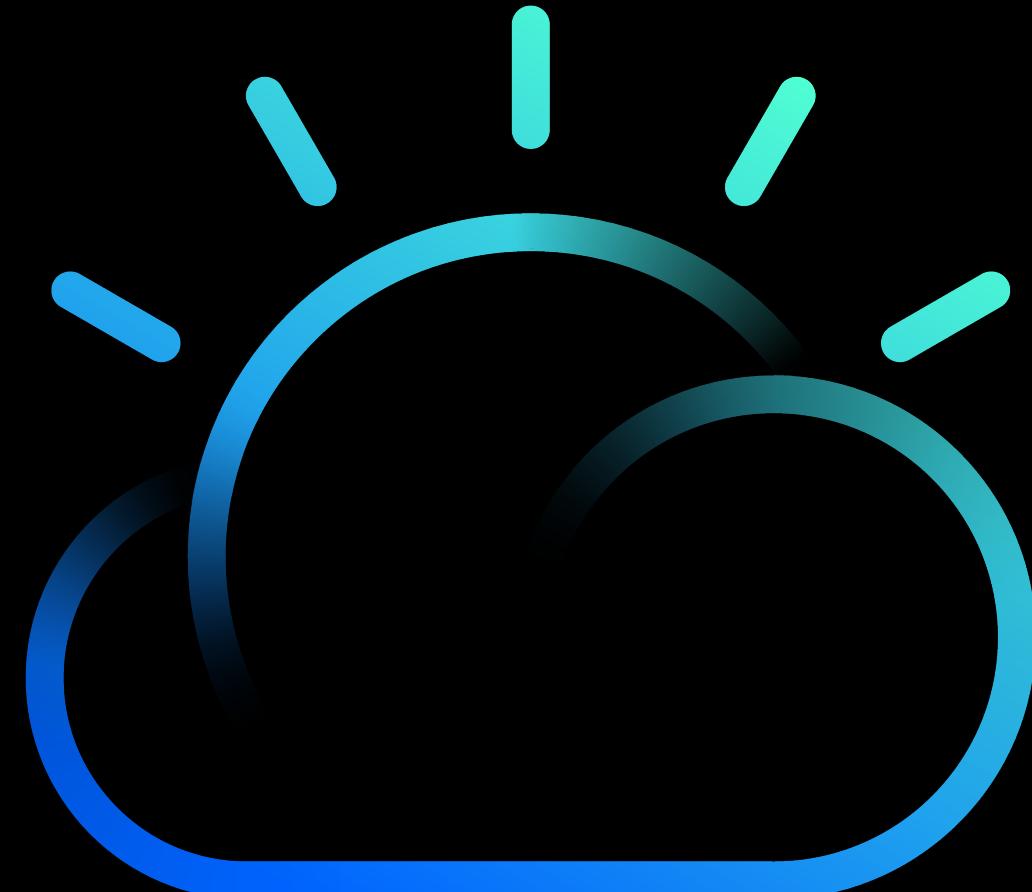
Sources and documentation will be available here:

https://github.com/niklaushirt/k8s_training_public

<https://github.com/niklaushirt/training>



°° The Journey to Cloud
Prepare the Labs



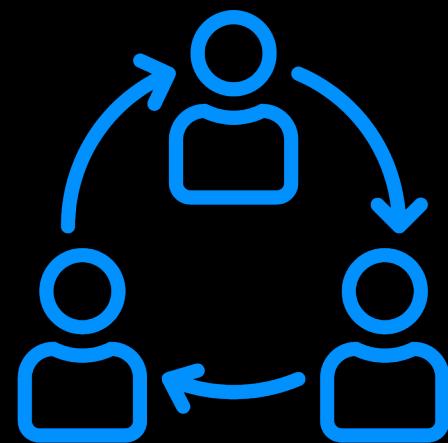
IBM Cloud

Session Objectives

Attendees will be grouped in **teams** wherever it makes sense to facilitate collaborative work.



Following some lectures will be **hands-on** work that each team collaborates to complete.



Teams

black 31701

olive 31711

peru 31715

white 31702

brown 31712

chocolate 31716

red 31703

lightblue 31713

orchid 31717

blue 31704

orange 31708

gold 31718

yellow 31705

purple 31709

pink 31719

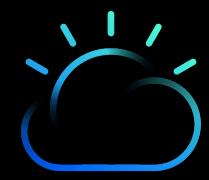
lime 31706

maroon 31710

violet 31720

cyan 31707

firebrick 31714



Collector - Accessing team web site

http://158.177.137.195:{port#}

Team name / color will be shown

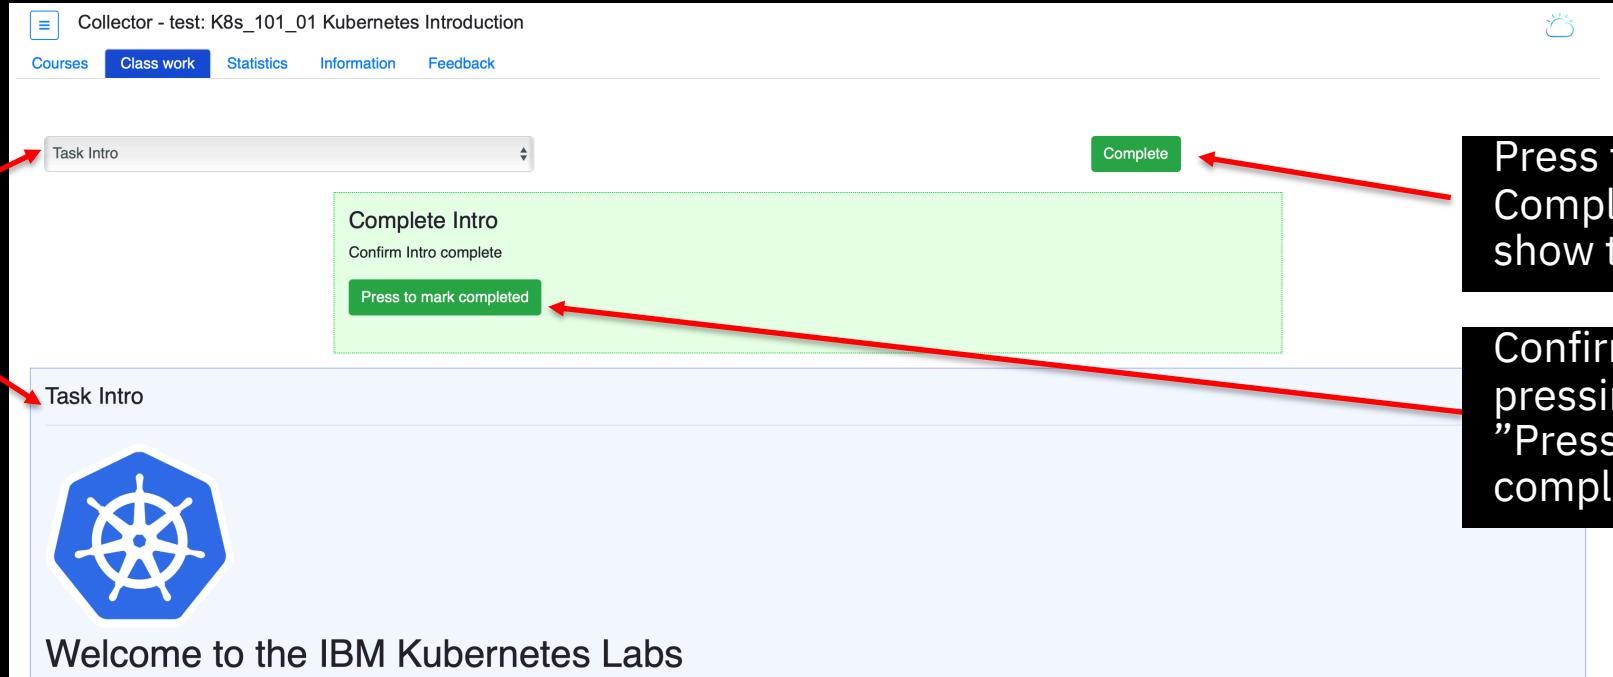
blue 31704

The screenshot shows a web browser window with the title "Collector - blue". The navigation bar includes links for "Courses", "Class work", "Statistics", "Information", and "Feedback". Below the navigation bar, a section titled "Catalog of courses" displays a list of courses. A dropdown menu is open over the first item, "KUB01 Lab Setup", with the text "select course" preceding the list. The list contains five items: "KUB01 Lab Setup", "KUB02 Kubernetes Introduction", "KUB03 Kubernetes Labs", and "KUBADV01 Istio". To the right of the dropdown, there is a button labeled "Begin course". A large callout box with a red border and white text, containing the instruction "Select course and press button to begin", points to the "Begin course" button.

Current course catalog

Collector – Class work

Select class work and the blue portion of the screen is shown



Press the green Complete button to show the green portion.

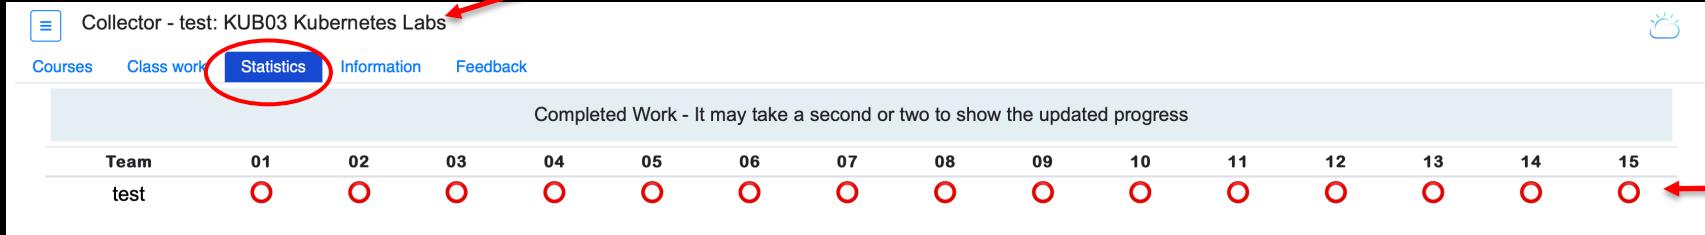
Confirm completion by pressing the green "Press to mark completed" button.



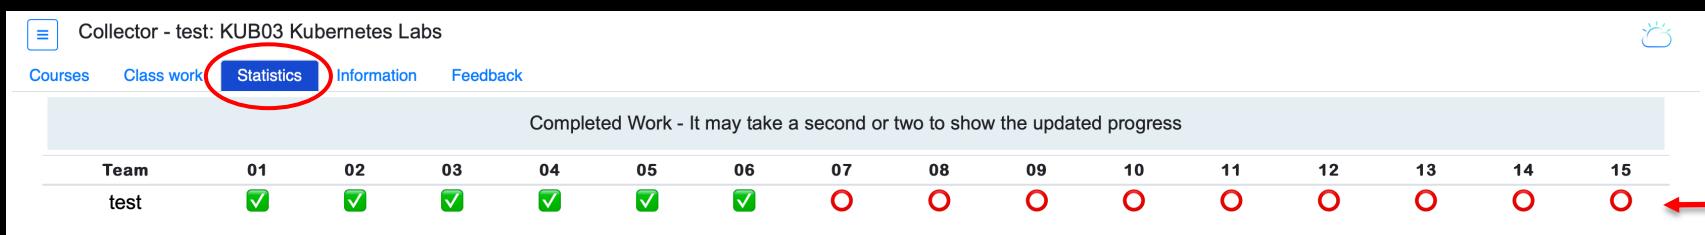
The Complete Button might not show instantly depending on the course settings

Collector – Track course completed work

Course title



The number of items tracked will change based on the current course selected.



Green checkmark - item is completed
Red circle - item is waiting to be completed

Collector – Instructor Dashboard

Remaining Time for the Lab

Collector - instructor: K8s_101_01 Kubernetes Introduction

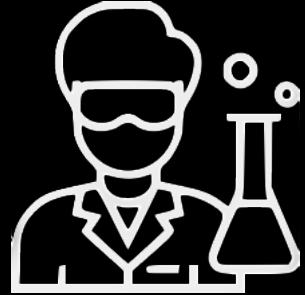
Remaining time: 0h 29m 50s

Courses Class work Statistics Information Feedback Insight

Completed Work - It may take a second or two to show the updated progress

Team	01	02	03	04	05	Cnt
instructor	✓	○	○	○	○	0
lightblue	✓	✓	✓	✓	✓	1
olive	✓	✓	○	○	○	2
peru	○	○	○	○	○	3
chocolate	○	○	○	○	○	4
pink	○	○	○	○	○	5
violet	○	○	○	○	○	6





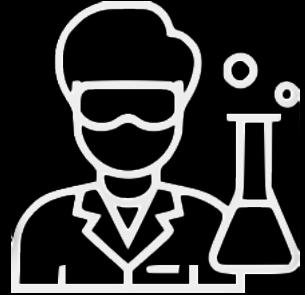
JTC90 Lab Setup

Task 1: Download Training VM

Task 2: Setup VirtualBox >6.14

Task 3: Start Training VM

Task 4: Login / Check



JTC90 Lab Setup

EVERYBODY

Task 2: Setup VirtualBox >6.14

OK

Task 3: Start Training VM

Task 4: Login ? Check

?

The Journey to Cloud **Security Introduction**

01



IBM Cloud

Kubernetes – Security – Back to basics

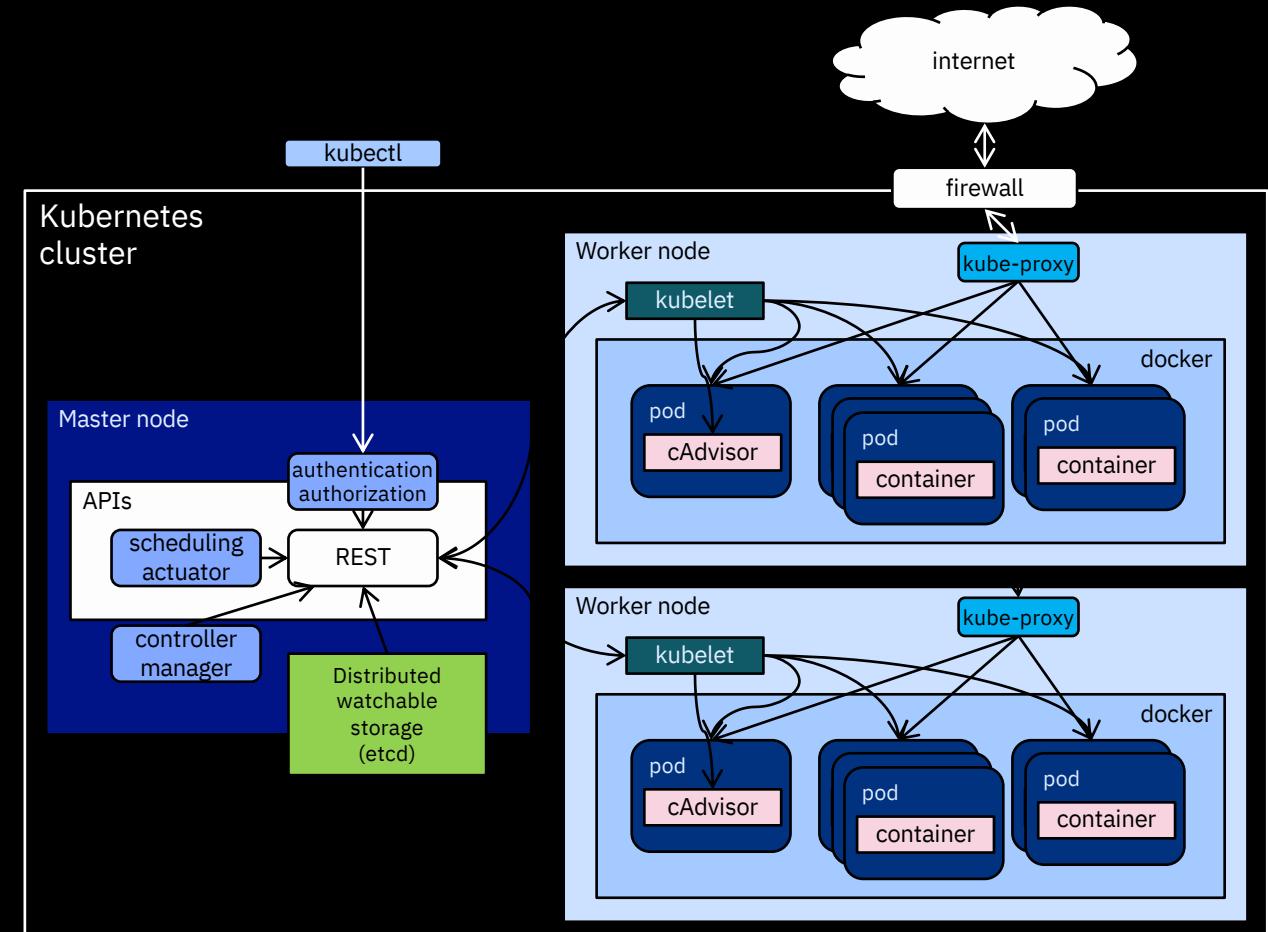


Master node

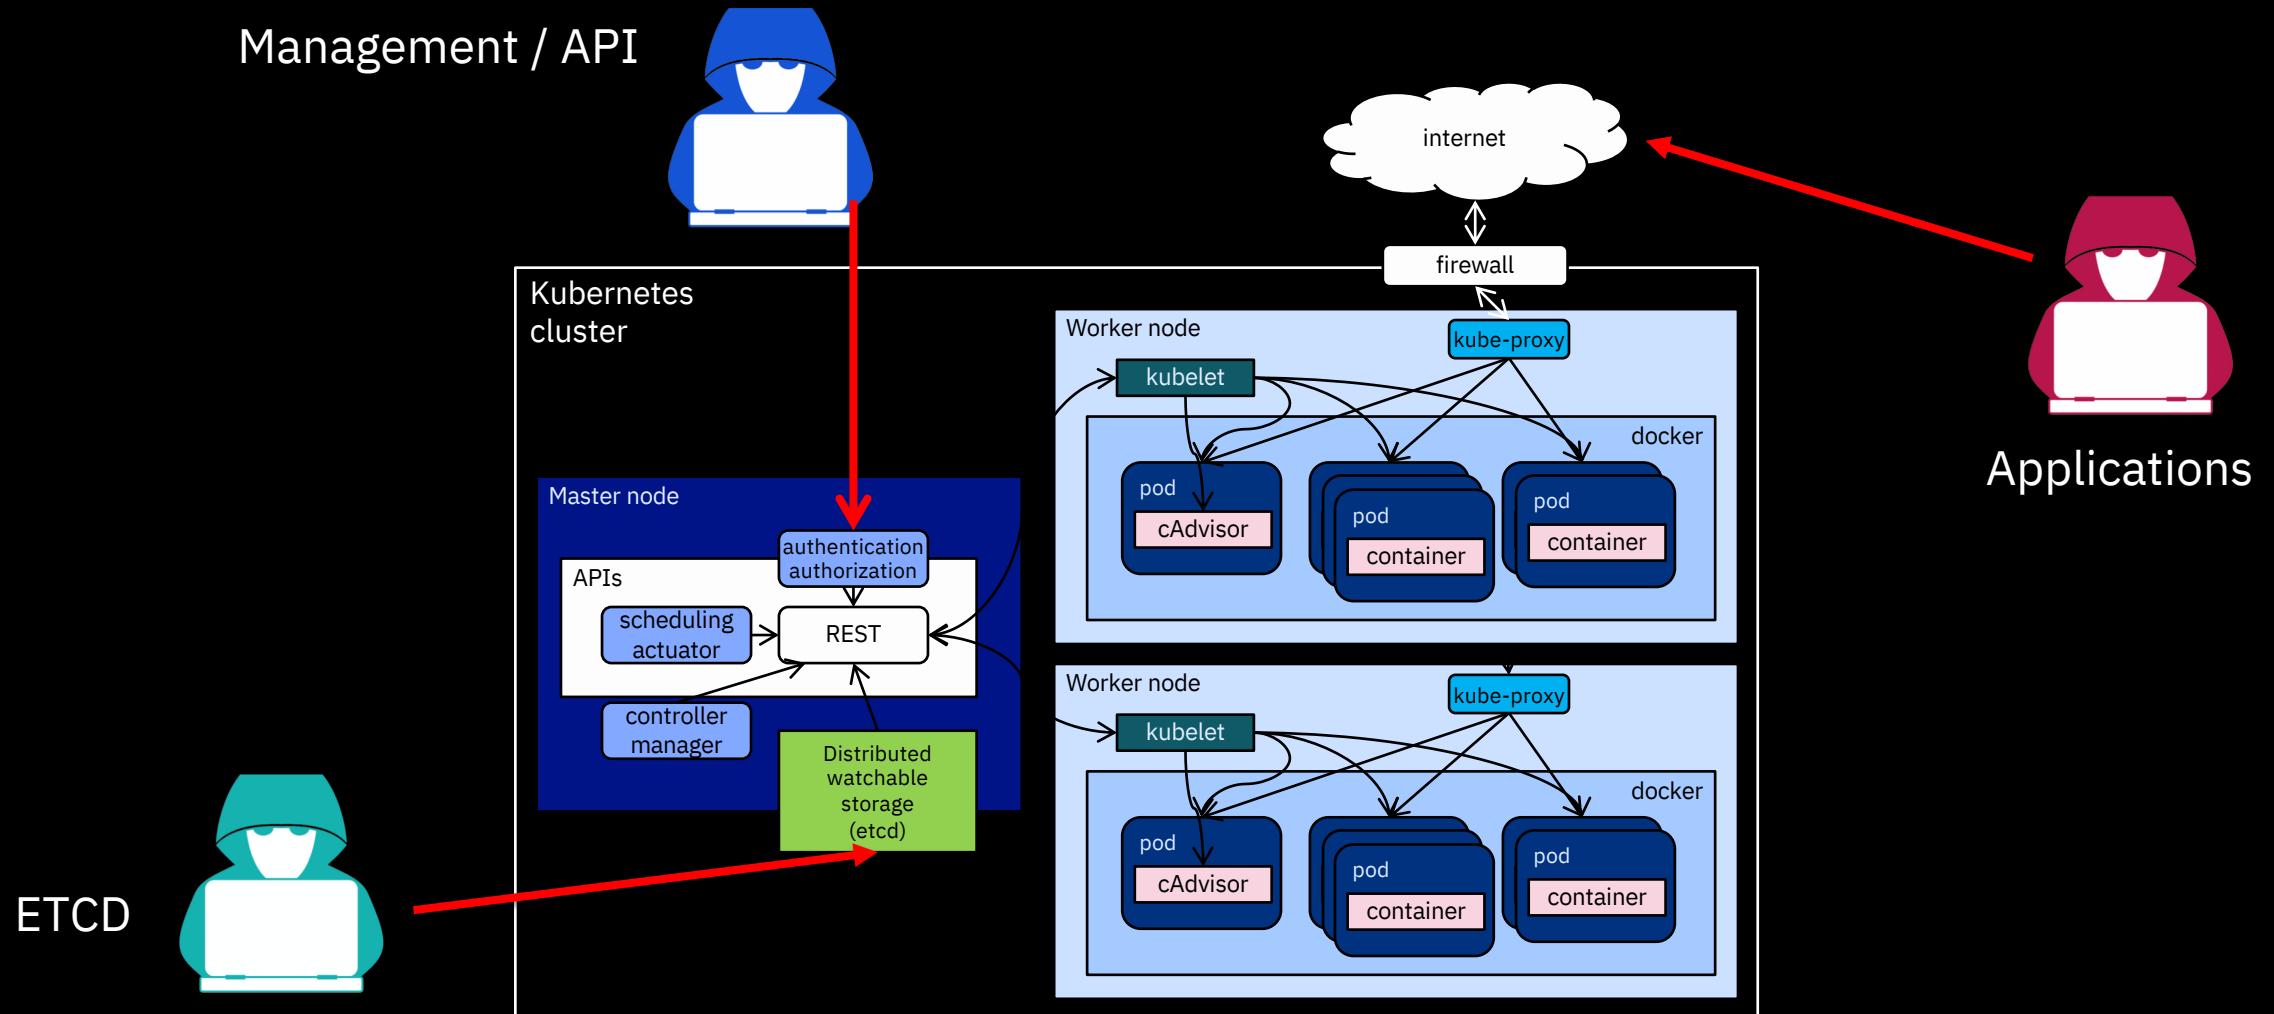
- Node that manages the cluster
- Scheduling, replication & control
- Multiple nodes for HA

Worker nodes

- Node where pods are run
- Docker engine
- kubelet agent accepts & executes commands from the master to manage pods
- cAdvisor – Container Advisor provides resource usage and performance statistics
- kube-proxy – routes inbound or ingress traffic



Kubernetes – Security – Attack surface



Kubernetes – Security Basic Topics

The **Billion Laughs** attack is a particular type of denial of service (DoS) attack which is aimed specifically at XML document parsers. This attack is also referred to as an XML bomb or an exponential entity expansion attack.

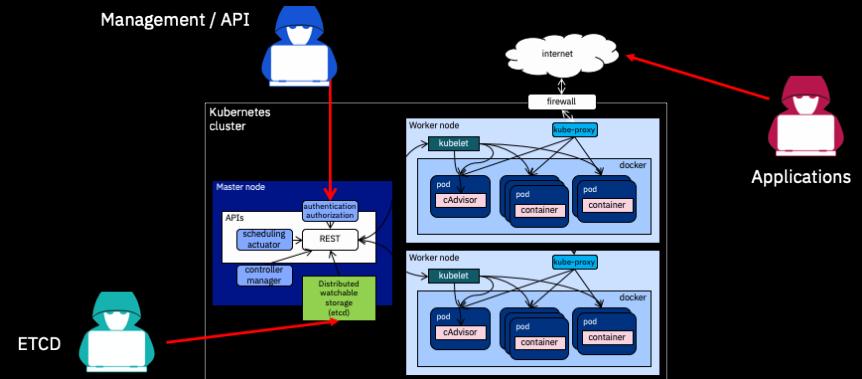
It exploits the fact that nested references to nodes can grow very large when expanded. Because the **kube-apiserver** doesn't perform validation on the manifest, it doesn't detect if those nested references will cause a problem. If the nesting references grow too large, excessive CPU and RAM usage can render the apiserver unresponsive to connections ... hence the Denial of Service.



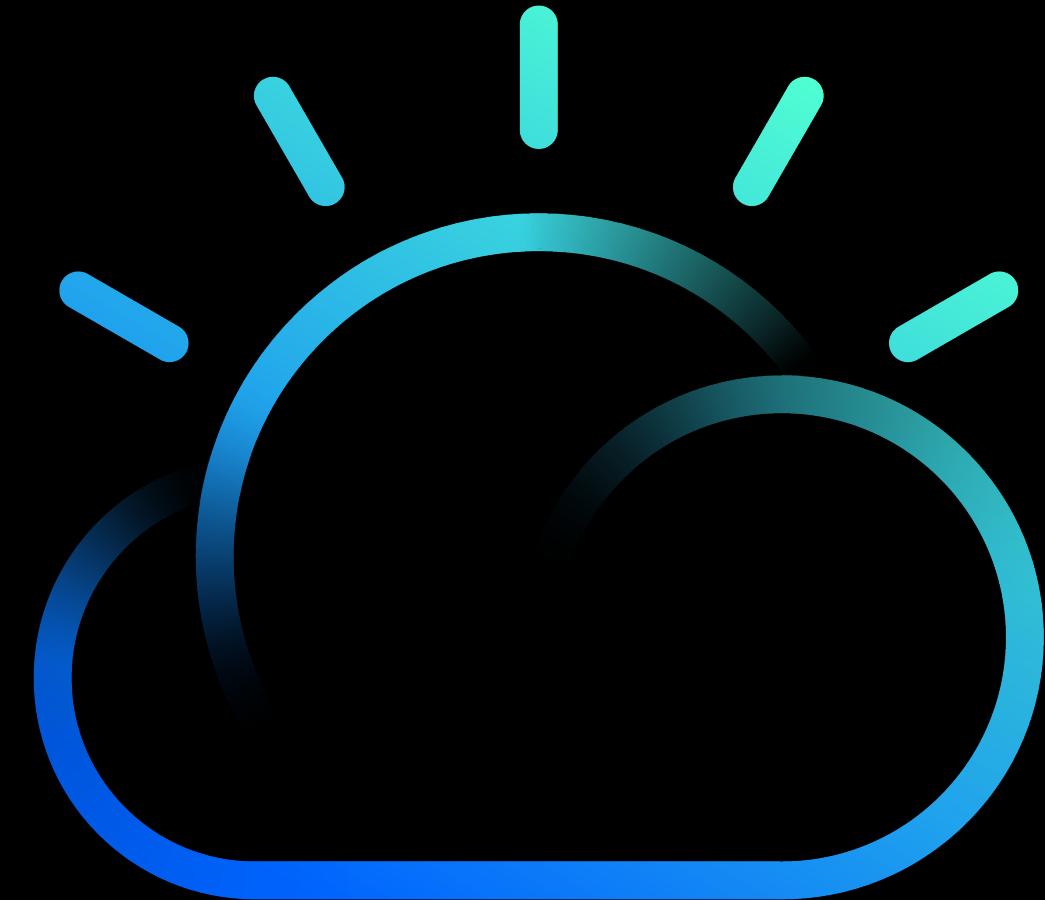
Kubernetes – Security Basic Topics

Reduce Kubernetes Attack Surfaces

- Restrict access to etcd
- Controlling access to the Kubernetes API
- Controlling access to the Kubelet
- Enforce resource usage limits for workloads
- Controlling what privileges containers run with
- Enable audit logging
- Enforcing Network Policies
- Rotate infrastructure credentials frequently
- Image Scanning of your containers
- Use Kubernetes secrets



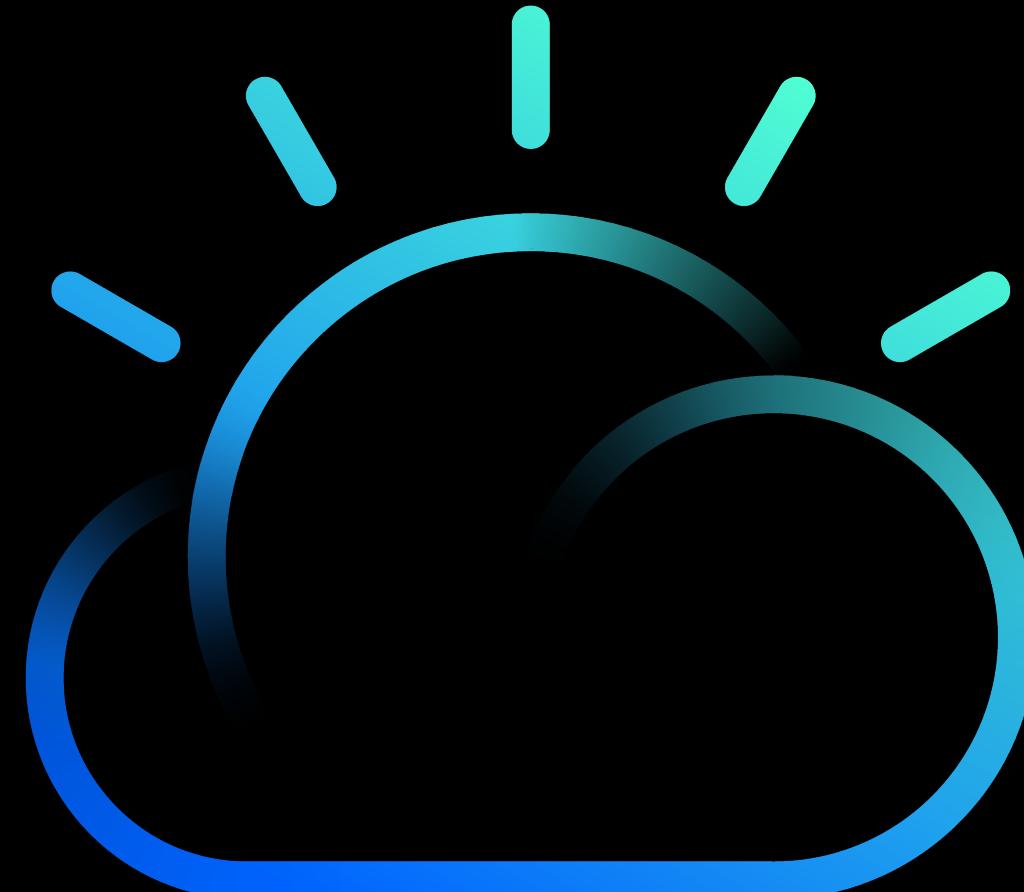
QUESTIONS?



IBM Cloud

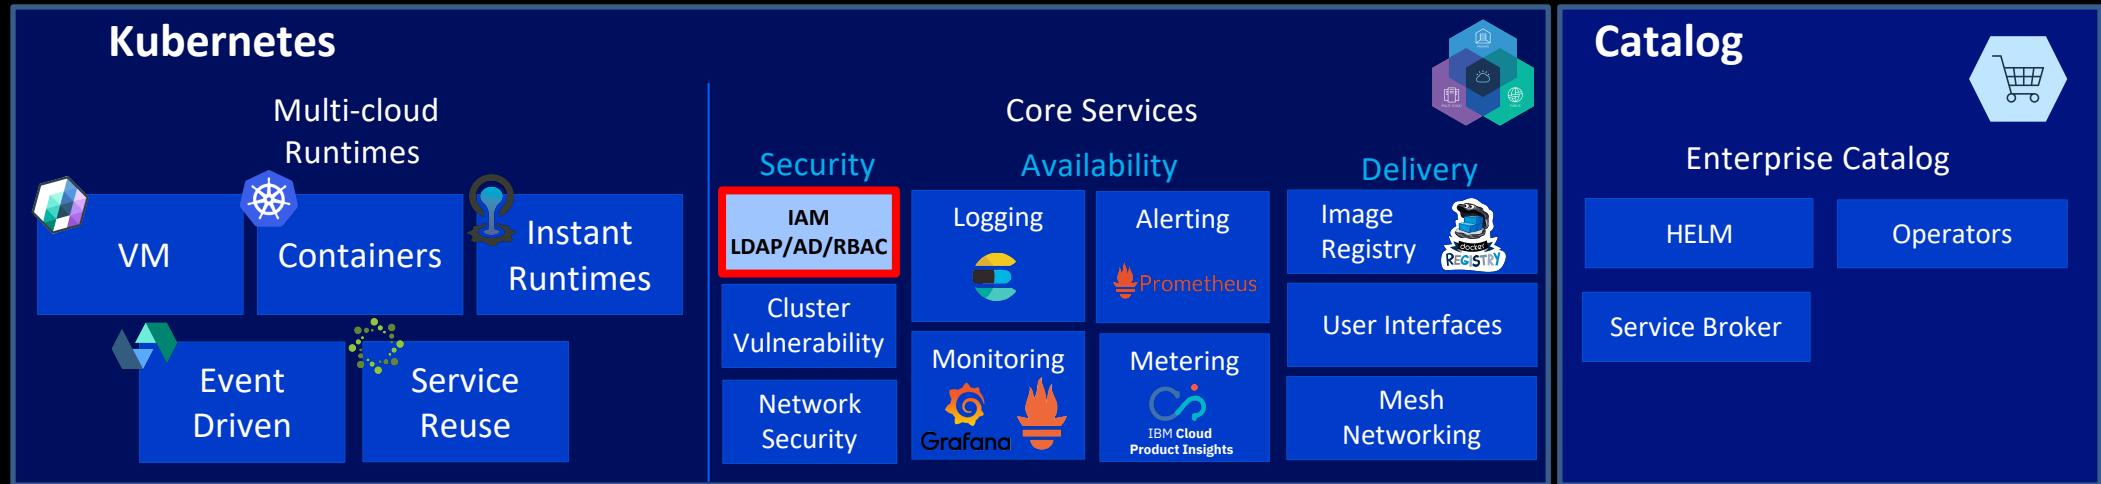
The Journey to Cloud **Security Elements**

02



IBM Cloud

Kubernetes – Core Services - Controlling K8s Access



Authentication – WHO am I - (token, certs, OpenID Connect Provider (OIDC)...)

Developer – create, view, get permission.

Tester – create, delete, update, get permission.

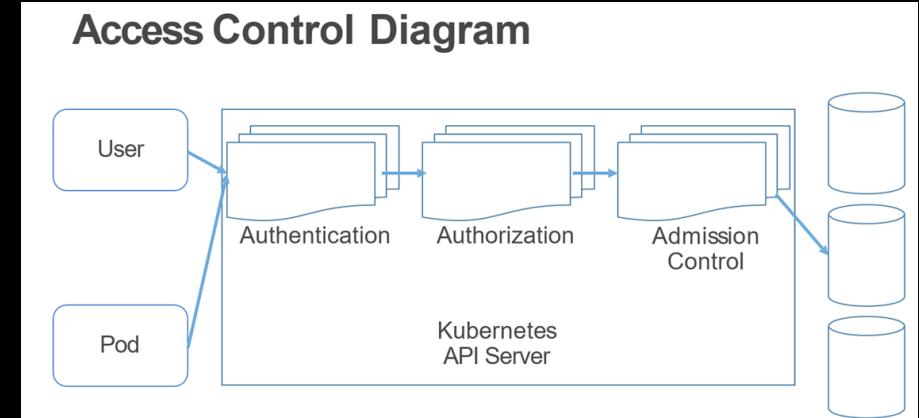
Administrator – All permission

Authorization – CAN I - (RBAC)

Is the user or application authorized or have permissions to access a K8s object?

Admission Control

Intercepts requests to the Kubernetes API server prior to persistence of the object, but after the request is authenticated and authorized





Integrating with LDAP/AD

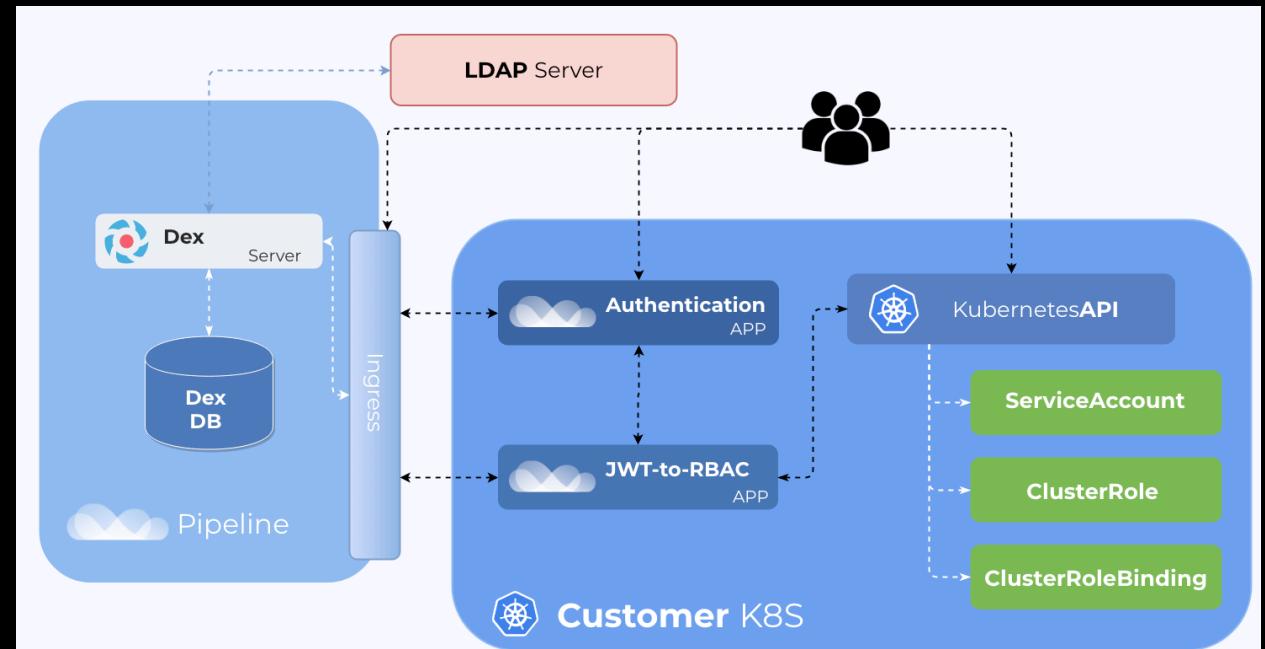
OpenID Connect

ID Tokens are JSON Web Tokens (JWTs)

dex

Dex acts as a portal to other identity providers through connectors

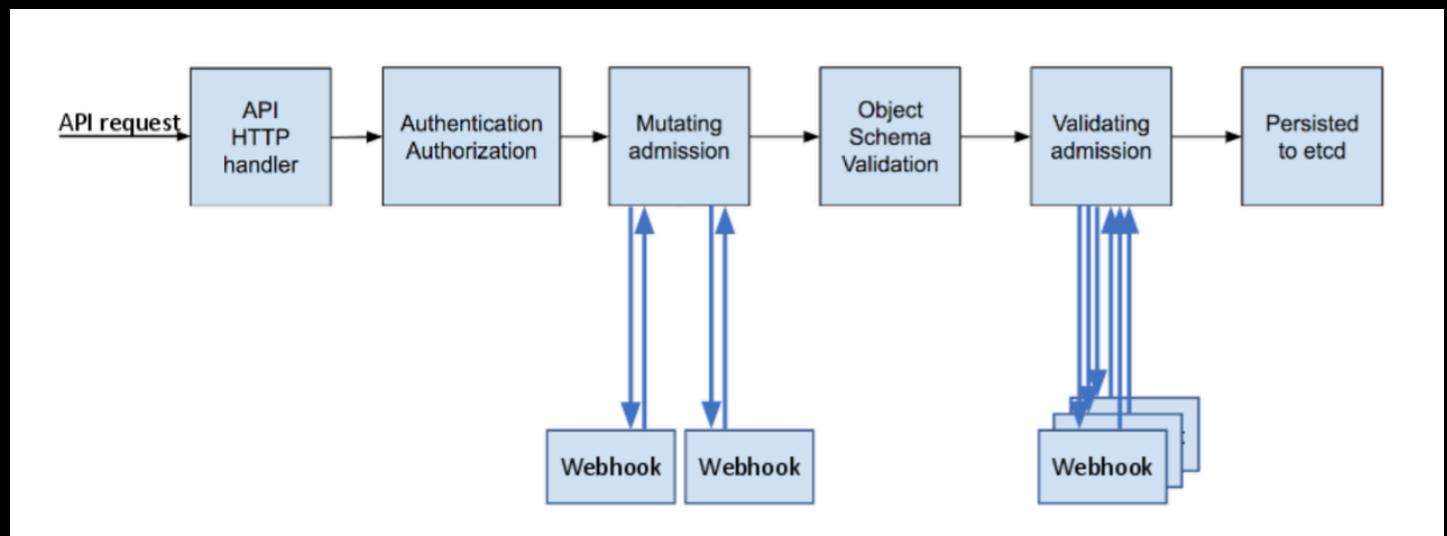
dex defer authentication to LDAP servers, SAML providers, or established identity providers like GitHub, Google, and Active Directory.



Kubernetes WebHooks

WebHooks Type

- **mutating**: to dynamically change incoming deployments on the fly (think automatic Istio sidecar injection), and
- **validating**: to accept or reject those same deployments based on the rules in your callback.



RBAC Basic Elements

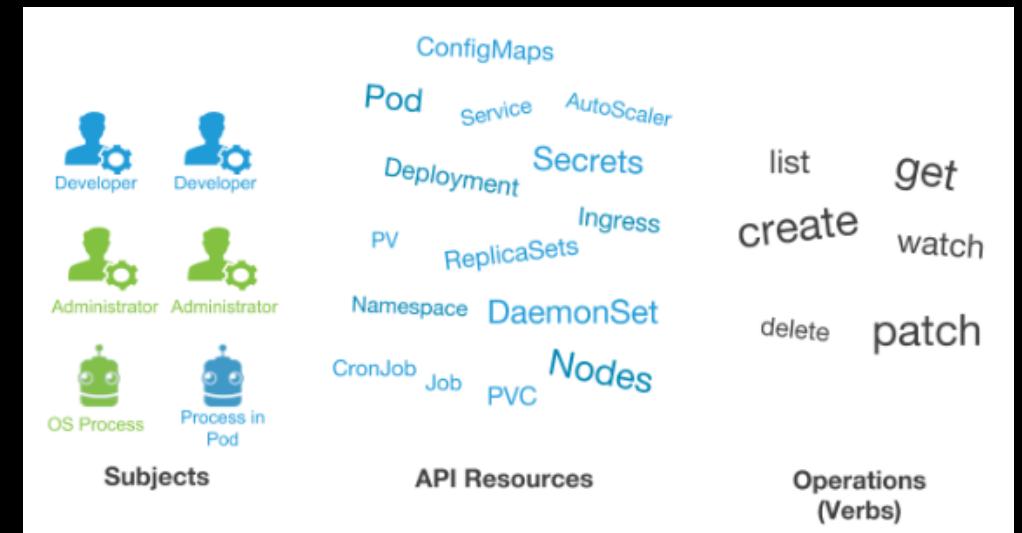
RBAC Building Blocks

Objects

- Pods
- PersistentVolumes
- ConfigMaps
- Deployments
- Nodes
- Secrets
- Namespaces

Verbs

- create
- get
- delete
- list
- update
- edit
- watch
- exec



RBAC Basic Elements

Rules

Operations (verbs) that can be carried out on a group of resources which belong to different API Groups.

Roles and ClusterRoles

Both consist of rules.

- Role: applicable to a single namespace
- ClusterRole: is cluster-wide

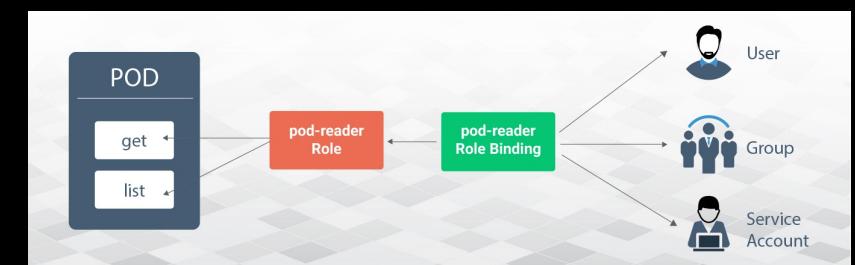
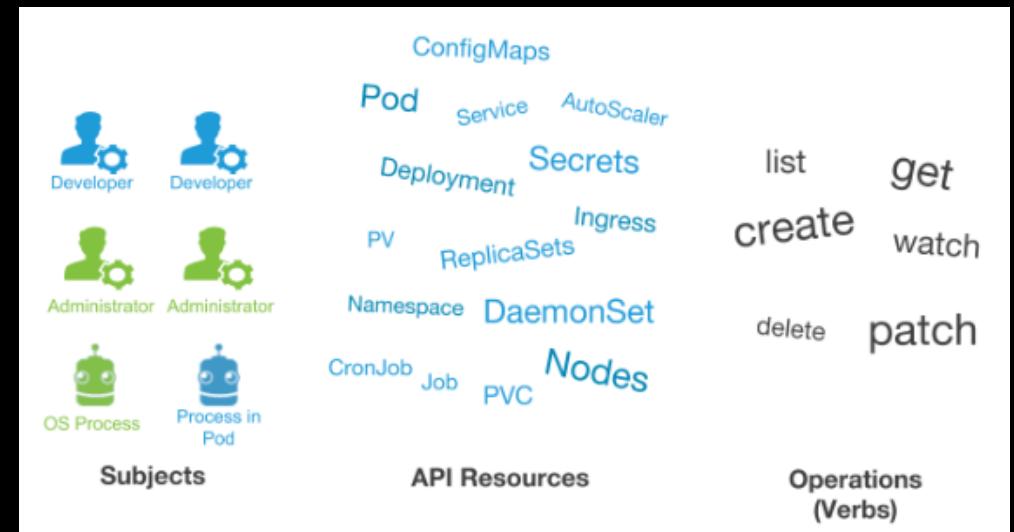
Subjects

Entity that attempts an operation in the cluster

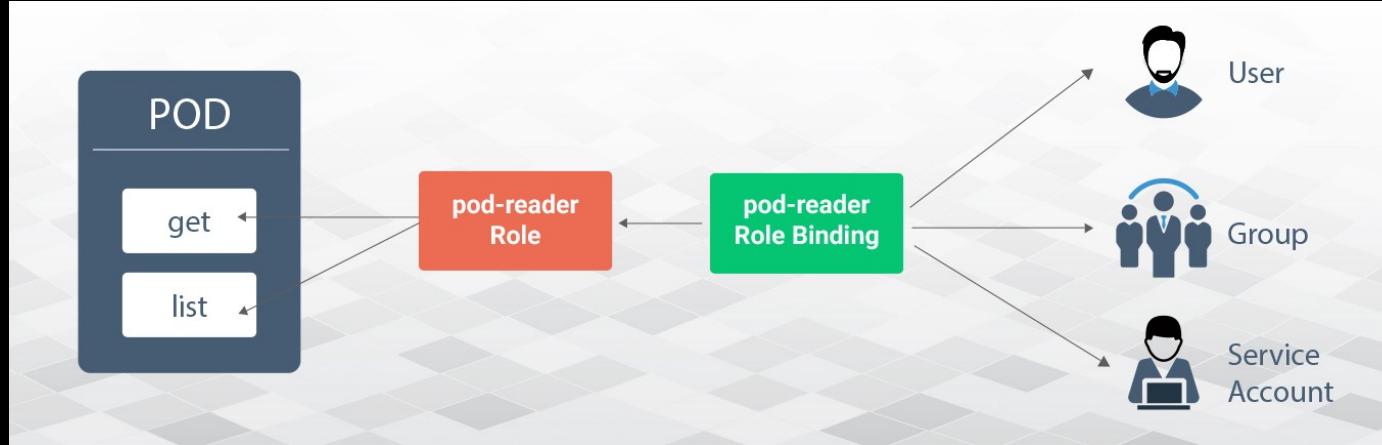
- User Accounts: Humans or processes living outside the cluster.
- Service Accounts: Namespaced account.
- Groups: Reference multiple accounts. (default groups like cluster-admin)

RoleBindings and ClusterRoleBindings

Bind subjects to roles



RBAC Example



```
kind: Role
metadata:
  name: pod-reader
rules:
- apiGroups: [""]
  resources: ["pods"]
  verbs: ["get", "list"]
```

```
kind: RoleBinding
metadata:
  name: pod-reader
subjects:
- kind: User
  name: John
roleRef:
  kind: Role
  name: pod-reader
```

Common RBAC Pitfalls

- **Cluster Administrator Role Granted Unnecessarily**
The built-in **cluster-admin** role grants effectively unlimited access to the cluster.
Should be granted only to the specific users that need it.
- **Improper Use of Role Aggregation**
Role Aggregation can be used to combine existing roles.
Review carefully or avoid altogether.
- **Duplicated Role Grant**
Role definitions may overlap with each other, making access revocation more difficult.
- **Unused Role**
Roles that are created but not granted to any subject can increase the complexity of RBAC management.
- **Grant of Missing Roles**
Role bindings can reference roles that do not exist.
If the role name is reused those role bindings can unexpectedly become active.

Always adapt
least privileged access
practices

Audit RBAC – rbac-lookup

Some tools to help you see more clearly

```
→ rbac-lookup openshift
```

SUBJECT	SCOPE
system:serviceaccounts:openshift	openshift
system:serviceaccounts:openshift-console	openshift-console
system:serviceaccounts:openshift-console	openshift-console
system:serviceaccounts:openshift-infra	openshift-infra
system:serviceaccounts:openshift-logging	openshift-logging
system:serviceaccounts:openshift-logging	openshift-logging
system:serviceaccounts:openshift-metrics-server	openshift-metrics-server
system:serviceaccounts:openshift-metrics-server	openshift-metrics-server
system:serviceaccounts:openshift-monitoring	openshift-monitoring
system:serviceaccounts:openshift-monitoring	openshift-monitoring
system:serviceaccounts:openshift-node	openshift-node
system:serviceaccounts:openshift-sdn	openshift-sdn
system:serviceaccounts:openshift-sdn	openshift-sdn
system:serviceaccounts:openshift-template-service-broker	openshift-template-service-broker
system:serviceaccounts:openshift-template-service-broker	openshift-template-service-broker
system:serviceaccounts:openshift-web-console	openshift-web-console
system:serviceaccounts:openshift-web-console	openshift-web-console

ROLE
ClusterRole/system:image-puller

```
→ rbac-lookup root
```

SUBJECT	SCOPE	ROLE
root	cluster-wide	ClusterRole/cluster-admin

Example Usage

- rbac-lookup root
- rbac-lookup openshift

Audit RBAC - rakkess

Some tools to help you see more clearly

# rakkess --namespace default	LIST	CREATE	UPDATE	DELETE
NAME				
bindings		x		
configmaps	✓	✓	✓	✓
controllerrevisions.apps	✓	x	x	x
cronjobs.batch	✓	✓	✓	✓
daemonsets.apps	✓	✓	✓	✓
daemonsets.extensions	✓	✓	✓	✓
deployments.apps	✓	✓	✓	✓
deployments.extensions	✓	✓	✓	✓
endpoints	✓	✓	✓	✓
events	✓	x	x	x
events.events.k8s.io	x	x	x	x
horizontalpodautoscalers.autoscaling	✓	✓	✓	✓
ingresses.extensions	✓	✓	✓	✓
jobs.batch	✓	✓	✓	✓
leases.coordination.k8s.io	x	x	x	x
limitranges	✓	x	x	x
localsubjectaccessreviews.authorization.k8s.io	x			
networkpolicies.extensions	✓	✓	✓	✓

Example Usage

- rakkess --namespace default
- rakkess --verbs get,delete,watch,patch

Audit RBAC – rbac-view

Some tools to help you see more clearly

The screenshot shows a web-based interface for the rbac-view tool. At the top, there's a header with a logo of an eye and the text "RBAC". Below the header, there are two tabs: "Cluster Roles" (which is selected) and "Roles". A legend at the top provides a key for the permission icons: green square (c) - create, red square (d) - delete, yellow square (g) - get, blue square (l) - list, purple square (w) - watch, orange square (p) - patch, brown square (u) - update, and grey square (*) - deletecollection.

The main area is titled "Roles" and contains a table with several rows. Each row represents a role and its details:

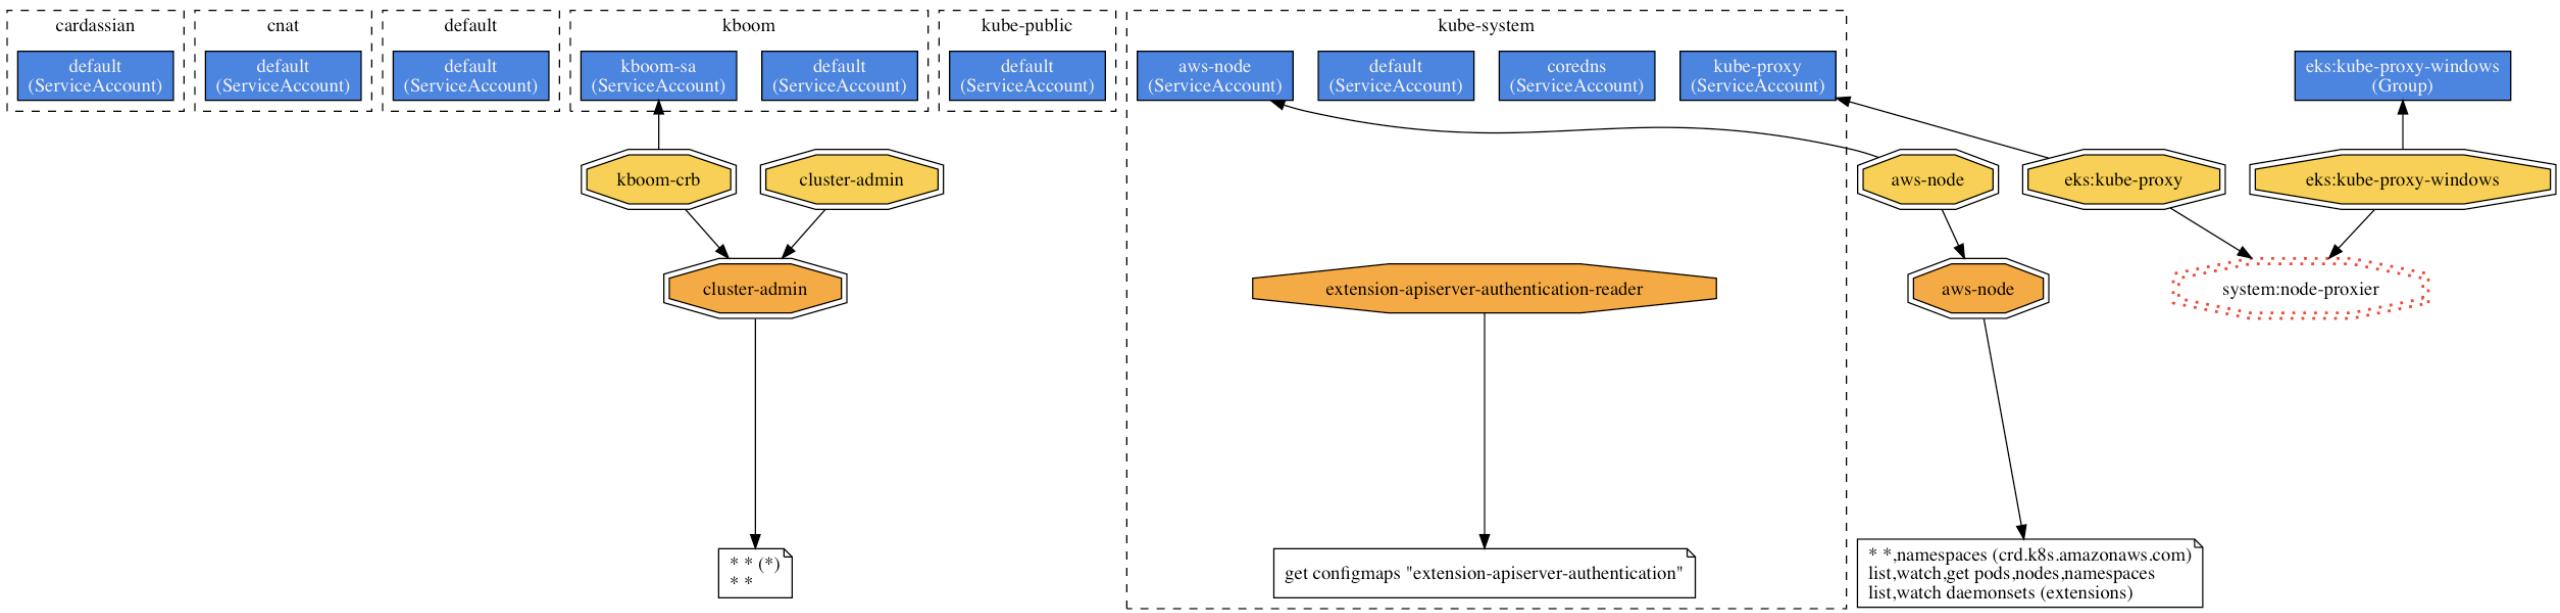
RoleName	Subjects	configmaps	events
system:controller:token-cleaner	Namespace: kube-system		(c)
kubeadm:bootstrap-signer-clusterinfo	Namespace: kube-public	(g)	
system:controller:bootstrap-signer	Namespace: kube-public	(u)	(c) (p) (u)
system:controller:bootstrap-signer	Namespace: kube-system		
system:leader-locking-kube-scheduler	Namespace: kube-system	(c)	(u)
system:controller:cloud-provider	Namespace: kube-system	(c) (g) (l) (w)	

Example Usage

```
./rbac-view --render html (default)  
./rbac-view --render json
```

Audit RBAC – rback

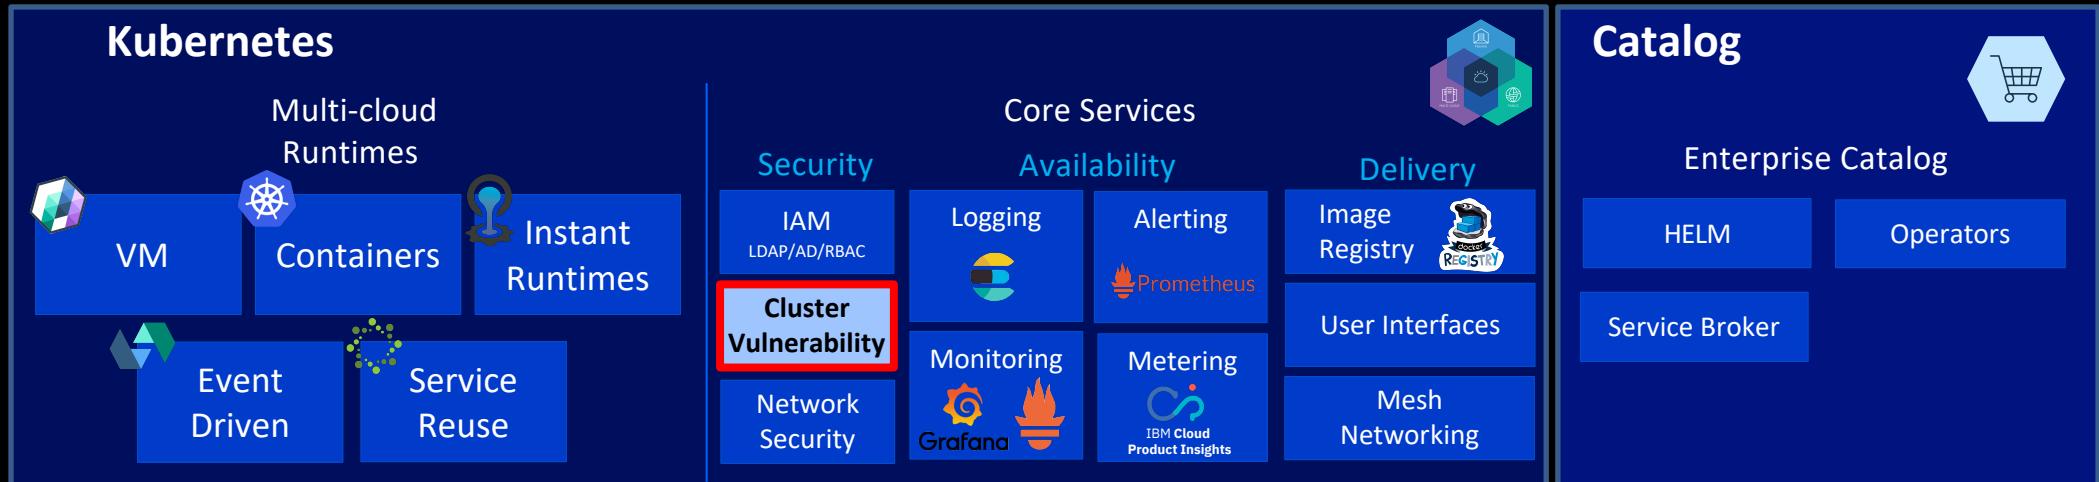
Some tools to help you see more clearly



Example Usage

```
./rbac-view --render html (default)  
./rbac-view --render json
```

Kubernetes – Core Services – Cluster Security



Kubernetes – Cluster Security – Secure Images

Scan images for vulnerabilities



One of the best-known Image scanning tools is Clair. It is an open source project for the static analysis of vulnerabilities in application containers (currently including appc and docker).

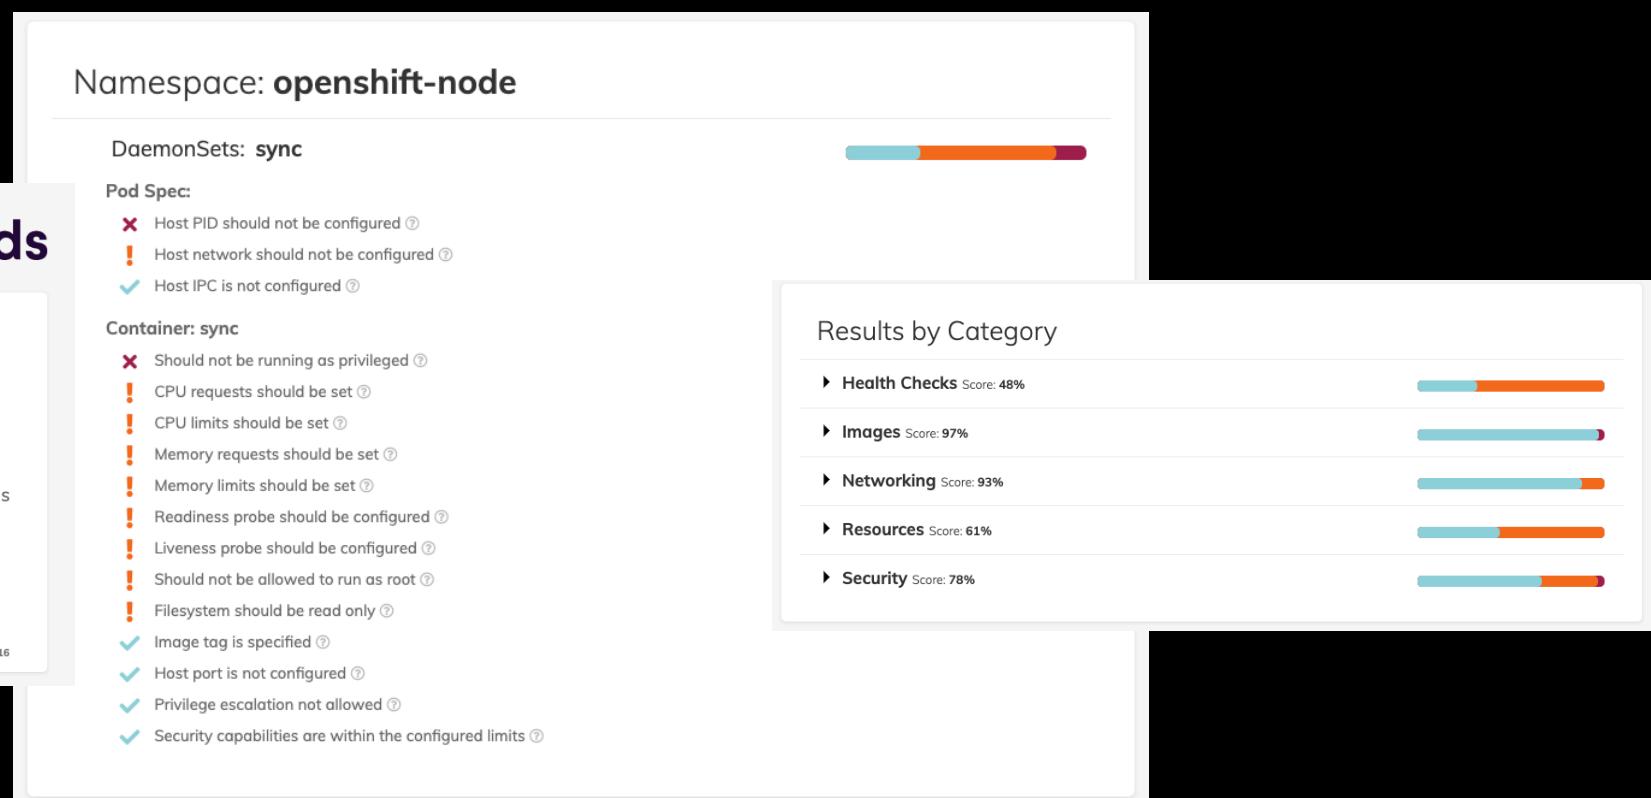
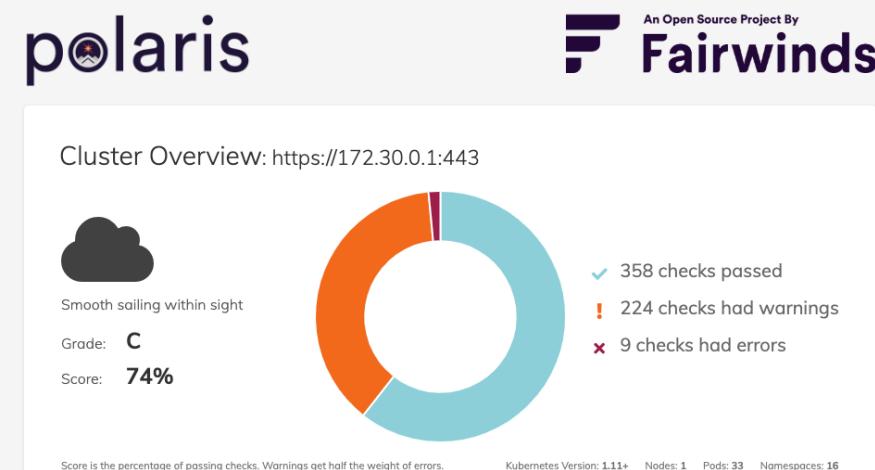
clair timeout 1m0s docker timeout: 1m0s no whitelist file Analysing 14 layers Got results from Clair API v1 Found 734 vulnerabilities Unknown: 223 Negligible: 345 Low: 164 Medium: 2						
SEVERITY	NAME	FEATURENAME	FEATUREVERSION	FIXEDBY	DESCRIPTION	LINK
Medium	CVE-2009-3546	libwmf	0.2.8.4-10.6		The _gdGetColors function in gd_gd.c in PHP 5.2.11 and 5.3.x before 5.3.1, and the GD Graphics Library 2.x, does not properly verify a certain colorsTotal structure member, which might allow remote attackers to conduct buffer overflow or buffer over-read attacks via a crafted GD file, a different vulnerability than CVE-2009-3293. NOTE: some of these details are obtained from third party information.	https://security-tracker.debian.org/tracker/CVE-2009-3546
Medium	CVE-2007-3996	libwmf	0.2.8.4-10.6		Multiple integer overflows in libgd in PHP before 5.2.4 allow remote attackers to cause a denial of service (application crash) and possibly execute arbitrary code via a large (1) srcW or (2) srcH value to the (a) gdImageCopyResized function, or a large (3) sy (height) or (4) sx (width) value to the (b) gdImageCreate or the (c) gdimageCreateTrueColor function.	https://security-tracker.debian.org/tracker/CVE-2007-3996

Kubernetes – Cluster Security – Secure Images

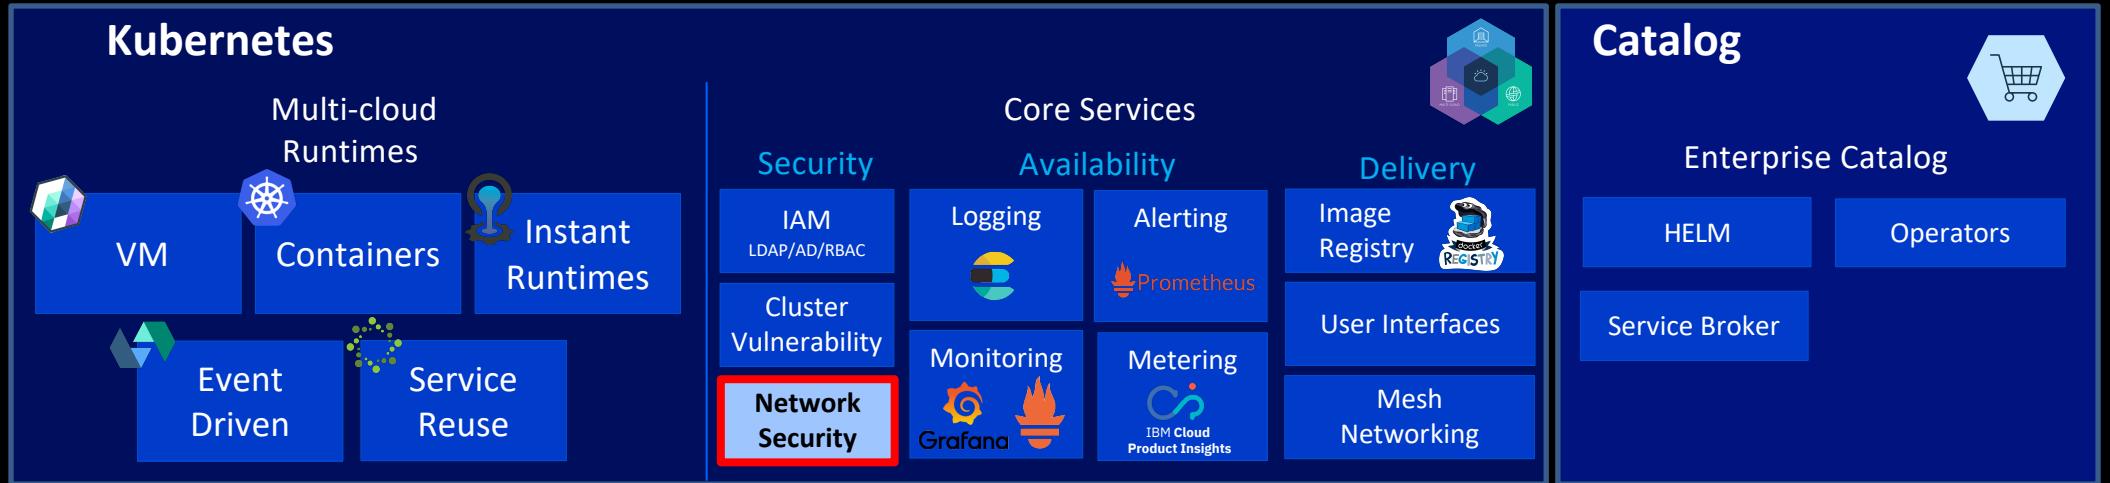
Scan for K8s Best Practices

Some tools to help you see more clearly

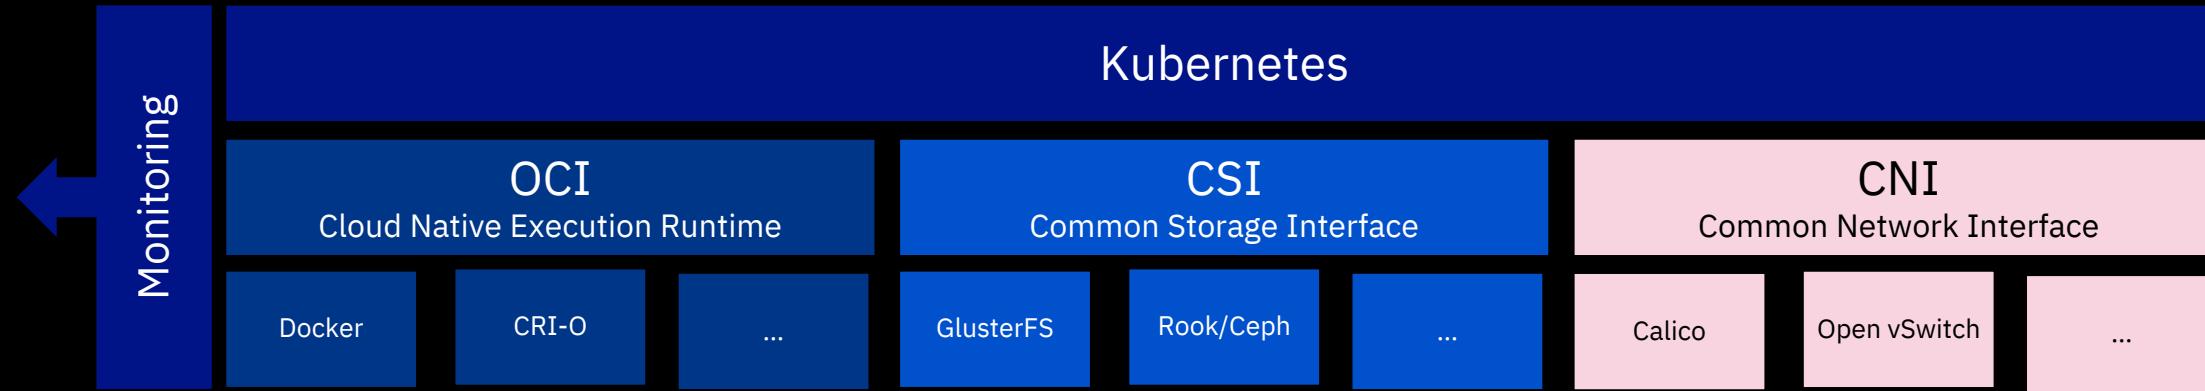
Polaris – Validation of best practices in your Kubernetes clusters



Kubernetes – Core Services – Network Security



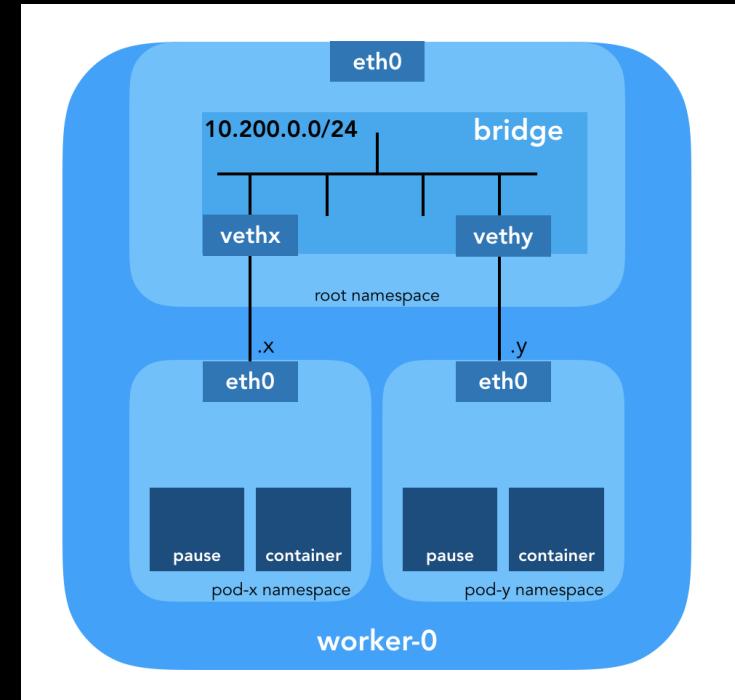
Kubernetes – Core Services – Network Security



CNI – Common Network Interface

Provides a rich set of security enforcement capabilities running on top of a highly scalable and efficient virtual network

- **Calico**
Virtual networking and network security for containers, VMs, and bare metal services.
- **Open vSwitch**
Production quality, multilayer virtual switch
- ...

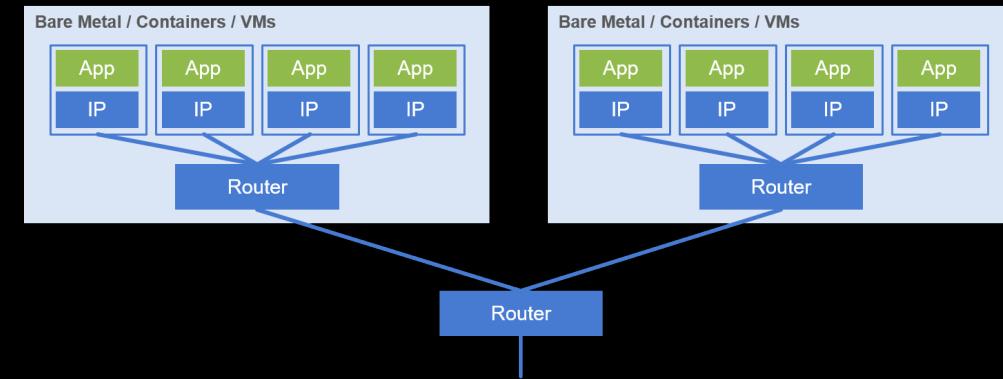


Kubernetes – Core Services – Network Security

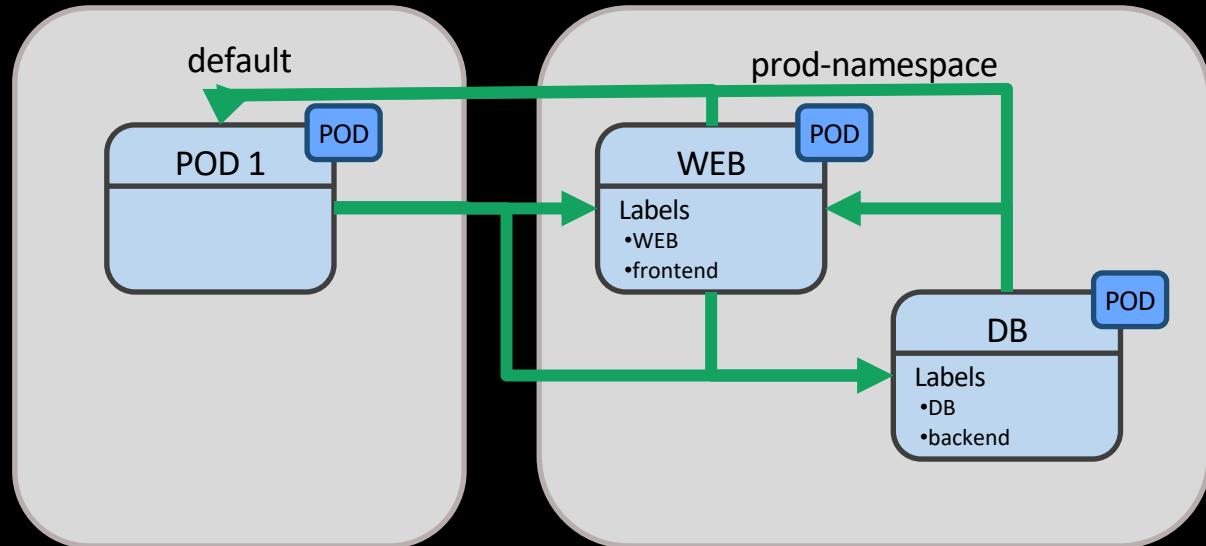


Operates at **Layer 3**, which is the network layer

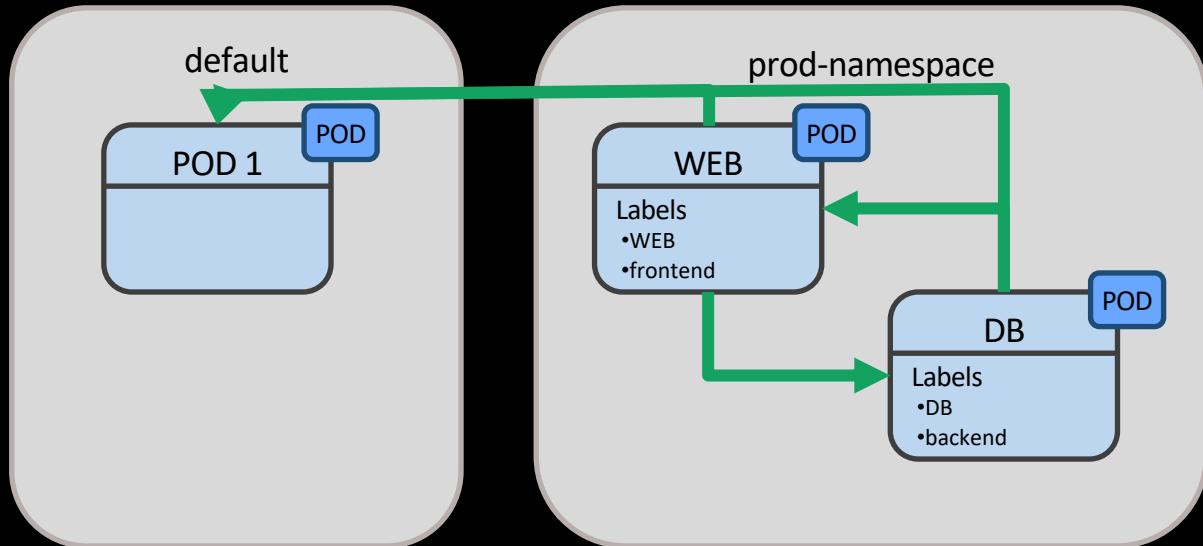
- Has the advantage of being **universal** (DNS, SQL, real-time streaming, ...)
- Can extend beyond the service mesh (including to **bare metal or VM** endpoints not under the control of Kubernetes).
- Calico's policy is enforced at the host node, outside the network namespace of the guest pods.
- Based on **iptables**, which are packet filters implemented in the standard Linux kernel, it is extremely fast.



Kubernetes – Core Services – Network Security

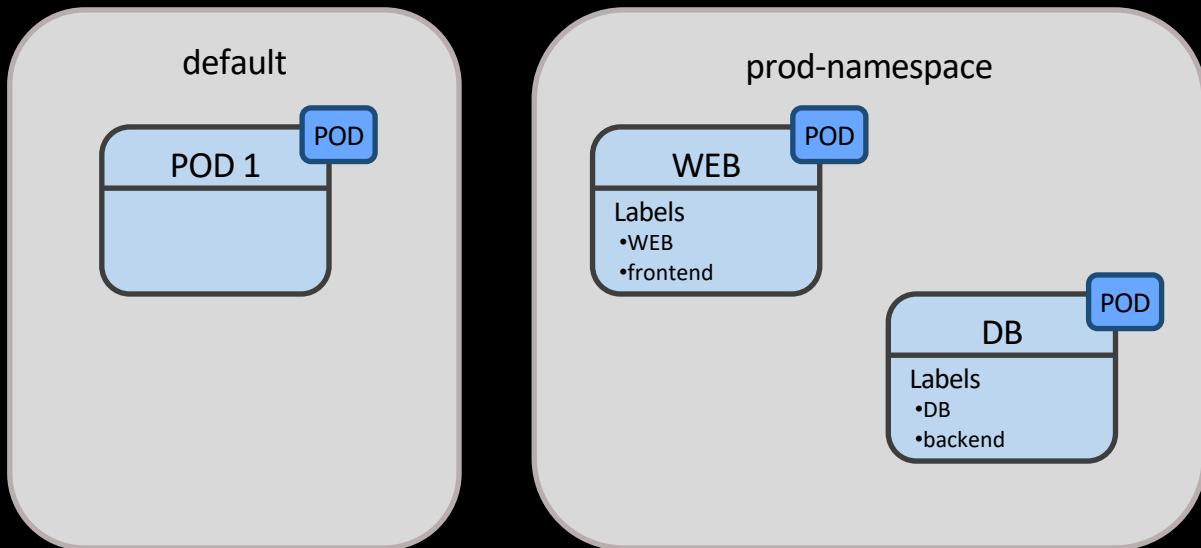


Kubernetes – Core Services – Network Security



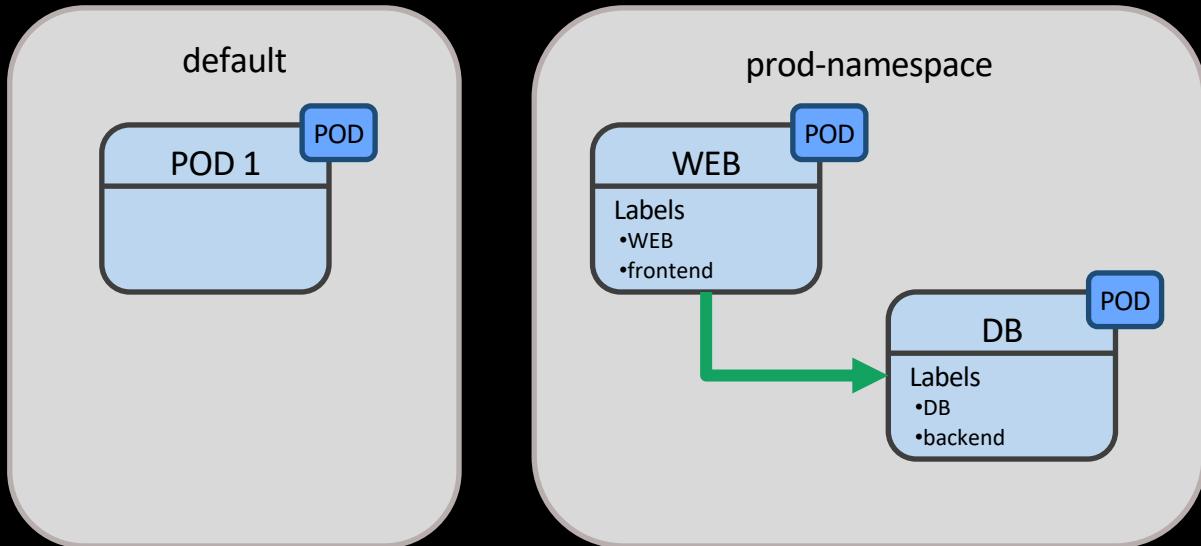
```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: default-deny
  namespace: demo-namespace
spec:
  podSelector:
    matchLabels: {}
```

Kubernetes – Core Services – Network Security



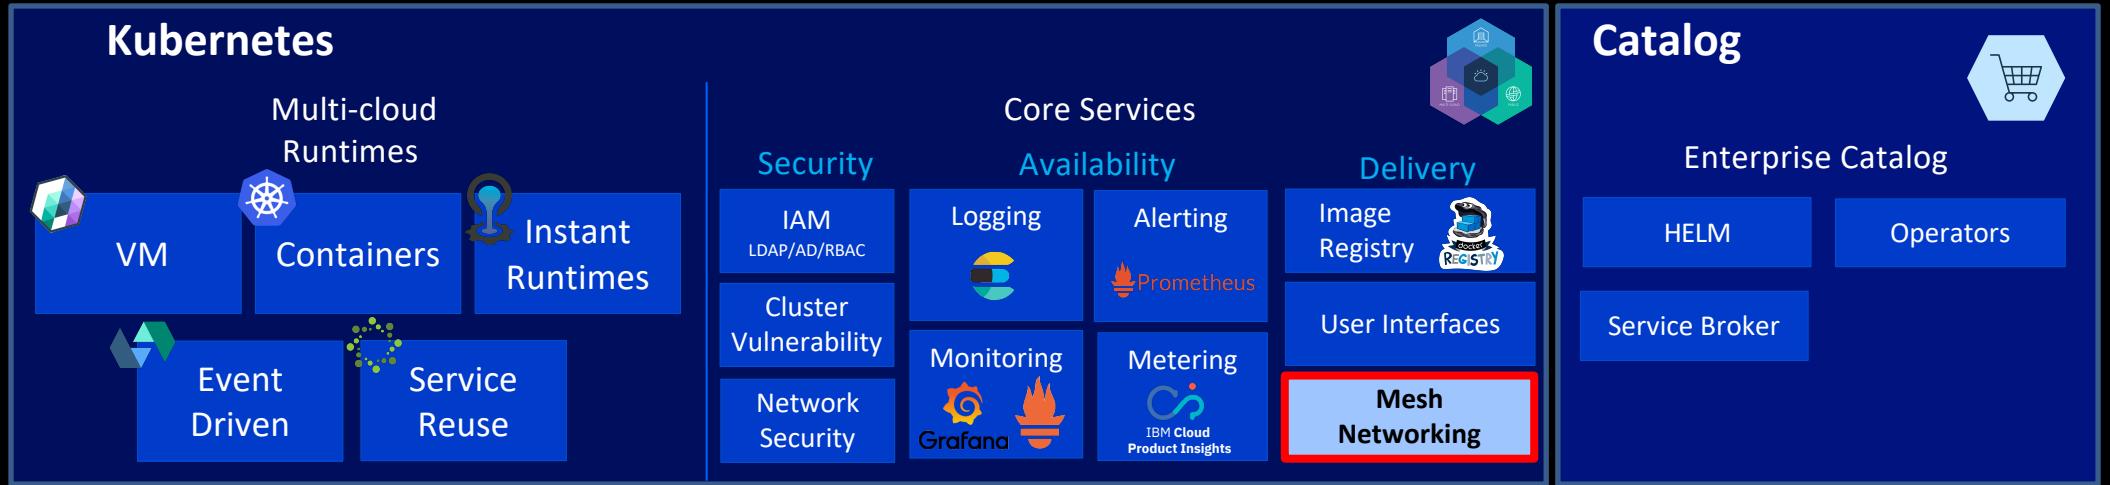
```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: default-deny-all
spec:
  podSelector:
    matchLabels: {}
```

Kubernetes – Core Services – Network Security



```
kind: NetworkPolicy
metadata:
  name: access-frontend-backend
  namespace: prod-namespace
spec:
  podSelector:
    matchLabels:
      run: DB
  ingress:
  - from:
    - podSelector:
        matchLabels:
          run: WEB
```

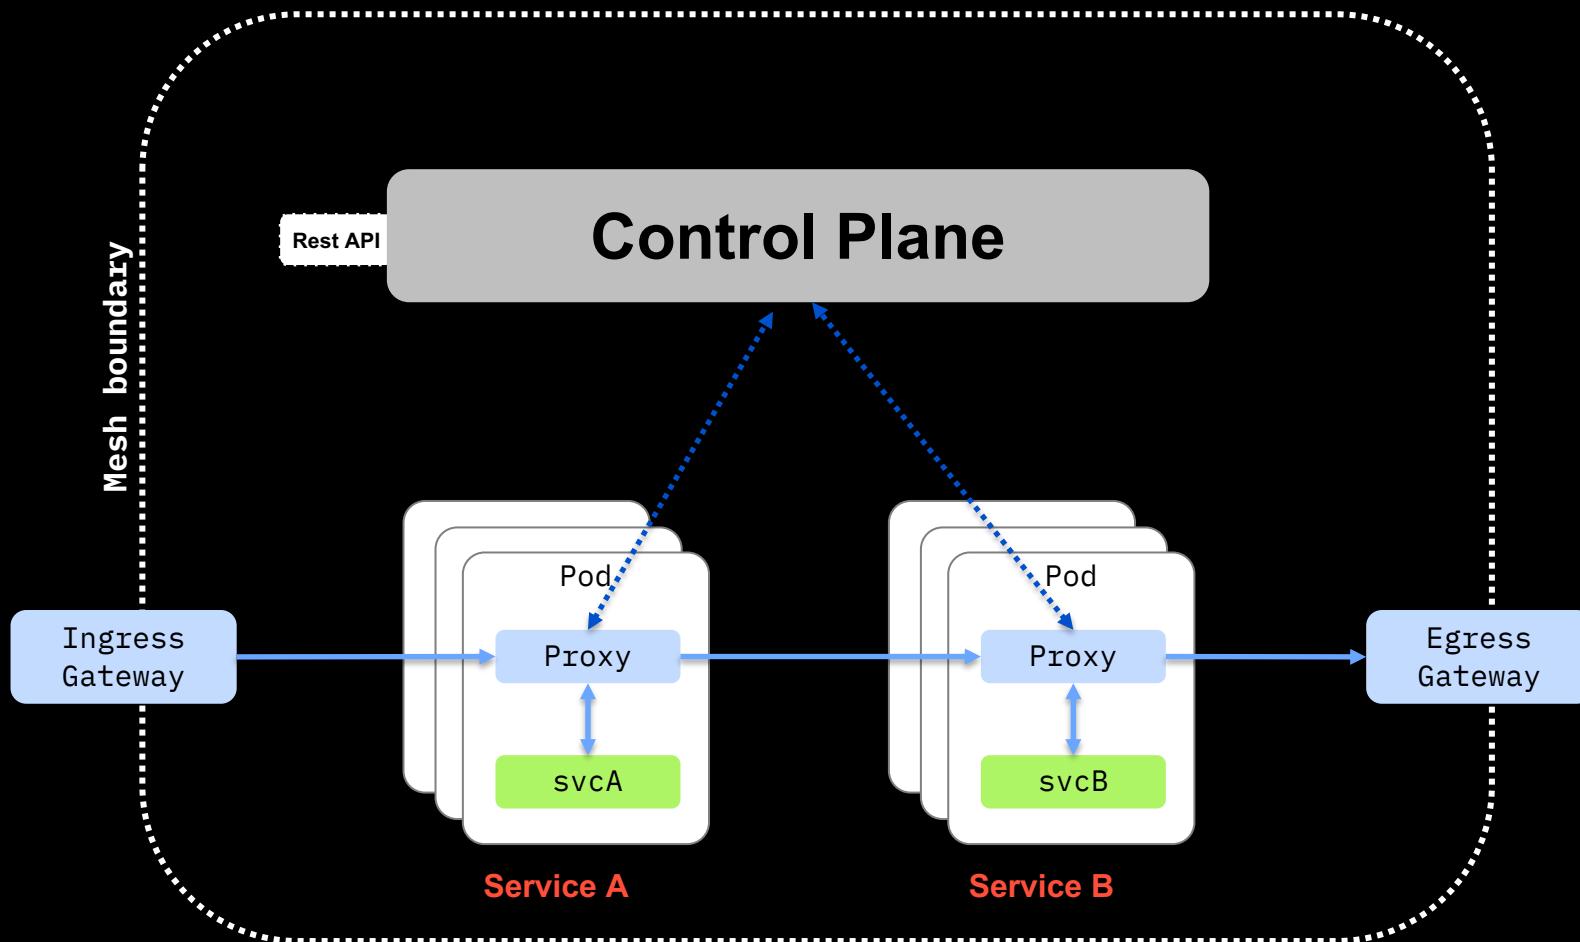
Kubernetes – Core Services – Mesh Networking



Service Mesh

**Dedicated infrastructure layer
to make
service-to-service communication
fast, safe and reliable**

ISTIO - Architecture



Addressing DevOps Challenges



ROLL OUT NEW VERSION WITHOUT DOWNTIME
OR CHANGING CODE

HOW TO DO **CANARY TESTING**

HOW TO DO **A/B TESTING**

THINGS DON'T ALWAYS GO CORRECTLY **IN PRODUCTION...**

HOW CAN I **LIMIT RATE** FOR SOME OF MY SERVICES?

I NEED TO **VIEW AND MONITOR** WHAT IS GOING ON WHEN CRISIS ARISES

HOW CAN I **SECURE** MY SERVICES?

TRAFFIC CONTROL

TRAFFIC SPLITTING

TRAFFIC STEERING

TRAFFIC MIRRORING
RESILIENCY
RESILIENCY TESTING

RATE LIMITING

TELEMETRY

AUTHENTICATION
AUTHORIZATION

Kubernetes – Core Services – Network Security

Certificate Management for TLS

cert-manager is a native Kubernetes certificate management controller.

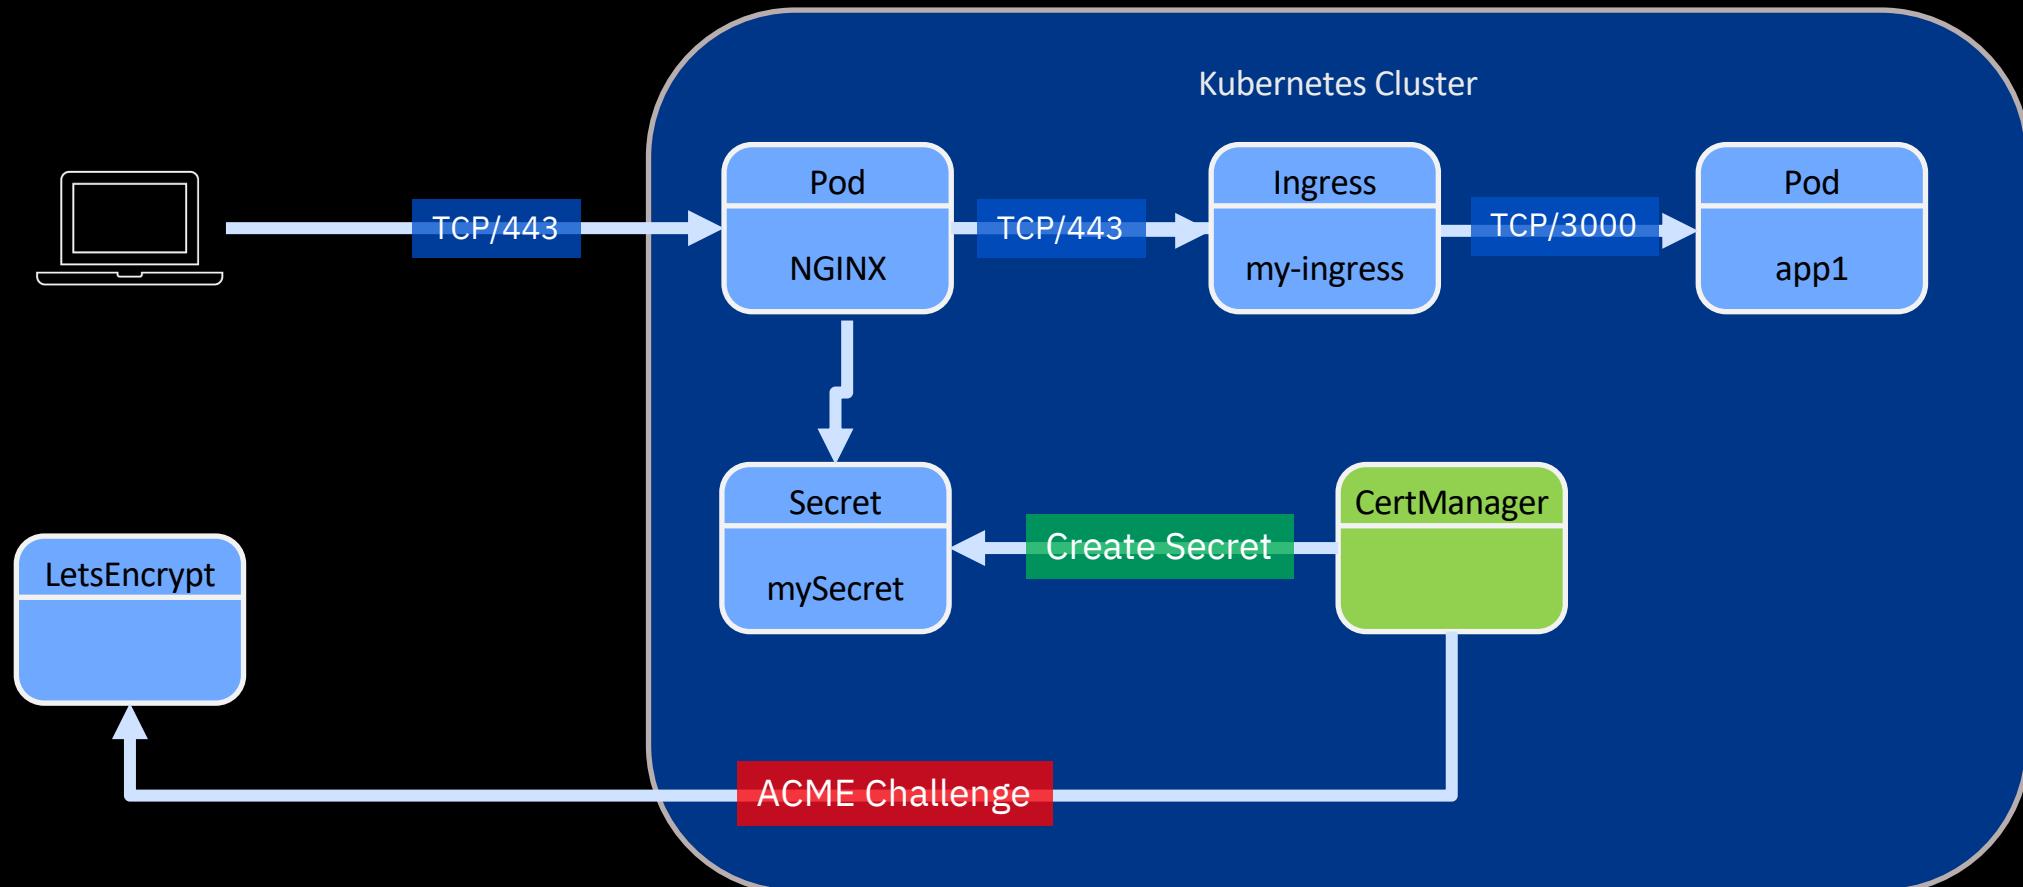
It can help with issuing certificates from a variety of sources, such as Let's Encrypt, HashiCorp Vault, Venafi, a simple signing keypair, or self signed.

It will ensure certificates are valid and up to date, and attempt to renew certificates at a configured time before expiry.



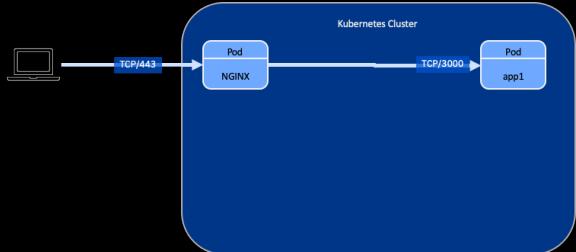
Tutorial: <https://docs.cert-manager.io/en/latest/tutorials/acme/quick-start/index.html#>

Kubernetes – Core Services – Network Security



Kubernetes – Core Services – Network Security

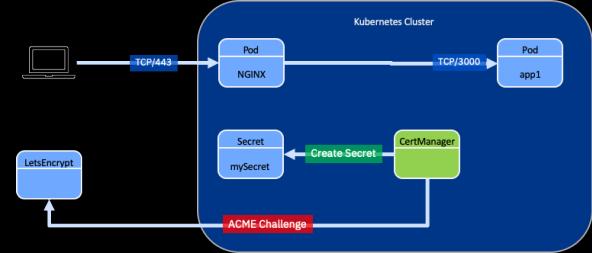
```
apiVersion: v1
kind: Service
metadata:
  name: app1
  labels:
    app: app1
    tier: frontend
spec:
  type: ClusterIP
  ports:
  - port: 80
    targetPort: 3000
  selector:
    app: app1
    tier: frontend
```



Kubernetes – Core Services – Network Security

```
apiVersion: v1
kind: Service
metadata:
  name: app1
  labels:
    app: app1
    tier: frontend
spec:
  type: ClusterIP
  ports:
  - port: 80
    targetPort: 3000
  selector:
    app: app1
    tier: frontend
```

```
apiVersion: certmanager.k8s.io/v1alpha1
kind: ClusterIssuer
metadata:
  name: letsencrypt-staging
spec:
  acme:
    server: https://acme-staging-v02.api.letsencrypt.org/directory
    email: <me@example.com>
    privateKeySecretRef:
      name: mySecret
    http01: {}
```

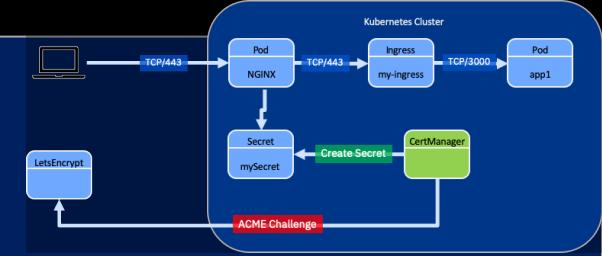


Kubernetes – Core Services – Network Security

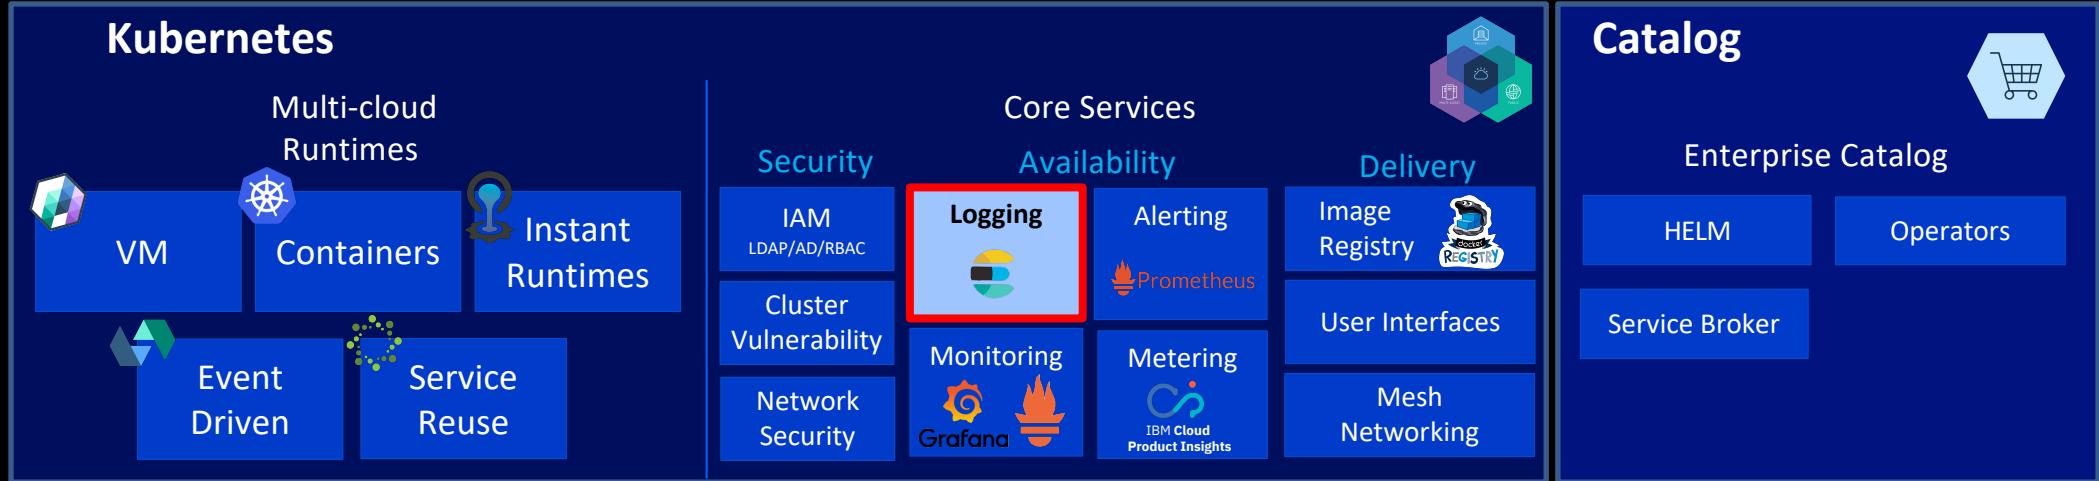
```
apiVersion: certmanager.k8s.io/v1alpha1
kind: ClusterIssuer
metadata:
  name: letsencrypt-staging
spec:
  acme:
    server: https://acme-staging-v02.api.letsencrypt.org/
    email: <me@example.com>
    privateKeySecretRef:
      name: mySecret
    http01: {}
```

```
apiVersion: v1
kind: Service
metadata:
  name: app1
  labels:
    app: app1
    tier: frontend
spec:
  type: ClusterIP
  ports:
  - port: 80
    targetPort: 3000
  selector:
    app: app1
    tier: frontend
```

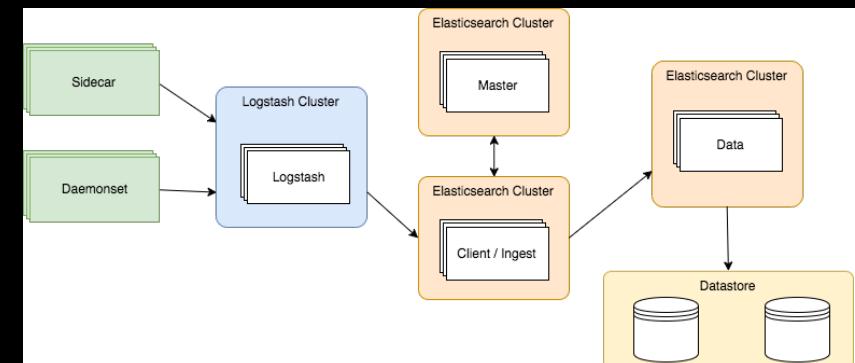
```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: my-ingress
  annotations:
    kubernetes.io/ingress.class: nginx
    certmanager.k8s.io/cluster-issuer: letsencrypt-staging
    kubernetes.io/tls-acme: "true"
spec:
  rules:
  - host: app.example.com
    http:
      paths:
      - path: /
        backend:
          serviceName: app1
          servicePort: 80
    tls:
    - secretName: mySecret
      hosts:
      - app.example.com
```



Kubernetes – Core Services - Logging



- **Logging (ELK):** The easiest and most embraced open-source logging solution for containerized applications



Logging components

The easiest and most embraced logging method for containerized applications is to write to standard out and standard error

Filebeat: A log data shipper for local files. Filebeat monitors the log directories or specific log files, tails the files, and forwards them either to [Elasticsearch](#) and/or [Logstash](#) for indexing.

Elasticsearch: An open source full-text search engine based on Lucene. It provides HTTP web interface and schema-free JSON documents.

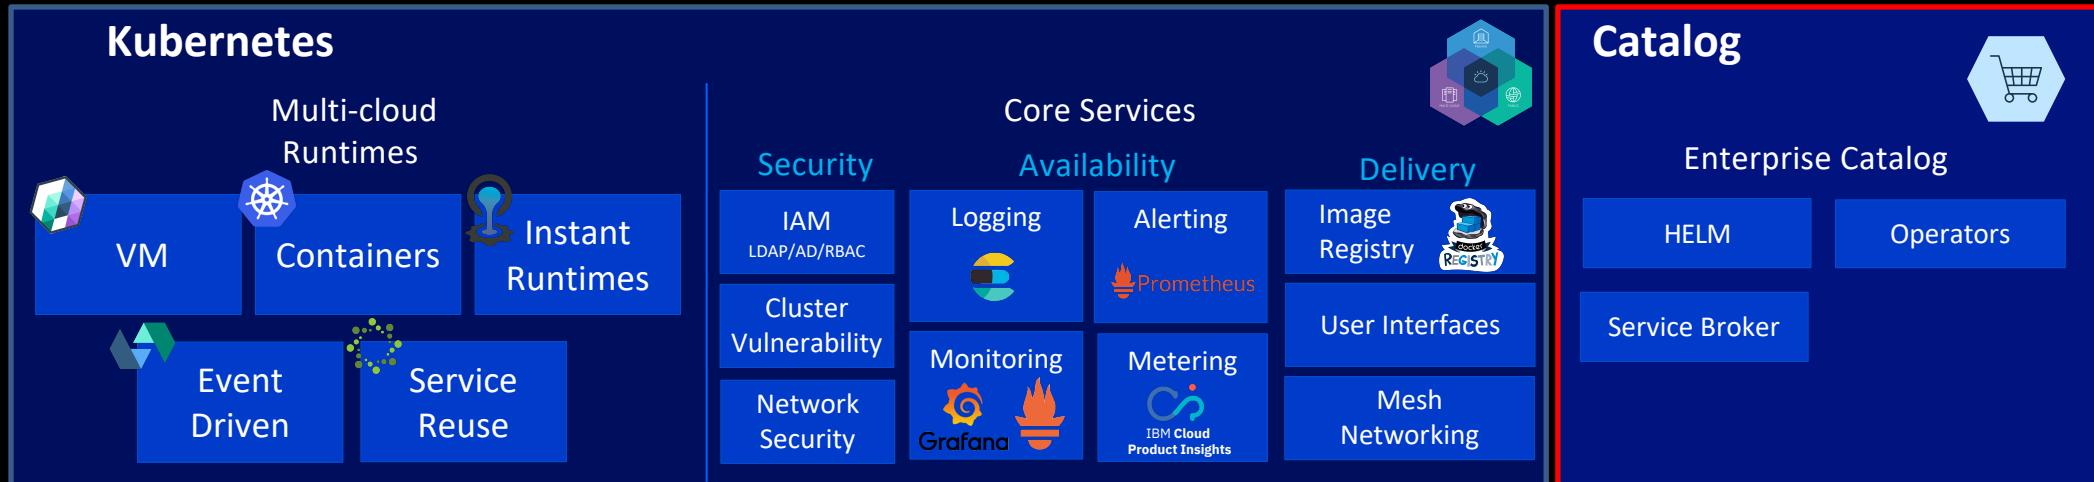
Logstash: A open source tool for collecting, parsing, and storing logs for future use.

Heapster: The Kubernetes network proxy runs on each node.

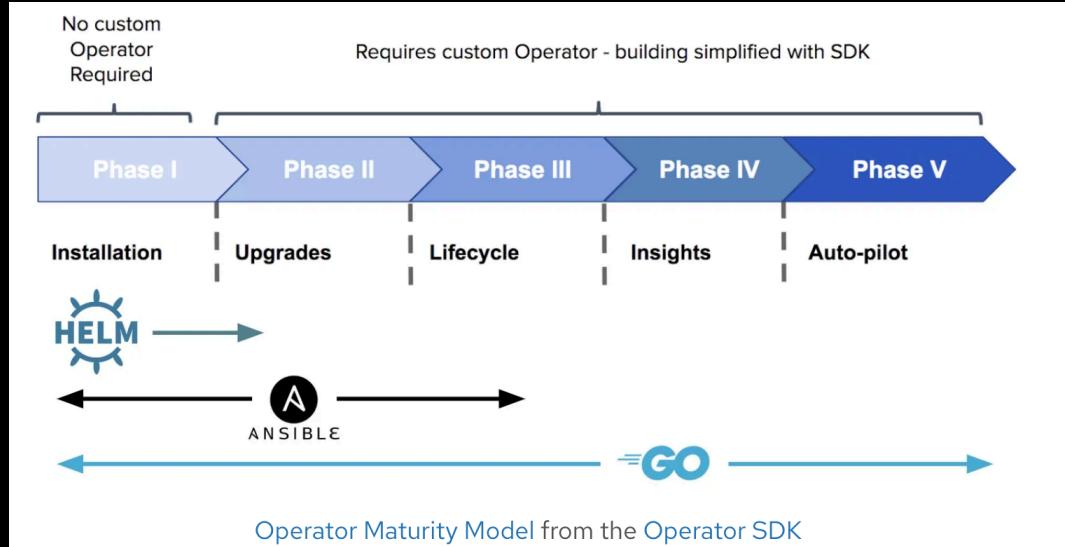
Kibana: An open source data visualization plugin for Elasticsearch. Users can create bar, line and scatter plots, or pie charts and maps on top of large volumes of data.



Kubernetes – Catalog – Helm & Operators

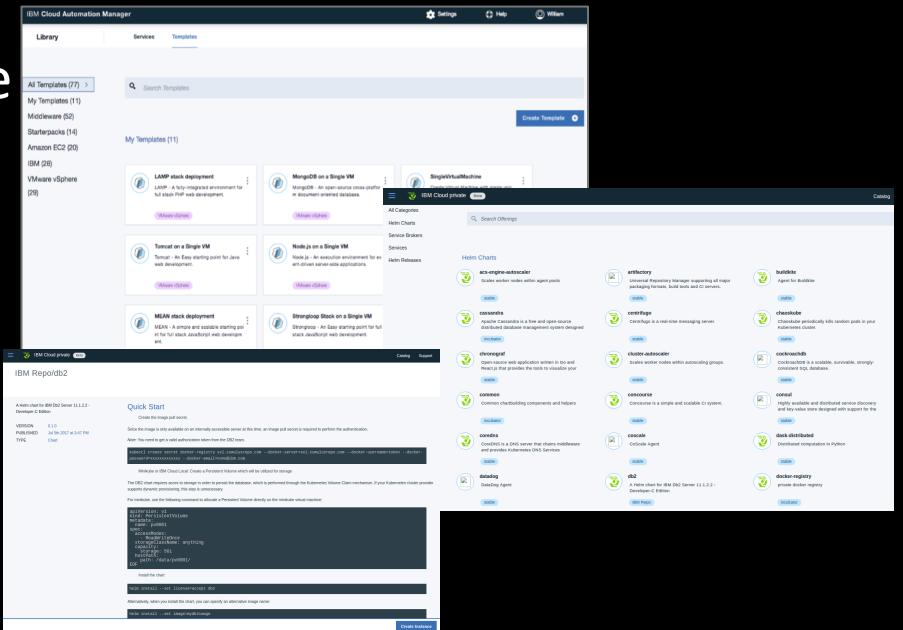


Kubernetes – Catalog – Helm & Operators



Provide a centralized location from which you can control the installation of packages into your cluster.

- Catalog of Open source content
 - Leverage pre-build curated Helm Chart, Terraform Configurations and Chef scripts
 - Reusable and customizable
 - Day 2 Operations

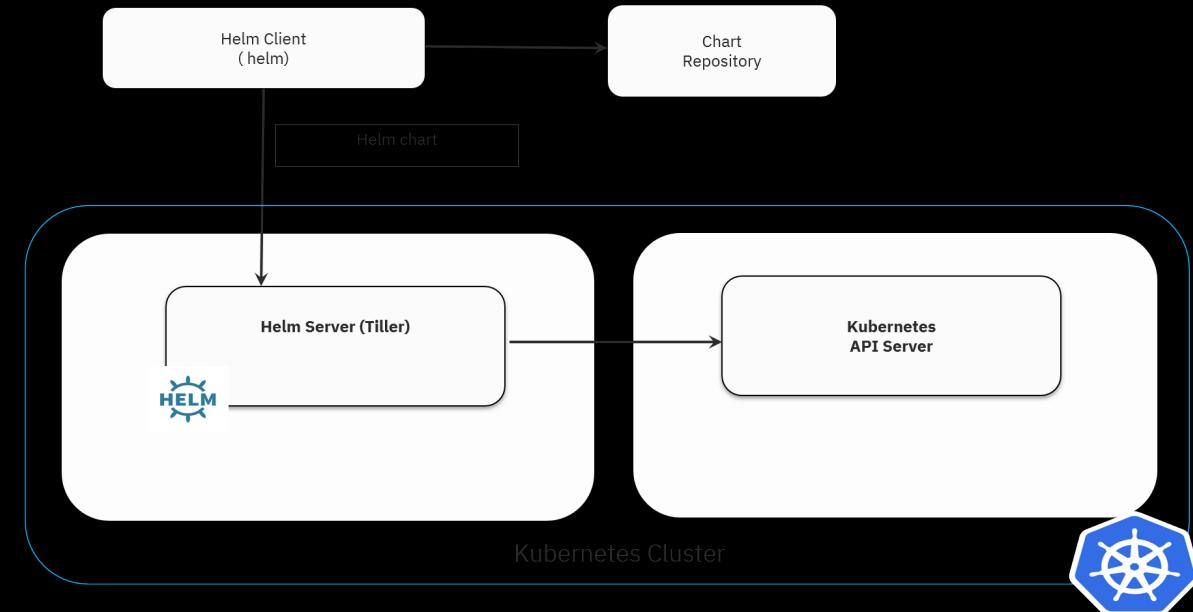


Helm Charts



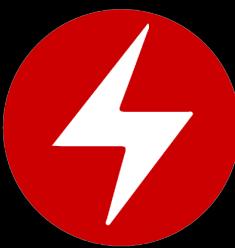
The Helm is a tool that streamlines installation and management of Kubernetes applications. It uses a packaging format called "charts", which are a collection of files that describe Kubernetes resources.

It has a client-server architecture with the client called helm and the server called Tiller



<https://helm.sh/>

Kubernetes Operators

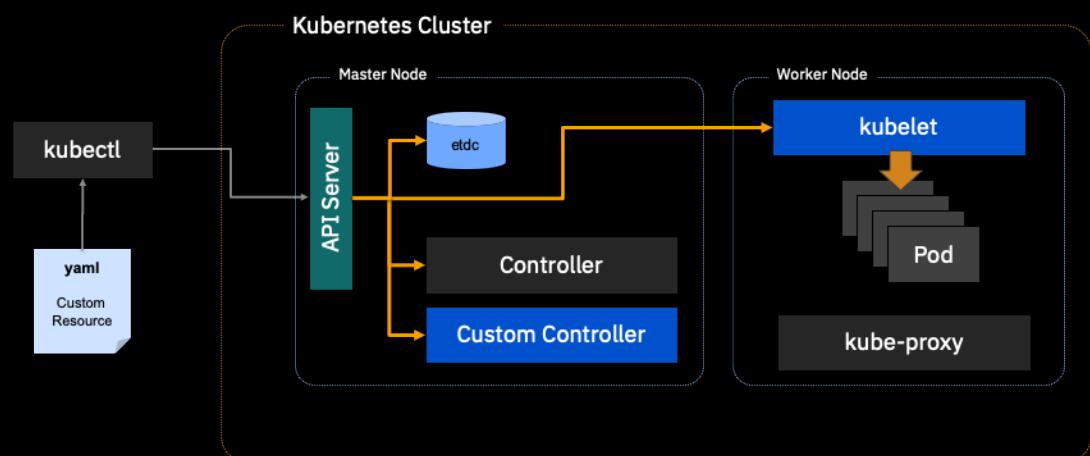


Operators are a **design pattern** made public in a 2016 CoreOS [blog](#) post.

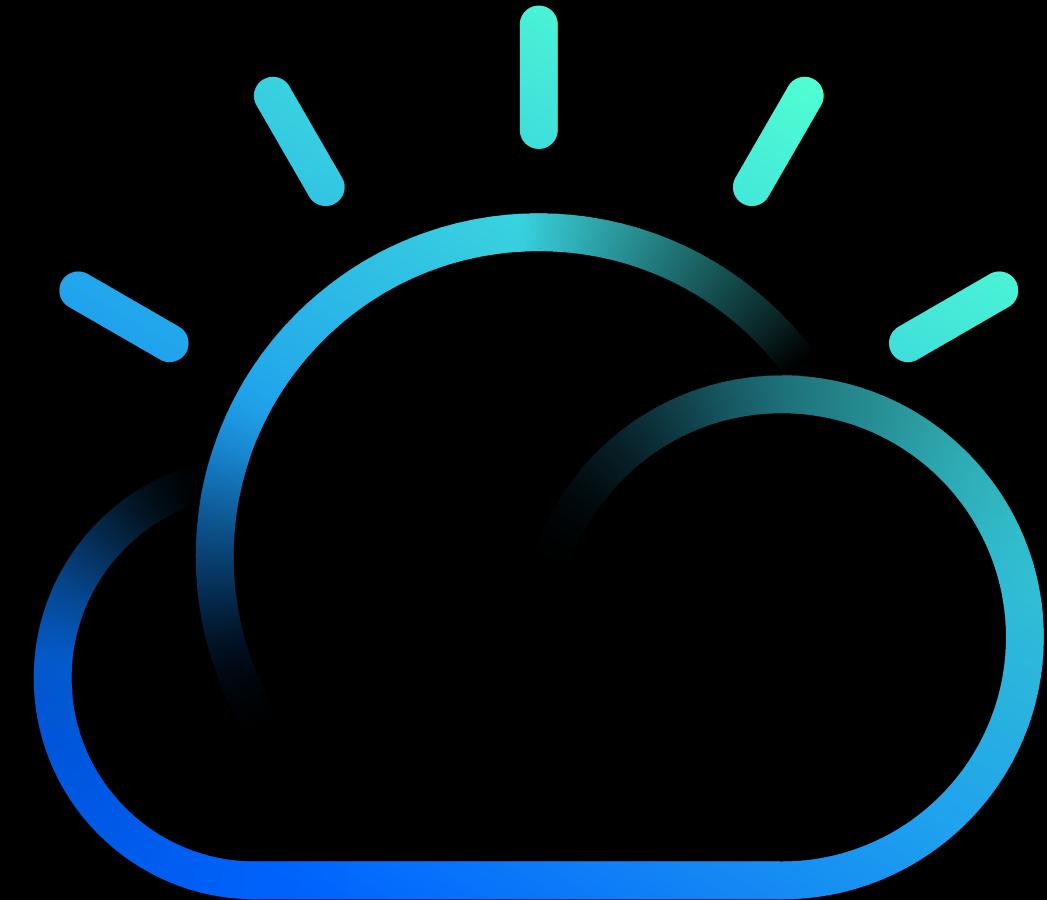
The goal of an Operator is to put **operational knowledge into software**.

Previously this knowledge only resided in the minds of administrators, various combinations of shell scripts or automation software like Ansible. It was outside of your Kubernetes cluster and hard to integrate.

Operators implement and automate common **Day-1** (installation, configuration, etc) and **Day-2** (re-configuration, update, backup, failover, restore, etc.) **activities** in a piece of software running inside your Kubernetes cluster, by integrating natively with Kubernetes concepts and APIs.



QUESTIONS?



IBM Cloud

The Journey to Cloud **Wrap-up and Tips**

03



IBM Cloud

Kubernetes – Tip 1

Use **Namespaces** Liberally

“Namespaces are cheap.”

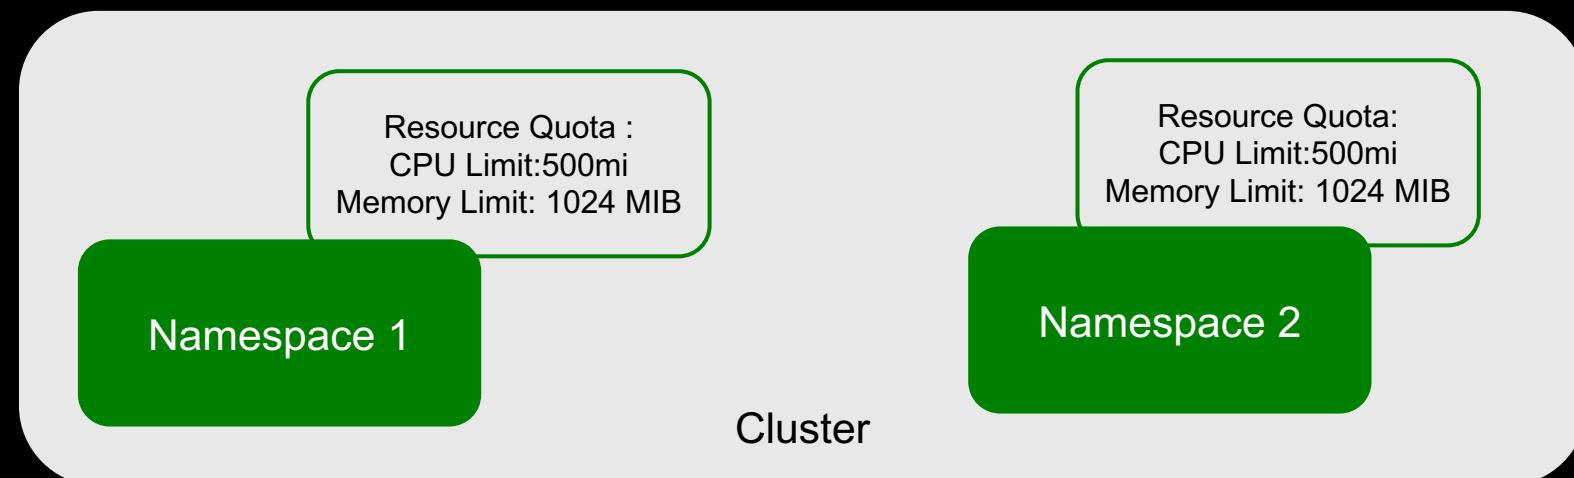
Use them to separate things like infrastructure tooling and applications.

This allows you to restrict access easily using **RBAC** and to limit the scope of applications. It will also make your **network policy** creation easier when you decide to do it.

Kubernetes – Tip 2

Set resource quota

Once quota in a namespace for compute resources set, the users are forced to set requests or limits for those values. Avoids starving of Cluster.



Kubernetes – Tip 3

Scan your containers for vulnerabilities

Known vulnerabilities account for a large portion of breaches. Use a tool to scan your containers for them and then mitigate them.

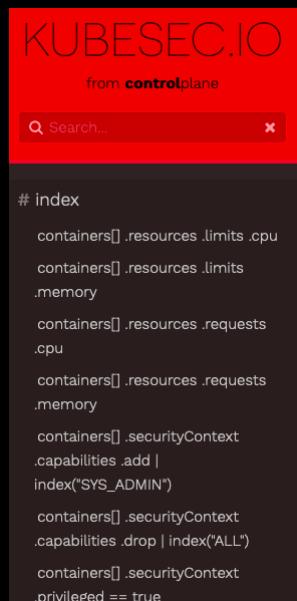
[Anchore](#), [Clair](#), and [Quay](#) are among my favorite tools, but there are others out there.



Kubernetes – Tip 4

Look At **KUBSEC.IO**

This tool will analyze your deployment yaml and tell you what you should change in order to improve the security of those pods. It even gives you a score that you can use to create a minimum standard. The score incorporates all of the best practices that I have outlined above, as well as several more.



KUBESEC.IO – V2

⚠ v1 API is deprecated, please read the [release notes](#) ⚠

Security risk analysis for Kubernetes resources



Kubernetes – Tip 5

Occam's razor (also **Ockham's razor** or Ocham's **razor**: Latin: novacula Occami; or law of parsimony: Latin: lex parsimoniae) is the problem-solving principle that states:

“Entities should not be multiplied without necessity.”

Or

“other things being equal, simpler explanations are generally better than more complex ones””

Kubernetes – Tip 6

Disable **Anonymous access**

By default, requests to the K8s API and kubelet's HTTPS endpoint that are not rejected by other configured authentication methods are treated as anonymous requests, and given a username of `system:anonymous` and a group of `system:unauthenticated`.

Anonymous authentication provides full access to the *kubelet* API, the only requirement being network access to the service.

To disable anonymous access pass `--anonymous-auth=false` to the API server and start the kubelet with the `--anonymous-auth=false` flag

If RBAC is enabled this risk is mitigated.

Kubernetes – Tip 6

Secure **kubelet**

Set `--anonymous-auth=false`

Set `--authorization-mode` to something other than `AlwaysAllow`

Set `--read-only-port=0`

Kubernetes – Tip 6

Secure **etcd**

Set --cert-file and --key-file to enable HTTPS connections

Set --client-cert-auth=true

Kubernetes – Tip 6

Secure API

The Kubernetes API server can serve requests on two ports:

- localhost port (8080) and
- secure port

The localhost port is intended for testing purposes and for other master components to talk to the API. Requests to the Localhost port **bypass authentication and authorization modules**.

Best practice is to set `--insecure-port flag=0` and remove the `--insecure-bind-address` flag from the API server manifest.

And remove the `--secure-port` flag from the API server spec to ensure that all requests to the secure port are authenticated and authorized.

Note: Has been deprecated in Kubernetes 1.10 and should be removed in the future

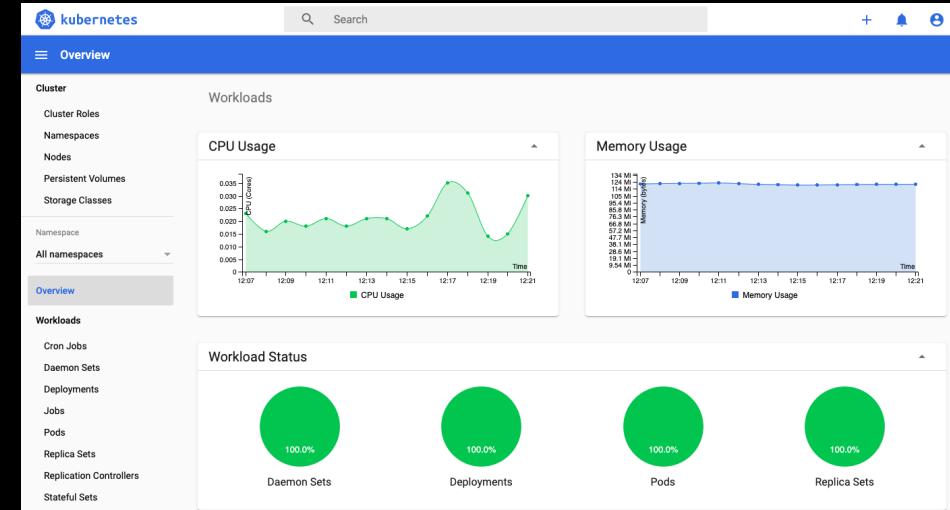
Kubernetes – Tip 7

Secure Kubernetes Dashboard

The Kubernetes Dashboard has often been used for attacks.
Easy to deploy with high privileges.

Best practice is to:

- Allow only authenticated access
- Use RBAC
- Ensure Dashboard service account has limited access
- Don't expose to the internet (use kubectl proxy for local access)



Kubernetes – Tip 8

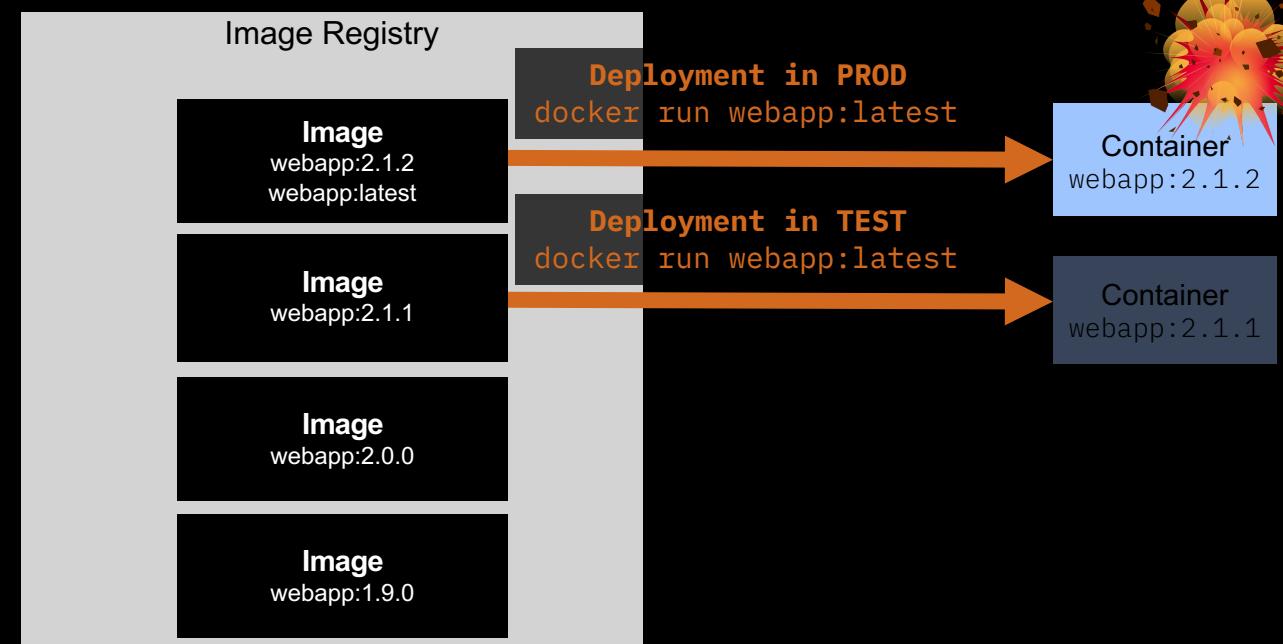
Avoid **Rolling Tags** at Any Cost

Avoid commands like `docker pull myregistry/myimage:latest`.

This “latest” is an example of a rolling tag (i.e. a tag that will point to different images over time).

If you want your deployments to be secure and maintainable, make sure that you use immutable images (for example: “`myregistry/myimage:1.1.2`”).

With this approach, every time you deploy or scale, you know what image you are using and you will have the guarantee that the deployed image has been tested and validated.



Kubernetes – Tip 9

Don't run **multiple processes per container**

Docker general best practices suggest a single process per container simply for usability and size reasons but it also **keeps your attack surface small** and limits the number of potential vulnerabilities.

Kubernetes – Tip 10

Don't run containers as the root user

By default, containers run as the root user inside the container. This is easy to avoid using the pod specification to set a high-numbered UID.

```
spec:  
  securityContext:  
    runAsUser: 10324
```

Kubernetes – Tip 11

Don't run containers as privileged

Running a container or pod as privileged gives it the ability to make modifications to the host. This is a large security issue that is easy to mitigate by just not doing it.

Kubernetes – Tip 12

Don't use the default list of capabilities

Docker runs containers with a significant set of Linux capabilities by default, many of which your app might not require.

The following config will drop all Linux capabilities by default, allowing you to add only the specific capabilities your app actually needs:

```
spec:  
  containers:  
    - name: foo  
      securityContext:  
        capabilities:  
          drop:  
            - ALL
```

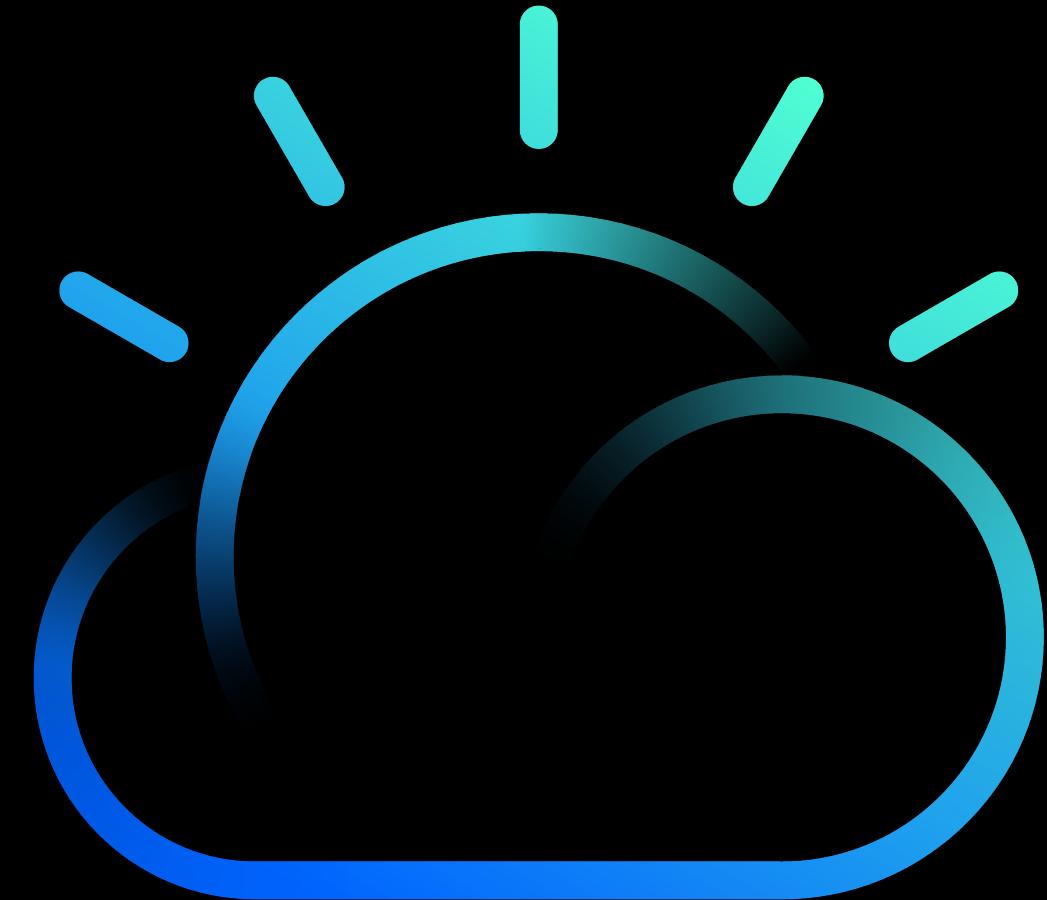
Kubernetes – Self Check

1. **Where** are images used in containers **coming from**?
2. How long ago were the images **scanned for vulnerabilities**?
3. Which of your containers are impacted by **known vulnerabilities**, and what's their severity?
4. Are any of these containers in **production** impacted by a known **vulnerability**?
5. Which vulnerable running containers or deployments should be **prioritized for remediation**?
6. Which deployments are using **privileged containers**, meaning they have full access to the host operating system?

Kubernetes – Self Check

7. What container **services are exposed** outside of the Kubernetes cluster?
8. Can we tell which **processes** are **running** in any container in any cluster?
9. Which **network communication** pathways are **active** but are not being used in production?
- 10.What **team** in the organization **owns** a particular **running application**?
- 11.How many of my clusters, namespaces, and nodes must **adhere to specific regulations**.

QUESTIONS?

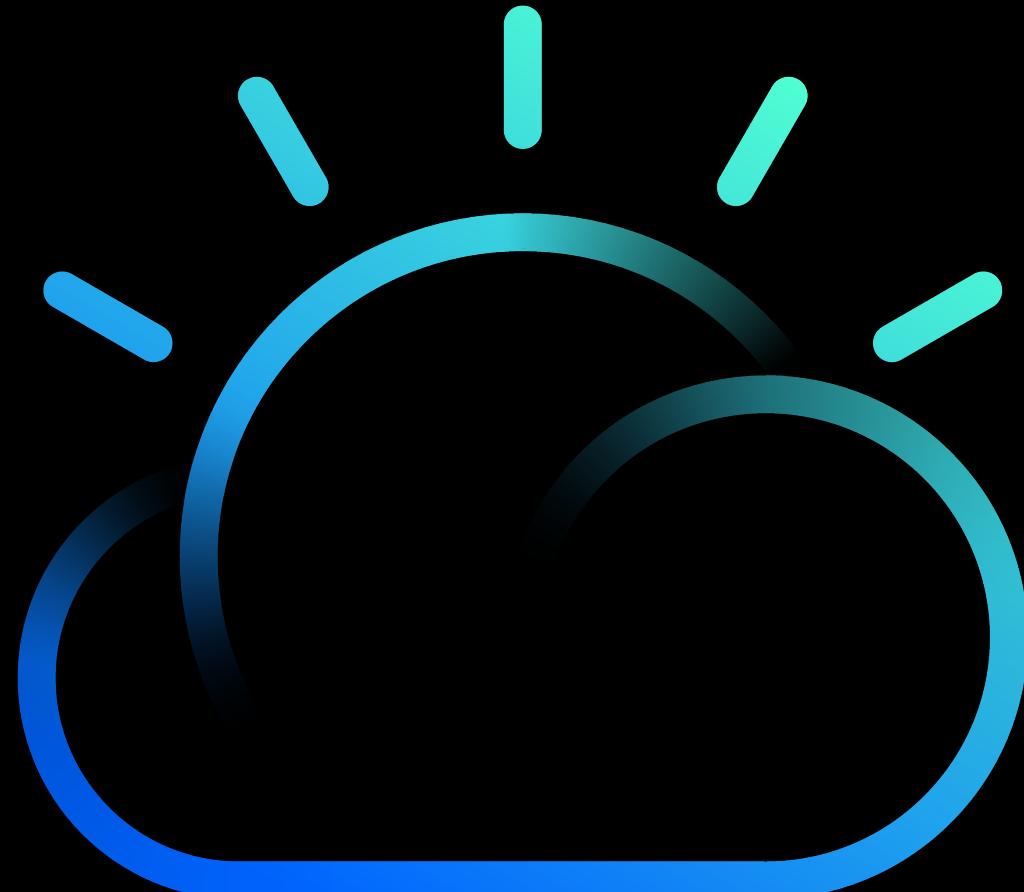


IBM Cloud



°° The Journey to Cloud
K8s Security - Hands-On

04



IBM Cloud

Remember your Team Color

black 31701

olive 31711

peru 31715

white 31702

brown 31712

chocolate 31716

red 31703

lightblue 31713

orchid 31717

blue 31704

orange 31708

gold 31718

yellow 31705

purple 31709

pink 31719

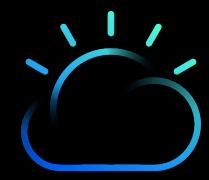
lime 31706

maroon 31710

violet 31720

cyan 31707

firebrick 31714



Collector - Accessing team web site

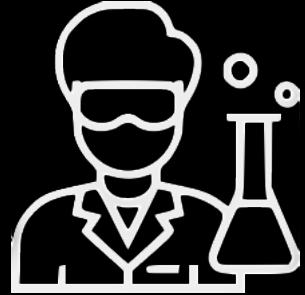
http://158.177.137.195:{port#}

Team name / color will be shown

blue 31704

The screenshot shows a web browser window with the title "Collector - blue". The navigation bar includes links for "Courses", "Class work", "Statistics", "Information", and "Feedback". Below the navigation bar is a section titled "Catalog of courses". A dropdown menu is open under the "Courses" link, listing five items: "select course", "KUB01 Lab Setup", "KUB02 Kubernetes Introduction", "KUB03 Kubernetes Labs", and "KUBADV01 Istio". To the right of the dropdown is a button labeled "Begin course". A large callout box with a red border and white text points to the "Begin course" button, containing the instruction "Select course and press button to begin".

Current course catalog –
Select **JTC16 Kubernetes Security Labs**



JTC16 Kubernetes Security Labs

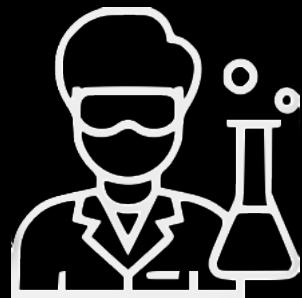
Lab 1 : Network Policies

Lab 2 : Role Based Access Control (RBAC)

Lab 3 : Service Accounts

Lab 4 : Security Tooling

Lab 5 : Image scanning

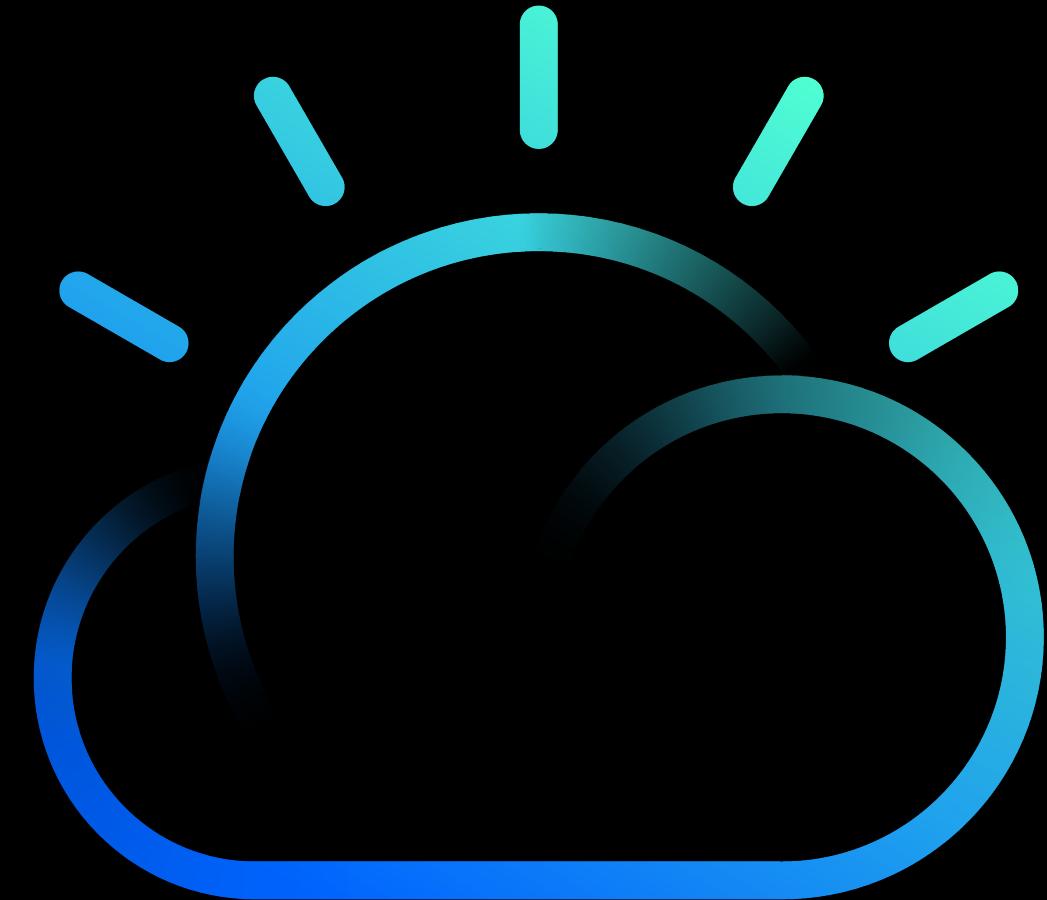


READY
SET
GO!!!!

<http://158.177.137.195:{port#}>

Duration: 60 mins

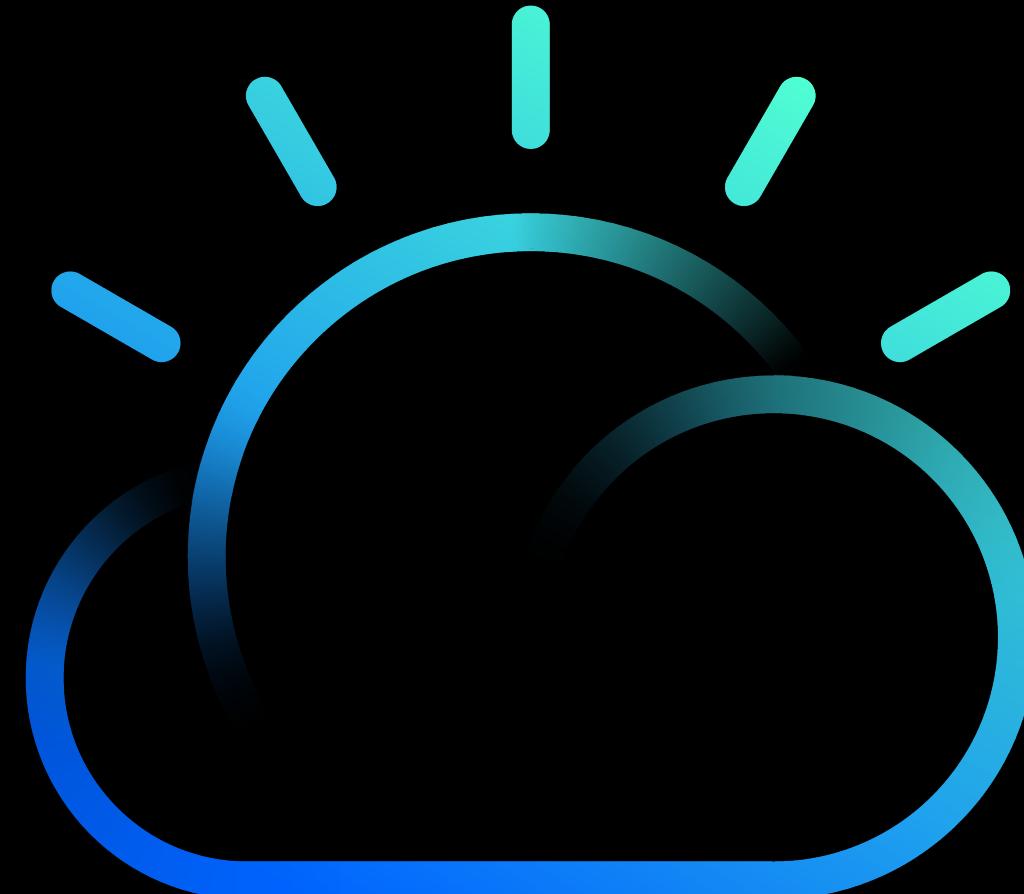
QUESTIONS?



IBM Cloud

The Journey to Cloud **Wrap Up**

Q9



IBM Cloud

Kubernetes Security – Some tools I use

RBAC

RBAC-lookup - <https://github.com/FairwindsOps/rbac-lookup>

rakkess - <https://github.com/corneliusweig/rakkess>

rbac-view – <https://github.com/jasonrichardsmith/rbac-view>

rback - <https://github.com/team-soteria/rback>

Scanning

Polaris - <https://github.com/FairwindsOps/polaris>

Clair - <https://github.com/coreos/clair>

Management

K9s - <https://github.com/derailed/k9s>

Kubernetes Security – Some Reading Tips

<https://kubernetespodcast.com/episode/065-attacking-and-defending-kubernetes/>

https://en.wikipedia.org/wiki/Red_team

<https://blog.ropnop.com/attacking-default-installs-of-helm-on-kubernetes/>

<https://kubernetes.io/docs/tasks/administer-cluster/encrypt-data/>

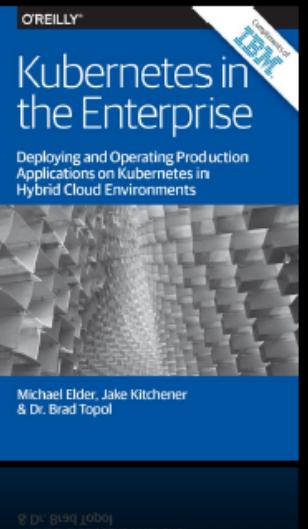
<https://kubernetes.io/blog/2019/03/21/a-guide-to-kubernetes-admission-controllers/>

<https://kubernetes.io/docs/tutorials/clusters/apparmor/>

<https://kubernetes.io/docs/concepts/policy/pod-security-policy/>

<https://kubernetes.io/docs/concepts/storage/volumes/#hostpath>

Kubernetes Security – Some Reading Tips

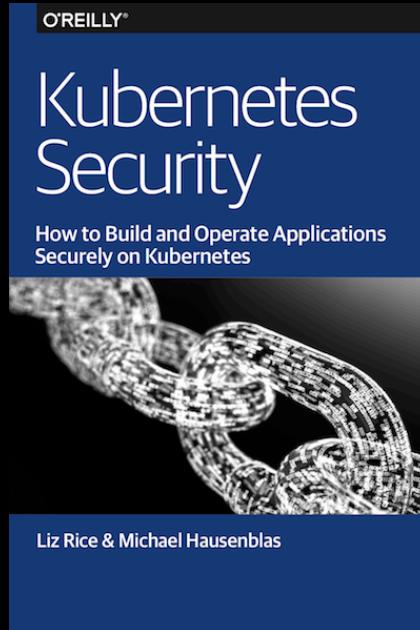


We wrote the books on
cloud adoption...

ibm.com/cloud/garage/adoption

...and Kubernetes in
the enterprise

<https://ibm.co/2LQketN>



<https://kubernetes-security.info/>



Sources and documentation will be available here:

https://github.com/niklaushirt/k8s_training_public

<https://github.com/niklaushirt/training>

THANK YOU!!!!

Niklaus Hirt

nikh@ch.ibm.com

@nhirt

IBM